



# Arm<sup>®</sup> Architecture Reference Manual for A-profile architecture

## Known issues in Issue I.a

**Non-Confidential**

Copyright © 2020, 2022 Arm Limited (or its affiliates).  
All rights reserved.

**Issue 02**

102105\_I.a\_02\_en



## Arm® Architecture Reference Manual for A-profile architecture

### Known issues in Issue I.a

Copyright © 2020, 2022 Arm Limited (or its affiliates). All rights reserved.

## Release information

### Document history

Issue	Date	Confidentiality	Change
F.c-04	18 December 2020	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue F.c, as of 18 December 2020
G.b-05	31 January 2022	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue G.b, as of 7 January 2022
H.a-06	22 July 2022	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue H.a, as of 22 July 2022
I.a-00	5 August 2022	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue I.a, as of 5 August 2022
I.a-01	30 September 2022	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue I.a, as of 23 September 2022
I.a-02	31 October 2022	Non-Confidential	Known Issues in Arm® Architecture Reference Manual, Issue I.a, as of 21 October 2022

## Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-

INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its affiliates) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <https://www.arm.com/company/policies/trademarks>.

Copyright © 2020, 2022 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.

110 Fulbourn Road, Cambridge, England CB1 9NJ.

(LES-PRE-20349|version 21.0)

## Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Unrestricted Access is an Arm internal classification.

## Product Status

The information in this document is Final, that is for a developed product.

## Feedback

Arm® welcomes feedback on this product and its documentation. To provide feedback on the product, create a ticket on <https://support.developer.arm.com>

To provide feedback on the document, fill the following survey: <https://developer.arm.com/documentation-feedback-survey>.

## Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used language that can be offensive. Arm strives to lead the industry and create change.

We believe that this document contains no offensive language. To report offensive language in this document, email [terms@arm.com](mailto:terms@arm.com).

# Contents

<b>1 Introduction.....</b>	<b>8</b>
1.1 Conventions.....	8
1.2 Useful resources.....	9
1.3 Other information.....	9
<b>2 Known issues.....</b>	<b>10</b>
2.1 C15788.....	10
2.2 C16212.....	10
2.3 D17015.....	10
2.4 D17119.....	11
2.5 R17661.....	11
2.6 E17792.....	11
2.7 C17811.....	14
2.8 E17996.....	14
2.9 D18330.....	16
2.10 D18465.....	16
2.11 R18485.....	16
2.12 D18520.....	16
2.13 D18736.....	17
2.14 R18746.....	18
2.15 D18800.....	19
2.16 D18823.....	20
2.17 C18842.....	24
2.18 C18843.....	24
2.19 D18853.....	24
2.20 D18889.....	25
2.21 C19027.....	26
2.22 C19047.....	27
2.23 D19116.....	27
2.24 D19121.....	27
2.25 D19162.....	28
2.26 D19178.....	28

2.27 C19202.....	31
2.28 D19239.....	32
2.29 D19275.....	32
2.30 R19370.....	32
2.31 D19372.....	33
2.32 E19440.....	33
2.33 D19451.....	33
2.34 D19452.....	33
2.35 D19494.....	34
2.36 R19519.....	35
2.37 D19521.....	35
2.38 D19549.....	36
2.39 D19560.....	36
2.40 D19561.....	36
2.41 D19581.....	37
2.42 D19642.....	37
2.43 C19644.....	38
2.44 D19647.....	38
2.45 C19649.....	39
2.46 D19680.....	39
2.47 E19713.....	40
2.48 D19741.....	40
2.49 D19753.....	41
2.50 C19772.....	41
2.51 C19793.....	42
2.52 D19800.....	45
2.53 D19804.....	45
2.54 R19810.....	45
2.55 D19817.....	45
2.56 D19829.....	46
2.57 E19831.....	46
2.58 D19833.....	47
2.59 C19835.....	47
2.60 D19887.....	47
2.61 E19892.....	48
2.62 D19917.....	49

2.63 D19918.....	49
2.64 D19928.....	50
2.65 D19936.....	50
2.66 C19956.....	51
2.67 D19961.....	51
2.68 C20009.....	51
2.69 D20011.....	52
2.70 D20053.....	53
2.71 D20128.....	53
2.72 R20165.....	54
2.73 D20171.....	54
2.74 D20210.....	55
2.75 C1186: SME.....	56
2.76 C1342: SME.....	56
2.77 D1386: SME.....	57
2.78 C215: SVE.....	58
2.79 C225: SVE.....	60
2.80 C256: SVE.....	60
2.81 C279: SVE.....	61
2.82 D302: SVE.....	61
2.83 D1461: Armv9 Debug.....	62
2.84 D1466: Armv9 Debug.....	62
2.85 D1493: Armv9 Debug.....	62

# 1 Introduction

## 1.1 Conventions

The following subsections describe conventions used in Arm documents.





### Glossary

The Arm Glossary is a list of terms used in Arm documentation, together with definitions for those terms. The Arm Glossary does not contain terms that are industry standard unless the Arm meaning differs from the generally accepted meaning.



See the Arm® Glossary for more information: [developer.arm.com/glossary](https://developer.arm.com/glossary).

### Typographic conventions

Arm documentation uses typographical conventions to convey specific meaning.

Convention	Use
<i>italic</i>	Citations.
<b>bold</b>	Interface elements, such as menu names.  Terms in descriptive lists, where appropriate.
monospace	Text that you can enter at the keyboard, such as commands, file and program names, and source code.
monospace <u>underline</u>	A permitted abbreviation for a command or option. You can enter the underlined text instead of the full command or option name.
<and>	Encloses replaceable terms for assembler syntax where they appear in code or code fragments.  For example:  <pre>MRC p15, 0, &lt;Rd&gt;, &lt;CRn&gt;, &lt;CRm&gt;, &lt;Opcode_2&gt;</pre>
<b>SMALL CAPITALS</b>	Terms that have specific technical meanings as defined in the <i>Arm® Glossary</i> . For example, <b>IMPLEMENTATION DEFINED</b> , <b>IMPLEMENTATION SPECIFIC</b> , <b>UNKNOWN</b> , and <b>UNPREDICTABLE</b> .
 Caution	Recommendations. Not following these recommendations might lead to system failure or damage.
 Warning	Requirements for the system. Not following these requirements might result in system failure or damage.
 Danger	Requirements for the system. Not following these requirements will result in system failure or damage.
 Note	An important piece of information that needs your attention.



Convention	Use
 Tip	A useful tip that might make it easier, better or faster to perform a task.
 Remember	A reminder of something important that relates to the information you are reading.

## 1.2 Useful resources

This document contains information that is specific to this product. See the following resources for other useful information.

Access to Arm documents depends on their confidentiality:

- Non-Confidential documents are available at [developer.arm.com/documentation](https://developer.arm.com/documentation). Each document link in the following tables goes to the online version of the document.
- Confidential documents are available to licensees only through the product package.

Arm product resources	Document ID	Confidentiality
Arm® Architecture Reference Manual for A-profile architecture, Issue I.a	DDI 0487I.a	Non-Confidential



Arm tests its PDFs only in Adobe Acrobat and Acrobat Reader. Arm cannot guarantee the quality of its documents when used with any other PDF reader.

Adobe PDF reader products can be downloaded at <http://www.adobe.com>

## 1.3 Other information

See the Arm website for other relevant information.

- [Arm® Developer](#).
- [Arm® Documentation](#).
- [Technical Support](#).
- [Arm® Glossary](#).

## 2 Known issues

This document records known issues in the Arm Architecture Reference Manual for A-profile architecture (DDI 0487), Issue I.a.

Key

- C = Clarification.
- D = Defect.
- R = Relaxation.
- E = Enhancement.

### 2.1 C15788

In section D8.13.5 (TLB maintenance instructions), in the subsection ‘TLB maintenance instructions that apply to a range of addresses’, the following text is added:

It is possible for a TLB range maintenance instruction for a translation regime that supports two VA ranges to be issued with an address in the TTBR1 half of the virtual address space, and SCALE and NUM values such that the range exceeds the top of the address space. In this scenario, the address is not considered to wrap on overflow and the PE is not required to invalidate any entries inserted for the TTBR0 half of the VA space.

### 2.2 C16212

In section D17.2.156 (VSTTBR\_EL2, Virtualization Secure Translation Table Base Register) and D17.2.158 (VTTBR\_EL2, Virtualization Translation Table Base Register), in the field ‘BADDR, bits [47:1]’, the references to:

stage 1 translation table base

are corrected to read:

stage 2 translation table base

### 2.3 D17015

Details of traps will be added through the use of new LDC and STC accessibility pseudocode in sections G8.3.17 (DBGDTRRXint, Debug Data Transfer Register, Receive) and G8.3.19 (DBGDTRTXint, Debug Data Transfer Register, Transmit). This accessibility pseudocode is the same as for the equivalent MRC and MCR instructions, except that:

- The reported exception syndrome value, if applicable, is 0x06.

- For LDC instructions the accessibility pseudocode loads the value to be written to the System register from 'MemA[address, 4]', where 'address' is the virtual address calculated by the LDC instruction.

## 2.4 D17119

In sections F3.1.10 (Advanced SIMD shifts and immediate generation), subsection 'Advanced SIMD two registers and shift amount' and F4.1.22 (Advanced SIMD shifts and immediate generation), subsection 'Advanced SIMD two registers and shift amount', the following constraints are added to VMOVL:

- 'L' must be '0'.
- 'imm3H' cannot be '000'.

## 2.5 R17661

In section D9.2 (Allocation Tags), the following Notes are removed:

Note: The value 0b1111 may incur a higher performance overhead than other Allocation Tag encodings.

Note: Arm recommends that software does not use instructions which write 0b1111 as an Allocation Tag to memory.

## 2.6 E17792

In section J1.3.3 (shared/functions), the AccType enumeration is refactored, such that the AccessDescriptor type is repurposed to hold information captured by the AccType enumeration and replaces the occurrences of AccType throughout the pseudocode in chapter J1 (Armv8 Pseudocode).

The enumeration AccType that reads:

```
enumeration AccType {AccType_NORMAL,           // Normal loads and stores
                     AccType_STREAM,           // Streaming loads and
stores                                     // stores
                     AccType_VEC,              // Vector loads
and stores                               // and stores
                     AccType_VECSTREAM,        // Streaming vector loads
and storesArmv8 Pseudocode              // and storesArmv8 Pseudocode
                     AccType_SVE,              // Scalable vector
loads and stores                       // loads and stores
                     AccType_SVESTREAM,        // Scalable vector
streaming loads and stores              // streaming loads and stores
                     AccType_SME,              // Scalable matrix
loads and stores                       // loads and stores
                     AccType_SMESTREAM,        // Scalable matrix
streaming loads and stores              // streaming loads and stores}
```

```

loads and stores      AccType_UNPRIVSTREAM,          // Streaming unprivileged
multiple              AccType_A32LSMD,              // Load and store
stores                AccType_ATOMIC,               // Atomic loads and
Release              AccType_ATOMICRW,
                     AccType_ORDERED,              // Load-Acquire and Store-
with atomic access    AccType_ORDEREDRW,
                     AccType_ORDEREDATOMIC,        // Load-Acquire and Store-Release
stores                AccType_ORDEREDATOMICRW,
                     AccType_ATOMICCLS64,          // Atomic 64-byte loads and
LORelease            AccType_LIMITEDORDERED,        // Load-LOAcquire and Store-
unprivileged          AccType_UNPRIV,              // Load and store
walk                  AccType_IFETCH,              // Instruction fetch
                     AccType_TTW,                 // Translation table
first element         AccType_NONFAULT,            // Non-faulting loads
                     AccType_CNOTFIRST,            // Contiguous FF load, not
EL1 and which is      AccType_NV2REGISTER,          // MRS/MSR instruction used at
                                                              // converted
to a memory access that uses the                          // EL2
translation regime    AccType_TRBE,                // TRBE memory access
                                                              // Other
operations            AccType_DC,                  // Data cache
maintenance           AccType_IC,                  // Instruction cache
maintenance           AccType_DCZVA,               // DC ZVA instructions
with PAN permission  AccType_ATPAN,                // Address translation
checks               AccType_AT};                  // Address translation

```

Is replaced with:

```

// AccessType
// =====

enumeration AccessType {
    AccessType_IFETCH, // Instruction FETCH
    AccessType_GPR,    // Software load/store to a General Purpose Register
    AccessType_ASIMD,  // Software ASIMD extension load/store instructions
    AccessType_SVE,    // Software SVE load/store instructions
    AccessType_SME,    // Software SME load/store instructions
    AccessType_IC,     // Sysop IC
    AccessType_DC,     // Sysop DC (not DC {Z,G,GZ}VA)
    AccessType_DCZero, // Sysop DC {Z,G,GZ}VA
    AccessType_AT,     // Sysop AT
    AccessType_NV2,    // NV2 memory redirected access
    AccessType_TRBE,   // Trace Buffer access
    AccessType_GPTW,   // Granule Protection Table Walk
    AccessType_TTW     // Translation Table Walk
};

```

The AccessDescriptor type that reads:

```
type AccessDescriptor is (
    boolean transactional,
    MPAMInfo mpam,
    AccType acctype)
```

Is updated to read:

```
// AccessDescriptor
// =====
// Memory access or translation invocation attributes that steer architectural
// behavior

type AccessDescriptor is (
    AccessType acctype,
    bits(2) el,                // Acting EL for the access
    SecurityState ss,          // Acting Security State for the access
    boolean acqsc,              // Acquire with Sequential Consistency
    boolean acqpc,              // FEAT_LRCPC: Acquire with Processor
    Consistency
    boolean relsc,              // Release with Sequential Consistency
    boolean limitedordered,     // FEAT_LOR: Acquire/Release with limited ordering
    boolean exclusive,          // Access has Exclusive semantics
    boolean atomicop,           // FEAT_LSE: Atomic read-modify-write access
    MemAtomicOp modop,          // FEAT_LSE: The modification operation in the 'atomicop'
    access
    boolean nontemporal,        // Hints the access is non-temporal
    boolean read,               // Read from memory or only require read
    permissions
    boolean write,              // Write to memory or only require write
    permissions
    CacheOp cacheop,            // DC/IC: Cache operation
    CacheOpScope opscope,       // DC/IC: Scope of cache operation
    CacheType cachetype,        // DC/IC: Type of target cache
    boolean pan,                // FEAT_PAN: The access is subject to
    PSTATE.PAN
    boolean transactional,       // FEAT_TME: Access is part of a transaction
    boolean nonfault,           // SVE: Non-faulting load
    boolean firstfault,         // SVE: First-fault load
    boolean first,              // SVE: First-fault load for the first
    active element
    boolean contiguous,          // SVE: Contiguous load/store not gather load/
    scatter store
    boolean streamingsve,        // SME: Access made by PE while in streaming SVE
    mode
    boolean ls64,                // FEAT_LS64: Accesses by accelerator support
    loads/stores
    boolean mops,                // FEAT_MOPS: Memory operation (CPY/SET)
    accesses
    boolean a32lsmd,             // A32 Load/Store Multiple Data access
    boolean tagchecked,          // FEAT_MTE2: Access is tag checked
    boolean tagaccess,           // FEAT_MTE: Access targets the tag bits
    MPAMInfo mpam                // FEAT_MPAM: MPAM information
)
```

## 2.7 C17811

In section I5.8.32 (ERR<n>STATUS, Error Record Primary Status Register, n = 0 - 65534), under the heading 'Accessing the ERR<n>STATUS', the text that reads:

To ensure correct and portable operation, when software is clearing the valid fields in the register to allow new errors to be recorded, Arm recommends that software:

- Read ERR<n>STATUS and determine which fields need to be cleared to zero.
- Write ones to all the W1C fields that are nonzero in the read value.
- Write zero to all the W1C fields that are zero in the read value.
- Write zero to all the RW fields.

is clarified to read:

To ensure correct and portable operation, when software is clearing the valid fields in the register to allow new errors to be recorded, Arm recommends that software:

- Read ERR<n>STATUS and determine which fields need to be cleared to zero.
- In a single write to ERR<n>STATUS:
  - Write ones to all the W1C fields that are nonzero in the read value.
  - Write zero to all the W1C fields that are zero in the read value.
  - Write zero to all the RW fields.
- Read back ERR<n>STATUS after the write to confirm no new fault has been recorded.

## 2.8 E17996

In section J1.2.3 (aarch32/functions) and J1.1.3 (aarch64/functions), the previous stub functions AArch32.PhysicalErrorSyndrome() and AArch64.PhysicalErrorSyndrome() respectively are now defined as:

```
// AArch32.PhysicalErrorSyndrome()
// =====
// Generate SError syndrome.

bits(16) AArch32.PhysicalErrorSyndrome()
    bits(32) syndrome = Zeros(32);
    FaultRecord fault = GetSavedFault();
    boolean long_format = TTBCR.EAE == '1';
    syndrome = AArch32.CommonFaultStatus(fault, long_format);
    return syndrome<15:0>;

// AArch64.PhysicalErrorSyndrome()
// =====
// Generate SError syndrome.

bits(25) AArch64.PhysicalErrorSyndrome(boolean implicit_esb)
    bits(25) syndrome = Zeros(25);
    FaultRecord fault = GetSavedFault();
    ErrorState errorstate = AArch64.PEErrorState(fault);
```

```

if errorstate == ErrorState_Uncategorized then
    syndrome = Zeros(25);
elsif errorstate == ErrorState_IMPDEF then
    syndrome<24> = '1'; // IDS
    syndrome<23:0> = bits(24) IMPLEMENTATION_DEFINED "IMPDEF ErrorState";
else
    syndrome<24> = '0'; // IDS
    syndrome<13> = (if implicit_esb then '1' else '0'); // IESB
    syndrome<12:10> = AArch64.EncodeAsyncErrorSyndrome(errorstate); // AET
    syndrome<5:0> = '010001'; // DFSC
return syndrome;

```

A new enumeration ErrorState is added in the same section, which is used instead of the errortype member of FaultRecord and PhysMemRetStatus:

```

enumeration ErrorState {ErrorState_UC, // Uncontainable
                        ErrorState_UEU, // Unrecoverable state
                        ErrorState_UEO, // Restartable state
                        ErrorState_UER, // Recoverable state
                        ErrorState_CE, // Corrected
                        ErrorState_Uncategorized,
                        ErrorState_IMPDEF};

```

A new function AArch32.CommonFaultStatus() is added to section J1.2.2 (aarch32/exceptions):

```

// AArch32.CommonFaultStatus()
// =====
// Return the common part of the fault status on reporting a Data
// or Prefetch Abort.

bits(32) AArch32.CommonFaultStatus(FaultRecord fault, boolean long_format)
    bits(32) target = Zeros(32);
    if HaveRASExt() && IsAsyncAbort(fault) then
        ErrorState errstate = AArch32.PEErrorState(fault);
        target<15:14> = AArch32.EncodeAsyncErrorSyndrome(errstate); // AET
    if IsExternalAbort(fault) then target<12> = fault.extflag; // ExT
    target<9> = if long_format then '1' else '0'; // LPAE
    if long_format then // Long-
descriptor format
        target<5:0> = EncodeLDFSC(fault.statuscode, fault.level); // STATUS
    else // Short-
descriptor format
        target<10,3:0> = EncodeSDFSC(fault.statuscode, fault.level); // FS
    return target;

```

A new function GetSavedFault() is added to section J1.3.3 (shared/functions):

```

// GetSavedFault()
// =====
// Return the saved asynchronous fault.

FaultRecord GetSavedFault();

```

## 2.9 D18330

*Arm® Architecture Reference Manual for A-profile architecture, Issue I.a* is somewhat inconsistent in its use of 'prefetch' and 'preload' to describe the bringing in of items into caches either by hardware prediction or as a result of some prefetch or preload instructions.

In future versions of *Arm® Architecture Reference Manual for A-profile architecture*, this will be cleaned up. The term 'prefetch' will be used for this functionality, with 'hardware prefetch' used where the prefetch is predicted by hardware, and 'software prefetch' used where the prefetch is prompted by particular instructions (such as the AArch64 PRFM or AArch32 PLD instructions).

## 2.10 D18465

In section D17.2.119 (SCTLR\_EL2, System Control Register (EL2)), for all of the bits that are described as having a function when `HCR_EL2.E2H==1 && HCR_EL2.TGE==1` and being **RESO** otherwise, it is clarified that these bits:

- Are **RESO** when `HCR_EL2.E2H==0`, so software should write the value 0.
- Are ignored by hardware when `HCR_EL2.E2H==1 && HCR_EL2.TGE==0`, but software doesn't have to set the value 0.
- Have their described effect when `HCR_EL2.E2H==1 && HCR_EL2.TGE==1`.

## 2.11 R18485

In section I5.8.8 (ERRDEVAFF, Device Affinity Register), the following text is added to the end of the Purpose section:

Depending on the **IMPLEMENTATION DEFINED** nature of the system, it might be possible that ERRDEVAFF is read before system firmware has configured the group of error records and/or the PE or group of PEs that the group of error records has affinity with. When this is the case, ERRDEVAFF reads as zero.

## 2.12 D18520

In section I5.8.31 (ERR<n>PFGF, Pseudo-fault Generation Feature Register, n = 0 - 65534), the text in MV, bit [12] that reads:

0b0 When an injected error is recorded, the node might update ERR<n>MISC<m>. If any syndrome is recorded by the node in ERR<n>MISC<m>, then ERR<n>STATUS.MV is set to 0b1. ERR<n>PFGCTL.MV is **RESO**.

is updated to read:



0b0 ERR<n>PFGCTL.MV not supported. When an injected error is recorded, the node might update ERR<n>MISC<m>. If any syndrome is recorded by the node in ERR<n>MISC<m>, then ERR<n>STATUS.MV is set to 0b1. If the node always sets ERR<n>.STATUS.MV to 0b1 when recording an injected error, then ERR<n>PFGCTL.MV might be RAO/WI. Otherwise, ERR<n>PFGCTL.MV is **RESO**.

Corresponding updates are made to section I5.8.30 (ERR<n>PFGCTL, Pseudo-fault Generation Control Register, n = 0 - 65534), for bit [12] 'when the node supports this control'. Similar corrections are made for the ERR<n>PFGF.AV and ERR<n>PFGCTL.AV controls.

## 2.13 D18736

In section I5.8.5 (ERRCRIC0, Critical Error Interrupt Configuration Register 0), under the heading 'Accessing the ERRCRIC0', the following text is added:

If the implementation does not use the recommended layout for the ERRIRQCR<n> registers, accesses to ERRCRIC0 are IMPLEMENTATION DEFINED.

ERRCRIC0 ignores writes if all of the following are true:

- The implementation uses the recommended layout for the ERRIRQCR<n> registers.
- ERRCRICR2.NSMI configures the physical address space for message signaled interrupts as Secure.
- Accessed as a Non-secure access.

The equivalent changes are made in the following sections:

- I5.8.6 (ERRCRICR1, Critical Error Interrupt Configuration Register 1).
- I5.8.7 (ERRCRICR2, Critical Error Interrupt Configuration Register 2).
- I5.8.11 (ERRERICR0, Error Recovery Interrupt Configuration Register 0).
- I5.8.12 (ERRERICR1, Error Recovery Interrupt Configuration Register 1).
- I5.8.13 (ERRERICR2, Error Recovery Interrupt Configuration Register 2).
- I5.8.14 (ERRFHICR0, Fault Handling Interrupt Configuration Register 0).
- I5.8.15 (ERRFHICR1, Faulting Handling Interrupt Configuration Register 1).
- I5.8.16 (ERRFHICR2, Faulting Handling Interrupt Configuration Register 2).

In section I5.8.7 (ERRCRICR2, Critical Error Interrupt Configuration Register 2), the text in the description of NSMSI, bit [6], that reads:

When the component supports configuring the Security attribute for messaged signaled interrupts and the component does not allow Non-secure writes to ERRCRICR2:

Security attribute. Defines the physical address space for message signaled interrupts.

0b0 Secure. 0b1 Non-secure.

The reset behavior of this field is:

- On a Error recovery reset, this field resets to an IMPLEMENTATION DEFINED VALUE.

When the component allows Non-secure writes to ERRCRICR2:

Reserved, **RES0**. Security attribute. Defines the physical address space for message signaled interrupts. The Security attribute used for message signaled interrupts is Non-secure.

is changed to read:

When the component supports configuring the physical address space for message signaled interrupts:

Non-secure message signaled interrupt. Defines the physical address space for message signaled interrupts.

0b0 Secure physical address space. 0b1 Non-secure physical address space.

The reset behavior of this field is:

- On an Error recovery reset, this field resets to an IMPLEMENTATION DEFINED VALUE.

Accessing this field has the following behavior:

- If accessed as a Non-secure access, access to this field is RES1.
- Otherwise, access to this field is RW.

The equivalent changes are made in the following sections:

- I5.8.13 (ERRERICR2, Error Recovery Interrupt Configuration Register 2).
- I5.8.16 (ERRFHICR2, Faulting Handling Interrupt Configuration Register 2).

## 2.14 R18746

In section B2.7.2 (Device memory), in the subsection 'Multi-register loads and stores that access Device memory', the following paragraph is added:

The architecture permits that the non-speculative execution of an instruction that loads or stores more than one general-purpose or SIMD and floating-point register might result in repeated accesses to the same address.

The equivalent edit is made in section E2.8.2 (Device Memory), in the subsection 'Multi-register loads and stores that access Device memory'.

## 2.15 D18800

In section D17.5.17 (PMUSERENR\_ELO, Performance Monitors User Enable Register), the EN, bit [0] description is updated to read:

Enable ELO accesses to the Performance Monitor registers. This applies to the following register accesses:

AArch64:

- MRS or MSR accesses to PMCR\_ELO, PMOVSLR\_ELO, PMSELR\_ELO, PMCCNTR\_ELO, PMXEVTYPER\_ELO, PMXEVCNTR\_ELO, PMCNTENSET\_ELO, PMCNTENCLR\_ELO, PMOVSSSET\_ELO, PMEVCNTR<n>\_ELO, PMEVTYPER<n>\_ELO, PMCCFILTR\_ELO.
- MSR accesses to PMSWINC\_ELO.
- MRS accesses to PMCEID0\_ELO, PMCEID1\_ELO.

AArch32:

- MRC and MCR accesses to PMCR, PMOVSRL, PMSELR, PMCCNTR, PMXEVTYPER, PMXEVCNTR, PMCNTENSET, PMCNTENCLR, PMOVSSSET, PMEVCNTR<n>, PMEVTYPER<n>, PMCCFILTR.
- MCR accesses to PMSWINC.
- MRC accesses to PMCEID0, PMCEID1.
- If FEAT\_PMUv3p1 is implemented, MRC accesses to PMCEID2, and PMCEID3.

0b0 ELO access to the specified registers is trapped, unless access is enabled by another field in this register.

0b1 ELO access to the specified registers is allowed, unless trapped by a higher priority exception.

When not enabled by any of the PMUSERENR\_ELO.{ER, CR, SW, EN} controls, an accesses to the register at ELO is trapped to EL1, or to EL2 when EL2 is implemented and enabled for the current Security state and HCR\_EL2.TGE is 1. Trapped MRS and MSR accesses are reported using EC syndrome value 0x18. Trapped MRC and MCR accesses are reported using EC syndrome value 0x03. Trapped MRRC and MCRR accesses are reported using EC syndrome value 0x04.

The reset behavior of this field is:

- On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

Equivalent changes are made to the {ER, CR, SW} fields, and to the PMUSERENR.{ER, CR, SW, EN} fields in section G8.4.18 (PMUSERENR, Performance Monitors User Enable Register).

## 2.16 D18823

In section J1.1.3 (aarch64/functions), the function CalculateBottomPACBit(), reading:

```
integer CalculateBottomPACBit(bit top_bit)
    integer tsz_field;
    boolean using64k;
    Constraint c;

    if PtrHasUpperAndLowerAddRanges() then
        assert S1TranslationRegime() IN {EL1, EL2};
        if S1TranslationRegime() == EL1 then
            // EL1 translation regime registers
            tsz_field = if top_bit == '1' then UInt(TCR_EL1.T1SZ) else
                UInt(TCR_EL1.T0SZ);
            using64k = if top_bit == '1' then TCR_EL1.TG1 == '11' else TCR_EL1.TG0
                == '01';
        else
            // EL2 translation regime registers
            assert HaveEL(EL2);
            tsz_field = if top_bit == '1' then UInt(TCR_EL2.T1SZ) else
                UInt(TCR_EL2.T0SZ);
            using64k = if top_bit == '1' then TCR_EL2.TG1 == '11' else TCR_EL2.TG0
                == '01';
        else
            tsz_field = if PSTATE.EL == EL2 then UInt(TCR_EL2.T0SZ) else
                UInt(TCR_EL3.T0SZ);
            using64k = if PSTATE.EL == EL2 then TCR_EL2.TG0 == '01' else TCR_EL3.TG0 ==
                '01';

        max_limit_tsz_field = (if !HaveSmallTranslationTableExt() then 39 else if
            using64k then 47 else 48);
        if tsz_field > max_limit_tsz_field then
            // TCR_ELx.TySZ is out of range
            c = ConstrainUnpredictable(Unpredictable_RESTnSZ);
            assert c IN {Constraint_FORCE, Constraint_NONE};
            if c == Constraint_FORCE then tsz_field = max_limit_tsz_field;
            tszmin = if using64k && AArch64.VAMax() == 52 then 12 else 16;
            if tsz_field < tszmin then
                c = ConstrainUnpredictable(Unpredictable_RESTnSZ);
                assert c IN {Constraint_FORCE, Constraint_NONE};
                if c == Constraint_FORCE then tsz_field = tszmin;
            return (64-tsz_field);
```

is updated to read:

```
integer CalculateBottomPACBit(bit top_bit)
    Regime regime;
    S1TTWParams walkparams;
    integer bottom_PAC_bit;

    // There is no distinction between AccType_NORMAL and AccType_IFETCH
    // when determining the translation regime
    regime = TranslationRegime(PSTATE.EL, AccType_NORMAL);

    walkparams = AArch64.GetS1TTWParams(regime, Replicate(top_bit, 64));
    bottom_PAC_bit = 64 - UInt(AArch64.PACEffectiveTxSZ(walkparams));

    return bottom_PAC_bit;
```

In section J1.1.3 (aarch64/functions), the function AArch64.PACEffectiveTxSZ() is added:

```
// AArch64.PACEffectiveTxSZ()
```

```
// =====
// Compute the effective value for TxSZ used to determine the placement of the PAC
// field

bits(6) AArch64.PACEffectiveTxSZ(S1TTWParams walkparams)
    constant integer slmaxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
    constant integer slmintxsz = AArch64.S1MinTxSZ(walkparams.ds, walkparams.tgx);

    if AArch64.S1TxSZFaults(walkparams) then
        if ConstrainUnpredictable(Unpredictable_RESTnSZ) == Constraint_FORCE then
            if UInt(walkparams.txsz) < slmintxsz then
                return slmintxsz<5:0>;
            if UInt(walkparams.txsz) > slmaxtxsz then
                return slmaxtxsz<5:0>;
        elseif UInt(walkparams.txsz) < slmintxsz then
            return slmintxsz<5:0>;
        elseif UInt(walkparams.txsz) > slmaxtxsz then
            return slmaxtxsz<5:0>;

    return walkparams.txsz;
```

In section J1.1.5 (aarch64/translation), the code within the function AArch64.GetS1TTWParams(), reading:

```
maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
mintxsz = AArch64.S1MinTxSZ(walkparams.ds, walkparams.tgx);
if UInt(walkparams.txsz) > maxtxsz then
    if !(boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value above maximum")
then
    walkparams.txsz = maxtxsz<5:0>;
elseif !Have52BitVAExt() && UInt(walkparams.txsz) < mintxsz then
    if !(boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value below minimum")
then
    walkparams.txsz = mintxsz<5:0>;
```

is removed.

In section J1.1.5 (aarch64/translation), the code within the function AArch64.GetS2TTWParams(), reading:

```
maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
mintxsz = AArch64.S2MinTxSZ(walkparams.ds, walkparams.tgx, slaarch64);
if UInt(walkparams.txsz) > maxtxsz then
    if !(boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value above maximum")
then
    walkparams.txsz = maxtxsz<5:0>;
elseif !Have52BitPAExt() && UInt(walkparams.txsz) < mintxsz then
    if !(boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value below minimum")
then
    walkparams.txsz = mintxsz<5:0>;
```

is removed.

In section J1.1.5 (aarch64/translation), the function AArch64.S1InvalidTxSZ(), reading:

```
boolean AArch64.S1InvalidTxSZ(S1TTWParams walkparams)
    mintxsz = AArch64.S1MinTxSZ(walkparams.ds, walkparams.tgx);
    maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);

    return UInt(walkparams.txsz) < mintxsz || UInt(walkparams.txsz) > maxtxsz;
```

is updated to read:

```
boolean AArch64.S1TxSZFaults(S1TTWParams walkparams)
{
    mintxs = AArch64.S1MinTxSZ(walkparams.ds, walkparams.tgx);
    maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);

    if UInt(walkparams.txsz) < mintxs then
        return (Have52BitVAExt() ||
                boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value below minimum");
    if UInt(walkparams.txsz) > maxtxsz then
        return boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value above maximum";

    return FALSE;
}
```

In section J1.1.5 (aarch64/translation), the function AArch64.S2InvalidTxSZ(), reading:

```
boolean AArch64.S2InvalidTxSZ(S2TTWParams walkparams, boolean slaarch64)
{
    mintxs = AArch64.S2MinTxSZ(walkparams.ds, walkparams.tgx, slaarch64);
    maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
    return UInt(walkparams.txsz) < mintxs || UInt(walkparams.txsz) > maxtxsz;
}
```

is updated to read:

```
boolean AArch64.S2TxSZFaults(S2TTWParams walkparams, boolean slaarch64)
{
    mintxs = AArch64.S2MinTxSZ(walkparams.ds, walkparams.tgx, slaarch64);
    maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);

    if UInt(walkparams.txsz) < mintxs then
        return (Have52BitPAExt() ||
                boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value below minimum");
    if UInt(walkparams.txsz) > maxtxsz then
        return boolean IMPLEMENTATION_DEFINED "Fault on TxSZ value above maximum";

    return FALSE;
}
```

In section J1.1.5 (aarch64/translation), the code within the function AArch64.S1Translate(), reading:

```
if (AArch64.S1InvalidTxSZ(walkparams) ||
    (!ispriv && walkparams.eopd == '1') ||
    (!ispriv && walkparams.nfd == '1' && IsDataAccess(acctype) &&
    TSTATE.depth > 0) ||
    (!ispriv && walkparams.nfd == '1' && acctype == AccType_NONFAULT) ||
    AArch64.VAIsOutOfRange(va, acctype, regime, walkparams)) then
    fault.statuscode = Fault_Translation;
    fault.level = 0;
    return (fault, AddressDescriptor UNKNOWN);
```

is updated to read:

```
constant integer slmintxs = AArch64.S1MinTxSZ(walkparams.ds, walkparams.tgx);
constant integer slmaxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
if AArch64.S1TxSZFaults(walkparams) then
    fault.statuscode = Fault_Translation;
    fault.level = 0;
    return (fault, AddressDescriptor UNKNOWN);
elseif UInt(walkparams.txsz) < slmintxs then
    walkparams.txsz = slmintxs<5:0>;
elseif UInt(walkparams.txsz) > slmaxtxsz then
    walkparams.txsz = slmaxtxsz<5:0>;
```

```

if AArch64.VAIsOutOfRange(va, acctype, regime, walkparams) then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

if !ispriv && walkparams.e0pd == '1' then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

if !ispriv && walkparams.nfd == '1' && IsDataAccess(acctype) && TSTATE.depth > 0
then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

if !ispriv && walkparams.nfd == '1' && acctype == AccType_NONFAULT then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

```

In section J1.1.5 (aarch64/translation), the code within the function AArch64.S2Translate(), reading:

```

if (AArch64.S2InvalidTxSZ(walkparams, slaarch64) ||
    AArch64.S2InvalidSL(walkparams) ||
    AArch64.S2InconsistentSL(walkparams) ||
    AArch64.IPAIsOutOfRange(ipa.paddress.address, walkparams)) then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

```

is updated to read:

```

constant integer s2mintxsz = AArch64.S2MinTxSZ(walkparams.ds, walkparams.tgx,
slaarch64);
constant integer s2maxtxsz = AArch64.MaxTxSZ(walkparams.tgx);
if AArch64.S2TxSZFaults(walkparams, slaarch64) then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);
elseif UInt(walkparams.txsz) < s2mintxsz then
    walkparams.txsz = s2mintxsz<5:0>;
elseif UInt(walkparams.txsz) > s2maxtxsz then
    walkparams.txsz = s2maxtxsz<5:0>;

if AArch64.S2InvalidSL(walkparams) || AArch64.S2InconsistentSL(walkparams) then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

if AArch64.IPAIsOutOfRange(ipa.paddress.address, walkparams) then
    fault.statuscode = Fault_Translation;
    fault.level      = 0;
    return (fault, AddressDescriptor UNKNOWN);

```

## 2.17 C18842

In section I5.5.14 (AMDEVARCH, Activity Monitors Device Architecture Register), the text in the ARCHID, bits [15:0] description that reads:

For AMU:

- Bits [15:12] are the architecture version, 0x0.
- Bits [11:0] are the architecture part number, 0xA66.

This corresponds to AMU architecture version AMUv1.

is changed to read:

For AMU:

- Bits [19:16] are the minor architecture version, 0x0.
- Bits [15:12] are the major architecture version, 0x0.
- Bits [11:0] are the architecture part number, 0xA66.

This corresponds to a generic AMU, version 1.0.

## 2.18 C18843

The current description of FEAT\_LPA2 in *Arm® Architecture Reference Manual for A-profile architecture, Issue I.a* lacks clarity between the ability to describe the size of the output address as having 52 bits, and there being 52 bits of physical address. This will be rectified in a future release of *Arm® Architecture Reference Manual for A-profile architecture*.

## 2.19 D18853

In section D17.2.107 (RGSR\_EL1, Random Allocation Tag Seed Register), the field descriptions are changed to read:

When GCR\_EL1.RRND == 0:

Bits [63:24] Reserved, RES0.

SEED, bits [23:8] Seed register used for generating values returned by RandomAllocationTag(). The reset behavior of this field is:

- On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

Bits [7:4] Reserved, RES0.



TAG, bits [3:0] Tag generated by the most recent IRG instruction. The reset behavior of this field is:

- On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

When GCR\_EL1.RRND == 1:

Bits [63:56] Reserved, RES0.

SEED, bits [55:8] IMPLEMENTATION DEFINED Note: Software is recommended to avoid writing SEED[15:0] with a value of zero, unless this has been generated by the PE in response to an earlier value with SEED being non-zero. The reset behavior of this field is:

- On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

Bits [7:4] Reserved, RES0.

TAG, bits [3:0] Tag generated by the most recent IRG instruction. The reset behavior of this field is:

- On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

## 2.20 D18889

In section C5.2.18 (SPSR\_EL1, Saved Program Status Register (EL1)), in the 'TCO, bit [25]' field, the text that reads:

When FEAT\_MTE is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether this field is **RES0** or behaves as if FEAT\_MTE is implemented.

is corrected to read:

When FEAT\_MTE2 is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether this field is **RES0** or behaves as if FEAT\_MTE2 is implemented.

The equivalent changes are made in the following sections:

- C5.2.19 (SPSR\_EL2, Saved Program Status Register (EL2)).
- C5.2.20 (SPSR\_EL3, Saved Program Status Register (EL3)).
- D13.3.14 (DPSR\_ELO, Debug Saved Program Status Register).

In section C5.2.26 (TCO, Tag Check Override), in the 'purpose' field, the text that reads:

When FEAT\_MTE is implemented, this register allows tag checks to be disabled globally.

When FEAT\_MTE is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether this register is **RES0** or behaves as if FEAT\_MTE is implemented.

is corrected to read:

Allows tag checks to be disabled globally.

When FEAT\_MTE2 is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether this register is **RES0** or behaves as if FEAT\_MTE2 is implemented.

In section D1.4.1 (PSTATE fields that are meaningful in AArch64 state) the text in the 'Additional details' column for the TCO entry of the table in R\_PCDTX that reads:

If FEAT\_MTE2 is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether the PSTATE.TCO bit is **RES0** or behaves as if FEAT\_MTE is implemented.

is corrected to read:

If FEAT\_MTE2 is not implemented, it is **CONSTRAINED UNPREDICTABLE** whether the PSTATE.TCO bit is **RES0** or behaves as if FEAT\_MTE2 is implemented.

In section H2.4.1 (PSTATE in Debug state), the text that reads:

When FEAT\_MTE is implemented, if Memory-access mode is enabled and PSTATE.TCO is 0, reads and writes to the external debug interface DTR registers are **CONSTRAINED UNPREDICTABLE**, with the following permitted behaviors:

- The PE behaves as if PSTATE.TCO is 0. That is, the load or store operation performs the tag check if required.
- The PE behaves as if PSTATE.TCO is 1. That is, the load or store operation does not perform the tag check.

is corrected to read:

When FEAT\_MTE2 is implemented, if Memory-access mode is enabled and PSTATE.TCO is 0, reads and writes to the external debug interface DTR registers are **CONSTRAINED UNPREDICTABLE**, with the following permitted behaviors:

- The PE behaves as if PSTATE.TCO is 0. That is, the load or store operation performs the tag check if required.
- The PE behaves as if PSTATE.TCO is 1. That is, the load or store operation does not perform the tag check.

## 2.21 C19027

In section D11.11.3 (Common event numbers), subsection 'Common microarchitectural events', the following text is added to the descriptions of MEM\_ACCESS\_CHECKED (0x4024), MEM\_ACCESS\_CHECKED\_RD (0x4025), and MEM\_ACCESS\_CHECKED\_WR (0x4026):

It is **IMPLEMENTATION DEFINED** whether the counter increments on a Tag Checked access made when Tag Check Faults are configured to be ignored by SCTLR\_ELx.TCF or SCTLR\_ELx.TCF0.

## 2.22 C19047

In section D17.2.27 (CLIDR\_EL1, Cache Level ID Register), the following Note is added to the descriptions of LoUU, bits [29:27], and LoUIS, bits [23:21]:

Note: This field does not describe the requirements for instruction cache invalidation. See CTR\_ELO.DIC.

The equivalent changes are made in section G8.2.27 (CLIDR, Cache Level ID Register).

## 2.23 D19116

In section D17.11.21 (CNTPS\_CTL\_EL1, Counter-timer Physical Secure Timer Control register), the following text is added under 'Configurations':

This register is present only when EL3 is implemented. Otherwise, direct accesses to CNTPS\_CTL\_EL1 are **UNDEFINED**.

Equivalent changes are made in the following sections:

- D17.11.23 (CNTPS\_CVAL\_EL1, Counter-timer Physical Secure Timer CompareValue register).
- D17.11.24 (CNTPS\_TVAL\_EL1, Counter-timer Physical Secure Timer TimerValue register).

## 2.24 D19121

In section D17.2.118 (SCTLR\_EL1, System Control Register (EL1)), in field 'C, bit [2]', the text that reads:

When the value of the HCR\_EL2.DC bit is 1, the PE ignores SCTLR.C. This means that Non-secure EL0 and Non-secure EL1 data accesses to Normal memory are Cacheable.

When FEAT\_VHE is implemented, and the value of HCR\_EL2.{E2H, TGE} is {1, 1}, this bit has no effect on the PE.

is changed to read:

When the Effective value of the HCR\_EL2.DC bit in the current Security state is 1, the PE ignores SCTLR\_EL1.C. This means that EL0 and EL1 data accesses to Normal memory are Cacheable.

When FEAT\_VHE is implemented, and the Effective value of HCR\_EL2.{E2H, TGE} is {1, 1}, this bit has no effect on the PE.

Similarly in field 'M, bit [0]', the text that reads:

If the value of HCR\_EL2.{DC, TGE} is not {0, 0} then in Non-secure state the PE behaves as if the value of the SCTLR\_EL1.M field is 0 for all purposes other than returning the value of a direct read of the field.

When FEAT\_VHE is implemented, and the value of HCR\_EL2.{E2H, TGE} is {1, 1}, this bit has no effect on the PE.

is changed to read:

If the Effective value of HCR\_EL2.{DC, TGE} in the current Security state is not {0, 0} then the PE behaves as if the value of the SCTLR\_EL1.M field is 0 for all purposes other than returning the value of a direct read of the field.

When FEAT\_VHE is implemented, and the Effective value of HCR\_EL2.{E2H, TGE} is {1, 1}, this bit has no effect on the PE.

The equivalent changes are made in section G8.2.126 (SCTLR, System Control Register).

## 2.25 D19162

In section B2.3.10 (Restrictions on the effects of speculation), in the subsection 'Restrictions on the effects of speculation from Armv8.5', the text that reads:

Any System register read under speculation to a register that is not architecturally accessible from the current Exception level cannot be used to form an address, to generate condition codes, or to generate SVE predicate values to be used by other instructions in the speculative sequence.

is updated to read:

Any read under speculation from a register that is not architecturally accessible from the current Exception level cannot be used to form an address, to generate condition codes, or to generate SVE predicate values to be used by other instructions in the speculative sequence.

The equivalent change is made in section E2.3.9 (Restrictions on the effects of speculation), in the subsection 'Further restrictions on the effects of speculation from Armv8.5'.

## 2.26 D19178

In section J1.1.3 (aarch64/functions), the function AddressSupportsLS64(), that reads as:

```
boolean AddressSupportsLS64(bits(64) address)
```

Is updated to read as:

```
boolean AddressSupportsLS64(bits(52) paddress);
```

The following changes are also made in the same section:

In MemStore64B(), the code that reads:

```
MemStore64B(bits(64) address, bits(512) value, AccType acctype)
    boolean iswrite = TRUE;
    constant integer size = 64;
    aligned = AArch64.CheckAlignment(address, size, acctype, iswrite);
    if !AddressSupportsLS64(address) then
        c = ConstrainUnpredictable(Unpredictable_LS64UNSUPPORTED);
        assert c IN {Constraint_LIMITED_ATOMICTY, Constraint_FAULT};
        if c == Constraint_FAULT then
            ...
        else
            // Accesses are not single-copy atomic above the byte level.
            for i = 0 to 63
                AArch64.MemSingle[address+8*i, 1, acctype, aligned] = value<7+8*i :
8*i>;
            else
                -- MemStore64BWithRet(address, value, acctype); // Return status is ignored
                by ST64B
    return;
```

Is updated to read:

```
MemStore64B(bits(64) address, bits(512) value, AccType acctype)
    boolean iswrite = TRUE;
    constant integer size = 64;
    aligned = AArch64.CheckAlignment(address, size, acctype, iswrite);
    AddressDescriptor memaddrdesc = AArch64.TranslateAddress(address, acctype,
iswrite,
                                                                    istagaccess, aligned,
size);

    // Check for aborts or debug exceptions
    if IsFault(memaddrdesc) then
        AArch64.Abort(address, memaddrdesc.fault);

    // Effect on exclusives
    if memaddrdesc.memattrs.shareability != Shareability_NSH then
        ClearExclusiveByAddress(memaddrdesc.paddress, ProcessorID(), 64);

    // Memory array access
    accdesc = CreateAccessDescriptor(acctype);
    if !AddressSupportsLS64(memaddrdesc.paddress.address) then
        c = ConstrainUnpredictable(Unpredictable_LS64UNSUPPORTED);
        assert c IN {Constraint_LIMITED_ATOMICTY, Constraint_FAULT};

        if c == Constraint_FAULT then
            ...
        else
            // Accesses are not single-copy atomic above the byte level.
            accdesc.acctype = AccType_ATOMIC;
            for i = 0 to size-1
                memstatus = PhysMemWrite(memaddrdesc, 1, accdesc, value<8*i+7:8*i>);
                if IsFault(memstatus) then
                    HandleExternalWriteAbort(memstatus, memaddrdesc, 1, accdesc);
                memaddrdesc.paddress.address = memaddrdesc.paddress.address+1;
            else
                memstatus = PhysMemWrite(memaddrdesc, size, accdesc, value);
                if IsFault(memstatus) then
                    HandleExternalWriteAbort(memstatus, memaddrdesc, size, accdesc);
    return;
```

In MemLoad64B(), the code that reads:

```
bits(512) MemLoad64B(bits(64) address, AccType acctype)
...

aligned = AArch64.CheckAlignment(address, size, acctype, iswrite);

if !AddressSupportsLS64(address) then
    c = ConstrainUnpredictable(Unpredictable_LS64UNSUPPORTED);
    assert c IN {Constraint_LIMITED_ATOMICTY, Constraint_FAULT};

    if c == Constraint_FAULT then
        // Generate a stage 1 Data Abort reported using the DFSC code of 110101.
        boolean secondstage = FALSE;
        boolean s2fslwalk = FALSE;
        FaultRecord fault = AArch64.ExclusiveFault(acctype, iswrite,
secondstage, s2fslwalk);
        AArch64.Abort(address, fault);
    else
        // Accesses are not single-copy atomic above the byte level
        for i = 0 to 63
            data<7+8*i : 8*i> = AArch64.MemSingle[address+8*i, 1, acctype,
aligned];
        return data;

    AddressDescriptor memaddrdesc;
    memaddrdesc = AArch64.TranslateAddress(address, acctype, iswrite, istagaccess,
aligned, size);

    // Check for aborts or debug exceptions
    if IsFault(memaddrdesc) then
        ...
        accdesc = CreateAccessDescriptor(acctype);
        PhysMemRetStatus memstatus;
        (memstatus, data) = PhysMemRead(memaddrdesc, size, accdesc);
        if IsFault(memstatus) then
            HandleExternalReadAbort(memstatus, memaddrdesc, size, accdesc);
        return data;
```

Is updated to read as:

```
bits(512) MemLoad64B(bits(64) address, AccType acctype)
...

aligned = AArch64.CheckAlignment(address, size, acctype, iswrite);

AddressDescriptor memaddrdesc;
memaddrdesc = AArch64.TranslateAddress(address, acctype, iswrite, istagaccess,
aligned, size);

// Check for aborts or debug exceptions
if IsFault(memaddrdesc) then
    ...
    accdesc = CreateAccessDescriptor(acctype);
    if !AddressSupportsLS64(memaddrdesc.paddress.address) then
        c = ConstrainUnpredictable(Unpredictable_LS64UNSUPPORTED);
        assert c IN {Constraint_LIMITED_ATOMICTY, Constraint_FAULT};

        if c == Constraint_FAULT then
            // Generate a stage 1 Data Abort reported using the DFSC code of 110101.
            boolean secondstage = FALSE;
            boolean s2fslwalk = FALSE;
            FaultRecord fault = AArch64.ExclusiveFault(acctype, iswrite,
secondstage, s2fslwalk);
            AArch64.Abort(address, fault);
        else
            // Accesses are not single-copy atomic above the byte level.
```

```

accdesc.acctype = AccType_ATOMIC;
for i = 0 to size-1
    PhysMemRetStatus memstatus;
    (memstatus, data<8*i+7:8*i>) = PhysMemRead(memaddrdesc, 1, accdesc);
    if IsFault(memstatus) then
        HandleExternalReadAbort(memstatus, memaddrdesc, 1, accdesc);
        memaddrdesc.paddress.address = memaddrdesc.paddress.address + 1;
else
    PhysMemRetStatus memstatus;
    (memstatus, data) = PhysMemRead(memaddrdesc, size, accdesc);
    if IsFault(memstatus) then
        HandleExternalReadAbort(memstatus, memaddrdesc, size, accdesc);
return data;

```

## 2.27 C19202

In section A2.2.1 (Additional functionality added to Armv8.0 in later releases), in the definition 'FEAT\_CSV2, FEAT\_CSV2\_2, and FEAT\_CSV2\_3, Cache Speculation Variant 2', the text that reads:

FEAT\_CSV2 adds a mechanism to identify if hardware cannot disclose information about whether branch targets trained in one hardware described context can control speculative execution in a different hardware described context.

is updated to read:

FEAT\_CSV2 adds a mechanism to identify if hardware cannot disclose information about whether branch targets, including those used by return instructions, trained in one hardware described context can control speculative execution in a different hardware described context.

In section B2.3.10 (Restrictions on the effects of speculation), in the subsection 'Restrictions on the effects of speculation from Armv8.5', the text that reads:

If FEAT\_CSV2 is implemented:

- Code running in one hardware-defined context (context1) cannot either exploitatively control, or predictively leak to, the speculative execution of code in a different hardware-defined context (context2), as a result of the behavior of any of the following resources:
  - Branch target prediction based on the branch targets used in context1.
    - This applies to both direct and indirect branches, but excludes the prediction of the direction of a conditional branch.

is updated to read:

If FEAT\_CSV2 is implemented:

- Code running in one hardware-defined context (context1) cannot either exploitatively control, or predictively leak to, the speculative execution of code in a different hardware-defined context (context2), as a result of the behavior of any of the following resources:
  - Branch target prediction based on the branch targets used in context1.
    - This applies to both direct and indirect branches, including those used by return instructions, but excludes the prediction of the direction of a conditional branch.

## 2.28 D19239

In section D17.2.49 (HCRX\_EL2, Extended Hypervisor Configuration Register), the text in the fields MSCEN, MCE2, CMOW, and SMPME that reads:

On a Warm reset, this field resets to an architecturally **UNKNOWN** value.

is corrected to read:

On a Warm reset:

- When EL3 is not implemented and EL2 is implemented, this field resets to 0.
- Otherwise, this field resets to an architecturally **UNKNOWN** value.

In the same register, the text in the fields VFNMI, VINMI, TALLINT, FGTnXS, FnXS, EnASR, EnALS, and EnAS0 that reads:

On a Warm reset, when EL3 is not implemented and EL2 is implemented, this field resets to 0.

is corrected to read:

On a Warm reset:

- When EL3 is not implemented and EL2 is implemented, this field resets to 0.
- Otherwise, this field resets to an architecturally **UNKNOWN** value.

## 2.29 D19275

In section D17.2.48 (HCR\_EL2, Hypervisor Configuration Register), in the description of FWB, bit [46], the following Note is removed:

When FEAT\_MTE2 is implemented, if the stage 1 page or block descriptor specifies the Tagged attribute, the final memory type is Tagged only if the final cacheable memory type is Inner and Outer Write-back cacheable and the final allocation hints are Read-Allocate, Write-Allocate.

## 2.30 R19370

In section E2.8.1 (Normal memory), after the text that reads:

Writes to a memory location with the Normal memory type that is either Non-cacheable or Write-Through cacheable for both the Inner and Outer Cacheability must reach the endpoint for that location in the memory system in finite time. Two writes to the same location, where at least one is using the Normal memory type, might be merged before they reach the endpoint unless there is an ordered-before relationship between the two writes.



The following text is added:

For the purposes of this requirement, the endpoint for a location in Conventional memory is the PoC.

## 2.31 D19372

In section D17.2.107 (RGSER\_EL1, Random Seed Allocation Tag Seed Register), the following text is added under 'Configurations':

When GCR\_EL1.RRND==0b0, direct and indirect reads and writes to the register appear to occur in program order relative to other instructions, without the need for any explicit synchronization.

## 2.32 E19440

In section H9.2.42 (EDSCR, External Debug Status and Control Register), in the fields RXfull, TXfull, RXO, TXU, TDA, SC2, HDE, and ERR, the following text is added:

When OLSR\_EL1.OSLK is 1, this bit can be indirectly read and written through the following System registers:

- MDSCR\_EL1.
- DBGDSCRext.

## 2.33 D19451

In section C6.2.378 (TLBI), in the 'Assembler symbols' subsection, the following statements are added to the definition of '<tlbi\_op>':

When FEAT\_RME is implemented, the following values are also valid:

PAALLOS	when op1 = 110, CRn = 1000, CRm = 0001, op2 = 100
RPAOS	when op1 = 110, CRn = 1000, CRm = 0100, op2 = 011
RPALOS	when op1 = 110, CRn = 1000, CRm = 0100, op2 = 111
PAALL	when op1 = 110, CRn = 1000, CRm = 0111, op2 = 100

## 2.34 D19452

Following the update communicated as D18736, in section I5.8.7 (ERRCRICR2, Critical Error Interrupt Configuration Register 2), the text that reads:

If accessed as a Non-secure access, access to this field is **RES1**.

is updated to read:

If accessed as a Non-secure or Realm access, access to this field is W1.

The equivalent changes are made in the following sections:

- I5.8.13 (ERRERICR2, Error Recovery Interrupt Configuration Register 2).
- I5.8.16 (ERRFHICR2, Faulting Handling Interrupt Configuration Register 2).

## 2.35 D19494

In section J1.3.3 (shared/functions/externalaborts) the function `IsSErrorEdgeTriggered()`, that reads as

```
boolean IsSErrorEdgeTriggered(bits(2) target_el, bits(25) syndrome)
    if HaveRASExt() then
        if HaveDoubleFaultExt() then
            return TRUE;

        if ELUsingAArch32(target_el) then
            if syndrome<11:10> != '00' then
                // AArch32 and not Uncontainable.
                return TRUE;
        else
            if syndrome<24> == '0' && syndrome<5:0> != '000000' then
                // AArch64 and neither IMPLEMENTATION_DEFINED syndrome nor
                Uncategorized.
                return TRUE;
    return boolean IMPLEMENTATION_DEFINED "Edge-triggered SError";
```

Is updated to read:

```
boolean IsSErrorEdgeTriggered()
    if HaveDoubleFaultExt() then
        return TRUE;
    else
        return boolean IMPLEMENTATION_DEFINED "Edge-triggered SError";
```

In section J1.1.2 (aarch64/exceptions/async), the function `AArch64.TakePhysicalSErrorException()`, that reads as:

```
AArch64.TakePhysicalSErrorException(boolean implicit_esb)
    ...

    bits(25) syndrome = Zeros(25);
    syndrome = AArch64.PhysicalSErrorSyndrome(implicit_esb);
    if IsSErrorEdgeTriggered(target_el, exception.syndrome) then
        ClearPendingPhysicalSError();
    ...
```

Is updated to read:

```
AArch64.TakePhysicalSErrorException(boolean implicit_esb)
    ...
```

```
bits(25) syndrome = AArch64.PhysicalErrorSyndrome(implicit_esb);  
if IsSErrorEdgeTriggered() then  
    ClearPendingPhysicalSError();  
...
```

In section J1.2.2 (aarch32/exceptions/async), similar changes are made to the function `AArch32.TakePhysicalSErrorException()`.

## 2.36 R19519

In section B2.3.10 (Restrictions on the effects of speculation), in the subsection 'Restrictions on the effects of speculation from Armv8.5', the sub-bullet point that reads:

- Data Value predictions based on data value from execution in context1.

is updated to include the following Note:

Note: `PSTATE.{N,Z,C,V}` values from context1 are not considered a data value for this purpose.

The equivalent change is made in section E2.3.9 (Restrictions on the effects of speculation), in the subsection 'Further restrictions on the effects of speculation from Armv8.5'.

In section C5.6.3 (DVP RCTX, Data Value Prediction Restriction by Context), the following Note is added:

Note: The prediction of the `PSTATE.{N,Z,C,V}` values is not considered a data value for this purpose.

The equivalent change is made in section G8.2.50 (DVPRCTX, Data Value Prediction Restriction by Context).

## 2.37 D19521

In section C5.2.25 (SVCR, Streaming Vector Control Register), for the field `ZA`, bit [1], the text that reads:

When a write to `SVCR.ZA` changes the value of `PSTATE.ZA`, the following applies:

When changed from 0 to 1, all implemented bits of the storage are set to zero. When changed from 1 to 0, there is no observable change to the storage.

Changes to this field do not have an affect on the SVE vector and predicate registers and FPSR.

is corrected to read:

When a write to `SVCR.ZA` changes the value of `PSTATE.ZA` from 0 to 1, all implemented bits of the storage are set to zero.

Changes to this field do not have an effect on the SVE vector and predicate registers and FPSR.

## 2.38 D19549

In section D11.11.3 (Common event numbers), in the subsection ‘Common microarchitectural events’, for each TRCEXTOUT<n> event, where <n> is 0 to 3, the text that reads:

This event must be implemented if FEAT\_ETE is implemented.

is updated to read:

This event must be implemented if FEAT\_ETE is implemented and the ETE implements External output <n>.

## 2.39 D19560

In section D17.2.26 (CCSIDR\_EL1, Current Cache Size Register), the text in LineSize, bits [2:0] when FEAT\_CCIDX is implemented that reads:

When FEAT\_MTE is implemented and enabled, where a cache only holds Allocation tags, this field is **RES0**.

is changed to read:

When FEAT\_MTE is implemented, where a cache only holds Allocation tags, this field is **RES0**.

The following text is added to LineSize, bits [2:0] when FEAT\_CCIDX is not implemented:

When FEAT\_MTE is implemented, where a cache only holds Allocation tags, this field is **RES0**.

## 2.40 D19561

In section D17.2.107 (RGSr\_EL1, Random Allocation Tag Seed Register), the text that reads:

When GCR\_EL1.RRND=0, direct and indirect reads and writes to the register appear to occur in program order relative to other instructions, without the need for any explicit synchronization.

is changed to read:

Direct and indirect reads and writes to the register appear to occur in program order relative to other instructions, without the need for any explicit synchronization.

## 2.41 D19581

In the function `AArch64.RestrictPrediction()` in section J1.1.4 (`aarch64/instrs`), the code that reads:

```
// If the instruction is executed at an EL lower than the specified
// level, it is treated as a NOP.
if UInt(target_el) > UInt(PSTATE.EL) then return;
```

Is updated to read:

```
// If the target EL is not implemented or the instruction is executed at an
// EL lower than the specified level, the instruction is treated as a NOP.
if !HaveEL(target_el) || UInt(target_el) > UInt(PSTATE.EL) then EndOfInstruction();
```

This affects the A64 System instructions in the following sections:

- C5.6.1 (CFP RCTX, Control Flow Prediction Restriction by Context).
- C5.6.2 (CPP RCTX, Cache Prefetch Prediction Restriction by Context).
- C5.6.3 (DVP RCTX, Data Value Prediction Restriction by Context).

An equivalent change is made in `AArch32.RestrictPrediction()` affecting the AArch32 System Registers in the following sections:

- G8.2.26 (CFPRCTX, Control Flow Prediction Restriction by Context).
- G8.2.34 (CPPRCTX, Cache Prefetch Prediction Restriction by Context).
- G8.2.50 (DVPRCTX, Data Value Prediction Restriction by Context).

## 2.42 D19642

In section D11.11.3 (Common event numbers), subsection ‘Common microarchitectural events’, the PMU events that read:

0x4025, MEM\_ACCESS\_RD\_CHECKED, Checked data memory access, read

0x4026, MEM\_ACCESS\_WR\_CHECKED, Checked data memory access, write

are corrected to read:

0x4025, MEM\_ACCESS\_CHECKED\_RD, Checked data memory access, read

0x4026, MEM\_ACCESS\_CHECKED\_WR, Checked data memory access, write

## 2.43 C19644

In section D11.11.3 (Common event numbers), subsection ‘Common microarchitectural events’, the text in the descriptions of MEM\_ACCESS\_CHECKED\_RD (0x4025) and MEM\_ACCESS\_CHECKED\_WR (0x4026) that reads:

Implementation of this optional event requires that FEAT\_MTE is implemented.

is corrected to read:

Implementation of this optional event requires that FEAT\_MTE2 is implemented.

This text is also added to the MEM\_ACCESS\_CHECKED (0x4024) event description.

## 2.44 D19647

In section D8.2.3 (Translation table base address register), the following text is added:

Direct writes to TTBR0\_ELx and TTBR1\_ELx occur in program order relative to one another, without the need for explicit synchronization. For any one translation, all indirect reads of TTBR0\_ELx and TTBR1\_ELx made as part of the translation observe only one point in that order of direct writes. Consistent with the general requirements for direct writes to System registers, direct writes to TTBRn\_ELx are not required to be observed by indirect reads until completion of a Context synchronization event.

A new subsection, ‘Example sequences for changing TTBRn\_ELx for AArch64’, is added after this text:

Example D8-1 Example software sequence for changing translation table base address and ASID value when TCR\_EL1.A1=1

Change TTBR0 to point to no valid entries Change TTBR1 (includes changing the ASID) Change TTBR0 to have valid entries in it ISB

Example D8-2 Example software sequence for changing translation table base address and ASID value when TCR\_EL1.A1=0

Change TTBR1 to point only at global entries Change TTBR0 (includes changing the ASID) Change TTBR1 to point at new tables, containing non-global entries ISB

## 2.45 C19649

In section B2.7.2 (Device Memory), in subsection ‘Reordering’, the bullet point in the note that reads:

The non-Reordering property is only required by the architecture to apply the order of arrival of accesses to a single memory-mapped peripheral of an **IMPLEMENTATION DEFINED** size, and is not required to have an impact on the order of observation of memory accesses to SDRAM. For this reason, there is no effect of the non-Reordering attribute on the ordering relations between accesses to different locations described in Ordering relations on page B2-165 as part of the formal definition of the memory model.

is updated to read:

The non-Reordering property is only required by the architecture to apply the order of arrival of accesses to a single memory-mapped peripheral of an **IMPLEMENTATION DEFINED** size, and is not required to have an impact on the order of observation of memory accesses to SDRAM. For this reason, there is no effect of the non-Reordering attribute on the ordering relations between accesses to different locations described in B2.3.3 Ordering relations on page B2-165 as part of the formal definition of the memory model. It does have an effect on the Peripheral Coherence Order described in section B2.3.7 (Completion and endpoint ordering).

## 2.46 D19680

In section C5.5.62 (TLBI VAE2, TLBI VAE2NXS, TLB Invalidate by VA, EL2), the accessibility pseudocode that reads:

```
elseif PSTATE.EL == EL2 then
    if HCR_EL2.E2H == '1' then
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL20, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
    else
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL2, VMID[],
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
elseif PSTATE.EL == EL3 then
    if !EL2Enabled() then
        UNDEFINED;
    elseif HCR_EL2.E2H == '1' then
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL20, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
    else
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL2, VMID[],
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
```

is corrected to read:

```
elseif PSTATE.EL == EL2 then
    if HCR_EL2.E2H == '1' then
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL20, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
    else
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL2, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
```

```
elseif PSTATE.EL == EL3 then
    if !EL2Enabled() then
        UNDEFINED;
    elseif HCR_EL2.E2H == '1' then
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL20, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
    else
        AArch64.TLBI_VA(SecurityStateAtEL(EL2), Regime_EL2, VMID_NONE,
        Shareability_NSH, TLBILevel_Any, TLBI_AllAttr, X[t, 64]);
```

The same change, from VMID[] to VMID\_NONE, is made in all the TLBI VAE2\*, TLBI VAE3\*, TLBI VALE2\*, and TLBI VALE3\* System instructions.

## 2.47 E19713

In section J1.3.3 (shared/functions), the contents of the HaveXXX() functions are updated to reflect the official feature names. For example:

```
boolean Have16bitVMID()
    return (HasArchVersion(ARMv8p1) && HaveEL(EL2) &&
    boolean IMPLEMENTATION_DEFINED "Has 16-bit VMID");
```

Is updated to read:

```
boolean Have16bitVMID()
    return IsFeatureImplemented(FEAT_VMID16);
```

## 2.48 D19741

In the function AArch64.WatchpointByteMatch() in section J1.1.1 (aarch64/debug), the code that reads:

```
if mask > bottom then
    ...
    if !IsOnes(DBGBVR_EL1[n]<63:top>) && !IsZero(DBGBVR_EL1[n]<63:top>) then
        if ConstrainUnpredictableBool(Unpredictable_DBGxVR_RESS) then
```

Is updated to read as:

```
if mask > bottom then
    ...
    if !IsOnes(DBGWVR_EL1[n]<63:top>) && !IsZero(DBGWVR_EL1[n]<63:top>) then
        if ConstrainUnpredictableBool(Unpredictable_DBGxVR_RESS) then
```

In the function AArch32.WatchpointByteMatch() in section J1.2.1 (aarch32/debug), the code that reads:

```
if mask > bottom then
    // If the DBGxVR<n>_EL1.RESS field bits are not a sign extension of the MSB
```



```
// of DBGWVR<n>_EL1.VA, it is UNPREDICTABLE whether they appear to be
// included in the match.
if !IsOnes(DBGWVR_EL1[n]<63:top>) && !IsZero(DBGWVR_EL1[n]<63:top>) then
    if ConstrainUnpredictableBool(Unpredictable_DBGWVR_RESS) then
        top = 63;
WVR_match = (vaddress<top:mask> == DBGWVR[n]<top:mask>);
```

Is updated to read as:

```
if mask > bottom then
    WVR_match = (vaddress<top:mask> == DBGWVR[n]<top:mask>);
```

## 2.49 D19753

In section J1.3.1 (shared/debug), the function Halt(), that reads as:

```
Halt(bits(6) reason, boolean is_async)

    CTI_SignalEvent(CrossTriggerIn_CrossHalt); // Trigger other cores to halt
    ...
```

Is updated to read:

```
Halt(bits(6) reason, boolean is_async)

    if HaveTME() && TSTATE.depth > 0 then
        FailTransaction(TMFailure_DBG, FALSE);

    CTI_SignalEvent(CrossTriggerIn_CrossHalt); // Trigger other cores to halt
    ...
```

## 2.50 C19772

In section C5.5.10 (TLBI ASIDE1, TLBI ASIDE1NXS, TLB Invalidate by ASID, EL1), in the subsection 'Executing TLBI ASIDE1, TLBI ASIDE1NXS instruction', the EL1 accessibility pseudocode that reads:

```
elseif EL2Enabled() && HCR_EL2.FB == '1' then
    if IsFeatureImplemented(FEAT_XS) && IsFeatureImplemented(FEAT_HCX) &&
        HCRX_EL2.FnXS == '1' then
        AArch64.TLBI_ASID(SecurityStateAtEL(EL1), Regime_EL10, VMID[],
            Shareability_ISH, TLBI_ExcludeXS, X[t, 64]);
```

is updated to read:

```
elseif EL2Enabled() && HCR_EL2.FB == '1' then
    if IsFeatureImplemented(FEAT_XS) && IsFeatureImplemented(FEAT_HCX) &&
        IsHCRXEL2Enabled() && HCRX_EL2.FnXS == '1' then
        AArch64.TLBI_ASID(SecurityStateAtEL(EL1), Regime_EL10, VMID[],
            Shareability_ISH, TLBI_ExcludeXS, X[t, 64]);
```

The same edits are made in the following sections:

- C5.5.29 (TLBI RVAAE1, TLBI RVAAE1NXS).
- C5.5.32 (TLBI RVAALE1, TLBI RAAVLE1NXS).
- C5.5.35 (TLBI RVAE1, TLBI RVAE1NXS).
- C5.5.44 (TLBI RVALE1, TLBI RAVLE1NXS).
- C5.5.53 (TLBI VAAE1, TLBI VAAE1NXS).
- C5.5.56 (TLBI VAALE1, TLBI VAALE1NXS).
- C5.5.59 (TLBI VAE1, TLBI VAE1NXS).
- C5.5.68 (TLBI VALE1, TLBI VALE1NXS).
- C5.5.77 (TLBI VMALLE1, TLBI VMALLE1NXS).
- G8.2.136 (TLBIALL, TLB Invalidate All).
- G8.2.142 (TLBIASID, TLB Invalidate by ASID match).
- G8.2.148 (TLBIMVA, TLB Invalidate by VA).
- G8.2.149 (TLBIMVAA, TLB Invalidate by VA, All ASID).
- G8.2.151 (TLBIMVAAL, TLB Invalidate by VA, All ASID, Last level).
- G8.2.156 (TLBIMVAL, TLB Invalidate by VA, Last level).

## 2.51 C19793

In section C5.5.25 (TLBI RIPAS2LE1IS, TLBI RIPAS2LE1ISNXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, Last level, EL1, Inner Shareable), in the subsection 'Purpose', the text that reads:

- The entry is a stage 2 only translation table entry, from the final level of the translation table walk.

is updated to read:

- The entry is a stage 2 only translation table entry, from the leaf level of the translation table walk, indicated by the TTL hint.

Equivalent changes are made in the following sections:

- C5.5.24 (TLBI RIPAS2LE1, TLBI RIPAS2LE1NXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, Last level, EL1).
- C5.5.26 (TLBI RIPAS2LE1OS, TLBI RIPAS2LE1OSNXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, Last level, EL1, Outer Shareable).
- C5.5.32 (TLBI RVAALE1, TLBI RVAALE1NXS, TLB Range Invalidate by VA, All ASID, Last level, EL1).
- C5.5.33 (TLBI RVAALE1IS, TLBI RVAALE1ISNXS, TLB Range Invalidate by VA, All ASID, Last Level, EL1, Inner Shareable).

- C5.5.34 (TLBI RVAALE1OS, TLBI RVAALE1OSNXS, TLB Range Invalidate by VA, All ASID, Last Level, EL1, Outer Shareable).

In section C5.5.35 (TLBI RVAE1, TLBI RVAE1NXS, TLB Range Invalidate by VA, EL1), in the subsection 'Purpose', the text that reads:

- The entry is a stage 1 translation table entry.

is updated to read:

- The entry is a stage 1 translation table entry, from any level of the translation table walk up to the level indicated in the TTL hint.

Equivalent changes are made in the following sections:

- C5.5.36 (TLBI RVAE1IS, TLBI RVAE1ISNXS, TLB Range Invalidate by VA, EL1, Inner Shareable).
- C5.5.37 (TLBI RVAE1OS, TLBI RVAE1OSNXS, TLB Range Invalidate by VA, EL1, Outer Shareable).
- C5.5.38 (TLBI RVAE2, TLBI RVAE2NXS, TLB Range Invalidate by VA, EL2).
- C5.5.39 (TLBI RVAE2IS, TLBI RVAE2ISNXS, TLB Range Invalidate by VA, EL2, Inner Shareable).
- C5.5.40 (TLBI RVAE2OS, TLBI RVAE2OSNXS, TLB Range Invalidate by VA, EL2, Outer Shareable).
- C5.5.44 (TLBI RVALE1, TLBI RVALE1NXS, TLB Range Invalidate by VA, Last level, EL1).
- C5.5.45 (TLBI RVALE1IS, TLBI RVALE1ISNXS, TLB Range Invalidate by VA, Last level, EL1, Inner Shareable).
- C5.5.46 (TLBI RVALE1OS, TLBI RVALE1OSNXS, TLB Range Invalidate by VA, Last level, EL1, Outer Shareable).
- C5.5.47 (TLBI RVALE2, TLBI RVALE2NXS, TLB Range Invalidate by VA, Last level, EL2).
- C5.5.48 (TLBI RVALE2IS, TLBI RVALE2ISNXS, TLB Range Invalidate by VA, Last level, EL2, Inner Shareable).
- C5.5.49 (TLBI RVALE2OS, TLBI RVALE2OSNXS, TLB Range Invalidate by VA, Last level, EL2, Outer Shareable).

In section C5.5.21 (TLBI RIPAS2E1, TLBI RIPAS2E1NXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, EL1), in the subsection 'Purpose', the text that reads:

- The entry is a stage 2 only translation table entry, from any level of the translation table walk.

is updated to read:

- The entry is a stage 2 only translation table entry, from any level of the translation table walk up to the level indicated in the TTL hint.

Equivalent changes are made in the following sections:

- C5.5.22 (TLBI RIPAS2E1IS, TLBI RIPAS2E1ISNXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, EL1, Inner Shareable).

- C5.5.23 (TLBI RIPAS2E1OS, TLBI RIPAS2E1OSNXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, EL1, Outer Shareable).
- C5.5.29 (TLBI RVAAE1, TLBI RVAAE1NXS, TLB Range Invalidate by VA, All ASID, EL1).
- C5.5.30 (TLBI RVAAE1IS, TLBI RVAAE1ISNXS, TLB Range Invalidate by VA, All ASID, EL1, Inner Shareable).
- C5.5.31 (TLBI RVAAE1OS, TLBI RVAAE1OSNXS, TLB Range Invalidate by VA, All ASID, EL1, Outer Shareable).
- C5.5.41 (TLBI RVAE3, TLBI RVAE3NXS, TLB Range Invalidate by VA, EL3).
- C5.5.42 (TLBI RVAE3IS, TLBI RVAE3ISNXS, TLB Range Invalidate by VA, EL3, Inner Shareable).
- C5.5.43 (TLBI RVAE3OS, TLBI RVAE3OSNXS, TLB Range Invalidate by VA, EL3, Outer Shareable).
- C5.5.50 (TLBI RVALE3, TLBI RVALE3NXS, TLB Range Invalidate by VA, Last level, EL3).
- C5.5.51 (TLBI RVALE3IS, TLBI RVALE3ISNXS, TLB Range Invalidate by VA, Last level, EL3, Inner Shareable).
- C5.5.52 (TLBI RVALE3OS, TLBI RVALE3OSNXS, TLB Range Invalidate by VA, Last level, EL3, Outer Shareable).

Also in section C5.5.25 (TLBI RIPAS2LE1IS, TLBI RIPAS2LE1ISNXS, TLB Range Invalidate by Intermediate Physical Address, Stage 2, Last level, EL1, Inner Shareable), in the field 'TTL, bits [38:37]', the text that reads:

TTL Level hint. The TTL hint is only guaranteed to invalidate entries in the range that match the level described by the TTL hint.

0b00 The entries in the range can be using any level for the translation table entries. 0b01 All entries to invalidate are Level 1 translation table entries. If FEAT\_LPA2 is not implemented, when using a 16KB translation granule, this value is reserved and hardware should treat this field as 0b00. 0b10 All entries to invalidate are Level 2 translation table entries. 0b11 All entries to invalidate are Level 3 translation table entries.

is updated to read:

TTL Level hint. The TTL hint is only guaranteed to invalidate:

- Non-leaf-level entries in the range up to but not including the level described by the TTL hint.
- Leaf-level entries in the range that match the level described by the TTL hint.

0b00 The entries in the range can be using any level for the translation table entries. 0b01 The TTL hint indicates level 1. If FEAT\_LPA2 is not implemented, when using a 16KB translation granule, this value is reserved and hardware should treat this field as 0b00. 0b10 The TTL hint indicates level 2. 0b11 The TTL hint indicates level 3.

Equivalent changes are made in all of the sections listed above.

## 2.52 D19800

In section J1.1.3 (aarch64/function), the function `IsHCRXEL2Enabled()`, that reads as:

```
boolean IsHCRXEL2Enabled()  
    assert(HaveFeatHCX());  
    ...
```

Is updated to read:

```
boolean IsHCRXEL2Enabled()  
    if !HaveFeatHCX() then return FALSE;  
    ...
```

## 2.53 D19804

In section D9.4.1 (Virtual address translation), the following text is added:

If a tag write by an STG instruction that does not also write data is translated by a writeable-clean descriptor, but the tag write effect is IGNORED due to a stage 1 descriptor not having the Tagged memory attribute, or because Allocation tag access is disabled for the instruction by SCR\_EL3.ATA, HCR\_EL2.ATA, SCTLR\_ELx.ATA or SCTLR\_ELx.ATA0, it is **CONSTRAINED UNPREDICTABLE** whether hardware updates the dirty state of that descriptor.

## 2.54 R19810

In section B2.3.3 (Ordering relations), the definition of 'Tag-ordered-before' is updated to read:

If FEAT\_MTE2 is implemented, a Memory Tag-Check-read R1 is Tag-ordered-before a Checked Memory Write effect W2 generated by the same instruction if and only if all of the following apply:

- There is an Intrinsic data dependency from R1 to a Conditional-Branching effect B3 generated by the same instruction as R1.
- There is an Intrinsic control dependency from the Conditional-Branching effect B3 to W2.

## 2.55 D19817

In section G8.3.33 (PMMIR, Performance Monitors Machine Identification Register) in the BUS\_SLOTS, bits [15:8] field, the text that reads:

Bus count. The largest value by which the BUS\_ACCESS event might increment in a single BUS\_CYCLES cycle. When this field is nonzero, the largest value by which the BUS\_ACCESS

event might increment in a single BUS\_CYCLES cycle is BUS\_SLOTS. This field has an **IMPLEMENTATION DEFINED** value. Access to this field is RO.

is corrected to read:

Bus count. The largest value by which the BUS\_ACCESS event might increment in a single BUS\_CYCLES cycle. When this field is nonzero, the largest value by which the BUS\_ACCESS event might increment in a single BUS\_CYCLES cycle is BUS\_SLOTS. If the information is not available, this field will read as zero. This field has an **IMPLEMENTATION DEFINED** value. Access to this field is RO.

The equivalent changes are made in section D17.5.12 (PMMIR\_EL1, Performance Monitors Machine Identification Register) and I5.3.30 (PMMIR, Performance Monitors Machine Identification Register).

## 2.56 D19829

In section D17.2.63 (ID\_AA64ISAR2\_EL1, AArch64 Instruction Set Attribute Register 2), in the 'RPRES, bits [7:4]' field, the following text is removed:

From Armv8.7, if Advanced SIMD and floating-point is implemented, the only permitted value is 0b0001.

## 2.57 E19831

In section K7.2 (Gray-count scheme for timer distribution scheme), the following pseudocode for Gray code encoding and decoding:

```
Gray[N] = Count[N]
Gray[i] = (XOR(Gray[N:i+1])) XOR Count[i] for N-1 >= i >= 0
Count[i] = XOR(Gray[N:i]) for N >= i >= 0
```

is updated to read:

```
Gray = Count EOR ('0':Count<N:1>)
Count<N> = Gray<N>
for i = N-1 downto 0
    Count<i> = Gray<i> EOR Count<i+1>
```

## 2.58 D19833

In section K7.2 (Gray-count scheme for timer distribution scheme) the following Note is removed:

This scheme has the advantage of being relatively simple to switch, in either direction, between operating with low-frequency and low-precision, and operating with high-frequency and high-precision. To achieve this, the ratio of the frequencies must be  $2^n$ , where  $n$  is an integer. A switch-over can occur only on the  $2^{n+1}$  boundary to avoid losing the Gray-coding property on a switch-over.

## 2.59 C19835

In section B2.3.12 (Limited ordering regions), after the following text:

A memory location lies within the LORegion identified by the LORegion Number if the PA lies between the Start Address and the End Address, inclusive. The Start Address must be defined to be aligned to 64KB and the End Address must be defined as the top byte of a 64KB block of memory.

the following statement is added:

It is permitted for multiple LORegion descriptors with non-overlapping address ranges to be configured with the same LORegion Number.

## 2.60 D19887

In section J1.1.3 (aarch64/functions), the write accessor Mem[] (assignment form) reading:

```
Mem[bits(64) address, integer size, AccType acctype, boolean ispair] = bits(size*8)
value_in
...
if !atomic && ispair && address == Align(address, halfsize) then
    single_is_aligned = TRUE;
    <highhalf, lowhalf> = value;
    AArch64.MemSingle[address, halfsize, acctype,
        single_is_aligned, ispair] = lowhalf;
    AArch64.MemSingle[address + halfsize, halfsize, acctype,
        single_is_aligned, ispair] = highhalf;
elseif atomic && ispair then
    AArch64.MemSingle[address, size, acctype, aligned, ispair] = value;
...
```

Is updated to read:

```
Mem[bits(64) address, integer size, AccType acctype, boolean ispair] = bits(size*8)
value_in
...
if !atomic && ispair && address == Align(address, halfsize) then
    single_is_pair = FALSE;
    single_is_aligned = TRUE;
```

```

    <highhalf, lowhalf> = value;
    AArch64.MemSingle[address, halfsize, acctype,
        single_is_aligned, single_is_pair] = lowhalf;
    AArch64.MemSingle[address + halfsize, halfsize, acctype,
        single_is_aligned, single_is_pair] = highhalf;
    elsif atomic && ispair then
        AArch64.MemSingle[address, size, acctype, aligned, ispair] = value;
    ...

```

## 2.61 E19892

In section J1.1.5 (aarch64/translation), the function S1HasPermissionsFault() that reads:

```

boolean AArch64.S1HasPermissionsFault(
    Regime regime,
    SecurityState ss,
    TTWState walkstate,
    S1TTWParams walkparams,
    boolean ispriv,
    AccType acctype,
    boolean iswrite
)

```

Is replaced by S1CheckPermissions():

```

FaultRecord AArch64.S1CheckPermissions(
    Regime regime,
    SecurityState ss,
    TTWState walkstate,
    S1TTWParams walkparams,
    boolean ispriv,
    AccType acctype,
    boolean iswrite,
    FaultRecord fault_in
)

```

In section J1.1.5 (aarch64/translation), the function S2HasPermissionsFault() that reads:

```

boolean AArch64.S2HasPermissionsFault(
    boolean s2fslwalk,
    TTWState walkstate,
    SecurityState ss,
    S2TTWParams walkparams,
    boolean ispriv,
    AccType acctype,
    boolean iswrite
)

```

Is replaced by S2CheckPermissions():

```

FaultRecord AArch64.S2CheckPermissions(
    boolean s2fslwalk,
    TTWState walkstate,
    SecurityState ss,
    S2TTWParams walkparams,
    boolean ispriv,
    AccType acctype,
)

```



```
        boolean iswrite,
        FaultRecord fault
    )
```

Appropriate changes are made in the pseudocode where these functions are called.

In section D8.15 (Pseudocode description of VMSSAv8-64 address translation), the subsection 'Fault detection' is updated to take these changes into account.

## 2.62 D19917

In section D17.2.36 (DCZID\_ELO, Data Cache Zero ID register), in the definition of 'BS, bits [3:0]', the following text is added:

■ If FEAT\_MTE2 is implemented, the minimum size supported is 16 bytes (value == 2).

## 2.63 D19918

In section J1.1.3 (aarch64/functions), in the AArch64.CheckAlignment() function, the code that reads:

```
    if SCTLRL[.A == '1' then check = TRUE;
    elseif HaveLSE2Ext() then
        check = (UInt(address<3:0>) + alignment > 16) && ((ordered && SCTLRL[.nAA ==
'0') || atomic);
    else check = atomic || ordered;
```

Is updated to read:

```
    if SCTLRL[.A == '1' then check = TRUE;
    elseif HaveLSE2Ext() then
        // For ordered pair operation check whether entire access is within 16-byte
        integer accsize = if ispair then alignment * 2 else alignment;
        check = (UInt(address<3:0>) + accsize > 16) && ((ordered && SCTLRL[.nAA ==
'0') || atomic);
    else check = atomic || ordered;
```

In section J1.1.3 (aarch64/functions), the Mem[] non-assignment (read) accessor function, the code that reads:

```
bits(size*8) Mem[...]
...
if ispair then
    // check alignment on size of element accessed, not overall access size
    aligned = AArch64.CheckAlignment(address, halfsize, acctype, iswrite);
else
    aligned = AArch64.CheckAlignment(address, size, acctype, iswrite);
```

Is updated to read:

```
bits(size*8) Mem[...]
...
integer align_size = if ispair then halfsize else size;
aligned = AArch64.CheckAlignment(address, align_size, acctype, iswrite, ispair);
```

Equivalent changes are made to the Mem[] assignment (write) accessor function.

## 2.64 D19928

In sections D17.2.118 (SCTLR\_EL1, System Control Register (EL1)) and D17.2.119 (SCTLR\_EL2, System Control Register (EL2)), in the 'EPAN, bit [57]' field, the text that reads:

Any speculative data accesses that would generate a Permission fault if the accesses were not speculative will not cause an allocation into a cache.

is corrected to read:

Any speculative data accesses that would generate a Permission fault as a result of PSTATE.PAN=1 if the accesses were not speculative will not cause an allocation into a cache.

## 2.65 D19936

In section D17.5.9 (PMEVTYPER<n>\_ELO, Performance Monitors Event Type Registers, n = 0 - 30), the description of 'T, bit [23]' that reads:

When FEAT\_TME is implemented:

Transactional state filtering bit. Controls counting in Transactional state.

0b0 This bit has no effect on filtering of events. 0b1 Do not count events in Transactional state.

is updated to read:

When FEAT\_TME is implemented:

Transactional state filtering bit. Controls counting of Attributable events in Non-transactional state.

0b0 This bit has no effect on filtering of events. 0b1 Do not count Attributable events in Non-transactional state.

For each Unattributable event, it is **IMPLEMENTATION DEFINED** whether the filtering applies.

Equivalent changes are made in the following sections:

- D17.5.1 (PMCCFILTR\_ELO, Performance Monitors Cycle Count Filter Register).
- I5.3.24 (PMEVTYPER<n>\_ELO, Performance Monitors Event Type Registers, n = 0 - 30).

The updated definition of 'T, bit [23]' is added to section I5.3.2 (PMCCFILTR\_ELO, Performance Monitors Cycle Counter Filter Register).

## 2.66 C19956

In section D11.11.3 (Common event numbers), in the description of PMU event '0x0012, BR\_PRED', the following text is added:

If no program-flow prediction resources are implemented, this event is optional, but Arm recommends that BR\_PRED counts all branches.

It is **IMPLEMENTATION DEFINED** when the branch is counted. Arm recommends that it is counted when the branch is resolved, that is, at the same point in the instruction pipeline as when the BR\_MIS\_PRED event would be counted if the branch resolves as mispredicted. This means that (BR\_PRED - BR\_MIS\_PRED) is the number of correctly predicted branches and the ratio (BR\_MIS\_PRED ÷ BR\_PRED) can be calculated in a meaningful way.

PMCEID0\_ELO[18] reads as 0b1 if this event is implemented and 0b0 otherwise.

## 2.67 D19961

In section C7.2.227 (SABDL, SABDL2), the text that reads:

This instruction subtracts the vector elements of the second source SIMD&FP register from the corresponding vector elements of the first source SIMD&FP register, places the absolute value of the results into a vector, and writes the vector to the lower or upper half of the destination SIMD&FP register.

is corrected to read:

This instruction subtracts the vector elements in the lower or upper half of the second source SIMD&FP register from the corresponding vector elements of the first source SIMD&FP register, places the absolute value of the results into a vector, and writes the vector to the destination SIMD&FP register.

## 2.68 C20009

In section D17.2.40 (FAR\_EL1, Fault Address Register (EL1)), the Note that reads:

The address held in this field is an address accessed by the instruction fetch or data access that caused the exception that actually gave rise to the instruction or data abort. It is the lower

address that gave rise to the fault. Where different faults from different addresses arise from the same instruction, such as for an instruction that loads or stores an unaligned address that crosses a page boundary, the architecture does not prioritize between those different faults.

is updated to read:

The address held in this field is an address accessed by the instruction fetch or data access that caused the exception that actually gave rise to the instruction or data abort. It is the lower address that gave rise to the fault that is reported. Where different faults from different addresses arise from the same instruction, such as for an instruction that loads or stores an unaligned address that crosses a page boundary, the architecture does not prioritize which fault is reported.

Equivalent changes are made in the following sections:

- D17.2.41 (FAR\_EL2, Fault Address Register (EL2)).
- D17.2.42 (FAR\_EL3, Fault Address Register (EL3)).
- D17.2.55 (HPFAR\_EL2, Hypervisor IPA Fault Address Register).

## 2.69 D20011

In section D11.11.3 (Common event numbers), subsection ‘Common microarchitectural events’, in the ‘0x0074, ASE\_SPEC, Operation speculatively executed, Advanced SIMD’ definition, the bullet points that read:

- Cryptographic operations other than PMULL, in AArch64 state.
- VMULL, in AArch32 state.

are changed to read:

- Cryptographic operations, other than PMULL, PMULL2 (1Q variant) in AArch64 state and VMULL (P64 variant) in AArch32 state.

In the same event definition, the text that reads:

In AArch64 state, PMULL, and in AArch32 state, VMULL are counted as Advanced SIMD operations.

is changed to read:

Advanced SIMD PMULL, PMULL2 (1Q variant) in AArch64 state and VMULL (P64 variant) in AArch32 state are counted as Advanced SIMD operations.

In the same section, in the ‘0x0077, CRYPTO\_SPEC, Operation speculatively executed, Cryptographic instruction’ definition, the text that reads:

The counter counts each operation counted by INST\_SPEC that is a cryptographic operation other than PMULL or VMULL.

See The Cryptographic Extension on page C3-333.

is changed to read:

The counter counts each operation counted by INST\_SPEC that is a cryptographic operation, other than Advanced SIMD PMULL, PMULL2 (1Q variant) and SVE2 PMULLB, PMULLT (Q variant) in AArch64 state, and Advanced SIMD VMULL (P64 variant) in AArch32 state.

See The Armv8 Cryptographic Extension on page A2-80 and SVE2 Crypto Extensions on page C4-485.

## 2.70 D20053

In section F2.11 (Advanced SIMD and floating-point load/store instructions), in Table F2-17 'SIMD and floating-point register file load/store instructions', the 'Operation' description for Vector Load Multiple that reads:

Load 1-16 consecutive 32-bit registers, floating-point only.

is corrected to read:

Load 1-32 consecutive 32-bit registers, floating-point only.

In the same table, the 'Operation' description for Vector Store Multiple that reads:

Store 1-16 consecutive 32-bit registers, floating-point only.

is corrected to read:

Store 1-32 consecutive 32-bit registers, floating-point only.

## 2.71 D20128

In section D13.6.3 (Additional information for each profiled memory access operation), the bullet list that reads:

The sampled data physical address packet is not output if any of the following are true:

- The PE does not translate the address, for example because it does not perform the access or the address translation generates a Translation fault.
- The sampled data virtual address packet is not output.
- Sampling of physical addresses is prohibited by System register controls.

is changed to read:

The sampled data physical address packet is not output if any of the following are true:

- The sampled operation operates on a virtual address and any of the following are true:
  - The PE does not translate the address, for example because it does not perform the access or the address translation generates a Translation fault.
  - The sampled data virtual address packet is not output.
- Sampling of physical addresses is prohibited by System register controls.

If `AArch64.ExclusiveMonitorPass()` or `AArch32.ExclusiveMonitorPass()` returns `FALSE` for a Store Exclusive instruction, it is **IMPLEMENTATION DEFINED** whether or not the physical address packet is output when permitted by the above rules.

## 2.72 R20165

In section D11.7.2 (Accuracy of event filtering), subsection ‘Software increment events’, the text that reads:

Software increment events must also be counted without the need for explicit synchronization. For example, two software increments executed without an intervening Context synchronization event must increment the event counter twice.

is updated to read:

If the PE performs two architecturally executed writes to the `PMSWINC_ELO` or `PMSWINC` register without an intervening Context synchronization event, then the counter is incremented twice.

## 2.73 D20171

In section C6.2.43 (CASH, CASAH, CASALH, CASLH), the bullet that reads:

- CAS has neither acquire nor release semantics.

is corrected to read:

- CASH has neither acquire nor release semantics.

In section C6.2.44 (CASP, CASPA, CASPAL, CASPL), the bullet that reads:

- CAS has neither acquire nor release semantics.

is corrected to read:

- CASP has neither acquire nor release semantics.

## 2.74 D20210

In section J1.1.3 (aarch64/functions), the function `AArch64.PhysicalErrorSyndrome()` that reads as:

```
bits(25) AArch64.PhysicalErrorSyndrome(boolean implicit_esb)
    bits(25) syndrome = Zeros(25);
    ...
    if errorstate == ErrorState_Uncategorized then
    ...
    elsif errorstate == ErrorState_IMPDEF then
    ...
    else
        syndrome<24> = '0';
        syndrome<13> = (if implicit_esb then '1' else '0');
        syndrome<12:10> = AArch64.EncodeAsyncErrorSyndrome(errorstate);
        syndrome<5:0> = '010001';
```

is changed to:

```
bits(25) AArch64.PhysicalErrorSyndrome(boolean implicit_esb)
    ...
    if errorstate == ErrorState_Uncategorized then
    ...
    elsif errorstate == ErrorState_IMPDEF then
    ...
    else
        syndrome<24> = '0';
        syndrome<13> = (if implicit_esb then '1' else '0');
        syndrome<12:10> = AArch64.EncodeAsyncErrorSyndrome(errorstate);
        syndrome<9> = fault.extflag;
        syndrome<5:0> = '010001';
```

Similarly in section J1.2.3 (aarch32/functions), the function `AArch32.PhysicalErrorSyndrome()` that reads as:

```
bits(16) AArch32.PhysicalErrorSyndrome()
    bits(32) syndrome = Zeros(32);
    FaultRecord fault = GetPendingPhysicalError();
    boolean long_format = TTBCR.EAE == '1';
    syndrome = AArch32.CommonFaultStatus(fault, long_format);
    return syndrome<15:0>;
```

is updated to:

```
bits(16) AArch32.PhysicalErrorSyndrome()
    bits(32) syndrome = Zeros(32);
    FaultRecord fault = GetPendingPhysicalError();
    if PSTATE.EL == EL2 then
        ErrorState errstate = AArch32.PEErrorState(fault);
        syndrome<11:10> = AArch32.EncodeAsyncErrorSyndrome(errstate);
        syndrome<9> = fault.extflag;
        syndrome<5:0> = '010001';
    else
        boolean long_format = TTBCR.EAE == '1';
        syndrome = AArch32.CommonFaultStatus(fault, long_format);
    return syndrome<15:0>;
```

## 2.75 C1186: SME

In section D17.2.70 (ID\_AA64ZFR0\_EL1, SVE Feature ID register 0), the following text:

Irrespective of the value of this field, when the PE is in Streaming SVE mode and it is not known whether FEAT\_SME\_FA64 is implemented and enabled at the current Exception level, software should not attempt to execute the instructions described by non-zero values of this field.

is added to the descriptions of the following fields:

- F64MM, bits [59:56].
- F32MM, bits [55:52].
- SM4, bits [43:40].
- SHA3, bits [35:32].
- BitPerm, bits [19:16].
- AES, bits [7:4].

The following text is added to the description of I8MM, bits [47:44]:

Irrespective of the value of this field, when the PE is in Streaming SVE mode and it is not known whether FEAT\_SME\_FA64 is implemented and enabled at the current Exception level, software should not attempt to execute the the SVE instructions SMMLA, UMMLA, and USMMLA.

The following text is added to the description of BF16, bits [23:20]:

Irrespective of the value of this field, when the PE is in Streaming SVE mode and it is not known whether FEAT\_SME\_FA64 is implemented and enabled at the current Exception level, software should not attempt to execute the SVE instruction BFMMMLA.

## 2.76 C1342: SME

In section D17.1.3 (Principles of the ID scheme for fields in ID registers), the following subsection is added:

Alternative ID scheme used for ID\_AA64SMFRO\_EL1

Apart from the ID\_AA64SMFRO\_EL1.SMEver field, which is a 4-bit unsigned integer conforming to the standard scheme, software must treat the other fields in this register as follows:

- A field value where the bit or all bits are 0 indicates that the SME feature or instructions described by this field are not implemented.
- A field value where the bit or all bits are 1 indicates that the SME feature or instructions described by this field are implemented.
- A multi-bit field value containing both 0 and 1 bits is a RESERVED encoding.



## 2.77 D1386: SME

In the following sections:

- C7.2.15 BFDOT (by element).
- C7.2.16 BFDOT (vector).
- C7.2.19 BFMMLA.
- C8.2.35 BFDOT (indexed).
- C8.2.36 BFDOT (vectors).
- C7.2.19 BFMMLA.

The text that reads:

If FEAT\_EBF16 is implemented and FPCR.EBF is 1, then this instruction:

- Performs a fused sum-of-products of each pair of adjacent BFloat16 elements in the first source vector with the specified pair of elements in the second source vector. The intermediate single-precision products are not rounded before they are summed, but the intermediate sum is rounded before accumulation into the single-precision destination element that overlaps with the corresponding pair of BFloat16 elements in the first source vector.
- Generates only the default NaN, as if FPCR.DN is 1.
- Follows all other floating-point behaviors that apply to single-precision arithmetic, as controlled by the effective value of the FPCR in the current execution mode, and captured in the FPSR.

is corrected to read:

If FEAT\_EBF16 is implemented and FPCR.EBF is 1, then this instruction:

- Performs a fused sum-of-products of each pair of adjacent BFloat16 elements in the first source vector with the specified pair of elements in the second source vector. The intermediate single-precision products are not rounded before they are summed, but the intermediate sum is rounded before accumulation into the single-precision destination element that overlaps with the corresponding pair of BFloat16 elements in the first source vector.
- Generates only the default NaN, as if FPCR.DN is 1.
- Does not modify the cumulative FPSR exception bits (IDC, IXC, UFC, OFC, DZC, and IOC).
- Disables trapped floating-point exceptions, as if the FPCR trap enable bits (IDE, IXE, UFE, OFE, DZE, and IOE) are all zero.
- Follows all other floating-point behaviors that apply to single-precision arithmetic, as governed by FPCR.RMode, FPCR.FZ, FPCR.AH, and FPCR.FIZ controls in the current execution mode.

## 2.78 C215: SVE

In section A1.4 (Supported data types), the text that reads:

- An SVE scalable vector register has an **IMPLEMENTATION DEFINED** width that is a multiple of 128 bits, up to a maximum of 2048 bits.

is changed to read:

- An SVE scalable vector register has an **IMPLEMENTATION DEFINED** width that is a power of two, from a minimum of 128 bits up to a maximum of 2048 bits.

Within the same section, the text that reads:

- An SVE predicate vector register has an **IMPLEMENTATION DEFINED** width that is a multiple of 16 bits, up to a maximum of 256 bits.

is changed to read:

- An SVE predicate vector register has an **IMPLEMENTATION DEFINED** width that is a power of two, from a minimum of 16 bits up to a maximum of 256 bits.

In section A1.4.2 (SVE vector format), in the subsection 'SVE configurable vector length', the following rule  $R_{RYQYY}$  is deleted:

An implementation is permitted to allow the vector length to be constrained to multiples of 128 that are not a power of two. It is **IMPLEMENTATION DEFINED** which of the permitted multiples of 128 are supported.

In section B1.2.2 (SVE vector registers), the rule  $R_{KCWQB}$  that reads:

The size of an SVE scalable vector register is an **IMPLEMENTATION DEFINED** multiple of 128 bits.

is changed to read:

The size of an SVE scalable vector register is an **IMPLEMENTATION DEFINED** power of two.

In section B1.2.3 (SVE predicate registers), the rule  $R_{NKRJV}$  that reads:

The size of an SVE predicate register is an **IMPLEMENTATION DEFINED** multiple of 16 bits.

is changed to read:

The size of an SVE predicate register is an **IMPLEMENTATION DEFINED** power of two.

In section D17.2.159 (ZCR\_EL1, SVE Control Register (EL1)), in the LEN, bits [3:0] field, the text that reads:

The Non-streaming SVE vector length can be any multiple of 128 bits, from 128 bits to 2048 bits inclusive.

is changed to read:

The Non-streaming SVE vector length can be any power of two from 128 bits to 2048 bits inclusive.

The same change is made in the following sections:

- D17.2.160 (ZCR\_EL2, SVE Control Register (EL2)).
- D17.2.161 (ZCR\_EL3, SVE Control Register (EL3)).

In section J1.1.3 (aarch64/functions), the code within the function `ImplementedSVEVectorLength()` that reads:

```
// Reduce SVE vector length to a supported value (e.g. power of two)
integer ImplementedSVEVectorLength(integer nbits_in)
    integer nbits = Min(nbits_in, MaxImplementedVL());

    assert 128 <= nbits && nbits <= 2048 && Align(nbits, 128) == nbits;
    while nbits > 128 do
        if IsPow2(nbits) || SupportedNonPowerTwoVL(nbits) then return nbits;
        nbits = nbits - 128;
    return nbits;
```

is changed to read:

```
// Reduce SVE vector length to a supported value (power of two)
integer ImplementedSVEVectorLength(integer nbits_in)
    integer maxbits = MaxImplementedVL();

    assert 128 <= maxbits && maxbits <= 2048 && IsPow2(maxbits);

    integer nbits = Min(nbits_in, maxbits);
    assert 128 <= nbits && nbits <= 2048 && Align(nbits, 128) == nbits;
    while nbits > 128 do
        if IsPow2(nbits) then return nbits;
        nbits = nbits - 128;
    return nbits;
```

Within the same section, the function `SupportedNonPowerTwoVL()` is removed.

In the Glossary, the definition of ‘Predicate register’ that reads:

An SVE predicate register, P0-P15, having a length that is a multiple of 16 bits, in the range 16 to 256, inclusive.

is changed to read:

An SVE predicate register, P0-P15, having a length that is a power of two, in the range 16 to 256, inclusive.

Also in the Glossary, the definition of ‘Scalable vector register’ that reads:

An SVE vector register, Z0-Z31, having a length that is a multiple of 128 bits, in the range 128 bits to 2048 bits, inclusive.

is changed to read:

An SVE vector register, Z0-Z31, having a length that is a power of two, in the range 128 bits to 2048 bits, inclusive.

## 2.79 C225: SVE

In section D7.2.1 (Virtual address space overflow), the following text is added:

The **UNKNOWN** virtual address behavior also applies to the set of bytes addressed by SVE and SME predicated, contiguous loads and stores that cross the `0xFFFF_FFFF_FFFF_FFFF` boundary, even if all of the virtual addresses below the boundary correspond to inactive elements. Conversely, for SVE gather loads and scatter stores, the **UNKNOWN** address behavior applies only to accesses corresponding to an individual Active element that crosses the boundary.

## 2.80 C256: SVE

In section H2.4.2 (Executing instructions in Debug state), in the subsection 'A64 instructions that are unchanged in Debug state', the list that reads:

SVE instructions

When FEAT\_SVE is implemented, these instructions are:

- CPY.
- DUP (scalar).
- EXT.
- INSR (scalar).
- PTRUE with ALL constraint and byte element size.
- RDNFR (unpredicated).
- RDVL.
- WRNFR.

is changed to read:

SVE instructions

When FEAT\_SVE is implemented, these instructions are:

- CPY.
- DUP (scalar).
- EXT, destructive variant.
- INSR (scalar).

- PTRUE with ALL constraint and byte element size.
- RDIFFR (unpredicated).
- RDVL.
- WRFFR.

## 2.81 C279: SVE

In section B1.2.4 (FFR, First Fault Register), rule  $R_{WZJVT}$  that reads:

Bits in the FFR are indirectly set to 0 as a result of a suppressed access or fault generated in response to an Active element of an SVE First-fault or Non-fault vector load.

is clarified to read:

Bits in the FFR are indirectly set to 0 as a result of a suppressed access or suppressed fault corresponding to an Active element of an SVE First-fault or Non-fault vector load.

## 2.82 D302: SVE

In section C1.2.6 (Register names), in the subsection ‘SIMD vector register list’, the text that reads:

Where an instruction operates on multiple SIMD and floating-point registers, for example vector load/store structure and table lookup operations, the registers are specified as a list enclosed by curly braces. This list consists of either a sequence of registers separated by commas, or a register range separated by a hyphen. The registers must be numbered in increasing order, modulo 32, in increments of one. The hyphenated form is preferred for disassembly if there are more than two registers in the list and the register number are increasing.

is updated to read:

Where an instruction operates on multiple SIMD&FP or SVE vector registers, for example vector load/store structure and table lookup operations, the registers are specified as a list enclosed by curly braces. This list consists of either a sequence of registers separated by commas, or a register range separated by a hyphen. The registers must be numbered in increasing order, modulo 32, in increments of one. The hyphenated form is preferred for disassembly if there are more than two registers in the list and the register numbers are increasing.

Similar updates are made throughout section C1.2 (Structure of the A64 assembler language) to account for the SVE assembler syntax.

## 2.83 D1461: Armv9 Debug

In section D4.6.12 (External Outputs), the statement  $I_{BZ HDF}$  that reads:

The ETE architecture supports between one and four External Outputs. The number of outputs that a trace unit has is **IMPLEMENTATION DEFINED**, but at least one output is always implemented.

is updated to read:

The ETE architecture supports between zero and four External Outputs. The number of outputs that a trace unit has is **IMPLEMENTATION DEFINED**, and Arm recommends that at least one output is implemented.

## 2.84 D1466: Armv9 Debug

In section D11.11.3 (Common event numbers), in the subsection 'Common microarchitectural events', the description for each CTI\_TRIGOUT<n> event, where <n> is in the range 4 to 7, that reads:

This event must be implemented if FEAT\_ETE is implemented.

is updated to read:

This event must be implemented if FEAT\_ETE is implemented and TRCIDR5.NUMEXTINSEL > (n - 4).

## 2.85 D1493: Armv9 Debug

In section D4.5.3 (Trace unit behavior while the PE is in Debug state), rule  $R_{DPKSC}$  that reads:

While the PE is in Debug state, the trace unit does not trace instructions that are executed.

is updated to read:

While the PE is in Debug state, the trace unit:

- Does not trace instructions that are executed.
- Does not trace the effects of instructions that are executed.
- Does not trace Exceptional occurrences.

Additionally, in section D4.5.8 (Filtering trace generation), in the subsection 'Rules for tracing Exceptional occurrences', rule  $R_{DPM BQ}$  that reads:

When an Exceptional occurrence occurs and TRCRSR.TA is 0b1, the Exceptional occurrence is traced.

is updated to read:

When an Exceptional occurrence occurs and the PE is not in Debug state and TRCRSR.TA is 0b1, the Exceptional occurrence is traced.