arm

Accessing memory-mapped peripherals with Arm DS

Version 1.0

Accessing memory-mapped peripherals with Arm DS

Release information

Document history

Issue	Date	Confidentiality	Change
0100-01	9 December 2022	Non-Confidential	Initial release

Proprietary Notice

Confidentiality Status

Product Status

The information in this document is for a product under development.

Feedback

Arm® welcomes feedback on this product and its documentation. To provide feedback on the product, create a ticket on https://support.developer.arm.com

To provide feedback on the document, fill the following survey: https://developer.arm.com/documentation-feedback-survey.

Page **3** of **19**

Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used language that can be offensive. Arm strives to lead the industry and create change.

We believe that this document contains no offensive language. To report offensive language in this document, email terms@arm.com.

Contents

1. Overview

In this tutorial, learn about mapping a C variable to each register of a memory-mapped peripheral, then using a pointer to that variable to read and write the register.

In most Arm embedded systems, peripherals are at specific addresses in memory. It is often convenient to map a C variable onto each register of a memory-mapped peripheral. Then, use a pointer to that variable to read and write the register. In your code, you must consider not only the size and address of the register, but also its alignment in memory.

Before you begin

*   Install and license Arm Development Studio (Arm DS). For more information, see the Arm Development Studio documentation.

Note You can test the examples in this tutorial by using the Cortex-A53x1 FVP, Base_A53x1, with Arm Compiler 6. Both Arm Compiler 6 and the Cortex-A53x1 FVP are included with Arm DS.

---

*   Download this Arm DS example project Arm DS example project) that includes the code from this tutorial and a debug launch configuration. Use this project to modify, build, and debug the examples from this tutorial.

Note Arm strongly recommends that you complete the Hello World Arm DS Tutorial before working with the examples in this tutorial.

Using the examples in this tutorial

The examples in this tutorial use a little-endian memory system.

*   For 32-bit registers, `unsigned int`.

*   For 16-bit registers, `unsigned short`.

*   For 8-bit registers, `unsigned char`.

The compiler generates the appropriate single load and store instructions, that is `LDR` and `STR` for 32-bit registers, `LDRH` and `STRH` for 16-bit registers, and `LDRB` and `STRB` for 8-bit registers.

You must also ensure that the memory-mapped registers lie on appropriate address boundaries, that is either all word-aligned, or aligned on their natural size boundaries. For example, 16-bit registers must be aligned on halfword addresses.

> **Note**
>
> Arm recommends that all registers, whatever their size, be aligned on word boundaries.

You can also use `#define` to simplify your code:

```c
#define PORTBASE 0xC0000000 /* Counter/Timer Base */
#define PortLoad ((volatile unsigned int *) PORTBASE) /* 32 bits */
#define PortValue ((volatile unsigned short *)(PORTBASE + 0x04)) /* 16 bits */
#define PortClear ((volatile unsigned char *)(PORTBASE + 0x08)) /* 8 bits */

void init_regs(void)
{
  unsigned int int_val;
  unsigned short short_val;
  unsigned char char_val;

  *PortLoad = (unsigned int) 0xF00FF00F;
  int_val = *PortLoad;

  *PortValue = (unsigned short) 0x0000;
  short_val = *PortValue;

  *PortClear = (unsigned char) 0x1F;
  char_val = *PortClear;
}
```

This code results in the following (interleaved) code:

```
;;;5       void init_regs(void)
000000  e59f1024 LDR r1,|L1.44|
;;;6       {
;;;7         unsigned int int_val;
;;;8         unsigned short short_val;
;;;9         unsigned char char_val;
;;;10        *PortLoad = (unsigned int) 0xF00FF00F;
000004  e3a00101 MOV r0,#0xC0000000
000008  e5801000 STR r1,[r0,#0]
;;;11        int_val = *PortLoad;
00000c  e5901000 LDR r1,[r0,#0]
;;;12        *PortValue = (unsigned short) 0x0000;
000010  e3a01000 MOV r1,#0
000014  e1c010b4 STRH r1,[r0,#4]
;;;13        short_val = *PortValue;
000018  e1d010b4 LDRH r1,[r0,#4]
;;;14        *PortClear = (unsigned char) 0x1F;
00001c  e3a0101f MOV r1,#0x1f
000020  e5c01008 STRB r1,[r0,#8]
;;;15        char_val = *PortClear;
000024  e5d00008 LDRB r0,[r0,#8]
;;;16      }
000028  e12fff1e BX lr
```

If you debug the previous code with the Cortex-A53 FVP in Arm DS, you can see the generated disassembly code in the Disassembly view.

Note

The instructions used might slightly differ from the previous interleaved code.

Figure 1-1: Code in Disassembly view

```
Console  Commands  Variables  Registers  Memory  Disassembly ⊠  +

    <Next Instruction>                                                    100

          Address            Opcode    Disassembly
                                       init_regs
⇨ EL3:0x00000000800000E8  D10043FF  SUB       sp,sp,#0x10
  EL3:0x00000000800000EC  D2B80008  MOV       x8,#0xc0000000
  EL3:0x00000000800000F0  32049FE9  ORR       w9,wzr,#0xf00ff00f
  EL3:0x00000000800000F4  B9000109  STR       w9,[x8,#0]
  EL3:0x00000000800000F8  B9400108  LDR       w8,[x8,#0]
  EL3:0x00000000800000FC  B9000FE8  STR       w8,[sp,#0xc]
  EL3:0x0000000080000100  D2800088  MOV       x8,#4
  EL3:0x0000000080000104  F2B80008  MOVK      x8,#0xc000,LSL #16
  EL3:0x0000000080000108  7900011F  STRH      wzr,[x8,#0]
  EL3:0x000000008000010C  79400108  LDRH      w8,[x8,#0]
  EL3:0x0000000080000110  790017E8  STRH      w8,[sp,#0xa]
  EL3:0x0000000080000114  D2800108  MOV       x8,#8
  EL3:0x0000000080000118  F2B80008  MOVK      x8,#0xc000,LSL #16
  EL3:0x000000008000011C  528003E9  MOV       w9,#0x1f
  EL3:0x0000000080000120  39000109  STRB      w9,[x8,#0]
  EL3:0x0000000080000124  39400108  LDRB      w8,[x8,#0]
  EL3:0x0000000080000128  390027E8  STRB      w8,[sp,#9]
  EL3:0x000000008000012C  910043FF  ADD       sp,sp,#0x10
  EL3:0x0000000080000130  D65F03C0  RET
  EL3:0x0000000080000134  00000000  UDF       #0
```

Note

The instructions LDR, STR, LDRH, STRH, LDRB, and STRB, the compiler generates as type casting (unsigned int, unsigned short, and unsigned char) which is used to store the wanted values in the peripheral port registers.

It is now possible to debug the code to check the wanted values are correctly written to the wanted memory addresses that correspondent to the peripheral registers. Before executing init_regs(), open the Memory view in Arm DS and then search for the port base address of your peripheral, in this case 0xC0000000. It is observed that the relevant memory addresses only contain uninitialized values:

Figure 1-2: Code in dissassmbly view

Stepping through the code, when `*PortLoad = (unsigned int) 0xF00FF00F;` is executed, the instruction `STR w9,[x8,#0]` is executed to store the wanted value, `0xF00FF00F`, into the peripheral register `PortLoad`, at the address `0xC0000000`:

Figure 1-3: Stepping through the code

```c
void init_regs(void)
{
  unsigned int int_val;
  unsigned short short_val;
  unsigned char char_val;

  *PortLoad = (unsigned int) 0xF00FF00F;
  int_val = *PortLoad;

  *PortValue = (unsigned short) 0x0000;
  short_val = *PortValue;

  *PortClear = (unsigned char) 0x1F;
  char_val = *PortClear;

  return;
}
```

Figure 1-4: Storing a wanted value

Open the Memory view to see that the wanted value is set to the peripheral register:

Figure 1-5: Wanted value set to peripheral register

Continue stepping until the instruction STRH wzr,[x8,#0] is reached. This instruction is executed to store the value 0x0000 in the PortValue register, mapped to the address 0xC0000004. In this case, the value set on the port register is half a word width (16-bits). When viewing memory at a word width, after the halfword write, the bottom 16-bits change and the top 16-bits remain the same. In this case, the bottom 16-bits change to 0x0000 and the top 16-bits remain 0xCFDF like the following:

Figure 1-6: Viewing memory at a word width after the halfword write

Continue stepping to the next relevant instruction of STRB w9,[x8,#0]. This instruction is executed to store the value 0x1F in the PortClear register, mapped to the address 0xC0000008. In this case, the value set on the port register is a byte wide (8-bits). When viewing memory at a word width, after the byte write, the bottom 8-bits change and the top 24 bits (3 bytes) remain the same. In this case, the bottom 8-bits change to 0x1F and the top 3 bytes remain 0xDFDFDF like the following:

Figure 1-7: Viewing memory at a word width after the byte write
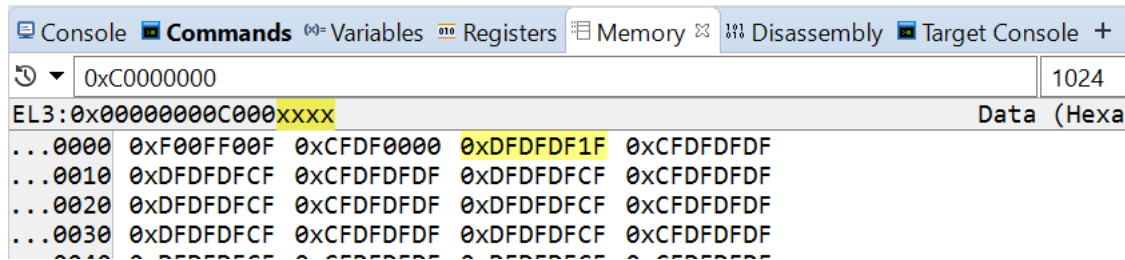
To check the variables are stored to the correct peripheral addresses, in the Variables view, look at the local variables `int_val`, `short_val`, and `char_val`. Look at the values, types, and sizes of the three variables.

Tip: Even if the sizes are different, word-alignment (32-bits) is respected when storing the values in the peripheral registers. Notice that the locations of these variables are not the locations of the peripheral registers. These variables are only local variables to check that the contents of the peripheral registers are the expected values.

Figure 1-8: Local variables to check that the contents of the peripheral registers

2. Arm recommendations

Arm recommends word alignment of peripheral registers even if they are 16-bit or 8-bit peripherals. In a little-endian system, the peripheral databus can connect directly to the least significant bits of the Arm databus. There is no need to multiplex or duplicate the peripheral databus onto the high bits of the Arm databus. In a big-endian system, the peripheral databus can connect directly to the most significant bits of the Arm databus. There is no need to multiplex or duplicate the peripheral databus onto the low bits of the Arm databus.

The Arm AMBA APB bridge uses the preceding recommendation to simplify the bridge design. The result of this is that only word-aligned addresses must be used for any width transfer. A read results in unused values on any bits which are not connected to the peripheral. So, if a 32-bit word is read from a 16-bit peripheral, the top 16 bits of the register value must be cleared before use.

For example, to access some 16-bit peripheral registers on a 16-bit alignment, you might write:

```
volatile unsigned short u16_IORegs[20];
```

This code works if your peripheral controller logic can route the peripheral databus to the high part (D31-D16) and low part (D15-D0) of the Arm databus. Which part is used depends on which address you are accessing. To use this code, check if this multiplexing logic exists in your design. The standard Arm APB bridge does not support this multiplexing logic.

3. Aligning registers

If you want to map 16-bit registers on 32-bit alignment as recommended, you could use on of these methods:

1. ```
   volatile unsigned short u16_IORegs[40];
   ```

   This code only allows accesses to even-numbered registers (each index in the array corresponds to 16-bits). You must double the register number. For example, to access the fourth register you could use index 8:

   ```
   x = u16_IORegs[8];
   u16_IORegs[8] = newval;
   ```

2. ```
   volatile unsigned int u32_IORegs[20];
   ```

   The registers are accessed as 32-bit values. But a simple peripheral controller, like an Arm AMBA APB bridge, reads unused values into the top bits from signals that are not connected to the peripheral. In a little-endian system, the used values map to D31–D16. So, when such a peripheral is read, it must be cast to an unsigned short to get the compiler to discard the upper 16-bits.

   This type casting effect was the case shown in the previous example in this documentation. This example used a casting to read the peripheral registers and save their contents in the variables `int_val`, `short_val`, and `char_val`.

   For example, to access peripheral register 4:

   ```
   x = (unsigned short)u32_IORegs[4];
   u32_IORegs[4] = newval;
   ```

3. Use a `struct`.

   Using a `struct` allows descriptive names to be used and can accommodate different peripheral register widths.

   > **Note**
   >
   > The padding is made explicit rather than relying on automatic padding added by the compiler. For example:

   ```
   struct PortRegs {
       unsigned short ctrlreg; /* offset 0 */
       unsigned short dummy1;
       unsigned short datareg; /* offset 4 */
       unsigned short dummy2;
       unsigned int data32reg; /* offset 8 */
   } iospace;

   x = iospace.ctrlreg;
   iospace.ctrlreg = newval;
   ```

The peripheral locations must not be accessed using:

Note

---

a. `__packed structs` where unaligned members are allowed and there is no internal padding

    OR

b. C bitfields, for example, by defining the variable unsigned int `isCorrect: 1`.

If the previous access methods are used, the compiler only uses 1-bit to store data, instead of 32-bits. This 1-bit data store is common when storing the data as a Boolean. However, if you use the previous methods, it is not possible to control the number and type of memory accesses the compiler performs. This can result in code that is non-portable, has undesirable side-effects, and does not work as intended. The recommended way of accessing peripherals is through explicit use of architecturally defined types such as `int`, `short`, and `char` on their natural alignments.

4. Mapping variables to specific addresses

Memory mapped registers can be accessed from C in two ways:

- Force an array or struct variable to a specific address

- Using a pointer to an array or struct

Both previous methods generate efficient code. Choose the method that you prefer.

Force an array or struct variable to a specific address

The array or struct variable is declared in a file on its own. When it is compiled, the object code for this file only contains data. If using the Arm Compiler, this data can be placed at a specified address using the Arm scatter-loading mechanism. Arm recommends this scatter-loading mechanism as the method for placing all scatter-loading AREAS at required locations in the memory map.

4.1 Test using a struct at a specific address

Test the recommended scatter-loading mechanism.

1. In the example project, open src > `iovar.c`.

    The `iovar.c` file contains a declaration of the array or `struct` variable:

    ```
    struct{
       volatile  unsigned  reg1;
       volatile  unsigned  reg2;
    } mem_mapped_reg;
    ```

2. Open `scatter.txt`.

    The `scatter.txt` file contains the following:

    ```
    ALL 0x80000000
    {
       ALL 0x80000000
       {
         *  (+RO,+RW,+ZI)
       }
    }
      IO 0xC0000000
       {
         IO 0xC0000000
          {
           iovar.o (+ZI)
          }
       }
    ```

    The scatter-loading description file must be specified at link time to the linker using the `--scatter scatter.txt` command-line option. This description creates two different load regions in your image: ALL and IO. The zero-initialized area (ZI) from `iovar.o` containing the `struct`,

goes into the I/O area at `0xC0000000`. All code (RO) and data areas (RW and ZI) from other object files go into the `ALL` region which starts at `0x80000000`.

4.2 Add the scatter-loading description easily to an Arm DS project

Use this method to add the scatter-loading description to a project.

1. In the Project Explorer view, right-click your project folder, and click Properties.

2. In C/C++ Build > Settings> Tool Settings, click Image Layout under Arm Linker 6.

3. Click Browse to select the scatter file, and click Apply and Close to save the changes.

Figure 4-1: Adding a scatter-loading file to a project
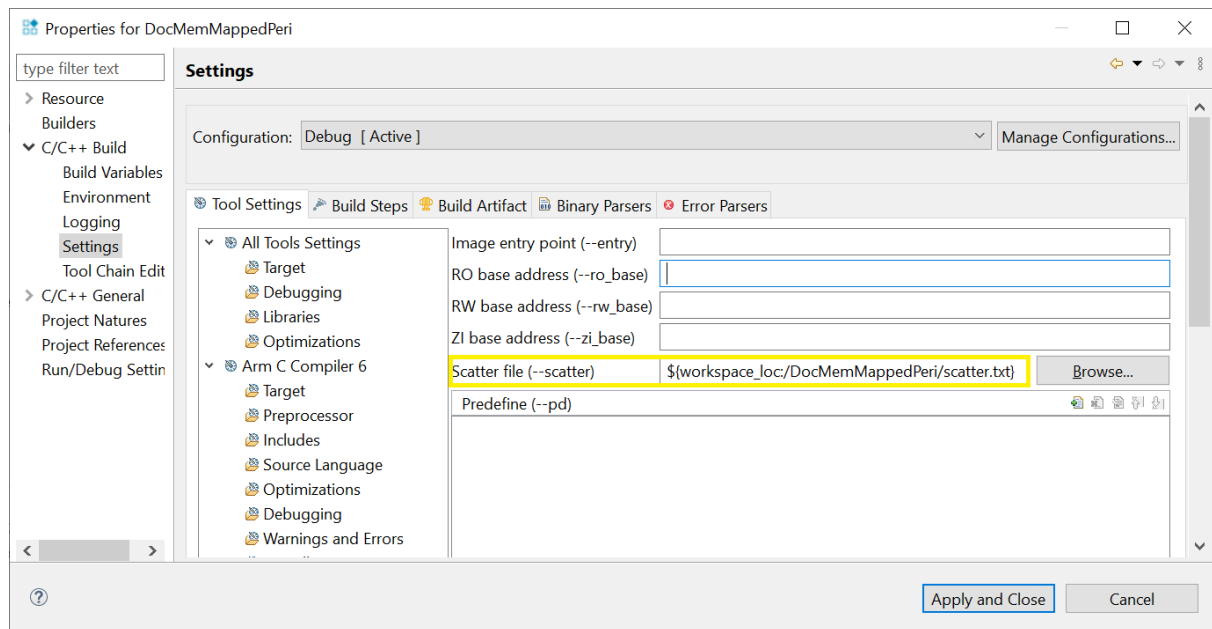
> **Note**
>
> In the example project, the scatter-loading file has already been added to the project.

If you have more than one set of memory mapped registers, you must define each group of variables as a separate execution region. It is possible that all the memory mapped registers could be within a single load region. To define each variable group as a separate execution region, each group of variables must be defined in a separate module.

The benefit of using a scatter-loading description file is that all the (target-specific) absolute addresses chosen for your devices, code, and data are in one file. Because everything is in one

file, maintenance is easy. Furthermore, if you decide to change your memory map, for example, if peripherals are moved, you do not need to rebuild your entire project. You just run the link step on the existing object files.

For more documentation on scatter-loading, check the following links:

- See the scatter-loading mechanism chapters in the Arm Compiler Reference Guide
- See the scatter-loading images with a simple memory map chapter in the Arm Compiler Reference Guide
- See the scatter File Syntax chapters in the Arm Compiler Reference Guide

Alternatively, use the `#pragma arm section` pragma to place the data into a specific section. Then, use scatter-loading to place that data at an explicit location. See the pragmas chapter on Arm Developer.

4. In the example project `DocMemMappedPeri.c`, uncomment the following code:

```
extern struct{
    volatile unsigned reg1;
    volatile unsigned reg2;
 } mem_mapped_reg;
int main(void) {
...
    mem_mapped_reg.reg1 = (unsigned int) 0xF00FF00F;
    mem_mapped_reg.reg2 = (unsigned int) 0x100CF00F;
...
}
```

And comment this line in `main()`:

```
//init_regs();
```

5. Clean and build the example project.

6. Launch the Cortex-A53 FVP.

7. Step through the code.

The fields `reg1` and `reg2` in the `struct` variable `mem_mapped_reg` are now mapped to the addresses of the peripheral registers.

When writing to these variables, you are directly writing to the peripheral registers. Enter the address of the peripheral registers, `0xC0000000`, into the Memory view. Observe how the data, `0xF00FF00F` and `0x100CF00F`, is correctly written by using the memory-mapped variables:

Figure 4-2: Memory view after struct write
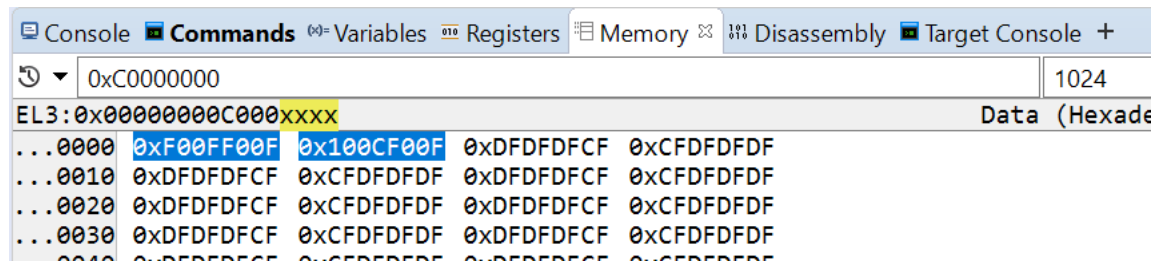
Using a pointer to an array or struct

```
struct PortRegs {
  unsigned short ctrlreg; /* offset 0 */
  unsigned short dummy1;
  unsigned short datareg; /* offset 4 */
  unsigned short dummy2;
  unsigned int data32reg; /* offset 8 */
};
volatile struct PortRegs *iospace = (struct PortRegs *)0xC0000000;
x = iospace->ctrlreg;
iospace->ctrlreg = newval;
```

The pointer can be either local or global. If global, to avoid having the base pointer reloaded after function calls, make iospace a constant pointer to the struct by changing its definition to:

```
volatile struct PortRegs * const iospace = (struct PortRegs *)0xC0000000;
```

5. Code efficiency

The Arm compiler normally uses a base register and the immediate offset field to compile `struct` members or specific array element accesses. The immediate offset is available in the load or store instructions. In the Arm instruction set, `LDR` and `STR` word and byte instructions have a 4KB range, but `LDRH` and `STRH` instructions have a smaller immediate offset of 256 bytes.

Equivalent 16-bit Thumb instructions are much more restricted. `LDR` and `STR` have a range of 32 words, `LDRH` and `STRH` have a range of 32 halfwords, and `LDRB` and `STRB` have a range of 32 bytes. However, 32-bit Thumb instructions offer a significant range improvement. Because of the range restrictions, it is important to group related peripheral registers near to each other if possible. The compiler is generally good at minimizing the number of instructions required to access array elements or structure members. To perform these accesses, the compiler uses base registers.

You can choose between using one large C `struct` or `array` for the whole I/O space and smaller per-peripheral structs. For the two methods, there is little difference in terms of code efficiency. Using a large `struct` might help if:

- For word and byte accesses, your code works with a base pointer with a 4KB range.

- Entire I/O space is <4KB.

But arguably, it is more elegant to have one `struct` for each peripheral. Smaller, per-peripheral structs are more maintainable.