



Debugging a Linux Symmetric Multi-Processing (SMP) Kernel using DS-5

Version 1.0

Non-Confidential

Copyright © 2020 Arm Limited (or its affiliates).
All rights reserved.

Issue 01

102608_0100_01_en



Debugging a Linux Symmetric Multi-Processing (SMP) Kernel using DS-5

Copyright © 2020 Arm Limited (or its affiliates). All rights reserved.

Release information

Document history

| Issue | Date | Confidentiality | Change |
|---------|----------------|------------------|---------------|
| 0100-01 | 1 January 2020 | Non-Confidential | First release |

Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED "AS IS". ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly

or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its affiliates) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <https://www.arm.com/company/policies/trademarks>.

Copyright © 2020 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.

110 Fulbourn Road, Cambridge, England CB1 9NJ.

(LES-PRE-20349|version 21.0)

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Unrestricted Access is an Arm internal classification.

Product Status

The information in this document is Final, that is for a developed product.

Feedback

Arm® welcomes feedback on this product and its documentation. To provide feedback on the product, create a ticket on <https://support.developer.arm.com>

To provide feedback on the document, fill the following survey: <https://developer.arm.com/documentation-feedback-survey>.

Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used language that can be offensive. Arm strives to lead the industry and create change.

We believe that this document contains no offensive language. To report offensive language in this document, email terms@arm.com.

Contents

| | |
|--|----|
| 1. Overview..... | 6 |
| 2. Prerequisites..... | 7 |
| 3. Creating a debug configuration..... | 10 |
| 4. Debugging the kernel - pre-MMU stage..... | 13 |
| 5. Debugging the kernel - post-MMU stage..... | 21 |
| 6. Viewing the structure of the kernel..... | 24 |
| 7. Multi-core features..... | 28 |
| 8. Logging features..... | 30 |
| 9. Single-stepping individual cores or threads..... | 31 |
| 10. Using expressions to view data..... | 32 |
| 11. Using the command line to view cores, thread, and process information..... | 33 |

1. Overview

You can develop, port, and debug the Linux kernel for a platform using DS-5 Debugger. You can set breakpoints, view registers, view memory, single step at source level, and perform all other standard debugging functions. You can debug the Linux kernel before the MMU is enabled (with a physical memory map), and after the MMU is enabled (with a virtual memory map). DS-5 allows you to do all this and more, not just for single-core platforms, but for SMP platforms too.

In DS-5, you configure a debugging session to a target using the **Debug Configuration** dialog. The Linux kernel debug options available in the dialog are primarily designed for post-MMU debug. But with some extra controls, you can use it for pre-MMU debug also. This makes it possible to debug the Linux kernel at source-level, all the way from its entry point, through the pre-MMU stages. You can then continue debugging seamlessly through the MMU enable stage to the post-MMU debug stage with full kernel awareness. This is achieved without disconnecting, reconfiguring, and reconnecting to your target.

This tutorial takes you through:

- [Creating a debug configuration](#), and loading the Linaro Linux 3.4 SMP kernel on a Cortex-A9x4 Fixed Virtual Platform (FVP) model supplied with DS-5.
- [Debugging the kernel - pre-MMU stage](#).
- [Debugging the kernel - post-MMU stage](#).
- [Viewing the structure of the kernel](#).
- [Multi-core features](#).

This tutorial also briefly looks at [Logging features](#), [Single-stepping individual cores or threads](#), and [Using the command line to view cores, thread, and process information](#).

2. Prerequisites

In the following, you can find the prerequisites to this tutorial:

- **DS-5 installation and licensing**

You need to have Arm DS-5 installed and the license to use it. If you do not have a license, see the [Getting Started tutorial](#) to see how to obtain a license.

- **Enable telnet on your Windows host**

The FVP model uses a telnet client for its simulated serial ports. Telnet is disabled by default in newer versions of Windows. You need to enable it. Check the documentation for your Windows version for details about enabling the telnet client.

- **DS-5 Linux Distribution Example**

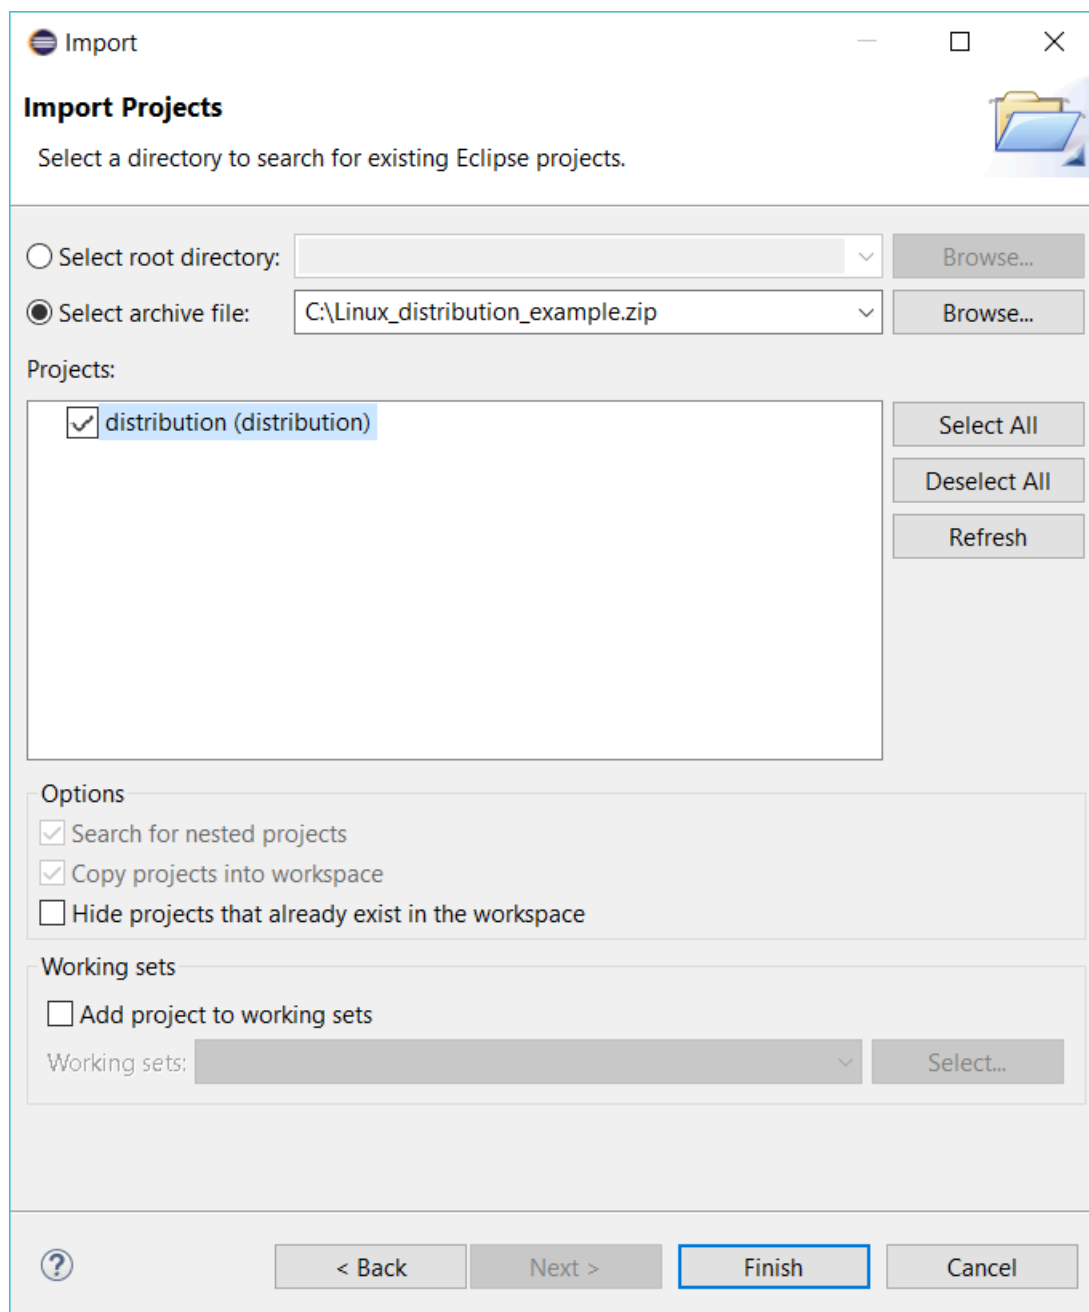
Download and import the **DS-5 Linux Distribution Example** into DS-5. This archive file contains the Linux kernel built with debug information, the `vmlinux` symbol file, file system, and the full source code.



If you are working on a Windows machine, you will encounter issues with file naming incompatibilities between Linux and Windows. These can be ignored for the purpose of this tutorial.

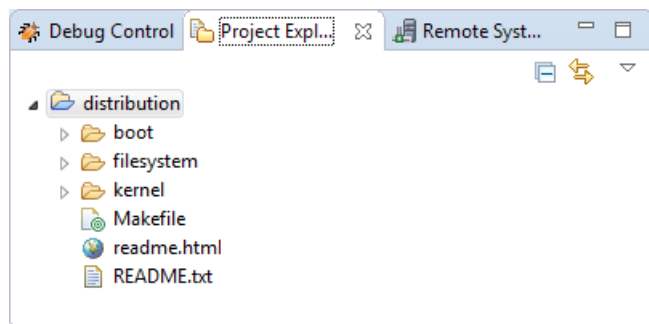
1. Go to the [DS-5 Downloads](#) page and download the `Linux_distribution_example.zip` file from the **DS-5 Extra Downloads** section.
2. In DS-5, select **File > Import** to display the **Import** dialog.
3. Under **General**, select **Existing Projects into Workspace** and click **Next**.
4. Select the **Select archive file** option, and **Browse** and select the file you downloaded from the Arm download center. The **distribution** project is available in the projects area.

Figure 2-1: Import projects window



5. Click **Finish** to import the files into the project folder. You can view the imported files in the **Project Explorer**.

Figure 2-2: Project Explorer folder view



- **Set up the sources**

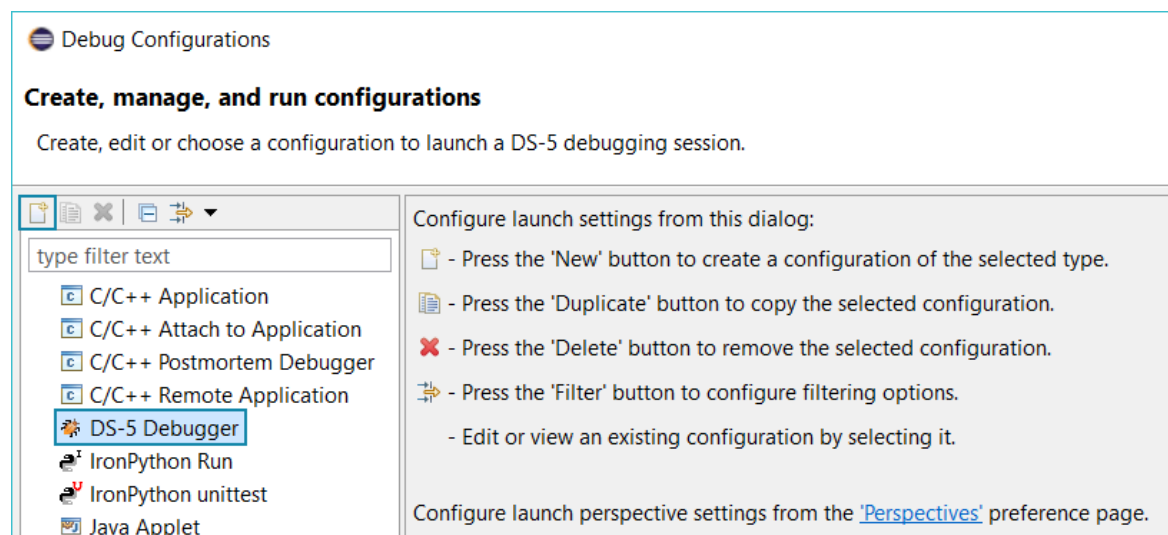
The DS-5 Linux Distribution Example archive file also contains the required sources. You can find it under <DS-5 Workspace>\distribution\kernel\linux-linaro-3.4-rc3-2012.04-0-patched\source.tar.bz2. Unzip this file into a location on your workstation.

3. Creating a debug configuration

Start by creating a **Debug Configuration** to configure connections to the Fixed Virtual Platform (FVP) model. When you create the **Debug Configuration**, you set up the kernel to load on the FVP.

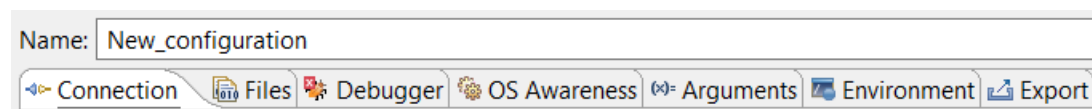
1. Launch DS-5.
2. From the main menu, select **Run > Debug Configurations** to open the **Debug Configurations** dialog.
3. In the **Debug Configurations** dialog:
 - a. Select **DS-5 Debugger**.
 - b. Click the **New launch configuration** button.

Figure 3-1: New launch configuration



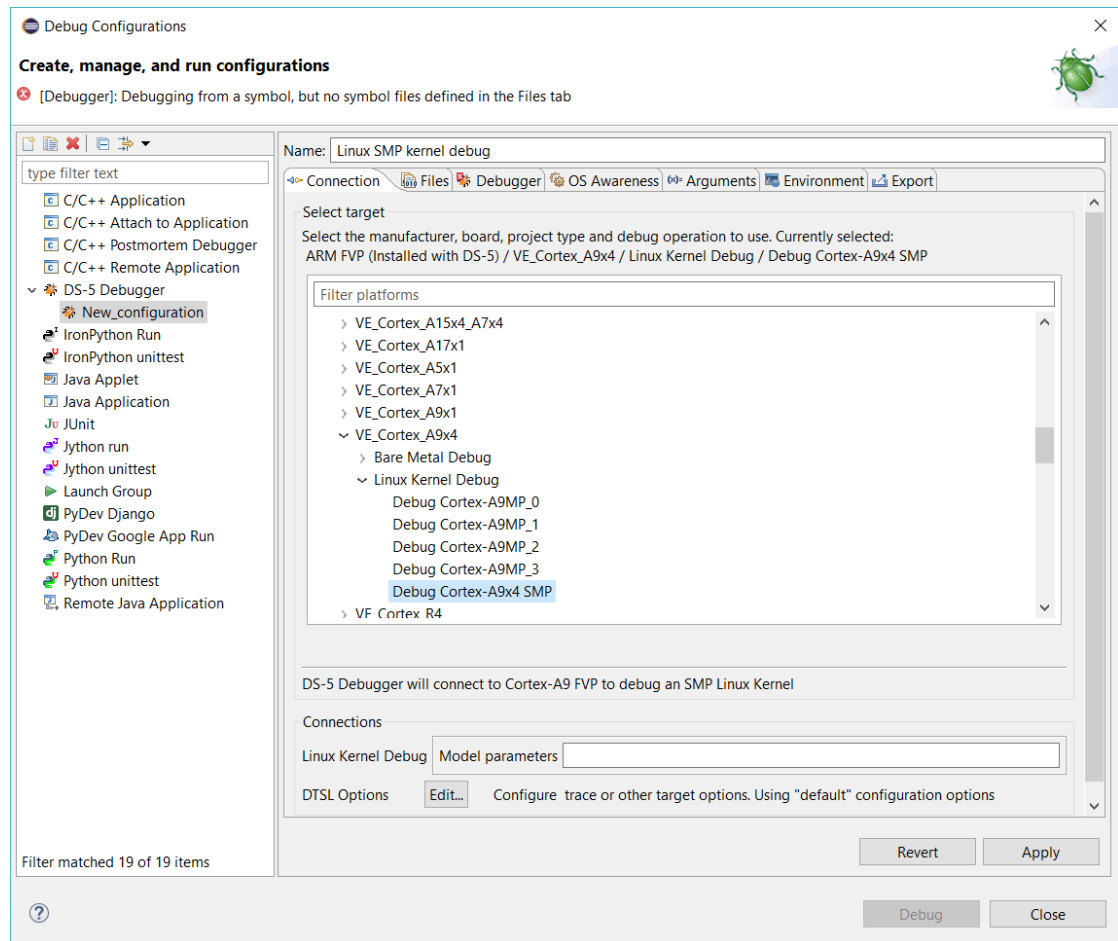
This creates a DS-5 debug configuration and displays the various tabs required to specify settings for loading your application on the target.

Figure 3-2: New configuration



- c. Give a name to the debug configuration. For example, **Linux SMP kernel debug**.
- d. In the **Connection** tab:
 1. Under **Select target**, browse and select **Arm FVP (Installed with DS-5) > VE_Cortex-A9x4 > Linux Kernel Debug > Debug Cortex-A9x4 SMP**.

Figure 3-3: Select target



2. Under **Connections > Linux Kernel Debug**, in the **Model parameters** field, enter the following as a single line:



- The code below is listed in separate lines for reading clarity. It must be in a single line with single spaces between the commands when copied to the Model parameters field.
- Replace <DS-5 Install folder> with the path to your DS-5 installation, and <DS-5 Workspace> with the path to your workspace.

```
--data "<DS-5 Install folder>\arm\linux_distribution\kernel_ve@0x80008000"

--data "<DS-5 Install folder>\arm\linux_distribution\rtsm_ve-
cortex_a9x4.dtb@0x80f00000"

-C motherboard.mmc.p_mmc_file="<DS-5 Install folder>\arm\linux_distribution
\rootfs.image"
```

```
-C motherboard.vfs2.mount="<DS-5 Workspace>"

-C motherboard.terminal_3.start_telnet=false

-C motherboard.terminal_3.mode=raw

-C motherboard.pl011_uart3.untimed_fifos=true

-C motherboard.terminal_3.start_port=5003

-C motherboard.smc_91c111.enabled=1

-C motherboard.hostbridge.userNetworking=1

-C motherboard.hostbridge.userNetPorts="8080=8080"
```

e. In the **Files** tab:

1. Under **Target Configuration** > **Application on host to download** field, click **File System**.
2. Browse to your DS-5 installation folder, and select <DS-5 Install folder>\arm\nlinux_distribution\bootwrapper_ve.axf to load the bootwrapper.

What is a boot wrapper?

The boot wrapper (bootwrapper_ve.axf) is a small program that must execute before starting the kernel.

The boot wrapper is responsible for setting up the execution environment for the kernel, including:

- Detecting the processor type.
- Initializing the vector table.
- Setting up the RAM.
- Initializing a serial port.
- Handling the device tree and kernel command line.
- And finally, starting the kernel.

3. Deselect the **Load symbols** option. The kernel debug information is loaded later.

f. In the **Debugger** tab, under **Run Control**, select **Debug from entry point**.

g. Click **Debug** to start debugging.

4. If the **Confirm Perspective Switch** dialog is displayed, click **Yes** to switch to the **DS-5 Debug Perspective**.

DS-5 connects to the model and displays the connection status in the **Debug Control** view and opens a window.

You have now created a debug configuration, loaded the files on the FVP, and started the debug connection. The next step is to start debugging the kernel.

4. Debugging the kernel - pre-MMU stage

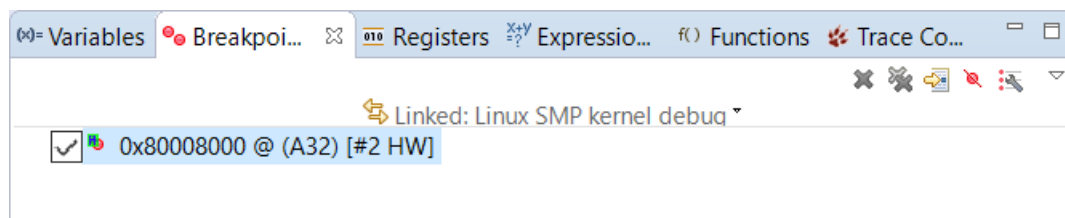
The following steps shows you how to debug the kernel in the pre-MMU stage:

Set a temporary hardware breakpoint on the entry point into the kernel

1. In the **Commands** view, enter `thbreak 0x80008000` to set a temporary hardware breakpoint on the entry point into the kernel. `0x80008000` is the base address in RAM where this Linux kernel is loaded, and is also the entry point for the kernel. It is the same address as specified in the `--data "<DS-5 Install folder>\arm\linux_distribution\kernel_ve@0x80008000"` parameter to the FVP model. This address is also hard-coded within the boot wrapper as the address to jump to when the boot wrapper has completed its setup tasks.

You can view the breakpoint in the Breakpoints view.

Figure 4-1: Breakpoint view



2. In the **Debug Control** view, click **Continue**  to start the target. The code is executed and stops at the breakpoint.

At this stage, you can:

- either continue with the next step - **Disable OS awareness before loading the debug symbols**
- or analyze the system registers in the **Registers** view:

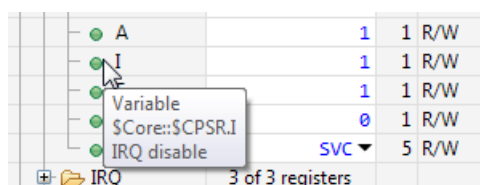
Analyzing the system registers in the Registers view:

1. You can use the **Registers** view to check if the **Core** and **CP15** system registers are set, as recommended by the [Booting Arm Linux reference page](#) (For Armv8 targets use the [AArch64 Linux reference page](#)).

For example, using the **Registers** view you can:

- Check the Core is in SVC (supervisor) mode with both **IRQ** and **FIQ** interrupts disabled.

Figure 4-2: IRQ and FIQ interrupts disabled



- Check R0 is 0.

Figure 4-3: R0 is 0

| Core | | 47 of 47 registers | | |
|------|------------|--------------------|-----|--|
| R0 | 0x00000000 | 32 | R/W | |
| R1 | 0xFFFFFFFF | 32 | R/W | |
| R2 | 0x80F01310 | 32 | R/W | |
| R3 | 0x00000000 | 32 | R/W | |
| R4 | 0x80000530 | 32 | R/W | |

- Check R1 contains 0xFFFFFFFF to indicate this is a “device tree” platform.

Figure 4-4: R1 contains 0xFFFFFFFF

| Core | | 47 of 47 registers | | |
|------|------------|--------------------|-----|--|
| R0 | 0x00000000 | 32 | R/W | |
| R1 | 0xFFFFFFFF | 32 | R/W | |
| R2 | 0x80F01310 | 32 | R/W | |
| R3 | 0x00000000 | 32 | R/W | |
| R4 | 0x80000530 | 32 | R/W | |

- Check R2 contains a pointer to the device tree. Right-mouse click R2 and select **Show Memory Pointed To By R2**.

Figure 4-5: R2 contains a pointer to the device tree

| Register | Value | Width | Access |
|----------|------------|-------|--------|
| R0 | 0x00000000 | 32 | R/W |
| R1 | 0xFFFFFFFF | 32 | R/W |
| R2 | 0x80F01310 | 32 | R/W |
| R3 | 0x00000000 | 32 | R/W |
| R4 | 0x80000660 | 32 | R/W |
| R5 | 0x00000000 | 32 | R/W |
| R6 | 0x00000000 | 32 | R/W |
| R7 | 0x00000000 | 32 | R/W |
| R8 | 0x00000000 | 32 | R/W |
| R9 | 0x00000000 | 32 | R/W |
| R10 | 0x00000000 | 32 | R/W |
| R11 | 0x00000000 | 32 | R/W |

Context Menu Options:

- Copy (Ctrl+C)
- Paste (Ctrl+V)
- Delete
- Select All (Ctrl+A)
- Show Memory Pointed To By R2
- Show Disassembly Pointed To By R2
- Translate Address in R2
- Send to:
 - Binary
 - Boolean
 - Float
 - Hexadecimal (selected)
 - Octal
 - Signed Decimal
 - Unsigned Decimal

- Expand CP15 > system > SCTLR > M and check that the MMU is off.

Figure 4-6: MMU is disabled

| | | | |
|--------|-------------------------|----|-----|
| CP15 | 134 of 134 registers | | |
| System | 28 of 28 registers | | |
| MIDR | 0x413FC090 | 32 | RO |
| SCTLR | 0x00C50078 | 32 | R/W |
| TE | Disabled | 1 | R/W |
| AFE | Full_access_permissions | 1 | R/W |
| TRE | Disabled | 1 | R/W |
| NMFI | Disabled | 1 | RO |
| EE | 0 | 1 | R/W |
| HA | 0 | 1 | R/W |
| V | Normal_vectors | 1 | R/W |
| I | Disabled | 1 | R/W |
| Z | Disabled | 1 | R/W |
| SW | Disabled | 1 | R/W |
| C | Disabled | 1 | R/W |
| A | Disabled | 1 | R/W |
| M | Disabled | 1 | R/W |

- Expand CP15 > system > SCTLR > c and check that the data cache is off.

Figure 4-7: Data cache is disabled

| | | | |
|--------|-------------------------|----|-----|
| CP15 | 134 of 134 registers | | |
| System | 28 of 28 registers | | |
| MIDR | 0x413FC090 | 32 | RO |
| SCTLR | 0x00C50078 | 32 | R/W |
| TE | Disabled | 1 | R/W |
| AFE | Full_access_permissions | 1 | R/W |
| TRE | Disabled | 1 | R/W |
| NMFI | Disabled | 1 | RO |
| EE | 0 | 1 | R/W |
| HA | 0 | 1 | R/W |
| V | Normal_vectors | 1 | R/W |
| I | Disabled | 1 | R/W |
| Z | Disabled | 1 | R/W |
| SW | Disabled | 1 | R/W |
| C | Disabled | 1 | R/W |
| A | Disabled | 1 | R/W |
| M | Disabled | 1 | R/W |

Variable
SCP15::\$System::\$SCTLR.C
Determines if data can be cached in a data or unified cache at any cache level

- Expand CP15 > system > SCTLR > I and check the status of the instruction cache. Depending on the status, this could be either on or off.

Figure 4-8: Check the instruction cache status

| | | | | |
|----------|--|----|-----|--|
| CP15 | 134 of 134 registers | | | |
| System | 28 of 28 registers | | | |
| MIDR | 0x413FC090 | 32 | RO | |
| SCTLR | 0x00C50078 | 32 | R/W | |
| TE | Disabled | 1 | R/W | |
| AFE | Full_access_permissions | 1 | R/W | |
| TRE | Disabled | 1 | R/W | |
| NMFI | Disabled | 1 | RO | |
| EE | 0 | 1 | R/W | |
| HA | 0 | 1 | R/W | |
| V | Normal_vectors | 1 | R/W | |
| I | Disabled | 1 | R/W | |
| Variable | SCP15::System::SCTLR.I | | | |
| | Determines if instructions can be cached in any instruction cache at any cache level | | | |

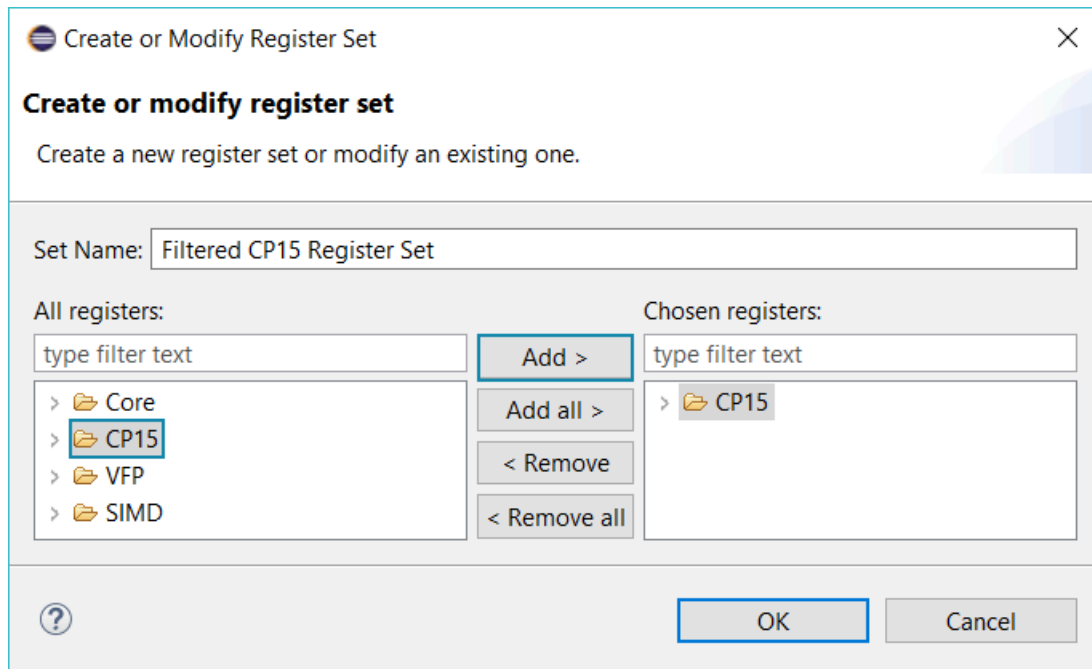
- You can set up the **Registers** view to filter the displayed additional information. For example, if you only want to view information present in the CP15 group of registers, you have to create a new **Register set**:
 - In the **Registers** view, use the drop down arrow to expand the viewed register sets and select **Create**.

Figure 4-9: Create registers

| | | | | |
|--|---------------|--------|-----|--|
| Variables Breakpoi... 010 Registers Expressio... f() Functions Trace Co... | | | | |
| Linked: Linux SMP kernel debug | | | | |
| Register Set: | All registers | | | |
| Name | All registers | | | |
| Core | | | | |
| R0 | Create | Manage | | |
| R1 | | | | |
| R2 | 0x80F01310 | 32 | R/W | |
| R3 | 0x00000000 | 32 | R/W | |
| R4 | 0x80000660 | 32 | R/W | |
| R5 | 0x00000000 | 32 | R/W | |
| R6 | 0x00000000 | 32 | R/W | |
| R7 | 0x00000000 | 32 | R/W | |
| R8 | 0x00000000 | 32 | R/W | |
| R9 | 0x00000000 | 32 | R/W | |
| R10 | 0x00000000 | 32 | R/W | |
| R11 | 0x00000000 | 32 | R/W | |

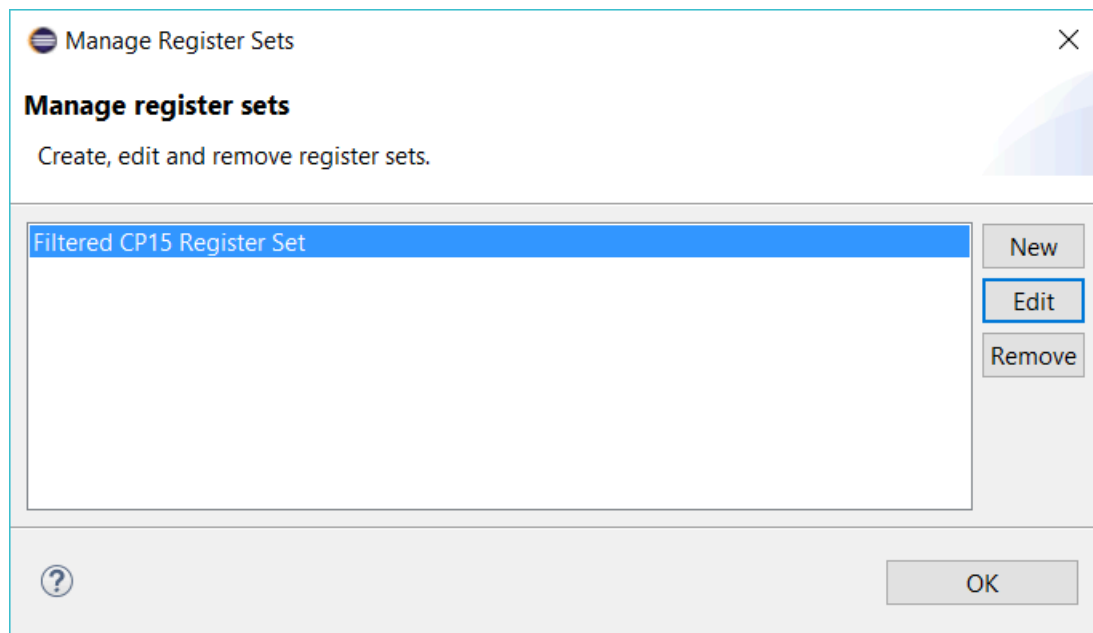
- In the **All registers** section, select CP15 and click **Add >**. Your selected registers appear under **Chosen registers**. Give the register set a name such as **Filtered CP15 Register Set**. You can create multiple register groups if needed. To finish creating your new register sets, click **OK**.

Figure 4-10: Filtered CP15 Register Set



- The **Registers** view displays the specific register group you selected.
- To switch between various register groups, click the All Registers drop down and select the group you want.
- To edit or remove existing register sets, in the Registers view, under the Register Set drop down, select Manage to open the Manage Register Set dialog.
 - To edit an existing register set, select a register set, and click Edit. Edit the register set in the Create or Modify Register Set dialog.
 - To remove a register set, select the register set you want to delete and click Remove.
 - Click OK to confirm your changes.

Figure 4-11: Manage register sets



Disable OS awareness before loading the debug symbols

Since the MMU is currently off, DS-5 Debugger's OS support must be disabled before the debug symbols in the `vmlinux` file are loaded. This is to avoid the debugger from trying to read the kernel structures before they are set up (and possibly resulting in data aborts).

Enter `set os enabled off` in the **Command** view and click **Submit** or press **Enter**. This disables OS awareness.

Calculate the offset

The debug symbols in the `vmlinux` file have virtual addresses. So, when the `vmlinux` file is loaded, the debugger assumes that the operating system is up and running with the MMU enabled. But this also can be used to debug pre-MMU at source-level by applying an offset when loading the symbol file.

To calculate the offset, calculate the difference between the physical (P) and virtual (V) addresses of the code. For example, if the kernel is linked at virtual address `0xC0008000` and is loaded at physical address `0x80008000`, the offset is `-0x40000000`, which is `0x80008000 - 0xC0008000` (P-V).

For the DS-5 Linux Distribution example used in this tutorial, the kernel is linked at virtual address `0x80008000` and loaded at physical address `0x80008000`, so the offset is `0x0`.

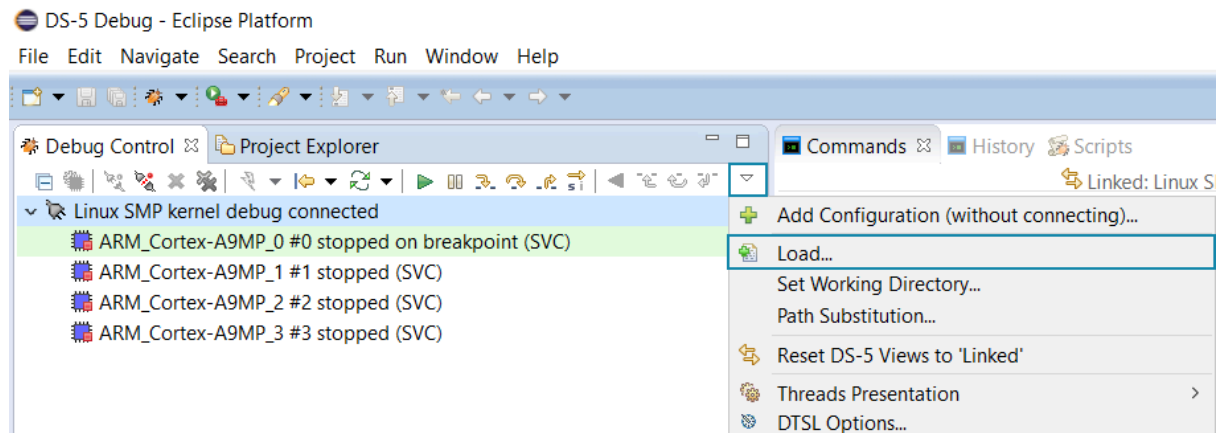
Load the vmlinux file

Now, load the `vmlinux` file with the kernel debug symbols required for this tutorial. The file is available in the Linux distribution example that you downloaded and imported into your workspace at the [start of this tutorial](#).

Loading the `vmlinux` file enables you to perform source-level debugging of the Linux kernel. To load the file using the user interface:

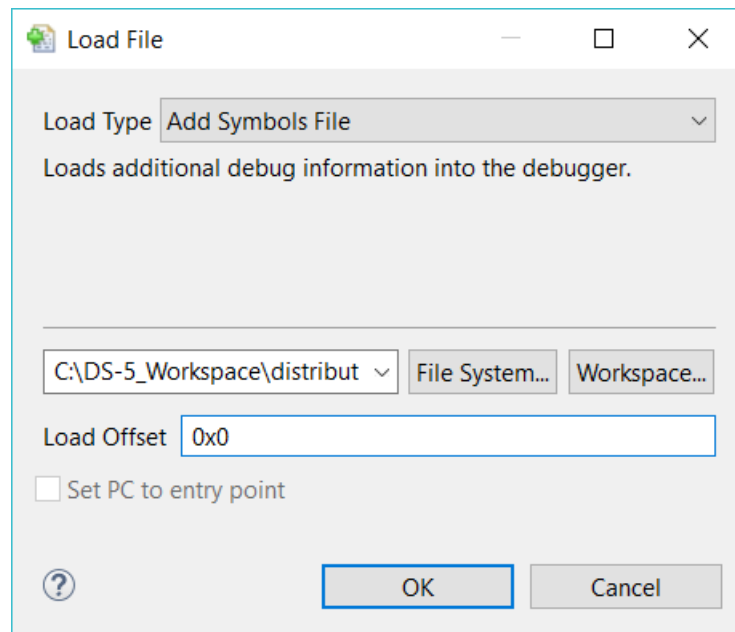
1. In the **Debug Control** view, select the drop-down menu and select [Load].

Figure 4-12: Debug control view



2. In the **Load File** dialog:
 - a. **Load Type** - Select **Add Symbols File**
 - b. **Workspace** - Click and select the `vmlinux` file. You can find this under `<DS-5 Workspace> \distribution\kernel\linux-linaro-3.4-rc3-2012.04-0-patched\built\VE_V7\vmlinux`
 - c. [Load Offset] - Enter `0x0`.

Figure 4-13: Load file window



- d. Click **OK**.

The `head.s` file opens in the **Editor** view.

You might need to single step once to open `head.S` in the editor. Also, if you get any **Source Not Found** alerts, you have to set the substitution path to the sources you set up as part of [Prerequisites](#). The necessary substitution path to the `head.s` source file should be `<extraction_directory>/arch/arm/kernel`.

The **Disassembly** view shows the symbol `stext`, at the entry point for the kernel.

Loading the `vmlinux` file using the command line

You can also load the `vmlinux` file using the command-line. Use the `add-symbol-file` command to add the symbol file.

For this tutorial, enter: `add-symbol-file "<Your Project>\distribution\kernel\linux-linaro-3.4-rc3-2012.04-0-patched\built\VE_V7\vmlinux" 0x0` in the **Commands** view.

Replace `<Your Project>` with the name of the project you created as part of [Prerequisites](#).

You can now, at source level, set breakpoints and watchpoints, view registers, view memory, single step, and other usual debug operations at this pre-MMU stage.

5. Debugging the kernel - post-MMU stage

Now that you have debugged the pre-MMU stage, we can move on to the post-MMU debugging stage.


1. In the **Commands** view, enter `thbreak __turn_mmu_on` to set a temporary hardware breakpoint to see when the MMU is turned on.
2. Click **Continue**  (or F8 on your keyboard).
3. When the breakpoint at `__turn_mmu_on` is reached, note the value of `SP` (Stack Pointer) under the `core` register group in the **Registers** view. It contains the virtual address of `__mmap_switched` and is the place that the code jumps to after the MMU is enabled. The value in this case is `0x804C005C`.

Figure 5-1: Stack Point value

| Core | | 47 of 47 registers | | |
|------|------|--------------------|----|-----|
| | R0 | 0x10C53C7D | 32 | R/W |
| | R1 | 0xFFFFFFFF | 32 | R/W |
| | R2 | 0x80F01310 | 32 | R/W |
| | R3 | 0x8000603C | 32 | R/W |
| | R4 | 0x8000406A | 32 | R/W |
| | R5 | 0x00000015 | 32 | R/W |
| | R6 | 0x10C03C7D | 32 | R/W |
| | R7 | 0x00011C0E | 32 | R/W |
| | R8 | 0x8000406A | 32 | R/W |
| | R9 | 0x413FC090 | 32 | R/W |
| | R10 | 0x00000000 | 32 | R/W |
| | R11 | 0x00000000 | 32 | R/W |
| | R12 | 0x80372850 | 32 | R/W |
| | SP | 0x804C005C | 32 | R/W |
| | LR | 0x80008044 | 32 | R/W |
| | PC | 0x8037AF00 | 32 | R/W |
| | CPSR | 0x400001D3 | 32 | R/W |

You can also view the status of the MMU. In the **Registers** view, expand **CP15 > SCTLR > M**. You can see that it is disabled at this point.

Figure 5-2: MMU status

| | | | | |
|--------|-------------------------|----|-----|--|
| Core | 47 of 47 registers | | | |
| CP15 | 134 of 134 registers | | | |
| System | 28 of 28 registers | | | |
| MIDR | 0x413FC090 | 32 | RO | |
| SCTLR | 0x00C50078 | 32 | R/W | |
| TE | Disabled | 1 | R/W | |
| AFE | Full_access_permissions | 1 | R/W | |
| TRE | Disabled | 1 | R/W | |
| NMFI | Disabled | 1 | RO | |
| EE | 0x0 | 1 | R/W | |
| HA | 0x0 | 1 | R/W | |
| V | Normal_vectors | 1 | R/W | |
| I | Disabled | 1 | R/W | |
| Z | Disabled | 1 | R/W | |
| SW | Disabled | 1 | R/W | |
| C | Disabled | 1 | R/W | |
| A | Disabled | 1 | R/W | |
| M | Disabled | 1 | R/W | |


- At the **Command** prompt, enter `thbreak *$SP` to place a temporary hardware breakpoint on the virtual address of `__mmap_switched`.
- Click **Continue**  to run the code. When the breakpoint at `__mmap_switched` is hit, the MMU is on. You can view the updated status of the MMU in the **Registers** view.

Figure 5-3: Updated status of the MMU

| | | | | |
|--------|-------------------------|----|-----|--|
| Core | 47 of 47 registers | | | |
| CP15 | 134 of 134 registers | | | |
| System | 28 of 28 registers | | | |
| MIDR | 0x413FC090 | 32 | RO | |
| SCTLR | 0x10C53C7D | 32 | R/W | |
| TE | Disabled | 1 | R/W | |
| AFE | Full_access_permissions | 1 | R/W | |
| TRE | Enabled | 1 | R/W | |
| NMFI | Disabled | 1 | RO | |
| EE | 0x0 | 1 | R/W | |
| HA | 0x0 | 1 | R/W | |
| V | High_vectors | 1 | R/W | |
| I | Enabled | 1 | R/W | |
| Z | Enabled | 1 | R/W | |
| SW | Enabled | 1 | R/W | |
| C | Enabled | 1 | R/W | |
| A | Disabled | 1 | R/W | |
| M | Enabled | 1 | R/W | |

- Enter `file` at the command prompt to discard currently loaded symbols.

This is done because the old set of symbols no longer apply, now that the MMU is on.

7. Now, reload the symbol file without the offset applied, so that post-MMU debugging will be done using unadjusted virtual addresses.

At the **Command** prompt, enter: `add-symbol-file "<DS-5 Workspace>\distribution\kernel\linux-linaro-3.4-rc3-2012.04-0-patched\built\VE_V7\vmlinux"` in the **Commands** view.

Replace `<DS-5 Workspace>` with the location of your workspace.

The source code for `__mmap_switched` in `head-common.s` appears in the **Editor**.


You can now set breakpoints, view registers, view memory, single step, and other usual debug operations at this post-MMU stage, all at source level.

6. Viewing the structure of the kernel

The following steps show you how to view the structure of the kernel:



1. After all the architecture-specific setup is done, the main C code entry into the kernel is in the `start_kernel()` function. This is available in the source at: `<DS-5 Workspace>\distribution\kernel\linux-linaro-3.4-rc3-2012.04-0-patched\source\init\main.c`.

Enter `thbreak start_kernel` at the **Command** prompt to set a temporary hardware breakpoint on it.

2. Click **Continue**  to run the code to reach the breakpoint.

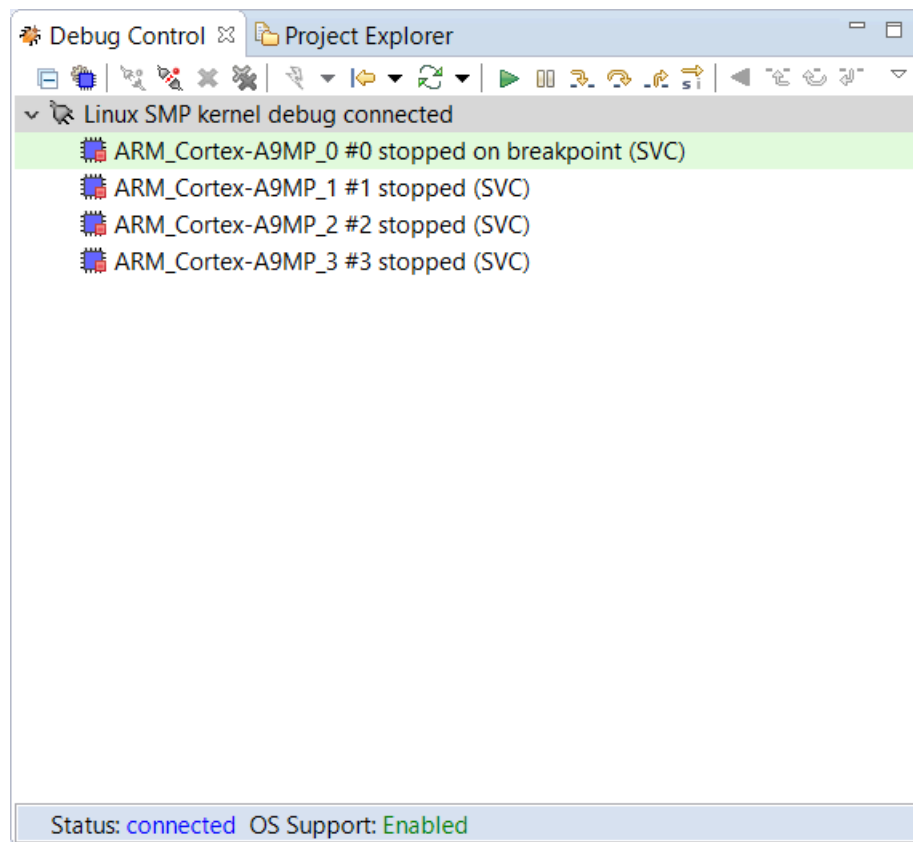
This displays the `main.c` source in the **Editor**.

Again, if you get any **Source Not Found** alerts, you have to set the substitution path to the sources you set up as part of [Prerequisites](#). The necessary substitution path to the `main.c` source file should be `<extraction_directory>/init/`.

3. Enter `set os enabled on` to enable OS support in the DS-5 Debugger.
 - When OS support is enabled, the debugger also displays the Linux kernel support details. For example, you should see something similar to: `Enabled Linux kernel support for version "Linux 3.4.0-rc3-ve-v7-ds5 #1 SMP Mon Oct 31 00:17:24 GMT 2016 arm"`.
 - You can click  in the **Debug Control** view to switch display to show threads. Similarly, you can click  to show cores.

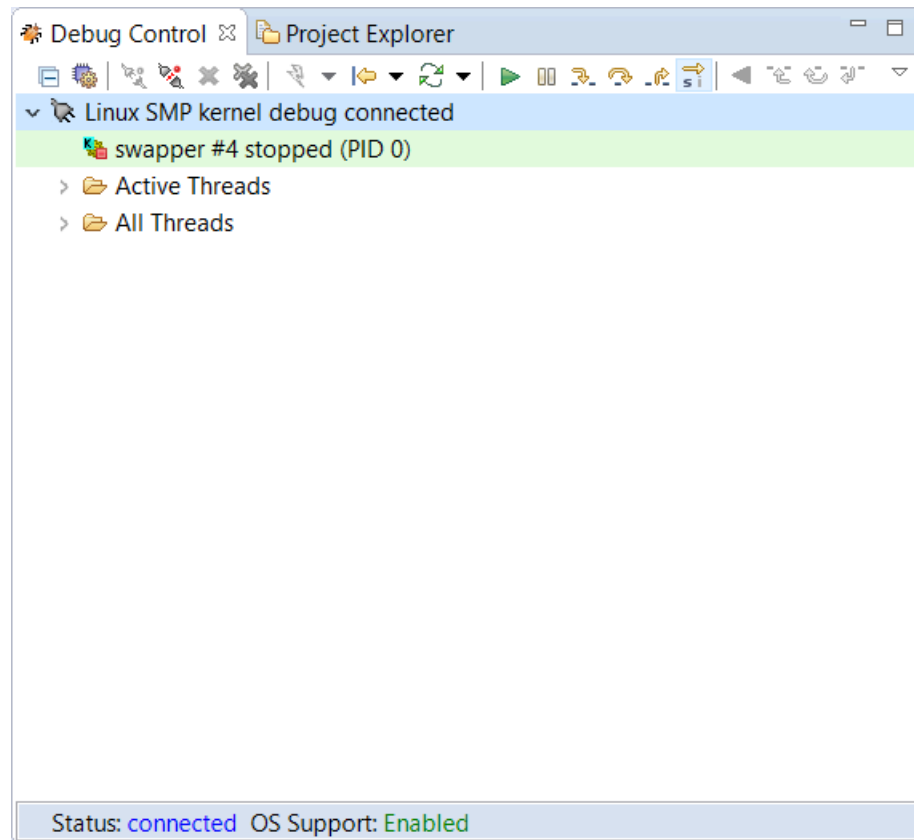
In this example, you can view four cores:

Figure 6-1: Core view



Similarly, when viewing threads, you can view **Active Threads** and **All Threads**:

Figure 6-2: Active and All Threads

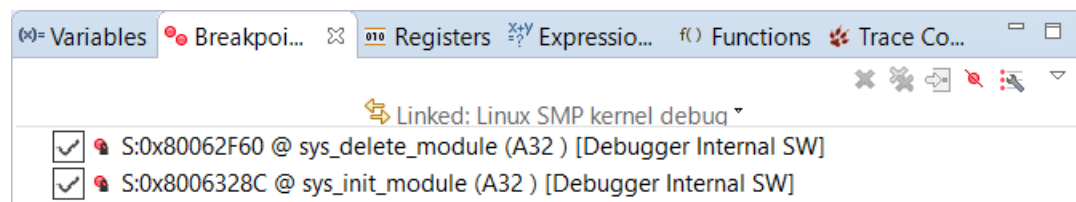


- When OS awareness is enabled and kernel symbols are loaded from the `vmlinux` file, DS-5 Debugger tries to access some locations in the kernel. For example, it tries to read `init_nsproxy.uts_ns->name` to get the kernel name and version. It also tries to set breakpoints automatically on `sys_init_module()` and `sys_delete_module()` to trap when kernel modules are inserted (`insmod`) and removed (`rmmmod`).

The debugger handles internal breakpoints automatically, you cannot enable or disable them, or control them in any way. When internal breakpoints are hit, the debugger automatically performs the corresponding action, and then restarts program execution by itself.

You can see these breakpoints appearing in the **Breakpoints** view.

Figure 6-3: Breakpoints view



4. To confirm the kernel symbols match the kernel image, enter `info os-version` in **Commands** view.

Your operating system version should match the Linux kernel version. For example, operating system version: `Linux 3.4.0-rc3-ve-v7-ds5 #1 SMP Mon Oct 31 00:17:24 GMT 2016 arm`.

This is similar to entering output `init_nsproxy.uts_ns->name`, which should display an output similar to: `{sysname = "Linux", nodename = "(none)", release = "3.4.0-rc3-ve-v7-ds5", version = "#1 SMP Mon Oct 31 00:17:24 GMT 2016", machine = "arm", domainname = "(none)"}`.

This might take a few moments to display because DS-5 Debugger has to process lots of debug symbols.

7. Multi-core features

So far, `cpu_0` performed all the tasks. The next steps explore using secondary cores for debugging.

While `cpu_0` is performing its tasks, the secondary cores are held in a standby state. Conceptually, you can think of the secondary cores being in a holding pen and being released one at a time, under the control of the primary core.

The secondary processors are released from their holding pen when the `platform_smp_prepare_cpus()` function is executed.

1. To observe when the function is executed, set a breakpoint on `platform_smp_prepare_cpus()` function.


At the **Command** prompt, enter `thbreak platform_smp_prepare_cpus` to set a temporary hardware breakpoint.

2. Click **Continue**  to run the code to reach the breakpoint.

The last call in this function writes the address of secondary startup into the system-wide flags register. This write releases the secondary CPUs, which then branch to the address (`&versatile_secondary_startup = 0x804C7480C`).

3. In the **Debug Control** view, notice that `core_0` is in C code with its MMU on, and has started some threads, but `core_1` is still in a holding pen (WFE/LDR/CMP/BNE) assembler loop with its MMU off.

Set a breakpoint on the `versatile_secondary_startup()` function. Enter `thbreak versatile_secondary_startup`.

4. Click **Continue**  to run to the breakpoint.
5. Expand `core_1` in the [Debug Control] view. You can see that it has entered `versatile_secondary_startup()`.

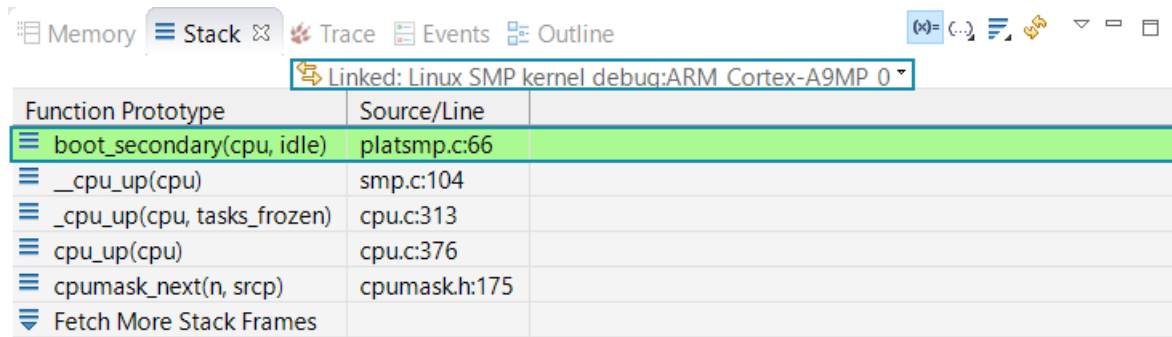
There are two levels of holding pens during an SMP kernel start-up. The first one holds all secondary cores until `core_0` reaches a known good state. Once `core_0` is in a known good state, the secondary cores are all released together. They then enter into a second-level holding pen. They are then released, one at a time by `core_0`.

All cores have their MMUs enabled at this point, so you can explore this second-level holding pen using symbols with virtual addresses by setting breakpoints.

- Enter `thbreak boot_secondary` at the **Command** prompt to set a breakpoint on `boot_secondary` function. This is executed by `core_0`.
- Enter `thbreak platform_secondary_init` at the **Command** prompt to set a breakpoint on the `platform_secondary_init` function. This is executed by `core_1`.

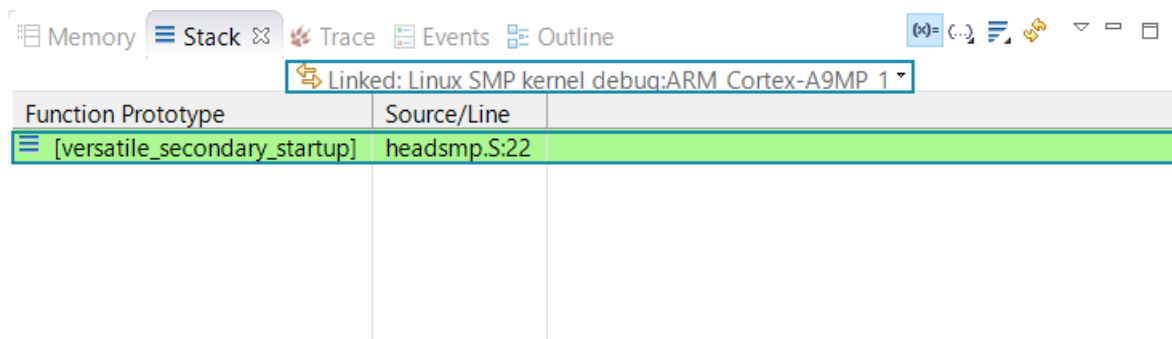
Run to each breakpoint. The `platsmp.c` file opens up in the **Editor**. As you hit each breakpoint, click on `core_0` and switch to the **Stack** window, then `core_1` and switch to the **Stack** window, to see functions under the respective cores.

Figure 7-1: Core_0 stack window



| Function Prototype | Source/Line |
|---|----------------------------|
| <code>boot_secondary(cpu, idle)</code> | <code>platsmp.c:66</code> |
| <code>_cpu_up(cpu)</code> | <code>smp.c:104</code> |
| <code>_cpu_up(cpu, tasks_frozen)</code> | <code>cpu.c:313</code> |
| <code>cpu_up(cpu)</code> | <code>cpu.c:376</code> |
| <code>cpumask_next(n, srcp)</code> | <code>cpumask.h:175</code> |
| Fetch More Stack Frames | |


Figure 7-2: Core_1 stack window



| Function Prototype | Source/Line |
|--|---------------------------|
| <code>[versatile_secondary_startup]</code> | <code>headsmp.S:22</code> |

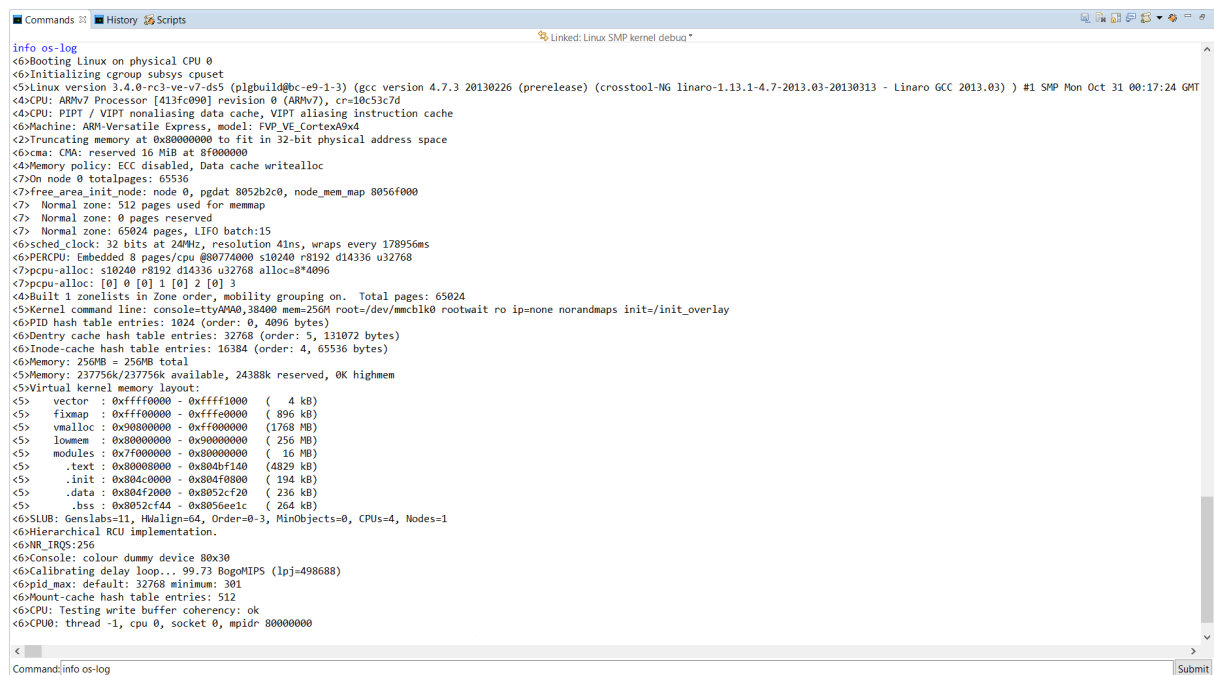
8. Logging features

It would be very useful to display early `printk` outputs in DS-5 Debugger's command window during kernel bring-up. Ideally, before the console is enabled, so no serial output is present.

For example, set a breakpoint on `console_init()` in `main.c`. Click **Continue**  to run to the breakpoint. No log messages are output to the console because it is not enabled yet. But you can view the entire log for activities so far.

- Enter `info os-log` in the **Command** prompt to view the log.

Figure 8-1: info os-log



```

info os-log
<>Booting Linux on physical CPU 0
<>Initializing cgroup subsys cpuset
<>Linux version 3.4.0-rc3-ve-v7-d55 (plgbuild@bc-e9-1-3) (gcc version 4.7.3 20130226 (prerelease) (crosstool-ng linaro-1.13.1-4.7-2013.03-20130313 - Linaro GCC 2013.03) ) #1 SMP Mon Oct 31 00:17:24 GMT
<>CPU: ARMv7 Processor [413fc090] revision 0 (ARMv7), cr=10c53c7d
<>CPU: PIPIT / VIPT nonaliasing data cache, VIPT aliasing instruction cache
<>Machine: ARM-Versatile Express, model: FVP_VE_CortexA9xd
<2>Truncating memory at 0x80000000 to fit in 32-bit physical address space
<>CMA: CMA: reserved 16 MiB at 8f000000
<>Memory policy: ECC disabled, Data cache writealloc
<>On node 0 totalpages: 65536
<>free_area_init_node: node 0, pgdat 8052b2c0, node_mem_map 8056f000
<> Normal zone: 512 pages used for memmap
<> Normal zone: 0 pages reserved
<> Normal zone: 65024 pages, LIFO batch:15
<>sched_clock: 32 bits at 240Hz, resolution 41ns, wraps every 178956ms
<>PERCPU: Embedded 8 pages/cpu @80774000 s10240 r8192 d14336 u32768
<>pcpu-alloc: s10240 r8192 d14336 u32768 alloc=8*4096
<>pcpu-alloc: [0] 0 [0] 1 [0] 2 [0] 3
<>Built 1 zonelists in Zone order, mobility grouping on. Total pages: 65024
<>Kernel command line: console=ttyAMA0,38400 mem=256M root=/dev/mmcblk0 rootwait ro ip=none norandmaps init=/init_overlay
<>PID hash table entries: 1024 (order: 0, 4096 bytes)
<>Dentry cache hash table entries: 32768 (order: 5, 131072 bytes)
<>Inode-cache hash table entries: 16384 (order: 4, 65536 bytes)
<>Memory: 256MB = 256MB total
<>Memory: 237756k/237756k available, 24388k reserved, 0K highmem
<>Virtual kernel memory layout:
<> vector : 0xffff0000 - 0xffff1000 ( 4 kB)
<> fixmap : 0xffff0000 - 0xffffe000 ( 896 kB)
<> vmalloc : 0x90800000 - 0xffff0000 (1768 MB)
<> lowmem : 0x80000000 - 0x90000000 ( 256 MB)
<> modules : 0x7f000000 - 0x80000000 ( 16 MB)
<> .text : 0x80000000 - 0x804bf160 (4829 kB)
<> .init : 0x804c0000 - 0x804f0800 ( 194 kB)
<> .data : 0x804f2000 - 0x8052cf20 ( 236 kB)
<> .bss : 0x8052cf44 - 0x8056ee1c ( 264 kB)
<>SLUB: Genslabs=11, Hwalign=64, Order=0-3, MinObjects=0, CPUs=4, Nodes=1
<>Hierarchical RCU implementation.
<>NR_IRQS:256
<>Console: colour dummy device 80x30
<>Calibrating delay loop... 99.73 BogoMIPS (lpj=498688)
<>pid_max: default: 32768 minimum: 301
<>Mount-cache hash table entries: 512
<>CPU: Testing write buffer coherency: ok
<>CPU0: thread -1, cpu 0, socket 0, mpidr 80000000

```

- If you want to view the log output line by line as the code is processed, enter `set os log-capture on` at the **Command** prompt.

9. Single-stepping individual cores or threads

You can single step through a single core or thread/process. Select either the core or the thread/process in the **Debug Control** view, then press **F5** on your keyboard.

When you single step through a process, it might be migrated to another core. If a breakpoint is set on a process, the debugger can track the migration of process-specific breakpoints to the other core.

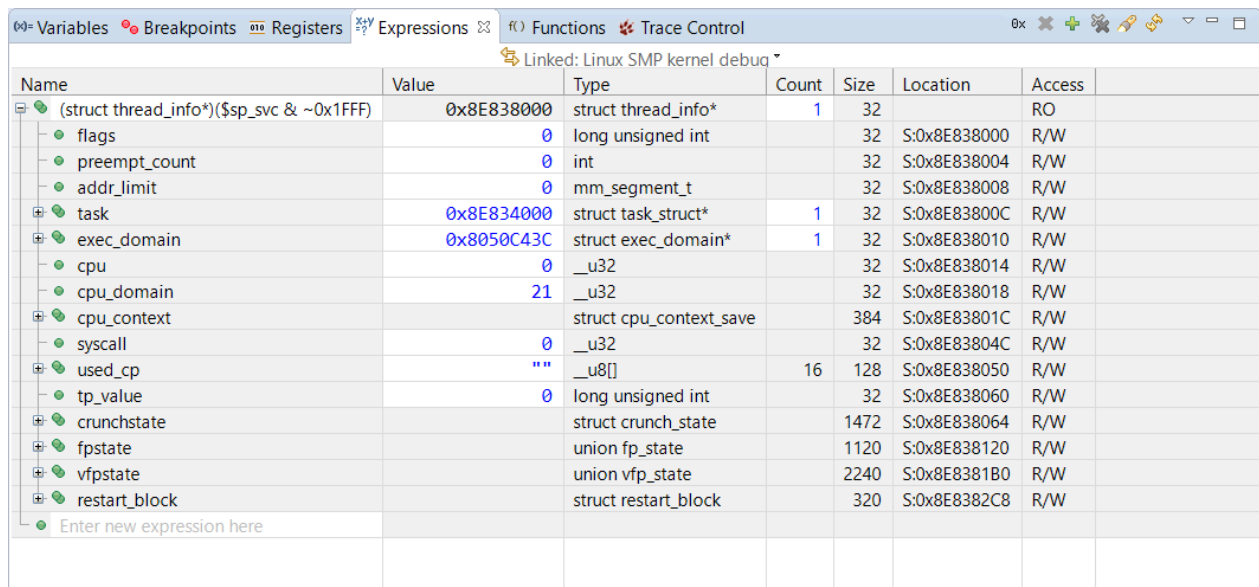
10. Using expressions to view data

You can use the **Expressions** view to create and work with expressions in DS-5. For example, to view the kernel's `thread_info` structure, in the **Expressions** view, enter:

```
(struct thread_info*)($sp_svc & ~0x1FFF).
```

When the view refreshes, expand the tree structure to explore its contents.

Figure 10-1: Expressions view



The screenshot shows the DS-5 Expressions view with the expression `(struct thread_info*)($sp_svc & ~0x1FFF)` entered. The view displays a tree structure of the `thread_info` structure, with fields expanded to show their values and types. The fields include `flags`, `preempt_count`, `addr_limit`, `task`, `exec_domain`, `cpu`, `cpu_domain`, `cpu_context`, `syscall`, `used_cp`, `tp_value`, `crunchstate`, `fpstate`, `vfpstate`, and `restart_block`.

| Name | Value | Type | Count | Size | Location | Access |
|---|------------|-------------------------|-------|------|--------------|--------|
| (struct thread_info*)(\$sp_svc & ~0x1FFF) | 0x8E838000 | struct thread_info* | 1 | 32 | | RO |
| flags | 0 | long unsigned int | | 32 | S:0x8E838000 | R/W |
| preempt_count | 0 | int | | 32 | S:0x8E838004 | R/W |
| addr_limit | 0 | mm_segment_t | | 32 | S:0x8E838008 | R/W |
| task | 0x8E834000 | struct task_struct* | 1 | 32 | S:0x8E83800C | R/W |
| exec_domain | 0x8050C43C | struct exec_domain* | 1 | 32 | S:0x8E838010 | R/W |
| cpu | 0 | _u32 | | 32 | S:0x8E838014 | R/W |
| cpu_domain | 21 | _u32 | | 32 | S:0x8E838018 | R/W |
| cpu_context | | struct cpu_context_save | | 384 | S:0x8E83801C | R/W |
| syscall | 0 | _u32 | | 32 | S:0x8E83804C | R/W |
| used_cp | "" | _u8[] | 16 | 128 | S:0x8E838050 | R/W |
| tp_value | 0 | long unsigned int | | 32 | S:0x8E838060 | R/W |
| crunchstate | | struct crunch_state | | 1472 | S:0x8E838064 | R/W |
| fpstate | | union fp_state | | 1120 | S:0x8E838120 | R/W |
| vfpstate | | union vfp_state | | 2240 | S:0x8E8381B0 | R/W |
| restart_block | | struct restart_block | | 320 | S:0x8E8382C8 | R/W |
| Enter new expression here | | | | | | |

11. Using the command line to view cores, thread, and process information

As you have seen, the **Debug Control** view is very useful to display details of individual cores, threads, and processes. You can also view the same information in the **Commands** view, by using individual commands.

- [info cores](#) - View information about the running processors.
- [info threads](#) - View information about the running threads.
- [info processes](#) - View information about the running processes.