



# Arm<sup>®</sup> DynamIQ<sup>™</sup> Shared Unit

Revision: r4p1

## Technical Reference Manual

### Non-Confidential

Copyright © 2016–2019, 2023 Arm Limited (or its affiliates).  
All rights reserved.

### Issue 04

100453\_0401\_04\_en



# Arm® DynamIQ™ Shared Unit

## Technical Reference Manual

Copyright © 2016–2019, 2023 Arm Limited (or its affiliates). All rights reserved.

## Release Information

### Document history

| Issue   | Date              | Confidentiality  | Change                  |
|---------|-------------------|------------------|-------------------------|
| 0000-00 | 30 September 2016 | Confidential     | First release for r0p0  |
| 0001-00 | 16 December 2016  | Confidential     | First release for r0p1  |
| 0002-00 | 20 June 2017      | Non-Confidential | First release for r0p2  |
| 0100-00 | 31 May 2017       | Confidential     | First release for r1p0  |
| 0200-00 | 16 October 2017   | Confidential     | First release for r2p0  |
| 0300-00 | 26 January 2018   | Non-Confidential | First release for r3p0  |
| 0300-01 | 27 April 2018     | Non-Confidential | Second release for r3p0 |
| 0400-00 | 27 November 2018  | Non-Confidential | First release for r4p0  |
| 0400-01 | 7 December 2018   | Non-Confidential | Second release for r4p0 |
| 0400-02 | 14 May 2019       | Non-Confidential | Third release for r4p0  |
| 0401-03 | 31 October 2019   | Non-Confidential | First release for r4p1  |
| 0401-04 | 30 June 2023      | Non-Confidential | Second release for r4p1 |

## Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its affiliates) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <https://www.arm.com/company/policies/trademarks>.

Copyright © 2016–2019, 2023 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.

110 Fulbourn Road, Cambridge, England CB1 9NJ.

(LES-PRE-20349|version 21.0)

## Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Unrestricted Access is an Arm internal classification.

## Product Status

The information in this document is Final, that is for a developed product.

## Feedback

Arm welcomes feedback on this product and its documentation. To provide feedback on the product, create a ticket on <https://support.developer.arm.com>.

To provide feedback on the document, fill the following survey: <https://developer.arm.com/documentation-feedback-survey>.

## Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used language that can be offensive. Arm strives to lead the industry and create change.

This document includes language that can be offensive. We will replace this language in a future issue of this document.

To report offensive language in this document, email [terms@arm.com](mailto:terms@arm.com).

# Contents

|  |           |
|--|-----------|
| <b>1. Introduction.....</b>                      | <b>12</b> |
| 1.1 Product revision status.....                 | 12        |
| 1.2 Intended audience.....                       | 12        |
| 1.3 Conventions.....                             | 12        |
| 1.4 Useful resources.....                        | 14        |
| <b>2. The DynamIQ Shared Unit.....</b>           | <b>16</b> |
| 2.1 About the DSU.....                           | 16        |
| 2.2 Features.....                                | 17        |
| 2.3 Implementation options.....                  | 18        |
| 2.3.1 Cluster configurations.....                | 22        |
| 2.4 Supported standards and specifications.....  | 22        |
| 2.5 Test features.....                           | 23        |
| 2.6 Design tasks.....                            | 23        |
| 2.7 Product revisions.....                       | 24        |
| <b>3. Technical overview.....</b>                | <b>25</b> |
| 3.1 Components.....                              | 25        |
| 3.2 Interfaces.....                              | 28        |
| 3.3 RAS support.....                             | 31        |
| 3.4 Page-based hardware attributes.....          | 31        |
| 3.5 L3 memory system variants.....               | 32        |
| <b>4. Clocks and resets.....</b>                 | <b>34</b> |
| 4.1 Clocks.....                                  | 34        |
| 4.2 Resets.....                                  | 35        |
| <b>5. Power management.....</b>                  | <b>37</b> |
| 5.1 About DSU power management.....              | 37        |
| 5.2 Power mode control.....                      | 38        |
| 5.3 Communication with the power controller..... | 39        |
| 5.4 L3 RAM power control.....                    | 39        |
| 5.4.1 L3 cache partial powerdown.....            | 39        |

|   |           |
|---|-----------|
| 5.4.2 L3 RAM retention.....                           | 40        |
| 5.5 Power modes.....                                  | 41        |
| 5.5.1 Power mode transitions.....                     | 41        |
| 5.5.2 Power mode transition behavior.....             | 46        |
| 5.5.3 Interlocks between core and DSU P-Channels..... | 47        |
| 5.5.4 Power mode encoding.....                        | 47        |
| 5.6 Power operating requirements.....                 | 48        |
| 5.6.1 Power control for DFT.....                      | 49        |
| 5.7 Wait For Interrupt and Wait For Event.....        | 50        |
| 5.8 Clock, voltage, and power domains.....            | 51        |
| 5.9 Cluster powerdown.....                            | 52        |
| 5.9.1 Transitioning in and out of coherency.....      | 52        |
| <b>6. L3 cache.....</b>                               | <b>54</b> |
| 6.1 About the L3 cache.....                           | 54        |
| 6.2 L3 cache allocation policy.....                   | 54        |
| 6.3 L3 cache partitioning.....                        | 55        |
| 6.4 Cache stashing.....                               | 56        |
| 6.5 L3 cache ECC and parity.....                      | 57        |
| 6.6 L3 cache data RAM latency.....                    | 60        |
| 6.7 Cache slices and portions.....                    | 62        |
| 6.7.1 Cache slice and master port selection.....      | 63        |
| 6.7.2 Default number of cache slices.....             | 63        |
| 6.7.3 Implementing a 1.5MB or 3MB L3 cache.....       | 64        |
| <b>7. ACE master interface.....</b>                   | <b>66</b> |
| 7.1 About the ACE master interface.....               | 66        |
| 7.1.1 Dual ACE interfaces.....                        | 66        |
| 7.2 ACE configurations.....                           | 67        |
| 7.3 ACE features.....                                 | 68        |
| 7.4 ACE master interface attributes.....              | 69        |
| 7.5 ACE channel properties.....                       | 72        |
| 7.6 ACE transactions.....                             | 73        |
| 7.7 Support for memory types.....                     | 75        |
| 7.8 Read response.....                                | 75        |
| 7.9 Write response.....                               | 76        |
| 7.10 Barriers.....                                    | 76        |

|  |           |
|--|-----------|
| 7.11 AXI compatibility mode.....                           | 76        |
| 7.11.1 Additional logic to support AXI compatibility.....  | 77        |
| 7.12 ACE privilege information.....                        | 77        |
| <b>8. CHI master interface.....</b>                        | <b>78</b> |
| 8.1 About the CHI master interface.....                    | 78        |
| 8.1.1 Dual CHI interfaces.....                             | 78        |
| 8.2 CHI version.....                                       | 79        |
| 8.3 CHI features.....                                      | 79        |
| 8.4 CHI configurations.....                                | 80        |
| 8.5 Attributes of the CHI master interface.....            | 81        |
| 8.6 CHI channel properties.....                            | 82        |
| 8.7 CHI transactions.....                                  | 83        |
| 8.8 Use of DataSource.....                                 | 86        |
| 8.9 Support for memory types.....                          | 87        |
| <b>9. ACP slave interface.....</b>                         | <b>88</b> |
| 9.1 About the ACP.....                                     | 88        |
| 9.2 ACP features.....                                      | 88        |
| 9.3 ACP ACE5-Lite subset.....                              | 89        |
| 9.4 ACP transaction types.....                             | 89        |
| 9.5 ACP performance.....                                   | 91        |
| <b>10. AXI master peripheral port.....</b>                 | <b>93</b> |
| 10.1 About the peripheral port.....                        | 93        |
| 10.2 Transaction ID encoding.....                          | 93        |
| <b>11. Control registers.....</b>                          | <b>95</b> |
| 11.1 About the control registers.....                      | 95        |
| 11.2 AArch32 control register summary.....                 | 95        |
| 11.3 AArch64 control register summary.....                 | 96        |
| 11.4 CLUSTERACPSID, Cluster ACP Scheme ID Register.....    | 97        |
| 11.5 CLUSTERACTLR, Cluster Auxiliary Control Register..... | 99        |
| 11.6 CLUSTERBUSQOS, Cluster Bus QoS Control Register.....  | 99        |
| 11.7 CLUSTERCFR, Cluster Configuration Register.....       | 101       |
| 11.8 CLUSTERECTLR, Cluster Extended Control Register.....  | 105       |
| 11.9 CLUSTERIDR, Cluster Main Revision ID Register.....    | 109       |

|  |            |
|--|------------|
| 11.10 CLUSTERL3HIT, Cluster L3 Hit Counter Register.....                   | 111        |
| 11.11 CLUSTERL3MISS, Cluster L3 Miss Counter Register.....                 | 113        |
| 11.12 CLUSTERPARTCR, Cluster Partition Control Register.....               | 115        |
| 11.13 CLUSTERPWRCTLR, Cluster Power Control Register.....                  | 119        |
| 11.14 CLUSTERPWRDN, Cluster Powerdown Register.....                        | 121        |
| 11.15 CLUSTERPWRSTAT, Cluster Power Status Register.....                   | 123        |
| 11.16 CLUSTERREVIDR, Cluster Revision ID Register.....                     | 125        |
| 11.17 CLUSTERSTASHSID, Cluster Stash Scheme ID Register.....               | 127        |
| 11.18 CLUSTERTHREADSID, Cluster Thread Scheme ID Register.....             | 129        |
| 11.19 CLUSTERTHREADSIDOVR, Cluster Thread Scheme ID Override Register..... | 130        |
| <b>12. Error system registers.....</b>                                     | <b>133</b> |
| 12.1 About the error system registers.....                                 | 133        |
| 12.2 Error system register summary.....                                    | 133        |
| 12.3 ERR1CTLR, Error Record Control Register.....                          | 134        |
| 12.4 ERR1FR, Error Record Feature Register.....                            | 136        |
| 12.5 ERR1MISC0, Error Record Miscellaneous Register 0.....                 | 138        |
| 12.6 ERR1MISC1, Error Record Miscellaneous Register 1.....                 | 140        |
| 12.7 ERR1PFGCDNR, Error Pseudo Fault Generation Count Down Register.....   | 141        |
| 12.8 ERR1PFGCTLR, Error Pseudo Fault Generation Control Register.....      | 141        |
| 12.9 ERR1PFGFR, Error Pseudo Fault Generation Feature Register.....        | 143        |
| 12.10 ERR1STATUS, Error Record Primary Status Register.....                | 145        |
| <b>13. PMU registers.....</b>  | <b>149</b> |
| 13.1 About the PMU registers.....  | 149        |
| 13.2 AArch32 PMU register summary.....                                     | 149        |
| 13.3 AArch64 PMU register summary.....                                     | 150        |
| 13.4 CLUSTERPMCR, Cluster Performance Monitors Control Register.....       | 151        |
| 13.5 CLUSTERPMCNTENSET, Cluster Count Enable Set Register.....             | 154        |
| 13.6 CLUSTERPMCNTENCLR, Cluster Count Enable Clear Register.....           | 156        |
| 13.7 CLUSTERPMOVSSSET, Cluster Overflow Flag Status Set Register.....      | 158        |
| 13.8 CLUSTERPMOVSCCLR, Cluster Overflow Flag Status Clear Register.....    | 160        |
| 13.9 CLUSTERPMSELR, Cluster Event Counter Selection Register.....          | 162        |
| 13.10 CLUSTERPMINTENSET, Cluster Interrupt Enable Set Register.....        | 164        |
| 13.11 CLUSTERPMINTENCLR, Cluster Interrupt Enable Clear Register.....      | 166        |
| 13.12 CLUSTERPMCCNTR, Cluster Performance Monitors Cycle Counter.....      | 168        |
| 13.13 CLUSTERPMXEVTYPER, Cluster Selected Event Type Register.....         | 170        |



|   |            |
|---|------------|
| 13.14 CLUSTERPMXEVCNTR, Cluster Selected Event Counter Register.....      | 171        |
| 13.15 CLUSTERPMMDCR, Cluster Monitor Debug Configuration Register.....    | 173        |
| 13.16 CLUSTERPMCEID0, Cluster Common Event Identification Register 0..... | 175        |
| 13.17 CLUSTERPMCEID1, Cluster Common Event Identification Register 1..... | 177        |
| 13.18 CLUSTERCLAIMSET, Cluster Claim Tag Set Register.....                | 179        |
| 13.19 CLUSTERCLAIMCLR, Cluster Claim Tag Clear Register.....              | 181        |
| 13.20 CLUSTERPMEVTYPEPER<n>, Cluster Event Type Register.....             | 183        |
| 13.21 CLUSTERPMEVCNTR<n>, Cluster Event Counter Register.....             | 184        |
| <b>14. Debug.....</b>   | <b>185</b> |
| 14.1 About debug methods.....   | 185        |
| 14.2 Terminology.....   | 186        |
| 14.3 About the DebugBlock.....  | 186        |
| 14.4 DebugBlock components.....   | 188        |
| 14.5 About the Embedded Cross Trigger.....                                | 189        |
| 14.5.1 Supported debug and trace trigger events.....                      | 190        |
| 14.6 CTI triggers.....  | 191        |
| <b>15. Performance Monitoring Unit.....</b>                               | <b>193</b> |
| 15.1 About the PMU.....   | 193        |
| 15.2 PMU functional description.....                                      | 193        |
| 15.3 PMU events.....  | 194        |
| 15.4 PMU interrupts.....  | 196        |
| <b>16. Debug registers.....</b>   | <b>197</b> |
| 16.1 Debug memory map.....  | 197        |
| 16.2 CTI register summary.....  | 204        |
| 16.3 CTIPIDR0, CTI Peripheral Identification Register 0.....              | 205        |
| 16.4 CTIPIDR1, CTI Peripheral Identification Register 1.....              | 207        |
| 16.5 CTIPIDR2, CTI Peripheral Identification Register 2.....              | 208        |
| 16.6 CTIPIDR3, CTI Peripheral Identification Register 3.....              | 209        |
| 16.7 CTIPIDR4, CTI Peripheral Identification Register 4.....              | 210        |
| 16.8 CTIITCTRL, CTI Integration Mode Control Register.....                | 211        |
| 16.9 CTIDEVAFF0, Cluster CTI Device Affinity register 0.....              | 212        |
| 16.10 CTIDEVID, CTI Device Identification Register.....                   | 213        |
| 16.11 External register access permissions.....                           | 215        |

|  |            |
|--|------------|
| <b>17. ROM table.....</b>                              | <b>216</b> |
| 17.1 About the ROM table.....                          | 216        |
| 17.2 ROM table register summary.....                   | 216        |
| <b>A Appendixes.....</b>                               | <b>219</b> |
| <b>A. Compatible Core Versions.....</b>                | <b>220</b> |
| A.1 Compatible Core Versions.....                      | 220        |
| <b>B. Signal descriptions.....</b>                     | <b>221</b> |
| B.1 Signal naming convention.....                      | 221        |
| B.2 DSU signals.....                                   | 222        |
| B.2.1 Clock and clock enable signals.....              | 222        |
| B.2.2 Reset signals.....                               | 223        |
| B.2.3 Configuration signals.....                       | 223        |
| B.2.4 GIC signals.....                                 | 224        |
| B.2.5 Generic Timer signals.....                       | 225        |
| B.2.6 Power management signals.....                    | 225        |
| B.2.7 Error signals.....                               | 227        |
| B.2.8 ACP interface signals.....                       | 228        |
| B.2.9 Peripheral port interface signals.....           | 229        |
| B.2.10 Broadcast signals for the memory interface..... | 231        |
| B.2.11 ACE interface signals.....                      | 231        |
| B.2.12 CHI interface signals.....                      | 235        |
| B.2.13 DebugBlock APB interface signals.....           | 239        |
| B.2.14 ATB interface signals.....                      | 240        |
| B.2.15 Timestamp signal.....                           | 241        |
| B.2.16 PMU interface signals.....                      | 241        |
| B.2.17 ELA signal.....                                 | 241        |
| B.2.18 DFT interface signals.....                      | 242        |
| B.2.19 MBIST interface signals.....                    | 242        |
| B.3 DebugBlock signals.....                            | 242        |
| B.3.1 Clock signal.....                                | 242        |
| B.3.2 Reset signal.....                                | 243        |
| B.3.3 Power and clock gate control signals.....        | 243        |
| B.3.4 Configuration signals.....                       | 243        |
| B.3.5 Debug signals.....                               | 244        |

B.3.6 CTI interface signals..... 246

B.3.7 DFT signals..... 246

**C. Revisions..... 247**

C.1 Revisions..... 247

# 1. Introduction

## 1.1 Product revision status

The  $r_xp_y$  identifier indicates the revision status of the product described in this manual, for example,  $r1p2$ , where:

|                         |  |
|-------------------------|--|
| <b><math>r_x</math></b> | Identifies the major revision of the product, for example, $r1$ .                        |
| <b><math>p_y</math></b> | Identifies the minor revision or modification status of the product, for example, $p2$ . |

## 1.2 Intended audience

This manual is written for system designers, system integrators, and programmers who are designing or programming a *System-on-Chip* (SoC) that uses the DSU.

## 1.3 Conventions

The following subsections describe conventions used in Arm documents.

### Glossary

The Arm® Glossary is a list of terms used in Arm documentation, together with definitions for those terms. The Arm Glossary does not contain terms that are industry standard unless the Arm meaning differs from the generally accepted meaning.

See the Arm Glossary for more information: [developer.arm.com/glossary](https://developer.arm.com/glossary).

| Convention                 | Use   |
|----------------------------|---|
| <i>italic</i>              | Citations.  |
| <b>bold</b>                | Terms in descriptive lists, where appropriate.  |
| monospace                  | Text that you can enter at the keyboard, such as commands, file and program names, and source code.   |
| monospace <u>underline</u> | A permitted abbreviation for a command or option. You can enter the underlined text instead of the full command or option name.   |
| <and>                      | Encloses replaceable terms for assembler syntax where they appear in code or code fragments.<br><br>For example:<br><br><pre>MRC p15, 0, &lt;Rd&gt;, &lt;CRn&gt;, &lt;CRm&gt;, &lt;Opcode_2&gt;</pre> |

| Convention            | Use  |
|-----------------------|--|
| <b>SMALL CAPITALS</b> | Terms that have specific technical meanings as defined in the <i>Arm® Glossary</i> . For example, <b>IMPLEMENTATION DEFINED</b> , <b>IMPLEMENTATION SPECIFIC</b> , <b>UNKNOWN</b> , and <b>UNPREDICTABLE</b> . |



Recommendations. Not following these recommendations might lead to system failure or damage.



Requirements for the system. Not following these requirements might result in system failure or damage.



Requirements for the system. Not following these requirements will result in system failure or damage.



An important piece of information that needs your attention.



A useful tip that might make it easier, better or faster to perform a task.




A reminder of something important that relates to the information you are reading.

## Timing diagrams

The following figure explains the components used in timing diagrams. Variations, when they occur, have clear labels. You must not assume any timing information that is not explicit in the diagrams.

Shaded bus and signal areas are undefined, so the bus or signal can assume any value within the shaded area at that time. The actual level is unimportant and does not affect normal operation.

**Figure 1-1: Key to timing diagram conventions**

## Signals

The signal conventions are:

### Signal level

The level of an asserted signal depends on whether the signal is active-HIGH or active-LOW. Asserted means:

- HIGH for active-HIGH signals.
- LOW for active-LOW signals.

### Lowercase n

At the start or end of a signal name, n denotes an active-LOW signal.

## 1.4 Useful resources

This document contains information that is specific to this product. See the following resources for other useful information.

Access to Arm documents depends on their confidentiality:

- Non-Confidential documents are available at [developer.arm.com/documentation](https://developer.arm.com/documentation). Each document link in the following tables goes to the online version of the document.
- Confidential documents are available to licensees only through the product package.

| Arm product resources                                      | Document ID    | Confidentiality |
|--|----------------|-----------------|
| Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide | 100454         | Confidential    |
| Arm® DynamIQ™ Shared Unit Integration Manual               | 100455         | Confidential    |
| Arm® DynamIQ™ Shared Unit Release Note                     | ARM-EPM-xxxxxx | Confidential    |

| Arm architecture and specifications    | Document ID | Confidentiality  |
|--|-------------|------------------|
| AMBA® 5 CHI Architecture Specification | IHI 0050    | Non-Confidential |

| Arm architecture and specifications  | Document ID | Confidentiality  |
|--|-------------|------------------|
| <a href="#">AMBA® APB Protocol Specification</a>   | IHI 0024    | Non-Confidential |
| <a href="#">AMBA® ATB Protocol Specification</a>   | IHI 0032    | Non-Confidential |
| <a href="#">AMBA® AXI and ACE Protocol Specification</a>   | IHI 0022    | Non-Confidential |
| <a href="#">AMBA® Low Power Interface Specification</a>  | IHI 0068    | Non-Confidential |
| <a href="#">Arm® Architecture Reference Manual for A-profile architecture</a>  | DDI 0487    | Non-Confidential |
| <a href="#">Arm® CoreSight™ Architecture Specification v3.0</a>  | IHI 0029    | Non-Confidential |
| <a href="#">Arm® CoreSight™ DAP-Lite2 Technical Reference Manual</a>   | 100572      | Non-Confidential |
| <a href="#">Arm® CoreSight™ ELA-500 Embedded Logic Analyzer Technical Reference Manual</a>                             | 100127      | Non-Confidential |
| <a href="#">Arm® CoreSight™ SoC-400 Technical Reference Manual</a>   | DDI 0480    | Non-Confidential |
| <a href="#">Arm® Embedded Trace Macrocell Architecture Specification ETMv4</a>   | IHI 0064    | Non-Confidential |
| <a href="#">Arm® Generic Interrupt Controller Architecture Specification, GIC architecture version 3 and version 4</a> | IHI 0069    | Non-Confidential |



Arm tests its PDFs only in Adobe Acrobat and Acrobat Reader. Arm cannot guarantee the quality of its documents when used with any other PDF reader.

Adobe PDF reader products can be downloaded at <http://www.adobe.com>.

## 2. The DynamIQ Shared Unit

This chapter introduces the *DynamIQ™ Shared Unit* (DSU) and its features.

### 2.1 About the DSU

The *DynamIQ™ Shared Unit* (DSU) comprises the L3 memory system, control logic, and external interfaces to support a DynamIQ™ cluster.

The DynamIQ™ cluster microarchitecture integrates one or more cores with the DSU to form a cluster that is implemented to a specified configuration.

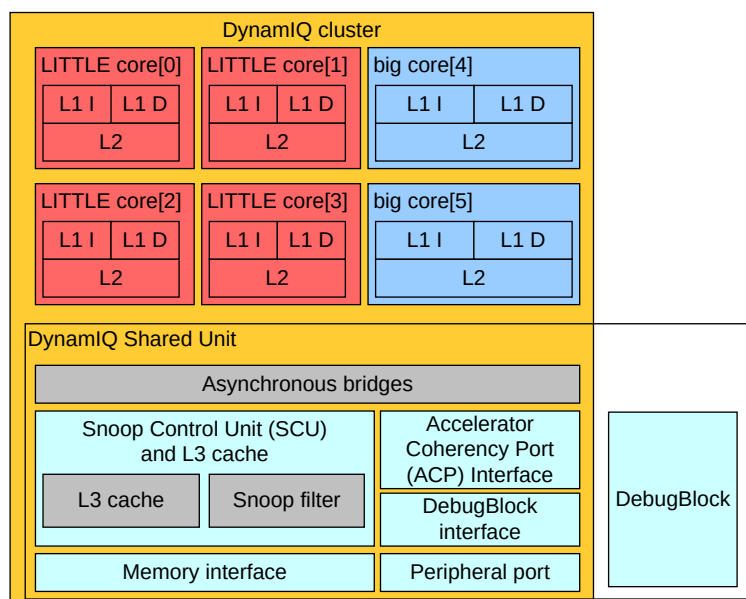
The cores are selected and configured during macrocell implementation.

A cluster can be implemented in one of three configurations:

- A set of cores having the same microarchitecture.
- Two sets of cores, where each set has a different microarchitecture. This configuration might be used in a DynamIQ™ big.LITTLE™ cluster.
- Three sets of cores, where each set has a different microarchitecture.

The following diagram shows a cluster that is composed of two sets of cores in a DynamIQ™ big.LITTLE™ configuration.

**Figure 2-1: DynamIQ™ cluster**





Within the DSU there is the L3 cache, the *Snoop Control Unit* (SCU), internal interfaces to the cores, and external interfaces to the SoC:

- The shared L3 cache simplifies process migration between the cores.



Some cores can be configured without L2 caches. To these cores, the shared L3 cache appears as an L2 cache. The term 'L3 cache' is used throughout this document to describe the shared cache.

- The *Snoop Control Unit* (SCU) maintains coherency between caches in the cores and L3. The SCU includes a snoop filter to optimize coherency maintenance operations.
- Internal interfaces to the cores are configured during macrocell implementation and are not directly visible.
- External interfaces are connected to the SoC.

Each core can be configured either to be run synchronously with the DSU, sharing the clock, or asynchronously, with an independent clock.

Microarchitecture features and system control registers that are specific to the implemented cores are described in separate *Technical Reference Manual* (TRMs) delivered with the cores.

A DebugBlock is provided with the DSU that integrates an Embedded Cross Trigger with debugging registers and supports debug over powerdown. The DebugBlock includes all functionality that is required in the debug power domain.

## 2.2 Features

The DSU includes the following features:

- AMBA ACE5 or AMBA 5 CHI main bus interface.
- Optional 128-bit wide I/O-coherent *Accelerator Coherency Port* (ACP).
- Optional 64-bit wide device peripheral port.
- Support for cores with 40-bit, 44-bit, or 48-bit physical addresses.
- Arm®v8.2-A debug logic or (for compatible cores) Arm®v8.4-A debug logic.
- *Reliability, Availability, and Serviceability* (RAS) support.
- Optional unified 16-way set-associative L3 cache.
- 64-byte cache lines throughout.
- Cache partitioning support.
- Partial L3 cache powerdown support.
- Optional cache protection in the form of *Error Correcting Code* (ECC) on L3 cache RAM instances.

- *Snoop Control Unit (SCU).*
- L3 memory system can be clocked at a rate synchronous to the external system interconnect or at integer multiples.
- Arm®v8.2-A or Arm®v8.4-A architecture cores supported.
- Support for three types of core.
- Cores can be clocked at different frequencies.

## 2.3 Implementation options

The DynamIQ™ Shared Unit can be implemented from a range of options. These options are specified during macrocell implementation.

**Table 2-1: Configuration parameters**

| Parameter name     | Permitted values                   | Description   |
|--------------------|------------------------------------|---|
| NUM_BIG_CORES      | 0, 1, 2, 3, or 4 cores             | The number of big cores that can be implemented depends on the number of LITTLE cores implemented.  |
| NUM_LITTLE_CORES   | 0, 1, 2, 3, 4, 5, 6, 7, or 8 cores | The number of LITTLE cores that can be implemented depends on the number of big cores implemented.  |
| NUM_OTHER_CORES    | 0, 1, 2, or 3 cores                | The number of 'other' cores that can be implemented depends on the number of big and LITTLE cores implemented.  |
| BIG_CORE_TYPE      | -                                  | The type of big core you can implement depends on your license.   |
| LITTLE_CORE_TYPE   | -                                  | The type of LITTLE core you can implement depends on your license.  |
| OTHER_CORE_TYPE    | -                                  | The type of 'other' core you can implement depends on your license.<br><b>Note:</b><br>This OTHER_CORE_TYPE is a different big core from the BIG_CORE_TYPE.                                   |
| MODULE             | -                                  | Name of the cluster top-level Verilog file. If you do not specify any top-level module name, then a name is autogenerated based on the number and types of cores present.                     |
| MODULE_DEBUG_BLOCK | -                                  | Name of the DebugBlock top-level Verilog file. If you do not specify any DebugBlock module name, then a name is autogenerated based on the top-level module name.                             |
| ACE                | TRUE,<br>FALSE                     | Main memory interface.<br><br><b>TRUE</b> Implement AMBA ACE5.<br><b>FALSE</b> Implement AMBA 5 CHI.  |
| MASTER_DATA_WIDTH  | 128, 256, or 512                   | Bus width for the main coherent master interface.<br><br><b>128</b> 1 x 128-bit wide ACE or CHI<br><b>256</b> 2 x 128-bit wide ACE, or 1x 256-bit wide CHI<br><b>512</b> 2 x 256-bit wide CHI |
| PORTER_SAM         | TRUE,<br>FALSE                     | System Address Map (SAM). Only applicable if the CHI interface is implemented.<br><br><b>TRUE</b> Include support for the CMN-600 interconnect SAM.<br><b>FALSE</b> Do not include.           |

| Parameter name       | Permitted values   | Description  |
|----------------------|--|--|
| ACP                  | TRUE,<br>FALSE   | <p><i>Accelerator Coherency Port (ACP)</i></p> <p><b>TRUE</b> Include ACP.<br/><b>FALSE</b> Do not include ACP.</p>  |
| PERIPH_PORT          | TRUE,<br>FALSE   | <p>Peripheral port</p> <p><b>TRUE</b> Include Peripheral port.<br/><b>FALSE</b> Do not include Peripheral port.</p>  |
| SCU_CACHE_PROTECTION | TRUE,<br>FALSE   | <p><i>Error Correcting Code (ECC) support</i></p> <p><b>TRUE</b> Protect the L3 cache RAMs and snoop filter RAMs with ECC.<br/><b>FALSE</b> No ECC on L3 cache RAMs and snoop filter RAMs.</p> <p><b>Note:</b><br/>The DSU ECC protection implementation is configured in conjunction with the cache protection parameter of your core (CORE_CACHE_PROTECTION). See the documentation for your core for more information.</p> <p>The valid combinations are:</p> <ul style="list-style-type: none"> <li>CORE_CACHE_PROTECTION = FALSE, SCU_CACHE_PROTECTION = FALSE.</li> <li>CORE_CACHE_PROTECTION = FALSE, SCU_CACHE_PROTECTION = TRUE.</li> <li>CORE_CACHE_PROTECTION = TRUE, SCU_CACHE_PROTECTION = TRUE.</li> </ul> |
| L3_CACHE             | TRUE,<br>FALSE   | <p>L3 cache present</p> <p><b>TRUE</b> Include L3 cache.<br/><b>FALSE</b> Do not include L3 cache.</p>   |
| L3_CACHE_SIZE        | 256KB<br><br>512KB<br><br>1024KB<br><br>1536KB<br><br>2048KB<br><br>3072KB<br><br>4096KB | <p>L3 cache size</p> <p><b>Note:</b></p> <ul style="list-style-type: none"> <li>Only some cores support 256KB.</li> <li>Software reports a cache size of: 2048KB when the 1536KB size is enabled and 4096KB when the 3072KB size is enabled. Since the non-power-of-two cache size is implemented as a power-of-two size with a quarter of the ways that are permanently powered down.</li> </ul>  |

| Parameter name      | Permitted values | Description   |
|---------------------|------------------|---|
| NUM_SLICES          | 1 or 2           | <p>Optional configuration for the number of L3 cache slices. See <a href="#">6.7 Cache slices and portions</a> on page 62 for more details on cache slices.</p> <p>If this parameter is not provided, then the default number of slices is used, based on the other configuration options.</p> <p>This parameter may be used to override the default number of slices if all of the following are true:</p> <ul style="list-style-type: none"> <li>The L3 cache size is configured as 256KB, 512KB or 1024KB, or no L3 cache is configured.</li> <li>There is only a single master port.</li> </ul> <p>If a larger cache size or two master ports are configured, then only 2 slices are supported.</p> |
| L3_DATA_WR_LATENCY  | 1, 2, or 2p      | <p>L3 cache data RAM input latency:</p> <ul style="list-style-type: none"> <li>1 cycle latency.</li> <li>2 cycles latency.</li> <li>2p, 2 cycle latency with an extra 1 cycle delay, which limits writes to one every 3 cycles.</li> </ul> <p>When L3_DATA_WR_LATENCY is set to 2p, the L3_DATA_RD_LATENCY must be set to 3.</p>  |
| L3_DATA_RD_LATENCY  | 2 or 3           | <p>L3 cache data RAM output latency:</p> <ul style="list-style-type: none"> <li>2 cycles latency.</li> <li>3 cycles latency.</li> </ul> <p><b>Note:</b><br/>L3_DATA_RD_LATENCY must be set to 3 when L3_DATA_WR_LATENCY is set to 2p.</p>   |
| L3_DATA_RD_SLICE    | TRUE, FALSE      | <p>L3 cache data RAM output register slice</p> <p><b>TRUE</b> Include L3 data RAMs output register slice.<br/><b>FALSE</b> Do not include register slice.</p>   |
| L3_DATA_STRETCH_CLK | TRUE, FALSE      | <p>Stretch the clock to the L3 data RAMs.</p> <p><b>TRUE</b> Stretch the clock so that the RAM clock pulse is HIGH for a whole SCLK cycle.<br/><b>FALSE</b> Use only gating, so that the RAM clock pulse is HIGH for half an SCLK cycle.</p> <p><b>Note:</b><br/>In either case, the L3 RAM will not be accessed on consecutive SCLK cycles. This option has no performance impact.</p>   |
| CORE_REG_SLICE      | TRUE, FALSE      | <p>For each core, include a register slice between the core and the <i>Snoop Control Unit</i> (SCU).</p> <p>For each core present, starting with core 0, set the value to <b>TRUE</b> or <b>FALSE</b>.</p> <p><b>TRUE</b> Include a register slice between the core and the SCU.<br/><b>FALSE</b> Do not include a register slice.</p>  |

| Parameter name      | Permitted values          | Description   |
|---------------------|---------------------------|---|
| ASYNC_BRIDGE        | TRUE,<br>FALSE            | <p>For each core, asynchronous bridge on core to L3 coherent interface.</p> <p>For each core present, starting with core 0, set the value to <b>TRUE</b> or <b>FALSE</b>.</p> <p><b>TRUE</b>      Include an asynchronous bridge for the main bus between core and L3.<br/> <b>FALSE</b>      Do not include an asynchronous bridge.</p> <p><b>Note:</b><br/> The choice affects the internal interface that is used for instruction and data fetch, evictions, and snoops. It does not affect the other interfaces such as debug, trace, and GIC which are always asynchronous.</p>  |
| INTERLEAVE_ADDR_BIT | 6, 7, 8, 9, 10, 11, or 12 | <p>Controls which physical address bit is used to interleave requests between cache slices and dual ACE or dual CHI masters. The default value is bit 6, which interleaves on cache line boundaries. Other values can only be configured when dual ACE or dual CHI interfaces are configured.</p> <p><b>Note:</b><br/> Interleaving on a larger granularity might help improve system performance on some SoC designs. However, it can also reduce performance under some circumstances on accesses that hit in the L3 cache, because the same interleave is used for both cache slices and ACE master. Therefore if changing this parameter, Arm® recommends performing benchmarking in your system to determine if the overall performance is acceptable.</p> |
| CORE_SYNC_LEVELS    | 2 or 3                    | Number of synchronizer stages in all asynchronous inputs into the core.   |
| SYNC_LEVELS         | 2 or 3                    | Number of synchronizer stages in all asynchronous inputs to the SCU and cluster logic.  |
| LEGACY_V7_DEBUG_MAP | TRUE,<br>FALSE            | <p>Legacy v7 debug memory map. Configure v7 or v8 Debug memory map.</p> <p><b>TRUE</b>                      v7 Debug memory map.<br/> <b>FALSE</b>                      v8 Debug memory map.</p>  |
| ELA                 | TRUE,<br>FALSE            | <p>Support for integrating the CoreSight ELA-500 Embedded Logic Analyzer.</p> <p><b>TRUE</b>                      Include one ELA-500 instance within the DSU.<br/> <b>FALSE</b>                      Do not include the ELA-500 within the DSU.</p> <p><b>Note:</b><br/> If enabled, to create a unique ELA instance for the DSU, either:</p> <ul style="list-style-type: none"> <li>• Use the <code>-ela</code> option with the generate script.</li> <li>• Run the unquify script that is supplied with the ELA.</li> </ul> <p>For more information, see <i>Configuring the execution testbench with the ELA-500</i> in the <i>Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide</i>.</p>   |
| ELA_RAM_ADDR_SIZE   | 2-25                      | The number of index bits in the ELA-500 RAM. For more information, see <i>Configuration parameters</i> section in the <i>Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide</i> .   |

| Parameter name | Permitted values | Description  |
|----------------|------------------|--|
| L3_PBHA        | TRUE, FALSE      | <p>Store the <i>Page-Based Hardware Attributes</i> (PBHA) in the L3 cache.</p> <p><b>TRUE</b> PBHA bits are stored in the L3 cache so are valid for L3 evictions to the master port.</p> <p><b>FALSE</b> PBHA bits are not stored in the L3 cache so are only valid for read transactions to the master port.</p> <p><b>Note:</b><br/>Not all cores provide PBHA bits. Only enable <code>L3_PBHA</code> when all cores support the PBHA bits. For more information about the PBHA bits, see <a href="#">3.4 Page-based hardware attributes</a> on page 31.</p> |
| POP_RAM        | TRUE, FALSE      | <p>Configure the RTL for Arm <i>Processor Optimization Pack</i> (POP) RAMs.</p> <p><b>TRUE</b> Configure the RTL for Arm POP RAMs.</p> <p><b>FALSE</b> Do not configure the RTL for Arm POP RAMs.</p> <p><b>Note:</b><br/>The RAM interface is different when POP_RAM is set to TRUE. Only set POP_RAM to TRUE when using Arm POP RAMs.</p>  |

### 2.3.1 Cluster configurations

A cluster can be created with one or more cores of the same type.

Alternatively, the cluster can be configured as a big.LITTLE cluster, with two different types of core in the same cluster.

You can also configure a cluster with three types of core in the same cluster, where the third type of core is designated as other. This other core can be any type of big core.

For information on the number and type of permissible cores in the DSU, see the appendix *Compatible core versions* in the *Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide*.

## 2.4 Supported standards and specifications

The DynamIQ™ Shared Unit complies with the Arm®v8.2-A or Arm®v8.4-A architecture.

**Table 2-2: Compliance with standards and specifications**

| Architecture specification or standard                  | Version   |
|---|---|
| Arm architecture  | Arm®v8.2-A or Arm®v8.4-A depending on core.   |
| <i>Advanced Microcontroller Bus Architecture</i> (AMBA) | <ul style="list-style-type: none"> <li>AMBA ACE5 (Issue F.b).</li> <li>AMBA 5 CHI (Issue B or Issue C, depending on configured cores).</li> </ul> |
| CoreSight™  | v3.   |

| Architecture specification or standard  | Version   |
|---|---|
| Debug   | Arm®v8.2-A or Arm®v8.4-A depending on core.   |
| Generic Interrupt Controller(GIC) architecture CPU interface and Stream Protocol interface. | v4.0<br><br><b>Note:</b><br>The DSU uses Affinity Level-1 to distinguish between different cores, which some interrupt controllers, such as GIC-500 do not support. |
| Performance Monitoring Unit (PMU)   | PMUv3.  |

## 2.5 Test features

The DSU provides interfaces for manufacturing test.

The following manufacturing test interfaces are supported:

|              |                    |
|--------------|--------------------|
| <b>DFT</b>   | For logic testing. |
| <b>MBIST</b> | For RAM testing.   |

## 2.6 Design tasks

The DynamIQ™ Shared Unit is delivered as a synthesizable *Register Transfer Level* (RTL) description in SystemVerilog HDL. Before you can use it, you must implement, integrate, and program it.

A different party can perform each of the following tasks. Each task can include implementation and integration choices that affect the behavior and features of the DSU and its associated cores.

### Implementation

The implementer configures and synthesizes the RTL to produce a hard macrocell. This task includes integrating RAMs into the design.

### Integration

The integrator connects the macrocell into a SoC. This task includes connecting it to a memory system and peripherals.

### Programming

The system programmer develops the software to configure and initialize the DSU and its associated cores and tests the application software.

The operation of the final device depends on the following:

### Build configuration

The implementer chooses the options that affect how the RTL source files are pre-processed. These options usually include or exclude logic that affects one or more of the area, maximum frequency, and features of the resulting macrocell.

## Configuration inputs

The integrator configures some features of the DSU by tying inputs to specific values. These configuration settings affect the start-up behavior before any software configuration is made. They can also limit the options available to the software.

## Software configuration

The programmer configures the DSU by programming particular values into registers. The configuration choices affect the behavior of the DSU and its associated cores.

## 2.7 Product revisions

This section describes the differences in functionality between product revisions.

|             |   |
|-------------|---|
| <b>r0p0</b> | First release.  |
| <b>r0p1</b> | CHI and ACP support added.  |
| <b>r0p2</b> | There are no functional changes in this release.  |
| <b>r1p0</b> | Support for no L3 cache, improved cache stashing support, cache partitioning support, and additional configuration options.   |
| <b>r2p0</b> | Support added for three types of core, CHI.C for Direct connect, and 1.5MB L3 cache size.   |
| <b>r3p0</b> | Support added for dual CHI masters.   |
| <b>r4p0</b> | Support added for CHI.D for Direct connect, 256kB cache size, and transport of <i>Page-Based Hardware Attributes</i> (PBHA).  |
| <b>r4p1</b> | Support added for additional core combinations, such as: 5 LITTLE cores; 4 big and 2 LITTLE cores, and transport of CHI REQSRCATTR signals for Direct connect, when required. |



## 3. Technical overview

This chapter describes the structure of the DSU.

### 3.1 Components

A DynamIQ™ cluster system comprises of two top-level modules:


- A module which includes the cores and the *DynamIQ™ Shared Unit* (DSU).
- The DebugBlock.

In this book, the DynamIQ™ cluster is referred to as the cluster.

Separating the debug components from the cluster allows the debug components to be implemented in a separate power domain, allowing debug over power down.

The following figure shows the major components in a DynamIQ™ cluster system.

**Figure 3-1: DynamIQ™ cluster components**



Throughout this book CN represents a set of cores, where CN has a value of the total number of cores -1. The Arm architecture allows for cores to be single, or multithreaded. A *Processing Element* (PE) performs a thread of execution. A single-threaded core has one PE and a multithreaded core has two or more PEs. Where a reference to a core is made, the core can be a single, or

multithreaded core. Signal names that are associated with PEs use the abbreviation PE, where PE has a value of the total number of PEs - 1.

### DynamIQ™ cluster compatible cores

The core types, and optional features are selected by the macrocell implementer. The cores are described in their respective TRMs. Because the cores are instantiated within the cluster, all interfacing between the cores and the DSU is implemented automatically. All external signal inputs and outputs pass through the DSU. The DSU buffers and resynchronizes many of these signals to allow the cores to be clocked at different speeds. The memory interface of each core is internally connected to the DSU L3 memory system. Where necessary, the DSU implements additional buffering to compensate for different clock rates of the core and DSU L3. Each core implements clock and power control interfaces. These interfaces are routed through the DSU to the respective core external power controller.

### CPU bridges

The CPU bridges control buffering and synchronization between the cores and DSU.



The CPU bridge RTL is included with the DSU deliverables, but is instantiated in the hierarchy of processor cores.

---

### Snoop Control Unit (SCU)

The SCU maintains coherency between all the data caches in the cluster.

The SCU contains buffers that can handle direct cache-to-cache transfers between cores without having to read or write data to the L3 cache. Cache line migration enables dirty cache lines to be moved between cores. Also, there is no requirement to write back transferred cache line data to the L3 cache.

### Clock and power management

The cluster supports a set of power-saving modes that are controlled by an external power controller. The modes are selected through power-mode requests on P-Channels, for each of the cores, and a separate P-Channel for the DSU.

Clock gating is supported through Q-Channel requests from an external clock controller to the DSU. The Q-Channels allow individual control of the SCLK, PCLK, ATCLK, and GICCLK clock inputs.

### L3 memory interfaces

#### Main memory master

The main memory interface supports up to two ACE or CHI master interfaces.

#### Accelerator Coherency Port

The *Accelerator Coherency Port* (ACP) is an optional slave interface. The ACP provides direct memory access to cacheable memory. The SCU maintains cache coherency by checking ACP

accesses for allocation in the core and L3 caches. The ACP implements a subset of the ACE-Lite protocol.

### Peripheral port

The peripheral port is an optional master interface and provides **DEVICE** accesses to tightly coupled accelerators. The port implements the AXI 4 master interface protocol.

### L3 cache

The cache size is implemented as either 256KB, 512KB, 1MB, 1.5MB, 2MB, 3MB, or 4MB. The 1.5MB and 3MB caches are 12-way set associative. The 256KB, 512KB, 1MB, 2MB, and 4MB caches are 16-way set associative. All caches have a 64-byte line length. ECC protection of data and tag RAMs is optionally implemented.

### DSU system control registers

The DSU implements a set of system control registers, which are common to all cores in the cluster. You can access these registers from any core in the cluster. These registers provide:

- Control for power management of the cluster.
- L3 cache partitioning control.
- CHI QoS bus control and scheme ID assignment.
- Information about the hardware configuration of the DSU.
- L3 cache hit and miss count information.

### Debug and trace components

Each core includes an *Embedded Trace Macrocell* (ETM) to allow program tracing while debugging.

Trigger events from the cores are combined and output to the Debug APB master. Trigger events to the cores, and debug register accesses, are received on the Debug APB slave.

### DebugBlock

The debug components are split between two architecturally defined logical power domains, the core power domain, and the debug power domain:

- The core power domain comprises one or more physical power domains for the cores and DSU.
- The debug power domain includes the DebugBlock.

The separate power domains allow the cores and the cluster to be powered down while maintaining essential state that is required to continue debugging. Separating the logical power domains into physical domains is optional and might not be available in individual systems.

### Cluster to DebugBlock APB

Trigger events from the cores are transferred to the DebugBlock as APB writes.

### DebugBlock to cluster APB

Trigger events to the cores are transferred as APB writes to the DSU. Register accesses from the system debug APB are transferred to the DSU.

### System debug APB

The system debug APB slave interface connects to external CoreSight™ components, such as the *Debug Access Port* (DAP).

### CTI and CTM

The DebugBlock implements an *Embedded Cross Trigger* (ECT). A *Cross Trigger Interface* (CTI) is allocated to each PE in the cluster, and an additional CTI is allocated to the cluster ELA when present. The CTIs are interconnected through the *Cross Trigger Matrix* (CTM). A single external channel interface is implemented to allow cross-triggering to be extended to the SoC.

### Debug ROM

The ROM table contains a list of components in the system. Debuggers can use the ROM table to determine which CoreSight components are implemented.

### Power management and clock gating

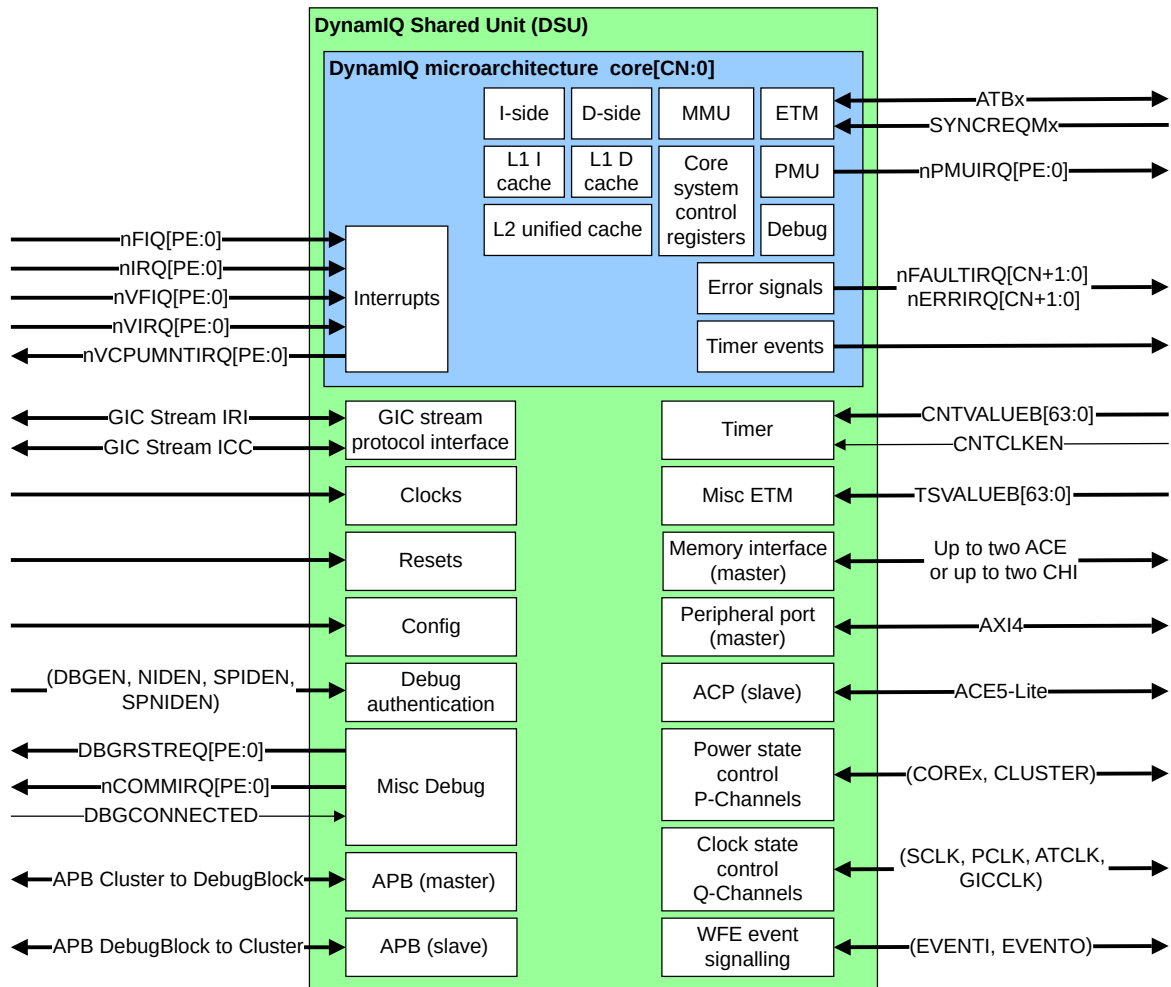
The DebugBlock implements two Q-Channel interfaces, one for requests to gate the PCLK clock, and a second for requests to control the debug power domain.

## 3.2 Interfaces

The external interfaces to connect to the SoC system.

### DSU interfaces

The following figure shows the major external interfaces of the *DynamIQ™ Shared Unit* (DSU).

**Figure 3-2: DSU interfaces**

The following table describes the major external interfaces.

**Table 3-1: DSU interfaces**

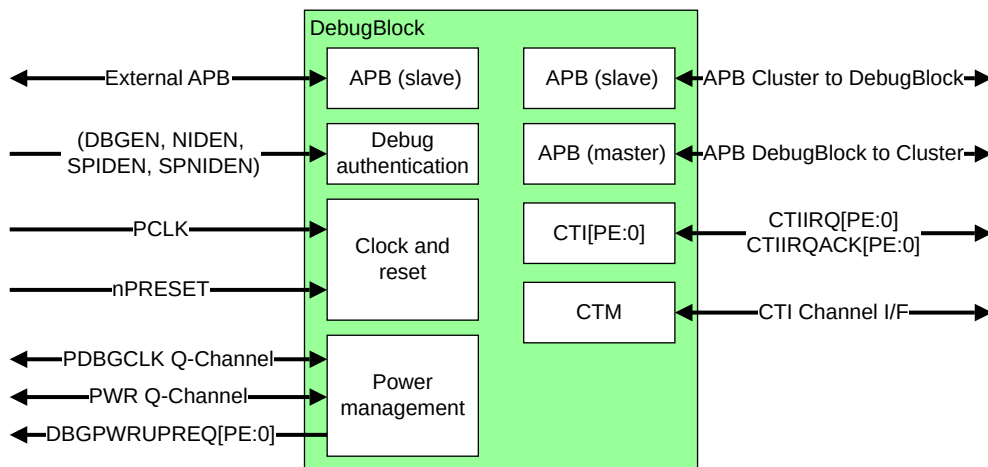
| Purpose                               | Protocol   | Notes   |
|---------------------------------------|------------|---|
| Trace                                 | ATB        | Master ATB interfaces. Each core has an ATB interface to output ETM trace information, see <a href="#">Appendix B.2.14 ATB Interface Signals</a> .                      |
| Memory                                | ACE or CHI | Master interface to main memory. You can configure the DSU with either 1 or 2 ACE, or 1 or 2 CHI interfaces. See <a href="#">2.3 Implementation options</a> on page 18. |
| Accelerator Coherency Port (optional) | ACE5-Lite  | Slave interface allowing an external master to make coherent requests to cacheable memory.  |
| Peripheral port (optional)            | AXI        | Low-latency master interface to external Device memory.   |
| Cluster to DebugBlock                 | APB        | APB interface from the cluster (master) to the DebugBlock (slave).  |
| DebugBlock to cluster                 | APB        | APB interface from the DebugBlock (master) to the cluster (slave).  |

| Purpose                                  | Protocol  | Notes   |
|--|-----------|---|
| Power state control                      | P-Channel | P-Channels for DSU and core power management.   |
| Clock state control                      | Q-Channel | Q-Channels for clock gating control.  |
| WFE event signaling                      | -         | Signals for <i>Wait For Event</i> (WFE) wake up events.   |
| Generic timer                            | -         | Input for the generic timer count value. The count value is distributed to all cores. Each core outputs timer events.                                   |
| GIC interfaces                           | -         | Interrupts to individual cores. A single GIC Stream Protocol interface is shared by all cores.  |
| <i>Design for Test</i> (DFT)             | -         | Interface to allow access for <i>Automatic Test Pattern Generation</i> (ATPG) scan-path testing.  |
| <i>Memory Built-In Self Test</i> (MBIST) | Arm MBIST | Internal interface that supports the manufacturing test of the L3 and SCU memories embedded in the DSU. Each core has its own internal MBIST interface. |

## DebugBlock interfaces

The following figure shows the major external interfaces.

**Figure 3-3: DebugBlock interfaces**



The following table describes the major external interfaces.

**Table 3-2: DebugBlock interfaces**

| Purpose                         | Protocol  | Notes  |
|---------------------------------|-----------|--|
| External debug                  | APB       | Slave interface to external debug component, for example a <i>Debug Access Port</i> (DAP). Allows access to Debug registers and resources. |
| Cluster to DebugBlock           | APB       | APB interface from the cluster (master) to the DebugBlock (slave).   |
| DebugBlock to cluster           | APB       | APB interface from the DebugBlock (master) to the cluster (slave).   |
| Cross-trigger channel interface | CTI       | Allows cross-triggering to be extended to external SoC components.   |
| Power management                | Q-Channel | Enables communication to an external power controller. To control clock gating and powerdown.  |

## 3.3 RAS support

The DSU supports *Reliability, Availability, and Serviceability* (RAS) features.

- Data poisoning on a 64-bit granule is supported in the DSU.
- Optional ECC protection on RAMs.
- Error recovery and fault handling interrupt outputs.
- Error record registers.

## 3.4 Page-based hardware attributes

The *Page-Based Hardware Attributes* (PBHA) bits are provided by the processor cores, and passed on or preserved by the DSU.

The PBHA bits are provided externally as part of the SRCATTR bus. PBHA affects the following:

### RAM sizes

To generate accurate PBHA bits on L3 cache evictions, store the bits in the L3 cache, by setting the `L3_PBHA` parameter. If this parameter is not set, the PBHA bits are only accurate for read transactions. If this parameter is set, the width of all L3 tag RAM instances is increased by two bits.

### ACP

All requests from ACP are given a fixed PBHA value of 0b00.

### Cache stash transactions on CHI

Cache stash transactions might be sent on the CHI interface. For these requests, the PBHA bits being used must be sent along with the stash snoop transaction. The `RXSNPFLIT[FwdTxnID]` field has two spare bits which are used for this functionality, `FwdTxnID[7:6]`. If the interconnect does not support PBHA, then these bits are driven to 0b00 by default.

### Transaction support

Transactions that do not have a physical address that is associated with them, for example DVM messages, do not provide the PBHA bits. Evict transactions that do not provide any data, for use in de-allocating a snoop filter, do not provide PBHA bits.

### Mismatched aliases

If the same physical address is accessed through more than one virtual address mapping, and the PBHA bits are different in the mappings, then the results are **UNPREDICTABLE**. The PBHA value sent on the bus could be for either mapping.



Not all cores provide the PBHA bits. When used with a core that does not provide the PBHA bits, the PBHA output of the DSU is not valid.

## 3.5 L3 memory system variants

The L3 cache might not be implemented in your system. There are three possible variants.

Possible L3 memory system implementations are:

### **L3 cache present**

This implementation is the default. It provides the most functionality and is suitable for general-purpose workloads.

### **L3 cache not present**

The L3 cache is not present, but snoop filter and SCU logic are present.

This variant allows multiple cores in the cluster, and manages the coherency between them. It supports other implementation options such as ACP, peripheral port, and ACE or CHI master ports. There is an area saving from not including the L3 cache RAMs, however performance of typical workloads is reduced. Therefore Arm recommends that this variant is only used in specialized use cases, or when there is a system cache present that the cores can use.

### **Direct connect**

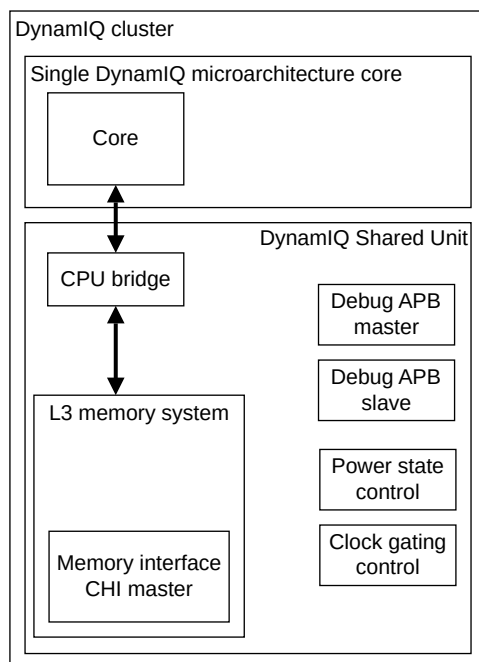
The L3 cache, snoop filter, and SCU logic are not present.

This variant is specifically for use with a CHI interconnect. It offers additional area savings and reduced latency when compared to the previous variants. Because there is no L3 cache in the cluster, this variant relies on the system cache in the interconnect for performance. Only some types of core support this variant. Of the cores that do support this variant, some only support this variant when used with the CMN-600 interconnect as they integrate its system address map component, while others do not have this restriction. See the Technical Reference Manual of your core to identify if your core supports this variant.

Because this variant does not include any coherency logic, it is only supported when there is a single core in the cluster. Optional interfaces such as ACP or the peripheral port are not supported. See [2.3 Implementation options](#) on page 18 for more information. The master port must be a single 256-bit CHI interface.

The following diagram shows the DSU implemented with L3 Direct connect.



**Figure 3-4: L3 Direct connect implementation**

## 4. Clocks and resets

This chapter describes the clocks and resets of the DSU.

### 4.1 Clocks

The DSU requires clock signals for each of the cores, internal logic, and external interfaces.

The following table describes the clocks.

**Table 4-1: DSU clock signals**

| Signal        | Description  |
|---------------|--|
| CORECLK[CN:0] | The per-core clocks for all core logic including L1 caches and L2 caches.  |
| SCLK          | The clock for the SCU and L3 memory system, including the ACE or CHI master interface. SCLK is also used for any cores that are configured to run synchronously to the DSU.  |
| PCLK          | The clock for the DebugBlock and DSU debug APB interfaces.<br><br><b>Note:</b><br>The DebugBlock and cluster both have PCLK inputs. You might choose to connect these to the same clock. Alternatively, you might choose to place an asynchronous bridge between the two clock inputs, in which case they might be different clocks. |
| ATCLK         | The clock for the ATB trace buses output from the DSU.<br><br><b>Note:</b><br>All ATB buses output from the DSU share the same clock.  |
| GICCLK        | The clock for the GIC AXI-stream interface between the DSU and an external GIC.  |
| PERIPHCLK     | The clock for peripheral logic inside the DSU such as timers, and clock and power management logic.  |

All clocks can be driven fully asynchronously to each other. The DSU contains all the necessary synchronizing logic for crossing between clock domains. There are no clock dividers and no latches in the design. The entire design is rising edge triggered.



Note

- It is possible to configure some or all the cores to run synchronously with the L3 memory system. If this option is chosen, the corresponding CORECLK signals are not present and the synchronous cores run with SCLK.
- The DebugBlock can be clocked by a different clock from the DSU PCLK. To allow this, the macrocell implementer can choose to add asynchronous bridges between the DSU and the DebugBlock.

Some external interfaces, such as the main ACE or CHI master interface, support a clock enable input to allow the external logic to run at a lower, synchronous, frequency.

While there is no functional requirement for the clocks to have any relationship with each other, the DSU is designed with the following expectations to achieve acceptable performance:

- CORECLK[CN:0] is dynamically scaled to match the performance requirements of each core.
- SCLK frequency affects the L3 hit latency and, therefore, it is important for achieving good performance. For best performance, Arm recommends running SCLK as close to CORECLK[CN:0] frequency as possible. However to reduce dynamic and leakage power, targeting a lower frequency might be required. Running SCLK at least approximately 75% of the CORECLK[CN:0] frequency might give an appropriate balance for many systems.
- SCLK can run at synchronous 1:1 or 2:1 frequencies with the external interconnect, avoiding the need for an asynchronous bridge between them.
- PCLK, ATCLK, and GICCLK can run at the same frequency as the relevant SoC components that they connect to. This would typically be approximately 25% of the maximum CORECLK[CN:0] frequency.
- PERIPHCLK contains the architectural timers, and software performance can be impacted if reads to these registers take too long. Therefore, Arm recommends that PERIPHCLK is run at least 25% of the maximum CORECLK[CN:0] frequency.

## 4.2 Resets

The DSU requires reset signals for each of the cores, internal logic, and external interfaces.

**Table 4-2: DSU reset signals**

| Signal            | Description  |
|-------------------|--|
| nCPUPORESET[CN:0] | The per-core primary Cold reset signal for all resettable registers in the CORECLK domain including debug registers, ETM registers, and RAS registers. |
| nCORERESET[CN:0]  | The per-core Warm reset signal for all resettable registers in the CORECLK domain excluding the debug registers, ETM registers, and RAS registers.     |
| nPRESET           | A single cluster-wide reset signal for all resettable registers in the PCLK domain.  |
| nSPORESET         | A single cluster-wide Cold reset signal for all resettable registers in the SCLK domain.   |
| nSRESET           | A Warm reset signal for all resettable registers in the SCLK domain excluding RAS registers.   |
| nATRESET          | A single cluster-wide reset signal for all resettable registers in the ATCLK domain.   |
| nGICRESET         | A single cluster-wide reset signal for all resettable registers in the GICCLK domain.  |
| nPERIPHRESET      | A single cluster-wide reset signal for all resettable registers in the PERIPHCLK domain.   |
| nMBISTRESET       | A single cluster-wide reset signal that acts on all resettable registers in the CORECLK and SCLK domains.  |

The nMBISTRESET is active low and must be driven high for functional mode. A single top level nMBISTRESET is utilized to reset the IP in preparation for MBIST operations.

All reset inputs can be asserted (HIGH to LOW) and deasserted (LOW to HIGH) asynchronously. Reset synchronization logic inside the DSU ensures that reset deassertion is synchronous for all resettable registers inside those reset domains. The core clock does not need to be present for reset assertion, but it must be present for reset deassertion to ensure reset synchronization.

## AArch32 (EL3) to AArch64 Execution state transition at reset

At *Exception level 3* (EL3), cores can only transition between AArch32 and AArch64 states at reset. The Execution state after reset is controlled by the AA64nAA32[PE:0] configuration signals. These signals are only sampled at reset.

To reset a core and change Execution state from software, a Warm reset request can be made by setting the RR bit of the RMR system register (from AArch32) or the RMR\_EL3 register (from AArch64). Following the register write and executing a `wfi` instruction, the cluster automatically resets the core without requiring any action by the external reset controller. The hardware automatically cleans and invalidates all the caches and safely disconnects the core from cluster before the reset is asserted.



Your implementation might include cores that do not support execution in AArch32 state at EL3.

---

## 5. Power management

This chapter describes the power domains and the power modes in the DSU.

### 5.1 About DSU power management

The DSU supports a range of low-power modes and cache RAM powerdown modes.

The DSU supports the following power modes:

#### **On mode**

On mode is the normal mode of operation where all the core and DSU functionality is available. The DSU individually disables internal clocks, and inputs to unused functional blocks. Only the logic that is in use consumes dynamic power.

#### **Functional retention mode**

Functional retention mode allows the L3 cache and snoop filter RAMs to be put temporarily in to a retention state while the L3 cache is not being accessed. The contents of the cache RAMs are retained.

#### **Memory retention mode**

Memory retention mode allows the L3 cache RAMs to be held in retention while the rest of the cluster is powered down. Keeping the RAMs in retention reduces the energy cost of writing dirty lines back to memory and reduces the cluster response time on powerup. It is not possible to snoop the cache in this mode. Therefore it is important that no other external coherent agents are active (for example, cores external to the cluster, or other coherent devices). In practice, this mode can only be used in a coherent system when the cluster is the only active agent.

#### **Off mode**

In Off mode, power is removed completely, and no state is retained. To avoid losing data, the DSU cleans and invalidates the L3 cache before taking the cluster out of coherence.

The DSU supports clock, voltage, and power domains that can be controlled by external logic.

The cluster, along with power management software, gives operating requirement hints to an external power controller. The power controller is responsible for:

- Coordinating power management with the rest of the SoC;
- Switching and isolating power and voltage domains; and
- Controlling clock gating cells.

## 5.2 Power mode control

Power management control is distributed between power management software, the cluster, and an external power controller.

The cluster supports a set of power operating requirements which act as hints to the external power controller. The operating requirements indicate: the required cache capacity, the RAM retention mode, and whether the cluster logic can be powered up or down. The power controller controls the logic clamps and power switches required to put the RAMs and logic into low-power operation.

Software sets the operating requirements by writing to the following system registers:

### **Cluster Power Control Register (CLUSTERPWRCTLR\_EL1)**

To request partial L3 cache powerup or powerdown, and to enable RAM retention capabilities.

### **Cluster Powerdown Register (CLUSTERPWRDN\_EL1)**

To request the power mode that the cluster is to enter, after all cores have powered off. For example, memory retention mode.

The operating requirements are signaled to the power controller through the cluster P-Channel interface. The power controller responds to a change of operating requirements by sequencing the transition between lower or higher power modes.

Power management algorithms can use *Cluster L3 Hit Counter Register* (CLUSTERL3HIT\_EL1) and *Cluster L3 Miss Counter Register* (CLUSTERL3MISS\_EL1) system registers to determine when to powerup or powerdown cache portions.

The status of the power settings is indicated in the *Cluster Power Status Register* (CLUSTERPWRSTAT\_EL1) system register.

The cluster receives power mode transition requests from the power controller and checks the validity of each transition. If the transition is supported, the cluster accepts the request. If the transition is not supported, the cluster denies the request. If the cluster accepts the request, the power controller can switch power domains off as appropriate.

The cluster automatically performs any internal operations required by a mode transition, before accepting the new mode. The internal actions performed by the cluster include: gating clocks, flushing caches, and disabling coherency.



If L3 RAM retention is not implemented, CLUSTERPWRCTLR\_EL1.L3 data RAM retention control must be left in the reset state.

---

## 5.3 Communication with the power controller

The cluster provides a P-Channel interface to allow the external power controller to set the operating power mode of the cluster, in response to requests from the cluster.

The cluster indicates the operating requirements on the CLUSTERPACTIVE bus. The power controller can then request a new power mode. The power controller indicates the requested mode on the CLUSTERPSTATE bus and asserts the CLUSTERPREQ handshake signal.

When the cluster has performed all the actions that are required in preparation for the power mode transition, the cluster accepts the request by asserting CLUSTERPACCEPT. If the request is not valid, the cluster denies the request by asserting CLUSTERPDENY. Reasons for denying the request might be, requesting an incorrect mode transition, or the requested mode is no longer appropriate because the L3 operating mode has changed.

At reset, the cluster reads the initial mode set by the power controller on the CLUSTERPSTATE bus.



Arm recommends that CLUSTERPREQ is asserted at reset. If CLUSTERPREQ is not asserted at reset, the power controller must wait for 40 PERIPHCLK cycles after reset is deasserted before it drives CLUSTERPREQ HIGH.

---

## 5.4 L3 RAM power control

Power savings can be made by putting parts, or all, of the L3 control logic and L3 cache RAM into retention modes, or powering down portions of the RAM.

### 5.4.1 L3 cache partial powerdown

Sections of the L3 cache that are known as portions, can be independently powered down to reduce RAM leakage power.

When all the L3 cache capacity is not required, then the L3 cache size can be reduced by powering down one or more of these portions. For instance, for small memory footprint workloads.

The L3 cache RAMs are organized to allow separate control of groups of cache ways. Each group has four cache ways. The L3 data RAMs are organized into two equal sized portions, where each portion consists of two groups of cache ways. The L3 tag RAMs are organized into four equal sized portions, with each portion corresponding to each group of cache ways.

Power control can be applied independently to each portion.

The operating cache capacity can be selected from: all,  $\frac{3}{4}$ ,  $\frac{1}{2}$ ,  $\frac{1}{4}$ , and none.

The following table shows how the available operating cache capacities relate to the RAM power enables for each portion.

**Table 5-1: L3 Cache capacity and RAM power enables**

| Cache capacity | Tag RAM |     |     |     | Data RAM |     |
|----------------|---------|-----|-----|-----|----------|-----|
|                | 0       | 1   | 2   | 3   | 0        | 1   |
| None           | Off     |     |     |     | Off      |     |
| $\frac{1}{4}$  | On      | Off |     |     | On       | Off |
| $\frac{1}{2}$  | On      |     | Off |     | On       | Off |
| $\frac{3}{4}$  | On      |     |     | Off | On       |     |
| All on         | On      |     |     |     | On       |     |

This table shows that one portion of the tag RAM for the  $\frac{1}{4}$  and  $\frac{3}{4}$  cache capacities are always powered down. Therefore, these cache capacities may not achieve the same power savings of the  $\frac{1}{2}$  and all L3 capacities, respectively.

The external power controller requests the required L3 cache capacity, through the cluster P-Channel, in response to power management software setting the operating requirements.



Confirm with your implementer whether RAM powerdown is supported. If RAM powerdown is supported, then ensure that the necessary logic clamps and power switches are implemented.

Memory transactions from the cores can still be processed when all of the portions are off, and while a portion is being powered on or off.

For more information on L3 cache slices and portions, see [6.7 Cache slices and portions](#) on page 62.

## 5.4.2 L3 RAM retention

The *DynamIQ™ Shared Unit* (DSU) supports the Functional retention and Memory retention power modes.

### Functional retention

Functional retention mode (SFONLY FUNC\_RET,  $\frac{1}{4}$ FUNC\_RET,  $\frac{1}{2}$ FUNC\_RET,  $\frac{3}{4}$ FUNC\_RET, or FULL FUNC\_RET mode) allows the L3 cache and snoop filter RAMs to be put temporarily into retention while the L3 cache is not being accessed.

When the L3 cache has not been accessed for a period of time, the DSU signals to the power controller that the L3 cache RAM can be put into retention. While in functional retention mode, core and snoop requests can still be received. On receiving a request, the DSU signals to the power controller to take the RAMs out of retention. The core or snoop request is stalled until the power controller signals to the DSU to enter an ON mode



(SFONLY ON,  $\frac{1}{4}$ ON,  $\frac{1}{2}$ ON,  $\frac{3}{4}$ ON, or FULL ON). When the request is completed, the DSU signals to the power controller that the RAMs can be put back into retention.

The *Cluster Power Control Register* (CLUSTERPWRCTLR) system register determines the duration of inactivity before the DSU requests the RAMs to be put into retention.

### Memory retention

Memory retention mode ( $\frac{1}{4}$ MEM\_RET,  $\frac{1}{2}$ MEM\_RET,  $\frac{3}{4}$ MEM\_RET, or FULL\_MEM\_RET mode) allows the L3 cache RAMs to be put into retention. Memory retention mode also disables the control logic in the DSU, and allows the snoop filter and *Long-Term Data Buffer* (LTDB) RAMs to be powered off.

Memory retention mode can be entered when the DSU is idle and all cores are OFF. In this mode, the L3 cache cannot process ACP or snoop requests.

Software can request the DSU to enter memory retention by setting the Memory retention required bit in CLUSTERPWDRN\_EL1 before that core is powered OFF. After all cores have transitioned to the OFF power mode, and the DSU becomes idle, the DSU indicates MEM\_RET on its P-Channel CLUSTERACTIVE signal.

Even if CLUSTERACTIVE indicates MEM\_RET, the DSU can still accept a P-Channel request to transition to the OFF mode.

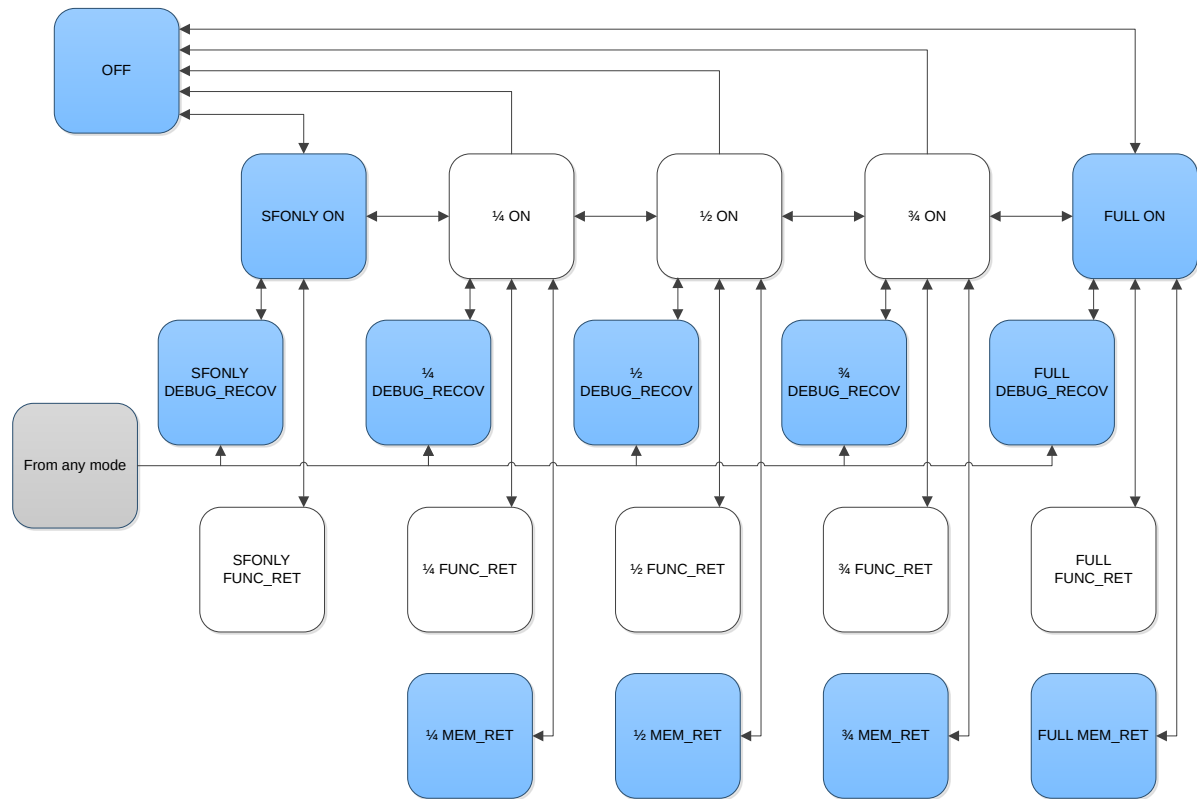
## 5.5 Power modes

The external power controller can request a new operating power mode in response to the operating requirements indicated by the DSU.

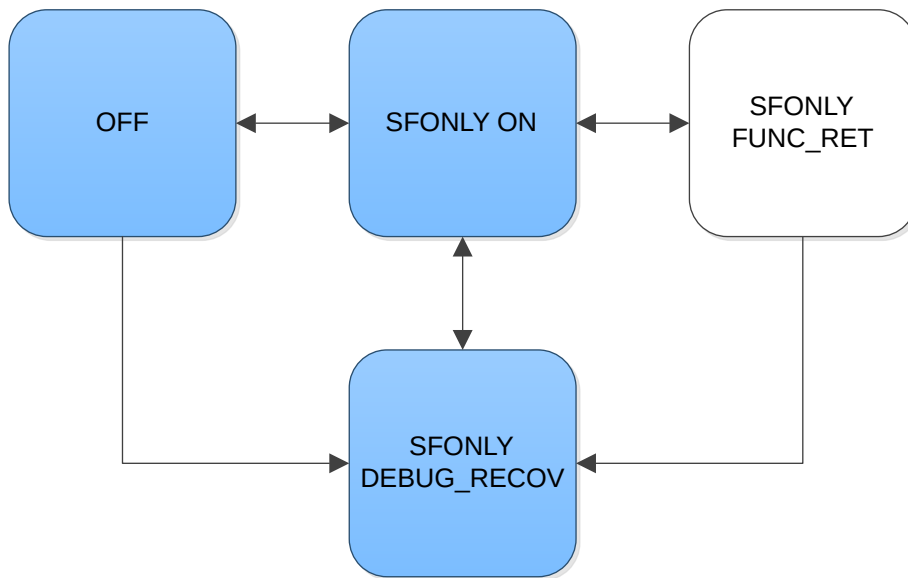
### 5.5.1 Power mode transitions

The *DynamIQ™ Shared Unit* (DSU) supports a set of power modes. The power controller can request the DSU to transition between modes.


The following diagram shows the supported DSU power modes, and the permitted transitions between them. The blue modes indicate the modes that the DSU can be initialized to at reset. See [5.3 Communication with the power controller](#) on page 38

**Figure 5-1: DSU power mode transitions**

The following diagram shows the supported DSU power modes when L3 is not implemented.

**Figure 5-2: DSU power mode transitions, no L3**

The following diagram shows the supported DSU power modes when Direct connect is implemented.

**Figure 5-3: DSU power mode transitions, Direct connect**



In a configuration that implements Direct connect, there is no snoop filter present. The power modes are named SFONLY for consistency with other configurations.

## FULL ON

In this mode, all DSU logic, snoop filter, and L3 cache RAMs are powered up and fully operational.

The DSU can be initialized into the FULL\_ON mode. In this case, it is treated as an implicit transition from the OFF mode so the L3 cache and snoop filter are invalidated. When a transition to the FULL ON mode is completed, the cache and snoop filter are accessible and coherent. The L3 cache does not require any configuration from software.

## SFONLY ON, ¼ON, ½ON, ¾ON

In these modes, the DSU logic and snoop filter RAMs are powered up but some of the L3 cache RAMs remain powered down.

The DSU can be initialized into the SFONLY ON mode. In this case, it is treated as an implicit transition from the OFF mode. When a transition to the SFONLY ON mode is completed, the snoop filter is accessible and coherent without needing any configuration from software.



Transitions between the ON modes are only allowed in incremental steps. For example, a transition directly from ½ON to FULL ON is not permitted. The CLUSTERACTIVE outputs reflect this, so that CLUSTERACTIVE does not indicate a required mode that cannot be directly reached from the current mode.

## SFONLY FUNC\_RET, ¼FUNC\_RET, ½FUNC\_RET, ¾FUNC\_RET, FULL FUNC\_RET

In these modes, the DSU is powered up, the snoop filter and L3 cache RAMs are in retention, meaning that the RAMs are inoperable but with state retained.

If a request from a core, or a snoop from the system, is required to be serviced:

- The DSU indicates that a transition to ON is required using CLUSTERACTIVE.
- The request is stalled until the DSU enters one of the ON modes.

When the RAMs are in retention, the clock to the RAMs is automatically gated outside of the retained domain.

## ¼MEM\_RET, ½MEM\_RET, ¾MEM\_RET, FULL MEM\_RET

In these modes, the DSU logic is powered down. The L3 cache RAMs are in retention and the snoop filter RAMs are powered down.

All cores in the cluster must be in the OFF mode to allow entry into any one of these modes.

Because the cache contains data that is accessible from the rest of the system, these modes must only be used if there are no other coherent agents in the system that are active and might need to access the data.

## OFF

In this mode, all DSU logic, snoop filter, and L3 cache RAMs are powered down.

The DSU can be initialized to this mode on Cold reset.

## Debug recovery mode

The Debug recovery mode can be used to assist debug of external watchdog-triggered reset events.

By default, the DSU invalidates the cache and snoop filter when there is a transition from OFF to an ON mode. In Debug recovery mode, cache invalidation is disabled allowing the contents of the L3 cache that were present before the reset to be observable after the reset. The contents of the L3 cache and snoop filter are preserved and are not altered on the transition back to the ON mode.

To enter Debug recovery mode, the P-Channel is initialized to `DEBUG_RECOV`, and the DSU is cycled through a reset.

To preserve the RAS state and cache contents, a transition to the Debug recovery mode can be made from any of the current states. Once in Debug recovery mode, a cluster-wide Warm reset must be applied externally. The RAS and cache state are preserved when the core is transitioned to the ON mode.



- Debug recovery mode is strictly for debug purposes. It must not be used for functional purposes, because correct operation of the cluster is not guaranteed when entering this mode.
- Debug recovery mode can occur at any time with no guarantee of the state of the cluster. A P-Channel request of this type is accepted immediately, therefore its effects on the core, cluster, or the wider system are unpredictable, and a wider system reset might be required. In particular, if there were outstanding memory system transactions at the time of the reset, then these transactions might complete after the reset when the cluster is not expecting them and cause a system deadlock. If the system sends a snoop to the cluster during this mode, then depending on the cluster state, the snoop might get a response and disturb the contents of the caches, or it might not get a response and cause a system deadlock.
- If the cluster is already in the middle of a power mode transition on the P-Channel, a clock gating transition on the SCLK Q-Channel, or the cluster is in Warm reset, then it might not be possible to enter `DEBUG_RECOV` without a Cold reset of the cluster.
- The `SFONLY`, `¼`, `½`, `¾` and `FULL` `DEBUG_RECOV` modes are all functionally equivalent. You can transition to any of these modes from any other mode, and to transition to any ON mode from any `DEBUG_RECOV` mode. However, you

must choose the correct ON mode corresponding to the L3 cache portions that were in use before Debug recovery mode.

## 5.5.2 Power mode transition behavior

When the power controller requests a transition between power modes, the *DynamIQ™ Shared Unit* (DSU) automatically performs some actions before accepting the new mode.

Your SoC might implement additional actions. When transitioning from a lower power to higher power mode, these additional actions are performed before the power controller requests a new mode. When transitioning from a higher power to lower power mode, these actions are performed after the DSU accepts the new mode.

The following table shows the permitted mode transitions and behavior that is associated with each transition.

**Table 5-2: Power mode transition behavior**

| Start mode  | End mode    | DSU behavior   | Partner implemented behavior                                    |
|-------------|-------------|--|---|
| OFF         | ON          | The L3 cache and snoop filter are initialized. The cluster is brought into coherency with the rest of the system.  | Power applied, isolation disabled.                              |
| MEM_RET     | ON          | The cluster is brought into coherency with the rest of the system.<br><br>The snoop filter RAMs are initialized.   | Power applied, isolation disabled.                              |
| ON          | FUNC_RET    | Waits for all memory transactions to complete. The clock to the retention domain is gated. L3 cache and snoop filter RAMs are put into retention.                      | RAM clamps and isolation enabled.                               |
| FUNC_RET    | ON          | L3 cache and snoop filter RAMs are taken out of retention.   | RAM clamps and isolation disabled.                              |
| FULL ON     | ¼ON         | Decreasing available cache ways. Relevant ways in L3 cache are cleaned and invalidated.  | RAM clamps and isolation are enabled for relevant ways.         |
| ¼ON         | ½ON         |  |   |
| ½ON         | ¾ON         |  |   |
| ¾ON         | SFONLY ON   |  |   |
| SFONLY ON   | ¼ON         | Increasing available cache ways. Relevant ways in L3 cache are initialized.  | RAM clamps and isolation are disabled for relevant ways.        |
| ¼ON         | ½ON         |  |   |
| ½ON         | ¾ON         |  |   |
| ¾ON         | FULL ON     |  |   |
| ON          | OFF         | Waits for all memory transactions to complete. L3 cache allocation is disabled. The L3 cache is cleaned and invalidated. The cluster is removed from system coherency. | DSU clamps and isolation to the rest of the system are enabled. |
| ON          | MEM_RET     | Waits for all memory transactions to complete. The cluster is removed from system coherency. This mode is only useful when the cluster is the only master active.      | DSU clamps and isolation to the rest of the system are enabled. |
| Any         | DEBUG_RECOV | -  | -   |
| DEBUG_RECOV | ON          | The cluster is brought into coherency with the rest of the system.   | Reset is applied, clamps and isolation are disabled.            |

### 5.5.3 Interlocks between core and DSU P-Channels

To ensure that correct operation is maintained, and to allow a cluster powerdown to be abandoned, there are interlocks between the core and DSU P-Channels:

- If a core P-Channel request is made while the DSU is not in an ON or FUNC\_RET mode, the core request stalls until the DSU has reached the appropriate mode.



If the DSU is in the MEM\_RET or OFF mode, the power controller must transition the DSU to the ON mode to avoid deadlock when signaling a core powerup request.

- If the DSU is requested to go to MEM\_RET or OFF while not all cores are OFF, the request is denied.
- If the DSU is in the process of transitioning from ON to OFF (particularly when flushing the L3 cache which can take a long time) and a core is requested to leave the OFF mode, the L3 cache flush is abandoned and the DSU P-Channel request is denied.

### 5.5.4 Power mode encoding

The power mode is encoded on the CLUSTERPSTATE P-Channel bus.

The following table shows the encoding of the power mode.

**Table 5-3: Cluster power domain CLUSTERPSTATE values**

| Power mode      | CLUSTERPSTATE Value | Logic      | Snoop filter and LTDB RAMs | L3 Tag ways 0-3, L3 Data portion 0, and L3 victim RAMs | L3 Tag ways 4-7 RAMs | L3 Tag ways 8-11 and L3 Data portion 1 RAMs | L3 Tag ways 12-15 RAMs |
|-----------------|---------------------|------------|----------------------------|--|----------------------|---|------------------------|
| FULL_ON         | 0b1001000           | Powered up | Powered up                 | Powered up   | Powered up           | Powered up                                  | Powered up             |
| ¾ ON            | 0b0111000           |            |                            |  |                      |   | Powered down           |
| ½ ON            | 0b0101000           |            |                            |  |                      |   |                        |
| ¼ ON            | 0b0011000           |            |                            |  |                      |   |                        |
| SFONLY ON       | 0b0001000           |            |                            | Powered down   |                      |   |                        |
| FULL_FUNC_RET   | 0b1000111           |            | Retention                  | Retention  | Retention            | Retention                                   | Retention              |
| ¾ FUNC_RET      | 0b0110111           |            |                            |  |                      |   | Powered down           |
| ½ FUNC_RET      | 0b0100111           |            |                            |  |                      |   |                        |
| ¼ FUNC_RET      | 0b0010111           |            |                            |  |                      |   |                        |
| SFONLY FUNC_RET | 0b0000111           |            |                            | Powered down   |                      |   |                        |

| Power mode         | CLUSTERPSTATE Value | Logic                           | Snoop filter and LTDB RAMs           | L3 Tag ways 0-3, L3 Data portion 0, and L3 victim RAMs | L3 Tag ways 4-7 RAMs                 | L3 Tag ways 8-11 and L3 Data portion 1 RAMs | L3 Tag ways 12-15 RAMs               |
|--------------------|---------------------|---------------------------------|--------------------------------------|--|--------------------------------------|---|--------------------------------------|
| FULL_MEM_RET       | 0b1000010           | Powered down                    | Powered down                         | Retention  | Retention                            | Retention                                   | Retention                            |
| ¾ MEM_RET          | 0b0110010           |                                 |                                      |  |                                      |   |                                      |
| ½ MEM_RET          | 0b0100010           |                                 |                                      |  | Powered down                         | Powered down                                | Powered down                         |
| ¼ MEM_RET          | 0b0010010           |                                 |                                      |  |                                      |   |                                      |
| OFF                | 0b0000000           |                                 |                                      | Powered down   | Powered down                         |   |                                      |
| FULL_DEBUG_RECOV   | 0b1001010           | Powered up or down <sup>1</sup> | Powered up or retention <sup>1</sup> | Powered up or retention <sup>1</sup>                   | Powered up or retention <sup>1</sup> | Powered up or retention <sup>1</sup>        | Powered up or retention <sup>1</sup> |
| ¾ DEBUG_RECOV      | 0b0111010           |                                 |                                      |  |                                      |   |                                      |
| ½ DEBUG_RECOV      | 0b0101010           |                                 |                                      |  | Powered down <sup>1</sup>            | Powered down <sup>1</sup>                   | Powered down <sup>1</sup>            |
| ¼ DEBUG_RECOV      | 0b0011010           |                                 |                                      |  |                                      |   |                                      |
| SFONLY_DEBUG_RECOV | 0b0001010           |                                 |                                      | Powered down <sup>1</sup>                              |                                      |   |                                      |

## 5.6 Power operating requirements

The DSU power operating requirements are indicated to the power controller by asserting bits on the CLUSTERPACTIVE bus.

The meaning of each bit is described in the following table.

**Table 5-4: L3 power domain CLUSTERPACTIVE bit positions**

| CLUSTERPACTIVE bit | Power operating requirement | Description  |
|--------------------|-----------------------------|--|
| [19]               | Cache ways 12-15            | <b>0</b> Cache ways are not required and can be requested to power down.<br><b>1</b> Cache ways are required to be active. |
| [18]               | Cache ways 8-11             |  |
| [17]               | Cache ways 4-7              |  |
| [16]               | Cache ways 0-3              |  |
| [8]                | On                          | DSU logic and RAMs are required to be powered up.  |
| [7]                | Functional retention        | DSU logic is active, RAMs can be put into retention.   |
| [2]                | Memory retention            | DSU logic can be off, RAMs required to be retained.  |
| [0]                | Off                         | Always 0. Request to power down DSU logic and RAMs should be accepted when the entire CLUSTERPACTIVE bus is zero.          |

<sup>1</sup> In DEBUG\_RECOV, the power mode does not need to accurately reflect the L3 portions that are powered up.



The following table shows how the CLUSTERPACTIVE bit combinations are mapped to power modes.

**Table 5-5: L3 memory system power domain requested modes**

| CLUSTERPACTIVE bit |                 |                |                |     |                      |                  |     | Requested mode         |
|--------------------|-----------------|----------------|----------------|-----|----------------------|------------------|-----|------------------------|
| [19]               | [18]            | [17]           | [16]           | [8] | [7]                  | [2]              | [0] |                        |
| Cache ways 12-15   | Cache ways 8-11 | Cache ways 4-7 | Cache ways 0-3 | On  | Functional retention | Memory retention | Off |                        |
| 1                  | X               | X              | X              | 1   | X                    | X                | X   | FULL_ON                |
| 0                  | 1               | X              | X              | 1   | X                    | X                | X   | $\frac{3}{4}$ ON       |
| 0                  | 0               | 1              | X              | 1   | X                    | X                | X   | $\frac{1}{2}$ ON       |
| 0                  | 0               | 0              | 1              | 1   | X                    | X                | X   | $\frac{1}{4}$ ON       |
| 0                  | 0               | 0              | 0              | 1   | X                    | X                | X   | SFONLY ON              |
| 1                  | X               | X              | X              | 0   | 1                    | X                | X   | FULL_FUNC_RET          |
| 0                  | 1               | X              | X              | 0   | 1                    | X                | X   | $\frac{3}{4}$ FUNC_RET |
| 0                  | 0               | 1              | X              | 0   | 1                    | X                | X   | $\frac{1}{2}$ FUNC_RET |
| 0                  | 0               | 0              | 1              | 0   | 1                    | X                | X   | $\frac{1}{4}$ FUNC_RET |
| 0                  | 0               | 0              | 0              | 0   | 1                    | X                | X   | SFONLY FUNC_RET        |
| 1                  | X               | X              | X              | 0   | 0                    | 1                | X   | FULL_MEM_RET           |
| 0                  | 1               | X              | X              | 0   | 0                    | 1                | X   | $\frac{3}{4}$ MEM_RET  |
| 0                  | 0               | 1              | X              | 0   | 0                    | 1                | X   | $\frac{1}{2}$ MEM_RET  |
| 0                  | 0               | 0              | 1              | 0   | 0                    | 1                | X   | $\frac{1}{4}$ MEM_RET  |
| X                  | X               | X              | X              | 0   | 0                    | 0                | X   | OFF                    |



Note

The CLUSTERPACTIVE and COREPACTIVE outputs are hints for the desired power mode. In some cases it might not be possible for the cluster to transition directly from the current mode to the requested mode. The power controller must be aware of the valid transitions and request transitions through intermediate modes if a direct transition is not valid. For example, if the cluster is in FUNC\_RET and CLUSTERPACTIVE indicates MEM\_RET, then the power controller must first request a transition from FUNC\_RET to ON, before requesting a transition from ON to MEM\_RET.

### 5.6.1 Power control for DFT

When DFT activities are being performed on the cluster, the P-Channel controls are not functional.

This means that if the cluster is in MBIST mode, or ATPG scan is in progress, then the CLUSTERPACTIVE or COREPACTIVE outputs can take any value and the P-Channel protocol will not be followed. The SoC power control components must be aware of the DFT activities and must ensure that the logic and RAMs are appropriately powered without using the P-Channel.

## 5.7 Wait For Interrupt and Wait For Event

To reduce dynamic power, each core can request entry in to a low-power state using the *Wait For Interrupt* (WFI) and *Wait For Event* (WFE) instructions.

In the low-power state, most of the clocks in a core are disabled while keeping the core powered up. This state reduces the power that is drawn to the static leakage current, leaving a small clock power overhead to enable the core to wake up.

In addition to the per-core WFI and WFE low-power states, the clock to (almost all) the SCU and L3 logic is automatically disabled when the cluster is sufficiently idle.

A `WFI` or `WFE` instruction completes when:

- All outstanding load instructions are completed.
- All store instructions are completed.
- All cache and TLB maintenance operations are completed.
- All bus traffic to the L3 is completed.

While a core is in the low-power state, the clocks in the core are temporarily enabled under the following conditions:

- A snoop request from the L3 cache that the L1 data cache or L2 unified cache must service.
- A cache or TLB maintenance operation that the core L1 caches or TLB must service.
- An APB access to the debug or trace registers residing in the core power domain.
- An access request from the GIC distributor to the GIC CPU interface.

While the clocks in the core are temporarily enabled, the core remains in the WFI or WFE low-power state.

### WFE wake up event signaling

- A *Send Event* (SEV) instruction signals a `WFE` wake up event to other clusters by asserting the `EVENTOREQ` output.
- The `EVENTIREQ` input indicates that another cluster or system component has signaled a `WFE` wake up event.

### System global exclusive monitor signaling

Any global exclusive monitor in the system must generate an event when it is cleared. This event must be signaled to the cluster using the `EVENTIREQ` input.

## 5.8 Clock, voltage, and power domains

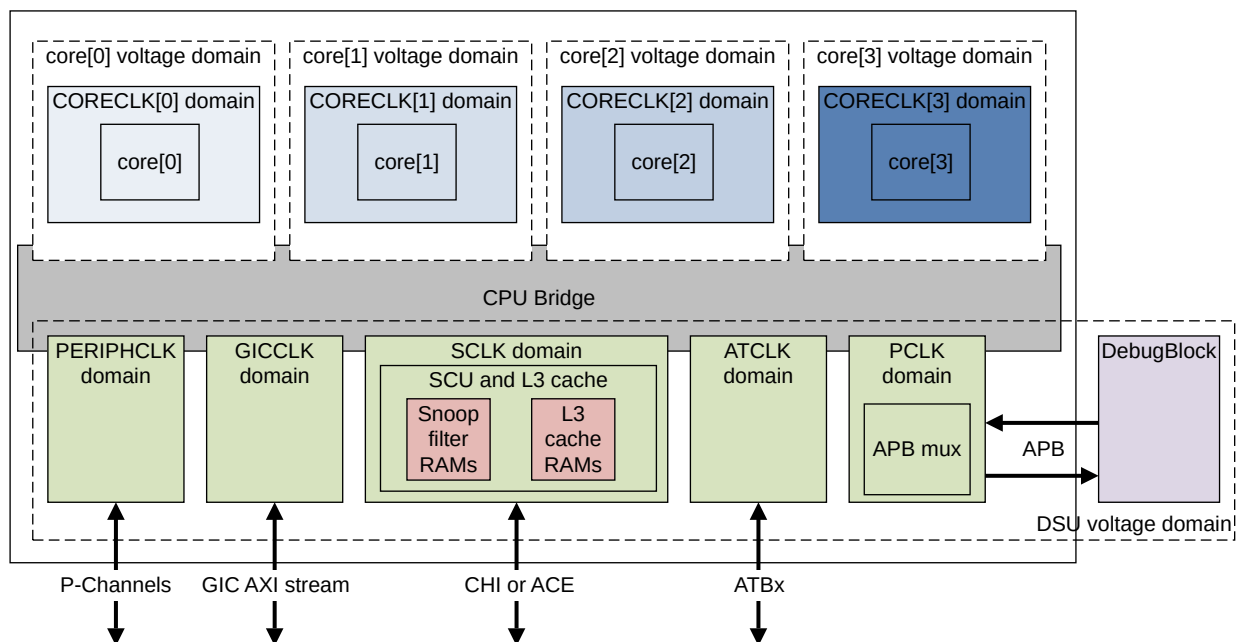
The DynamIQ™ cluster microarchitecture supports multiple clock, voltage, and power domains.

The number of domains that are implemented depends on the choices made by the SoC implementer. There might be fewer in your SoC.

The following diagram shows the clock, voltage, and power domains supported by the DSU and cores:

- Voltage domains are indicated by dashed outlines.
- Blocks that are in the same power domain have the same color.

**Figure 5-4: DSU Clock, voltage, and power domains**



### Clock domains

Each core can be implemented in a separate clock domain. The DSU has multiple clock domains.

The CPU Bridge contains all asynchronous bridges for crossing clock domains, and is split with one half of each bridge in the core clock domain and the other half in the relevant cluster domain. Each core can be implemented with or without an asynchronous bridge. If the asynchronous bridge is not implemented, the core is in the SCLK clock domain.

The DebugBlock can be implemented in the PCLK domain. However, the DebugBlock might be implemented in a different domain. In this case, asynchronous bridges must be implemented on the APB interfaces between the DebugBlock and the cluster.

Each clock domain, except the PERIPHCLK domain and CORECLK domains, has an associated Q-Channel, which allows the DSU to request the external clock controller to gate the clock on or off.

## Voltage domains

Each core can be implemented in a separate voltage domain. The DSU has a single separate voltage domain. This allows, for example, the DSU to be in the same voltage domain as the SoC interconnect and other system components.

## Power domains

Each core can be implemented in one, or more, separate power domains.

Additional power domains that can be implemented include:

- L3 cache RAM portions.
- SCU.
- DebugBlock.

## 5.9 Cluster powerdown

The cluster is taken out of coherence with the interconnect automatically when it is powered down. No software sequence is required.

After receiving the request to enter power off mode from the power controller, the DSU cleans and invalidates the L3 cache, and communicates with the interconnect to disable snoops into the cluster. All cores must be in the OFF mode before the cluster is powered down.



To support automated removal from coherency, the interconnect must support the SYSCOREQ and SYSCOACK handshake protocol. If not, the SoC is responsible for programming the interconnect to remove the cluster from coherency.


---

### 5.9.1 Transitioning in and out of coherency

The *DynamIQ™ Shared Unit* (DSU) provides a hardware mechanism for taking the cluster in and out of coherence with the system interconnect. The cluster enables coherency during powerup and disables it during powerdown.

The system interconnect can use the SYSCOREQ and SYSCOACK signals, to take the cluster in and out of coherence. If the system interconnect supports these signals, they can be connected directly. If the system interconnect does not support these signals, Arm recommends that they are connected to the power controller. In this case, the power controller must take any actions necessary to make the transition.

The following diagram shows the timing of SYSCOREQ and SYSCOACK.

**Figure 5-5: SYSCOREQ SYSCOACK four-phase coherency handshake**

To enable coherency, the cluster always asserts SYSCOREQ during powerup. When the system interconnect has enabled coherency, it asserts SYSCOACK and can then start sending snoop requests to the cluster. The cluster accepts snoop requests whenever either signal is asserted.

The cluster disables coherency during powerdown. The cluster deasserts SYSCOREQ and waits for the system interconnect to deassert SYSCOACK. The system interconnect must not deassert SYSCOACK until it can guarantee that: there are no further snoop requests to be sent, and that all snoop requests it has already sent have fully completed.

The signals must obey the following four-phase handshake rules:

- SYSCOREQ can only change when SYSCOACK is at the same level.
- SYSCOACK can only change when SYSCOREQ is at the opposite level.

#### 5.9.1.1 Coherency signals naming convention

The *DynamIQ™ Shared Unit* (DSU) supports up to two CHI or ACE interfaces. Each set of coherency signals has a unique suffix.

The coherency signals are named as follows:

##### **CHI master 0**

SYSCOREQ and SYSCOACK.

##### **CHI master 1**

SYSCOREQM1 and SYSCOACKM1.

##### **ACE master 0**

SYSCOREQM0 and SYSCOACKM0.

##### **ACE master 1**

SYSCOREQM1 and SYSCOACKM1.

## 6. L3 cache

This chapter describes the optional L3 cache.

### 6.1 About the L3 cache

All the cores in the cluster share the L3 cache.

The shared L3 cache of the DSU provides the following functionality:

- A dynamically optimized cache allocation policy.
- Groups of cache ways can be partitioned and assigned to individual processes<sup>2</sup>. Cache partitioning ensures that processes do not dominate the use of the cache to disadvantage other processes.
- Support for stashing requests from the ACP and CHI interfaces. These stashing requests can also target any of the L2 caches of cores within the cluster.
- *Error Correcting Code (ECC)* protection is supported on the L3 cache data and tag RAMs.
- The cache can be implemented with either one or two cache slices. Each cache slice consists of data, tag, victim, and snoop filter RAMs and associated logic.



- On powerdown, the DSU automatically performs cache cleaning, eliminating the need for software-controlled cache cleaning.
  - The ACE master interface does not support cache stashing.
- 

### 6.2 L3 cache allocation policy

The L3 cache data allocation policy changes depending on the pattern of data usage.

Exclusive allocation is used when data is allocated in only one core. Inclusive allocation is used when data is shared between cores.

For example, an initial request from core 0 allocates data in the L1 or L2 caches but not in the L3 cache. When data is evicted from core 0, the evicted data is allocated in the L3 cache. The allocation policy of this cache line is still exclusive. If core 0 refetches the line, it is allocated in the L1 or L2 caches of core 0 and removed from the L3 cache. The allocation policy of this cache line is still exclusive. If core 1 accesses this line for reading, it remains allocated in core 0 and is also allocated in both the core 1 and L3 caches. In this case, this line has an inclusive allocation, because it is being shared between cores.

---

<sup>2</sup> A process is an instance of a computer program.

## 6.3 L3 cache partitioning

The L3 cache supports a partitioning scheme that alters the victim selection policy to prevent a process from utilizing the entire L3 cache to the disadvantage of other processes.

Cache partitioning is intended for specialized software where there are distinct classes of processes running with different cache accessing patterns. For example, two processes A and B run on separate cores in the same cluster and therefore share the L3 cache. If process A is more data-intensive than process B, then process A can cause all the cache lines that process B allocates to be evicted. Evicting these allocated cache lines can reduce the performance of process B.

L3 cache partitioning is achieved by partition scheme IDs and groups of cache ways, where:

- Each group contains four ways.
- Each group can either be assigned as private to one or more partition scheme IDs, or be left unassigned.
- Each unassigned group can be shared between all eight partition scheme IDs.

Each core in the cluster must be assigned to at least one of the eight partition scheme IDs. L3 cache accesses from a given core can allocate into:

- Any cache way that belongs to a group that is assigned as private to the partition scheme ID of this core.
- Any cache way that belongs to an unassigned group that is shared by the entire cluster.

Up to four private L3 cache partitions can be created.



If some cache ways are powered down, the number of ways in each L3 cache partition are reduced. This reduction in cache ways can result in insufficient ways being made available to cores, which may degrade their performance. Therefore Arm recommends that caution is used when powering down cache ways while using cache partitioning.

---

### Partitioning setup example

The following example illustrates how the control registers can be programmed to partition the L3 cache into these three separate partitions:

- One partition that the hypervisor owns and uses for ACP and stashes.
  - Two partitions that the OS assigns to processes running on the cores.
1. Software, running at EL3, sets ACTLR\_EL3[10] to 0b1 and ACTLR\_EL3[11] to 0b1 to delegate control of the partitioning to EL2.
  2. The hypervisor, running at EL2, sets the CLUSTERPARTCR\_EL1 to 0x00008601 to configure these three L3 cache partitions:
    - Scheme ID 0 (1/4 of the cache).
    - Scheme ID 2 (1/2 of the cache).

- Scheme ID 3 (1/4 of the cache).
3. The hypervisor sets the CLUSTERACPSID\_EL1 to 0x0 and the CLUSTERSTASHSID\_EL1 to 0x0, to make ACP requests and stashes use partition scheme ID 0.
  4. The hypervisor sets the CLUSTERTHREADSIDOVR\_EL1 to 0x00060002. This setting indicates that:
    - The upper two bits of the scheme ID are under the control of the hypervisor.
    - The lowest bit of the scheme ID can be controlled by the OS.
  5. The hypervisor sets ACTLR\_EL2[10] to 0b1 to delegate control of the CLUSTERTHREADSID\_EL1 register to EL1.
  6. The OS, running at EL1, can set the CLUSTERTHREADSID\_EL1 to 0x0 or 0x1 to select between the two L3 cache partitions allocated to it. These would be mapped to scheme IDs 2 and 3 by the CLUSTERTHREADSIDOVR\_EL1 register. The OS can update this register on context switches to select which partition each process has access to.

## 6.4 Cache stashing

Cache stashing allows an external agent to request that a line is brought (or stashed) into a cache.

Cache stashing can either be performed over the ACP interface, or the CHI master interface. Stash requests can target the L3 cache, or any of the L2 caches of cores within the cluster. The available stashing bandwidth is likely to be higher when stashing to the L3 cache.

On the CHI interface, stash requests (snoops) into both the L2 and L3 caches are supported by default. The field, StashLPIDValid, indicates the target of the stash, as follows:

- If the field is clear, then the stash is directed to the L3 cache.
- If this field is set, then the stash is directed to an L2 cache and the StashLPID field indicates which core is targeted.

On the ACP, by default, accesses are implicit stash requests into the L3 cache. Signal AWSTASHLPIDENS indicates a stash into L2. In this case, signal AWSTASHLPIDS[3:1] indicates which core to target, and AWSTASHLPIDS[0] indicates the thread. The choice of thread has no effect on the stash operation.

On the ACE master interface, cache stashing is not supported.

The cluster always attempts to allocate a stash request, unless it is heavily utilized and does not have any free buffers. In this case, the cluster drops a stash request to avoid a potential system deadlock.



Note

If there is a stashing snoop targeting an L3 (rather than a core) and there is no L3 cache, then the DSU will not perform the data pull, so the DSU will not send a SnpResp\*\_Read response. In this case, the DSU can send different responses for the different stashing snoops as the follows:



- For SnpStashShared, SnpStashUnique or SnpMakeInvalidStash, the DSU returns SnpResp\_I response.
  - For SnpUniqueStash, the DSU returns SnpRespData\_I or SnpResp\_I.
- 

## 6.5 L3 cache ECC and parity

Error Correcting Code (ECC) and parity protection is supported on L3 cache RAM instances.

The ECC algorithm that is used is *Single Error Correct, Double Error Detect* (SECEDED). It allows detection and correction of any 1-bit error and detection of any 2-bit error in all protected RAMs.

L3 cache RAMs implement ECC, as follows:

- SECEDED ECC on the L3 data RAMs, per 64-bit.
- SECEDED ECC on the L3 tag RAMs, per entry.
- SECEDED ECC on the snoop filter RAMs, per entry.

If an error has no functional effect and only results in a minor change in performance, then the affected RAM is not protected. For example, victim selection RAMs are not protected with ECC or parity.

### Error correction

When a correctable error is detected in the L3 cache data RAMs, the data is corrected inline before returning to the requester.

When a correctable error is detected in the L3 cache tag RAMs or the snoop filter RAMs the following correction mechanism is used:

- The value is corrected and written back to the source address (Read-Correct-Write).
- The lookup is replayed.

The DSU includes hardware that provides limited support for hard error correction. A hard error is a physical error in the RAM that prevents the correct value from being written. Hard errors can be corrected, but might cause a degradation in performance when the locations with errors are accessed.

### Uncorrectable errors and data poisoning

If an error is detected as having two bits in error in a RAM protected by ECC, then this error is not correctable. In this case, the behavior depends on the type of RAM, as follows:

#### Data RAM or Long-Term Data Buffer RAM

When an uncorrectable error is detected in an L3 data RAM or *Long-Term Data Buffer* (LTDB) RAM, the chunk of data with the error is marked as poisoned. This poison status is then transferred with the data and stored:

- In the cache, if the data is allocated back into a cache.

- In the LTDB RAM, if the data is moved there.

The poison status is stored for every 64 bits of data.

If the interconnect supports poisoning, then the poison status is transferred with the data when the line is evicted or snooped from the cluster. No abort is generated when a line is poisoned. The abort is deferred until a load or instruction fetch consumes the poisoned data.

If the interconnect does not support poisoning and a poisoned cache line is evicted or snooped from the cluster, then the DSU generates an interrupt, nERRIRQ, to notify software that data has potentially been lost.



Software can indicate if the interconnect supports poisoning or not by setting the interconnect data poisoning support bit in the Extended Control Register of the cluster. For details, see [11.8 CLUSTERECTLR, Cluster Extended Control Register](#) on page 105.

## Tag RAM

When an uncorrectable error is detected in an L3 tag RAM, then either the address or coherency state of the line is unknown, so the data cannot be poisoned. In this case, the line is invalidated and the DSU generates an interrupt, nERRIRQ, to notify software that data has potentially been lost.

## Snoop filter tag RAM

When an uncorrectable error is detected in a snoop filter tag RAM, either the address or coherency state of the line is unknown, so the data cannot be poisoned. In this case, the snoop filter entry is invalidated, but the line remains present in one or more of the cores. The DSU generates an interrupt, nERRIRQ, to notify software that data has potentially been lost.



Arm recommends that a system reset is performed as soon as possible, in response to this interrupt. Because the core caches and the snoop filter are inconsistent after this error, which can lead to **UNPREDICTABLE** behavior. The effect of the error depends on the type of core, but it could result in further data corruption, or deadlocks, making it impossible to cleanly recover from such an error.

## Error reporting

Any detected error is reported in the *Error Record Primary Status Register* (ERR<n>STATUS) and the *Error Record Miscellaneous Register 0* (ERR<n>MISCO). These detected errors include errors that are successfully corrected and errors that cannot be corrected. If multiple errors occur on the same clock cycle, then only one error is reported and the OF (overflow) bit of ERR<n>STATUS is set.

The ERRSELR\_EL1 register provides the following two error records:

- Record 0 is private to the core and is updated on any error in the core RAMs including L1 caches, TLB, and L2 cache.

- Record 1 is for the cluster, which is shared between all cores in the cluster, and records any error in the L3 and snoop filter RAMs.

If enabled in the ERR<n>CTLR register (by setting one or more of the UI, FI, or CFI bits), all the detected errors generate a fault handling interrupt on:

- The nFAULTIRQ[0] pin (for L3 and snoop filter errors)
- The nFAULTIRQ[n+1] pin (for core n L1 and L2 errors)

Errors that cannot be corrected, also cause an abort or error handling interrupt because they can result in data corruption. The interrupt alerts software to this error so that it can either attempt to recover or restart the system. Some errors can be deferred by poisoning the data. These errors do not cause an abort at the time of the error, but only when the erroneous data is consumed.

The following describes the different types of errors that can occur and their effects:

- Uncorrectable errors in the L3 data RAMs when read by a core can cause a precise or imprecise Data Abort or Prefetch Abort, depending on the implementation of the core.
- Uncorrectable errors in the L3 data RAMs in a line when this line is being evicted from a cache, cause the data to be poisoned. The eviction might be because of a natural eviction, a line fill from a higher level of cache, a cache maintenance operation, or a snoop. If the poisoned line is evicted from the cluster for any reason and the interconnect does not support data poisoning, then the nERRIRQ[0] pin is asserted.
- Uncorrectable errors in the L1 tag or dirty RAMs, or in the L2 tag RAMs, cause the nERRIRQ[n+1] pin to be asserted for core<n>.
- Uncorrectable errors in the L3 tag RAMs or SCU snoop filter RAMs cause the nERRIRQ[0] pin to be asserted.



Arm recommends that the nERRIRQ pins are connected to the interrupt controller, so that an interrupt or system error is generated when the pin is asserted.

---

The fault and error interrupt pins can be cleared by writing to the ERR<n>STATUS registers.

When a dirty cache line with an error on the data RAMs is evicted from the cluster, the write on the master interface still takes place. However, if the error is uncorrectable then:

- On ACE, the uncorrected data is written.
- On CHI, the uncorrected data is written but the data poison field indicates that there is a data error.

When a snoop hits on a line with an uncorrectable data error, the data is returned (if the snoop requires the data) and:

- On ACE, the nERRIRQ[0] pin is asserted.



The snoop response does not indicate the error.

- On CHI, the snoop response indicates that either the data is poisoned, when the interconnect supports poisoning, or that there is an error.

If a snoop hits on a tag that has an uncorrectable error, then it is treated as a snoop miss. Because the error means that it is not known whether the cache line is valid.

If an ACP access reads a cache line with an uncorrectable error, then it returns an ACP response to indicate a slave error.

Sometimes an error can be counted multiple times. For example, multiple accesses might read the location with the error before the line is evicted.

### Error injection

Error injection is used to test out the error detection reporting and recording structure by deliberately inserting errors into the error reporting logic.

The injected errors are only pseudo-errors. They cause an error report to be signaled but the error injection does not corrupt the target location. Therefore, an injected pseudo error does not cause any automatic error correction logic to be activated.

Error injection uses the error injection and reporting registers to insert errors. Pseudo-errors are injected using the ERR1PFGCTLR, Error Pseudo Fault Generation Control Register.

Corrected and deferred pseudo-errors are triggered with accesses to *Long-Term Data Buffer* (LTDB) RAMs. Uncontainable pseudo-errors are triggered with accesses to the level 3 tag RAM instances. If the cluster is configured without a level 3 cache it will not be possible to trigger pseudo-errors in the DSU cluster.

### Related information

[Control registers](#) on page 95

## 6.6 L3 cache data RAM latency

The L3 data RAM interface can be implemented with a configurable latency on the input and output paths.

The following options are available:

- Either a 1-cycle (the default) or 2-cycles write latency on the input path to the L3 data RAMs.
- Either a 2-cycles (the default) or 3-cycles read latency on the output path from the L3 data RAMs.

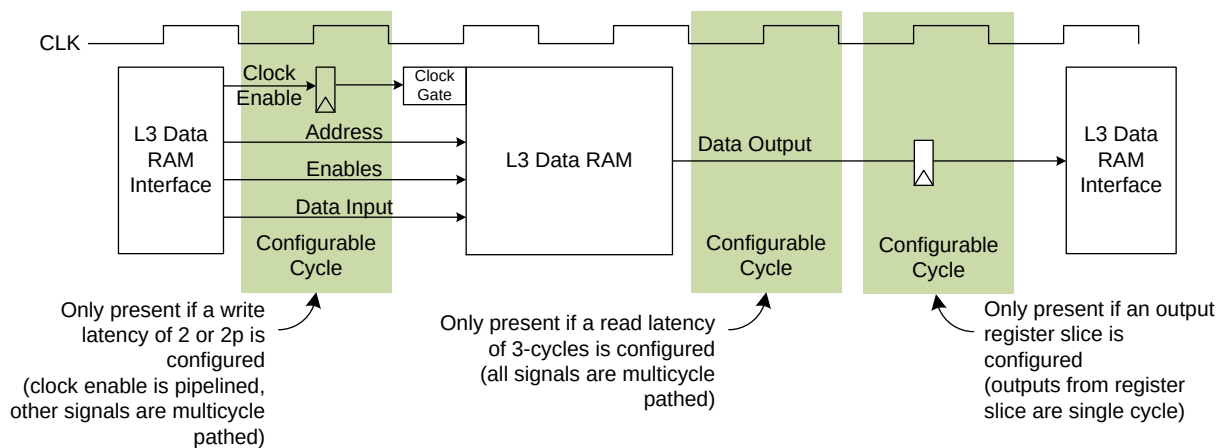
- An optional register slice on the output of the L3 data RAMs.

On the input paths, the 2p write latency keeps the RAM input signals stable for an extra cycle, which allows an extra cycle of hold timing on the RAM inputs. If a 2 or 2p write latency is requested on the input paths, then the RAM clock enable is pipelined and a multicycle path is applied to all other RAM input signals.

On the output paths, the 2-cycles read latency and 3-cycles read latency applies a multicycle path to all RAM output signals. The output of the optional register slice is single cycle and must never have a multicycle path applied.

The following diagram shows the L3 data RAM timing.

**Figure 6-1: L3 cache data RAM latency**



An increase in RAM latency increases the L3 hit latency, which reduces performance.



Note

- Only use the 3-cycles read latency option if the RAM cannot meet the timing requirement of the 2-cycles latency. When the read latency is set to 3-cycles, the decrease in lookup throughput to one access every three core clock cycles is possibly more significant. Because when there are a series of L3 data RAM accesses close together, then the memory system could begin to back up.
- Use the register slice if the wire routing delay from the RAM to the SCU logic cannot meet the timing requirements.

The following table describes the impact on L3 data RAM performance with the different latency configuration parameters.

**Table 6-1: L3 data RAM performance with different latency configurations**

| L3_DATA_WR_LATENCY | L3_DATA_RD_LATENCY | L3_DATA_RD_SLICE | L3 hit latency cycles | L3 lookup bandwidth   | L3 write bandwidth    |
|--------------------|--------------------|------------------|-----------------------|-----------------------|-----------------------|
| 1                  | 2                  | No               | 3                     | Access every 2-cycles | Access every 2-cycles |


| L3_DATA_WR_LATENCY | L3_DATA_RD_LATENCY | L3_DATA_RD_SLICE | L3 hit latency cycles | L3 lookup bandwidth   | L3 write bandwidth    |
|--------------------|--------------------|------------------|-----------------------|-----------------------|-----------------------|
| 1                  | 3                  | No               | 4                     | Access every 3-cycles | Access every 2-cycles |
| 1                  | 2                  | Yes              | 4                     | Access every 2-cycles | Access every 2-cycles |
| 1                  | 3                  | Yes              | 5                     | Access every 3-cycles | Access every 2-cycles |
| 2                  | 2                  | No               | 4                     | Access every 2-cycles | Access every 2-cycles |
| 2                  | 3                  | No               | 5                     | Access every 3-cycles | Access every 2-cycles |
| 2                  | 2                  | Yes              | 5                     | Access every 2-cycles | Access every 2-cycles |
| 2                  | 3                  | Yes              | 6                     | Access every 3-cycles | Access every 2-cycles |
| 2p                 | 3                  | No               | 5                     | Access every 3-cycles | Access every 3-cycles |
| 2p                 | 3                  | Yes              | 6                     | Access every 3-cycles | Access every 3-cycles |

## 6.7 Cache slices and portions

The DSU is implemented as either having one or two L3 cache slices. A cache slice consists of data, tag, victim, and snoop filter RAMs and associated logic. A portion is a further subdivision of RAM in a cache slice.

When two L3 cache slices are implemented, the overall cache is effectively divided in two. Each cache slice has their own associated logic. But they are not independent because there is some shared logic between them. For each cache slice, the data RAM is subdivided into two portions, while the tag RAM is subdivided into four portions.

The following figure shows the differences between a single and a dual cache slice configuration.

**Figure 6-2: Comparison between a single and dual L3 cache slice configuration**

Dividing the L3 cache into two slices provides the following advantages:

- Improving the physical floorplan when implementing the macrocell, particularly for larger cache sizes.
- Increasing the bandwidth because the two slices can be accessed in parallel.

### 6.7.1 Cache slice and master port selection

For a dual cache slice implementation, requests are sent to a particular slice depending on the address and the memory attributes:

- For Normal Non-cacheable requests, the behavior depends on CLUSTERECTLR[0]. See [11.8 CLUSTERECTLR, Cluster Extended Control Register](#) on page 105.
- For Cacheable requests, addresses are interleaved between slice 0 and slice 1, based on an address bit set by the INTERLEAVE\_ADDR\_BIT configuration parameter.
- Device requests are always sent to slice 0.

In a configuration with dual master ports, the slices directly correspond to the master ports, so an access sent to slice 0 uses master port 0.

See [7.1.1 Dual ACE interfaces](#) on page 66 for more information.

### 6.7.2 Default number of cache slices

The configuration parameters that determine the default number of L3 cache slices.

Two cache slices are implemented, by default, when any of the following configuration options are chosen:

- More than four LITTLE cores are configured.
- One or more big cores are configured.

- The L3 cache size is greater than 1MB.
- A second ACE master port is configured.
- Either one or two 256-bit CHI master ports is configured.

A single slice is implemented by default for all other configurations.



When there is either no L3 cache or an L3 cache that is 256KB, 512KB, or 1MB in size and there is only a single master port, then you can override the default number of L3 cache slices by using the `NUM_SLICES` configuration parameter. For more information, see [2.3 Implementation options](#) on page 18.

### 6.7.3 Implementing a 1.5MB or 3MB L3 cache

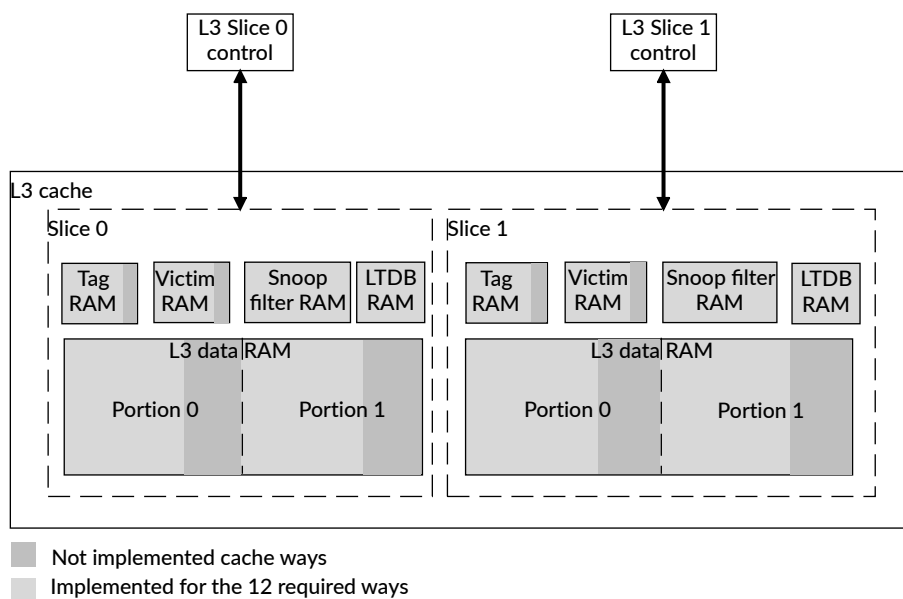
When selecting a non-power-of-two L3 cache size of 1.5MB or 3MB, each cache slice is only implemented with 12 ways. Unlike the power-of-two 256KB, 512KB, 1MB, 2MB, or 4MB L3 cache sizes that implement 16 ways.

For a non-power-of-two L3 cache size:

- The L3 cache size that is reported to software is the next higher power-of-two cache size. For instance, software reports the 1.5MB L3 cache size as being 2MB and the 3MB L3 cache size as being 4MB.
- The last  $\frac{1}{4}$  of the tag, victim, and data RAMs are not implemented in each cache slice. Therefore, partially powering down the 1.5MB or 3MB L3 cache from Full to  $\frac{3}{4}$  has no effect on the power consumption of the DSU because these ways are not implemented.

The following diagram shows the 1.5MB or 3MB L3 cache implementation that is compared with a 2MB or 4MB cache implementation respectively.



**Figure 6-3: 1.5MB or 3MB L3 cache implementation**

For more information on the cache partial powerdown, see [5.4.1 L3 cache partial powerdown](#) on page 39.

## 7. ACE master interface

This chapter describes the ACE master memory interface.

### 7.1 About the ACE master interface

You can configure the DynamIQ™ Shared Unit to use the ACE5 protocol for the master memory interface.

#### 7.1.1 Dual ACE interfaces

The DSU can be implemented with one or two ACE interfaces. Two interfaces give greater bandwidth for memory transactions.

Transactions generated by the cluster are routed to either interface based on the transaction type, memory type, and transaction address:

- All DVM transactions are routed to interface 0.
- All Device transactions are routed to interface 0.

For Cacheable transactions, the configuration parameter `INTERLEAVE_ADDR_BIT` controls which transaction address bit is used to select the routing between interface 0 and interface 1. The default value is to select bit 6 of the transaction address, which interleaves on cache line boundaries. If the selected transaction address bit has a value 0, then interface 0 is used, otherwise interface 1 is used. For more information on `INTERLEAVE_ADDR_BIT` and the impact on performance, see [2.3 Implementation options](#) on page 18.

For Normal Non-cacheable transactions, routing is dependent on the value of `CLUSTERECTLR.Non-cacheable behavior control, bit[0]`:

- If `CLUSTERECTLR[0]` is set to 1, Normal Non-cacheable transactions are routed to both interface 0 and 1 in the same way as Cacheable transactions, using the same bit of the transaction address.
- If `CLUSTERECTLR[0]` is set to 0, Normal Non-cacheable transactions are routed to interface 0.



Setting `CLUSTERECTLR.Non-cacheable behavior control` has other implications for the system. See [11.8 CLUSTERECTLR, Cluster Extended Control Register](#) on page 105 for more details.

---

When the external memory system sends snoops, it must either:

- Send the snoop to both interfaces.

- Send the snoop only to the interface that is relevant for the address of that snoop. This behavior is normal operation for an external memory system that contains a snoop filter. The snoop filter indicates that the line is present in one of the two masters.

The second method is more efficient, and if two masters are implemented, Arm recommends that the external memory system includes a snoop filter.

### DVM messages

DVM messages can be received on both interfaces, however they are only required on interface 0, and any DVM message sent to interface 1 is treated as a no-op. Therefore for best performance, Arm recommends that, when possible, your interconnect is configured to avoid sending DVM messages to interface 1.

## 7.2 ACE configurations

The following table shows the supported ACE configurations.

**Table 7-1: Supported ACE configurations**

| Signal              | Feature                                 |                           |   |                           |
|---------------------|---|---------------------------|---|---------------------------|
|                     | ACE non-coherent                        |                           | ACE coherent                            |                           |
|                     | With no cache or invisible system cache | With visible system cache | With no cache or invisible system cache | With visible system cache |
| BROADCASTCACHEMAINT | 0                                       | 1                         | 0                                       | 1                         |
| BROADCASTOUTER      | 0                                       | 0                         | 1                                       | 1                         |



Note

- ACE non-coherent mode, no system cache, can be used to connect to an AXI interconnect.
- A visible system cache requires cache maintenance transactions to ensure that a write is visible to all observers.
- An invisible system cache is one that does not require cache maintenance transactions to ensure that a write is visible to all observers. This is true even if those observers use different memory attributes.

The following table shows the key features in each of the supported ACE configurations.

**Table 7-2: Supported features in the ACE configurations**

| Features                             | DSU Configuration |                  |              |
|--------------------------------------|-------------------|------------------|--------------|
|                                      | ACE non-coherent  | ACE non-coherent | ACE coherent |
|                                      | No system cache   | System cache     |              |
| AXI3 or AXI4 interconnect compliance | Yes               | No               | No           |
| ACE interconnect compliance          | Yes               | Yes              | Yes          |
| Barriers on AR and AW channels       | No                | No               | No           |

| Features                                 | DSU Configuration |                  |              |
|--|-------------------|------------------|--------------|
|  | ACE non-coherent  | ACE non-coherent | ACE coherent |
|  | No system cache   | System cache     |              |
| Cache maintenance requests on AR channel | No                | Yes              | Yes          |
| Snoops on AC channel                     | No                | No               | Yes          |
| Coherent requests on AR or AW channel    | No                | No               | Yes          |
| DVM requests on AR channel               | No                | No               | Yes          |

## 7.3 ACE features

AMBA defines a set of interface properties for the ACE interconnect. The following table shows which of these properties the DSU supports, or requires the cluster interconnect and system to support.

**Table 7-3: ACE interconnect properties for the DSU**

| ACE property                    | Supported by the DSU  | Interconnect support required  |
|---------------------------------|---|--|
| Continuous_Cache_Line_Read_Data | Not applicable  | Yes  |
| Multi_Copy_Atomicity            | Yes   | Yes  |
| Ordered_Write_Observation       | Not applicable  | No   |
| WriteEvict_Transaction          | Yes if CLUSTERECTLR Cache UniqueClean eviction control is programmed to 1.  | Yes if CLUSTERECTLR Cache UniqueClean eviction control is programmed to 1. |
| DVM_v8                          | Yes   | Yes if BROADCASTOUTER is HIGH.   |
| Atomic_Transactions             | Yes, for atomic instructions that access cacheable memory.<br><br>No, for atomic instructions that access Device or non-cacheable memory. | No   |
| DVM_v8.1                        | Yes   | Yes if BROADCASTOUTER is HIGH.   |
| Cache_Stash_Transactions        | No  | No   |
| DeAllocation_Transactions       | No  | No   |
| Persistent_CMO                  | No  | No   |
| Poison                          | No  | No   |
| Data_Check                      | No  | No   |
| QoS_Accept                      | No  | No   |
| Trace_Signals                   | No  | No   |
| Loopback_Signals                | No  | No   |
| Low_Power_Signals               | Yes   | Yes  |
| Untranslated_Transactions       | No  | No   |
| NSAccess_Identifiers            | No  | No   |

## 7.4 ACE master interface attributes

This section describes the read and write issuing capabilities and ID encoding.

The following table lists the read and write issuing capabilities.

**Table 7-4: ACE master interface attributes**

| Attribute                          | Value                      | Comments   |
|------------------------------------|----------------------------|--|
| Write issuing capability           | Configuration dependent    | <p>The maximum number of writes is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul> <p>See <a href="#">6.7 Cache slices and portions</a> on page 62.</p> <p>Device and Normal Non-cacheable transactions are limited to a total of 15 write transactions by default. This value can be used by system components to size buffers when bridging to other interface protocols, for example PCIe. Normal Non-cacheable transactions can be removed from this limit by setting the CLUSTERECTLR.Noncacheable behavior control, bit[0] to 1. See <a href="#">11.8 CLUSTERECTLR, Cluster Extended Control Register</a> on page 105.</p>   |
| Read issuing capability            | Configuration dependent    | <p>The maximum number of reads is:</p> <ul style="list-style-type: none"> <li>98, if two slices are present.</li> <li>34, if one slice is present.</li> </ul> <p><b>Note:</b></p> <ul style="list-style-type: none"> <li>Two-part <i>Distributed Virtual Memory</i> (DVM) messages use the same ID for both parts, and therefore can have two outstanding transactions on the same ID.</li> <li>For Device and Normal Non-cacheable reads, the read issuing capability is limited by the combined issuing capability. Because the combined issuing capability is always lower than the read issuing capability.</li> </ul>   |
| Combined issuing capability        | Configuration dependent    | <p>The combined issuing capability is:</p> <ul style="list-style-type: none"> <li>98, if two slices are present.</li> <li>34, if one slice is present.</li> </ul> <p>The Device combined issuing capability is limited to:</p> <ul style="list-style-type: none"> <li>39, if two slices are present.</li> <li>23, if one slice is present.</li> </ul> <p>The Device and Normal Non-cacheable combined issuing capability is limited to:</p> <ul style="list-style-type: none"> <li>78, if two slices are present and the CLUSTERECTLR.Noncacheable behavior control, bit[0] is set to 1.</li> <li>39, if two slices are present and the CLUSTERECTLR.Noncacheable behavior control, bit[0] is set to 0.</li> <li>23, if one slice is present.</li> </ul> |
| Exclusive access thread capability | Number of hardware threads | Each hardware thread can have 1 exclusive access sequence in progress.   |

| Attribute           | Value                   | Comments  |
|---------------------|-------------------------|---|
| Write ID capability | Configuration dependent | <p>The maximum write ID capability is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul> <p>Only Device memory types with nGnRnE or nGnRE can have more than one outstanding transaction with the same AXI ID. All other memory types use a unique AXI ID for every outstanding transaction.</p>   |
| Write ID width      | 8                       | The ID encodes the source of the memory transaction. See the Encodings for AWIDM0[7:0] table.   |
| Read ID capability  | Configuration dependent | <p>The maximum read ID capability is:</p> <ul style="list-style-type: none"> <li>98, if two slices are present.</li> <li>34, if one slice is present.</li> </ul> <p>Only Device memory types with nGnRnE or nGnRE can have more than one outstanding transaction with the same AXI ID. All other memory types use a unique AXI ID for every outstanding transaction.</p> <p>Two part DVMs use the same ID for both parts, and therefore can have two outstanding transactions on the same ID.</p> |
| Read ID width       | 9                       | The ID encodes the source of the memory transaction. See the Encodings for ARIDM0[8:0] table.   |



Note

The issuing capability described here is the maximum possible for the whole cluster. This can be used to size interconnect capabilities if you want to achieve the maximum performance available. However, this maximum may not be able to be reached by a single core on its own. It may need multiple cores generating heavy memory traffic simultaneously to reach it. The capabilities will vary by core type, for example big cores will typically generate more transactions than LITTLE cores. It can also vary by memory type, with typically a significantly lower limit for Device or Non-cacheable transactions than for Cacheable transactions.

The following table shows the encodings for AWIDM0[7:0], WIDM0[7:0]. When two ACE masters are configured, the maximum number of reads and writes are unchanged. The reads and writes can be distributed between the two masters, or all send to one of the masters, depending on the memory type and address.

**Table 7-5: Encodings for AWIDM0[7:0] and WIDM0[7:0]**

| Attribute | Value                   | Issuing capability per ID | Comments   |
|-----------|-------------------------|---------------------------|--|
| Write ID  | 0b000t0nnn              | 1                         | Core nnn <sup>3</sup> , thread t <sup>4</sup> , system domain store exclusives (except for those that are Device non-reorderable). |
|           | 0b001t0nnn              | 15                        | Core nnn <sup>3</sup> thread t <sup>4</sup> , non-reorderable Device writes.   |
|           | 0b1xxxxxxx <sup>5</sup> | 1                         | All other types of write.  |
|           | Other encodings         | -                         | Not used   |

The following table shows the Encodings for ARIDM0[8:0].

**Table 7-6: Encodings for ARIDM0[8:0]**

| Attribute | Value                    | Issuing capability per ID | Comments   |
|-----------|--------------------------|---------------------------|--|
| Read ID   | 0b0000t0nnn              | 1                         | Core nnn <sup>3</sup> , thread t <sup>4</sup> , load exclusives (except for those that are Device non-reorderable), and Cacheable Shareable store exclusives (sent as exclusive CleanUnique transactions). |
|           | 0b0001t0nnn              | 17                        | Core nnn <sup>3</sup> , thread t <sup>4</sup> , non-reorderable Device reads.  |
|           | 0b001000000              | 1                         | DVM Sync   |
|           | 0b001000001              | 256                       | DVM Complete   |
|           | 0b01xxxxxxx <sup>5</sup> | 1                         | All other types of read  |
|           | 0b1xxxxxxx <sup>5</sup>  |                           |  |
|           | Other encodings          | -                         | Not used   |



- These ID and transaction details are provided for information only. Arm strongly recommends that all interconnects and peripherals are designed to support any type and number of transactions on any ID, to ensure compatibility with future products.
- The Device and Normal Non-cacheable transaction limits that are specified in [Table 7-4: ACE master interface attributes](#) on page 69 apply.

For more information about the ACE and AXI signals that are described in this manual, see the [AMBA® AXI and ACE Protocol Specification](#).

## Related information

[Cache slices and portions](#) on page 62

<sup>3</sup> nnn is the core number 0b000-0b111 in binary.

<sup>4</sup> t is the hardware thread number and is 0 if the core does not support multiple hardware threads.

<sup>5</sup> x is a do not care value, can be 0 or 1.

## 7.5 ACE channel properties

The following table shows the properties of the ACE channels.

**Table 7-7: ACE channel properties**

| Property                    | Value                  | Comment   |
|-----------------------------|------------------------|---|
| Snoop acceptance capability | 9 per master interface | The SCU can accept and process a maximum of nine snoop requests from the system for each ACE master interface. It counts requests from the request being accepted on the AC channel to the response being accepted on the CR channel.   |
| Snoop latency               | Hit                    | When there is a hit in L3 cache, the best case for response and data is 10 SCLK cycles. When there is a miss in the L3 cache but a hit in an L1 or L2 cache in a core, then the latency varies depending on the type and configuration of the core.<br><br><b>Note:</b><br>Latencies can be higher if hazards occur or if there are not enough buffers to accept requests.  |
|                             | Miss                   | Best case six SCLK cycles when the snoop filter and L3 cache tags indicate the miss.  |
|                             | DVM                    | The cluster takes a minimum of six SCLK cycles to provide a response to DVM packets.  |
| Snoop filter                | Supported              | The cluster provides support for an external snoop filter in an interconnect. It indicates when clean lines are evicted from the cluster by sending Evict transactions on the write channel.<br><br>However there are some cases that can prevent an Evict transaction from being sent. Therefore you must ensure that you build any external snoop filter to handle a capacity overflow. When exceeding capacity, the snoop filter should send a back-invalidation to the cluster.<br><br>Examples of cases where evicts are not produced include: <ul style="list-style-type: none"> <li>• Linefills that take external aborts.</li> <li>• Store exclusives that fail.</li> <li>• Mis-matched aliases.</li> </ul> |
| Supported transactions      | -                      | All transactions described by the ACE protocols: <ul style="list-style-type: none"> <li>• Are accepted on the master interface from the system.</li> <li>• Can be produced on the ACE master interface except: <ul style="list-style-type: none"> <li>◦ Barriers</li> <li>◦ MakeInvalid</li> <li>◦ ReadShared</li> <li>◦ ReadOnceCleanInvalid</li> <li>◦ ReadOnceMakeInvalid</li> <li>◦ StashOnceShared</li> <li>◦ StashOnceUnique</li> <li>◦ StashTranslation</li> <li>◦ WriteUniquePtl.<sup>6</sup></li> <li>◦ WriteUniqueFull.<sup>6</sup></li> <li>◦ WriteUniquePtlStash</li> <li>◦ WriteUniqueFullStash</li> </ul> </li> </ul>   |



See the [AMBA® AXI and ACE Protocol Specification](#) for more information about the ACE channel.

## 7.6 ACE transactions

The DSU does not generate any FIXED bursts and a burst does not cross a cache line boundary.

The cache linefill fetch length is always 64 bytes.

The DSU generates only a subset of all possible ACE transactions on the master interface.

For WriteBack Cacheable transfers, the supported transfers are:

- WRAP 4 128-bit for read transfers (linefills).
- INCR 4 128-bit for write transfers (evictions).
- INCR 4 128-bit for read transfers (linefills).
- INCR 1 128-bit for read transfers if ACP is configured.

For Normal Non-cacheable or Device transactions:

- INCR N (N:2 or 4) 128-bit read transfers.
- INCR N (N:2 or 4) 128-bit write transfers.
- WRAP N (N:2 or 4) 128-bit read transfers.
- INCR 1 8-bit, 16-bit, 32-bit, 64-bit, and 128-bit read transfers.
- INCR 1 8-bit, 16-bit, 32-bit, 64-bit, and 128-bit write transfers.
- INCR 1 8-bit, 16-bit, 32-bit, 64-bit, and 128-bit exclusive read transfers.
- INCR 1 8-bit, 16-bit, 32-bit, 64-bit, and 128-bit exclusive write transfers.

The following points apply to ACE transactions:

- WRAP bursts are only 128-bit size.
- INCR burst, more than one transfer, are only 128-bit size.
- No transaction is marked as FIXED.
- Write transfers with none, some, or all byte strobes LOW can occur.

The following table shows the ACE transactions that can be generated and some typical operations that might cause these transactions to be generated.

<sup>6</sup> The AMBA ACE5 transaction types WriteUniqueFull and WriteUniquePtl were known in AMBA4 ACE as WriteLineUnique and WriteUnique, respectively.



This table does not provide an exhaustive list of operations that generate each type of transaction, because there are many possibilities.

**Table 7-8: ACE transactions**

| Transaction          | Operation   |
|----------------------|---|
| ReadNoSnoop          | Non-cacheable loads or instruction fetches.<br>Linefills of non-shareable cache lines into L1, L2, or L3 caches.  |
| ReadOnce             | Cacheable loads that are not allocating into the cache.   |
| ReadClean            | Cache data linefills started by a load instruction.<br>Cache linefills started by an instruction fetch.   |
| ReadNotSharedDirty   | Cache data linefills started by a load instruction.<br>Cache linefills started by an instruction fetch.   |
| ReadUnique           | Data linefills started by a store instruction.  |
| CleanUnique          | Store instructions that hit in the cache but the line is not in a unique coherence state.<br>Store instructions that are not allocating into the caches, for example when streaming writes. |
| MakeUnique           | Store instructions of a full cache line of data, that miss in the caches.   |
| CleanShared          | Cache maintenance instructions.   |
| CleanSharedPersist   | Not used.   |
| CleanInvalid         | Cache maintenance instructions.   |
| DVM                  | TLB and instruction cache maintenance instructions.   |
| DVM Complete         | DVM Sync snoops received from the interconnect.   |
| WriteNoSnoop         | Non-cacheable store instructions.<br>Evictions of non-shareable cache lines from L1, L2, and L3 caches.   |
| WriteBack            | Evictions of dirty lines from the L1, L2, or L3 cache.<br>Streaming writes that are not allocating into the cache.  |
| WriteClean           | Evictions of dirty lines from the L3 cache, when the line is still present in an L1 or L2 cache.<br>Some cache maintenance instructions.  |
| WriteEvict           | Evictions of unique clean lines, when configured in the CLUSTERECTLR.   |
| Evict                | Evictions of clean lines, when configured in the CLUSTERECTLR.  |
| Barriers             | Not used.   |
| MakeInvalid          | Not used.   |
| ReadShared           | Not used.   |
| ReadOnceCleanInvalid | Not used.   |
| ReadOnceMakeInvalid  | Not used.   |
| StashOnceShared      | Not used.   |

| Transaction                  | Operation |
|------------------------------|-----------|
| StashOnceUnique              | Not used. |
| StashTranslation             | Not used. |
| WriteUniquePtl               | Not used. |
| WriteUniqueFull <sup>7</sup> | Not used. |
| WriteUniquePtlStash          | Not used. |
| WriteUniqueFullStash         | Not used. |

## 7.7 Support for memory types

The cores in the DynamIQ™ cluster simplify the coherency logic by downgrading some memory types.

Normal memory that is marked as both Inner Write-Back Cacheable and Outer Write-Back Cacheable is cached in the core data caches and the L3 cache.

All other Normal memory types are treated as Non-cacheable and are sent on the master interface as Normal Non-cacheable.

## 7.8 Read response

The ACE master can delay accepting a read data channel transfer by holding RREADY LOW for an indeterminate number of cycles.

RREADY can be deasserted LOW between read data channel transfers that form part of the same transaction.

The ACE master asserts the read acknowledge signal RACK HIGH in the ACLK cycle following acceptance of the last read data channel transfer for a transaction. RACK is asserted in AXI compatibility mode in addition to ACE configurations.



Note

- For interoperability of system components, Arm recommends that components interfacing with the ACE master are fully ACE-compliant with no reliance on the subset of permitted RACK behavior that is described for the DSU.
- If the interconnect does not perform hazarding between coherent and non-coherent requests, then, after it has returned the first transfer of read data for a non-coherent read, it must return all the remaining read transfers in the transaction.

<sup>7</sup> The AMBA5 ACE transaction types WriteUniqueFull and WriteUniquePtl were known in AMBA4 ACE as WriteLineUnique and WriteUnique, respectively.

The completion of the read transfers must not depend on either of the following:

- Snoop requests being sent to the core.
- The core needing to respond to a snoop request that could be to the same address.

---

## 7.9 Write response

The ACE master requires that the slave does not return a write response until it has received the write address.

The ACE master always accepts write responses without delay by holding BREADY HIGH.

The ACE master asserts the write acknowledge signal WACK HIGH in the ACLK cycle following acceptance of a write response. WACK is asserted in AXI compatibility mode in addition to ACE configurations.



For interoperability reasons, Arm recommends that system components fully comply with the ACE specification and do not rely on the DSU behavior described here.

---

## 7.10 Barriers

The DSU does not support sending barrier transactions to the interconnect. Barriers are always terminated within the cluster.

You must ensure that your interconnect and any peripherals connected that are to it do not return a write response for a transaction until that transaction would be considered complete by a later barrier. This means that the write must be observable to all other masters in the system. Arm expects most peripherals to meet this requirement.

## 7.11 AXI compatibility mode

The DSU implements an AXI compatibility mode that enables you to use the DSU in a standalone environment where the AMBA ACE5 interface is not required.

To enable this mode, you must ensure that the BROADCASTOUTER, BROADCASTCACHEMAINT, and BROADCASTPERSIST input signals are set to LOW. You must also tie the ACVALIDMx, ACWAKEUPMx, CDREADYMx, and CRREADYMx input signals LOW.

The AXI3 protocol supports write interleaving which is not used by the DSU. To allow compatibility with AXI3 components, the DSU provides WIDMx output signals, which can be connected to an AXI3 device.

If using AXI4 and ACE components, you must leave the WIDMx output signals unconnected as these signals do not exist in these protocols.

For single master implementations, WIDM0 is provided. For dual master implementations, WIDM0 and WIDM1 are provided.

### 7.11.1 Additional logic to support AXI compatibility

To support AXI compatibility, the DSU requires additional logic.

The DSU implements a handshake for system coherency using the SYSCOREQ\* and SYSCOACK\* signals. In AXI compatibility mode, the DSU does not support coherency. However, because this interface uses a handshake protocol, transitions on SYSCOREQ\* request must be responded to by a corresponding acknowledge on SYSCOACK\*.

Arm recommends that you implement the following logic in your system:

- Add a single-bit register with the input connected to SYSCOREQ\*.
- Connect the output of the register to SYSCOACK\*.
- The register must be reset LOW, and clocked and reset using the clock and reset from your AXI system.

For single ACE master implementations, SYSCOREQM0 and SYSCOACKM0 are provided. For dual ACE master implementations, SYSCOREQM1 and SYSCOACKM1 are also provided.

## 7.12 ACE privilege information

ACE provides information about the privilege level of accesses on the ARPROTM[0] and AWPROM[0] signals. This information is not available from cores within the cluster. Therefore these signals are always driven to HIGH indicating that the access could be a privileged access.

## 8. CHI master interface

This chapter describes the AMBA 5 CHI master memory interface.

### 8.1 About the CHI master interface

You can configure the DSU to use the AMBA 5 CHI protocol for the master memory interface.

#### 8.1.1 Dual CHI interfaces

The DSU can be implemented with one or two CHI interfaces.

Transactions generated by the cluster are routed to either interface based on the transaction type, memory type, and transaction address:

- All DVM transactions are routed to interface 0.
- All Device transactions are routed to interface 0.

For Cacheable transactions, the configuration parameter `INTERLEAVE_ADDR_BIT` controls which transaction address bit is used to select the routing between interface 0 and interface 1. The default value is to select bit 6 of the transaction address, which interleaves on cache line boundaries. If the selected transaction address bit has a value 0, then interface 0 is used, otherwise interface 1 is used. For more information on `INTERLEAVE_ADDR_BIT` and the impact on performance, see [2.3 Implementation options](#) on page 18.

For Normal Non-cacheable transactions, routing is dependent on the value of `CLUSTERECTLR.Non-cacheable behavior control, bit[0]`:

- If `CLUSTERECTLR[0]` is set to 1, Normal Non-cacheable transactions are routed to both interface 0 and 1 in the same way as Cacheable transactions, using the same bit of the transaction address.
- If `CLUSTERECTLR[0]` is set to 0, Normal Non-cacheable transactions are routed to interface 0.

When the external memory system sends snoops, it must either:

- Send the snoop to both interfaces.
- Send the snoop only to the interface that is relevant for the address of that snoop. This behavior is normal operation for an external memory system that contains a snoop filter. The snoop filter indicates that the line is present in one of the two masters.

The second method is more efficient, and if two masters are implemented, Arm recommends that the external memory system includes a snoop filter.

#### DVM messages

DVM messages can be received on both interfaces, however they are only required on interface 0, and any DVM message sent to interface 1 is treated as a no-op. Therefore for best performance,

Arm recommends that, when possible, your interconnect is configured to avoid sending DVM messages to interface 1.

### System address map

If the DSU is configured with the integrated CMN-600 SAM and two CHI interfaces, the DSU will only contain one SAM, and the node\_id field in the CMN-600 por\_rnsam\_node\_info register will contain the node ID of CHI interface 0.

## 8.2 CHI version

The DSU supports CHI Issue B, Issue C, or Issue D, depending on configuration.

Direct Connect configurations support CHI Issue C or Issue D, depending on core type, and all other configurations only support CHI Issue B.

The widths of the TXDAT and RXDAT channels are increased by one bit when CHI Issue C is enabled. CHI Issue D affects the widths of all channels.

## 8.3 CHI features

AMBA defines a set of interface properties for the CHI interconnect. The following table shows which of these properties the DSU supports, or requires the interconnect and system to support.

**Table 8-1: CHI interconnect properties for the DSU**

| CHI property             | Supported by the DSU   | Interconnect support required   |
|--------------------------|--|---|
| Atomic_Transactions      | Yes if BROADCASTATOMIC is HIGH.  | Yes if BROADCASTATOMIC is HIGH.   |
| Cache_Stash_Transactions | Yes  | Yes   |
| Direct_Memory_Transfer   | Yes  | Optional. The DSU supports this feature if it is implemented by the interconnect. |
| Direct_Cache_Transfer    | Yes  | Optional. The DSU supports this feature if it is implemented by the interconnect. |
| Data_Poison              | Yes if cache protection (SCU_CACHE_PROTECTION) is enabled.                                       | Yes if cache protection (SCU_CACHE_PROTECTION) is enabled.                        |
| Data_Check               | No   | No  |
| CCF_Wrap_Order           | Yes. The DSU always sends Data packets in critical chunk first wrap order.                       | No  |
| Req_Addr_Width           | 44. If a core with a 48-bit physical address width is configured inside the cluster, this is 48. | Not applicable  |
| NodeID_Width             | 11   | Not applicable  |
| Data_Width               | User configurable: 128 bits or 256 bits.   | Not applicable  |
| Barrier_Transactions     | No   | No. The DSU does not use these transaction types.                                 |

| CHI property                                    | Supported by the DSU             | Interconnect support required                     |
|---|----------------------------------|---|
| Data return from SC state.                      | Yes                              | Not applicable                                    |
| I/O de-allocation transactions (ROMI and ROCI). | No                               | No. The DSU does not use these transaction types. |
| ReadNotSharedDirty transactions                 | Yes                              | Yes   |
| CleanSharedPersist transactions                 | Yes if BROADCASTPERSIST is HIGH. | Yes if BROADCASTPERSIST is HIGH.                  |

For more information on these features, see [AMBA® 5 CHI Architecture Specification](#).

## 8.4 CHI configurations

You can change the coherency configurations to suit your system configuration using the BROADCASTCACHEMAINT and BROADCASTOUTER input signals.

The following table shows the permitted combinations of these signals and the supported configurations in the DSU, with a CHI bus.

**Table 8-2: Supported CHI configurations**

| Signal              | Feature                                 |                           |                             |                           |
|---------------------|---|---------------------------|-----------------------------|---------------------------|
|                     | CHI non-coherent                        |                           | CHI coherent                |                           |
|                     | With no cache or invisible system cache | With visible system cache | With invisible system cache | With visible system cache |
| BROADCASTCACHEMAINT | 0                                       | 1                         | 0                           | 1                         |
| BROADCASTOUTER      | 0                                       | 0                         | 1                           | 1                         |



- A visible system cache requires cache maintenance transactions to ensure that a write is visible to all observers.
- An invisible system cache is one that does not require cache maintenance transactions to ensure that a write is visible to all observers. This is true even if those observers use different memory attributes.

The following table shows the key features in each of the supported CHI configurations.

**Table 8-3: Supported features in the CHI configurations**

| Features                                    | Configuration                           |                           |              |
|---|---|---------------------------|--------------|
|   | CHI non-coherent                        |                           | CHI coherent |
|   | With no cache or invisible system cache | With visible system cache |              |
| Cache maintenance requests on TXREQ channel | No                                      | Yes                       | Yes          |
| Snoops on RXSNP channel                     | No                                      | No                        | Yes          |
| Coherent requests on TXREQ channel          | No                                      | No                        | Yes          |
| DVM requests on TXREQ channel               | No                                      | No                        | Yes          |



## 8.5 Attributes of the CHI master interface

The following table lists the possible values for the read and write issuing capabilities.

**Table 8-4: Attributes of the CHI master memory interface**

| Attribute                                   | Value                       | Comments   |
|---|-----------------------------|--|
| Write issuing capability                    | Configuration dependent     | <p>The maximum number of writes is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul>   |
| Read issuing capability                     | Configuration dependent     | <p>The maximum number of reads is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul> <p><b>Note:</b><br/>For Device and Normal Non-cacheable reads, the read issuing capability is limited by the combined issuing capability. Because the combined issuing capability is always lower than the read issuing capability.</p>  |
| Combined issuing capability                 | Configuration dependent     | <p>The combined issuing capability is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul> <p>The Device combined issuing capability is limited to:</p> <ul style="list-style-type: none"> <li>39, if two slices are present.</li> <li>23, if one slice is present.</li> </ul> <p>The Device and Normal Non-cacheable combined issuing capability is limited to:</p> <ul style="list-style-type: none"> <li>78, if two slices are present with one main master port.</li> <li>78, if two slices are present with two main master ports and the CLUSTERECTLR.Noncacheable behavior control, bit[0] is set to 1. See <a href="#">11.8 CLUSTERECTLR, Cluster Extended Control Register</a> on page 105.</li> <li>39, if two slices are present with two main master ports and the CLUSTERECTLR.Noncacheable behavior control, bit[0] is set to 0.</li> <li>23, if one slice is present.</li> </ul> |
| Exclusive hardware access thread capability | Number of hardware threads. | Each hardware thread can have one exclusive access sequence in progress.   |
| Transaction ID width                        | 8 bits                      | <p>There is no fixed mapping between CHI transaction IDs and cores. Transaction IDs can be used for either reads or writes.</p> <p><b>Note:</b><br/>The source of the transaction is encoded in the LPID field, see <a href="#">Table 8-6: CHI LPID assignment</a> on page 84.</p>   |

| Attribute                 | Value                   | Comments   |
|---------------------------|-------------------------|--|
| Transaction ID capability | Configuration dependent | <p>The ID capability is:</p> <ul style="list-style-type: none"> <li>96, if two slices are present.</li> <li>32, if one slice is present.</li> </ul> <p><b>Note:</b><br/>Unlike in an AMBA ACE5 configuration, there is never any ID reuse in CHI implementations, regardless of the memory type.</p> |
| NodeID widths             | 11 bits                 | -  |
| TXREQFLIT.RSVDC           | 0 bits                  | -  |
| TXDATFLIT.RSVDC           | 0 bits                  | -  |
| TXDATFLIT.DataCheck       | 0 bits                  | -  |



Note

The issuing capability described here is the maximum possible for the whole cluster, and can be used to size interconnect capabilities if you want to achieve the maximum performance available. However this maximum may not be able to be reached by a single core on its own, it may need multiple cores generating heavy memory traffic simultaneously to reach it. The capabilities will vary by core type, for example big cores will typically generate more transactions than LITTLE cores. It can also vary by memory type, with typically a significantly lower limit for Device or Non-cacheable transactions than for Cacheable transactions.

## Related information

[Cache slices and portions](#) on page 62

## 8.6 CHI channel properties

The following table shows the snoop capabilities and other CHI channel properties for the DSU.

**Table 8-5: CHI channel properties**

| Property                    | Value                   | Comment   |
|-----------------------------|-------------------------|---|
| Snoop acceptance capability | Configuration dependent | <p>For dual 256-bit CHI masters, the SCU can accept and process a maximum of 11 snoop requests from each master port.</p> <p>For a single 256-bit CHI master, the SCU can accept and process a maximum of 14 snoop requests from the system. For a 128-bit CHI master, the SCU can accept and process a maximum of 11 snoop requests.</p> |
| DVM acceptance capability   | 4                       | <p>The SCU can accept and process a maximum of four DVM transactions from the system. Each of these four transactions can be a two part DVM message.</p> <p>The interconnect must be configured to never send more than four DVM messages to the cluster, otherwise the system might deadlock.</p>  |

| Property               | Value     | Comment  |
|------------------------|-----------|--|
| Snoop latency          | Hit       | When there is a hit in L3 cache, the best case for response and data is 10 SCLK cycles. When there is a miss in the L3 cache but a hit in an L1 or L2 cache in a core, then the latency varies. This latency variation depends on the type and configuration of that core.<br><br>Latencies can be higher if hazards occur or if there are not enough buffers to absorb requests.  |
|                        | Miss      | Best case for latency is six SCLK cycles when the snoop filter and L3 cache tags indicate the miss.  |
|                        | DVM       | The cluster takes a minimum of six SCLK cycles to provide a response to DVM packets.   |
| Snoop filter           | Supported | The cluster provides support for an external snoop filter in an interconnect. It indicates when clean lines are evicted from the cluster by sending Evict transactions on the CHI write channel.<br><br>However there are some cases that can prevent an Evict transaction from being sent. Therefore you must ensure that you build any external snoop filter to handle a capacity overflow. When exceeding capacity, the snoop filter should send a back-invalidation to the cluster.<br><br>Examples of case where evicts are not produced include: <ul style="list-style-type: none"> <li>• Linefills that take external aborts.</li> <li>• Store exclusives that fail.</li> <li>• Mis-matched aliases.</li> </ul> |
| Supported transactions | -         | All transactions that are described by the CHI protocol: <ul style="list-style-type: none"> <li>• Are accepted on the CHI master interface from the system.</li> <li>• Can be produced on the CHI master interface except: <ul style="list-style-type: none"> <li>◦ ReadShared.</li> <li>◦ MakeInvalid.</li> <li>◦ EOBarrier.</li> <li>◦ ECBBarrier.</li> <li>◦ WriteCleanPtl.</li> <li>◦ WriteUniquePtl.</li> <li>◦ WriteBackPtl.</li> <li>◦ WriteUniqueFullStash.</li> <li>◦ WriteUniquePtlStash.</li> <li>◦ ReadOnceCleanInvalid.</li> <li>◦ ReadOnceMakeInvalid.</li> </ul> </li> </ul>  |

## 8.7 CHI transactions

CHI transactions are sent to a specific node in the interconnect that is based on the following criteria:

- Type of access.
- Address of the access.
- Settings of the System Address Map.

Addresses that map to an HN-F node can be marked as cacheable memory in the translation tables, and can take part in the cache coherency protocol. Addresses that map to an HN-I or MN must be marked as device or Non-cacheable memory.

CHI TXREQ transactions include the *Logical Processor ID* (LPID) field. This field uniquely identifies the logical core that generated the request transaction. The following table shows CHI LPID assignment.

**Table 8-6: CHI LPID assignment**

| LPID      | Description  |
|-----------|--|
| 0x0-0xF   | Bits [3:0] are encoded:<br><div style="display: flex; justify-content: space-between; padding: 0 10px;"> <div style="border-left: 1px solid black; padding-left: 5px;">[0]</div> <div>Thread number.</div> </div> <div style="display: flex; justify-content: space-between; padding: 0 10px;"> <div style="border-left: 1px solid black; padding-left: 5px;">[3:1]</div> <div>CPUID.</div> </div> |
| 0x10-0x1D | Reserved.  |
| 0x1E      | ACP request.   |
| 0x1F      | Cache copy back.   |

The following table shows the CHI transactions that can be generated and some typical operations that might cause these transactions to be generated.



This table does not provide an exhaustive list of operations that generate each type of transaction, because there are many possibilities.

**Table 8-7: CHI transaction types**

| Transaction        | Operation   |
|--------------------|---|
| ReadNoSnp          | Non-cacheable loads or instruction fetches.<br><br>Linefills of Non-shareable cache lines into L1, L2, or L3 caches.  |
| ReadOnce           | Cacheable loads that are not allocating into the cache.   |
| ReadClean          | Cache data linefills started by a load instruction.<br><br>Cache linefills started by an instruction fetch.   |
| ReadShared         | Not used.   |
| ReadNotSharedDirty | Cache data linefills started by a load instruction.<br><br>Cache linefills started by an instruction fetch.   |
| ReadUnique         | Cache data linefills started by a store instruction.  |
| CleanUnique        | Store instructions that hit in the cache but the line is not in a unique coherence state.   |
| MakeUnique         | Store instructions of a full cache line of data, that miss in the caches.   |
| CleanShared        | Cache maintenance instructions.   |
| CleanSharedPersist | Cache maintenance instructions. The <i>Data Cache Clean to the Point of Persistence</i> (DC_CVAP) cache maintenance instruction only generates this transaction when the BROADCASTPERSIST input signal is HIGH. |
| CleanInvalid       | Cache maintenance instructions.   |

| Transaction          | Operation  |
|----------------------|--|
| MakeInvalid          | Not used.  |
| DVMOp                | TLB and instruction cache maintenance instructions.  |
| EOBarrier            | Not used.  |
| ECBarrier            | Not used.  |
| PrefetchTgt          | Hardware prefetch hint to the memory controller.   |
| StashOnceShared      | Cache prefetch when there is no L3 cache present.  |
| StashOnceUnique      | Cache prefetch when there is no L3 cache present.  |
| WriteNoSnpPtl        | Non-cacheable store instructions.  |
| WriteNoSnpFull       | Non-cacheable store instructions.<br><br>Evictions of Non-shareable cache lines.   |
| WriteUniqueFull      | Cacheable writes of a full cache line, that are not allocating into L1, L2, or L3 caches, for example streaming writes.                      |
| WriteUniquePtl       | Not used.  |
| WriteBackFull        | Evictions of dirty lines from the L1, L2, or L3 caches.  |
| WriteBackPtl         | Not used.  |
| WriteCleanFull       | Evictions of dirty lines from the L3 cache, when the line is still present in an L1 or L2 cache.<br><br>Some cache maintenance instructions. |
| WriteCleanPtl        | Not used.  |
| WriteEvictFull       | Evictions of unique clean lines, when configured in the CLUSTERECTLR.  |
| Evict                | Evictions of clean lines, when configured in the CLUSTERECTLR.   |
| AtomicStore          | Atomic instruction.  |
| AtomicLoad           | Atomic instruction.  |
| AtomicSwap           | Atomic instruction.  |
| AtomicCompare        | Atomic instruction.  |
| WriteUniqueFullStash | Not used.  |
| WriteUniquePtlStash  | Not used.  |
| ReadOnceCleanInvalid | Not used.  |
| ReadOnceMakeInvalid  | Not used.  |

External memory accesses generate the following transactions in an implementation configured with a CHI master interface.

**Table 8-8: CHI transaction usage**

| Attributes  |                 | CHI transaction |           |            |                                 |                                  |
|-------------|-----------------|-----------------|-----------|------------|---------------------------------|----------------------------------|
| Memory type | Shareability    | SnpAttr         | Load      | Store      | Load exclusive                  | Store exclusive                  |
| Device      | Outer Shareable | Non-snooperable | ReadNoSnp | WriteNoSnp | ReadNoSnp and Excl set to HIGH. | WriteNoSnp and Excl set to HIGH. |

| Attributes  |                 | CHI transaction |                                 |   |   |   |
|---|-----------------|-----------------|---------------------------------|---|---|---|
| Memory type   | Shareability    | SnpAttr         | Load                            | Store   | Load exclusive  | Store exclusive   |
| Normal, Inner Non-cacheable, Outer Non-cacheable  | Non-shareable   | Non-snoopable   | ReadNoSnp                       | WriteNoSnp  | ReadNoSnp and Excl set to HIGH.                         | WriteNoSnp and Excl set to HIGH.  |
|   | Inner Shareable |                 |                                 |   |   |   |
|   | Outer Shareable |                 |                                 |   |   |   |
| Normal, Inner Non-cacheable, Outer Write-Back or Write-Through, or Normal, Inner Write-Through, Outer Write-Back, Write-Through or Non-cacheable, or Normal Inner Write-Back Outer Non-cacheable or Write-Through | Non-shareable   | Non-snoopable   | ReadNoSnp                       | WriteNoSnp  | ReadNoSnp and Excl set to HIGH.                         | WriteNoSnp and Excl set to HIGH.  |
|   | Inner Shareable |                 |                                 |   |   |   |
|   | Outer Shareable |                 |                                 |   |   |   |
| Normal, Inner Write-Back, Outer Write-Back  | Non-shareable   | Non-snoopable   | ReadNoSnp                       | WriteNoSnp when the line is evicted or if not allocating into the cache.  | ReadNoSnp   | WriteNoSnp when the line is evicted.  |
|   | Inner Shareable | Snoopable       | ReadNotSharedDirty or ReadClean | ReadUnique, CleanUnique, or MakeUnique if allocating into the cache, then a WriteBackFull when the line is evicted. | ReadNotSharedDirty or ReadClean, with Excl set to HIGH. | CleanUnique with Excl set to HIGH if required, then a WriteBackFull when the line is evicted. |
|   | Outer Shareable | Snoopable       |                                 |   |   |   |
|   |                 |                 |                                 | WriteUniqueFull if not allocating into the cache.   |   |   |

## 8.8 Use of DataSource

Some CHI responses from the interconnect include a DataSource field indicating where the data was supplied from. When making use of the DataSource field, Arm recommends providing this information as accurately as possible using the encodings recommended in the *Suggested DataSource value encodings* provided in the [AMBA® 5 CHI Architecture Specification](#).

The value of this field is used to calculate some PMU events, and can also be used by some cores to tune the performance of their data prefetchers.

## 8.9 Support for memory types

The cores in the DynamIQ™ cluster simplify the coherency logic by downgrading some memory types.

Normal memory that is marked as both Inner Write-Back Cacheable and Outer Write-Back Cacheable is cached in the core data caches and the L3 cache.

All other Normal memory types are treated as Non-cacheable and are sent on the master interface as Normal Non-cacheable.

## 9. ACP slave interface

This chapter describes the ACP slave interface.

### 9.1 About the ACP

The *Accelerator Coherency Port* (ACP) is an optional slave interface, conforming to a subset of the ACE5-Lite specification.

The ACP slave interface allows an external master to access memory through the main memory interface of the DSU. Only access to Cacheable memory is permitted.

The read and write data buses of the ACP are 128 bits. Accesses are optimized for cache line length.

To maintain cache coherency, accesses are checked in all cached locations in the cluster. That is, the L3 cache, and the data caches in each core.

By default, ACP write accesses are implicit stash requests to the L3 cache. Alternatively, implicit stash requests can target the L2 cache of a selected core.

### 9.2 ACP features

The ACP interface for the DSU supports the following properties.

**Table 9-1: ACP interface properties for the DSU**

| ACP property                    | Supported by the DSU |
|---------------------------------|----------------------|
| Port_Type                       | Accelerator          |
| Continuous_Cache_Line_Read_Data | Yes                  |
| Multi_Copy_Atomicity            | Yes                  |
| Ordered_Write_Observation       | No                   |
| WriteEvict_Transaction          | No                   |
| DVM_v8                          | No                   |
| Atomic_Transactions             | No                   |
| DVM_v8.1                        | No                   |
| Cache_Stash_Transactions        | Yes                  |
| DeAllocation_Transactions       | No                   |
| Persistent_CMO                  | No                   |
| Poison                          | No                   |
| Data_Check                      | No                   |
| QoS_Accept                      | No                   |



| ACP property              | Supported by the DSU |
|---------------------------|----------------------|
| Trace_Signals             | No                   |
| Loopback_Signals          | No                   |
| Low_Power_Signals         | Yes                  |
| Untranslated_Transactions | No                   |
| NSAccess_Identifiers      | No                   |

## 9.3 ACP ACE5-Lite subset

The ACP conforms to a subset of the ACE5-Lite specification.

The ACP ACE5-Lite subset is described in [AMBA® AXI and ACE Protocol Specification](#). The DSU has the following additional restrictions:

- The values of ARCACHES and AWCACHES are restricted to:
  - 0b0111.
  - 0b1011.
  - 0b1111.
- All transactions can be Secure or Non-secure.
- Exclusive accesses are not supported. ARLOCK and AWLOCK signals are not present.
- All requests can specify Outer Shareable and Non-shareable using the AWDOMAINS and ARDOMAINS signals.
- Barriers are not supported. The BRESP response for any write transaction indicates global observability for the transaction.
- ARSIZE and AWSIZE signals are not present. A value of 4 (16 bytes) is assumed.
- The values of ARLENS and AWLENS are restricted to:
 

|   |             |
|---|-------------|
| 0 | One beat.   |
| 3 | Four beats. |
- ARBURST and AWBURST signals are not present. A value of 0b01 (INCR) is assumed.
- ARSNOOP signals are not present. A value of 0b0000 is assumed.
- ARQOS and AWQOS signals are not present.

## 9.4 ACP transaction types

The ACP supports transaction types having the following transfer size and length combinations:

- 16-byte INCR read transaction:
  - ARLENS is 0 (one beat).

- Address aligned to 16-byte boundary (ARADDRS[3:0] is 0b0000).
- 64-byte INCR read transaction:
  - ARLENS is 3 (four beats).
  - Address aligned to 64-byte boundary (ARADDRS[5:0] is 0b000000).

The ACP supports the following write transaction transfer size and length combinations:

- 16-byte INCR write transaction:
  - WSTRBS, any combination of bytes, including no bytes, are valid.
  - AWLENS is 0 (one beat).
  - AWSNOOPS is WriteUniquePtl or WriteUniquePtlStash.
  - Address aligned to 16-byte boundary (AWADDRS[3:0] is 0b0000).
- 64-byte INCR write transaction:
  - WSTRBS, any combination of bytes, including no bytes, are valid.



When AWSNOOPS is WriteUniqueFull, all bytes must be valid.

---

- AWLENS is 3 (four beats).
  - AWSNOOPS is WriteUniquePtl, WriteUniquePtlStash, or WriteUniqueFull.
  - Address aligned to 64-byte boundary (AWADDRS[5:0] is 0b000000).
- 



- The AMBA 5 ACE-Lite transaction types WriteUniqueFull and WriteUniquePtl were known in AMBA 4 ACE-Lite as WriteLineUnique and WriteUnique, respectively.
  - The DSU treats WriteUniquePtlStash as a WriteUniquePtl and does not perform a stash operation for this transaction type.
- 

The ACP supports the following Cache Stash Transaction transfer size and length:

- 64-byte INCR write stash transaction:
  - WSTRBS, all bytes are valid.
  - AWLENS is 3 (four beats).
  - AWSNOOPS is WriteUniqueFullStash.
  - Address aligned to 64-byte boundary (AWADDRS[5:0] is 0b000000).
- Dataless 64-byte INCR write stash transaction:
  - No W-Channel transfers.
  - AWLENS is 3.
  - AWSNOOPS is StashOnceShared or StashOnceUnique.

- Address aligned to 64-byte boundary (AWADDRS[5:0] is 0b000000).

Stash requests can target the L2 cache of a selected core by asserting signal AWSTASHLPIDENS and indicating the selected core number on AWSTASHLPIDS[3:1]. The signal AWSTASHLPIDS[0] is reserved for the thread number, but this does not affect the stash request.



Requests not meeting these restrictions cause a SLVERR response on RRESPS or BRESPS.

The following table lists the ACP supported transactions:

**Table 9-2: ACP supported transactions**

| Transaction          | Notes  |
|----------------------|--|
| ReadOnce             | -  |
| WriteUniqueFull      | -  |
| WriteUniquePtl       | -  |
| WriteUniquePtlStash  | Treated as a WriteUniquePtl. The DSU does not perform a stash. |
| WriteUniqueFullStash | -  |
| StashOnceUnique      | -  |
| StashOnceShared      | -  |

## 9.5 ACP performance

For optimum performance, use the following guidelines for ACP transactions.

The master must avoid sending more than one outstanding transaction on the same AXI ID, to prevent the second transaction stalling the interface until the first has completed. If the master requires explicit ordering between two transactions, Arm recommends that it waits for the response to the first transaction before sending the second transaction.

Writes are higher performance when they use WriteUniqueFull or WriteUniqueFullStash transactions.

WriteUniquePtl or WriteUniquePtlStash transactions always incur a read-modify write sequence.

Some L3 resources are shared between the ACP interface and the cores. Therefore, heavy traffic on the ACP interface might, in some cases, reduce the performance of the cores.

Write transactions use the Write-Allocate bit of the memory type (AWCACHES[3]) to decide whether to allocate to the L3 cache, as follows:

- If the stash request does not target a core (AWSTASHLPIDENS is LOW) and AWCACHES[3] is HIGH, then the cache line is allocated to the L3 cache.

- If the stash request does not target a core (AWSTASHLPIDENS is LOW) and AWCACHES[3] is LOW, then the cache line is not allocated to the L3 cache and it will be written out on the master port instead.
- When the stash request does not target a core (AWSTASHLPIDENS is LOW), then the WriteUniqueFullStash transaction performs the same operation as WriteUniqueFull.
- Stash requests that target a core (AWSTASHLPIDENS is HIGH) always attempt to allocate to the core L2 cache. In this case, Arm recommends that AWCACHES[3] is HIGH. Since, if AWCACHES[3] is LOW, then the line will not initially be allocated to the cache. Instead the line will be written out on the master port before being fetched back into the core, which is inefficient.

The following table describes the ACP acceptance capabilities.

**Table 9-3: ACP acceptance capabilities**

| Attribute                      | Value | Description  |
|--------------------------------|-------|--|
| Write acceptance capability    | 33    | The ACP can accept up to 33 write transactions.  |
| Read acceptance capability     | 33    | The ACP can accept up to 33 read transactions.   |
| Combined acceptance capability | 34    | The ACP can accept up to 34 transactions. There is no performance benefit above 32 outstanding transactions. |

## 10. AXI master peripheral port

This chapter describes the AXI master peripheral port.

### 10.1 About the peripheral port

The peripheral port supports Device accesses to tightly coupled accelerators.

The peripheral port can be used for low-latency access to peripherals local to the cluster. It has the same latency as the main master port. However, the overall system latency to devices that are connected to the main master port is greater because of the higher latency of the system interconnect.

The peripheral port is optionally implemented. It is a 64-bit AXI4 master interface.

The peripheral port supports access to only Device-nGRE, nGnRE, and nGnRnE memory types:

- All accesses must be aligned load or store instructions of 64 bits or less. Unaligned or larger accesses are not supported and generate an external abort.
- Atomic instructions are not supported and generate an external abort.
- Load and store exclusive instructions are not supported. Store exclusive instructions will fail and the result register will reflect this, but the memory location might be updated.

Accesses to the port using other memory types are unpredictable.

The peripheral port address range is defined by two configuration input buses.

ASTARTMP[PA-1:20] for the start of the address range, which is inclusive, and AENDMP[PA-1:20] for the end of the address range, which is exclusive. PA is the largest physical address width of any connected core. These signals are only captured at reset.

### 10.2 Transaction ID encoding

The AXI interface provides several read and write issuing capabilities and requires a specific format for AXI transaction IDs.

The following table describes the read and write issuing capabilities of the AXI interface.

**Table 10-1: AXI issuing capabilities**

| Attribute                   | Value | Comments   |
|-----------------------------|-------|--|
| Write issuing capability    | 5     | There can be up to 5 outstanding write transactions.     |
| Read issuing capability     |       | There can be up to 5 outstanding read transactions.      |
| Combined issuing capability | 10    | There can be up to 10 outstanding transactions.          |
| Write ID capability         | 5     | Each ID can have up to 5 outstanding write transactions. |
| Read ID capability          |       | Each ID can have up to 5 outstanding read transactions.  |

The following table lists the encoding for AXI transaction IDs.

**Table 10-2: AXI transaction ID encoding**

| Attribute | Value | Comments                                      |
|-----------|-------|---|
| All IDs   | tnnn  | Thread t <sup>8</sup> , core nnn <sup>9</sup> |



These ID and transaction details are provided for information only. Arm strongly recommends that all interconnects and peripherals are designed to support any type and number of transactions on any ID, to ensure compatibility with future products.

See the [AMBA® AXI and ACE Protocol Specification](#) for more information about the ACE and AXI signals described in this manual.

<sup>8</sup> t is the hardware thread number and is 0 when the core does not support multiple hardware threads.

<sup>9</sup> nnn is the core number 0b000-0b111 in binary.

# 11. Control registers

This chapter describes the control registers for the DSU.

## 11.1 About the control registers

The DSU contains system control registers in the SCU and L3 logic to control the functionality of the cluster. Most of these registers are shared between all the cores in the cluster, but a few are private to each core.

The chapter is presented as follows:

### AArch32 control register summary

This section lists the AArch32 control registers by access encoding.

### AArch64 control register summary

This section lists the AArch64 control registers by access encoding.

### Register descriptions

The remainder of the chapter provides generic register descriptions, that apply to both AArch32 and AArch64 registers. They are listed in alphabetical order.

## 11.2 AArch32 control register summary

This section lists the AArch32 control registers implemented in the DSU, sorted by access encoding.

**Table 11-1: DynamIQ™ Shared Unit AArch32 control registers**

| Register mnemonic | Copro | CRn | Opc1 | CRm | Opc2 | Width | Register name and description  |
|-------------------|-------|-----|------|-----|------|-------|--|
| CLUSTERCFR        | cp15  | c15 | 0    | c3  | 0    | 32    | <a href="#">11.7 CLUSTERCFR, Cluster Configuration Register</a> on page 101      |
| CLUSTERIDR        | cp15  | c15 | 0    | c3  | 1    | 32    | <a href="#">11.9 CLUSTERIDR, Cluster Main Revision ID Register</a> on page 109   |
| CLUSTERREVIDR     | cp15  | c15 | 0    | c3  | 2    | 32    | <a href="#">11.16 CLUSTERREVIDR, Cluster Revision ID Register</a> on page 125    |
| CLUSTERACTLR      | cp15  | c15 | 0    | c3  | 3    | 32    | <a href="#">11.5 CLUSTERACTLR, Cluster Auxiliary Control Register</a> on page 99 |
| CLUSTERECTLR      | cp15  | c15 | 0    | c3  | 4    | 32    | <a href="#">11.8 CLUSTERECTLR, Cluster Extended Control Register</a> on page 105 |
| CLUSTERPWRCTLR    | cp15  | c15 | 0    | c3  | 5    | 32    | <a href="#">11.13 CLUSTERPWRCTLR, Cluster Power Control Register</a> on page 119 |
| CLUSTERPWRDN      | cp15  | c15 | 0    | c3  | 6    | 32    | <a href="#">11.14 CLUSTERPWRDN, Cluster Powerdown Register</a> on page 121       |
| CLUSTERPWRSTAT    | cp15  | c15 | 0    | c3  | 7    | 32    | <a href="#">11.15 CLUSTERPWRSTAT, Cluster Power Status Register</a> on page 123  |

| Register mnemonic   | Copro | CRn | Opc1 | CRm | Opc2 | Width | Register name and description  |
|---------------------|-------|-----|------|-----|------|-------|--|
| CLUSTERTHREADSID    | cp15  | c15 | 0    | c4  | 0    | 32    | <a href="#">11.18 CLUSTERTHREADSID</a> , Cluster Thread Scheme ID Register on page 129             |
| CLUSTERACPSID       | cp15  | c15 | 0    | c4  | 1    | 32    | <a href="#">11.4 CLUSTERACPSID</a> , Cluster ACP Scheme ID Register on page 97                     |
| CLUSTERSTASHSID     | cp15  | c15 | 0    | c4  | 2    | 32    | <a href="#">11.17 CLUSTERSTASHSID</a> , Cluster Stash Scheme ID Register on page 127               |
| CLUSTERPARTCR       | cp15  | c15 | 0    | c4  | 3    | 32    | <a href="#">11.12 CLUSTERPARTCR</a> , Cluster Partition Control Register on page 114               |
| CLUSTERBUSQOS       | cp15  | c15 | 0    | c4  | 4    | 32    | <a href="#">11.6 CLUSTERBUSQOS</a> , Cluster Bus QoS Control Register on page 99                   |
| CLUSTERL3HIT        | cp15  | c15 | 0    | c4  | 5    | 32    | <a href="#">11.10 CLUSTERL3HIT</a> , Cluster L3 Hit Counter Register on page 111                   |
| CLUSTERL3MISS       | cp15  | c15 | 0    | c4  | 6    | 32    | <a href="#">11.11 CLUSTERL3MISS</a> , Cluster L3 Miss Counter Register on page 113                 |
| CLUSTERTHREADSIDOVR | cp15  | c15 | 0    | c4  | 7    | 32    | <a href="#">11.19 CLUSTERTHREADSIDOVR</a> , Cluster Thread Scheme ID Override Register on page 130 |

## 11.3 AArch64 control register summary

This section lists the AArch64 control registers implemented in the DSU, sorted by access encoding.

**Table 11-2: DynamIQ™ Shared Unit AArch64 control registers**

| Register mnemonic    | Op0 | CRn | Op1 | CRm | Op2 | Width | Register name and description  |
|----------------------|-----|-----|-----|-----|-----|-------|--|
| CLUSTERCFR_EL1       | 3   | c15 | 0   | c3  | 0   | 32    | <a href="#">11.7 CLUSTERCFR</a> , Cluster Configuration Register on page 101           |
| CLUSTERIDR_EL1       | 3   | c15 | 0   | c3  | 1   | 32    | <a href="#">11.9 CLUSTERIDR</a> , Cluster Main Revision ID Register on page 109        |
| CLUSTERREVIDR_EL1    | 3   | c15 | 0   | c3  | 2   | 32    | <a href="#">11.16 CLUSTERREVIDR</a> , Cluster Revision ID Register on page 125         |
| CLUSTERACTLR_EL1     | 3   | c15 | 0   | c3  | 3   | 32    | <a href="#">11.5 CLUSTERACTLR</a> , Cluster Auxiliary Control Register on page 99      |
| CLUSTERECTLR_EL1     | 3   | c15 | 0   | c3  | 4   | 32    | <a href="#">11.8 CLUSTERECTLR</a> , Cluster Extended Control Register on page 105      |
| CLUSTERPWRCTLR_EL1   | 3   | c15 | 0   | c3  | 5   | 32    | <a href="#">11.13 CLUSTERPWRCTLR</a> , Cluster Power Control Register on page 119      |
| CLUSTERPWRDN_EL1     | 3   | c15 | 0   | c3  | 6   | 32    | <a href="#">11.14 CLUSTERPWRDN</a> , Cluster Powerdown Register on page 121            |
| CLUSTERPWRSTAT_EL1   | 3   | c15 | 0   | c3  | 7   | 32    | <a href="#">11.15 CLUSTERPWRSTAT</a> , Cluster Power Status Register on page 123       |
| CLUSTERTHREADSID_EL1 | 3   | c15 | 0   | c4  | 0   | 32    | <a href="#">11.18 CLUSTERTHREADSID</a> , Cluster Thread Scheme ID Register on page 129 |
| CLUSTERACPSID_EL1    | 3   | c15 | 0   | c4  | 1   | 32    | <a href="#">11.4 CLUSTERACPSID</a> , Cluster ACP Scheme ID Register on page 97         |
| CLUSTERSTASHSID_EL1  | 3   | c15 | 0   | c4  | 2   | 32    | <a href="#">11.17 CLUSTERSTASHSID</a> , Cluster Stash Scheme ID Register on page 127   |



| Register mnemonic       | Op0 | CRn | Op1 | CRm | Op2 | Width | Register name and description  |
|-------------------------|-----|-----|-----|-----|-----|-------|--|
| CLUSTERPARTCR_EL1       | 3   | c15 | 0   | c4  | 3   | 32    | <a href="#">11.12 CLUSTERPARTCR</a> , Cluster Partition Control Register on page 114               |
| CLUSTERBUSQOS_EL1       | 3   | c15 | 0   | c4  | 4   | 32    | <a href="#">11.6 CLUSTERBUSQOS</a> , Cluster Bus QoS Control Register on page 99                   |
| CLUSTERL3HIT_EL1        | 3   | c15 | 0   | c4  | 5   | 32    | <a href="#">11.10 CLUSTERL3HIT</a> , Cluster L3 Hit Counter Register on page 111                   |
| CLUSTERL3MISS_EL1       | 3   | c15 | 0   | c4  | 6   | 32    | <a href="#">11.11 CLUSTERL3MISS</a> , Cluster L3 Miss Counter Register on page 113                 |
| CLUSTERTHREADSIDOVR_EL1 | 3   | c15 | 0   | c4  | 7   | 32    | <a href="#">11.19 CLUSTERTHREADSIDOVR</a> , Cluster Thread Scheme ID Override Register on page 130 |

## 11.4 CLUSTERACPSID, Cluster ACP Scheme ID Register


The CLUSTERACPSID register provides the scheme ID for ACP transactions.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERACPSID) and AArch64 (CLUSTERACPSID\_EL1) registers.

CLUSTERACPSID is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-1: CLUSTERACPSID bit assignments**



### RAZ, [31:3]

Read-As-Zero.

### Scheme ID for ACP transactions, [2:0]

These bits reset to 0b000.

### Configurations

The AArch32 CLUSTERACPSID register is architecturally mapped to the AArch64 CLUSTERACPSID\_EL1 register.

## Usage constraints

### Accessing the CLUSTERACPSID

#### In AArch64 state (CLUSTERACPSID\_EL1):

To read this register in AArch64 state (CLUSTERACPSID\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_1; Read CLUSTERACPSID_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERACPSID\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_1, <Xt>; Write Xt into CLUSTERACPSID_EL1
```

#### In AArch32 state (CLUSTERACPSID):

To read this register in AArch32 state (CLUSTERACPSID) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 1; Read CLUSTERACPSID into Rt
```

To write this register in AArch32 state (CLUSTERACPSID) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 1; Write Rt into CLUSTERACPSID
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERACPSID | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERACPSID | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERACPSID | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.SMEN is 1 and ACTLR\_EL2.SMEN is 1, or ACTLR\_EL3.SMEN is 1 and SCR.NS is 0.

If write access is not permitted, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.5 CLUSTERACTLR, Cluster Auxiliary Control Register

The CLUSTERACTLR register is Reserved.

### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.ACTLREN is 1 and ACTLR\_EL2.ACTLREN is 1, or ACTLR\_EL3.ACTLREN is 1 and SCR.NS is 0.

If write access is not permitted, then trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.6 CLUSTERBUSQOS, Cluster Bus QoS Control Register

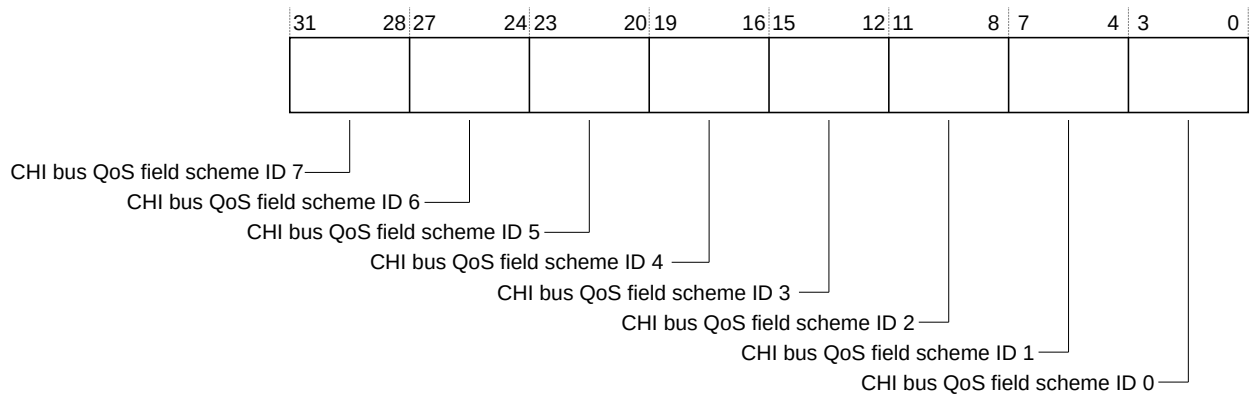
The CLUSTERBUSQOS provides control for the CHI *Quality of Service* (QoS) fields for different scheme IDs.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERBUSQOS) and AArch64 (CLUSTERBUSQOS\_EL1) registers.

CLUSTERBUSQOS is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-2: CLUSTERBUSQOS bit assignments**



### CHI bus QoS field scheme ID 7, [31:28]

Value driven on the CHI bus QoS field for scheme ID 7.

These bits reset to 0xE.

### CHI bus QoS field scheme ID 6, [27:24]

Value driven on the CHI bus QoS field for scheme ID 6.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 5, [23:20]**

Value driven on the CHI bus QoS field for scheme ID 5.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 4, [19:16]**

Value driven on the CHI bus QoS field for scheme ID 4.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 3, [15:12]**

Value driven on the CHI bus QoS field for scheme ID 3.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 2, [11:8]**

Value driven on the CHI bus QoS field for scheme ID 2.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 1, [7:4]**

Value driven on the CHI bus QoS field for scheme ID 1.

These bits reset to 0xE.

**CHI bus QoS field scheme ID 0, [3:0]**

Value driven on the CHI bus QoS field for scheme ID 0.

These bits reset to 0xE.

**Configurations**

The AArch32 CLUSTERBUSQOS register is architecturally mapped to the AArch64 CLUSTERBUSQOS\_EL1 register.

**Usage constraints****Accessing the CLUSTERBUSQOS****In AArch64 state (CLUSTERBUSQOS\_EL1):**

To read this register in AArch64 state (CLUSTERBUSQOS\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_4; Read CLUSTERBUSQOS_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERBUSQOS\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_4, <Xt>; Write Xt into CLUSTERBUSQOS_EL1
```

**In AArch32 state (CLUSTERBUSQOS):**

To read this register in AArch32 state (CLUSTERBUSQOS) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 4; Read CLUSTERBUSQOS into Rt
```

To write this register in AArch32 state (CLUSTERBUSQOS) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 4; Write Rt into CLUSTERBUSQOS
```

**Accessibility**

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERBUSQOS | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERBUSQOS | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERBUSQOS | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

**Traps and enables**

This register is write accessible in EL1 if ACTLR\_EL3.SMEN is 1 and ACTLR\_EL2.SMEN is 1, or ACTLR\_EL3.SMEN is 1 and SCR.NS is 0.

If write access is not permitted, then trap to the lowest Exception level that denied access (EL2 or EL3).

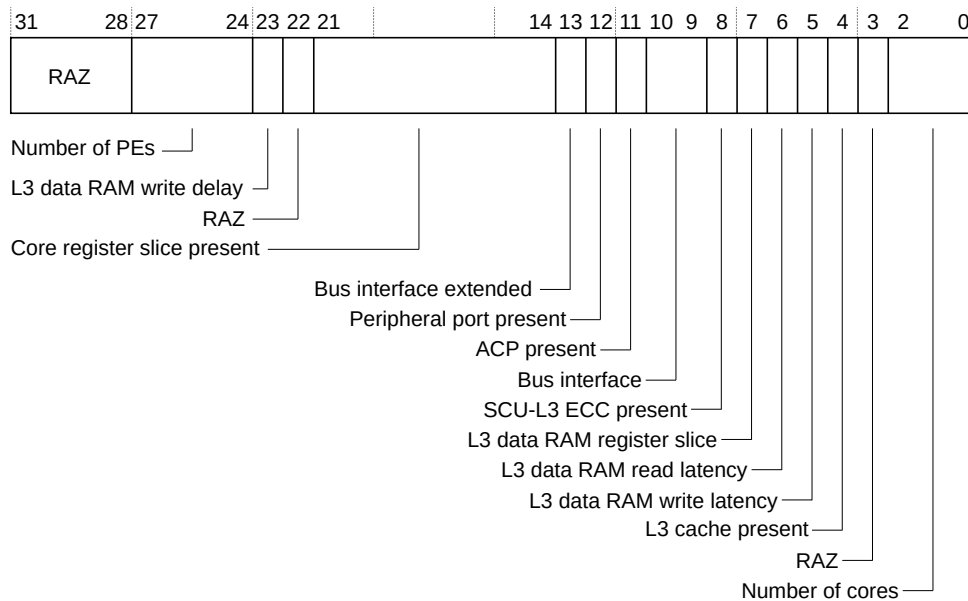
## 11.7 CLUSTERCFR, Cluster Configuration Register

The CLUSTERCFR register contains details of the hardware configuration of the cluster. This register is read-only write-ignores RO (WI), and is common to all execution threads.

**Bit field descriptions**

This description applies to both the AArch32 (CLUSTERCFR) and AArch64 (CLUSTERCFR\_EL1) registers.

CLUSTERCFR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-3: CLUSTERCFR bit assignments****RAZ, [31:28]**

Read-As-Zero

**Number of PEs, [27:24]**

PE - 1, where PE is the number of processing elements in the cluster. Each core contains either one or two PEs.

**L3 data RAM write delay, [23]**

The possible values are:

- 0 Writes are not limited.
- 1 Writes are limited to one write every three cycles.

**RAZ, [22]**

Read-As-Zero

**Core register slice present, [21:14]**

Core register slice is present. Each bit represents a core, with bit[14] for core 0 up to bit[21] for core 7:

- 0 No register slice is present.
- 1 Register slice is present.

**Bus interface extended, [13]**

See Bus interface (bits[10:9]).

**Peripheral port present, [12]**

Peripheral port is present:

|   |                                |
|---|--------------------------------|
| 0 | No peripheral port is present. |
| 1 | Peripheral port is present.    |

**ACP present, [11]**

ACP interface is present:

|   |                              |
|---|------------------------------|
| 0 | No ACP interface is present. |
| 1 | ACP interface is present.    |

**Bus interface, [10:9]**

Bus interface configuration:

|      |   |
|------|---|
| 0b00 | Single 128-bit ACE.   |
| 0b01 | Dual 128-bit ACE.   |
| 0b10 | Single 128-bit CHI.   |
| 0b11 | If Bus interface extended (bit [13]) is 1, then Dual 256-bit CHI.   |
|      | If Bus interface extended (bit [13]) is 0, then Single 256-bit CHI. |

**SCU-L3 ECC present, [8]**

SCU-L3 is configured with ECC:

|   |                                    |
|---|------------------------------------|
| 0 | SCU-L3 is not configured with ECC. |
| 1 | SCU-L3 is configured with ECC.     |

**L3 data RAM register slice, [7]**

L3 data RAM read register slice:

|   |                               |
|---|-------------------------------|
| 0 | No register slice is present. |
| 1 | Register slice is present.    |

**L3 data RAM read latency, [6]**

L3 data RAM read latency:

|   |   |
|---|---|
| 0 | Two-cycle output delay from the L3 data RAMs. |
| 1 | Three-cycle output delay from L3 data RAMs.   |

**L3 data RAM write latency, [5]**

L3 data RAM write latency:

|   |  |
|---|--|
| 0 | One cycle input delay from the L3 data RAMs. |
| 1 | Two cycle input delay from the L3 data RAMs. |

**L3 cache present, [4]**

L3 cache is present:

|   |                         |
|---|-------------------------|
| 0 | No L3 cache is present. |
| 1 | L3 cache is present.    |

**RAZ, [3]**

Read-As-Zero

**Number of cores, [2:0]**

Number of cores present in the cluster:

|       |                       |
|-------|-----------------------|
| 0b000 | One core, core0.      |
| 0b001 | Two cores, core0-1.   |
| 0b010 | Three cores, core0-2. |
| 0b011 | Four cores, core0-3.  |
| 0b100 | Five cores, core0-4.  |
| 0b101 | Six cores, core0-5.   |
| 0b110 | Seven cores, core0-6. |
| 0b111 | Eight cores, core0-7. |

**Configurations**

The AArch32 CLUSTERCFR register is architecturally mapped to the AArch64 CLUSTERCFR\_EL1 register.

**Usage constraints****Accessing the CLUSTERCFR****In AArch64 state (CLUSTERCFR\_EL1):**

To read this register in AArch64 state (CLUSTERCFR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_0; Read CLUSTERCFR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERCFR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_0, <Xt>; Write Xt into CLUSTERCFR_EL1
```

**In AArch32 state (CLUSTERCFR):**

To read this register in AArch32 state (CLUSTERCFR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 0; Read CLUSTERCFR into Rt
```



To write this register in AArch32 state (CLUSTERCFR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 0; Write Rt into CLUSTERCFR
```

### Accessibility

This register is accessible in software as follows:

| <systemreg> | Control |     |    | Accessibility |       |       |       |
|-------------|---------|-----|----|---------------|-------|-------|-------|
|             | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERCFR  | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERCFR  | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERCFR  | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

There are no special traps and enables.

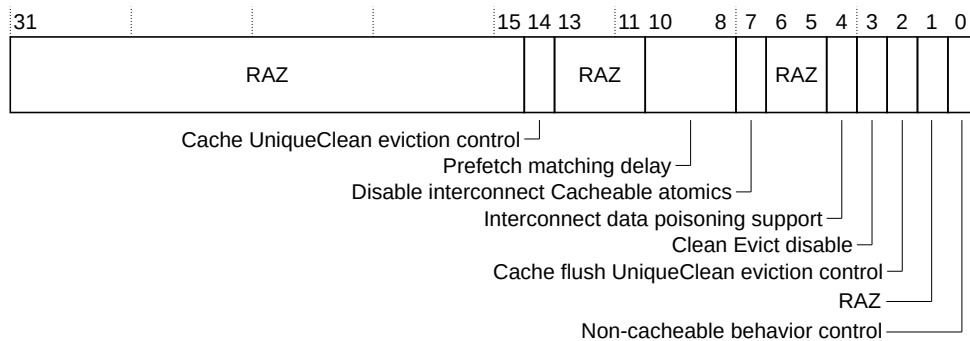
## 11.8 CLUSTERECTLR, Cluster Extended Control Register

The CLUSTERECTLR register provides implementation-specific control of the microarchitecture. It must only be written to as part of the initial system configuration setup following reset. This register is RW, and is common to all execution threads.

### Bit field descriptions

CLUSTERECTLR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

This description applies to both the AArch32 (CLUSTERECTLR) and AArch64 (CLUSTERECTLR\_EL1) registers.

**Figure 11-4: CLUSTERECTLR bit assignments****RAZ, [31:15]**

Read-As-Zero

**Cache UniqueClean eviction control, [14]**

Enables sending WriteEvict transactions on the ACE or CHI master for UniqueClean evictions. WriteEvict transactions update downstream caches that are outside the cluster. Enable WriteEvict transactions only if there is an L4 or system cache implemented in the system. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | Disables sending data with UniqueClean evictions. |
| <b>1</b> | Enables sending data with UniqueClean evictions.  |

For ACE, this bit resets to 0.

For CHI, this bit resets to 1.

**RAZ, [13:11]**

Read-As-Zero

**Prefetch matching delay, [10:8]**

Prefetch matching delay. Controls the amount of time a prefetch waits for a possible match with a later read. Encoded as powers of 2, from 1-128.

These bits reset to 0x5.

**Disable interconnect Cacheable atomics, [7]**

Disable Cacheable atomics being sent to the interconnect. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | If BROADCASTATOMIC is set HIGH, Cacheable atomics are sent to the interconnect. |
| <b>1</b> | Cacheable atomics are handled inside the cluster.                               |

This bit has no effect for ACE configurations.

This bit resets to 0.

#### **RAZ, [6:5]**

Read-As-Zero

#### **Interconnect data poisoning support, [4]**

Interconnect data poisoning support. This bit is RAZ for ACE configurations. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | Interconnect does not support data poisoning. Therefore nERRIRQ is asserted when poisoned data is evicted from the cluster or returned to a snoop.              |
| <b>1</b> | Interconnect supports data poisoning. Therefore no error recovery interrupt is generated when poisoned data is evicted from the cluster or returned to a snoop. |

For ACE, this bit resets to 0.

For CHI, this bit resets to 1.

#### **Cache Evict disable, [3]**

Disables sending of Evict transactions on the ACE or CHI master for clean cache lines that are evicted from the cluster. Evict transactions are required only if the external interconnect contains a snoop filter that requires notification when the cluster evicts the cache line. The possible values are:

- |          |                                      |
|----------|--------------------------------------|
| <b>0</b> | Enables sending Evict transactions.  |
| <b>1</b> | Disables sending Evict transactions. |

This bit resets to 0.

#### **Cache flush UniqueClean eviction control, [2]**

Disables the sending of WriteEvict requests on the ACE or CHI master when powering down part or all of the L3 cache. The possible values are:

- |          |  |
|----------|--|
| <b>0</b> | Evictions during L3 cache powerdown behave like normal evictions. This is the reset value. |
| <b>1</b> | Disables sending data with UniqueClean evictions caused by powering down the L3 cache.     |

#### **RAZ, [1]**

Read-As-Zero

#### **Non-cacheable behavior control, [0]**

Enable Normal Non-cacheable writes to all master interfaces and, when ACE is configured, also disable the limit on the number of Normal Non-cacheable writes. The possible values are:

- 0** All Normal Non-cacheable and Device transactions are sent to interface 0.
- When ACE is configured, Device and Normal Non-cacheable writes are limited to 15 outstanding transactions.
- 1** If dual master interfaces are configured, Normal Non-cacheable transactions are interleaved between master interfaces like Cacheable transactions. Device transactions are still sent to interface 0.
- When ACE is configured, only Device writes are limited to 15 outstanding transactions. The limit is removed for Normal Non-cacheable transactions.

If this bit is set to 1, then further writes to this bit are ignored.



Setting this bit might have implications on the behavior of system components, for example a CPE-425 Coherent PCIe Extension in an ACE system.

## Configurations

The AArch32 CLUSTERECTLR register is architecturally mapped to the AArch64 CLUSTERECTLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERECTLR

#### In AArch64 state (CLUSTERECTLR\_EL1):

To read this register in AArch64 state (CLUSTERECTLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_4; Read CLUSTERECTLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERECTLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_4, <Xt>; Write Xt into CLUSTERECTLR_EL1
```

#### In AArch32 state (CLUSTERECTLR):

To read this register in AArch32 state (CLUSTERECTLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 4; Read CLUSTERECTLR into Rt
```

To write this register in AArch32 state (CLUSTERECTLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 4; Write Rt into CLUSTERECTLR
```

### Accessibility

This register is accessible in software as follows:

| <systemreg>  | Control |     |    | Accessibility |     |     |     |
|--------------|---------|-----|----|---------------|-----|-----|-----|
|              | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERECTLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERECTLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERECTLR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.ECTLREN is 1 and ACTLR\_EL2.ECTLREN is 1, or ACTLR\_EL3.ECTLR\_EN is 1 and SCR.NS is 0.

If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.9 CLUSTERIDR, Cluster Main Revision ID Register


The CLUSTERIDR register contains the revision and patch level of the DSU. This register is read-only write-ignores RO (WI), and is common to all execution threads.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERIDR) and AArch64 (CLUSTERIDR\_EL1) registers.

CLUSTERIDR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-5: CLUSTERIDR bit assignments**



**RAZ, [31:8]**

Read-As-Zero.

**Variant, [7:4]**

Indicates the variant of the DSU. This is the major revision number x in the rx part of the rxy description of the product revision status. This value is:

**0x4**                      r4

**Revision, [3:0]**

Indicates the minor revision number of the DSU. This is the minor revision number y in the py part of the rxy description of the product revision status. This value is:

**0x1**                      p1

**Configurations**

The AArch32 CLUSTERIDR register is architecturally mapped to the AArch64 CLUSTERIDR\_EL1 register.

**Usage constraints****Accessing the CLUSTERIDR****In AArch64 state (CLUSTERIDR\_EL1):**

To read this register in AArch64 state (CLUSTERIDR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_1; Read CLUSTERIDR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERIDR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_1, <Xt>; Write Xt into CLUSTERIDR_EL1
```

**In AArch32 state (CLUSTERIDR):**

To read this register in AArch32 state (CLUSTERIDR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 1; Read CLUSTERIDR into Rt
```

To write this register in AArch32 state (CLUSTERIDR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 1; Write Rt into CLUSTERIDR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg> | Control |     |    | Accessibility |       |       |       |
|-------------|---------|-----|----|---------------|-------|-------|-------|
|             | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERIDR  | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERIDR  | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERIDR  | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

There are no traps and enables that affect this register.

# 11.10 CLUSTERL3HIT, Cluster L3 Hit Counter Register

The CLUSTERL3HIT register contains a count of the L3 cache hits. This register is intended for use in algorithms for determining when to power up or power down portions. This register is RW, and is common to all execution threads.


## Bit field descriptions

This description applies to both the AArch32 (CLUSTERL3HIT) and AArch64 (CLUSTERL3HIT\_EL1) registers.

If CLUSTERPMMDCR.SPME == 0, this counter does not trigger for secure transactions.

CLUSTERL3HIT is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-6: CLUSTERL3HIT bit assignments**



## L3 cache hit count, [31:0]

Count of number of L3 cache hits, for use in portion control calculations. The counter saturates to its maximum value on overflow.

On a write, CLUSTERL3HIT gets set to the specified value.

These bits reset to 0x00000000.

## Configurations

The AArch32 CLUSTERL3HIT register is architecturally mapped to the AArch64 CLUSTERL3HIT\_EL1 register.

## Usage constraints

### Accessing the CLUSTERL3HIT

#### In AArch64 state (CLUSTERL3HIT\_EL1):

To read this register in AArch64 state (CLUSTERL3HIT\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_5; Read CLUSTERL3HIT_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERL3HIT\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_5, <Xt>; Write Xt into CLUSTERL3HIT_EL1
```

#### In AArch32 state (CLUSTERL3HIT):

To read this register in AArch32 state (CLUSTERL3HIT) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 5; Read CLUSTERL3HIT into Rt
```

To write this register in AArch32 state (CLUSTERL3HIT) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 5; Write Rt into CLUSTERL3HIT
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>  | Control |     |    | Accessibility |     |     |     |
|--------------|---------|-----|----|---------------|-----|-----|-----|
|              | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERL3HIT | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERL3HIT | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERL3HIT | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.PWREN is 1 and ACTLR\_EL2.PWREN is 1, or ACTLR\_EL3.PWREN is 1 and SCR.NS is 0.



If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.11 CLUSTERL3MISS, Cluster L3 Miss Counter Register

The CLUSTERL3MISS register contains a count of the L3 cache misses. This register is intended for use in algorithms for determining when to power up or power down portions. This register is RW, and is common to all execution threads.


### Bit field descriptions

This description applies to both the AArch32 (CLUSTERL3MISS) and AArch64 (CLUSTERL3MISS\_EL1) registers.

If CLUSTERPMMDCCR.SPME == 0, this counter does not trigger for secure transactions.

CLUSTERL3MISS is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-7: CLUSTERL3MISS bit assignments**



### L3 cache miss count, [31:0]

Count of number of L3 cache misses, for use in portion control calculations. The counter saturates to its maximum value on overflow.

On a write, CLUSTERL3MISS gets set to the specified value.

These bits reset to 0x00000000.

### Configurations

The AArch32 CLUSTERL3MISS register is architecturally mapped to the AArch64 CLUSTERL3MISS\_EL1 register.

### Usage constraints

#### Accessing the CLUSTERL3MISS

##### In AArch64 state (CLUSTERL3MISS\_EL1):

To read this register in AArch64 state (CLUSTERL3MISS\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_6; Read CLUSTERL3MISS_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERL3MISS\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_6, <Xt>; Write Xt into CLUSTERL3MISS_EL1
```

#### In AArch32 state (CLUSTERL3MISS):

To read this register in AArch32 state (CLUSTERL3MISS) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 6; Read CLUSTERL3MISS into Rt
```

To write this register in AArch32 state (CLUSTERL3MISS) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 6; Write Rt into CLUSTERL3MISS
```

#### Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERL3MISS | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERL3MISS | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERL3MISS | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

#### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.PWREN is 1 and ACTLR\_EL2.PWREN is 1, or ACTLR\_EL3.PWREN is 1 and SCR.NS is 0.

If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.12 CLUSTERPARTCR, Cluster Partition Control Register


The CLUSTERPARTCR register controls a group of ways to be marked as private to a scheme ID. This register is RW.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPARTCR) and AArch64 (CLUSTERPARTCR\_EL1) registers.

CLUSTERPARTCR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-8: CLUSTERPARTCR bit assignments**



Each bit, if set, indicates that a group of four ways is allocated as private to that scheme ID. If more than one scheme ID assigns the same group of ways as private, then those ways are shared between the scheme IDs that have assigned them as private. All ways not assigned to any scheme

ID are treated as shared between all scheme IDs. If a scheme ID does not have any private ways allocated, and there are no remaining shared ways, then any use of the scheme ID will allocate to way group 0, as this is considered a programming error.

**Way group 3 is assigned as private to scheme ID 7, [31]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 7, [30]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 7, [29]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 7, [28]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 6, [27]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 6, [26]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 6, [25]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 6, [24]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 5, [23]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 5, [22]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 5, [21]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 5, [20]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 4, [19]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 4, [18]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 4, [17]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 4, [16]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 3, [15]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 3, [14]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 3, [13]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 3, [12]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 2, [11]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 2, [10]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 2, [9]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 2, [8]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 1, [7]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 1, [6]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 1, [5]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 1, [4]**

This bit resets to 0.

**Way group 3 is assigned as private to scheme ID 0, [3]**

This bit resets to 0.

**Way group 2 is assigned as private to scheme ID 0, [2]**

This bit resets to 0.

**Way group 1 is assigned as private to scheme ID 0, [1]**

This bit resets to 0.

**Way group 0 is assigned as private to scheme ID 0, [0]**

This bit resets to 0.

**Configurations**

The AArch32 CLUSTERPARTCR register is architecturally mapped to the AArch64 CLUSTERPARTCR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPARTCR

#### In AArch64 state (CLUSTERPARTCR\_EL1):

To read this register in AArch64 state (CLUSTERPARTCR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_3; Read CLUSTERPARTCR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPARTCR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_3, <Xt>; Write Xt into CLUSTERPARTCR_EL1
```

#### In AArch32 state (CLUSTERPARTCR):

To read this register in AArch32 state (CLUSTERPARTCR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 3; Read CLUSTERPARTCR into Rt
```

To write this register in AArch32 state (CLUSTERPARTCR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 3; Write Rt into CLUSTERPARTCR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPARTCR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPARTCR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPARTCR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.SMEN is 1 and ACTLR\_EL2.SMEN is 1, or ACTLR\_EL3.SMEN is 1 and SCR.NS is 0.

If write access is not permitted, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.13 CLUSTERPWRCTLR, Cluster Power Control Register


The CLUSTERPWRCTLR register controls power features of the cluster. This register is RW, and is common to all execution threads.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPWRCTLR) and AArch64 (CLUSTERPWRCTLR\_EL1) registers.

CLUSTERPWRCTLR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

Figure 11-9: CLUSTERPWRCTLR bit assignments



#### RAZ, [31:8]

Read-As-Zero.

#### Cache portion power request, [7:4]

These bits are output on CLUSTERACTIVE[19:16] to indicate to the power controller which cache portions must be powered. This is an advisory status to the power controller, and does not cause the P-Channel to deny any requests that are based on this status.

The reset values of these bits depend on the cluster P-Channel initialization state.

#### RAZ, [3]

Read-As-Zero.

#### Functional retention control, [2:0]

Duration of inactivity before the DSU uses CLUSTERACTIVE to request functional retention. The possible values of these bits are:

|       |  |
|-------|--|
| 0b000 | Disable the retention circuit. This is the default condition.        |
| 0b001 | Two Architectural Timer ticks are required before retention entry.   |
| 0b010 | Eight Architectural Timer ticks are required before retention entry. |
| 0b011 | 32 Architectural Timer ticks are required before retention entry.    |

|              |  |
|--------------|--|
| <b>0b100</b> | 64 Architectural Timer ticks are required before retention entry.  |
| <b>0b101</b> | 128 Architectural Timer ticks are required before retention entry. |
| <b>0b110</b> | 256 Architectural Timer ticks are required before retention entry. |
| <b>0b111</b> | 512 Architectural Timer ticks are required before retention entry. |

These bits reset to 0b000.

## Configurations

The AArch32 CLUSTERPWRCTLR register is architecturally mapped to the AArch64 CLUSTERPWRCTLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPWRCTLR

#### In AArch64 state (CLUSTERPWRCTLR\_EL1):

To read this register in AArch64 state (CLUSTERPWRCTLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_5; Read CLUSTERPWRCTLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPWRCTLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_5, <Xt>; Write Xt into CLUSTERPWRCTLR_EL1
```

#### In AArch32 state (CLUSTERPWRCTLR):

To read this register in AArch32 state (CLUSTERPWRCTLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 5; Read CLUSTERPWRCTLR into Rt
```

To write this register in AArch32 state (CLUSTERPWRCTLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 5; Write Rt into CLUSTERPWRCTLR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>    | Control |     |    | Accessibility |     |     |     |
|----------------|---------|-----|----|---------------|-----|-----|-----|
|                | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPWRCTLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPWRCTLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPWRCTLR | x       | 1   | 1  | -             | n/a | RW  | RW  |



'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.PWREN is 1 and ACTLR\_EL2.PWREN is 1, or ACTLR\_EL3.PWREN is 1 and SCR.NS is 0.

If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.14 CLUSTERPWRDN, Cluster Powerdown Register


The CLUSTERPWRDN register controls powerdown requirements of the cluster. This register is RW, and is banked for each thread of execution.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPWRDN) and AArch64 (CLUSTERPWRDN\_EL1) registers.

CLUSTERPWRDN is a 32-bit register, and is part of SCU and L3 cache configuration registers.

### Figure 11-10: CLUSTERPWRDN bit assignments



## RAZ, [31:2]

Read-As-Zero.

## Memory retention required, [1]

The possible values are:

- |   |  |
|---|--|
| 0 | Indicates on CLUSTERPACTIVE that memory retention is not required when all cores are powered down. |
| 1 | Indicates on CLUSTERPACTIVE that memory retention is required when all cores are powered down.     |

This signal provides an advisory status to the power controller. It does not deny a P-Channel request to powerdown the cluster.

This bit resets to 0.

## Cluster power required, [0]

The possible values are:

- |          |  |
|----------|--|
| <b>0</b> | Indicates on CLUSTERPACTIVE that cluster power is not required when all cores are powered down.  |
| <b>1</b> | Indicates on CLUSTERPACTIVE that cluster power is required even when all cores are powered down. |

This signal provides an advisory status to the power controller. It does not deny a P-Channel request to powerdown the cluster.

This bit resets to 0.

## Configurations

The AArch32 CLUSTERPWRDN register is architecturally mapped to the AArch64 CLUSTERPWRDN\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPWRDN

#### In AArch64 state (CLUSTERPWRDN\_EL1):

To read this register in AArch64 state (CLUSTERPWRDN\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_6; Read CLUSTERPWRDN_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPWRDN\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_6, <Xt>; Write Xt into CLUSTERPWRDN_EL1
```

#### In AArch32 state (CLUSTERPWRDN):

To read this register in AArch32 state (CLUSTERPWRDN) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 6; Read CLUSTERPWRDN into Rt
```

To write this register in AArch32 state (CLUSTERPWRDN) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 6; Write Rt into CLUSTERPWRDN
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>  | Control |     |    | Accessibility |     |     |     |
|--------------|---------|-----|----|---------------|-----|-----|-----|
|              | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPWRDN | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPWRDN | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPWRDN | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.PWREN is 1 and ACTLR\_EL2.PWREN is 1, or ACTLR\_EL3.PWREN is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.15 CLUSTERPWRSTAT, Cluster Power Status Register


The CLUSTERPWRSTAT register contains the status of the power features. This register is read-only write-ignores RO (WI), and is common to all execution threads.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPWRSTAT) and AArch64 (CLUSTERPWRSTAT\_EL1) registers.

CLUSTERPWRSTAT is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-11: CLUSTERPWRSTAT bit assignments**



### RAZ, [31:8]

Read-As-Zero.

### Cache portion power status, [7:4]

These bits indicate which cache portions are currently powered up and available. These bits can be used to determine when the state requested in bits [7:4] of the CLUSTERPWRCTLR has taken effect. The possible values are:

**0b1111** Ways 0-15 are powered up.

|               |                           |
|---------------|---------------------------|
| <b>0b0111</b> | Ways 0-11 are powered up. |
| <b>0b0011</b> | Ways 0-7 are powered up.  |
| <b>0b0001</b> | Ways 0-3 are powered up.  |
| <b>0b0000</b> | No ways are powered up.   |

The reset value for these bits depends on the P-Channel initialization state.

### RAZ, [3:2]

Read-As-Zero.

### Retention when powered down, [1]

Enabled memory retention when all cores are powered down. This bit is a combined version of all banked bits for each execution thread from the CLUSTERPWRDN register. The possible values are:

|          |  |
|----------|--|
| <b>0</b> | Disables memory retention when all cores are powered down. |
| <b>1</b> | Enables memory retention when all cores are powered down.  |

This bit resets to 0.

### Disable cluster powerdown, [0]

Disabled cluster powerdown when all cores are powered down. This bit is a combined version of all banked per-thread bits from the CLUSTERPWRDN register. The possible values are:

|          |   |
|----------|---|
| <b>0</b> | Enables cluster powerdown when all cores are powered down.  |
| <b>1</b> | Disables cluster powerdown when all cores are powered down. |

This bit resets to 0.

## Configurations

The AArch32 CLUSTERPWRSTAT register is architecturally mapped to the AArch64 CLUSTERPWRSTAT\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPWRSTAT

#### In AArch64 state (CLUSTERPWRSTAT\_EL1):

To read this register in AArch64 state (CLUSTERPWRSTAT\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_7; Read CLUSTERPWRSTAT_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPWRSTAT\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_7, <Xt>; Write Xt into CLUSTERPWRSTAT_EL1
```

**In AArch32 state (CLUSTERPWRSTAT):**

To read this register in AArch32 state (CLUSTERPWRSTAT) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 7; Read CLUSTERPWRSTAT into Rt
```

To write this register in AArch32 state (CLUSTERPWRSTAT) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 7; Write Rt into CLUSTERPWRSTAT
```

**Accessibility**

This register is accessible in software as follows:

| <systemreg>    | Control |     |    | Accessibility |       |       |       |
|----------------|---------|-----|----|---------------|-------|-------|-------|
|                | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERPWRSTAT | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERPWRSTAT | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERPWRSTAT | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

**Traps and enables**

This register is write accessible in EL1 if ACTLR\_EL3.PWREN is 1 and ACTLR\_EL2.PWREN is 1, or ACTLR\_EL3.PWREN is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).


## 11.16 CLUSTERREVIDR, Cluster Revision ID Register

The CLUSTERREVIDR register enables ECO patches applied to the cluster level to be identified by the operating system. This register is read-only write-ignores RO (WI), and is common to all execution threads.

**Bit field descriptions**

This description applies to both the AArch32 (CLUSTERREVIDR) and AArch64 (CLUSTERREVIDR\_EL1) registers.

CLUSTERREVIDR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-12: CLUSTERREVIDR bit assignments****ID number, [31:0]**

Implementation-specific revision information. The reset value is determined by the specific DSU implementation.

**0x00000000** Revision code is zero.

**Configurations**

The AArch32 CLUSTERREVIDR register is architecturally mapped to the AArch64 CLUSTERREVIDR\_EL1 register.

**Usage constraints****Accessing the CLUSTERREVIDR****In AArch64 state (CLUSTERREVIDR\_EL1):**

To read this register in AArch64 state (CLUSTERREVIDR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C3_2; Read CLUSTERREVIDR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERREVIDR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C3_2, <Xt>; Write Xt into CLUSTERREVIDR_EL1
```

**In AArch32 state (CLUSTERREVIDR):**

To read this register in AArch32 state (CLUSTERREVIDR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c3, 2; Read CLUSTERREVIDR into Rt
```

To write this register in AArch32 state (CLUSTERREVIDR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c3, 2; Write Rt into CLUSTERREVIDR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |       |       |       |
|---------------|---------|-----|----|---------------|-------|-------|-------|
|               | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERREVIDR | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERREVIDR | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERREVIDR | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

There are no traps and enables affecting this register.

## 11.17 CLUSTERSTASHSID, Cluster Stash Scheme ID Register


The CLUSTERSTASHSID register provides the scheme ID for stash requests received from the interconnect.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERSTASHID) and AArch64 (CLUSTERSTASHID\_EL1) registers.

CLUSTERSTASHSID is a 32-bit register, and is part of SCU and L3 cache configuration registers.

### Figure 11-13: CLUSTERSTASHSID bit assignments



## RAZ, [31:3]

Read-As-Zero.

## Scheme ID for stash requests, [2:0]

Scheme ID for stash requests received from the interconnect.

These bits reset to 0b000.

## Configurations

The AArch32 CLUSTERSTASHSID is architecturally mapped to the AArch64 CLUSTERSTASHSID\_EL1 register.

## Usage constraints

### Accessing the CLUSTERSTASHSID

#### In AArch64 state (CLUSTERSTASHSID\_EL1):

To read this register in AArch64 state (CLUSTERSTASHSID\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_2; Read CLUSTERSTASHSID_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERSTASHSID\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_2, <Xt>; Write Xt into CLUSTERSTASHSID_EL1
```

#### In AArch32 state (CLUSTERSTASHSID):

To read this register in AArch32 state (CLUSTERSTASHSID) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 2; Read CLUSTERSTASHSID into Rt
```

To write this register in AArch32 state (CLUSTERSTASHSID) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 2; Write Rt into CLUSTERSTASHSID
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>     | Control |     |    | Accessibility |     |     |     |
|-----------------|---------|-----|----|---------------|-----|-----|-----|
|                 | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERSTASHSID | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERSTASHSID | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERSTASHSID | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.SMEN is 1 and ACTLR\_EL2.SMEN is 1, or ACTLR\_EL3.SMEN is 1 and SCR.NS is 0.



If write access is not permitted, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.18 CLUSTERTHREADSID, Cluster Thread Scheme ID Register


The CLUSTERTHREADSID register is banked for each thread of execution. For multithreaded cores the scheme ID is generated by a logical OR of the thread ID registers for that core.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERTHREADSID) and AArch64 (CLUSTERTHREADSID\_EL1) registers.

CLUSTERTHREADSID is a 32-bit register, and is part of SCU and L3 cache configuration registers.

### Figure 11-14: CLUSTERTHREADSID bit assignments



## RAZ, [31:3]

Read-As-Zero.

## Scheme ID for current thread, [2:0]

Scheme ID for this thread.

These bits reset to 0b000.

## Configurations

The AArch32 CLUSTERTHREADSID is architecturally mapped to the AArch64 CLUSTERTHREADSID\_EL1 register.

## Usage constraints

## Accessing the CLUSTERTHREADSID

In AArch64 state (CLUSTERTHREADSID\_EL1):

To read this register in AArch64 state (CLUSTERTHREADSID\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3 0 C15 C4 0; Read CLUSTERTHREADSID EL1 into Xt
```

To write this register in AArch64 state (CLUSTERTHREADSID\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_0, <Xt>; Write Xt into CLUSTERTHREADSID_EL1
```

#### In AArch32 state (CLUSTERTHREADSID):

To read this register in AArch32 state (CLUSTERTHREADSID) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 0; Read CLUSTERTHREADSID into Rt
```

To write this register in AArch32 state (CLUSTERTHREADSID) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 0; Write Rt into CLUSTERTHREADSID
```

#### Accessibility

This register is accessible in software as follows:

| <systemreg>      | Control |     |    | Accessibility |     |     |     |
|------------------|---------|-----|----|---------------|-----|-----|-----|
|                  | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERTHREADSID | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERTHREADSID | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERTHREADSID | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

#### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.TSIDEN is 1 and ACTLR\_EL2.TSIDEN is 1, or ACTLR\_EL3.TSIDEN is 1 and SCR.NS is 0.

If write access is not permitted, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 11.19 CLUSTERTHREADSIDOVR, Cluster Thread Scheme ID Override Register

The CLUSTERTHREADSIDOVR register can be used by more privileged software, for example a hypervisor, to partition the scheme IDs. Therefore allowing the less privileged software to use only


a subset of scheme IDs. This register is RW. The CLUSTERTHREADSIDOVR register is banked for each thread of execution.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERTHREADSIDOVR) and AArch64 (CLUSTERTHREADSIDOVR\_EL1) registers.

CLUSTERTHREADSIDOVR is a 32-bit register, and is part of SCU and L3 cache configuration registers.

**Figure 11-15: CLUSTERTHREADSIDOVR bit assignments**



#### RAZ, [31:19]

Read-As-Zero.

#### Scheme ID mask, [18:16]

A bit set in the mask causes the matching bit to be taken from this register rather than from the CLUSTERTHREADSID register.

These bits reset to 0b000.

#### RAZ, [15:3]

Read-As-Zero.

#### Scheme ID for this thread if masked, [2:0]

Scheme ID for this thread if masked.

These bits reset to 0b000.

### Configurations

The AArch32 CLUSTERTHREADSIDOVR register is architecturally mapped to the AArch64 CLUSTERTHREADSIDOVR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERTHREADSIDOVR

#### In AArch64 state (CLUSTERTHREADSIDOVR\_EL1):

To read this register in AArch64 state (CLUSTERTHREADSIDOVR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C4_7; Read CLUSTERTHREADSIDOVR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERTHREADSIDOVR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C4_7, <Xt>; Write Xt into CLUSTERTHREADSIDOVR_EL1
```

#### In AArch32 state (CLUSTERTHREADSIDOVR):

To read this register in AArch32 state (CLUSTERTHREADSIDOVR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c4, 7; Read CLUSTERTHREADSIDOVR into Rt
```

To write this register in AArch32 state (CLUSTERTHREADSIDOVR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c4, 7; Write Rt into CLUSTERTHREADSIDOVR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>         | Control |     |    | Accessibility |     |     |     |
|---------------------|---------|-----|----|---------------|-----|-----|-----|
|                     | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERTHREADSIDOVR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERTHREADSIDOVR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERTHREADSIDOVR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3.SMEN is 1 and ACTLR\_EL2.SMEN is 1, or ACTLR\_EL3.SMEN is 1 and SCR.NS is 0.

If write access is not permitted, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 12. Error system registers

This chapter describes the ERR1\* error registers for the DSU.

### 12.1 About the error system registers

The error system registers are accessed by both the AArch32 ERX\* error registers and the AArch64 ERX\* error registers. The ERR1\* registers are agnostic to the architectural state. For example, if ERRSELR==1 and ERRSELR\_EL1==1, then ERR1PFGF is accessed by both ERXPFGFR and ERXPFGFR\_EL1.

The chapter is presented as follows:

#### Error system register summary

This section identifies the ERR1\* error registers and lists the corresponding ERX\* registers that can access them.

#### Register descriptions

The remainder of the chapter provides register descriptions. They are listed in alphabetical order.



The contents of the error system registers are preserved over a Warm reset.

---

See appropriate core Technical Reference Manual for more information about the ERX\* registers.

### 12.2 Error system register summary

This section lists the DSU ERR1\* error record registers that are accessed from the core AArch32 and AArch64 ERX\* error registers.

The ERR1\* registers are agnostic to the architectural state. For example, if ERRSELR==1 and ERRSELR\_EL1==1, then ERR1PFGF is accessed by both ERXPFGFR and ERXPFGFR\_EL1.

For a description of the ERX\* error registers see the related core documentation.

The following table describes the architectural error record registers in the DSU.

**Table 12-1: Architectural error system register summary**

| Register mnemonic | Size | Register name and description  | Access aliases from AArch32 and AArch64                                     |
|-------------------|------|--|---|
| ERR1CTLR          | 64   | <a href="#">12.3 ERR1CTLR, Error Record Control Register</a> on page 134           | ERXCTLR, ERXCTLR2, ERXCTLR_EL1  |
| ERR1FR            | 64   | <a href="#">12.4 ERR1FR, Error Record Feature Register</a> on page 136             | ERXFR, ERXFR2, ERXFR_EL1  |
| ERR1MISCO         | 64   | <a href="#">12.5 ERR1MISCO, Error Record Miscellaneous Register 0</a> on page 138  | ERXMISCO, ERXMISC1, ERXMISCO_EL1  |
| ERR1MISC1         | 64   | <a href="#">12.6 ERR1MISC1, Error Record Miscellaneous Register 1</a> on page 140  | ERXMISC2 accesses bits [31:0], ERXMISC3 accesses bits [63:32], ERXMISC1_EL1 |
| ERR1STATUS        | 32   | <a href="#">12.10 ERR1STATUS, Error Record Primary Status Register</a> on page 145 | ERXSTATUS, ERXSTATUS_EL1  |

The following table describes the error record registers that are **IMPLEMENTATION DEFINED** in the DSU.

**Table 12-2: Implementation defined error system register summary**

| Register mnemonic | Size | Register name   | Access aliases from AArch32 and AArch64 |
|-------------------|------|---|---|
| ERR1PFGCDNR       | 32   | <a href="#">12.7 ERR1PFGCDNR, Error Pseudo Fault Generation Count Down Register</a> on page 141 | ERXPFGCDNR, ERXPFGCDNR_EL1              |
| ERR1PFGCTLR       | 32   | <a href="#">12.8 ERR1PFGCTLR, Error Pseudo Fault Generation Control Register</a> on page 141    | ERXPFGCTLR, ERXPFGCTLR_EL1              |
| ERR1PFGFR         | 32   | <a href="#">12.9 ERR1PFGFR, Error Pseudo Fault Generation Feature Register</a> on page 143      | ERR1PFGFR, ERXPFGFR_EL1                 |

## 12.3 ERR1CTLR, Error Record Control Register


The ERR1CTLR register contains enable bits for the node that writes to this record:

- Enabling error detection and correction.
- Enabling an error recovery interrupt.
- Enabling a fault handling interrupt.
- Enabling error recovery reporting as a read or write error response.

### Bit field descriptions

ERR1CTLR is a 64-bit register and is part of the RAS registers functional group.

Figure 12-1: ERR1CTLR bit assignments



RES0, [63:9]


RES0      Reserved.

CFI, [8]

Fault handling interrupt for corrected errors enable.  
The fault handling interrupt is generated when one of the standard CE counters on ERR1MISCO overflows and the overflow bit is set. The possible values are:

- |   |   |
|---|---|
| 0 | Fault handling interrupt is not generated for corrected errors. |
| 1 | Fault handling interrupt is generated for corrected errors.     |

The interrupt is generated even if the error status is overwritten because the error record already records a higher priority error. If the node does not support this control, this bit is RES0.



Note

This condition applies to both reads and writes.

RES0, [7:4]

RES0      Reserved.

FI, [3]

Fault handling interrupt enable.  
The fault handling interrupt is generated for all detected Deferred errors and Uncorrected errors. The possible values are:

- |   |                                    |
|---|------------------------------------|
| 0 | Fault handling interrupt disabled. |
| 1 | Fault handling interrupt enabled.  |

UI, [2]

Uncorrected error recovery interrupt enable. When enabled, the error recovery interrupt is generated for all detected Uncorrected errors that are not deferred. The possible values are:

|   |                                       |
|---|---------------------------------------|
| 0 | Error recovery interrupt is disabled. |
| 1 | Error recovery interrupt enabled.     |



Applies to both reads and writes.

## RES0, [1]

|      |           |
|------|-----------|
| RES0 | Reserved. |
|------|-----------|

## ED, [0]

Error reporting and logging enable. When reporting is disabled, the DSU will not record or signal errors, but it will continue to detect and correct errors. The possible values of this bit are:

|   |  |
|---|--|
| 0 | Error reporting is disabled. In this state, uncorrectable errors might result in corrupt data being silently propagated. |
| 1 | Error reporting is enabled.  |

Correct error detection and correction codes are written for writes even when reporting is disabled.

## Configurations

ERR1CTLR resets to 0x0000000000000000.

This register is accessible from the following registers when ERRSELR.SEL==1:

- [31:0]: ERXCTLR.
- [63:32]: ERXCTLR2.
- ERXCTLR\_EL1.

See the appropriate core Technical Reference Manual for information about these ERX\* registers.

# 12.4 ERR1FR, Error Record Feature Register

The ERR1FR register defines which of the common architecturally defined features are implemented and, of the implemented features, which are software programmable.


## Bit field descriptions

ERR1FR is a 64-bit register, and is part of the *Reliability, Availability, Serviceability* (RAS) registers functional group.

The register is Read Only.



### Figure 12-2: ERR1FR bit assignments



## RAZ/WI, [63:20]

Read-as-zero/Write ignore.

CEO, [19:18]

Corrected Error Overwrite. The value is:

**0b00** Counts CE if a counter is implemented and keeps the previous error status. If the counter overflows, or no counter is implemented, ERROSTATUS.OF is set to 1.

## DUI, [17:16]

Error recovery interrupt for deferred errors. The value is:

**0b00** The core or cluster does not support this feature.

RP, [15]

Repeat counter. The value is:

- 1 A first repeat counter and a second other counter are implemented. The repeat counter is the same size as the primary error counter.

## CEC, [14:12]

Corrected Error Counter. The value is:

**0b010** The node implements an 8-bit standard CE counter in ERRORMISC0[39:32].

## CFI, [11:10]

Fault handling interrupt for corrected errors. The value is:

**0b10** The node implements a control for enabling fault handling interrupts on corrected errors.

UE, [9:8]

In-band uncorrected error reporting. The value is:

**0b01** The node implements in-band uncorrected error reporting, that is external aborts.

#### FI, [7:6]

Fault handling interrupt. The value is:

**0b10** The node implements a fault handling interrupt and implements controls for enabling and disabling.

#### UI, [5:4]

Error recovery interrupt for uncorrected errors. The value is:

**0b10** The node implements an error recovery interrupt and implements controls for enabling and disabling.

#### DE, [3:2]

Defers Errors enable. The value is:

**0b01** Defers Errors are always enabled.

#### ED, [1:0]

Error detection and correction. The value is:

**0b10** Error detection is controllable.

### Configurations

This register is accessible from the following registers when ERRSELR.SEL==1:

- [31:0]: ERXFR.
- [63:32]: ERXFR2.
- ERXFR\_EL1.


See the appropriate core Technical Reference Manual for information about these ERX\* registers.

## 12.5 ERR1MISC0, Error Record Miscellaneous Register 0

The ERR1MISC0 register is an error syndrome register. It contains information on the corrected error count and information to identify where the error was detected. It also contains other state information not present in the corresponding status and address error record registers.

### Bit field descriptions

ERR1MISC0 is a 64-bit register, and is part of the *Reliability, Availability, Serviceability* (RAS) registers functional group.

**Figure 12-3: ERR1MISC0 bit assignments****RES0, [63:48]**

**RES0** Reserved.

**OFO, [47]**

Sticky overflow bit, other. The possible values are:

|          |                                   |
|----------|-----------------------------------|
| <b>0</b> | Other counter has not overflowed. |
| <b>1</b> | Other counter has overflowed.     |

The fault handling interrupt is generated when the corrected fault handling interrupt is enabled and either overflow bit is set to 1.

**CECO, [46:40]**

Corrected error count, other. Incremented for each Corrected error that does not match the recorded syndrome.

**OFR, [39]**

Sticky overflow bit, repeat. The possible values are:

|          |                                    |
|----------|------------------------------------|
| <b>0</b> | Repeat counter has not overflowed. |
| <b>1</b> | Repeat counter has overflowed.     |

The fault handling interrupt is generated when the corrected fault handling interrupt is enabled and either overflow bit is set to 1.

**CECR, [38:32]**

Corrected error count, repeat. Incremented for the first recorded error, which also records other syndrome, and then again for each Corrected error that matches the recorded syndrome.

**WAY, [31:28]**

Indicates the way in the L3 tag RAM, L3 data RAM, or snoop filter RAM that contained the error. If the error is in the Long-Term Data Buffer (LTDB) RAM, then this indicates the RAM instance that contained the error.

- For the L3 tag RAM, L3 data RAM, and LTDB RAM all four bits are used.
- For the snoop filter RAM, only bits[30:28] are used and bit[31] is set **RES0**.

**RES0, [27:19]**

|             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**INDX, [18:6]**

Indicates the index that contained the error.

Upper bits of the index are unused depending on the cache size.

**RES0, [5:4]**

|             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**LVL, [3:1]**

Indicates the level that contained the error. The value is:

|              |          |
|--------------|----------|
| <b>0b010</b> | Level 3. |
|--------------|----------|

**IND, [0]**

Indicates the type of cache that contained the error. The value is:

|          |           |
|----------|-----------|
| <b>0</b> | L3 cache. |
|----------|-----------|

**Configurations**

ERR1MISCO resets to 0x0000000000000000.

This register is accessible from the following registers when ERRSELR.SEL==1:

- [31:0]: ERXMISCO.
- [63:32]: ERXMISC1.
- ERXMISCO\_EL1.

See the appropriate core Technical Reference Manual for information about these ERX\* registers.

## 12.6 ERR1MISC1, Error Record Miscellaneous Register 1

The ERR1MISC1 register is not used in the DSU and marked as **RES0**.

**Configurations**

ERR1MISC1 is accessible from the following registers when ERRSELR.SEL==1:

- [31:0]: ERXMISC2.
- [63:32]: ERXMISC3.
- ERXMISC1\_EL1.

See the appropriate core Technical Reference Manual for information about these ERX\* registers.


## 12.7 ERR1PFGCDNR, Error Pseudo Fault Generation Count Down Register

ERR1PFGCDNR is the DSU node register that generates one of the errors enabled in the corresponding ERR1PFGCTL register.

### Bit field descriptions

ERR1PFGCDNR is a 32-bit read/write register.

**Figure 12-4: ERR1PFGCDNR bit assignments**



### CDN, [31:0]

Count Down value. The reset value of the Error Generation Counter is used for the countdown.

### Configurations

ERR1PFGCDNR resets to 0x00000000.

There are no configuration options.

ERR1PFGCDNR is accessible from the following registers when ERRSELR.SEL==1:

- ERXPFGCDNR.
- ERXPFGCDNR\_EL1.


See the appropriate core Technical Reference Manual for information about these ERX\* registers.

## 12.8 ERR1PFGCTLR, Error Pseudo Fault Generation Control Register

The ERR1PFGCTLR register is the DSU node register that enables controlled fault generation.

### Bit field descriptions

ERR1PFGCTLR is a 32-bit read/write register.

**Figure 12-5: ERR1PFGCTLR bit assignments****CDNEN, [31]**

Count down enable. This bit controls transfers from the value that is held in the ERR1PFGCDNR into the Error Generation Counter and enables this counter to start counting down. The possible values are:

- |   |  |
|---|--|
| 0 | The Error Generation Counter is disabled.  |
| 1 | The value held in the ERR1PFGCDNR register is transferred into the Error Generation Counter. The Error Generation Counter counts down. |

**R, [30]**

Restartable bit. When it reaches 0, the Error Generation Counter restarts from the ERR1PFGCDNR value or stops. The possible values are:

- |   |  |
|---|--|
| 0 | When it reaches 0, the counter stops.  |
| 1 | When it reaches 0, the counter reloads the value that is stored in ERR1PFGCDNR and starts counting down again. |

**RES0, [29:7]**

- |      |           |
|------|-----------|
| RES0 | Reserved. |
|------|-----------|

**CE, [6]**

Corrected error generation enable. The possible values are:

- |   |  |
|---|--|
| 0 | No corrected error is generated.   |
| 1 | A corrected error might be generated when the Error Generation Counter is triggered. |

**DE, [5]**

Deferred Error generation enable. The possible values are:

- |   |   |
|---|---|
| 0 | No deferred error is generated.   |
| 1 | A deferred error might be generated when the Error Generation Counter is triggered. |

**RES0, [4:2]**

|             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**UC, [1]**

Uncontainable error generation enable. The possible values are:

|          |   |
|----------|---|
| <b>0</b> | No uncontainable error is generated.  |
| <b>1</b> | An uncontainable error might be generated when the Error Generation Counter is triggered. |

**RES0, [0]**

|             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**Configurations**

ERR1PFGCTLR resets to 0x00000000.

There are no configuration notes.

ERR1PFGCTLR is accessible from the following registers when ERRSELR.SEL==1:

- ERXPFGCTLR.
- ERXPFGCTLR\_EL1.


See the appropriate core Technical Reference Manual for information about these ERX\* registers.

## 12.9 ERR1PFGFR, Error Pseudo Fault Generation Feature Register

The ERR1PFGFR register is the DSU node register that defines which fault generation features are implemented.

**Bit field descriptions**

ERR1PFGFR is a 32-bit read-only register.

**Figure 12-6: ERR1PFGFR bit assignments****PFG, [31]**

Pseudo Fault Generation. The possible values are:

- |          |  |
|----------|--|
| <b>0</b> | The node does not support fault injection.       |
| <b>1</b> | The node implements a fault injection mechanism. |

**R, [30]**

Restartable bit. When it reaches zero, the Error Generation Counter restarts from the ERR0PFGCDN value or stops. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**RES0, [29:7]**

- |             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**CE, [6]**

Corrected Error generation. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**DE, [5]**

Deferred Error generation. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**UEO, [4]**

Latent or Restartable Error generation. The possible values are:

- |          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |



**UER, [3]**

Signaled or Recoverable Error generation. The possible values are:

|          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**UEU, [2]**

Unrecoverable Error generation. The possible values are:

|          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**UC, [1]**

Uncontainable Error generation. The possible values are:

|          |   |
|----------|---|
| <b>0</b> | The node does not support this feature. |
| <b>1</b> | This feature is controllable.           |

**RES0, [0]**

|             |           |
|-------------|-----------|
| <b>RES0</b> | Reserved. |
|-------------|-----------|

**Configurations**

ERR1PFGFR resets to 0xC0000062.

There are no configuration notes.

ERR1PFGFR is accessible from the following registers when ERRSELR.SEL==1:

- ERXPFGFR.
- ERXPFGFR\_EL1.

See the appropriate core Technical Reference Manual for information about these ERX\* registers.

## 12.10 ERR1STATUS, Error Record Primary Status Register

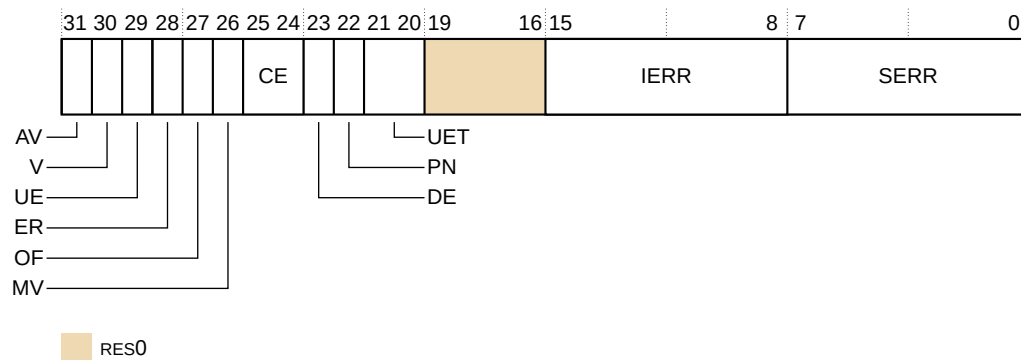
The ERR1STATUS register contains the following information about the error record:

- Whether any error has been detected.
- Whether any detected error was not corrected and returned to a master.
- Whether any detected error was not corrected and deferred.
- Whether a second error of the same type was detected before software handled the first error.
- Whether any error has been reported.
- Whether the other error record registers contain valid information.

## Bit field descriptions

ERR1STATUS is a 32-bit register.

**Figure 12-7: ERR1STATUS bit assignments**



### AV, [31]

Address Valid. The value is:

0 ERR1ADDR is not valid.

### V, [30]

Status Register valid. The possible values are:

0 ERR1STATUS is not valid.  
1 ERR1STATUS is valid. At least one error has been recorded.

### UE, [29]

Uncorrected error. The possible values are:

0 No error that could not be corrected or deferred has been detected.  
1 At least one error that could not be corrected or deferred has been detected. If error recovery interrupts are enabled, then the interrupt signal is asserted until this bit is cleared.

### ER, [28]

Error reported. The value is:

0 No external abort has been reported.

### OF, [27]

Overflow. The possible values are:

0

- If UE == 1, then no error status for an Uncorrected error has been discarded.
- If UE == 0 and DE == 1, then no error status for a Deferred error has been discarded.

- If UE == 0, DE == 0, and CE != 0b00, then:
  - If a Corrected error counter is implemented, it has not overflowed.
  - If no Corrected error counter is implemented, no error status for a Corrected error has been discarded.

**1** More than one error has occurred and so details of the other error have been discarded.

## MV, [26]

Miscellaneous Registers Valid. The possible values are:

**0** ERR1MISC0 and ERR1MISC1 are not valid.  
**1** Indicates that ERR1MISC0 contains additional information about any error recorded that is by this record.

## CE, [25:24]

Corrected error. The possible values are:

**0b00** No corrected errors are recorded.  
**0b10** At least one corrected error recorded.

## DE, [23]

Deferred error. The possible values are:

**0** No errors were deferred.  
**1** At least one error was not corrected and deferred by poisoning.

## PN, [22]

Poison. The possible values are:

**0** The DSU cannot distinguish a poisoned value from a corrupted value.  
**1** There is an uncorrected error due to data that was earlier poisoned.

## UET, [21:20]

Uncorrected Error Type. The value is:

**0b00** Uncontainable.

## RES0, [19:16]

**RES0** Reserved.

## IERR, [15:8]

An **IMPLEMENTATION DEFINED** error code. The possible values are:

**0x0** No error, or error on other RAMs.  
**0x2** Error on a L3 snoop filter RAM.

**SERR, [7:0]**

Primary error code. The possible values are:

|             |                                      |
|-------------|--------------------------------------|
| <b>0x0</b>  | No error.                            |
| <b>0x2</b>  | ECC error from internal data buffer. |
| <b>0x6</b>  | ECC error on cache data RAM.         |
| <b>0x7</b>  | ECC error on cache tag or dirty RAM. |
| <b>0x12</b> | Bus error.                           |

**Configurations**

ERR1STATUS resets to 0x00000000.

There are no configuration notes.

ERR1STATUS is accessible from the following registers when ERRSELR.SEL==1:

- ERXSTATUS.
- ERXSTATUS\_EL1.

See the appropriate core Technical Reference Manual for information about these ERX\* registers.

# 13. PMU registers

This chapter describes the PMU registers for the DSU.

## 13.1 About the PMU registers

The chapter is presented as follows:

### AArch32 PMU register summary

This section lists the AArch32 PMU registers by access encoding.

### AArch64 PMU register summary

This section lists the AArch64 PMU registers by access encoding.

### Register descriptions

The remainder of the chapter provides generic register descriptions, that apply to both AArch32 and AArch64 registers. They are listed in alphabetical order.

## 13.2 AArch32 PMU register summary

This section lists the AArch32 PMU registers implemented in the DSU, sorted by access encoding.

**Table 13-1: DynamIQ™ Shared Unit AArch32 32-bit PMU registers**

| Register mnemonic | Copro | CRn | Opc1 | CRm | Opc2 | Width   | Register name and description  |
|-------------------|-------|-----|------|-----|------|---|--|
| CLUSTERPMCR       | cp15  | c15 | 0    | c5  | 0    | 32  | See <a href="#">13.4 CLUSTERPMCR, Cluster Performance Monitors Control Register</a> on page 151    |
| CLUSTERPMCNTENSET | cp15  | c15 | 0    | c5  | 1    | 32  | See <a href="#">13.5 CLUSTERPMCNTENSET, Cluster Count Enable Set Register</a> on page 154.         |
| CLUSTERPMCNTENCLR | cp15  | c15 | 0    | c5  | 2    | 32  | See <a href="#">13.6 CLUSTERPMCNTENCLR, Cluster Count Enable Clear Register</a> on page 156.       |
| CLUSTERPMOVSSET   | cp15  | c15 | 0    | c5  | 3    | 32  | See <a href="#">13.7 CLUSTERPMOVSSET, Cluster Overflow Flag Status Set Register</a> on page 158.   |
| CLUSTERPMOVSCLR   | cp15  | c15 | 0    | c5  | 4    | 32  | See <a href="#">13.8 CLUSTERPMOVSCLR, Cluster Overflow Flag Status Clear Register</a> on page 160. |
| CLUSTERPMSELR     | cp15  | c15 | 0    | c5  | 5    | 32  | See <a href="#">13.9 CLUSTERPMSELR, Cluster Event Counter Selection Register</a> on page 162.      |
| CLUSTERPMINTENSET | cp15  | c15 | 0    | c5  | 6    | 32  | See <a href="#">13.10 CLUSTERPMINTENSET, Cluster Interrupt Enable Set Register</a> on page 164.    |
| CLUSTERPMINTENCLR | cp15  | c15 | 0    | c5  | 7    | 32  | See <a href="#">13.11 CLUSTERPMINTENCLR, Cluster Interrupt Enable Clear Register</a> on page 166.  |
| CLUSTERPMCCNTR    | cp15  | c15 | 0    | c6  | 0    | 64-bit register with only bits [31:0] accessible. | See <a href="#">13.12 CLUSTERPMCCNTR, Cluster Performance Monitors Cycle Counter</a> on page 168.  |

| Register mnemonic | Copro | CRn | Opc1 | CRm | Opc2 | Width | Register name and description   |
|-------------------|-------|-----|------|-----|------|-------|---|
| CLUSTERPMXEVTYPER | cp15  | c15 | 0    | c6  | 1    | 32    | See <a href="#">13.13 CLUSTERPMXEVTYPER</a> , Cluster Selected Event Type Register on page 169.       |
| CLUSTERPMXEVCNTR  | cp15  | c15 | 0    | c6  | 2    | 32    | See <a href="#">13.14 CLUSTERPMXEVCNTR</a> , Cluster Selected Event Counter Register on page 171.     |
| CLUSTERPMMDCR     | cp15  | c15 | 6    | c6  | 3    | 32    | See <a href="#">13.15 CLUSTERPMMDCR</a> , Cluster Monitor Debug Configuration Register on page 173    |
| CLUSTERPMCEID0    | cp15  | c15 | 0    | c6  | 4    | 32    | See <a href="#">13.16 CLUSTERPMCEID0</a> , Cluster Common Event Identification Register 0 on page 175 |
| CLUSTERPMCEID1    | cp15  | c15 | 0    | c6  | 5    | 32    | See <a href="#">13.17 CLUSTERPMCEID1</a> , Cluster Common Event Identification Register 1 on page 177 |
| CLUSTERPMCLAIMSET | cp15  | c15 | 0    | c6  | 6    | 32    | See <a href="#">13.18 CLUSTERCLAIMSET</a> , Cluster Claim Tag Set Register on page 179.               |
| CLUSTERPMCLAIMCLR | cp15  | c15 | 0    | c6  | 7    | 32    | See <a href="#">13.19 CLUSTERCLAIMCLR</a> , Cluster Claim Tag Clear Register on page 181.             |

The following table describes the DynamIQ™ Shared Unit AArch32 64-bit PMU register.

**Table 13-2: DynamIQ™ Shared Unit AArch32 64-bit PMU register**

| Register mnemonic | Copro | Opc1 | CRm | Width | Register name and description  |
|-------------------|-------|------|-----|-------|--|
| CLUSTERPMCCNTR    | cp15  | 15   | c15 | 64    | See <a href="#">13.12 CLUSTERPMCCNTR</a> , Cluster Performance Monitors Cycle Counter on page 168. |

## 13.3 AArch64 PMU register summary

This section lists the AArch64 PMU registers implemented in the DSU, sorted by access encoding.

**Table 13-3: DynamIQ™ Shared Unit AArch64 PMU registers**

| Register mnemonic     | Op0 | CRn | Op1 | CRm | Op2 | Width | Register name and description  |
|-----------------------|-----|-----|-----|-----|-----|-------|--|
| CLUSTERPMCR_EL1       | 3   | c15 | 0   | c5  | 0   | 32    | See <a href="#">13.4 CLUSTERPMCR</a> , Cluster Performance Monitors Control Register on page 151     |
| CLUSTERPMCNTENSET_EL1 | 3   | c15 | 0   | c5  | 1   | 32    | See <a href="#">13.5 CLUSTERPMCNTENSET</a> , Cluster Count Enable Set Register on page 154.          |
| CLUSTERPMCNTENCLR_EL1 | 3   | c15 | 0   | c5  | 2   | 32    | See <a href="#">13.6 CLUSTERPMCNTENCLR</a> , Cluster Count Enable Clear Register on page 156.        |
| CLUSTERPMOVSSET_EL1   | 3   | c15 | 0   | c5  | 3   | 32    | See <a href="#">13.7 CLUSTERPMOVSSET</a> , Cluster Overflow Flag Status Set Register on page 158.    |
| CLUSTERPMOVSCCLR_EL1  | 3   | c15 | 0   | c5  | 4   | 32    | See <a href="#">13.8 CLUSTERPMOVSCCLR</a> , Cluster Overflow Flag Status Clear Register on page 160. |
| CLUSTERPMSELR_EL1     | 3   | c15 | 0   | c5  | 5   | 32    | See <a href="#">13.9 CLUSTERPMSELR</a> , Cluster Event Counter Selection Register on page 162.       |
| CLUSTERPMINTENSET_EL1 | 3   | c15 | 0   | c5  | 6   | 32    | See <a href="#">13.10 CLUSTERPMINTENSET</a> , Cluster Interrupt Enable Set Register on page 164.     |
| CLUSTERPMINTENCLR_EL1 | 3   | c15 | 0   | c5  | 7   | 32    | See <a href="#">13.11 CLUSTERPMINTENCLR</a> , Cluster Interrupt Enable Clear Register on page 166.   |

| Register mnemonic     | Op0 | CRn | Op1 | CRm | Op2 | Width | Register name and description   |
|-----------------------|-----|-----|-----|-----|-----|-------|---|
| CLUSTERPMCCNTR_EL1    | 3   | c15 | 0   | c6  | 0   | 64    | See 13.12 CLUSTERPMCCNTR, Cluster Performance Monitors Cycle Counter on page 168.     |
| CLUSTERPMXEVTYPER_EL1 | 3   | c15 | 0   | c6  | 1   | 32    | See 13.13 CLUSTERPMXEVTYPER, Cluster Selected Event Type Register on page 169.        |
| CLUSTERPMXVCNTR_EL1   | 3   | c15 | 0   | c6  | 2   | 32    | See 13.14 CLUSTERPMXVCNTR, Cluster Selected Event Counter Register on page 171.       |
| CLUSTERPMMDCR_EL3     | 3   | c15 | 6   | c6  | 3   | 32    | See 13.15 CLUSTERPMMDCR, Cluster Monitor Debug Configuration Register on page 173     |
| CLUSTERPMCEID0_EL1    | 3   | c15 | 0   | c6  | 4   | 32    | See 13.16 CLUSTERPMCEID0, Cluster Common Event Identification Register 0 on page 175. |
| CLUSTERPMCEID1_EL1    | 3   | c15 | 0   | c6  | 5   | 32    | See 13.17 CLUSTERPMCEID1, Cluster Common Event Identification Register 1 on page 177. |
| CLUSTERPMCLAIMSET_EL1 | 3   | c15 | 0   | c6  | 6   | 32    | See 13.18 CLUSTERCLAIMSET, Cluster Claim Tag Set Register on page 179.                |
| CLUSTERPMCLAIMCLR_EL1 | 3   | c15 | 0   | c6  | 7   | 32    | See 13.19 CLUSTERCLAIMCLR, Cluster Claim Tag Clear Register on page 181.              |

## 13.4 CLUSTERPMCR, Cluster Performance Monitors Control Register


The CLUSTERPMCR register provides details of the Performance Monitors implementation, including the number of counters implemented. CLUSTERPMCR also provides configurations and controls to the counters.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMCR) and AArch64 (CLUSTERPMCR\_EL1) registers.

CLUSTERPMCR is a 32-bit register, and is part of the PMU registers.

**Figure 13-1: CLUSTERPMCR bit assignments**



RES1  
RES0

### IMP, [31:24]

Indicates the implementer code. This value is:

**0x41** ASCII character 'A' - implementer is Arm Limited.

This is a read-only field.

#### **IDCODE, [23:16]**

Identification code. This value is:

**0x41** DSU

This is a read-only field.

#### **N, [15:11]**

Identifies the number of event counters implemented. The value is:

**0b00110** Six event counters are implemented together with the cycle counter CLUSTERPMCCNTR.

#### **RES0, [10:7]**

**RES0** Reserved.

#### **RES1, [6]**

**RES1** Reserved.

#### **RES0, [5]**

**RES0** Reserved.

#### **X, [4]**

**RAZ/  
WI** Read-As-Zero/Writes ignored.

#### **RES0, [3]**

**RES0** Reserved.

#### **C, [2]**

Clock counter reset. This bit is WO. The effects of writing to this bit are:

**0** No action.  
**1** Reset CLUSTERPMCCNTR to zero.

This bit is always RAZ.

#### **P, [1]**

Event counter reset. This bit is WO. The effects of writing to this bit are:

**0** No action.  
**1** Reset all event counters, not including CLUSTERPMCCNTR, to zero.

Resetting the event counters does not clear any overflow bits to 0.



This bit is always RAZ.

## E, [0]

Counters enable. The possible values are:

- |   |   |
|---|---|
| 0 | All counters including CLUSTERPMCCNTR are disabled. |
| 1 | All counters are enabled by CLUSTERPMCNTENSET.      |

## Configurations

The AArch32 CLUSTERPMCR register is architecturally mapped to the AArch64 CLUSTERPMCR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMCR

#### In AArch64 state (CLUSTERPMCR\_EL1):

To read this register in AArch64 state (CLUSTERPMCR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_0; Read CLUSTERPMCR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_0, <Xt>; Write Xt into CLUSTERPMCR_EL1
```

#### In AArch32 state (CLUSTERPMCR):

To read this register in AArch32 state (CLUSTERPMCR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 0; Read CLUSTERPMCR into Rt
```

To write this register in AArch32 state (CLUSTERPMCR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 0; Write Rt into CLUSTERPMCR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg> | Control |     |    | Accessibility |     |     |     |
|-------------|---------|-----|----|---------------|-----|-----|-----|
|             | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMCR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMCR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMCR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

### 13.5 CLUSTERPMCNTENSET, Cluster Count Enable Set Register


The CLUSTERPMCNTENSET register enables the Cycle Count Register, CLUSTERPMCCNTR, and any implemented event counters CLUSTERPMEVCNTR<n>. Reading this register shows which counters are enabled.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMCNTENSET) and AArch64 (CLUSTERPMCNTENSET\_EL1) registers.

CLUSTERPMCNTENSET is a 32-bit register, and is part of the PMU registers.

### Figure 13-2: CLUSTERPMCNTENSET bit assignments



## C, [31]

CLUSTERPMCCNTR enable bit. Enables the cycle counter register. The possible values are:

- |   |  |
|---|--|
| 0 | When read, means that the cycle counter is disabled. When written, has no effect.            |
| 1 | When read, means that the cycle counter is enabled. When written, enables the cycle counter. |

P<n>, bit [n], for n = 0-30

Event counter enable bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- 0 When read, means that CLUSTERPMEVCNTR<n> is disabled. When written, has no effect.

- 1 When read, means that CLUSTERPMEVCNTR<n> event counter is enabled. When written, enables CLUSTERPMEVCNTR<n>.

## Configurations

The AArch32 CLUSTERPMCNTENSET register is architecturally mapped to the AArch64 CLUSTERPMCNTENSET\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMCNTENSET

#### In AArch64 state (CLUSTERPMCNTENSET\_EL1):

To read this register in AArch64 state (CLUSTERPMCNTENSET\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_1; Read CLUSTERPMCNTENSET_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCNTENSET\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_1, <Xt>; Write Xt into CLUSTERPMCNTENSET_EL1
```

#### In AArch32 state (CLUSTERPMCNTENSET):

To read this register in AArch32 state (CLUSTERPMCNTENSET) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 1; Read CLUSTERPMCNTENSET into Rt
```

To write this register in AArch32 state (CLUSTERPMCNTENSET) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 1; Write Rt into CLUSTERPMCNTENSET
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>       | Control |     |    | Accessibility |     |     |     |
|-------------------|---------|-----|----|---------------|-----|-----|-----|
|                   | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMCNTENSET | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMCNTENSET | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMCNTENSET | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.6 CLUSTERPMCNTENCLR, Cluster Count Enable Clear Register


The CLUSTERPMCNTENCLR register disables the Cycle Count Register, CLUSTERPMCCNTR, and any implemented event counters CLUSTERPMEVCNTR<n>. Reading this register shows which counters are enabled. CLUSTERPMCNTENCLR is used along with CLUSTERPMCNTENSET register.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMCNTENCLR) and AArch64 (CLUSTERPMCNTENCLR\_EL1) registers.

CLUSTERPMCNTENCLR is a 32-bit register, and is part of the PMU registers.

### Figure 13-3: CLUSTERPMCNTENCLR bit assignments



## C, [31]

CLUSTERPMCCNTR enable bit. Enables the cycle counter register. The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that the cycle counter is disabled. When written, has no effect.             |
| 1 | When read, means that the cycle counter is enabled. When written, disables the cycle counter. |

P<n>, bit [n], for n = 0-30

Event counter disable bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that CLUSTERPMEVCNTR<n> is disabled. When written, has no effect.                            |
| 1 | When read, means that CLUSTERPMEVCNTR<n> event counter is enabled. When written, disables CLUSTERPMEVCNTR<n>. |

## Configurations

The AArch32 CLUSTERPMCNTENCLR register is architecturally mapped to the AArch64 CLUSTERPMCNTENCLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMCNTENCLR

#### In AArch64 state (CLUSTERPMCNTENCLR\_EL1):

To read this register in AArch64 state (CLUSTERPMCNTENCLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_2; Read CLUSTERPMCNTENCLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCNTENCLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_2, <Xt>; Write Xt into CLUSTERPMCNTENCLR_EL1
```

#### In AArch32 state (CLUSTERPMCNTENCLR):

To read this register in AArch32 state (CLUSTERPMCNTENCLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 2; Read CLUSTERPMCNTENCLR into Rt
```

To write this register in AArch32 state (CLUSTERPMCNTENCLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 2; Write Rt into CLUSTERPMCNTENCLR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>       | Control |     |    | Accessibility |     |     |     |
|-------------------|---------|-----|----|---------------|-----|-----|-----|
|                   | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMCNTENCLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMCNTENCLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMCNTENCLR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.7 CLUSTERPMOVSSET, Cluster Overflow Flag Status Set Register


The CLUSTERPMOVSET register sets the state of the overflow bit for the Cycle Count Register, CLUSTERPMCCNTR, and each of the implemented event counters CLUSTERPMEVCNTR<n>.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMOVSSET) and AArch64 (CLUSTERPMOVSSET\_EL1) registers.

CLUSTERPMOVSSET is a 32-bit register, and is part of the PMU registers.

### Figure 13-4: CLUSTERPMOVSSET bit assignments



## C, [31]

CLUSTERPMCCNTR overflow bit. The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that the cycle counter has not overflowed. When written, has no effect.          |
| 1 | When read, means that the cycle counter has overflowed. When written, sets the overflow bit to 1. |

P<n>, bit [n], for n = 0-30

Event counter overflow set bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that CLUSTERPMEVCNTR<n> has not overflowed. When written, has no effect.                             |
| 1 | When read, means that CLUSTERPMEVCNTR<n> has overflowed. When written, sets the CLUSTERPMEVCNTR<n> overflow bit to 1. |

## Configurations

The AArch32 CLUSTERPMOVSSET register is architecturally mapped to the AArch64 CLUSTERPMOVSSET\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMOVSSET

#### In AArch64 state (CLUSTERPMOVSSET\_EL1):

To read this register in AArch64 state (CLUSTERPMOVSSET\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_3; Read CLUSTERPMOVSSET_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMOVSSET\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_3, <Xt>; Write Xt into CLUSTERPMOVSSET_EL1
```

#### In AArch32 state (CLUSTERPMOVSSET):

To read this register in AArch32 state (CLUSTERPMOVSSET) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 3; Read CLUSTERPMOVSSET into Rt
```

To write this register in AArch32 state (CLUSTERPMOVSSET) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 3; Write Rt into CLUSTERPMOVSSET
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>     | Control |     |    | Accessibility |     |     |     |
|-----------------|---------|-----|----|---------------|-----|-----|-----|
|                 | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMOVSSET | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMOVSSET | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMOVSSET | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.8 CLUSTERPMOVSLR, Cluster Overflow Flag Status Clear Register


The CLUSTERPMOVSLR register contains the state of the overflow bit for the Cycle Counter Register, CLUSTERPMCCNTR, and each of the implemented event counters CLUSTERPMEVCNTR<n>. Writing to this register clears these bits.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMOVSLR) and AArch64 (CLUSTERPMOVSLR\_EL1) registers.

CLUSTERPMOVSLR is a 32-bit register, and is part of the PMU registers.

### Figure 13-5: CLUSTERPMOVSLR bit assignments



## C, [31]

CLUSTERPMCCNTR overflow bit. The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that the cycle counter has not overflowed. When written, has no effect.            |
| 1 | When read, means that the cycle counter has overflowed. When written, clears the overflow bit to 0. |

CLUSTERPMCCR.LC controls whether an overflow is detected from CLUSTERPMCCNTR [31] or from CLUSTERPMCCNTR.

P<n>, bit [n], for n = 0-30

Event counter overflow clear bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that CLUSTERPMEVCNTR<n> has not overflowed. When written, has no effect.                               |
| 1 | When read, means that CLUSTERPMEVCNTR<n> has overflowed. When written, clears the CLUSTERPMEVCNTR<n> overflow bit to 0. |



## Configurations

The AArch32 CLUSTERPMOVSLR register is architecturally mapped to the AArch64 CLUSTERPMOVSLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMOVSLR

#### In AArch64 state (CLUSTERPMOVSLR\_EL1):

To read this register in AArch64 state (CLUSTERPMOVSLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_4; Read CLUSTERPMOVSLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMOVSLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_4, <Xt>; Write Xt into CLUSTERPMOVSLR_EL1
```

#### In AArch32 state (CLUSTERPMOVSLR):

To read this register in AArch32 state (CLUSTERPMOVSLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 4; Read CLUSTERPMOVSLR into Rt
```

To write this register in AArch32 state (CLUSTERPMOVSLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 4; Write Rt into CLUSTERPMOVSLR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>    | Control |     |    | Accessibility |     |     |     |
|----------------|---------|-----|----|---------------|-----|-----|-----|
|                | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMOVSLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMOVSLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMOVSLR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.9 CLUSTERPMSEL, Cluster Event Counter Selection Register

The CLUSTERPMSEL register selects the current event counter CLUSTERPMEVCNTR<n>. CLUSTERPMSEL is used along with CLUSTERPMXEVTYPER to determine the value of a selected event counter. It is also used along with CLUSTERPMXEVCNTR, to determine the value of a selected event counter.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMSELR) and AArch64 (CLUSTERPMSELR\_EL1) registers.

CLUSTERPMSELR is a 32-bit register, and is part of the PMU registers.

**Figure 13-6: CLUSTERPMSELR bit assignments**



RES0

## RES0, [31:5]

|      |           |
|------|-----------|
| RES0 | Reserved. |
|------|-----------|

## SEL, [4:0]

Selects event counter, CLUSTERPMEVCNTR<n>, where n is the value held in this field. This value identifies which event counter is accessed when a subsequent access to CLUSTERPMXEVTYPER or CLUSTERPMXVCNTR occurs.

This field can take any value from 0b00000 to (CLUSTERPMCR.N - 1).

If this field is set to a value greater than or equal to the number of implemented counters, the results of access to CLUSTERPMXEVTYPER or CLUSTERPMXEVCNTR are as if the register is RAZ/WI.

Direct reads of this field return an UNKNOWN value.

## Configurations

The AArch32 CLUSTERPMSELR register is architecturally mapped to the AArch64 CLUSTERPMSELR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMSELR

#### In AArch64 state (CLUSTERPMSELR\_EL1):

To read this register in AArch64 state (CLUSTERPMSELR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_5; Read CLUSTERPMSELR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMSELR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_5, <Xt>; Write Xt into CLUSTERPMSELR_EL1
```

#### In AArch32 state (CLUSTERPMSELR):

To read this register in AArch32 state (CLUSTERPMSELR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 5; Read CLUSTERPMSELR into Rt
```

To write this register in AArch32 state (CLUSTERPMSELR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 5; Write Rt into CLUSTERPMSELR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMSELR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMSELR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMSELR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.10 CLUSTERPMINTENSET, Cluster Interrupt Enable Set Register


The CLUSTERPMINTENSET register enables the generation of interrupt requests on overflows from the Cycle Count Register, CLUSTERPMC CNTR, and the event counters, CLUSTERPMEVCNTR<n>. Reading the register shows which overflow interrupt requests are enabled.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMINTENSET) and AArch64 (CLUSTERPMINTENSET\_EL1) registers.

CLUSTERPMINTENSET is a 32-bit register, and is part of the PMU registers.

### Figure 13-7: CLUSTERPMINTENSET bit assignments



## C, [31]

CLUSTERPMCCNTR overflow interrupt request enable bit. The possible values are:

- |   |  |
|---|--|
| 0 | When read, means that the cycle counter overflow interrupt request is disabled. When written, has no effect.                                     |
| 1 | When read, means that the cycle counter overflow interrupt request is enabled. When written, enables the cycle count overflow interrupt request. |

P<n>, bit [n], for n = 0-30

Event counter overflow interrupt request enable bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- 0** When read, means that CLUSTERPMEVCNTR<n> event counter interrupt request is disabled. When written, has no effect.

- 1 When read, means that CLUSTERPMEVCNTR<n> event counter interrupt request is enabled. When written, enables the CLUSTERPMEVCNTR<n> interrupt request.

## Configurations

The AArch32 CLUSTERPMINTENSET register is architecturally mapped to the AArch64 CLUSTERPMINTENSET\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMINTENSET

#### In AArch64 state (CLUSTERPMINTENSET\_EL1):

To read this register in AArch64 state (CLUSTERPMINTENSET\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_6; Read CLUSTERPMINTENSET_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMINTENSET\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_6, <Xt>; Write Xt into CLUSTERPMINTENSET_EL1
```

#### In AArch32 state (CLUSTERPMINTENSET):

To read this register in AArch32 state (CLUSTERPMINTENSET) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 6; Read CLUSTERPMINTENSET into Rt
```

To write this register in AArch32 state (CLUSTERPMINTENSET) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 6; Write Rt into CLUSTERPMINTENSET
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>       | Control |     |    | Accessibility |     |     |     |
|-------------------|---------|-----|----|---------------|-----|-----|-----|
|                   | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMINTENSET | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMINTENSET | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMINTENSET | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

### 13.11 CLUSTERPMINTENCLR, Cluster Interrupt Enable Clear Register


The CLUSTERPMINTENCLR register disables the generation of interrupt requests on overflows from the Cycle Count Register, CLUSTERPMCNCNTR, and the event counters, CLUSTERPMEVCNTR<n>. Reading the register shows which overflow interrupt requests are enabled.

## Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMINTENCLR) and AArch64 (CLUSTERPMINTENCLR\_EL1) registers.

CLUSTERPMINTENCLR is a 32-bit register, and is part of the PMU registers.

**Figure 13-8: CLUSTERPMINTENCLR bit assignments**



## C, [31]

CLUSTERPMCCNTR overflow interrupt request disable bit. The possible values are:

- |   |   |
|---|---|
| 0 | When read, means that the cycle counter overflow interrupt request is disabled. When written, has no effect.                                      |
| 1 | When read, means that the cycle counter overflow interrupt request is enabled. When written, disables the cycle count overflow interrupt request. |

P<n>, bit [n], for n = 0-30

Event counter overflow interrupt request disable bit for CLUSTERPMEVCNTR<n>.

Bits [30:N] are RAZ/WI. N is the value in CLUSTERPMCR.N.

The possible values are:

- 0** When read, means that CLUSTERPMEVCNTR<n> event counter interrupt request is disabled. When written, has no effect.

- 1 When read, means that CLUSTERPMEVCNTR<n> event counter interrupt request is enabled. When written, disables the CLUSTERPMEVCNTR<n> interrupt request.

## Configurations

The AArch32 CLUSTERPMINTENCLR register is architecturally mapped to the AArch64 CLUSTERPMINTENCLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERPMINTENCLR

#### In AArch64 state (CLUSTERPMINTENCLR\_EL1):

To read this register in AArch64 state (CLUSTERPMINTENCLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C5_7; Read CLUSTERPMINTENCLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMINTENCLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C5_7, <Xt>; Write Xt into CLUSTERPMINTENCLR_EL1
```

#### In AArch32 state (CLUSTERPMINTENCLR):

To read this register in AArch32 state (CLUSTERPMINTENCLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c5, 7; Read CLUSTERPMINTENCLR into Rt
```

To write this register in AArch32 state (CLUSTERPMINTENCLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c5, 7; Write Rt into CLUSTERPMINTENCLR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>       | Control |     |    | Accessibility |     |     |     |
|-------------------|---------|-----|----|---------------|-----|-----|-----|
|                   | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMINTENCLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMINTENCLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMINTENCLR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

**Traps and enables**

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.12 CLUSTERPMCCNTR, Cluster Performance Monitors Cycle Counter

The CLUSTERPMCCNTR register holds the value of the Cluster Cycle Counter, CLUSTERCCNT, that counts cluster clock cycles.

**Bit field descriptions**

This description applies to both the AArch32 (CLUSTERPMCCNTR) and AArch64 (CLUSTERPMCCNTR\_EL1) registers.

CLUSTERPMCCNTR is a 64-bit register, and is part of the PMU registers.

**Figure 13-9: CLUSTERPMCCNTR bit assignments**

**CLUSTERCCNT, [63:0]**

Cluster cycle count. This field increments every cluster clock cycle.

Writing 1 to CLUSTERPMCR.C sets this field to 0.

**Configurations**

The AArch32 CLUSTERPMCCNTR register is architecturally mapped to the AArch64 CLUSTERPMCCNTR\_EL1[63:0] register.

All counters are subject to any changes in clock frequency, including clock stopping caused by entering quiescent states. This means that it is CONSTRAINED UNPREDICTABLE whether or not CLUSTERPMCCNTR continues to increment when clocks are stopped.

This register is in the Warm reset domain. On a Warm or Cold reset RW fields in this register reset to architecturally UNKNOWN values.



## Usage constraints

### Accessing the CLUSTERPMCCNTR

#### In AArch64 state (CLUSTERPMCCNTR\_EL1):

To read this register in AArch64 state (CLUSTERPMCCNTR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_0; Read CLUSTERPMCCNTR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCCNTR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_0, <Xt>; Write Xt into CLUSTERPMCCNTR_EL1
```

#### In AArch32 state (CLUSTERPMCCNTR):

To read this register in AArch32 state (CLUSTERPMCCNTR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 0; Read CLUSTERPMCCNTR into Rt
```

To write this register in AArch32 state (CLUSTERPMCCNTR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 0; Write Rt into CLUSTERPMCCNTR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>    | Control |     |    | Accessibility |     |     |     |
|----------------|---------|-----|----|---------------|-----|-----|-----|
|                | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMCCNTR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMCCNTR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMCCNTR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

## 13.13 CLUSTERPMXEVTYPER, Cluster Selected Event Type Register


When CLUSTERPMSELR.SEL selects an event counter, CLUSTERPMXEVTYPER accesses a CLUSTERPMEVTYPER<n> register.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMXEVTYPER) and AArch64 (CLUSTERPMXEVTYPER\_EL1) registers.

CLUSTERPMXEVTYPER is a 32-bit register, and is part of the PMU registers.

**Figure 13-10: CLUSTERPMXEVTYPER bit assignments**



### Bits, [31:0]

Event type register.

This register accesses CLUSTERPMEVTYPER<n> where n is the value in CLUSTERPMSELR.SEL.

### Configurations

The AArch32 CLUSTERPMXEVTYPER register is architecturally mapped to the AArch64 CLUSTERPMXEVTYPER\_EL1 register.

This register is in the Warm reset domain. On a Warm or Cold reset RW fields in this register reset to architecturally UNKNOWN values.

### Usage constraints

#### Accessing the CLUSTERPMXEVTYPER

##### In AArch64 state (CLUSTERPMXEVTYPER\_EL1):

To read this register in AArch64 state (CLUSTERPMXEVTYPER\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_1; Read CLUSTERPMXEVTYPER_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMXEVTYPER\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_1, <Xt>; Write Xt into CLUSTERPMXEVTYPER_EL1
```

**In AArch32 state (CLUSTERPMXEVTYPYPER):**

To read this register in AArch32 state (CLUSTERPMXEVTYPYPER) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 1; Read CLUSTERPMXEVTYPYPER into Rt
```

To write this register in AArch32 state (CLUSTERPMXEVTYPYPER) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 1; Write Rt into CLUSTERPMXEVTYPYPER
```

**Accessibility**

This register is accessible in software as follows:

| <systemreg>         | Control |     |    | Accessibility |     |     |     |
|---------------------|---------|-----|----|---------------|-----|-----|-----|
|                     | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMXEVTYPYPER | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMXEVTYPYPER | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMXEVTYPYPER | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

**Traps and enables**

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

**Related information**

[CLUSTERPMXEVTYPYPER<n>, Cluster Event Type Register](#) on page 183


## 13.14 CLUSTERPMXEVCNTR, Cluster Selected Event Counter Register

The CLUSTERPMXEVCNTR register reads or writes the values of the selected event counter, CLUSTERPMXEVCNTR<n>. The register field CLUSTERPMSELR.SEL determines which event counter is selected.

**Bit field descriptions**

This description applies to both the AArch32 (CLUSTERPMXEVCNTR) and AArch64 (CLUSTERPMXEVCNTR\_EL1) registers.

CLUSTERPMXEVCNTR is a 32-bit register, and is part of the PMU registers.

**Figure 13-11: CLUSTERPMXEVCNTR bit assignments****CLUSTERPMXEVCNTR<n>, bits [31:0]**

Value of the selected event counter, CLUSTERPMXEVCNTR<n>, where n is the value stored in CLUSTERPMSELR.SEL.

**Configurations**

The AArch32 CLUSTERPMXEVCNTR register is architecturally mapped to the AArch64 CLUSTERPMXEVCNTR\_EL1 register.

This register is in the Warm reset domain. On a Warm or Cold reset RW fields in this register reset to architecturally UNKNOWN values.

**Usage constraints****Accessing the CLUSTERPMXEVCNTR****In AArch64 state (CLUSTERPMXEVCNTR\_EL1):**

To read this register in AArch64 state (CLUSTERPMXEVCNTR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_2; Read CLUSTERPMXEVCNTR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMXEVCNTR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_2, <Xt>; Write Xt into CLUSTERPMXEVCNTR_EL1
```

**In AArch32 state (CLUSTERPMXEVCNTR):**

To read this register in AArch32 state (CLUSTERPMXEVCNTR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 2; Read CLUSTERPMXEVCNTR into Rt
```

To write this register in AArch32 state (CLUSTERPMXEVCNTR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 2; Write Rt into CLUSTERPMXEVCNTR
```

**Accessibility**

This register is accessible in software as follows:

| <systemreg>     | Control |     |    | Accessibility |     |     |     |
|-----------------|---------|-----|----|---------------|-----|-----|-----|
|                 | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMXEVNTR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERPMXEVNTR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERPMXEVNTR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

### Related information

CLUSTERPMXEVNTR<n>, [Cluster Event Counter Register](#) on page 184

## 13.15 CLUSTERPMMDCR, Cluster Monitor Debug Configuration Register


The CLUSTERPMMDCR register defines which common architectural and common microarchitectural feature events are implemented.

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMMDCR) and AArch64 (CLUSTERPMMDCR\_EL3) registers.

CLUSTERPMMDCR is a 32-bit register, and is part of the PMU registers.

**Figure 13-12: CLUSTERPMMDCR bit assignments**



**RES0, [31:1]**

**RES0**

Reserved.

## SPME, [0]

Secure Performance Monitors enable. This allows the counting of Secure events. The possible values are:

- |   |                                       |
|---|---------------------------------------|
| 0 | Counting of secure events prohibited. |
| 1 | Counting of secure events allowed.    |

This bit resets to zero.

## Configurations

The AArch32 CLUSTERPMMDCR register is architecturally mapped to the AArch64 CLUSTERPMMDCR\_EL3 register.

## Usage constraints

### Accessing the CLUSTERPMMDCR

#### In AArch64 state (CLUSTERPMMDCR\_EL1):

To read this register in AArch64 state (CLUSTERPMMDCR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_6_C15_C6_3; Read CLUSTERPMMDCR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMMDCR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_6_C15_C6_3, <Xt>; Write Xt into CLUSTERPMMDCR_EL1
```

#### In AArch32 state (CLUSTERPMMDCR):

To read this register in AArch32 state (CLUSTERPMMDCR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 6, <Rt>, c15, c6, 3; Read CLUSTERPMMDCR into Rt
```

To write this register in AArch32 state (CLUSTERPMMDCR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 6, <Rt>, c15, c6, 3; Write Rt into CLUSTERPMMDCR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMMDCR | x       | x   | 0  | -             | -   | n/a | RW  |
| CLUSTERPMMDCR | x       | 0   | 1  | -             | -   | -   | RW  |

| <systemreg>   | Control |     |    | Accessibility |     |     |     |
|---------------|---------|-----|----|---------------|-----|-----|-----|
|               | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERPMMDCR | x       | 1   | 1  | -             | n/a | -   | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

This register is only accessible from EL3.


## 13.16 CLUSTERPMCEID0, Cluster Common Event Identification Register 0

The CLUSTERPMCEID0 register defines which common architectural and common microarchitectural feature events are implemented. This register is read-only write-ignores RO (WI).

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMCEID0) and AArch64 (CLUSTERPMCEID0\_EL1) registers.

**Figure 13-13: CLUSTERPMCEID0 bit assignments**



See the PMCEID0 register description in the accompanying core for more information.

### CE[31:0], [31:0]

Common architectural and microarchitectural feature events that can be counted by the PMU event counters.

The following table shows the CLUSTERPMCEID0 bit assignments with event implemented or not implemented when the associated bit is set to 1 or 0. See the [Arm® Architecture Reference Manual for A-profile architecture](#) for more information about these events.

**Table 13-16: PMU events**

| Bit  | Event number | Event mnemonic | Description  |
|------|--------------|----------------|--|
| [30] | 0x1E         | CHAIN          | Chain. For odd-numbered counters, counts when for each overflow of the preceding even-numbered counter. For even-numbered counters, does not count:<br><br><b>1</b> This event is implemented. |

| Bit  | Event number | Event mnemonic | Description  |
|------|--------------|----------------|--|
| [29] | 0x1D         | BUS_CYCLES     | Bus cycle:<br><b>1</b> This event is implemented.          |
| [26] | 0x1A         | MEMORY_ERROR   | Local memory error:<br><b>1</b> This event is implemented. |
| [25] | 0x19         | BUS_ACCESS     | Bus access:<br><b>1</b> This event is implemented.         |
| [17] | 0x11         | CPU_CYCLES     | Cycle:<br><b>1</b> This event is implemented.              |

### Configurations

The AArch32 CLUSTERPMCEID0 register is architecturally mapped to the AArch64 CLUSTERPMCEID0\_EL1 register.

### Usage constraints

#### Accessing the CLUSTERPMCEID0

##### In AArch64 state (CLUSTERPMCEID0\_EL1):

To read this register in AArch64 state (CLUSTERPMCEID0\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_4; Read CLUSTERPMCEID0_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCEID0\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_4, <Xt>; Write Xt into CLUSTERPMCEID0_EL1
```

##### In AArch32 state (CLUSTERPMCEID0):

To read this register in AArch32 state (CLUSTERPMCEID0) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 4; Read CLUSTERPMCEID0 into Rt
```

To write this register in AArch32 state (CLUSTERPMCEID0) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 4; Write Rt into CLUSTERPMCEID0
```

### Accessibility

This register is accessible in software as follows:



| <systemreg>    | Control |     |    | Accessibility |       |       |       |
|----------------|---------|-----|----|---------------|-------|-------|-------|
|                | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERPMCEID0 | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERPMCEID0 | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERPMCEID0 | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

### Traps and enables

There are no traps or enables affecting this register.


## 13.17 CLUSTERPMCEID1, Cluster Common Event Identification Register 1

The CLUSTERPMCEID1 register defines which common architectural and common microarchitectural feature events are implemented. This register is read-only write-ignores RO (WI).

### Bit field descriptions

This description applies to both the AArch32 (CLUSTERPMCEID1) and AArch64 (CLUSTERPMCEID1\_EL1) registers.

**Figure 13-14: CLUSTERPMCEID1 bit assignments**



See the PMCEID1 register description in the accompanying core for more information.

### RES0, [31:17]

RES0 Reserved.

### CE[48:32], [16:0]

Common architectural and microarchitectural feature events that can be counted by the PMU event counters.

For each bit described in the following table, the event is implemented if the bit is set to 1, or not implemented if the bit is set to 0.

**Table 13-18: PMU common event**

| Bit  | Event number | Event mnemonic     | Description  |
|------|--------------|--------------------|--|
| [12] | 0x2C         | L3D_CACHE_WB       | Attributable Level 3 unified cache writeback.<br><br><b>1</b> This event is implemented.                 |
| [11] | 0x2B         | L3D_CACHE          | Attributable Level 3 unified cache access.<br><br><b>1</b> This event is implemented.                    |
| [10] | 0x2A         | L3D_CACHE_REFILL   | Attributable Level 3 unified cache refill.<br><br><b>1</b> This event is implemented.                    |
| [9]  | 0x29         | L3D_CACHE_ALLOCATE | Attributable Level 3 unified cache allocation without refill.<br><br><b>1</b> This event is implemented. |

### Configurations

The AArch32 CLUSTERPMCEID1 register is architecturally mapped to the AArch64 CLUSTERPMCEID1\_EL1 register.

### Usage constraints

#### Accessing the CLUSTERPMCEID1

##### In AArch64 state (CLUSTERPMCEID1\_EL1):

To read this register in AArch64 state (CLUSTERPMCEID1\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_5; Read CLUSTERPMCEID1_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERPMCEID1\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_5, <Xt>; Write Xt into CLUSTERPMCEID1_EL1
```

##### In AArch32 state (CLUSTERPMCEID1):

To read this register in AArch32 state (CLUSTERPMCEID1) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 5; Read CLUSTERPMCEID1 into Rt
```

To write this register in AArch32 state (CLUSTERPMCEID1) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 5; Write Rt into CLUSTERPMCEID1
```

### Accessibility

This register is accessible in software as follows:

Copyright © 2016–2019, 2023 Arm Limited (or its affiliates). All rights reserved.  
Non-Confidential

| <systemreg>    | Control |     |    | Accessibility |       |       |       |
|----------------|---------|-----|----|---------------|-------|-------|-------|
|                | E2H     | TGE | NS | EL0           | EL1   | EL2   | EL3   |
| CLUSTERPMCEID1 | x       | x   | 0  | -             | RO/WI | n/a   | RO/WI |
| CLUSTERPMCEID1 | x       | 0   | 1  | -             | RO/WI | RO/WI | RO/WI |
| CLUSTERPMCEID1 | x       | 1   | 1  | -             | n/a   | RO/WI | RO/WI |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

There are no traps or enables affecting this register.

### 13.18 CLUSTERCLAIMSET, Cluster Claim Tag Set Register

The CLUSTERCLAIMSET register provides various bits that can be separately set to indicate whether functionality is in use by a debug agent. All debug agents must implement a common protocol to use these bits.


## Bit field descriptions

For examples on how these bits can be used, see the CLAIMSET register description in the [Arm® CoreSight™ Architecture Specification v3.0](#).

This description applies to both the AArch32 (CLUSTERCLAIMSET) and AArch64 (CLUSTERCLAIMSET\_EL1) registers.

CLUSTERCLAIMSET is a 32-bit register, and is part of the PMU registers.

### Figure 13-15: CLUSTERCLAIMSET bit assignments



## RAZ/WI, [31:4]

Read-As-Zero, Writes Ignored.

SET, bits [3:0]

Each bit in this field is a SET bit. Permitted values of SET[n] are:

|                          |                                   |
|--------------------------|-----------------------------------|
| <b>Write</b><br><b>0</b> | No effect.                        |
| <b>Write</b><br><b>1</b> | Set the claim tag bit for bit[n]. |

|             |   |
|-------------|---|
| <b>Read</b> | The claim tag that is represented by bit[n] is not implemented. |
| <b>0</b>    |   |
| <b>Read</b> | The claim tag that is represented by bit[n] is implemented.     |
| <b>1</b>    |   |

## Configurations

The AArch32 CLUSTERCLAIMSET register is architecturally mapped to the AArch64 CLUSTERCLAIMSET\_EL1 register.

## Usage constraints

### Accessing the CLUSTERCLAIMSET

#### In AArch64 state (CLUSTERCLAIMSET\_EL1):

To read this register in AArch64 state (CLUSTERCLAIMSET\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_6; Read CLUSTERCLAIMSET_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERCLAIMSET\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_6, <Xt>; Write Xt into CLUSTERCLAIMSET_EL1
```

#### In AArch32 state (CLUSTERCLAIMSET):

To read this register in AArch32 state (CLUSTERCLAIMSET) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 6; Read CLUSTERCLAIMSET into Rt
```

To write this register in AArch32 state (CLUSTERCLAIMSET) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 6; Write Rt into CLUSTERCLAIMSET
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>     | Control |     |    | Accessibility |     |     |     |
|-----------------|---------|-----|----|---------------|-----|-----|-----|
|                 | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERCLAIMSET | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERCLAIMSET | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERCLAIMSET | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

## Traps and enables

This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

### 13.19 CLUSTERCLAIMCLR, Cluster Claim Tag Clear Register

The CLUSTERCLAIMCLR register provides various bits that can be separately cleared to indicate whether functionality is in use by a debug agent. All debug agents must implement a common protocol to use these bits.


## Bit field descriptions

For examples on how these bits can be used, see [Arm® CoreSight™ Architecture Specification v3.0](#).

This description applies to both the AArch32 (CLUSTERCLAIMCLR) and AArch64 (CLUSTERCLAIMCLR\_EL1) registers.

CLUSTERCLAIMCLR is a 32-bit register, and is part of the PMU registers.

**Figure 13-16: CLUSTERCLAIMCLR bit assignments**



## RAZ/WI, [31:4]

Read-As-Zero, Writes Ignored.

## CLR, bits [3:0]

Each bit in this field is a SET bit. Permitted values of SET[n] are:

|                   |   |
|-------------------|---|
| <b>Write</b><br>0 | No effect.  |
| <b>Write</b><br>1 | Set the claim tag bit for bit[n].                               |
| <b>Read</b><br>0  | The claim tag that is represented by bit[n] is not implemented. |
| <b>Read</b><br>1  | The claim tag that is represented by bit[n] is implemented.     |

## Configurations

The AArch32 CLUSTERCLAIMCLR register is architecturally mapped to the AArch64 CLUSTERCLAIMCLR\_EL1 register.

## Usage constraints

### Accessing the CLUSTERCLAIMCLR

#### In AArch64 state (CLUSTERCLAIMCLR\_EL1):

To read this register in AArch64 state (CLUSTERCLAIMCLR\_EL1) into a general-purpose register, use the MRS instruction with the following syntax:

```
MRS <Xt>, S3_0_C15_C6_7; Read CLUSTERCLAIMCLR_EL1 into Xt
```

To write this register in AArch64 state (CLUSTERCLAIMCLR\_EL1) from a general-purpose register, use the MSR instruction with the following syntax:

```
MSR S3_0_C15_C6_7, <Xt>; Write Xt into CLUSTERCLAIMCLR_EL1
```

#### In AArch32 state (CLUSTERCLAIMCLR):

To read this register in AArch32 state (CLUSTERCLAIMCLR) into a general-purpose register, use the MRC (or MRC2) instruction with the following syntax:

```
MRC p15, 0, <Rt>, c15, c6, 7; Read CLUSTERCLAIMCLR into Rt
```

To write this register in AArch32 state (CLUSTERCLAIMCLR) from a general-purpose register, use the MCR (or MCR2) instruction with the following syntax:

```
MCR p15, 0, <Rt>, c15, c6, 7; Write Rt into CLUSTERCLAIMCLR
```

## Accessibility

This register is accessible in software as follows:

| <systemreg>     | Control |     |    | Accessibility |     |     |     |
|-----------------|---------|-----|----|---------------|-----|-----|-----|
|                 | E2H     | TGE | NS | EL0           | EL1 | EL2 | EL3 |
| CLUSTERCLAIMCLR | x       | x   | 0  | -             | RW  | n/a | RW  |
| CLUSTERCLAIMCLR | x       | 0   | 1  | -             | RW  | RW  | RW  |
| CLUSTERCLAIMCLR | x       | 1   | 1  | -             | n/a | RW  | RW  |

'n/a' Not accessible. The PE cannot be executing at this Exception level, so this access is not possible.

Traps and enables


This register is write accessible in EL1 if ACTLR\_EL3[12] is 1 and ACTLR\_EL2[12] is 1, or ACTLR\_EL3[12] is 1 and SCR.NS is 0. If write access is not possible, then Trap to the lowest Exception level that denied access (EL2 or EL3).

13.20 CLUSTERPMEVTYPER<n>, Cluster Event Type Register

The CLUSTERPMEVTYPER<n> register configures event counter n to count the specified event and its associated security filtering.

Bit field descriptions

Figure 13-17: CLUSTERPMEVTYPER<n> bit assignments



**S, [31]**  
Secure events filtering bit. Controls counting of events that are generated by Secure transactions. The possible values are:

- |   |                             |
|---|-----------------------------|
| 0 | Count Secure events.        |
| 1 | Do not count Secure events. |

**RES0, [30]**  
**RES0** Reserved.

**NS, [29]**  
Non-secure events filtering bit. Controls counting of events that are generated by Non-secure transactions. The possible values are:

- |                |  |
|----------------|--|
| <b>NS == S</b> | If the value of this bit equals the value of S,[31] bit then count Non-secure events.                |
| <b>NS != S</b> | If the value of this bit does not equal the value of S,[31] bit then do not count Non-secure events. |

**RES0, [28:16]**

|      |           |
|------|-----------|
| RES0 | Reserved. |
|------|-----------|

**evtCount, [15:0]**

Event number. For the list of implemented events, see [15.3 PMU events - DSU R3 and R3 ADP](#) on page 193.

If evtCount is programmed to an event that is not implemented, no events are counted.

**Usage constraints****Accessing the CLUSTERPMEVTYPER<n> register**


This register is accessible using the CLUSTERPMXEVTYPER register. It is not directly accessible as a system or CP15 register.

## 13.21 CLUSTERPMEVCNTR<n>, Cluster Event Counter Register

The CLUSTERPMEVCNTR<n> register holds the count for event counter n.

**Bit field descriptions**

**Figure 13-18: CLUSTERPMEVCNTR<n> bit assignments**

**Bits, [31:0]**

Event counter n. Value of event counter n, where n is the number of this register.

**Usage constraints****Accessing the CLUSTERPMEVCNTR<n> register**

This register is accessible using the CLUSTERPMXEVTYPER register. It is not directly accessible as a system or CP15 register.



# 14. Debug


This chapter describes the debug features of the DSU and the associated DebugBlock component.

## 14.1 About debug methods

The DSU along with its associated cores is part of a debug system that supports both self-hosted and external debug.

The following figure shows a typical external debug system.

**Figure 14-1: External debug system**



### Debug host

A computer, for example a personal computer, that is running a software debugger such as the DS-5 Debugger. With the debug host, you can issue high-level commands, such as setting a breakpoint at a certain location or examining the contents of a memory address.

### Protocol converter

The debug host sends messages to the debug target using an interface such as Ethernet. However, the debug target typically implements a different interface protocol. A device such as DSTREAM is required to convert between the two protocols.

### Debug target

The lowest level of the system implements system support for the protocol converter to access the debug unit using the *Advanced Peripheral Bus* (APB) slave interface. An example of a debug target is a development system with a test chip or a silicon part with a DSU.

### Debug unit

Helps debugging software that is running on the core:

- DSU and external hardware based around the core.
- Operating systems.
- Application software.

With the debug unit, you can:

- Stop program execution.
- Examine and alter process and coprocessor state.
- Examine and alter memory and the state of the input or output peripherals.

- Restart the PE.

For self-hosted debug, the debug target runs additional debug monitor software that runs on the core in the cluster. This way, it does not require expensive interface hardware to connect a second host computer.

## 14.2 Terminology

The debug system supports both single and multi-threaded cores.

The Arm architecture allows for cores to be single, or multi-threaded. A *Processing Element* (PE) performs a thread of execution. A single-threaded core has one PE and a multi-threaded core has two or more PEs. Because the debugging system allows individual threads to be debugged, the term PE is used throughout this chapter. Where a reference to a core is made, the core can be a single, or multi-threaded core.

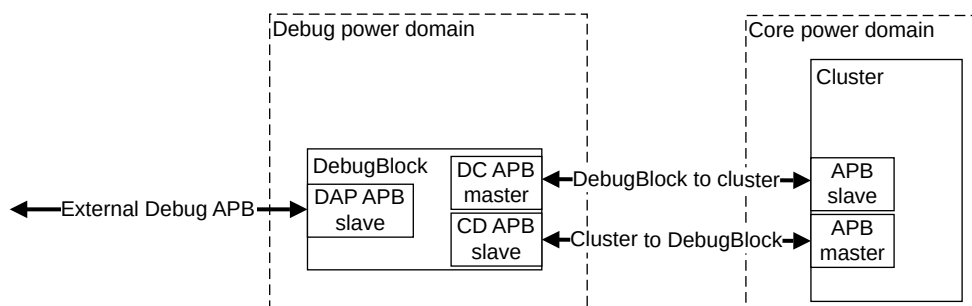
## 14.3 About the DebugBlock

The DebugBlock combines the functions, registers, and interfaces that are required for debug over powerdown.

The DebugBlock is provided as a separate component to allow implementation in a separate power domain from the cluster. Having a separate debug power domain allows the connection to a debugger be maintained while the cores and cluster are powered down.

The following diagram shows how the DebugBlock is connected to the cluster.

**Figure 14-2: Debug APB connections**



The DebugBlock has three APB interfaces:

### External Debug APB (DAP APB)

An APB slave interface, allowing communication with an external debugger, for example through a CoreSight *Debug Access Port* (DAP).

All debug register read and write requests from an external debugger are received on this bus.

#### **DebugBlock to cluster (DC APB)**

An APB master interface that is connected to the cluster. It sends all debug register read and write requests to the cluster.

CTI output trigger events are sent to the cluster as trigger requests on this bus.

#### **Cluster to DebugBlock (CD APB)**

An APB slave interface that is connected to the cluster. It receives CTI input trigger event requests from the cluster.

### **Debug register reads and writes**

The DebugBlock holds all the debug registers that are implemented in the Debug power domain. Registers implemented in the Debug power domain are specified in the [Arm® Architecture Reference Manual for A-profile architecture](#).

Accesses through the DAP APB interface to Debug domain registers are handled internally by the DebugBlock. Accesses through the DAP APB interface to Core power domain registers are passed on to the cluster through the DC APB interface.

### **CTI trigger events**

Trigger events are transferred between the DebugBlock and cluster through the CD APB and DC APB interfaces.

#### **Input trigger events**

Input trigger events are sent from the cluster to the CTIs through the CD APB as write transactions.

#### **Output trigger events**

Output trigger events are sent from the CTIs to the cluster through the DC APB as write transactions.

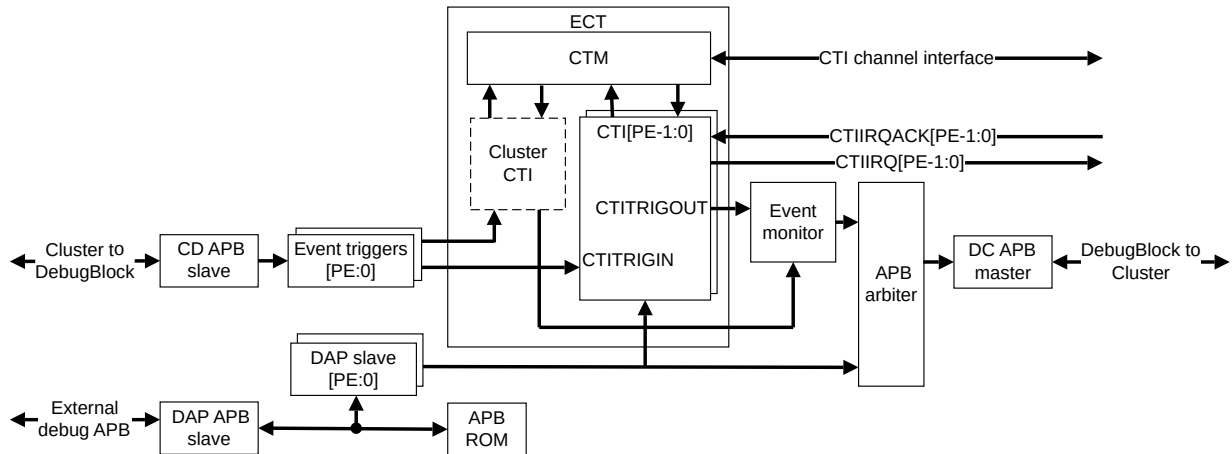
### **DebugBlock power states**

The DebugBlock supports two power modes: ON and OFF. These power modes are controlled using the power Q-Channel interface, which due to an erratum does not function correctly. Therefore, Arm recommends that the DebugBlock power Q-Channel is not used and that PWRQREQn is tied HIGH. To power down the DebugBlock, Arm recommends that the DebugBlock is put in reset before powering down. This causes any transactions, on the external Debug APB interface, that have not completed to complete with a SLVERR.

## 14.4 DebugBlock components

The components are:

**Figure 14-3: DebugBlock block diagram**



The CTIs shown in the diagram includes both the CTIs attached to each of the PEs [0:PE-1] and the cluster CTI. The cluster CTI is present only when the cluster ELA is present.

### ECT

The DebugBlock implements the *Embedded Cross Trigger* (ECT).

### APB ROM

The APB ROM table holds the address decoding for each debug component in the DebugBlock and the cluster. The APB ROM table complies with the [Arm® CoreSight™ Architecture Specification v3.0](#). Both v7 and v8 debug address maps are supported.

### Event monitor

The event monitor converts changes in CTI output triggers to APB write transactions.

### Event triggers

The event triggers convert APB write transactions to CTI input triggers.

### APB arbiter

The DC APB transfers both register accesses and CTI output trigger events. The APB arbiter multiplexes the two sources of transactions.

### DAP slave

The DAP slave holds copies of registers in the debug power domain.

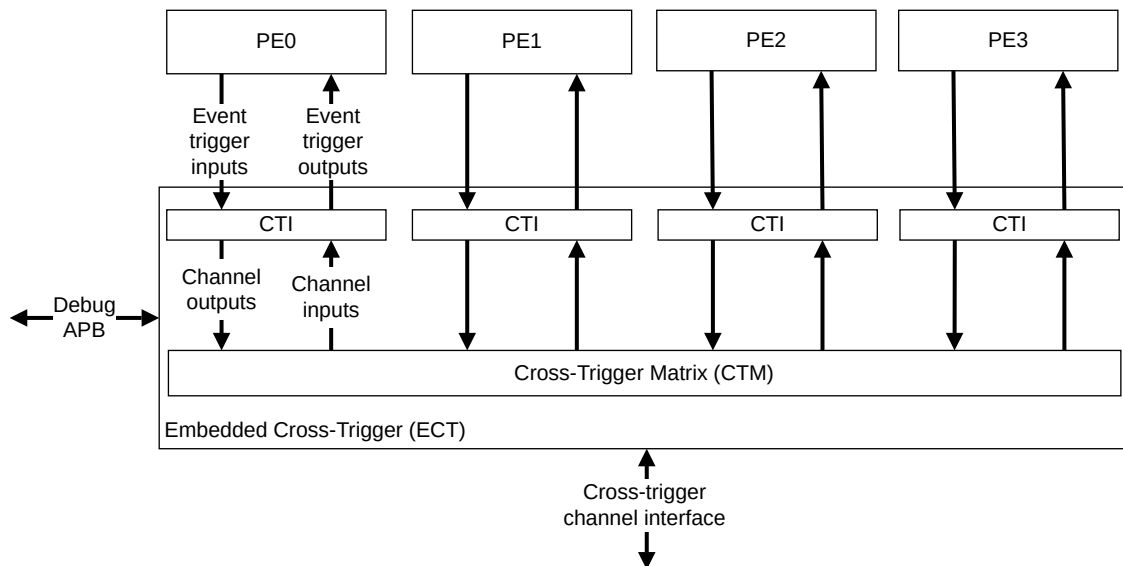
## 14.5 About the Embedded Cross Trigger

The *Embedded Cross Trigger* (ECT) allows debug events to be sent between PEs.

The ECT provides a *Cross Trigger Interface* (CTI) for each PE in the cluster. The CTIs are interconnected through a *Cross Trigger Matrix* (CTM) to send debug and trace events between PEs.

The following diagram shows a conceptual view of the trigger event inputs and outputs between the PEs and ECT.

**Figure 14-4: Embedded Cross Trigger concept**



The CTIs selectively send trigger events to the CTM on their respective channel outputs. The CTIs receive trigger events from the CTM on their channel inputs.

Trigger events are transferred between CTIs over the channel interface. The CTM connects the channel interface to the channel inputs and channel outputs of the CTIs.

### External interfaces

The external cross-trigger channel interface, from the CTM, allows cross-triggering between SoC external devices.

The Debug APB provides access to the CTI registers. This allows an external debugger to configure the trigger event routing, and send events to PEs, for example, to put a PE into Debug state.

### CTI registers

Registers in the CTI:

- Control the mapping of the input trigger events to channel outputs.

- Control the mapping of the channel inputs to output trigger events.
- Capture the state of input and output trigger events.
- Set, clear, or pulse output trigger events.

### 14.5.1 Supported debug and trace trigger events

The CTIs each have nine input and output trigger events that are mapped onto the debug and trace events in the PEs and ELAs.

The debug and trace trigger events from the CTI to the PE are:

#### **Debug request trigger event**

A trigger event sent from the CTI to the PE to force the PE into Debug state.

#### **Restart request trigger event**

A trigger event sent from the CTI to the PE to request the PE to exit Debug state.

#### **Generic CTI interrupt trigger event**

A trigger event sent from the CTI to the GIC.

#### **ETM trace input trigger events**

Four trigger events sent from the CTI to the ETM trace in the PE.

#### **ELA input trigger events**

Two trigger events sent from the CTI to the ELA attached to the PE.

The debug and trace events from the PE to the CTI are:

#### **Cross-halt trigger event**

A trigger event sent from the PE to the CTI when the PE enters Debug state.

#### **Performance Monitors overflow trigger event**

A trigger event sent from the PE to the CTI when a PMU event counter overflows.

#### **ETM trace output trigger events**

Four trigger events sent from the ETM in the PE to the CTI.

#### **ELA output trigger events**

Two trigger events sent from the ELA (attached to the PE) to the CTI.

#### **Profiling sample trigger event**

A trigger event sent from the PE to the CTI when a profiling sample is written out.

The cluster CTI has two input and output trigger events that are mapped onto the trigger events in the cluster ELA. The trigger events from the cluster CTI to the cluster ELA are:

#### **Cluster ELA input trigger events**

Two trigger events sent from the cluster CTI to the cluster ELA.

The trigger events from the cluster ELA to the cluster CTI are:

### Cluster ELA output trigger events

Two trigger events from the cluster ELA to the cluster CTI.

## 14.6 CTI triggers

Events are mapped onto CTI input and output triggers. All PEs in the cluster have the same mapping.

### PE CTI input trigger events

The following table shows how events are mapped onto PE CTI input triggers.

**Table 14-1: Allocation of PE CTI trigger inputs**

| Trigger number | Source | Destination | Type  | Event description                           |
|----------------|--------|-------------|-------|---|
| 0              | PE     | CTI         | Pulse | Cross-halt trigger event                    |
| 1              |        |             |       | Performance Monitors Overflow trigger event |
| 2              | PE     | CTI         | Pulse | Profiling sample trigger event              |
| 3              | -      | -           | -     | Reserved                                    |
| 4-7            | ETM    | CTI         | Pulse | ETM Trace Output trigger events             |
| 8-9            | ELA    | CTI         |       | ELA CTTRIGOUT[1:0] trigger events           |

### PE CTI output trigger events

The following table shows how events are mapped onto PE CTI output triggers.

**Table 14-2: Allocation of PE CTI trigger outputs**

| Trigger number | Source | Destination | Type  | Event description                           |
|----------------|--------|-------------|-------|---|
| 0              | CTI    | PE          | Level | Debug Request trigger event                 |
| 1              |        |             | Pulse | Restart Request trigger event               |
| 2              | CTI    | GIC         | Pulse | Generic CTI Interrupt trigger event         |
| 3              | -      | -           | -     | Reserved                                    |
| 4-7            | CTI    | ETM         | Pulse | Generic Trace External Input trigger events |
| 8-9            | CTI    | ELA         | Pulse | ELA CTTRIGIN[1:0] trigger events            |

### Allocation of cluster CTI trigger inputs

The following table shows how events are mapped onto the cluster CTI input triggers.

**Table 14-3: Allocation of cluster CTI trigger inputs**

| Trigger number | Source      | Destination | Type  | Event description          |
|----------------|-------------|-------------|-------|----------------------------|
| 0-7            | -           | -           | -     | Reserved                   |
| 8-9            | Cluster ELA | Cluster CTI | Pulse | Cluster ELA CTTRIGOUT[1:0] |

## Allocation of cluster CTI trigger outputs

The following table shows how events are mapped onto the cluster CTI output triggers.

**Table 14-4: Allocation of cluster CTI trigger inputs**

| Trigger number | Source      | Destination | Type  | Event description            |
|----------------|-------------|-------------|-------|------------------------------|
| 0-7            | -           | -           | -     | Reserved                     |
| 8-9            | Cluster CTI | Cluster ELA | Pulse | Cluster ELA<br>CTTRIGIN[1:0] |



# 15. Performance Monitoring Unit

This chapter describes the *Performance Monitoring Unit* (PMU).

## 15.1 About the PMU

The DynamIQ™ Shared Unit includes performance monitors that enable you to gather various statistics on the operation of the memory of the cluster during runtime. These provide useful information about the behavior of the cluster that you can use when debugging or profiling code.

The PMU provides six counters. Each counter can count any of the events available in the cluster. The absolute counts that are recorded might vary because of pipeline effects. This has negligible effect except in cases where the counters are enabled for a very short time.

## 15.2 PMU functional description

This section describes the functionality of the PMU.

The PMU includes the following interfaces and counters:

### Event interface

Events from all other units from across the design are provided to the PMU.

### System register

You can program the PMU registers using the system registers.

### Counters

The PMU has 32-bit counters that increment when they are enabled, based on events, and a 64-bit cycle counter.

### PMU register interfaces

The DynamIQ™ Shared Unit supports access to the performance monitor registers from the internal system register interface.

### Related information

[PMU registers](#) on page 149

## 15.3 PMU events

The following table shows the events that are generated and the numbers that the PMU uses to reference the events.

**Table 15-1: PMU events**

| PMU event number | Event mnemonic        | Event description   |
|------------------|-----------------------|---|
| 0x0011           | CYCLES                | Cycle counter.  |
| 0x0019           | BUS_ACCESS            | Bus access counter. Counts every beat of data that is transferred over the data channels between the SCU and the interconnect.<br><br>This event counts the sum of BUS_ACCESS_RD and BUS_ACCESS_WR. |
| 0x001A           | MEMORY_ERROR          | Local memory error counter. Counts every Correctable or Uncorrectable memory error (ECC or parity) in the protected RAMs.   |
| 0x001D           | BUS_CYCLES            | ACE or CHI bus cycle counter.   |
| 0x001E           | CHAIN                 | Odd performance counter chain mode.   |
| 0x0029           | L3D_CACHE_ALLOCATE    | Level 3 unified cache allocation without refill counter. Counts every full cache line write into the L3 cache which does not cause a linefill.  |
| 0x002A           | L3D_CACHE_REFILL      | Level 3 unified cache refill counter. Counts every cacheable read transaction issued to the interconnect.<br><br>This event counts the sum of L3D_CACHE_REFILL_RD and L3D_CACHE_REFILL_WR.          |
| 0x002B           | L3D_CACHE             | Level 3 unified cache access counter. Counts every cacheable read or write transaction issued to the SCU.<br><br>This event counts the sum of L3D_CACHE_RD and L3D_CACHE_WR.                        |
| 0x002C           | L3D_CACHE_WB          | Level 3 unified cache write-back counter. Counts every write-back from the L3 cache.  |
| 0x0060           | BUS_ACCESS_RD         | Bus access, read counter. Counts every beat of data transferred over the read data channel between the SCU and the interconnect.  |
| 0x0061           | BUS_ACCESS_WR         | Bus access, write counter. Counts every beat of data transferred over the write data channel between the SCU and the interconnect.  |
| 0x0062           | BUS_ACCESS_SHARED     | Bus access, shared counter. Counts every beat of shared data transferred over the data channels between the SCU and the interconnect.   |
| 0x0063           | BUS_ACCESS_NOT_SHARED | Bus access, not shared counter. Counts every beat of not shared data transferred over the write data channel between the SCU and the interconnect.  |
| 0x0064           | BUS_ACCESS_NORMAL     | Bus access, normal counter. Counts every beat of normal data transferred over the write data channel between the SCU and the interconnect.  |
| 0x0065           | BUS_ACCESS_PERIPH     | Bus access,Periph counter. Counts every beat of device data transferred over the write data channel between the SCU and the interconnect.   |
| 0x00A0           | L3D_CACHE_RD          | Level 3 unified cache access, read counter. Counts every cacheable read transaction that is issued to the SCU. Prefetches and stashes are not counted.  |
| 0x00A1           | L3D_CACHE_WR          | Level 3 unified cache access, write counter. Counts every cacheable write transaction issued to the SCU.  |
| 0x00A2           | L3D_CACHE_REFILL_RD   | Level 3 unified cache refill, read counter. Counts every cacheable read transaction issued to the interconnect caused by a read transaction. Prefetches and stashes are not counted.                |

| PMU event number | Event mnemonic       | Event description   |
|------------------|----------------------|---|
| 0x00A3           | L3D_CACHE_REFILL_WR  | Level 3 unified cache refill, write counter. Counts every cacheable read transaction issued to the interconnect caused by a write transaction.  |
| 0x0119           | ACP_ACCESS           | ACP access counter. Counts every beat of data transferred over the data channels between the SCU and the accelerated coherency port.<br><br>This event counts the sum of ACP_ACCESS_RD and ACP_ACCESS_WR.   |
| 0x011D           | ACP_CYCLES           | ACP cycle counter.  |
| 0x0160           | ACP_ACCESS_RD        | ACP access, read counter. Counts every beat of data transferred over the read data channel between the SCU and the peripheral port.   |
| 0x0161           | ACP_ACCESS_WR        | ACP access, write counter. Counts every beat of data transferred over the write data channel between the SCU and the peripheral port.   |
| 0x0219           | PP_ACCESS            | Peripheral port access counter. Counts every beat of data transferred over the data channels between the SCU and the peripheral port.<br><br>This event counts the sum of PP_ACCESS_RD and PP_ACCESS_WR.  |
| 0x021D           | PP_CYCLES            | Peripheral port cycle counter.  |
| 0x0260           | PP_ACCESS_RD         | Peripheral port access, read counter. Counts every beat of data transferred over the read data channel between the SCU and the peripheral port.   |
| 0x0261           | PP_ACCESS_WR         | Peripheral port access, write counter. Counts every beat of data transferred over the write data channel between the SCU and the peripheral port.   |
| 0x00C0           | SCU_SNP_ACCESS       | SNP access counter. Counts every external snoop request.  |
| 0x00C1           | SCU_SNP_EVICT        | SNP evictions counter. Counts every invalidating external snoop request that causes an L3 cache eviction.   |
| 0x00C2           | SCU_SNP_NO_CPU_SNP   | SNP, no CPU snoop counter. Counts every external snoop request that completes without needing to snoop a core.  |
| 0x0500           | SCU_PFTCH_CPU_ACCESS | Prefetch access, CPU counter. Counts every L3 prefetch transaction originating from a core.   |
| 0x0501           | SCU_PFTCH_CPU_MISS   | Prefetch data miss, CPU counter. Counts every L3 prefetch transaction originating from a core where data was read in from outside the cluster.  |
| 0x0502           | SCU_PFTCH_CPU_HIT    | Prefetch data hit, CPU counter. Counts every L3 prefetch transaction originating from a core where the L3 prefetch hit in the cluster.  |
| 0x0503           | SCU_PFTCH_CPU_MATCH  | Prefetch match, CPU counter. Counts every completed L3 prefetch transaction originating from a core that is matched by a compatible read request. This includes one caused by a L3 prefetch to the core, before the L3 prefetch times-out and is allocated into the L3 cache. |
| 0x0504           | SCU_PFTCH_CPU_KILL   | Prefetch terminate, CPU counter. Counts every killed L3 prefetch transaction originating from a core that is terminated due to an incompatible match.   |
| 0x0510           | SCU_STASH_ICN_ACCESS | Stash access, ICN counter. Counts every stash transaction originating from the interconnect.  |
| 0x0511           | SCU_STASH_ICN_MISS   | Stash data miss, ICN counter. Counts every stash transaction originating from the interconnect which utilizes a data pull, or is added to the stash queue and later issues a read.  |
| 0x0512           | SCU_STASH_ICN_HIT    | Stash data hit, ICN counter. Counts every non-invalidating stash transaction originating from the interconnect which hits in the cluster.   |
| 0x0513           | SCU_STASH_ICN_MATCH  | Stash match, ICN counter. Counts every completed stash transaction originating from the interconnect which is matched by a compatible read request. This includes one caused by a stash to the core, before the stash times out and is allocated into the L3 cache.           |
| 0x0514           | SCU_STASH_ICN_KILL   | Stash terminated, ICN counter. Counts every killed stash transaction originating from the interconnect that is terminated due to an incompatible match.   |
| 0x00D0           | SCU_HZD_ADDRESS      | Arbitration hazard, address counter. Counts every flush caused by an address hazard.  |

## 15.4 PMU interrupts

The DSU asserts the nCLUSTERPMUIRQ signal when the PMU generates an interrupt.

You can route this signal to an external interrupt controller for prioritization and masking. This is the only mechanism that signals this interrupt to a core.

# 16. Debug registers

This chapter describes the debug registers for the DSU.

## 16.1 Debug memory map

The debug memory map holds the base addresses for each debug component that is connected to the Debug APB.

Each component in the table requires 4KB, and uses the bottom 4KB of each 64KB region. The remaining 60KB of each region is reserved.

In the tables, individual *Processing Elements* (PE) are identified. For a single-threaded core, the PE number is the same as the core number. Only eight cores are supported.

The following table shows the address mapping for the DSU and the individual core debug APB components when configured for v8 Debug memory map. If an address range is not mapped to a component, it is indicated as reserved. For more information, see [Arm® Architecture Reference Manual for A-profile architecture](#).

**Table 16-1: Address mapping for APB components on 64KB pages**

| Address           | Component            |
|-------------------|----------------------|
| 0x000000          | ROM table            |
| 0x010000          | PE0 Debug            |
| 0x020000          | PE0 CTI              |
| 0x030000          | PE0 PMU              |
| 0x040000          | PE0 ETM              |
| 0x050000-0x0BFFFF | Reserved, RES0       |
| 0x0C0000          | PE0 ELA              |
| 0x0D0000          | Cluster ELA          |
| 0x0E0000          | Cluster CTI          |
| 0x0F0000          | PE0 Activity Monitor |
| 0x100000-0x10FFFF | Reserved, RES0       |
| 0x110000          | PE1 Debug            |
| 0x120000          | PE1 CTI              |
| 0x130000          | PE1 PMU              |
| 0x140000          | PE1 ETM              |
| 0x150000-0x1BFFFF | Reserved, RES0       |
| 0x1C0000          | PE1 ELA              |
| 0x1D0000-0x1EFFFF | Reserved, RES0       |
| 0x1F0000          | PE1 Activity Monitor |
| 0x200000-0x20FFFF | Reserved, RES0       |

| Address           | Component            |
|-------------------|----------------------|
| 0x210000          | PE2 Debug            |
| 0x220000          | PE2 CTI              |
| 0x230000          | PE2 PMU              |
| 0x240000          | PE2 ETM              |
| 0x250000-0x2BFFFF | Reserved, RES0       |
| 0x2C0000          | PE2 ELA              |
| 0x2D0000-0x2EFFFF | Reserved, RES0       |
| 0x2F0000          | PE2 Activity Monitor |
| 0x300000-0x30FFFF | Reserved, RES0       |
| 0x310000          | PE3 Debug            |
| 0x320000          | PE3 CTI              |
| 0x330000          | PE3 PMU              |
| 0x340000          | PE3 ETM              |
| 0x350000-0x3BFFFF | Reserved, RES0       |
| 0x3C0000          | PE3 ELA              |
| 0x3D0000-0x3EFFFF | Reserved, RES0       |
| 0x3F0000          | PE3 Activity Monitor |
| 0x400000-0x40FFFF | Reserved, RES0       |
| 0x410000          | PE4 Debug            |
| 0x420000          | PE4 CTI              |
| 0x430000          | PE4 PMU              |
| 0x440000          | PE4 ETM              |
| 0x450000-0x4BFFFF | Reserved, RES0       |
| 0x4C0000          | PE4 ELA              |
| 0x4D0000-0x4EFFFF | Reserved, RES0       |
| 0x4F0000          | PE4 Activity Monitor |
| 0x500000-0x50FFFF | Reserved, RES0       |
| 0x510000          | PE5 Debug            |
| 0x520000          | PE5 CTI              |
| 0x530000          | PE5 PMU              |
| 0x540000          | PE5 ETM              |
| 0x550000-0x5BFFFF | Reserved, RES0       |
| 0x5C0000          | PE5 ELA              |
| 0x5D0000-0x5EFFFF | Reserved, RES0       |
| 0x5F0000          | PE5 Activity Monitor |
| 0x600000-0x60FFFF | Reserved, RES0       |
| 0x610000          | PE6 Debug            |
| 0x620000          | PE6 CTI              |
| 0x630000          | PE6 PMU              |
| 0x640000          | PE6 ETM              |

| Address           | Component             |
|-------------------|-----------------------|
| 0x650000-0x6BFFFF | Reserved, RES0        |
| 0x6C0000          | PE6 ELA               |
| 0x6D0000-0x6EFFFF | Reserved, RES0        |
| 0x6F0000          | PE6 Activity Monitor  |
| 0x700000-0x70FFFF | Reserved, RES0        |
| 0x710000          | PE7 Debug             |
| 0x720000          | PE7 CTI               |
| 0x730000          | PE7 PMU               |
| 0x740000          | PE7 ETM               |
| 0x750000-0x7BFFFF | Reserved, RES0        |
| 0x7C0000          | PE7 ELA               |
| 0x7D0000-0x7EFFFF | Reserved, RES0        |
| 0x7F0000          | PE7 Activity Monitor  |
| 0x800000-0x80FFFF | Reserved, RES0        |
| 0x810000          | PE8 Debug             |
| 0x820000          | PE8 CTI               |
| 0x830000          | PE8 PMU               |
| 0x840000          | PE8 ETM               |
| 0x850000-0x8BFFFF | Reserved, RES0        |
| 0x8C0000          | PE8 ELA               |
| 0x8D0000-0x8EFFFF | Reserved, RES0        |
| 0x8F0000          | PE8 Activity Monitor  |
| 0x900000-0x90FFFF | Reserved, RES0        |
| 0x910000          | PE9 Debug             |
| 0x920000          | PE9 CTI               |
| 0x930000          | PE9 PMU               |
| 0x940000          | PE9 ETM               |
| 0x950000-0x9BFFFF | Reserved, RES0        |
| 0x9C0000          | PE9 ELA               |
| 0x9D0000-0x9EFFFF | Reserved, RES0        |
| 0x9F0000          | PE9 Activity Monitor  |
| 0xA00000-0xA0FFFF | Reserved, RES0        |
| 0xA10000          | PE10 Debug            |
| 0xA20000          | PE10 CTI              |
| 0xA30000          | PE10 PMU              |
| 0xA40000          | PE10 ETM              |
| 0xA50000-0xABFFFF | Reserved, RES0        |
| 0xAC0000          | PE10 ELA              |
| 0xAD0000-0xAEFFFF | Reserved, RES0        |
| 0xAF0000          | PE10 Activity Monitor |

| Address           | Component             |
|-------------------|-----------------------|
| 0xB00000-0xB0FFFF | Reserved, RES0        |
| 0xB10000          | PE11 Debug            |
| 0xB20000          | PE11 CTI              |
| 0xB30000          | PE11 PMU              |
| 0xB40000          | PE11 ETM              |
| 0xB50000-0xBBFFFF | Reserved, RES0        |
| 0xBC0000          | PE11 ELA              |
| 0xBD0000-0xBEFFFF | Reserved, RES0        |
| 0xBF0000          | PE11 Activity Monitor |
| 0xC00000-0xC0FFFF | Reserved, RES0        |
| 0xC10000          | PE12 Debug            |
| 0xC20000          | PE12 CTI              |
| 0xC30000          | PE12 PMU              |
| 0xC40000          | PE12 ETM              |
| 0xC50000-0xCBFFFF | Reserved, RES0        |
| 0xCC0000          | PE12 ELA              |
| 0xCD0000-0xCEFFFF | Reserved, RES0        |
| 0xCF0000          | PE12 Activity Monitor |
| 0xD00000-0xD0FFFF |                       |
| 0xD10000          | PE13 Debug            |
| 0xD20000          | PE13 CTI              |
| 0xD30000          | PE13 PMU              |
| 0xD40000          | PE13 ETM              |
| 0xD50000-0xDBFFFF | Reserved, RES0        |
| 0xDC0000          | PE13 ELA              |
| 0xDD0000-0xDEFFFF | Reserved, RES0        |
| 0xDF0000          | PE13 Activity Monitor |
| 0xE00000-0xE0FFFF | Reserved, RES0        |
| 0xE10000          | PE14 Debug            |
| 0xE20000          | PE14 CTI              |
| 0xE30000          | PE14 PMU              |
| 0xE40000          | PE14 ETM              |
| 0xE50000-0xEBFFFF | Reserved, RES0        |
| 0xEC0000          | PE14 ELA              |
| 0xED0000-0xEEFFFF | Reserved, RES0        |
| 0xEF0000          | PE14 Activity Monitor |
| 0xF00000-0xF0FFFF | Reserved, RES0        |
| 0xF10000          | PE15 Debug            |
| 0xF20000          | PE15 CTI              |
| 0xF30000          | PE15 PMU              |



| Address           | Component             |
|-------------------|-----------------------|
| 0xF40000          | PE15 ETM              |
| 0xF50000-0xFBFFFF | Reserved, RES0        |
| 0xFC0000          | PE15 ELA              |
| 0xFD0000-0xFEFFFF | Reserved, RES0        |
| 0xFF0000          | PE15 Activity Monitor |

The following table shows the address mapping for the DSU and the individual core debug APB components when configured for v7 Debug memory map. If an address range is not mapped to a component, it is indicated as reserved.

**Table 16-2: Address mapping for APB components on 4KB pages**

| Address            | Component            |
|--------------------|----------------------|
| 0x000000           | ROM table            |
| 0x001000           | Cluster ELA          |
| 0x002000           | Cluster CTI          |
| 0x003000 - 0x3FFFF | Reserved             |
| 0x004000           | PE0 Debug            |
| 0x005000           | PE1 Debug            |
| 0x006000           | PE2 Debug            |
| 0x007000           | PE3 Debug            |
| 0x008000           | PE0 ELA              |
| 0x009000           | PE1 ELA              |
| 0x00A000           | PE2 ELA              |
| 0x00B000           | PE3 ELA              |
| 0x00C000           | PE0 PMU              |
| 0x00D000           | PE1 PMU              |
| 0x00E000           | PE2 PMU              |
| 0x00F000           | PE3 PMU              |
| 0x010000-0x013FFF  | Reserved             |
| 0x014000           | PE0 CTI              |
| 0x015000           | PE1 CTI              |
| 0x016000           | PE2 CTI              |
| 0x017000           | PE3 CTI              |
| 0x018000           | PE0 Activity Monitor |
| 0x019000           | PE1 Activity Monitor |
| 0x01A000           | PE2 Activity Monitor |
| 0x01B000           | PE3 Activity Monitor |
| 0x01C000           | PE0 ETM              |
| 0x01D000           | PE1 ETM              |
| 0x01E000           | PE2 ETM              |
| 0x01F000           | PE3 ETM              |

| Address           | Component            |
|-------------------|----------------------|
| 0x020000-0x023FFF | Reserved             |
| 0x024000          | PE4 Debug            |
| 0x025000          | PE5 Debug            |
| 0x026000          | PE6 Debug            |
| 0x027000          | PE7 Debug            |
| 0x028000          | PE4 ELA              |
| 0x029000          | PE5 ELA              |
| 0x02A000          | PE6 ELA              |
| 0x02B000          | PE7 ELA              |
| 0x02C000          | PE4 PMU              |
| 0x02D000          | PE5 PMU              |
| 0x02E000          | PE6 PMU              |
| 0x02F000          | PE7 PMU              |
| 0x030000-0x033FFF | Reserved             |
| 0x034000          | PE4 CTI              |
| 0x035000          | PE5 CTI              |
| 0x036000          | PE6 CTI              |
| 0x037000          | PE7 CTI              |
| 0x038000          | PE4 Activity Monitor |
| 0x039000          | PE5 Activity Monitor |
| 0x03A000          | PE6 Activity Monitor |
| 0x03B000          | PE7 Activity Monitor |
| 0x03C000          | PE4 ETM              |
| 0x03D000          | PE5 ETM              |
| 0x03E000          | PE6 ETM              |
| 0x03F000          | PE7 ETM              |
| 0x040000-0x043FFF | Reserved             |
| 0x044000          | PE8 Debug            |
| 0x045000          | PE9 Debug            |
| 0x046000          | PE10 Debug           |
| 0x047000          | PE11 Debug           |
| 0x048000          | PE8 ELA              |
| 0x049000          | PE9 ELA              |
| 0x04A000          | PE10 ELA             |
| 0x04B000          | PE11 ELA             |
| 0x04C000          | PE8 PMU              |
| 0x04D000          | PE9 PMU              |
| 0x04E000          | PE10 PMU             |
| 0x04F000          | PE11 PMU             |
| 0x050000-0x053FFF | Reserved             |

| Address           | Component             |
|-------------------|-----------------------|
| 0x054000          | PE8 CTI               |
| 0x055000          | PE9 CTI               |
| 0x056000          | PE10 CTI              |
| 0x057000          | PE11 CTI              |
| 0x058000          | PE8 Activity Monitor  |
| 0x059000          | PE9 Activity Monitor  |
| 0x05A000          | PE10 Activity Monitor |
| 0x05B000          | PE11 Activity Monitor |
| 0x05C000          | PE8 ETM               |
| 0x05D000          | PE9 ETM               |
| 0x05E000          | PE10 ETM              |
| 0x05F000          | PE11 ETM              |
| 0x060000-0x063FFF | Reserved              |
| 0x064000          | PE12 Debug            |
| 0x065000          | PE13 Debug            |
| 0x066000          | PE14 Debug            |
| 0x067000          | PE15 Debug            |
| 0x068000          | PE12 ELA              |
| 0x069000          | PE13 ELA              |
| 0x06A000          | PE14 ELA              |
| 0x06B000          | PE15 ELA              |
| 0x06C000          | PE12 PMU              |
| 0x06D000          | PE13 PMU              |
| 0x06E000          | PE14 PMU              |
| 0x06F000          | PE15 PMU              |
| 0x070000-0x073FFF | Reserved              |
| 0x074000          | PE12 CTI              |
| 0x075000          | PE13 CTI              |
| 0x076000          | PE14 CTI              |
| 0x077000          | PE15 CTI              |
| 0x078000          | PE12 Activity Monitor |
| 0x079000          | PE13 Activity Monitor |
| 0x07A000          | PE14 Activity Monitor |
| 0x07B000          | PE15 Activity Monitor |
| 0x07C000          | PE12 ETM              |
| 0x07D000          | PE13 ETM              |
| 0x07E000          | PE14 ETM              |
| 0x07F000          | PE15 ETM              |

## 16.2 CTI register summary

This section describes the CTI registers in the *DynamIQ™ Shared Unit* (DSU). These registers are accessed through the external debug interface.

The following table gives a summary of the CTI registers. For registers that are not described in this chapter, see the [Arm® Architecture Reference Manual for A-profile architecture](#).

**Table 16-3: CTI register summary**

| Offset      | CTI              | Description  |
|-------------|------------------|--|
| 0x000       | CTICONTROL       | CTI Control register                                 |
| 0x004-0x00C | -                | Reserved   |
| 0x010       | CTIINTACK        | CTI Output Trigger Acknowledge Register              |
| 0x014       | CTIAPPSET        | CTI Application Trigger Set Register                 |
| 0x018       | CTIAPPCLEAR      | CTI Application Trigger Clear Register               |
| 0x01C       | CTIAPPULSE       | CTI Application Pulse Register                       |
| 0x020       | CTIINEN0         | CTI Input Trigger to Output Channel Enable Registers |
| 0x024       | CTIINEN1         |  |
| 0x028       | CTIINEN2         |  |
| 0x02C       | CTIINEN3         |  |
| 0x030       | CTIINEN4         |  |
| 0x034       | CTIINEN5         |  |
| 0x038       | CTIINEN6         |  |
| 0x03C       | CTIINEN7         |  |
| 0x040       | CTIINEN8         |  |
| 0x044       | CTIINEN9         |  |
| 0x048-0x09C | -                | Reserved   |
| 0x0A0       | CTIOUTEN0        | CTI Input Channel to Output Trigger Enable Registers |
| 0x0A4       | CTIOUTEN1        |  |
| 0x0A8       | CTIOUTEN2        |  |
| 0x0AC       | CTIOUTEN3        |  |
| 0x0B0       | CTIOUTEN4        |  |
| 0x0B4       | CTIOUTEN5        |  |
| 0x0B8       | CTIOUTEN6        |  |
| 0x0BC       | CTIOUTEN7        |  |
| 0x0C0       | CTIOUTEN8        |  |
| 0x0C4       | CTIOUTEN9        |  |
| 0x0C8-0x12C | -                | Reserved   |
| 0x130       | CTITRIGINSTATUS  | CTI Trigger In Status Register                       |
| 0x134       | CTITRIGOUTSTATUS | CTI Trigger Out Status Register                      |
| 0x138       | CTICHINSTATUS    | CTI Channel In Status Register                       |
| 0x13C       | CTICHOUTSTATUS   | CTI Channel Out Status Register                      |

| Offset      | CTI           | Description   |
|-------------|---------------|---|
| 0x140       | CTIGATE       | CTI Channel Gate Enable Register  |
| 0x144-0xF97 | -             | Reserved  |
| 0xFA0       | CTICLAIMSET   | CTI Claim Tag Set Register  |
| 0xFA4       | CTICLAIMCLR   | CTI Claim Tag Clear Register  |
| 0xFA8       | CTIDEVAFF0    | <a href="#">16.9 CTIDEVAFF0, Cluster CTI Device Affinity register 0</a> on page 212 <sup>10</sup> |
|             |               | CTI Device Affinity Register 0 <sup>11</sup>  |
| 0xFAC       | CTIDEVAFF1    | CTI Device Affinity Register 1  |
| 0xFB0-0xFB4 | -             | Reserved  |
| 0xFB8       | CTIAUTHSTATUS | CTI Authentication Status Register  |
| 0xFBC       | CTIDEVARCH    | CTI Device Architecture Register  |
| 0xFC0       | CTIDEVID2     | CTI Device Identification Register 2  |
| 0xFC4       | CTIDEVID1     | CTI Device Identification Register 1  |
| 0xFC8       | CTIDEVID      | <a href="#">16.10 CTIDEVID, CTI Device Identification Register</a> on page 213                    |
| 0xFCC       | CTIDEVTYPE    | CTI Device Type Register  |
| 0xFD0       | CTIPIDR4      | <a href="#">16.7 CTIPIDR4, CTI Peripheral Identification Register 4</a> on page 210               |
| 0xFD4-0xFDC | -             | Reserved  |
| 0xFE0       | CTIPIDR0      | <a href="#">16.3 CTIPIDR0, CTI Peripheral Identification Register 0</a> on page 205               |
| 0xFE4       | CTIPIDR1      | <a href="#">16.4 CTIPIDR1, CTI Peripheral Identification Register 1</a> on page 206               |
| 0xFE8       | CTIPIDR2      | <a href="#">16.5 CTIPIDR2, CTI Peripheral Identification Register 2</a> on page 208               |
| 0xFEC       | CTIPIDR3      | <a href="#">16.6 CTIPIDR3, CTI Peripheral Identification Register 3</a> on page 209               |
| 0xFF0       | CTICIDR0      | CTI Component Identification Register 0   |
| 0xFF4       | CTICIDR1      | CTI Component Identification Register 1   |
| 0xFF8       | CTICIDR2      | CTI Component Identification Register 2   |
| 0xFFC       | CTICIDR3      | CTI Component Identification Register 3   |

## 16.3 CTIPIDR0, CTI Peripheral Identification Register 0

The CTIPIDR0 register provides information to identify a CTI component.

### Bit field descriptions

CTIPIDR0 is a 32-bit register.

This register is Read Only.

<sup>10</sup> The CTIDEVAFF0 register for the cluster has a different format.

<sup>11</sup> The CTIDEVAFF0 register for the PEs is described in [Arm® Architecture Reference Manual for A-profile architecture](#).

### Figure 16-1: CTIPIDR0 bit assignments



RES0

## RES0, [31:8]

RES0 Reserved.

## Part\_0, [7:0]

Least significant byte of the part number.

For the cluster, the value is:

**0xE3** If v7 debug memory map is implemented.

**0xE4** If v8 debug memory map is implemented.

For the PEs, the value is a copy of bits [11:4] of the Main ID Register (MIDR) of the corresponding PE.

## Configurations

The CTIPIDR0 is in the Debug power domain and is optional to implement in the external register interface.

## Usage constraints

## Accessing the CTIPIDR0

The CTIPIDR0 can be accessed through the external debug interface with offset 0xFE0.

## Accessibility

The accessibility of the CTIPIDR0 by condition code is:

|         |
|---------|
| Default |
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

## 16.4 CTIPIDR1, CTI Peripheral Identification Register 1


The CTIPIDR1 register provides information to identify a CTI component.

## Bit field descriptions

CTIPIDR1 is a 32-bit register.

This register is Read Only.

### Figure 16-2: CTIPIDR1 bit assignments



RES0

RES0, [31:8]

RES0 Reserved.

DES\_0, [7:4]

This is the least significant nibble of JEP106 ID code.

**0xB** Arm Limited.

## Part\_1, [3:0]

Most significant nibble of the part number:

- For the cluster, the value is 0x4.
- For the PEs, the value is a copy of bits [15:12] of the MIDR of the corresponding PE.

## Configurations

The CTIPIDR1 is in the Debug power domain and is optional to implement in the external register interface.

## Usage constraints

## Accessing the CTIPIDR1

The CTIPIDR1 can be accessed through the external debug interface with offset 0xFE4.

## Accessibility

The accessibility of the CTIPIDR1 by condition code is:

|         |
|---------|
| Default |
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

## 16.5 CTIPIDR2, CTI Peripheral Identification Register 2


The CTIPIDR2 register provides information to identify a CTI component.

### Bit field descriptions

CTIPIDR2 is a 32-bit register.

This register is Read Only.

**Figure 16-3: CTIPIDR2 bit assignments**



 RES0

#### RES0, [31:8]

**RES0** Reserved.

#### Revision, [7:4]

**0x7** r4p1

#### JEDEC, [3]

**1** **RES1.** Indicates that a JEP106 identity code is used.

#### DES\_1, [2:0]

**0b011** Arm Limited. This is the most significant nibble of JEP106 ID code.

### Configurations

The CTIPIDR2 is in the Debug power domain and is optional to implement in the external register interface.

### Usage constraints

#### Accessing the CTIPIDR2

The CTIPIDR2 can be accessed through the external debug interface with offset **0xFE8**.



Accessibility

The accessibility of the CTIPIDR2 by condition code is:

| Default |
|---------|
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

16.6 CTIPIDR3, CTI Peripheral Identification Register 3

The CTIPIDR3 register provides information to identify a CTI component.

Bit field descriptions

CTIPIDR3 is a 32-bit register.

This register is Read Only.

Figure 16-4: CTIPIDR3 bit assignments



RES0

RES0, [31:8]

RES0      Reserved.

REVAND, [7:4]

0x0      Part minor revision.

CMOD, [3:0]

0x0      Customer modified.

Configurations

The CTIPIDR3 is in the Debug power domain and is optional to implement in the external register interface.

Usage constraints

Accessing the CTIPIDR3

The CTIPIDR3 can be accessed through the external debug interface with offset 0xFEC.

Accessibility

The accessibility of the CTIPIDR3 by condition code is:

| Default |
|---------|
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

16.7 CTIPIDR4, CTI Peripheral Identification Register 4

The CTIPIDR4 register provides information to identify a CTI component.

Bit field descriptions

CTIPIDR4 is a 32-bit register.

This register is Read Only.

Figure 16-5: CTIPIDR4 bit assignments



RES0

RES0, [31:8]

RES0      Reserved.

Size, [7:4]

0x0      Size of the component. Log<sub>2</sub> of the number of 4KB pages from the start of the component to the end of the component ID registers.

DES\_2, [3:0]

0x4      Arm Limited. This is the least significant nibble of the JEP106 continuation code.

Configurations

The CTIPIDR4 is in the Debug power domain and is optional to implement in the external register interface.

Usage constraints

Accessing the CTIPIDR4

The CTIPIDR4 can be accessed through the external debug interface with offset 0xFD0.

Accessibility

The accessibility of the CTIPIDR4 by condition code is:

| Default |
|---------|
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

16.8 CTIITCTRL, CTI Integration Mode Control Register


The CTIITCTRL register shows that the DSU does not implement an integration mode.

Bit field descriptions

CTIPIDR4 is a 32-bit register.

This register is Read Only.

Figure 16-6: CTIITCTRL bit assignments



RES0, [31:1]

RES0      Reserved.

IME, [0]

Integration mode enable. The value is:

0      Normal operation.

Configurations

The CTIITCTRL register is in the Debug power domain.

Usage constraints

Accessing the CTIITCTRL

The CTIITCTRL register can be accessed through external debug interface with offset 0x $\text{F00}$ .

Accessibility

The accessibility of the CTIITCTRL register by condition code is:

| Default |
|---------|
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

16.9 CTIDEVAFF0, Cluster CTI Device Affinity register 0


The CTIDEVAFF0 register identifies that the CTI relates to the DSU cluster.

Bit field descriptions

CTIDEVAFF0 is a 32-bit register.

This register is Read Only.

Figure 16-7: CTIDEVAFF0 bit assignments



RES0, [31:24]

RES0      Reserved.

Aff2, [23:16]

ClusterID Affinity Level-2 field.

The value is set by the CLUSTERIDAFF2[7:0] configuration input bus.

RES1, [15]

RES1      Reserved.

RES0, [14:0]

RES0                      Reserved.

Configurations

The CTIDEVAFF0 is in the Debug power domain and is optional to implement in the external register interface.

Usage constraints

Accessing the CTIDEVAFF0

The CTIDEVAFF0 can be accessed through the external debug interface with offset 0xFA8.

Accessibility

The accessibility of the CTIDEVAFF0 by condition code is:

|         |
|---------|
| Default |
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

16.10 CTIDEVID, CTI Device Identification Register


The CTIDEVID register describes the CTI component to the debugger.

Bit field descriptions

CTIDEVID is a 32-bit register.

This register is Read Only.

Figure 16-8: CTIDEVID bit assignments



RES0, [31:26]

RES0                      Reserved.

INOUT, [25:24]

Input and output options. Indicates the presence of an input gate. This value is:

**0b01** CTIGATE masks propagation of input events from external channels.

#### RES0, [23:22]

**RES0** Reserved.

#### NUMCHAN, [21:16]

Number of channels implemented. This value is:

**0b000100** Four channels implemented.

#### RES0, [15:14]

**RES0** Reserved.

#### NUMTRIG, [13:8]

Number of triggers implemented. This value is:

**0b001010** Ten triggers implemented.

#### RES0, [7:5]

**RES0** Reserved.

#### EXTMAXNUM, [4:0]

Maximum number of external triggers implemented. This value is:

**0b000000** No external triggers implemented.

### Configurations

The CTIDEVID register is in the Debug power domain.

### Usage constraints

#### Accessing the CTIDEVID

The CTIDEVID register can be accessed through the external debug interface with offset 0xFC8.

#### Accessibility

The accessibility of the CTIDEVID register by condition code is:

| Default |
|---------|
| RO      |

See [16.11 External register access permissions](#) on page 214 for the condition codes.

## 16.11 External register access permissions

External access permission to the CTI registers is subject to the conditions at the time of the access. The following table describes the response of the core to accesses through the external debug and memory-mapped interfaces.

**Table 16-12: External register conditions**

| Name    | Condition  | Description  |
|---------|--|--|
| Off     | EDPRSR.PU is 0                                   | Core power domain is completely off, or in a low-power state where the core power domain registers cannot be accessed.   |
| DLK     | EDPRSR.DLK is 1                                  | OS Double Lock is locked.  |
| OSLK    | OSLSR_EL1.OSLK is 1                              | OS Lock is locked.   |
| EDAD    | <code>AllowExternalDebugAccess() == FALSE</code> | External debug access is disabled. When an error is returned because of an EDAD condition code, and this is the highest priority error condition, EDPRSR.SDAD is set to 1. Otherwise EDPRSR.SDAD is unchanged. |
| Default | -  | None of the conditions apply, normal access.   |

The following table shows an example of external register condition codes for access to a CTI register. To determine the access permission for the register, scan the columns from left to right. Stop at the first column a condition is true, the entry gives the access permission of the register and scanning stops.

**Table 16-13: External register condition code example**

| Off | DLK | OSLK | EDAD | Default |
|-----|-----|------|------|---------|
| -   | -   | -    | -    | RO      |

# 17. ROM table

This chapter describes the CoreSight™ ROM Table component.

## 17.1 About the ROM table

The ROM table holds the locations of debug components.

The ROM table complies with the [Arm® CoreSight™ Architecture Specification v3.0](#). This table contains a list of components such as PE debug units, *Cross Trigger Interfaces* (CTIs), *PE Performance Monitoring Units* (PMUs), and *Embedded Trace Macrocells* (ETMs). Debuggers can use the ROM table to determine which components are implemented.

If a component is not included in your implementation, the corresponding ROM table entry indicates that the component is not present.

## 17.2 ROM table register summary

The specific implementation of your macrocell determines the contents of the ROM table.

The following table shows the offsets from the physical base address of the ROM table. The register formats are described in the [Arm® CoreSight™ Architecture Specification v3.0](#).

**Table 17-1: ROM table registers**

| Offset      | Name         | Reset value  | Description                   |
|-------------|--------------|--|-------------------------------|
| 0x000-0x144 | ROMENTRY0-81 | _12  | ROM entries                   |
| 0x148-0x9FC | -            | 0x00000000   | Reserved                      |
| 0xA00       | DBGPCR0      | <ul style="list-style-type: none"> <li>0x00000000 if Cluster ELA is not implemented.</li> <li>0x00000001 if Cluster ELA is implemented.</li> </ul> | Debug Power Control Registers |
| 0xA04       | DBGPCR1      | <ul style="list-style-type: none"> <li>0x00000000 if Core0 ELA is not implemented.</li> <li>0x00000001 if Core0 ELA is implemented.</li> </ul>     |                               |
| 0xA08       | DBGPCR2      | <ul style="list-style-type: none"> <li>0x00000000 if Core1 ELA is not implemented.</li> <li>0x00000001 if Core1 ELA is implemented.</li> </ul>     |                               |



| Offset      | Name    | Reset value  | Description                           |
|-------------|---------|--|---------------------------------------|
| 0xA0C       | DBGPCR3 | <ul style="list-style-type: none"> <li>0x00000000 if Core2 ELA is not implemented.</li> <li>0x00000001 if Core2 ELA is implemented.</li> </ul> |                                       |
| 0xA10       | DBGPCR4 | <ul style="list-style-type: none"> <li>0x00000000 if Core3 ELA is not implemented.</li> <li>0x00000001 if Core3 ELA is implemented.</li> </ul> |                                       |
| 0xA14       | DBGPCR5 | <ul style="list-style-type: none"> <li>0x00000000 if Core4 ELA is not implemented.</li> <li>0x00000001 if Core4 ELA is implemented.</li> </ul> |                                       |
| 0xA18       | DBGPCR6 | <ul style="list-style-type: none"> <li>0x00000000 if Core5 ELA is not implemented.</li> <li>0x00000001 if Core5 ELA is implemented.</li> </ul> |                                       |
| 0xA1C       | DBGPCR7 | <ul style="list-style-type: none"> <li>0x00000000 if Core6 ELA is not implemented.</li> <li>0x00000001 if Core6 ELA is implemented.</li> </ul> |                                       |
| 0xA20       | DBGPCR8 | <ul style="list-style-type: none"> <li>0x00000000 if Core7 ELA is not implemented.</li> <li>0x00000001 if Core7 ELA is implemented.</li> </ul> |                                       |
| 0xA24-0xA7C | -       | 0x00000000   | Reserved                              |
| 0xA80       | DBGPSR0 | 0x00000000   | Debug Power Status Registers          |
| 0xA84       | DBGPSR1 | 0x00000000   |                                       |
| 0xA88       | DBGPSR2 | 0x00000000   |                                       |
| 0xA8C       | DBGPSR3 | 0x00000000   |                                       |
| 0xA90       | DBGPSR4 | 0x00000000   |                                       |
| 0xA94       | DBGPSR5 | 0x00000000   |                                       |
| 0xA98       | DBGPSR6 | 0x00000000   |                                       |
| 0xA9C       | DBGPSR7 | 0x00000000   |                                       |
| 0xAA0       | DBGPSR8 | 0x00000000   |                                       |
| 0xAA4-0xBF0 | -       | 0x00000000   | Reserved                              |
| 0xC00       | PRIDR0  | 0x00000001   | Power Reset Identification Register 0 |

| Offset      | Name       | Reset value              | Description                          |
|-------------|------------|--------------------------|--------------------------------------|
| 0xC04-0xFB4 | -          | 0x00000000               | Reserved                             |
| 0xFB8       | AUTHSTATUS | 0x00000008               | Authentication Status Register       |
| 0xFBC       | DEVARCH    | 0x47700AF7               | Device Architecture Register         |
| 0xFC0-0xFC4 | -          | 0x00000000               | Reserved                             |
| 0xFC8       | DEVID      | 0x00000020               | Device ID Register                   |
| 0xFCC       | DEVTYPE    | 0x00000000               | Device Type Register                 |
| 0xFD0       | PIDR4      | 0x00000004               | Peripheral Identification Register 4 |
| 0xFD4-0xFDC | -          | 0x00000000               | Reserved                             |
| 0xFE0       | PIDR0      | 0x000000E3 <sup>13</sup> | Peripheral Identification Register 0 |
|             |            | 0x000000E4 <sup>14</sup> |                                      |
| 0xFE4       | PIDR1      | 0x000000B4               | Peripheral Identification Register 1 |
| 0xFE8       | PIDR2      | 0x0000007B               | Peripheral Identification Register 2 |
| 0xFEC       | PIDR3      | 0x00000000               | Peripheral Identification Register 3 |
| 0xFF0       | CIDR0      | 0x0000000D               | Component Identification Register 0  |
| 0xFF4       | CIDR1      | 0x00000090               | Component Identification Register 1  |
| 0xFF8       | CIDR2      | 0x00000005               | Component Identification Register 2  |
| 0xFFC       | CIDR3      | 0x000000B1               | Component Identification Register 3  |

<sup>12</sup> The ROMENTRY entry values depend on the number and type of cores implemented.

<sup>13</sup> If v7 debug memory map is implemented.

<sup>14</sup> If v8 debug memory map is implemented.

# A Appendixes

This section contains the appendixes for the signal descriptions and compatible core information for the *DynamIQ™ Shared Unit* (DSU).

# Appendix A Compatible Core Versions

This appendix provides the location of where to obtain information about the permissible combinations of cores.

## A.1 Compatible Core Versions

For information on the number and type of permissible cores in the DSU, see the appendix *Compatible core versions* in the *Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide*.



The *Arm® DynamIQ™ Shared Unit Configuration and Sign-off Guide* is a confidential document only available to licensees.

---

# Appendix B Signal descriptions

This appendix describes the DSU signals.

## B.1 Signal naming convention

Signals and buses are named using the following convention:

- Some signals or buses are per-core or per-thread. For single-threaded cores, there is a one-to-one mapping between cores and threads.
- Signals specified with a width of [CN:0] (<signal>[CN:0]) have 1 bit per core, where CN is the number of cores minus 1 (for example, CN =3 for a quad-core system).
- Signals named <signal>x have an instance of the signal per core, where x takes values from 0 to the number of cores minus 1.
- Signals specified with a width of [PE:0] (<signal>[PE:0]) have 1 bit per thread, where PE is the total number of threads minus one (for example PE=5 for a cluster with two single-threaded cores and two dual-threaded cores).
- Signals named <signal>y have an instance of the signal per thread, where y takes values from 0 to the total number of threads minus 1.

### Signal widths

The DSU supports cores that might have different PA widths. The PA width determines how much physical memory the core can access and is a fixed value for each type of core.

There are signals with variable widths that depend on the PA size. In this case, the width is given either as [p:0] or [q:0]. If there are different cores in the cluster with different PA widths, then the PA width of the cluster matches the largest PA width of the configured cores.

The following table shows the PA widths that the DSU supports and the corresponding p and q values.

**Table A-1: p and q values for signal widths**

| PA width | p  | q  |
|----------|----|----|
| 40       | 39 | 43 |
| 44       | 43 | 43 |
| 48       | 47 | 47 |
| 52       | 51 | 51 |



Note

The cluster configuration script outputs the maximum core PA width when the `--verbose` option is used.

In some cases, there are signals with variable widths that depend on the configuration. In this case, the width is given as [i:0], [j:0], [k:0], or [d:0]. The following tables show the respective values for various configurations.

**Table A-2: Signal width i, j, k values for CHI.B, CHI.C, and CHI.D**

| CHI   | i  | j   | k  |
|-------|----|-----|----|
| CHI.B | 58 | 85  | 48 |
| CHI.C | 58 | 85  | 48 |
| CHI.D | 65 | 100 | 63 |

**Table A-3: CHI signal widths d values**

| Interface | Cache Protection | d for CHI.B | d for CHI.C | d for CHI.D |
|-----------|------------------|-------------|-------------|-------------|
| 128-bit   | N                | 212         | 213         | -           |
|           | Y                | 214         | 215         | -           |
| 256-bit   | N                | 356         | 357         | 365         |
|           | Y                | 360         | 361         | 369         |

## B.2 DSU signals

This section shows the *DynamIQ™ Shared Unit* (DSU) signals.

### B.2.1 Clock and clock enable signals

This section describes the clock and clock enable signals.

**Table A-4: Clock signals**

| Signal        | Direction | Description   |
|---------------|-----------|---|
| CORECLK[CN:0] | Input     | Clock for each core.<br><b>Note:</b><br>If a core is synchronous to the cluster, then the corresponding bit of this signal is not present and the core uses SCLK instead. If all cores are synchronous to the cluster, then this signal is not present. |
| SCLK          | Input     | Clock for the SCU/L3 and the AMBA interface.  |
| PCLK          | Input     | Clock for the debug APB interface.  |
| ATCLK         | Input     | Clock for the ATB trace interface.  |
| GICCLK        | Input     | Clock for the GIC interface.  |
| PERIPHCLK     | Input     | Clock for the timers, power management, and other miscellaneous logic.  |

**Table A-5: Clock enable signals**

| Signal  | Direction | Description  |
|---------|-----------|--|
| ACLKENM | Input     | ACE Master bus clock enable. This pin is only present when the DSU is configured with the ACE interface. |

| Signal   | Direction | Description  |
|----------|-----------|--|
| ACLKENMP | Input     | AXI Master peripheral port clock enable. This pin is only present when the DSU is configured with the peripheral port interface. |
| ACLKENS  | Input     | AXI Slave bus clock enable. This pin is only present when the DSU is configured with the ACP interface.                          |
| CNTCLKEN | Input     | Counter clock enable.  |
| SCLKENM  | Input     | CHI Master bus clock enable. This pin is only present when the DSU is configured with the CHI interface.                         |
| TSCLKEN  | Input     | Timestamp clock enable.  |

## B.2.2 Reset signals

This section describes the reset signals.

**Table A-6: Reset signals**

| Signal            | Direction | Description  |
|-------------------|-----------|--|
| nATRESET          | Input     | Active-LOW reset for all registers in the ATCLK domain.                              |
| nCORERESET[CN:0]  | Input     | Active-LOW core Warm reset, excluding Debug, RAS, and ETM.                           |
| nCPUPORESET[CN:0] | Input     | Active-LOW core Cold reset.  |
| nGICRESET         | Input     | Active-LOW reset for all registers in the GICCLK domain.                             |
| nMBISTRESET       | Input     | Active-LOW reset for all MBIST logic.  |
| nPERIPHERSET      | Input     | Active-LOW reset for all registers in the PERIPHCLK domain.                          |
| nPRESET           | Input     | Active-LOW reset for all registers in the PCLK domain.                               |
| nSPORESET         | Input     | Active-LOW Cold reset for all registers in the SCLK domain.                          |
| nSRESET           | Input     | Active-LOW Warm reset for all registers in the SCLK domain, excluding RAS registers. |

## B.2.3 Configuration signals

This section describes the configuration signals.

**Table A-7: Configuration signals**

| Signal          | Direction       | Description  |   |                |   |                 |
|-----------------|-----------------|--|---|----------------|---|-----------------|
| AA64nAA32[PE:0] | Input           | <p>Register width state. The options are:</p> <table><tr><td>0</td><td>AArch32</td></tr><tr><td>1</td><td>AArch64</td></tr></table> <p><b>Note:</b><br/>This signal is only present for cores which support AArch32 execution at EL3. For example, if there are two single-threaded cores, and only core 1 supports executing AArch32 at EL3, then only bit[1] of this signal is present. If none of the cores support AArch32 execution at EL3, this signal is not present.</p> | 0 | AArch32        | 1 | AArch64         |
| 0               | AArch32         |  |   |                |   |                 |
| 1               | AArch64         |  |   |                |   |                 |
| CFGEND[PE:0]    | Input           | <p>Endianness configuration, controls the reset value of the SCTLR_EL3/SCTLR EE bit. The options are:</p> <table><tr><td>0</td><td>EE bit is LOW.</td></tr><tr><td>1</td><td>EE bit is HIGH.</td></tr></table>   | 0 | EE bit is LOW. | 1 | EE bit is HIGH. |
| 0               | EE bit is LOW.  |  |   |                |   |                 |
| 1               | EE bit is HIGH. |  |   |                |   |                 |

| Signal             | Direction | Description   |
|--------------------|-----------|---|
| CFGTE[PE:0]        | Input     | <p>Enable Thumb exceptions, controls the reset value of the SCTLR.TE bit. The options are:</p> <p><b>0</b>      TE is LOW.<br/><b>1</b>      TE is HIGH.</p> <p><b>Note:</b><br/>This signal is only present for cores which support AArch32 execution at EL3. For example, if there are two single-threaded cores, and only core 1 supports executing AArch32 at EL3, then only bit[1] of this signal is present. If none of the cores support AArch32 execution at EL3, this signal is not present.</p>   |
| CLUSTERIDAFF2[7:0] | Input     | Value read in ClusterID Affinity Level-2 field, MPIDR bits[23:16].  |
| CLUSTERIDAFF3[7:0] | Input     | Value read in ClusterID Affinity Level-3 field, MPIDR bits[39:32].  |
| CRYPTODISABLE      | Input     | <p>Disables the Cryptographic Extensions.</p> <p><b>Note:</b><br/>If Cryptographic Extensions are not enabled, this signal is not present.</p>  |
| GICCDISABLE        | Input     | Globally disables the CPU interface logic and routes the external interrupt signals directly to the cores.  |
| RVBARADDRy[p:2]    | Input     | Reset Vector Base Address for executing in 64-bit state.  |
| VINITHI[PE:0]      | Input     | <p>Enable high exception vectors, controls the reset value of the SCTLR.V bit. The options are:</p> <p><b>0</b>      Starts exceptions vectors at address 0x00000000.<br/><b>1</b>      Starts exceptions vectors at address 0xFFFF0000.</p> <p><b>Note:</b><br/>This signal is only present for cores which support AArch32 execution at EL3. For example, if there are two single-threaded cores, and only core 1 supports executing AArch32 at EL3, then only bit[1] of this signal is present. If none of the cores support AArch32 execution at EL3, this signal is not present.</p> |

## B.2.4 GIC signals

This section describes the GIC and AXI4 Stream Protocol signals.

The following table shows the GIC signals.

**Table A-8: GIC signals**

| Signal            | Direction | Description   |
|-------------------|-----------|---|
| nFIQ[PE:0]        | Input     | Active-LOW, level-sensitive fast interrupt request.                                 |
| nIRQ[PE:0]        | Input     | Active-LOW, level-sensitive interrupt request.                                      |
| nVFIQ[PE:0]       | Input     | Active-LOW, level-sensitive virtual fast interrupt request.                         |
| nVIRQ[PE:0]       | Input     | Active-LOW, level-sensitive virtual interrupt request.                              |
| nVCPUMNTIRQ[PE:0] | Output    | Active-LOW, level-sensitive virtual CPU interface maintenance interrupt PPI output. |

The following table shows the AXI4 Stream Protocol signals.



**Table A-9: AXI4 Stream Protocol signals**

| Signal         | Direction | Description   |
|----------------|-----------|---|
| IRITVALID      | Input     | Distributor to GIC CPU Interface messages. TVALID indicates that the master is driving a valid transfer.  |
| IRITREADY      | Output    | Distributor to GIC CPU Interface messages. TREADY indicates that the slave can accept a transfer in the current cycle.  |
| IRITDATA[15:0] | Input     | Distributor to GIC CPU Interface messages. TDATA is the primary payload that is used to provide the data that is passing across the interface.  |
| IRITLAST       | Input     | Distributor to GIC CPU Interface messages. TLAST indicates the boundary of a packet.  |
| IRITDEST[3:0]  | Input     | Distributor to GIC CPU Interface messages. TDEST provides routing information for the data stream. Depending on the cluster configuration, not all bits of this bus are used. Any unused upper bits should be tied LOW. |
| IRITWAKEUP     | Input     | AXI4 stream protocol activity indicator.  |
| ICCTVALID      | Output    | GIC CPU Interface to distributor messages. TVALID indicates that the master is driving a valid transfer.  |
| ICCTREADY      | Input     | GIC CPU Interface to distributor messages. TREADY indicates that the slave can accept a transfer in the current cycle.  |
| ICCTDATA[15:0] | Output    | GIC CPU Interface to distributor messages. TDATA is the primary payload that is used to provide the data that is passing across the interface.  |
| ICCTLAST       | Output    | GIC CPU Interface to distributor messages. TLAST indicates the boundary of a packet.  |
| ICCTID[3:0]    | Output    | GIC CPU Interface to distributor. TID is the data stream identifier that indicates different streams of data.   |
| ICCTWAKEUP     | Output    | AXI4 Stream Protocol activity indicator.  |

## B.2.5 Generic Timer signals

This section describes the Generic Timer signals.

**Table A-10: Generic Timer signals**

| Signal           | Direction | Description  |
|------------------|-----------|--|
| nCNTHPIRQ[PE:0]  | Output    | Active-LOW, level-sensitive Hypervisor physical timer event. |
| nCNTSNSIRQ[PE:0] | Output    | Active-LOW, level-sensitive Non-secure physical timer event. |
| nCNTPSIRQ[PE:0]  | Output    | Active-LOW, level-sensitive Secure physical timer event.     |
| nCNTVIRQ[PE:0]   | Output    | Active-LOW, level-sensitive virtual timer event.             |
| nCNTHVIRQ[PE:0]  | Output    | Active-LOW, level-sensitive Hypervisor virtual timer event.  |
| CNTVALUEB[63:0]  | Input     | Counter value in binary encoding.                            |

## B.2.6 Power management signals

This section describes the power management signals.



The PPMCTLx[3:0] signal is only present for certain cores. See the associated core Technical Reference Manual for details.

**Table A-11: Power management signals**

| Signal               | Direction | Description   |
|----------------------|-----------|---|
| ATCLKQACCEPTn        | Output    | Active-LOW signal that indicates that the cluster accepts the clock controller request.   |
| ATCLKQACTIVE         | Output    | Indicates that the cluster requires ATCLK to be active.   |
| ATCLKQDENY           | Output    | Active-LOW signal that indicates that the cluster denies the clock controller request.  |
| ATCLKQREQn           | Input     | Active-LOW signal that indicates that the clock controller wants to gate the clock.   |
| CLUSTERPACCEPT       | Output    | Indicates that the cluster accepts the power controller request.  |
| CLUSTERPACTIVE[19:0] | Output    | Indicates if the cluster is active in various power states.   |
| CLUSTERPDENY         | Output    | Indicates that the cluster denies the power controller request.   |
| CLUSTERPREQ          | Input     | Indicates that the power controller wants the cluster to move to a new power state.   |
| CLUSTERPSTATE[6:0]   | Input     | Power state that the power controller requires the cluster to move to.  |
| COREPACCEPTx         | Output    | Indicates that the referenced core accepts the power controller request.  |
| COREPACTIVEx[17:0]   | Output    | Indicates if the referenced core is active in various power states.   |
| COREPDENYx           | Output    | Indicates that the referenced core denies the power controller request.   |
| COREPREQx            | Input     | Indicates that the power controller wants the referenced core to move to a new power state.   |
| COREPSTATEx[5:0]     | Input     | Power state that the power controller requires the core to move to.   |
| EVENTIACK            | Output    | Event input request acknowledge. It is not asserted until EVENTIREQ is HIGH, and then remains asserted until after EVENTIREQ goes LOW.                            |
| EVENTIREQ            | Input     | Event input request for wake up from WFE state. It must remain asserted until EVENTIACK is asserted, and must not be re-asserted until EVENTIACK is LOW.          |
| EVENTOACK            | Input     | Event output request acknowledge. It must not be asserted until EVENTOREQ is HIGH, and then must remain asserted until after EVENTOREQ goes LOW.                  |
| EVENTOREQ            | Output    | Event output request for wake up, triggered by SEV instruction. It is only asserted when EVENTOACK is LOW, and then remains HIGH until after EVENTOACK goes HIGH. |
| GICCLKQACCEPTn       | Output    | Active-LOW signal that indicates that the cluster accepts the clock controller request.   |
| GICCLKQACTIVE        | Output    | Indicates that the cluster requires GICCLK to be active.  |
| GICCLKQDENY          | Output    | Indicates that the cluster denies the clock controller request.   |
| GICCLKQREQn          | Input     | Active-LOW signal that indicates that the clock controller wants the gate the clock.  |
| PCLKQACCEPTn         | Output    | Active-LOW signal that indicates that the cluster accepts the clock controller request.   |
| PCLKQACTIVE          | Output    | Indicates that the cluster requires PCLK to be active.  |
| PCLKQDENY            | Output    | Indicates that the cluster denies the clock controller request.   |
| PCLKQREQn            | Input     | Active-LOW signals that indicate that the clock controller wants the gate the clock.  |
| SCLKQACCEPTn         | Output    | Active-LOW signal that indicates that the cluster accepts the clock controller request.   |
| SCLKQACTIVE          | Output    | Indicates that the cluster requires SCLK to be active.  |
| SCLKQDENY            | Output    | Indicates that the cluster denies the clock controller request.   |
| SCLKQREQn            | Input     | Active-LOW signal that indicates that the clock controller wants the gate the clock.  |

| Signal          | Direction | Description  |
|-----------------|-----------|--|
| MPMMEN[PE:0]    | Input     | <p>Maximum Power Mitigation Mechanism (MPMM) enable signal. The options are:</p> <p><b>0</b> MPMM disabled.<br/><b>1</b> MPMM enabled.</p> <p>See the associated core Technical Reference Manual for details of the MPMM behavior if supported on your core.</p> <p><b>Note:</b><br/>This signal is only present for cores that support MPMM. For example, if there are two single-threaded cores, and only core 1 supports MPMM, then only one bit of this signal is present. If none of the cores support MPMM, then this signal is not present.</p> |
| MPMMSTATEx[1:0] | Input     | <p>MPMM state signal. If MPMM is enabled, selects which MPMM configuration to apply.</p> <p>If MPMM is supported on your core, see the associated core Technical Reference Manual for details of the MPMM configuration selection.</p> <p><b>Note:</b><br/>This signal is only present for cores that support MPMM configuration selection based on input pins. Not all cores that support MPMM support this.</p>  |
| PPMCTLx[3:0]    | Input     | <p>Power policy control signal.</p> <p><b>Note:</b><br/>This signal is only present for certain cores. See the associated core Technical Reference Manual for details.</p>   |

## B.2.7 Error signals

This section describes the error signals.

**Table A-12: Timing constraints for the error signals**

| Signal            | Direction | Description  |
|-------------------|-----------|--|
| nFAULTIRQ[CN+1:1] | Output    | Active-LOW, level-sensitive fault indicator for a detected 1-bit or 2-bit ECC or Parity error in the RAMs. Bits [CN+1:1] are for the L1 and L2 RAMs in each core. The bits are output from the PERIPHCLK domain.   |
| nFAULTIRQ[0]      | Output    | Active-LOW, level-sensitive fault indicator for a detected 1-bit or 2-bit ECC or Parity error in the RAMs. Bit [0] is for L3 or snoop filter RAMs. It is output from the SCLK domain.  |
| nERRIRQ[CN+1:1]   | Output    | Active-LOW, level-sensitive error indicator for an ECC error that causes potential data corruption or loss of coherency. Bits [CN+1:1] are for the L1 and L2 RAMs in each core. The bits are output from the PERIPHCLK domain.   |
| nERRIRQ[0]        | Output    | <p>Active-LOW, level-sensitive error indicator for an ECC error that causes potential data corruption or loss of coherency.</p> <p>Bit [0] is for L3 or snoop filter RAMs or ACE or CHI write transactions with a write response condition. It is output from the SCLK domain.</p> |

## B.2.8 ACP interface signals

This interface exists only if the DSU is configured to have the ACP interface.

**Table A-13: ACP wakeup signal**

| Signal   | Direction | Description                   |
|----------|-----------|-------------------------------|
| AWAKEUPS | Input     | ACP Slave activity indicator. |

**Table A-14: ACP write address channel signals**

| Signal            | Direction | Description                |
|-------------------|-----------|----------------------------|
| AWREADYS          | Output    | Write address ready.       |
| AWVALIDS          | Input     | Write address valid.       |
| AWIDS[7:0]        | Input     | Write address ID.          |
| AWADDRS[p:0]      | Input     | Write address.             |
| AWLENS[7:0]       | Input     | Write burst length.        |
| AWCACHES[3:0]     | Input     | Write cache type.          |
| AWDOMAINS[1:0]    | Input     | Write attributes.          |
| AWSNOOPS[3:0]     | Input     | Write request type.        |
| AWPROTS[2:0]      | Input     | Write protection type.     |
| AWSTASHLPIDS[3:0] | Input     | Write stash target core.   |
| AWSTASHLPIDENS    | Input     | Write stash target enable. |

**Table A-15: ACP write data channel signals**

| Signal        | Direction | Description                          |
|---------------|-----------|--------------------------------------|
| WREADYS       | Output    | Write data ready.                    |
| WVALIDS       | Input     | Write data valid.                    |
| WDATAS[127:0] | Input     | Write data.                          |
| WSTRBS[15:0]  | Input     | Write byte-lane strobes.             |
| WLASTS        | Input     | Write data last transfer indication. |

**Table A-16: ACP write response channel signals**

| Signal      | Direction | Description           |
|-------------|-----------|-----------------------|
| BREADYS     | Input     | Write response ready. |
| BVALIDS     | Output    | Write response valid. |
| BIDS[7:0]   | Output    | Write response ID.    |
| BRESPS[1:0] | Output    | Write response.       |

**Table A-17: ACP read address channel signals**

| Signal     | Direction | Description         |
|------------|-----------|---------------------|
| ARREADYS   | Output    | Read address ready. |
| ARVALIDS   | Input     | Read address valid. |
| ARIDS[7:0] | Input     | Read address ID.    |

| Signal         | Direction | Description           |
|----------------|-----------|-----------------------|
| ARADDRS[p:0]   | Input     | Read address.         |
| ARLENS[7:0]    | Input     | Read burst length.    |
| ARCACHES[3:0]  | Input     | Read cache type.      |
| ARDOMAINS[1:0] | Input     | Read attributes.      |
| ARPROTS[2:0]   | Input     | Read protection type. |

**Table A-18: ACP read data channel signals**

| Signal        | Direction | Description                         |
|---------------|-----------|-------------------------------------|
| RREADYS       | Input     | Read data ready.                    |
| RVALIDS       | Output    | Read data valid.                    |
| RIDS[7:0]     | Output    | Read data ID.                       |
| RDATAS[127:0] | Output    | Read data.                          |
| RRESPS[1:0]   | Output    | Read data response.                 |
| RLASTS        | Output    | Read data last transfer indication. |

## B.2.9 Peripheral port interface signals

This interface exists only if the DSU is configured to have a peripheral port interface.

**Table A-19: Peripheral port wakeup signals**

| Signal    | Direction | Description             |
|-----------|-----------|-------------------------|
| AWAKEUPMP | Output    | AXI activity indicator. |

**Table A-20: Peripheral port configuration signals**

| Signal         | Direction | Description                                      |
|----------------|-----------|--|
| ASTARTMP[p:20] | Input     | Start address for peripheral port address range. |
| AENDMP[p:20]   | Input     | End address for peripheral port address range.   |

**Table A-21: Peripheral port write address channel signals**

| Signal          | Direction | Description            |
|-----------------|-----------|------------------------|
| AWREADYMP       | Input     | Write address ready.   |
| AWVALIDMP       | Output    | Write address valid.   |
| AWIDMP[3:0]     | Output    | Write address ID.      |
| AWADDRMP[p:0]   | Output    | Write address.         |
| AWCACHEMP[3:0]  | Output    | Write cache type.      |
| AWDOMAINMP[1:0] | Output    | Write protection type. |
| AWPROTMP[2:0]   | Output    | Write protection type. |
| AWSIZEMP[2:0]   | Output    | Write size.            |

**Table A-22: Peripheral port write data channel signals**

| Signal        | Direction | Description                          |
|---------------|-----------|--------------------------------------|
| WREADYMP      | Input     | Write data ready.                    |
| WVALIDMP      | Output    | Write data valid.                    |
| WDATAMP[63:0] | Output    | Write data.                          |
| WSTRBMP[7:0]  | Output    | Write byte-lane strobes.             |
| WLASTMP       | Output    | Write data last transfer indication. |

**Table A-23: Peripheral port write response channel signals**

| Signal       | Direction | Description           |
|--------------|-----------|-----------------------|
| BREADYMP     | Output    | Write response ready. |
| BVALIDMP     | Input     | Write response valid. |
| BIDMP[3:0]   | Input     | Write response ID.    |
| BRESPMP[1:0] | Input     | Write response.       |

**Table A-24: Peripheral port read address channel signals**

| Signal          | Direction | Description           |
|-----------------|-----------|-----------------------|
| ARREADYMP       | Input     | Read address ready.   |
| ARVALIDMP       | Output    | Read address valid.   |
| ARIDMP[3:0]     | Output    | Read address ID.      |
| ARADDRMP[p:0]   | Output    | Read address.         |
| ARCACHEMP[3:0]  | Output    | Read cache type.      |
| ARPROTMP[2:0]   | Output    | Read protection type. |
| ARSIZEMP[2:0]   | Output    | Read size.            |
| ARDOMAINMP[1:0] | Output    | Write attributes.     |

**Table A-25: Peripheral port read data channel signals**

| Signal        | Direction | Description         |
|---------------|-----------|---------------------|
| RREADYMP      | Output    | Read data ready.    |
| RVALIDMP      | Input     | Read data valid.    |
| RIDMP[3:0]    | Input     | Read data ID.       |
| RDATAMP[63:0] | Input     | Read data.          |
| RRESPMP[1:0]  | Input     | Read data response. |

## B.2.10 Broadcast signals for the memory interface

The signals in the following table are common to the memory interface, irrespective of whether the ACE or CHI protocol is used.

**Table A-26: Broadcast signals**

| Signal                 | Direction | Description  |
|------------------------|-----------|--|
| BROADCASTCACHEMAINT    | Input     | Enable broadcasting of cache maintenance operations to downstream caches. The options are:<br><br><b>0</b> Cache maintenance operations are not broadcast to downstream caches.<br><b>1</b> Cache maintenance operations are broadcast to downstream caches.   |
| BROADCASTCACHEMAINTPOU | Input     | Enable broadcasting of cache maintenance operations to the point of unification. The options are:<br><br><b>0</b> Cache maintenance operations DCCMVAU and DC CVAU are not broadcast to other clusters. This is more efficient if all other clusters are Arm Cortex processors.<br><b>1</b> Cache maintenance operations DCCMVAU and DC CVAU are broadcast to other clusters.  |
| BROADCASTPERSIST       | Input     | Enable broadcasting of cache clean to the point of persistence operations. The options are:<br><br><b>0</b> DC CVAP instructions are treated the same as DC CVAC.<br><b>1</b> DC CVAP instructions are sent as a clean to the point of persistence transaction externally.   |
| BROADCASTOUTER         | Input     | Enable broadcasting of Outer Shareable transactions. The options are:<br><br><b>0</b> Outer Shareable transactions are not broadcast externally.<br><b>1</b> Outer Shareable transactions are broadcast externally.<br><br><b>Note:</b><br>Within the cluster, Inner Shareable and Outer Shareable memory are treated identically. Therefore, this signal also controls the broadcast of Inner Shareable transactions. |

## B.2.11 ACE interface signals

This interface exists only if the DSU is configured to have the ACE master interface.

The following tables show the signals present when one ACE interface is implemented. The signals have a suffix of M0. If a second ACE interface is implemented, then there is an additional set of ACE signals, suffixed by M1.

**Table A-27: ACE write address channel signals**

| Signal                      | Direction | Description       |
|-----------------------------|-----------|-------------------|
| AWADDRM0[q:0] <sup>15</sup> | Output    | Write address.    |
| AWBURSTM0[1:0]              | Output    | Write burst type. |

<sup>15</sup> The value of q is 43 when p is 39. For any other value of p, q is the same as p.

| Signal          | Direction | Description   |
|-----------------|-----------|---|
| AWCACHEM0[3:0]  | Output    | Write cache type.   |
| AWDOMAINM0[1:0] | Output    | Write shareability domain type.   |
| AWIDM0[7:0]     | Output    | Write address ID.   |
| AWLENM0[7:0]    | Output    | Write burst length.   |
| AWLOCKM0        | Output    | Write lock type.  |
| AWPROTM0[2:0]   | Output    | Write protection type.  |
| AWREADYM0       | Input     | Write address ready.  |
| AWSIZEM0[2:0]   | Output    | Write burst size.   |
| AWSNOOPM0[2:0]  | Output    | Write snoop request type.   |
| AWUNIQUEM0      | Output    | For WriteBack, WriteClean and WriteBackUC transactions. Indicates that the write is Unique (0b1) or shared (0b0). |
| AWVALIDM0       | Output    | Write address valid.  |

**Table A-28: ACE write data channel signals**

| Signal         | Direction | Description   |
|----------------|-----------|---|
| WDATAM0[127:0] | Output    | Write data.   |
| WIDM0[7:0]     | Output    | Write data ID. Only used when connecting to an AXI3 interconnect. |
| WLASTM0        | Output    | Write data last transfer indication.                              |
| WREADYM0       | Input     | Write data ready.   |
| WSTRBM0[15:0]  | Output    | Write byte-lane strobes.  |
| WVALIDM0       | Output    | Write data valid.   |

**Table A-29: ACE write data response channel signals**

| Signal       | Direction | Description           |
|--------------|-----------|-----------------------|
| BIDM0[7:0]   | Input     | Write response ID.    |
| BREADYM0     | Output    | Write response valid. |
| BRESPM0[1:0] | Input     | Write response.       |
| BVALIDM0     | Input     | Write response valid. |

**Table A-30: ACE read address channel signals**

| Signal                      | Direction | Description                    |
|-----------------------------|-----------|--------------------------------|
| ARADDRM0[q:0] <sup>15</sup> | Output    | Read address.                  |
| ARBURSTM0[1:0]              | Output    | Read burst type.               |
| ARCACHEM0[3:0]              | Output    | Read cache type.               |
| ARDOMAINM0[1:0]             | Output    | Read shareability domain type. |
| ARIDM0[8:0]                 | Output    | Read address ID.               |
| ARLENM0[7:0]                | Output    | Read burst length.             |
| ARLOCKM0                    | Output    | Read lock type.                |
| ARPROTM0[2:0]               | Output    | Read protection type.          |
| ARREADYM0                   | Input     | Read address ready.            |
| ARSIZEM0[2:0]               | Output    | Read burst size.               |



| Signal           | Direction | Description                            |
|------------------|-----------|--|
| ARSNOOPM0[3:0]   | Output    | Read snoop request type.               |
| ARVALIDM0        | Output    | Read address valid.                    |
| ARVMIDEXTM0[3:0] | Output    | Additional VMID bits for DVM messages. |

**Table A-31: ACE read data channel signals**

| Signal         | Direction | Description                         |
|----------------|-----------|-------------------------------------|
| RDATAM0[127:0] | Input     | Read data.                          |
| RIDM0[8:0]     | Input     | Read data ID.                       |
| RLASTM0        | Input     | Read data last transfer indication. |
| RREADYM0       | Output    | Read data ready.                    |
| RRESPM0[3:0]   | Input     | Read data response.                 |
| RVALIDM0       | Input     | Read data valid.                    |

**Table A-32: ACE coherency address channel signals**

| Signal                      | Direction | Description                            |
|-----------------------------|-----------|--|
| ACADDRM0[q:0] <sup>15</sup> | Input     | Snoop address.                         |
| ACPROTM0[2:0]               | Input     | Snoop protection type.                 |
| ACREADYM0                   | Output    | Master ready to accept snoop address.  |
| ACVMIDEXTM0[3:0]            | Input     | Additional VMID bits for DVM messages. |
| ACSNOOPM0[3:0]              | Input     | Snoop request type.                    |
| ACVALIDM0                   | Input     | Snoop address valid.                   |

**Table A-33: ACE coherency response channel signals**

| Signal        | Direction | Description                           |
|---------------|-----------|---------------------------------------|
| CRREADYM0     | Input     | Slave ready to accept snoop response. |
| CRVALIDM0     | Output    | Snoop response valid.                 |
| CRRESPM0[4:0] | Output    | Snoop response.                       |

**Table A-34: ACE coherency data channel handshake signals**

| Signal          | Direction | Description                       |
|-----------------|-----------|-----------------------------------|
| CDDATAM0[127:0] | Output    | Snoop data.                       |
| CDLASTM0        | Output    | Snoop data last transfer.         |
| CDREADYM0       | Input     | Slave ready to accept snoop data. |
| CDVALIDM0       | Output    | Snoop data valid.                 |

**Table A-35: ACE read and write acknowledge signals**

| Signal | Direction | Description        |
|--------|-----------|--------------------|
| RACKM0 | Output    | Read acknowledge.  |
| WACKM0 | Output    | Write acknowledge. |

**Table A-36: ACE low-power wakeup signals**

| Signal     | Direction | Description                     |
|------------|-----------|---------------------------------|
| AWAKEUPM0  | Output    | ACE pending activity indicator. |
| ACWAKEUPM0 | Input     | ACE snoop activity indicator.   |

**Table A-37: ACE coherency signals**

| Signal     | Direction | Description  |
|------------|-----------|--|
| SYSCOREQM0 | Output    | The cluster requests to be part of the coherency domain and can safely receive snoops. |
| SYSCOACKM0 | Input     | The system might send snoops to the cluster.   |

**Table A-38: ACE SRCATTRM0/1 bus signals**

| Signal            | Direction | Description   |
|-------------------|-----------|---|
| RDSRCATTRM0[11:0] | Output    | These signals provide additional information about the transactions.<br><br>The functionality of the individual bits is described in the following table. |
| WRSRCATTRM0[11:0] | Output    |   |
| RDSRCATTRM1[11:0] | Output    |   |
| WRSRCATTRM1[11:0] | Output    |   |

**Table A-39: ACE SRCATTRM0/1 bus signal functionality**

| Bits   | Field           | Description   | Encoding  |
|--------|-----------------|---|---|
| [11:7] | Source          | Source of request.<br><br>This field has the same encoding as TXREQ[LPID] field on CHI. | The possible values are:<br><b>0x0-0xF</b><br>Core ID/Thread ID<br><br>Bit 0 is the thread ID, or 0 for single threaded processors.<br><b>0x1E</b><br>ACP<br><b>0x1F</b><br>L3 cache eviction |
| [6]    | Outer cacheable | Outer cacheability of access.   | <b>0</b> Outer non-cacheable (or device).<br><b>1</b> Outer cacheable (WT/WB)   |

| Bits  | Field               | Description  | Encoding  |
|-------|---------------------|--|---|
| [5]   | Stream_or_transient | <p>The exact conditions for this bit depend on the core, which indicates that this access was either:</p> <ul style="list-style-type: none"> <li>Generated by a non-temporal instruction, or</li> <li>Generated as part of a detected stream, or</li> <li>Marked as transient in the translation tables.</li> </ul> <p><b>Note:</b><br/>This bit only applies to accesses that are directly part of the stream, and not to cache evictions that are a side effect of the stream.</p> | <p><b>0</b> Normal</p> <p><b>1</b> Streaming or transient</p>   |
| [4]   | Prefetch            | Demand or Prefetch.  | <p><b>0</b> Demand</p> <p><b>1</b> Prefetch</p>   |
| [3:2] | Subsource           | Instruction, or Data or Table walk.  | <p>The possible values are:</p> <p><b>0b00</b> Instruction</p> <p><b>0b01</b> Data</p> <p>This includes all miscellaneous transaction types, for example CMOs and DVMs.</p> <p><b>0b10</b> Table walk</p> |
| [1:0] | PBHA                | Page-based hardware attributes.  |   |

## B.2.12 CHI interface signals

This interface exists only if the DSU is configured to have the CHI master interface.

### Common signals for single or dual CHI interfaces

The following signal is present in both single or dual CHI configurations.



For information on the signal width naming convention values, see [B.1 Signal naming convention](#) on page 221.

**Table A-40: CHI common configuration signal**

| Signal                          | Direction | Description   |
|---------------------------------|-----------|---|
| BROADCASTATOMIC <sup>16</sup> . | Input     | Enable broadcasting of atomic transactions. The options are:<br><br><div> <div>0</div> <div>Atomic transactions are not sent externally.</div> </div> <div> <div>1</div> <div>Atomic transactions are sent externally.</div> </div> |

### Per-interface signals for single or dual CHI interfaces

The following tables show the signals present for either a single CHI interface, or the signals for CHI interface 0 in a dual CHI implementation. If a second CHI interface is implemented, then there is an additional set of CHI signals, suffixed by M1. For example, in a dual CHI configuration, the signals for the DSU cluster CHI Node Identifiers for the two CHI interfaces are NODEID[10:0] and NODEIDM1[10:0].

**Table A-41: CHI configuration signals**

| Signal       | Direction | Description                      |
|--------------|-----------|----------------------------------|
| NODEID[10:0] | Input     | DSU cluster CHI Node Identifier. |

**Table A-42: CHI activity signals**

| Signal          | Direction | Description                          |
|-----------------|-----------|--------------------------------------|
| RXSACTIVE       | Input     | Receive pending activity indicator.  |
| TXSACTIVE       | Output    | Transmit pending activity indicator. |
| RXLINKACTIVEREQ | Input     | Receive link active request.         |
| RXLINKACTIVEACK | Output    | Receive link active acknowledge.     |
| TXLINKACTIVEREQ | Output    | Transmit link active request.        |
| TXLINKACTIVEACK | Input     | Transmit link active acknowledge.    |

**Table A-43: CHI transmit request virtual channel signals**

| Signal           | Direction | Description                               |
|------------------|-----------|---|
| TXREQFLITPEND    | Output    | Transmit Request Flit pending.            |
| TXREQFLITV       | Output    | Transmit Request Flit valid.              |
| TXREQFLIT[j+q:0] | Output    | Transmit Request Flit payload.            |
| TXREQLCRDV       | Input     | Transmit Request link-layer credit valid. |

**Table A-44: CHI transmit response virtual channel signals**

| Signal         | Direction | Description                                |
|----------------|-----------|--|
| TXRSPFLITPEND  | Output    | Transmit response flit pending.            |
| TXRSPFLITV     | Output    | Transmit response flit valid.              |
| TXRSPFLIT[i:0] | Output    | Transmit response flit.                    |
| TXRSPLCRDV     | Input     | Transmit response link-layer credit valid. |

<sup>16</sup> This signal is sampled on reset like the other broadcast signals that are listed in [B.2.10 Broadcast signals for the memory interface](#) on page 230

**Table A-45: CHI transmit data virtual channel signals**

| Signal         | Direction | Description                            |
|----------------|-----------|--|
| TXDATFLITPEND  | Output    | Transmit Data Flit pending.            |
| TXDATFLITV     | Output    | Transmit Data Flit valid.              |
| TXDATFLIT[d:0] | Output    | Transmit Data Flit.                    |
| TXDATLCRDV     | Input     | Transmit Data link-layer credit valid. |

**Table A-46: CHI receive snoop virtual channel signals**

| Signal           | Direction | Description                            |
|------------------|-----------|--|
| RXSNPFLITPEND    | Input     | Receive Snoop Flit pending.            |
| RXSNPFLITV       | Input     | Receive Snoop Flit valid.              |
| RXSNPFLIT[k+q:0] | Input     | Receive Snoop Flit.                    |
| RXSNPLCRDV       | Output    | Receive Snoop link-layer credit valid. |

**Table A-47: CHI receive response virtual channel signals**

| Signal         | Direction | Description                               |
|----------------|-----------|---|
| RXRSPFLITPEND  | Input     | Receive Response Flit pending.            |
| RXRSPFLITV     | Input     | Receive Response Flit valid.              |
| RXRSPFLIT[i:0] | Input     | Receive Response Flit.                    |
| RXRSPLCRDV     | Output    | Receive Response link-layer credit valid. |

**Table A-48: CHI receive data virtual channel signals**

| Signal         | Direction | Description                           |
|----------------|-----------|---------------------------------------|
| RXDATFLITPEND  | Input     | Receive Data Flit pending.            |
| RXDATFLITV     | Input     | Receive Data Flit valid.              |
| RXDATFLIT[d:0] | Input     | Receive Data Flit.                    |
| RXDATLCRDV     | Output    | Receive Data link-layer credit valid. |

**Table A-49: CHI coherency signals**

| Signal   | Direction | Description  |
|----------|-----------|--|
| SYSCOREQ | Output    | The cluster requests to be part of the coherency domain and can safely receive snoops. |
| SYSCOACK | Input     | The system might send snoops to the cluster.   |

**Table A-50: CHI REQSRCATTR bus signals**

| Signal          | Direction | Description   |
|-----------------|-----------|---|
| REQSRCATTR[6:0] | Output    | These signals provide additional information about the transactions.<br><br>The functionality of the individual bits is described in the following table. |

**Table A-51: CHI REQSRCATTR bus signal functionality**

| Bits  | Field               | Description   | Encoding  |
|-------|---------------------|---|---|
| [6]   | Outer cacheable     | Outer cacheability of access.   | <b>0</b> Outer non-cacheable (or device).<br><b>1</b> Outer cacheable (WT/WB)   |
| [5]   | Stream_or_transient | The exact conditions for this bit depend on the core, which indicates that this access was either: <ul style="list-style-type: none"> <li>Generated by a non-temporal instruction, or</li> <li>Generated as part of a detected stream, or</li> <li>Marked as transient in the translation tables.</li> </ul> <b>Note:</b><br>This bit only applies to accesses that are directly part of the stream, and not to cache evictions that are a side effect of the stream. | <b>0</b> Normal<br><b>1</b> Streaming or transient  |
| [4]   | Prefetch            | Demand or Prefetch.   | <b>0</b> Demand<br><b>1</b> Prefetch  |
| [3:2] | Subsource           | Instruction, or Data or Table walk.   | The possible values are:<br><b>0b00</b> Instruction<br><b>0b01</b> Data<br><br>This includes all miscellaneous transaction types, for example, CMOs and DVMs.<br><b>0b10</b> Table walk |
| [1:0] | PBHA                | Page-based hardware attributes.   |   |

### B.2.12.1 System address map signals

This interface exists only if the DSU is configured to have the CHI interface and the CMN-600 SAM option is enabled. In a dual CHI configuration, there is only one set of system address map signals.

**Table A-52: CHI system address map signals**

| Signal         | Direction | Description                                   |
|----------------|-----------|---|
| RXPUBCFGACTIVE | Input     | SAM configuration pending activity indicator. |
| RXPUBFLITPEND  | Input     | Receive Data Flit pending.                    |

| Signal             | Direction | Description                            |
|--------------------|-----------|--|
| RXPUBFLITV         | Input     | Receive Data Flit valid.               |
| RXPUBFLIT[34:0]    | Input     | Receive Data Flit.                     |
| RXPUBLICRDV_RP1    | Output    | Receive Data link-layer credit valid.  |
| RXPUBLINKACTIVEREQ | Input     | Receive link active request.           |
| RXPUBLINKACTIVEACK | Output    | Receive link active acknowledge.       |
| RXPUBLINKFLIT      | Input     | Receive link-layer flit.               |
| TXPUBFLITPEND      | Output    | Transmit Data Flit pending.            |
| TXPUBFLITV         | Output    | Transmit Data Flit valid.              |
| TXPUBFLIT[34:0]    | Output    | Transmit Data Flit.                    |
| TXPUBLCRDV_RP1     | Input     | Transmit Data link-layer credit valid. |
| TXPUBLINKACTIVEREQ | Output    | Transmit link active request.          |
| TXPUBLINKACTIVEACK | Input     | Transmit link active acknowledge.      |
| TXPUBLINKFLIT      | Output    | Transmit link-layer flit.              |

**Table A-53: CHI system address map configuration signals**

| Signal                    | Direction | Description                              |
|---------------------------|-----------|--|
| RNSAM_CFGM_DESTID[15:0]   | Input     | CMN-600 CFG master destination ID.       |
| RNSAM_CFGS_NODEID[15:0]   | Input     | Source ID for the cluster CFG slave.     |
| RNSAM_DEFAULT_TGTID[10:0] | Input     | Default target ID of the boot node.      |
| RNSAM_DN_TGTID[10:0]      | Input     | Target ID for the DN that handles DVMs.  |
| RNSAM_PGMED_RDY           | Input     | Reset value of the Programmed Ready bit. |

### B.2.13 DebugBlock APB interface signals

The interface between the DebugBlock and the cluster consists of a pair of APB interfaces, one in each direction. The signals on the cluster interface should be connected to the equivalently named signal on the DebugBlock.

The following table shows the signals between the cluster APB master and the DebugBlock APB slave.

**Table A-54: Cluster to DebugBlock APB interface signals**

| Signal        | Direction | Description                                 |
|---------------|-----------|---|
| PSELCD        | Output    | APB CD select.                              |
| PADDRCD[8:2]  | Output    | APB address.                                |
| PENABLECD     | Output    | APB enable.                                 |
| PWDATACD[8:0] | Output    | APB write data.                             |
| PREADYCD      | Input     | APB slave ready, used to extend a transfer. |
| PSLVERRCD     | Input     | APB slave transfer error.                   |
| PWAKEUPCD     | Output    | APB activity indicator.                     |

The following table shows the signals between the DebugBlock APB master and the cluster APB slave.

**Table A-55: DebugBlock to cluster APB interface signals**

| Signal         | Direction | Description   |
|----------------|-----------|---|
| PSELDC         | Input     | APB DC select.  |
| PADDRDC[19:2]  | Input     | APB address.  |
| PENABLEDC      | Input     | APB enable.   |
| PWRITEDC       | Input     | APB read/write indicator. When asserted, indicates a write. |
| PWDATADC[31:0] | Input     | APB write data.   |
| PRDATADC[31:0] | Output    | APB read data.  |
| PREADYDC       | Output    | APB slave ready, used to extend a transfer.                 |
| PSLVERRDC      | Output    | APB slave transfer error.                                   |
| PWAKEUPDC      | Input     | APB activity indicator.                                     |

The following table shows the miscellaneous debug signals.

**Table A-56: Miscellaneous debug signals**

| Signal          | Direction | Description  |
|-----------------|-----------|--|
| DBGRSTREQ[PE:0] | Output    | Debug reset request.   |
| nCOMMIRQ[PE:0]  | Output    | Active-LOW, level-sensitive comms channel receive or transmit interrupt request. |
| DBGGEN          | Input     | Invasive debug enable.   |
| NIDEN           | Input     | Non-invasive debug enable.   |
| SPIDEN          | Input     | Secure privilege invasive debug enable.  |
| SPNIDEN         | Input     | Secure privilege non-invasive debug enable.                                      |
| DBGCONNECTED    | Input     | A debugger is connected and so the DebugBlock can be accessed on boot.           |

## B.2.14 ATB interface signals

The following table shows the ATB interface signals.

**Table A-57: ATB interface signals**

| Signal         | Direction | Description              |
|----------------|-----------|--------------------------|
| ATREADYMx      | Input     | ATB device ready.        |
| AFVALIDMx      | Input     | FIFO flush request.      |
| ATDATAMx[31:0] | Output    | Data.                    |
| ATVALIDMx      | Output    | Data valid.              |
| ATBYTESMx[1:0] | Output    | Data size.               |
| AFREADYMx      | Output    | FIFO flush finished      |
| ATIDMx[6:0]    | Output    | Trace source ID.         |
| ATWAKEUPMx     | Output    | ATB activity indicator.  |
| SYNCREQMx      | Input     | Synchronization request. |



## B.2.15 Timestamp signal

This section describes the timestamp signal.

The TSVALUEB input provides the incrementing count value that is required for generating the trace timestamp for each core. The DSU includes an interpolator on the TSVALUEB input for each of the cores.

**Table A-58: Timestamp signal**

| Signal         | Direction | Description                   |
|----------------|-----------|-------------------------------|
| TSVALUEB[63:0] | Input     | Timestamp in binary encoding. |

## B.2.16 PMU interface signals

This section describes the PMU interface signals.

**Table A-59: PMU interface signals**

| Signal             | Direction | Description  |
|--------------------|-----------|--|
| COREINSTRRET[PE:0] | Output    | Indicates that a core has retired at least one instruction recently. This signal is output from the PERIPHCLK domain.  |
| COREINSTRRUN[PE:0] | Output    | Indicates that a core is in a running state (powered up and not in WFI or WFE low-power states). This signal is output from the PERIPHCLK domain.  |
| nCLUSTERPMUIRQ     | Output    | Active-LOW, level-sensitive Cluster PMU interrupt request. This signal is output from the SCLK domain.   |
| nPMBIRQ[PE:0]      | Output    | Active-LOW, level-sensitive SPE interrupt request. This signal only exists for cores that support the architectural <i>Statistical Profiling Extension</i> (SPE). This signal is output from the PERIPHCLK domain. |
| nPMUIRQ[PE:0]      | Output    | Active-LOW, level-sensitive PMU interrupt request.   |
| PMUSNAPSHOTACK     | Output    | Acknowledge a snapshot request. This signal forms a four-phase handshake with PMUSNAPSHOTREQ.  |
| PMUSNAPSHOTREQ     | Input     | Request for a snapshot of the PMU counters. When asserted, this signal must remain HIGH until PMUSNAPSHOTACK is asserted. It must not be reasserted until PMUSNAPSHOTACK is deasserted.                            |

## B.2.17 ELA signal

This signal is present only if ELA-500 support is configured.

**Table A-60: ELA signal**

| Signal    | Direction | Description  |
|-----------|-----------|--|
| STOPCLOCK | Output    | ELA request for the system to stop the clocks. This signal is generated from multiple clock domains and must be synchronized in the system before use. |

## B.2.18 DFT interface signals

This section describes the DFT signals.

**Table A-61: DFT interface signals**

| Signal                  | Direction | Description   |
|-------------------------|-----------|---|
| DFTRAMHOLD              | Input     | Disables the RAM chip select during scan shift.   |
| DFTRSTDISABLE[1:0]      | Input     | Disables internal synchronized reset during scan shift.   |
| DFTCGEN                 | Input     | Forces on the clock gates during scan shift.  |
| DFTMCPHOLD              | Input     | Disables Multicycle Paths on RAM interfaces.  |
| DFTCORECLKDISABLE[CN:0] | Input     | Disables specific CORECLK signals from toggling during scan, to reduce test power if only part of the design is being scanned.  |
| DFTSCLKBYPASS           | Input     | Bypasses clock stretching logic on L3 data RAMs.<br><br><b>Note:</b><br>This signal is only present when the <code>L3_DATA_STRETCH_CLK</code> configuration option is enabled. See <a href="#">2.3 Implementation options</a> on page 18. |

## B.2.19 MBIST interface signals

This section describes the MBIST interface signals.

**Table A-62: MBIST interface signals**

| Signal   | Direction | Description         |
|----------|-----------|---------------------|
| MBISTREQ | Input     | MBIST test request. |

## B.3 DebugBlock signals

This section shows the DebugBlock signals.

### B.3.1 Clock signal

The following table shows the clock signal.

**Table A-63: Clock signal**

| Signal | Direction | Description                                 |
|--------|-----------|---|
| PCLK   | Input     | Clock for the DebugBlock and APB interface. |

### B.3.2 Reset signal

The following table shows the reset signal.

**Table A-64: Reset signal**

| Signal  | Direction | Description                                |
|---------|-----------|--|
| nPRESET | Input     | Reset for all registers in the DebugBlock. |

### B.3.3 Power and clock gate control signals

The following table shows the power and clock gate control signals.

**Table A-65: Power and clock gate signals**

| Signal          | Direction | Description   |
|-----------------|-----------|---|
| PDBGCLKQREQn    | Input     | Active-LOW signal that indicates that the clock controller wants to gate the clock.   |
| PWRQREQn        | Input     | Active-LOW signal that indicates that the power controller wants to power down the DebugBlock.  |
| PDBGCLKQACTIVE  | Output    | Indicates that the DebugBlock requires PCLK to be active.   |
| PDBGCLKQACCEPTn | Output    | Active-LOW signal that indicates that the DebugBlock accepts the clock controller request.  |
| PDBGCLKQDENY    | Output    | Indicates that the DebugBlock denies the clock controller request.  |
| PWRQACTIVE      | Output    | Indicates that the DebugBlock requires power.   |
| PWRQACCEPTn     | Output    | Active-LOW signal that indicates that the DebugBlock accepts the power controller request. It causes all future APB accesses to the DebugBlock to receive a PSLVERR response. |
| PWRQDENY        | Output    | Indicates that the DebugBlock denies the power controller request because of an ongoing transaction.  |

### B.3.4 Configuration signals

The following table shows the configuration signals. All these configuration pins are sampled only on reset.



The following signals must be tied to the same values as the corresponding cluster signals.

**Table A-66: Configuration signals**

| Signal             | Direction | Description   |
|--------------------|-----------|---|
| CLUSTERIDAFF2[7:0] | Input     | Value read in ClusterID Affinity Level-2 field, MPIDR bits [23:16]. |
| CLUSTERIDAFF3[7:0] | Input     | Value read in ClusterID Affinity Level-3 field, MPIDR bits [39:32]. |

| Signal      | Direction | Description  |
|-------------|-----------|--|
| ELADISABLE  | Input     | Disables the ELA logic in the cluster.<br><br>This signal is not present if ELA support is not implemented.<br><br><b>Note:</b><br>The ELA is intended only for silicon debug of sample devices. The ELADISABLE pin must be driven HIGH in any production devices. |
| GICCDISABLE | Input     | Globally disables the CPU interface logic and routes the external signals directly to the cores.   |

### B.3.5 Debug signals

This section describes the debug interface signals.

The following table shows the external debug APB interface signals.



The PPROTDBG[2:0] signal is only included when the DSU is used with a core that implements the Arm®v8.4-A debug architecture.

**Table A-67: External debug APB interface signals**

| Signal          | Direction | Description  |
|-----------------|-----------|--|
| PSELDBG         | Input     | APB select.  |
| PADDRDBG[23:2]  | Input     | APB address.   |
| PADDRDBG31      | Input     | APB address bit[31].<br><br>Controls the ETM OS Lock mechanism.<br><br><b>Note:</b><br>Software lock is not supported. |
| PENABLEDBG      | Input     | APB enable.  |
| PPROTDBG[2:0]   | Input     | APB protection type. Only present when a core is used that implements the Arm®v8.4-A debug architecture.               |
| PWRITEDBG       | Input     | APB read/write indicator. When asserted, indicates a write.  |
| PWDATADBG[31:0] | Input     | APB write data.  |
| PWAKEUPDBG      | Input     | APB activity indicator.  |
| PRDATADBG[31:0] | Output    | APB read data.   |
| PREADYDBG       | Output    | APB slave ready, used to extend a transfer.  |
| PSLVERRDBG      | Output    | APB slave transfer error.  |

The following table shows the debug authentication signals.



The signals NIDEN and SPNIDEN are not present if the core implements the Arm®v8.4-A Debug architecture.

**Table A-68: Debug authentication signals**

| Signal             | Direction | Description  |
|--------------------|-----------|--|
| DBGGEN             | Input     | Invasive debug enable.   |
| NIDEN              | Input     | Non-invasive debug enable. Not present if core implements Arm®v8.4-A debug architecture.                   |
| SPIDEN             | Input     | Secure privilege invasive debug enable.  |
| SPNIDEN            | Input     | Secure privilege non-invasive debug enable. Not present if core implements Arm®v8.4-A Debug.               |
| DBGPWRUPREQ[PE:0]  | Output    | Request to power up a PE for debug.  |
| CLUSTERDBGPWRUPREQ | Output    | Request power to the cluster domain. Supports the powerup of Cluster-level ELA from the external debugger. |

The following table shows the signals between the cluster APB master and the DebugBlock APB slave.

**Table A-69: Cluster to DebugBlock APB signals**

| Signal        | Direction | Description                                 |
|---------------|-----------|---|
| PADDRCD[8:2]  | Input     | APB address.                                |
| PENABLECD     | Input     | APB enable.                                 |
| PREADYCD      | Output    | APB slave ready, used to extend a transfer. |
| PSELCD        | Input     | APB select.                                 |
| PSLVERRCD     | Output    | APB slave transfer error.                   |
| PWAKEUPCD     | Input     | APB activity indicator.                     |
| PWDATACD[8:0] | Input     | APB write data.                             |

The following table shows the signals between the DebugBlock APB master and the cluster APB slave.



The PPROTDC[1:1] signal is only included when the DSU is used with a core that implements the Arm®v8.4-A debug architecture.

**Table A-70: DebugBlock to cluster APB signals**

| Signal         | Direction | Description   |
|----------------|-----------|---|
| PADDRDC[19:2]  | Output    | APB address.  |
| PENABLEDC      | Output    | APB enable.   |
| PPROTDC[1:1]   | Output    | APB security state. Only present when a core is used that implements the Arm®v8.4-A debug architecture. |
| PRDATADC[31:0] | Input     | APB read data.  |
| PREADYDC       | Input     | APB slave ready, used to extend a transfer.   |

| Signal         | Direction | Description   |
|----------------|-----------|---|
| PSELDC         | Output    | APB select.   |
| PSLVERRDC      | Input     | APB slave transfer error.                                   |
| PWAKEUPDC      | Output    | APB activity indicator.                                     |
| PWDATADC[31:0] | Output    | APB write data.   |
| PWRITEDC       | Output    | APB read/write indicator. When asserted, indicates a write. |

### B.3.6 CTI interface signals

This section describes the CTI interface signals.

**Table A-71: CTI interface signals**

| Signal           | Direction | Description                                |
|------------------|-----------|--|
| CTICHIN[3:0]     | Input     | Channel In.                                |
| CTICHOUTACK[3:0] | Input     | Channel Out acknowledge.                   |
| CTICHOUT[3:0]    | Output    | Channel Out.                               |
| CTICHINACK[3:0]  | Output    | Channel In acknowledge.                    |
| CISBYPASS        | Input     | Channel interface sync bypass.             |
| CIHSBYPASS[3:0]  | Input     | Channel interface H/S bypass.              |
| CTIIRQ[PE:0]     | Output    | Active-HIGH, edge-sensitive CTI interrupt. |
| CTIIRQACK[PE:0]  | Input     | CTI interrupt acknowledge.                 |

### B.3.7 DFT signals

The following table shows the DFT signals.

**Table A-72: DFT interface signals**

| Signal             | Direction | Description   |
|--------------------|-----------|---|
| DFTCGEN            | Input     | Forces on the clock gates during scan shift.            |
| DFTRSTDISABLE[1:0] | Input     | Disables internal synchronized reset during scan shift. |

# Appendix C Revisions

This appendix describes the technical changes between released issues of this book.

## C.1 Revisions

This section describes the technical changes between released issues of this document.

**Table A-73: Issue 0000-00**

| Change        | Location |
|---------------|----------|
| First release | -        |

**Table A-74: Differences between issue 0000-00 and issue 0001-00**

| Change                         | Location           |
|--------------------------------|--------------------|
| Added support for ACE and CHI. | The whole document |

**Table A-75: Differences between issue 0001-00 and issue 0002-00**

| Change  | Location  |
|---|---|
| Change of terminology, <i>Flexible Cluster Microarchitecture</i> (FCM) to <i>DynamIQ™ Shared Unit</i> (DSU) | The whole document  |
| Updated CLUSTERPSTATE value   | <a href="#">5.5.4 Power mode encoding</a> on page 47                                |
| Added description of error protection of snoop filter tag RAMs  | <a href="#">6.5 L3 cache ECC and parity</a> on page 57                              |
| Added traps and enables descriptions to AARCH64 registers   | <a href="#">11. Control registers</a> on page 95                                    |
| Updated register description  | <a href="#">12.3 ERR1CTLR, Error Record Control Register</a> on page 134            |
| Updated revision field  | <a href="#">16.5 CTIPIDR2, CTI Peripheral Identification Register 2</a> on page 208 |
| Updated ROM table Peripheral Identification Register 2 address  | <a href="#">17.2 ROM table register summary</a> on page 216                         |

**Table A-76: Differences between issue 0001-00 and issue 0100-00**

| Change  | Location                                   |
|---|--|
| Added L3_CACHE implementation option            | Configuration parameters table             |
| Added L3_DATA_STRETCH_CLK implementation option | Configuration parameters table             |
| Added INTERLEAVE_ADDR_BIT implementation option | Configuration parameters table             |
| Added ELA_RAM_ADDR_SIZE implementation option   | Configuration parameters table             |
| Added new L3 cache operating capacity level     | Section L3 cache partial powerdown         |
| Added new section                               | L3 cache partitioning                      |
| Updated L3 cache stashing information           | L3 cache stashing                          |
| Updated ACE interface information               | Dual ACE interfaces                        |
| Updated ACP interface information               | About the ACP                              |
| Changes to Cluster Configuration Register       | CLUSTERCFR, Cluster Configuration Register |

| Change  | Location                          |
|---|-----------------------------------|
| Changes to allocation of CTI input trigger events | CTI triggers                      |
| Changes to PMU events                             | PMU events                        |
| Changes to ROM table reset values                 | ROM table register summary        |
| Changes to DFT interface signals                  | DFT interface signals             |
| Changes to Cluster to DebugBlock APB signals      | Cluster to DebugBlock APB signals |

**Table A-77: Differences between issue 0100-00 and 0200-00**

| Change   | Location   |
|--|--|
| Support added for three sets of cores of different microarchitecture.            | <a href="#">2.1 About the DSU</a> on page 16   |
| ARMv8.2-A architecture cores supported.  | <a href="#">2.2 Features</a> on page 17  |
| NUM_LITTLE_CORES parameter includes 5 cores.                                     | <a href="#">2.3 Implementation options</a> on page 18  |
| NUM_OTHER_CORES parameter added.   | <a href="#">2.3 Implementation options</a> on page 18  |
| OTHER_CORE_TYPE parameter added.   | <a href="#">2.3 Implementation options</a> on page 18  |
| L3_CACHE_SIZE parameter extended to include 1.5MB cache size option.             | <a href="#">2.3 Implementation options</a> on page 18  |
| NUM_SLICES parameter added.  | <a href="#">2.3 Implementation options</a> on page 18  |
| Added note to clarify the use of L3_DATA_STRETCH_CLK parameter on performance.   | <a href="#">2.3 Implementation options</a> on page 18  |
| Added note to clarify the use of ASYNC_BRIDGE parameter.                         | <a href="#">2.3 Implementation options</a> on page 18  |
| Added note to ELA parameter about the 2 methods of creating unique instances.    | <a href="#">2.3 Implementation options</a> on page 18  |
| New section Cluster Configurations added.  | <a href="#">2.3.1 Cluster configurations</a> on page 22                                      |
| AMBA 5 Issue C support added.  | <a href="#">2.4 Supported standards and specifications</a> on page 22                        |
| Summarized r2p0 main technical changes.  | <a href="#">2.7 Product revisions</a> on page 24   |
| Added note about CPU bridge RTL.   | <a href="#">3.1 Components</a> on page 25  |
| L3 cache 1.5MB cache size.   | <a href="#">3.1 Components</a> on page 25  |
| Changes between power mode transitions MEM_RET and FUNC_RET.                     | <a href="#">5. Power management</a> on page 37   |
| Updated note on L3 cache slices.   | <a href="#">6.7 Cache slices and portions</a> on page 62                                     |
| Updated description for Read ID Attribute in table Encodings for ARIDM[8:0].     | <a href="#">7.4 ACE master interface attributes</a> on page 68                               |
| Updated Snoop acceptance capability description in ACE channel properties table. | <a href="#">7.5 ACE channel properties</a> on page 71  |
| Added section CHI version.   | <a href="#">8.2 CHI version</a> on page 79   |
| Updated Snoop acceptance capability description.                                 | <a href="#">8.6 CHI channel properties</a> on page 82.                                       |
| Updated Cluster Main RevisionID Register (CLUSTERIDR) description for r2p0.      | <a href="#">11.9 CLUSTERIDR, Cluster Main Revision ID Register</a> on page 109               |
| Added Cluster Count Enable Set Register (CLUSTERPMCNTENSET) description.         | <a href="#">13.5 CLUSTERPMCNTENSET, Cluster Count Enable Set Register</a> on page 154        |
| Added Cluster Count Enable Clear Register (CLUSTERPMCNTENCLR) description.       | <a href="#">13.6 CLUSTERPMCNTENCLR, Cluster Count Enable Clear Register</a> on page 156      |
| Added Cluster Overflow Flag Status Set Register (CLUSTERPMOVSSSET) description.  | <a href="#">13.7 CLUSTERPMOVSSSET, Cluster Overflow Flag Status Set Register</a> on page 158 |



| Change   | Location   |
|--|--|
| Added Cluster Overflow Flag Status Clear Register (CLUSTERPMOVSCCLR) description.    | <a href="#">13.8 CLUSTERPMOVSCCLR, Cluster Overflow Flag Status Clear Register</a> on page 160   |
| Added Cluster Event Counter Selection Register (CLUSTERPMSELR) description.          | <a href="#">13.9 CLUSTERPMSELR, Cluster Event Counter Selection Register</a> on page 162         |
| Added Cluster Interrupt Enable Set Register (CLUSTERPMINTENSET) description.         | <a href="#">13.10 CLUSTERPMINTENSET, Cluster Interrupt Enable Set Register</a> on page 164       |
| Added Cluster Interrupt Enable Clear Register (CLUSTERPMINTENCLR) description.       | <a href="#">13.11 CLUSTERPMINTENCLR, Cluster Interrupt Enable Clear Register</a> on page 166     |
| Added Cluster Performance Monitors Cycle Counter (CLUSTERPMCCNTR) description.       | <a href="#">13.12 CLUSTERPMCCNTR, Cluster Performance Monitors Cycle Counter</a> on page 168     |
| Added Cluster Selected Event Type Register (CLUSTERPMXEVTYPER) description.          | <a href="#">13.13 CLUSTERPMXEVTYPER, Cluster Selected Event Type Register</a> on page 169        |
| Added Cluster Selected Event Counter Register (CLUSTERPMXVCNTR) description.         | <a href="#">13.14 CLUSTERPMXVCNTR, Cluster Selected Event Counter Register</a> on page 171       |
| Updated Cluster Monitor Debug Configuration Register (CLUSTERPMMDCR) description.    | <a href="#">13.15 CLUSTERPMMDCR, Cluster Monitor Debug Configuration Register</a> on page 173    |
| Updated Cluster Common Event Identification Register 0 (CLUSTERPMCEID0) description. | <a href="#">13.16 CLUSTERPMCEID0, Cluster Common Event Identification Register 0</a> on page 175 |
| Updated Cluster Common Event Identification Register 1 (CLUSTERPMCEID1) description. | <a href="#">13.17 CLUSTERPMCEID1, Cluster Common Event Identification Register 1</a> on page 177 |
| Added Cluster Claim Tag Set Register (CLUSTERCLAIMSET) description.                  | <a href="#">13.18 CLUSTERCLAIMSET, Cluster Claim Tag Set Register</a> on page 179                |
| Added Cluster Claim Tag Clear Register (CLUSTERCLAIMCLR) description.                | <a href="#">13.19 CLUSTERCLAIMCLR, Cluster Claim Tag Clear Register</a> on page 181              |
| Updated Cluster Event Type Register (CLUSTERPMEVTYPER<n>) description.               | <a href="#">13.20 CLUSTERPMEVTYPER&lt;n&gt;, Cluster Event Type Register</a> on page 183         |
| Added Cluster Event Counter Register (CLUSTERPMEVCNTR<n>) description.               | <a href="#">13.21 CLUSTERPMEVCNTR&lt;n&gt;, Cluster Event Counter Register</a> on page 184       |
| Added signal MPMEN   | <a href="#">B.2.6 Power management signals</a> on page 225                                       |

**Table A-78: Differences between issue 0200-00 and issue 0300-00**

| Change   | Location  |
|--|---|
| Updated Implementation options table.  | <a href="#">2.3 Implementation options</a> on page 18       |
| Updated Product revisions.   | <a href="#">2.7 Product revisions</a> on page 24            |
| Updated L3 cache description to qualify which cache sizes are 12-way set associative and which are 16-way set associative. | <a href="#">3.1 Components</a> on page 25                   |
| Updated DSU interfaces table to include two CHI interfaces.  | <a href="#">3.2 Interfaces</a> on page 28                   |
| Qualified resets as either Cold reset, or Warm reset.  | Throughout document   |
| Updated description of Off power mode.   | <a href="#">5.1 About DSU power management</a> on page 37   |
| Updates description of Power mode control.   | <a href="#">5.2 Power mode control</a> on page 37           |
| Updated description of L3 cache partial powerdown.   | <a href="#">5.4.1 L3 cache partial powerdown</a> on page 39 |
| Updated the description of SFONLY FUNC_RET, ¼FUNC_RET, ½FUNC_RET, ¾FUNC_RET, FULL FUNC_RET.                                | <a href="#">5.5.1 Power mode transitions</a> on page 41     |
| Added note about the SFONLY, ¼, ½, ¾ and FULL DEBUG_RECOV modes.   | <a href="#">5.5.1 Power mode transitions</a> on page 41     |
| Added Victim RAM and LTDB RAMs.  | <a href="#">5.5.4 Power mode encoding</a> on page 47        |

| Change   | Location  |
|--|---|
| Added note about CLUSTERACTIVE and COREACTIVE outputs.                       | <a href="#">5.6 Power operating requirements</a> on page 48                         |
| Added section about power control for DFT.                                   | <a href="#">5.6.1 Power control for DFT</a> on page 49                              |
| Described the signal naming convention for dual CHI.                         | <a href="#">5.9.1.1 Coherency signals naming convention</a> on page 53              |
| Qualified how AWSTASHLPIDS signal indicates core and thread.                 | <a href="#">6.4 Cache stashing</a> on page 56                                       |
| Updated description of L3 cache data RAM latency.                            | <a href="#">6.6 L3 cache data RAM latency</a> on page 60                            |
| Updated L3 cache slices section to include cache portions description.       | <a href="#">6.7 Cache slices and portions</a> on page 62                            |
| Added combined issuing capability to ACE master interface attributes.        | <a href="#">7.4 ACE master interface attributes</a> on page 68                      |
| Updated note.  | <a href="#">7.9 Write response</a> on page 76                                       |
| Added dual CHI interface description.  | <a href="#">8.1.1 Dual CHI interfaces</a> on page 78                                |
| Updated CHI channel properties to account for dual CHI.                      | <a href="#">8.6 CHI channel properties</a> on page 82                               |
| Updated CLUSTERCFR register description.                                     | <a href="#">11.7 CLUSTERCFR, Cluster Configuration Register</a> on page 101         |
| Updated CLUSTERECTLR register description.                                   | <a href="#">11.8 CLUSTERECTLR, Cluster Extended Control Register</a> on page 105    |
| Updated CLUSTERIDR register description.                                     | <a href="#">11.9 CLUSTERIDR, Cluster Main Revision ID Register</a> on page 109      |
| Updated CLUSTERPWRCTLR register description.                                 | <a href="#">11.13 CLUSTERPWRCTLR, Cluster Power Control Register</a> on page 119    |
| Updated PMU events table description to align with architecture description. | <a href="#">15.3 PMU events - DSU R3 and R3 ADP</a> on page 193                     |
| Updated CTIPIDR2 register description.                                       | <a href="#">16.5 CTIPIDR2, CTI Peripheral Identification Register 2</a> on page 208 |
| Added MPMSTATEx[1:0] signal description.                                     | <a href="#">B.2.6 Power management signals</a> on page 225                          |
| Updated CHI interface signals for the dual CHI interface.                    | <a href="#">B.2.12 CHI interface signals</a> on page 235                            |

**Table A-79: Differences between issue 0300-00 and issue 0300-01**

| Change  | Location   |
|---|--|
| Qualified the use of L3_DATA_WR_LATENCY value of 2p.  | <a href="#">2.3 Implementation options</a> on page 18  |
| Clarified how CTIs are allocated.   | <a href="#">3.1 Components</a> on page 25  |
| Clarified the Cluster to DebugBlock interface.  | <a href="#">3.2 Interfaces</a> on page 28  |
| Clarified the Direct connect connection.  | <a href="#">3.5 L3 memory system variants</a> on page 32   |
| Updated ACE transactions.   | <a href="#">7.6 ACE transactions</a> on page 73  |
| Updated the description for the WAY, [31:28] bit field in the ERR1MISC0 register description. | <a href="#">12.5 ERR1MISC0, Error Record Miscellaneous Register 0</a> on page 138  |
| Updated DebugBlock diagram and description to clarify the Cluster CTI.                        | <a href="#">14.4 DebugBlock components</a> on page 187, <a href="#">14.5 About the Embedded Cross Trigger</a> on page 188, <a href="#">14.5.1 Supported debug and trace trigger events</a> on page 190 |
| Added Compatible Core Versions Appendix.  | <a href="#">A. Compatible Core Versions</a> on page 220  |

**Table A-80: Differences between issue 0400-00 and issue 0300-01**

| Change  | Location                                 |
|---|--|
| Fixed various typographical errors.                     | The whole document.                      |
| Updated Features section to include Arm®v8.4-A support. | <a href="#">2.2 Features</a> on page 17. |

| Change   | Location   |
|--|--|
| Updated Implementation options section to include new L3 cache sizes of 256KB and 3MB, and to document new <code>L3_PBHA</code> parameter. | <a href="#">2.3 Implementation options</a> on page 18  |
| Added Arm®v8.4-A core and debug support. Added GIC v4.0 support.   | <a href="#">2.4 Supported standards and specifications</a> on page 22  |
| Updated product revisions for r4p0 to include support for CHI.D, 256KB L3 cache size, and PBHA.  | <a href="#">2.7 Product revisions</a> on page 24   |
| Updated Components to include new cache sizes.   | <a href="#">3.1 Components</a> on page 25  |
| Added Page Based Hardware Attribute description.   | <a href="#">3.4 Page-based hardware attributes</a> on page 31  |
| Updated description of Direct connect.   | <a href="#">3.5 L3 memory system variants</a> on page 32   |
| Updated CLUSTERPSTATE value  | <a href="#">5.5.4 Power mode encoding</a> on page 47   |
| Updated memory retention description to explain Snoop Filter RAMs cannot be placed into retention.   | <a href="#">5.4.2 L3 RAM retention</a> on page 40  |
| Updated section for 256KB and 3MB cache size.  | <a href="#">6.7.3 Implementing a 1.5MB or 3MB L3 cache</a> on page 64  |
| Added new note to qualify issuing capability for the cluster.  | <a href="#">7.4 ACE master interface attributes</a> on page 68 and <a href="#">8.6 CHI channel properties</a> on page 82 |
| Added new section describing support for memory types.   | <a href="#">7.7 Support for memory types</a> on page 75  |
| Updated description to explain the additional signals to tie LOW for AXI compatibility mode.   | <a href="#">7.11 AXI compatibility mode</a> on page 76   |
| Added new section.   | <a href="#">8.8 Use of DataSource</a> on page 86   |
| Added new section.   | <a href="#">7.7 Support for memory types</a> on page 75  |
| Added note about WriteUniquePtlStash.  | <a href="#">9.4 ACP transaction types</a> on page 89   |
| Updated section.   | <a href="#">9.5 ACP performance</a> on page 91   |
| Updated Variant bit-field in CLUSTERIDR register description.  | <a href="#">11.9 CLUSTERIDR, Cluster Main Revision ID Register</a> on page 109   |
| Updated Revision bit-field in CTIPIDR2 register description.   | <a href="#">16.5 CTIPIDR2, CTI Peripheral Identification Register 2</a> on page 208                                      |
| Added new signals PPROTDBG[2:0], PPROTDC[1:1]  | <a href="#">B.3.5 Debug signals</a> on page 244  |
| Added note about NIDEN and SPNIDEN signals   | <a href="#">B.3.5 Debug signals</a> on page 244  |

**Table A-81: Differences between issue 0400-00 and issue 0400-01**

| Change                          | Location |
|---------------------------------|----------|
| There are no technical changes. | -        |

**Table A-82: Differences between issue 0400-01 and 0400-02**

| Change                                | Location  |
|---------------------------------------|---|
| Fixed various typographical errors.   | The whole document.   |
| Changed ACE-Lite to ACE5-Lite.        | <a href="#">3.2 Interfaces</a> on page 28                     |
| Corrected two misspelt PBHA acronyms. | <a href="#">3.4 Page-based hardware attributes</a> on page 31 |

| Change  | Location   |
|---|--|
| Corrected the first 4 rows of the 'L3 write bandwidth' column of the 'L3 data RAM performance' table.                                   | <a href="#">6.6 L3 cache data RAM latency</a> on page 60   |
| Changed ACE to ACE5.  | <a href="#">7.1 About the ACE master interface</a> on page 66  |
| Amended the 'CHI transaction usage' table.  | <a href="#">8.7 CHI transactions</a> on page 83  |
| Changed ACE-Lite to ACE5-Lite.  | <a href="#">9.1 About the ACP</a> on page 88   |
| Changed ACE-Lite to ACE5-Lite.  | <a href="#">9.3 ACP ACE5-Lite subset</a> on page 89  |
| Better explained the Write-Allocate bit.  | <a href="#">9.5 ACP performance</a> on page 91   |
| Added read/write access examples for both the AArch64 and AArch32 states to the 'Usage constraints' section of these Control registers. | <a href="#">11.4 CLUSTERACPSID</a> , Cluster ACP Scheme ID Register on page 97, <a href="#">11.6 CLUSTERBUSQOS</a> , Cluster Bus QoS Control Register on page 99, <a href="#">11.7 CLUSTERCFR</a> , Cluster Configuration Register on page 101, <a href="#">11.8 CLUSTERECTLR</a> , Cluster Extended Control Register on page 105, <a href="#">11.9 CLUSTERIDR</a> , Cluster Main Revision ID Register on page 109, <a href="#">11.10 CLUSTERL3HIT</a> , Cluster L3 Hit Counter Register on page 111, <a href="#">11.11 CLUSTERL3MISS</a> , Cluster L3 Miss Counter Register on page 113, <a href="#">11.12 CLUSTERPARTCR</a> , Cluster Partition Control Register on page 114, <a href="#">11.13 CLUSTERPWRCtrlR</a> , Cluster Power Control Register on page 119, <a href="#">11.14 CLUSTERPWRDN</a> , Cluster Powerdown Register on page 121, <a href="#">11.15 CLUSTERPWRSTAT</a> , Cluster Power Status Register on page 123, <a href="#">11.16 CLUSTERREVIDR</a> , Cluster Revision ID Register on page 125, <a href="#">11.17 CLUSTERSTASHID</a> , Cluster Stash Scheme ID Register on page 127, <a href="#">11.18 CLUSTERTHREADSID</a> , Cluster Thread Scheme ID Register on page 129, and <a href="#">11.19 CLUSTERTHREADSIDOVR</a> , Cluster Thread Scheme ID Override Register on page 130   |
| Updated the <b>WAY</b> , [31:28] description to include the LTDB RAM.   | <a href="#">12.5 ERR1MISC0</a> , Error Record Miscellaneous Register 0 on page 138   |
| Added read/write access examples for both the AArch64 and AArch32 states to the 'Usage constraints' section of these PMU registers.     | <a href="#">13.4 CLUSTERPMCR</a> , Cluster Performance Monitors Control Register on page 151, <a href="#">13.5 CLUSTERPMCNTENSET</a> , Cluster Count Enable Set Register on page 154, <a href="#">13.6 CLUSTERPMCNTENCLR</a> , Cluster Count Enable Clear Register on page 156, <a href="#">13.7 CLUSTERPMOVSET</a> , Cluster Overflow Flag Status Set Register on page 158, <a href="#">13.8 CLUSTERPMOVCLR</a> , Cluster Overflow Flag Status Clear Register on page 160, <a href="#">13.9 CLUSTERPMSELR</a> , Cluster Event Counter Selection Register on page 162, <a href="#">13.10 CLUSTERPMINTENSET</a> , Cluster Interrupt Enable Set Register on page 164, <a href="#">13.11 CLUSTERPMINTENCLR</a> , Cluster Interrupt Enable Clear Register on page 166, <a href="#">13.12 CLUSTERPMCCNTR</a> , Cluster Performance Monitors Cycle Counter on page 168, <a href="#">13.13 CLUSTERPMXEVTYP</a> , Cluster Selected Event Type Register on page 169, <a href="#">13.14 CLUSTERPMXEVCNTR</a> , Cluster Selected Event Counter Register on page 171, <a href="#">13.15 CLUSTERPMMDCR</a> , Cluster Monitor Debug Configuration Register on page 173, <a href="#">13.16 CLUSTERPMCEID0</a> , Cluster Common Event Identification Register 0 on page 175, <a href="#">13.17 CLUSTERPMCEID1</a> , Cluster Common Event Identification Register 1 on page 177, <a href="#">13.18 CLUSTERCLAIMSET</a> , Cluster Claim Tag Set Register on page 179, and <a href="#">13.19 CLUSTERCLAIMCLR</a> , Cluster Claim Tag Clear Register on page 181 |

| Change  | Location   |
|---|--|
| Amended the 'ACE SRCATTRM0/1 bus signal functionality' table. | <a href="#">B.2.11 ACE interface signals</a> on page 231 |
| Amended the 'CHI REQSRCATTR bus signal functionality' table.  | <a href="#">B.2.12 CHI interface signals</a> on page 235 |

**Table A-83: Differences between issue 0400-02 and 0401-03**

| Change   | Location   |
|--|--|
| Summarized the r4p1 main technical changes.  | <a href="#">2.7 Product revisions</a> on page 24   |
| Updated the revised revision-related values.   | <a href="#">11.9 CLUSTERIDR, Cluster Main Revision ID Register</a> on page 109, <a href="#">16.5 CTIPIDR2, CTI Peripheral Identification Register 2</a> on page 208, and <a href="#">17.2 ROM table register summary</a> on page 216   |
| Updated the ACE and CHI signals due to the Direct connect changes.   | <a href="#">B.2.11 ACE interface signals</a> on page 231 and <a href="#">B.2.12 CHI interface signals</a> on page 235  |
| Updated the descriptions of the PMU events, from 0x0500 to 0x0504, to specify the L3 prefetch.                   | <a href="#">15.3 PMU events - DSU R3 and R3 ADP</a> on page 193  |
| Removed all references to the L3 data cache and replaced them the more accurate designation of L3 unified cache. | <a href="#">5.7 Wait For Interrupt and Wait For Event</a> on page 50 and <a href="#">13.17 CLUSTERPMCEID1, Cluster Common Event Identification Register 1</a> on page 177  |
| Added more specific connection information for interrupt signals.  | <a href="#">B.2.4 GIC signals</a> on page 224, <a href="#">B.2.5 Generic Timer signals</a> on page 225, <a href="#">B.2.7 Error signals</a> on page 227, <a href="#">B.2.13 DebugBlock APB interface signals</a> on page 239, <a href="#">B.2.16 PMU interface signals</a> on page 241, and <a href="#">B.3.6 CTI interface signals</a> on page 246  |
| Correctly specified the PPMCTLx signal.  | <a href="#">B.2.6 Power management signals</a> on page 225   |
| Describe how an interpolator is included on the TSVALUEB input signal.   | <a href="#">B.2.15 Timestamp signal</a> on page 241  |
| Amended the L3 cache section.  | <a href="#">6.1 About the L3 cache</a> on page 54, <a href="#">6.3 L3 cache partitioning</a> on page 54, <a href="#">6.4 Cache stashing</a> on page 56, <a href="#">6.5 L3 cache ECC and parity</a> on page 57, <a href="#">6.6 L3 cache data RAM latency</a> on page 60, <a href="#">6.7 Cache slices and portions</a> on page 62, <a href="#">6.7.1 Cache slice and master port selection</a> on page 63, <a href="#">6.7.2 Default number of cache slices</a> on page 63, and <a href="#">6.7.3 Implementing a 1.5MB or 3MB L3 cache</a> on page 64 |
| Ensured the power modes use a consistent use of the same case.   | <a href="#">5.5.1 Power mode transitions</a> on page 41  |
| Added the POP_RAM configuration parameter, which was missing from the table.                                     | <a href="#">2.3 Implementation options</a> on page 18  |

| Change  | Location  |
|---|---|
| Amended the DebugBlock power states description to describe a workaround due to an erratum. | <a href="#">14.3 About the DebugBlock</a> on page 186 |

**Table A-84: Differences between Issue 0401-03 and Issue 0401-04**

| Change  | Location  |
|---|---|
| Added a note to clarify what the <code>OTHER_CORE_TYPE</code> is.   | <a href="#">2.3 Implementation options</a> on page 18   |
| Updated master interface to main memory. The DSU AE can be configured with either 1 or 2 ACE, or 1 or 2 CHI interfaces. | <a href="#">3.2 Interfaces</a> on page 28   |
| Updated the nMBISTRESET signal description.   | <a href="#">4.2 Resets</a> on page 35   |
| Updated L3 power domain CLUSTERACTIVE 0 bit description.  | <a href="#">5.6 Power operating requirements</a> on page 48                                       |
| Added text about stashing to an L3 when no L3 present.  | <a href="#">6.4 Cache stashing</a> on page 56   |
| Added new section about injectable errors.  | <a href="#">6.5 L3 cache ECC and parity</a> on page 57  |
| Reworked paragraph about peripheral port address range.   | <a href="#">10.1 About the peripheral port</a> on page 93   |
| Updated CLUSTERTHREADSIDOVR, Cluster Thread Scheme ID Override Register is banked for each thread of execution.         | <a href="#">11.19 CLUSTERTHREADSIDOVR, Cluster Thread Scheme ID Override Register</a> on page 130 |