

RealView™ Compilation Tools for BREW

Version 1.2

Linker and Utilities Guide



RealView Compilation Tools for BREW

Linker and Utilities Guide

Copyright © 2002 ARM Limited. All rights reserved.

Release Information

Change History		
Date	Issue	Change
June 2002	A	Release 1.2

Proprietary Notice

Words and logos marked with ® or ™ are registered trademarks or trademarks owned by ARM Limited. Other brands and names mentioned herein may be the trademarks of their respective owners.

Neither the whole nor any part of the information contained in, or the product described in, this document may be adapted or reproduced in any material form except with the prior written permission of the copyright holder.

The product described in this document is subject to continuous developments and improvements. All particulars of the product and its use contained in this document are given by ARM in good faith. However, all warranties implied or expressed, including but not limited to implied warranties of merchantability, or fitness for purpose, are excluded.

This document is intended only to assist the reader in the use of the product. ARM Limited shall not be liable for any loss or damage arising from the use of any information in this document, or any error or omission in such information, or any incorrect use of the product.

Contents

RealView Compilation Tools for BREW Linker and Utilities Guide

	Preface	
	About this book	vi
	Feedback	ix
Chapter 1	Introduction	
	1.1 About the linker and utilities	1-2
Chapter 2	The armlink Command Syntax	
	2.1 About armlink	2-2
	2.2 armlink command syntax	2-8
Chapter 3	Using the Basic Linker Functionality	
	3.1 Specifying the image structure	3-2
	3.2 Section placement	3-8
	3.3 Optimizations and modifications	3-11
	3.4 Using command-line options to create simple images	3-14
Chapter 4	Accessing Image Symbols	
	4.1 Accessing linker-defined symbols	4-2
	4.2 Accessing symbols in another image	4-5

4.3	Hiding and renaming global symbols	4-9
4.4	Using \$Super\$\$ and \$Sub\$\$ to override symbol definitions	4-12

Chapter 5

Using libraries

5.1	About libraries	5-2
5.2	Library searching, selection, and scanning	5-3

Chapter 6

Using fromELF

6.1	About fromELF	6-2
6.2	fromELF command-line options	6-3
6.3	Examples of fromELF usage	6-9

Glossary

Preface

This preface introduces the *RealView Compilation Tools for BREW* (RVCT for BREW) linker and utilities documentation. It contains the following sections:

- *About this book* on page vi
- *Feedback* on page ix.

About this book

This book provides reference information for RealView Compilation Tools for BREW (RVCT for BREW) v1.2. It describes the command-line options to the linker and other ARM tools in RVCT for BREW.

Intended audience

This book is written for all developers who are producing applications using RVCT for BREW. It assumes that you are an experienced software developer.

Using this book

This book is organized into the following chapters and appendixes:

Chapter 1 *Introduction*

Read this chapter for an introduction to the linker and related utilities in RVCT for BREW v1.2.

Chapter 2 *The armlink Command Syntax*

Read this chapter for an explanation of all command-line options accepted by the linker.

Chapter 3 *Using the Basic Linker Functionality*

Read this chapter for details on creating simple images.

Chapter 4 *Accessing Image Symbols*

Read this chapter for details on accessing symbols in images.

Chapter 5 *Using libraries*

Read this chapter for an explanation of the procedures involved in creating and accessing library objects.

Chapter 6 *Using fromELF*

Read this chapter for a description of the fromELF utility program and how you can use it to change image format.

Typographical conventions

The following typographical conventions are used in this book:

monospace	Denotes text that can be entered at the keyboard, such as commands, file and program names, and source code.
-----------	--

<u>monospace</u>	Denotes a permitted abbreviation for a command or option. The underlined text can be entered instead of the full command or option name.
<i>monospace italic</i>	Denotes arguments to commands and functions where the argument is to be replaced by a specific value.
monospace bold	Denotes language keywords when used outside example code.
<i>italic</i>	Highlights important notes, introduces special terminology, denotes internal cross-references, and citations.
bold	Highlights interface elements, such as menu names. Also used for emphasis in descriptive lists, where appropriate, and for ARM processor signal names.

Further reading

This section lists publications from both ARM Limited and third parties that provide additional information on developing code for the ARM family of processors.

ARM periodically provides updates and corrections to its documentation. See <http://www.arm.com> for current errata sheets and addenda, and the ARM Frequently Asked Questions.

ARM publications

This book contains reference information that is specific to development tools supplied with RVCT for BREW. Other publications included in the suite are:

- *RealView Compilation Tools for BREW v1.2 Installation Guide* (ARM DSI 0024)
- *RealView Compilation Tools for BREW v1.2 Assembler Guide* (ARM DUI 0170)
- *RealView Compilation Tools for BREW v1.2 Compilers and Libraries Guide* (ARM DUI 0171).

The *RealView Compilation Tools for BREW v1.2 Assembler Guide* and *RealView Compilation Tools for BREW v1.2 Compilers and Libraries Guide* are supplied in PDF format in `install_directory\PDF`.

For license management information, see the *ARM FLEXlm License Management Guide* (ARM DUI 0209). This is supplied in PDF format in `install_directory\PDF\ARM_FLEXlmGuide.pdf`.

In addition, refer to the following documentation for specific information relating to ARM products:

- *ARM ELF specification* (SWS ESPC 0003)
- *TIS DWARF 2 specification*
- *ARM-Thumb Procedure Call Standard specification* (SWS ESPC 0002)
- *ARM Architecture Reference Manual* (ARM DDI 0100)
- *ARM Reference Peripheral Specification* (ARM DDI 0062)
- the ARM datasheet or technical reference manual for your hardware device.

Additional code examples and information about writing code can be found in the *ARM Firmware Suite* (AFS) example code and documentation. Refer to the following ARM publications for more details on AFS:

- *ARM Firmware Suite User Guide* (ARM DUI 0136)
- *ARM Firmware Suite Reference Guide* (ARM DUI 0102).

Other publications

This book is not intended to be an introduction to the ARM assembly language, C, or C++ programming languages. Other books provide general information about programming.

The following book gives general information about the ARM architecture:

- *ARM System-on-chip Architecture* (second edition), Furber, S., (2000). Addison Wesley. ISBN 0-201-67519-6.

Feedback

ARM Limited welcomes feedback on both RealView Compilation Tools for BREW and the documentation.

Feedback on RealView Compilation Tools for BREW

If you have any problems with RealView Compilation Tools for BREW, contact your supplier. To help them provide a rapid and useful response, give:

- your name and company
- the serial number of the product
- details of the release you are using
- details of the platform you are running on, such as the hardware platform, operating system type and version
- a small standalone sample of code that reproduces the problem
- a clear explanation of what you expected to happen, and what actually happened
- the commands you used, including any command-line options
- sample output illustrating the problem
- the version string of the tools, including the version number and build numbers.

Feedback on this book

If you have any problems with this book, send an email to errata@arm.com giving:

- the document title
- the document number
- the page number(s) to which your comments apply
- a concise explanation of the problem.

General suggestions for additions and improvements are also welcome.

Chapter 1

Introduction

This chapter introduces the ARM linker and utility programs provided with RVCT for BREW. It contains the following sections:

- *About the linker and utilities* on page 1-2
- *The linker* on page 1-2
- *The utilities* on page 1-3.

1.1 About the linker and utilities

RVCT for BREW consists of a suite of tools, together with supporting documentation, that enable you to write applications for the ARM family of RISC processors. You can use RVCT for BREW to build C, C++, and ARM assembly language programs.

The RVCT for BREW toolkit consists of the following major components:

- command-line development tools
- utilities
- supporting software.

This book describes the ARM linker, `armlink`, and the command-line utility tools provided with RVCT for BREW. See *ARM publications* on page vii for a list of the other books in the RVCT for BREW documentation suite that give information on the ARM assemblers, compilers, and supporting software.

1.1.1 The linker

`armlink` combines the contents of one or more object files with selected parts of one or more object libraries to produce an ELF executable image, or a partially linked ELF object.

The linker can link ARM code and Thumb® code, and automatically generates interworking veneers to switch processor state when required. The linker also automatically generates long branch veneers, where required, to extend the range of branch instructions.

The linker supports command-line options that enable you to specify separate locations for code and data within the system memory map.

The linker can perform common section elimination and unused section elimination to reduce the size of your output image. In addition, the linker enables you to:

- produce debug and reference information about linked files
- generate a static callgraph and list the stack usage over it
- control the contents of the symbol table in output images.

The linker automatically selects the appropriate standard C or C++ library variants to link with, based on the build attributes of the objects it is linking.

The linker does not generate output formats other than ELF. To convert ELF images to other format, such as plain binary for loading into ROM, use the `fromELF` utility. See *The utilities* on page 1-3.

See Chapter 2 *The armlink Command Syntax* for detailed information on the ARM linker.

1.1.2 The utilities

This section gives an overview of the fromELF utility tool that is provided to support the main development tools.

fromELF

fromELF is the ARM image conversion utility. It accepts ELF format input files and converts them to a variety of output formats, including:

- plain binary
- Motorola 32-bit S-record format
- Intel Hex-32 format
- Byte Oriented (Verilog Memory Model) Hex format
- resave as ELF format.

The utility can also produce textual information about the input file, and disassemble code. See Chapter 6 *Using fromELF* for detailed information.

Chapter 2

The armlink Command Syntax

This chapter describes the full command syntax for armlink. This chapter contains the following sections:

- *About armlink* on page 2-2
- *armlink command syntax* on page 2-8.

2.1 About armlink

armlink, the RealView Compilation Tools for BREW linker, enables you to:

- link a collection of objects and libraries (in either ARM or Thumb code) into an executable image
- partially link a collection of objects into an object that can be used as input to a future link step
- specify where the code and data are to be located in memory
- produce debug and reference information about the linked files.

Objects consist of input sections that contain code, initialized data, or the locations of memory that must be set to zero. Input sections can be *Read-Only* (RO), *Read-Write* (RW), or *Zero-Initialized* (ZI). These attributes are used by armlink to group input sections into bigger building blocks called output sections, regions, and images. Load regions are equivalent to ELF segments.

Load regions typically exist in the system memory map at reset or after the image is loaded into the target by a debugger. Before you can execute the image, you might have to move some of its regions to their execution addresses. The memory map of an image therefore has the following distinct views:

- the load view of the memory when the program and data are first loaded
- the execution view of the memory after code is moved to its normal execution location.

The term *root region* is used to describe regions that execute at their load address.

See *Specifying the image structure* on page 3-2 for more information on the image hierarchy.

2.1.1 Input to armlink

Input to armlink consists of:

- One or more object files in ELF Object Format. This format is described in the ARM ELF specification. See *ARM publications* on page vii for more information.
- One or more libraries as described in Chapter 5 *Using libraries*.
- A symbol definitions file.

Note

For backward compatibility, armlink also accepts object files in AOF format and libraries in ALF format. These formats are obsolete and will not be supported in the future.

2.1.2 Output from armlink

Output from a successful invocation of armlink is one of the following:

- an executable image in ELF executable format
- a partially-linked object in ELF object format.

For simple images, ELF executable files contain segments that are approximately equivalent to RO and RW output sections in the image. An ELF executable file also has ELF sections that contain the image output sections.

You can use the fromELF utility to convert an executable image in ELF executable format to other file formats. See Chapter 6 *Using fromELF* for more information.

Constructing an executable image

When you use armlink to construct an executable image, it:

- resolves symbolic references between the input object files
- extracts object modules from libraries to satisfy otherwise unsatisfied symbolic references
- sorts input sections according to their attributes and names, and merges similarly attributed and named sections into contiguous chunks
- eliminates duplicate copies of debug sections
- relocates relocatable values
- generates an executable image.

Constructing a partially-linked object

When you use armlink to construct a partially-linked object, it:

- eliminates duplicate copies of debug sections
- minimizes the size of the symbol table
- leaves unresolved references unresolved
- generates an object that can be used as an input to a subsequent link step.

2.1.3 Summary of armlink options

This section gives a brief overview of each armlink command-line option. The options are arranged into functional groups.

Accessing help and information

To get information on the available command-line options use:

`-help`

To get the tool version number use:

`-vsn`

Specifying the output type and the output file name

Use the following option to create a partially-linked object instead of an executable image:

`-partial`

Name the output file using the following option:

`-output`

Specify the output file format using the following option:

`-elf`

ELF is the default, and only, output format. Use `fromELF` to convert ELF output to other formats. See Chapter 6 *Using fromELF* for more information.

Using a via file

Use the following option to specify a via file containing additional command-line arguments to the linker:

`-via`

See the section on via files in *RealView Compilation Tools for BREW Compilers and Libraries Guide* for more information.

Specifying memory map information for the image

Use the following options to specify simple memory maps:

- reloc
- ro-base
- rw-base
- ropi
- rwp

The simple memory map options cannot be used for partial linking because they specify the memory map of an executable image.

Controlling image contents

These options control various miscellaneous factors affecting the image contents:

- debug | -nodebug
- entry
- first
- keep
- last
- libpath
- edit
- locals | -nolocal
- remove | -noremove
- scanlib | -noscanlib
- match

Generating image-related information

These options control how you extract and present information about the image:

- callgraph
- info
- map
- symbols
- symdefs
- xref
- xreffrom
- xrefto

With the exception of -callgraph, by default the linker prints the information you request on the standard output stream, stdout. You can redirect the information to a text file using the -list command-line option.

For -callgraph, the information is saved as an HTML file. An HTML file named *output_name.htm* is saved in the same directory as the generated image.

Controlling armlink diagnostics

These options control how armlink emits diagnostics:

- errors
- list
- verbose
- strict
- unresolved
- mangled
- unmangled

2.2 armlink command syntax

The complete linker command syntax is:

```
armlink [-hhelp] [-vsn] [-partial] [-output file] [-elf] [-reloc]
[[-ro-base address]] [-ropi] [-rw-base address] [-rwp] [-split]
[[-debug]|[-nodebug]] [[-remove (RO/RW/ZI/DBG)]|[-noremove]] [-entry location]
[-keep section-id] [-first section-id] [-last section-id] [-libpath pathlist]
[[-scanlib]|[-noscanlib]] [[-locals]|[-nolocals]] [-callgraph] [[-info topics]]
[-map] [-symbols] [[-symdefs file]] [[-edit file]] [-xref]
[[-xreffrom object(section)]] [[-xrefto object(section)]] [[-errors file]]
[[-list file]] [-verbose] [[-unmangled]|[-mangled]] [-match crossmangled]
[[-via file]] [-strict] [[-unresolved symbol]] [-MI|-LI|-BI] [input-file-list]
```

where:

-hhelp	This option prints a summary of some commonly used command-line options.				
-vsn	This option displays the armlink version information.				
-partial	This option creates a partially-linked object instead of an executable image.				
-output <i>file</i>	<p>This option specifies the name of the output file. The file can be either a partially-linked object or executable image. If the output file name is not specified, armlink uses the following defaults:</p> <table> <tr> <td>__image.axf</td><td>if the output is an executable image</td></tr> <tr> <td>__object.o</td><td>if the output is a partially-linked object.</td></tr> </table> <p>If <i>file</i> is specified without path information, it is created in the current working directory. If path information is specified, then that directory becomes the default output directory.</p>	__image.axf	if the output is an executable image	__object.o	if the output is a partially-linked object.
__image.axf	if the output is an executable image				
__object.o	if the output is a partially-linked object.				
-elf	This option generates the image in ELF format. This is the only output format supported by armlink. This is the default.				
-reloc	<p>This option makes relocatable images.</p> <p>A relocatable image has a dynamic segment that contains relocations that can be used to relocate the image post link-time. Examples of post link-time relocation include advanced ROM construction and dynamic loading at run time.</p> <p>If the image is loaded at its link-time address, the relocatable image produced by armlink does not require the relocations to be processed and debug data for the image is valid. Loading the</p>				

image at a different address to the link-time address and processing the relocations, however, invalidates any debug data present in the image.

Used on its own, `-reloc` makes an image similar to Simple type 1, but the single load region has the RELOC attribute. See *Type 1: one load region and contiguous output regions* on page 3-14 for more information on type 1 images.

The combination `-reloc -split` makes an image similar to Simple type 3, but the two load regions now have the RELOC attribute. See *Type 3: two load regions and noncontiguous output regions* on page 3-17 for more information on type 3 images.

Using `-reloc -rw-base` without also using `-split` causes an error.

- | | |
|-------------------------------|--|
| <code>-ro-base address</code> | This option sets both the load and execution addresses of the region containing the RO output section at <i>address</i> . <i>address</i> must be word-aligned. If this option is not specified, the default RO base address is <code>0x8000</code> . |
| <code>-ropi</code> | <p>This option makes the load and execution region containing the RO output section position-independent. If this option is not used the region is marked as absolute. Usually each read-only input section must be read-only position-independent. If this option is selected, armlink:</p> <ul style="list-style-type: none"> • checks that relocations between sections are valid • ensures that any code generated by armlink itself, such as interworking veneers, is read-only position-independent. |

———— **Note** ————

The ARM tools cannot determine if the final output image will be read-only position-independent until armlink finishes processing input sections. This means that armlink might emit ROPI error messages, even though you have selected the ROPI option for the compiler and assembler.

- | | |
|-------------------------------|---|
| <code>-rw-base address</code> | <p>This option sets the execution addresses of the region containing the RW output section at <i>address</i>. <i>address</i> must be word-aligned.</p> <p>If this option is used with <code>-split</code>, it sets both the load and the execution address of the region containing the RW output sections at <i>address</i>.</p> |
|-------------------------------|---|

-rwpi This option makes the load and execution region containing the RW and ZI output section position-independent. If this option is not used the region is marked as absolute. This option requires a value for **-rw-base**. If **-rw-base** is not specified, **-rw-base 0** is assumed. Usually each writable input section must be read-write position-independent.

If this option is selected, armlink:

- checks that the PI attribute is set on input sections to any read-write execution regions
- checks that relocations between sections are valid
- generates static base-relative entries in the tables that are used when RO and RW regions are copied or initialized.

Note

The compiler does not force your writable data to be position-independent. This means that armlink might diagnose that data is not RWPI, even though you have selected the RWPI option for the compiler and assembler.

-split This option splits the default load region, that contains the RO and RW output sections, into the following load regions:

- One containing the RO output section. The default load address is **0x8000**, but a different address can be specified with the **-ro-base** option.
- One containing the RW output section. The load address is specified with the **-rw-base** option. This option requires a value for **-rw-base**. If **-rw-base** is not specified, **-rw-base 0** is assumed.

-debug This option includes debug information in the output file. The debug information includes debug input sections and the symbol and string table. This is the default.

-nodebug This option turns off the inclusion of debug information in the output file. The image is smaller, but you cannot debug it at the source level. armlink discards any debug input section it finds in the input objects and library members, and does not include the symbol and string table in the image as loaded into the debugger. This only affects the image size as loaded into the debugger and has no effect on the size of any resulting binary image that is downloaded to the target.

If you are creating a partially-linked object rather than an image, armlink discards the debug input sections it finds in the input objects, but does produce the symbol and string table in the partially-linked object.

Note

Do not use `-nodebug` if a `fromELF` step is required. If your image is produced without debug information:

- `fromELF` cannot translate the image into other file formats
 - `fromELF` cannot produce a meaningful disassembly listing.
-

`-remove (RO/RW/ZI/DBG)`

This option performs unused section elimination on the input sections to remove unused sections from the image. An input section is considered to be used if it contains the image entry point, or if it is referred to from a used section. See also *Unused section elimination* on page 3-11.

Note

You must take care not to remove exception handlers when using `-remove`. Use the `-keep` option to identify exception handlers or use the `ENTRY` directive to label them as entry points.

You can use section attribute qualifiers for more precise control of the unused section elimination process. If a qualifier is used, it can be one or more of the following:

- | | |
|-----|---|
| RO | Remove all unused sections of type RO. |
| RW | Remove all unused sections of type RW. |
| ZI | Remove all unused sections of type ZI. |
| DBG | Remove all unused sections of type DEBUG. |

The qualifiers can appear in any case and order, but must be enclosed in parentheses (), and must be separated by a slash /.

The default is `-remove (RO/RW/ZI/DBG)`.

If no section attribute qualifiers are specified, all unused sections are eliminated. `-remove` is equivalent to `-remove (RO/RW/ZI/DBG)`.

`-noremove`

This option does not perform unused section elimination on the input sections. This retains all input sections in the final image even if they are unused.

-entry *location* This option specifies the unique initial entry point of the image. The image can contain multiple entry points, but the initial entry point specified using this command is stored in the executable file header for use by the loader. There can be only one occurrence of this option on the command line. ARM debuggers use this entry address to initialize the PC when an image is loaded. The initial entry point must meet the following conditions:

- the image entry point must lie within an execution region
- the execution region must be non-overlay, and must be a root execution region (load address == execution address).

Replace *location* with one of the following:

entry_address

A numerical value, for example:

-entry 0x0

symbol This option specifies an image entry point as the address of *symbol*. For example:

-entry int_handler

If multiple definitions of *symbol* exist, armlink generates an error message.

offset+object(section)

This option specifies an image entry point as an *offset* inside a *section* within a particular *object*. For example:

-entry 8+startup(startupseg)

There must be no spaces within the argument to **-entry**. The *input section* and object names are matched without case-sensitivity. You can use the following simplified notation:

- *object(section)* if offset is zero
- *object* if there is only one input section. If this form is used and there is more than one non-debug input section in object, armlink generates an error message.

`-keep section-id` Specifies input sections that are not to be removed by unused section elimination. See *Specifying an image memory map* on page 3-5. All forms of *section-id* argument to `-keep` can contain the `*` and `?` wildcards. Replace *section-id* with one of the following:

symbol This option specifies that the input section defining *symbol* is to be retained during unused section elimination. If multiple definitions of *symbol* exist, then all input sections that define *symbol* are treated similarly. For example:

`-keep int_handler`

To keep all sections that define a symbol ending in `_handler`, use:

`-keep *_handler`

There can be multiple occurrences of this command on the command line.

object(section)

This option specifies that *section* from *object* is to be retained during unused section elimination. The input section and object names are matched without case-sensitivity. For example, to keep the `vect` section from the `vectors.o` object use:

`-keep vectors.o(vect)`

To keep all sections from the `vectors.o` object where the first three characters of the name of the section is `vec`, use:

`-keep vectors.o(vec*)`

There can be multiple occurrences of this command on the command line.

object This option specifies that the single input section from *object* is to be retained during unused section elimination. The object name is matched without case-sensitivity. If you use this short form and there is more than one input section in *object*, armlink generates an error message. For example:

`-keep dspdata.o`

To keep the single input section from each of the objects that have a name starting with `dsp`, use:

`-keep dsp*.o`

There can be multiple occurrences of this option on the command line.

-first *section-id* This option places the selected input section first in its execution region. This can, for example, place the section containing the reset and interrupt vector addresses first in the image. Replace *section-id* with one of the following:

symbol Selects the section that defines *symbol*. You must not specify a symbol that has more than one definition, because more than one section cannot be placed first.
For example:
-first reset

object(section)
Selects *section* from *object*. There must be no space between *object* and the following open parenthesis. For example:
-first init.o(init)

object Selects the single input section in *object*. If you use this short form and there is more than one input section, armlink generates an error message. For example:
-first init.o

Using -first cannot override the basic attribute sorting order for output sections in regions that places RO first, RW second, and ZI last. If the region has an RO section, an RW or a ZI section cannot be placed first. If the region has an RO or RW section, a ZI section cannot be placed first.

Two different sections cannot both be placed first in the same execution region, so only one instance of this option is permitted.

-last *section-id* This option places the selected input section last in its execution region. For example, this can force an input section that contains a checksum to be placed last in the RW section. Replace *section-id* with one of the following:

symbol Selects the section that defines *symbol*. You must not specify a symbol that has more than one definition, because more than one section cannot be placed last.
For example:
-last checksum

object(section)
Selects the *section* from *object*. There must be no space between *object* and the following open parenthesis. For example:
-last checksum.o(check)

object Selects the single input section from *object*. If there is more than one input section in *object*, armlink generates an error message.

Using `-last` cannot override the basic attribute sorting order for output sections in regions that places RO first, RW second, and ZI last. If the region has a ZI section, an RW section cannot be placed last. If the region has an RW or ZI section, an RO section cannot be placed last.

Two different sections cannot both be placed last in the same execution region, so only one instance of this option is permitted.

`-libpath pathlist` This option specifies a list of paths that are used to search for the ARM standard C and C++ libraries.

Note

This option does not affect searches for user libraries.

These paths override the path specified by the ARMLIB environment variable. *pathlist* is a comma-separated list of paths *path1*, *path2*, ... *pathn* that are only used to search for required ARM libraries. The default path for the directory containing the ARM libraries is specified by the ARMLIB environment variable. See *Library searching, selection, and scanning* on page 5-3 for more information on including libraries.

`-scanlib` This option enables scanning of default libraries (the standard ARM C and C++ libraries) to resolve references. This is the default.

`-noscanlib` This option prevents the scanning of default libraries in a link step.

`-locals` This option instructs the linker to add local symbols to the output symbol table when producing an executable image. This is the default.

`-nolocals` This option instructs the linker not to add local symbols to the output symbol table when producing an executable image. This is a useful optimization if you want to reduce the size of the output symbol table.

-callgraph

This option creates a static callgraph of functions in HTML format. The callgraph gives definition and reference information for all functions in the image.

Note

Any assembler files must contain appropriate FRAME directives, so that the linker can use the debug information subsequently generated to work out stack usage.

Also, the stack usage cannot take into account, for example, recursive functions and function pointers.

For each function func it lists:

- the processor state for which the function is compiled (ARM or Thumb)
- the set of functions that call func
- the set of functions that are called by func
- the number of times the address of func is used in the image.

In addition, the callgraph identifies functions that are:

- called through interworking veneers
- defined outside the image
- allowed to remain undefined (weak references).

The static callgraph also provides stack usage information. It lists:

- the size of the stack frame used by each function
- the maximum size of the stack used by the function over any call sequence, that is, over any acyclic chain of function calls in the callgraph.

-info *topics*

This option prints information about specified topics, where *topics* is a comma-separated list of topic keywords. A topic keyword can be one of the following:

- | | |
|--------|--|
| sizes | Gives a list of the Code and Data (RO Data, RW Data, ZI Data, and Debug Data) sizes for each input object and library member in the image. Using this option implies -info sizes,totals. |
| totals | Gives totals of the Code and Data (RO Data, RW Data, ZI Data, and Debug Data) sizes for input objects and libraries. |

- veneers Gives details of armlink-generated veneers. For more information on veneers see *Veneer generation* on page 3-12.
- unused Lists all unused sections that were eliminated from the image as a result of using `-remove`.

Note

Spaces are not permitted between keywords in a list. For example, you can enter:

`-info sizes,totals`

but not:

`-info sizes, totals`

- `-map` This option creates an image map. The image map contains the address and the size of each load region, execution region, and input section in the image, including debugging and linker-generated input sections.
- `-symbols` This option lists each local and global symbol used in the link step, and its value. This includes linker-generated symbols.
- `-symdefs file` This option creates a symbol definition file containing the global symbol definitions from the output image.
- By default, all global symbols are written to the symdefs file. If *file* already exists, the linker restricts its output to the symbols listed in the existing symdefs file.
- If *file* is specified without path information, the linker searches for it in the directory where the output image is being written. If it is not found, it is created in that directory.
- You can use the symbol definition file as input when linking another image. See *Accessing symbols in another image* on page 4-5 for more information.
- `-edit file` This option enables you to specify a *steering file* containing commands to edit the symbol tables in the output binary. You can specify commands in the steering file to:
- Hide global symbols. Use this option to hide specific global symbols in object files. The hidden symbols are not publicly visible.
 - Rename global symbols. Use this option to resolve symbol naming conflicts.

See *Hiding and renaming global symbols* on page 4-9 for more information on steering file syntax.

- xref This option lists all cross-references between input sections.

- xreffrom *object(section)*

This option lists cross-references from input *section* in *object* to other input sections. This is a useful subset of the listing produced by using -xref if you are interested in references from a specific input section. You can have multiple occurrences of this option to list references from more than one input section.

- xrefto *object(section)*

This option lists cross-references to input *section* in *object* from other input sections. This is a useful subset of the listing produced by using -xref if you are interested in references to a specific input section. You can have multiple occurrences of this option in order to list references to more than one input section.

- errors *file* Redirects the diagnostics from the standard error stream to *file*.

- list *file* This option redirects the diagnostics from output of the -info, -map, -symbols, -xref, -xreffrom, and -xrefto commands to *file*.
 If *file* is specified without path information, it is created in the output directory, that is the directory the output image is being written to.

- verbose This option prints detailed information about the link operation, including the objects that are included and the libraries from which they are taken.

- unmangled This option instructs the linker to display unmangled C++ symbol names in diagnostic messages, and in listings produced by the -xref, -xreffrom, -xrefto, and -symbols options.
 If this option is selected, the linker unmangles C++ symbol names so that they are displayed as they appear in your source code. This is the default.

- mangled This option instructs the linker to display mangled C++ symbol names in diagnostic messages, and in listings produced by the -xref, -xreffrom, -xrefto, and -symbols options. If this option is selected, the linker does not unmangle C++ symbol names. The symbol names are displayed as they appear in the object symbol tables.

-match crossmangled

This tells the linker to match the following combinations together:

- a reference to an unmangled symbol with the mangled definition
- a reference to a mangled symbol with the unmangled definition.

Libraries and matching operate as follows:

- if the library members define a mangled definition, and there is an unresolved unmangled reference, the member is loaded to satisfy it
- if the library members define an unmangled definition, and there is an unresolved mangled reference, the member is loaded to satisfy it.

Note

This option has no effect if used with partial linking. The partial object contains all the unresolved references to unmangled symbols, even if the mangled definition exists. Matching is done only in the final link step.

-via file

This option reads a further list of input filenames and linker options from *file*.

You can enter multiple -via options on the armlink command line. The -via options can also be included within a via file. See *RealView Compilation Tools for BREW Compilers and Libraries Guide* for more information on writing via files.

-strict

This option instructs the linker to report conditions that might result in failure as errors, rather than warnings.

-unresolved symbol

This option matches each reference to an undefined symbol to the global definition of *symbol*. *symbol* must be both defined and global, otherwise it appears in the list of undefined symbols and the link step fails. This option is particularly useful during top-down development, because it enables you to test a partially-implemented system by matching each reference to a missing function to a dummy function.

This option does not display warnings.

input-file-list

This is a space-separated list of objects or libraries.

A special type of object file, the symdef file, can be included in the list to provide global symbol values for a previously generated image file. See *Accessing symbols in another image* on page 4-5 for more information.

You can use libraries in the input file list in the following ways:

- Specify particular members to be extracted from a library and added to the image as individual objects. For example, specify `mystring.lib(strcmp.o)` in the input file list.
- Specify a library to be added to the list of libraries that is used to extract members if they resolve any non-weak unresolved references. For example, specify `mystring.lib` in the input file list. Other libraries are added to this list implicitly by armlink when it scans the default library directories and selects the closest matching library variants available. Members from the libraries in this list are added to the image only when they resolve an unresolved non-weak reference. For more information see *Library searching, selection, and scanning* on page 5-3.

armlink processes the input file list in the following order:

1. Objects are added to the image unconditionally.
2. Members selected from libraries using patterns are added to the image unconditionally, as if they were objects. For example:
`armlink main.o mylib(stdio.o) mylib(a*.o).`
unconditionally adds all `a*.o` objects, and `stdio.o`.
3. Libraries are added to the list of libraries that are later used to resolve any remaining non-weak unresolved references.

Chapter 3

Using the Basic Linker Functionality

This chapter describes the basic armlink functionality. This chapter contains the following sections:

- *Specifying the image structure* on page 3-2
- *Section placement* on page 3-8
- *Optimizations and modifications* on page 3-11
- *Using command-line options to create simple images* on page 3-14.

For information about advanced linker functionality, refer to the following chapters:

Symbol access Chapter 4 *Accessing Image Symbols*.

Libraries Chapter 5 *Using libraries*.

Changing format Chapter 6 *Using fromELF*.

3.1 Specifying the image structure

The structure of an image is defined by:

- the number of its constituent regions and output sections
- the positions in memory of these regions and sections when the image is loaded
- the positions in memory of these regions and sections when the image executes.

3.1.1 Building blocks for objects and images

An image, as stored in an executable file, is constructed from a hierarchy of images, regions, output sections, and input sections:

- An image consists of one or more regions. Each region consists of one or more output sections.
- Each output section contains one or more input sections.
- Input sections are the code and data information in an object file.

Figure 3-1 on page 3-3 shows the relationship between regions, output sections, and input sections.




Figure 3-1 Building blocks for an image

Input sections

An input section contains code or initialized data, or describes a fragment of memory that is not initialized or that must be set to zero before the image can execute. Input sections can have the attributes RO, RW, or ZI. These three attributes are used by armlink to group input sections into bigger building blocks called output sections and regions.

Output sections

An output section is a contiguous sequence of input sections that have the same RO, RW, or ZI attribute. An output section has the same attributes as its constituent input sections. Within an output section, the input sections are sorted according to the rules described in *Section placement* on page 3-8.

Regions

A region is a contiguous sequence of one to three output sections. The output sections in a region are sorted according to their attributes. The RO output section is first, then the RW output section, and finally the ZI output section. A region typically maps onto a physical memory device, such as ROM, RAM, or peripheral.

3.1.2 Load view and execution view of an image

Image regions are placed in the system memory map at load time. Before you can execute the image, you might have to move some of its regions to their execution addresses and create the ZI output sections. For example, initialized RW data might have to be copied from its load address in ROM to its execution address in RAM.

The memory map of an image has the following distinct views as shown in Figure 3-2.

- Load view

This view describes each image region and section in terms of the address it is located at when the image is loaded into memory, that is the location before the image starts executing
- Execution view

This view describes each image region and section in terms of the address it is located at while the image is executing.

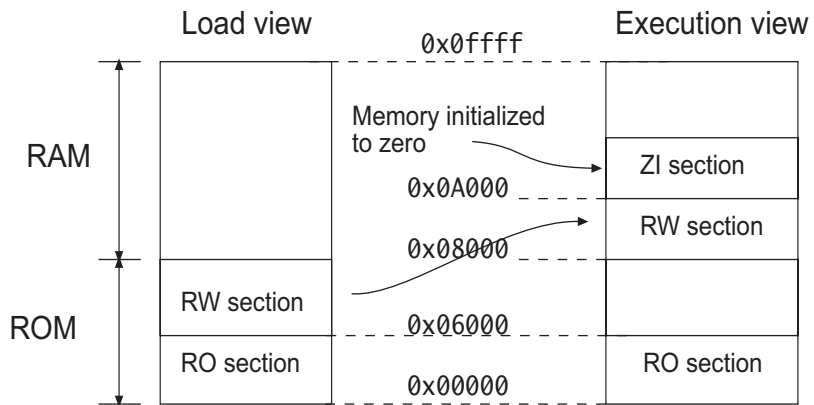


Figure 3-2 Load and execution memory maps

Table 3-1 compares the load and execution views.

Table 3-1

Load	Description	Execution	Description
Load address	The address where a section, or region is loaded into memory before the image containing it starts executing. The load address of a section or a non-root region can differ from its execution address.	Execution address	The address where a section or region is located while the image containing it is being executed
Load region	A region in the load address space.	Execution region	A region in the execution address space

3.1.3 Specifying an image memory map

An image can consist of any number of regions and output sections. Any number of these regions can have different load and execution addresses. To construct the memory map of an image, armlink must have information about:

Grouping How input sections are grouped into output sections and regions.

Placement Where image regions are to be located in the memory maps.

Depending on the complexity of the memory maps of the image, there are two ways to pass this information to armlink:

Using command-line options

The following options can be used for simple cases where an image has only one or two load regions and up to three execution regions:

- -ro-base
- -rw-base
- -split
- -ropi
- -rwpi

For more information, see *Using command-line options to create simple images* on page 3-14.

3.1.4 Image entry points

An entry point in an image is a location where program execution can start. There are two distinct types of entry point:

Initial entry point

The *initial* entry point for an image is a single value that is stored in the ELF header file. For programs loaded into RAM by an operating system or boot loader, the loader starts the image execution by transferring control to the initial entry point in the image.

There can be only one initial entry point in an image. The initial entry point can be, but is not required to be, one of the entry points set by the ENTRY directive.

Entry points set by the ENTRY directive

These are entry points that are set in the assembly language sources with the ENTRY directive. In embedded systems, this directive is typically used to mark code that is entered through the exception vectors.

Specifying an initial entry point

You can specify an initial entry point with the `-entry` linker option. You can specify the `-entry` option only once. See the description in *armlink command syntax* on page 2-8 for more information.

An image must contain an initial entry point if it is to be executed by a loader, such as a boot loader or operating system program loader. For example, an image that is an OS is loaded by the boot loader and entered at the initial entry point specified in the executable file header. After the image is loaded it overwrites the boot loader and becomes the OS. In the case of an operating system, the image contains additional entry points, such as the exception handler addresses.

For embedded applications with ROM at zero use `-entry 0x0` (or `0xffff0000` for CPUs that have high vectors).

The initial entry point must meet the following conditions:

- the image entry point must always lie within an execution region
- the execution region must be non-overlay, and must be a root region (the load address is the same as the execution address).

If you do not use the `-entry` option to specify the initial entry point the linker attempts to set one if it can:

- If you do not use the `-entry` option and the input objects contain only one entry point set by the `ENTRY` directive, the linker uses that entry point as the initial entry point for the image.
- The linker generates an output object that does not contain an initial entry point if you do not use the `-entry` option and either:
 - more than one entry point has been specified by using the `ENTRY` directive
 - no entry point has been specified by using the `ENTRY` directive.

If the output object does not contain an initial entry point, the linker issues the following warning:

L6305W: Image does not have an entry point. (Not specified or not set due to multiple choices)

Specifying multiple entry points

An embedded image can have multiple entry points. For example, an embedded image typically has entry points that are used by the Reset, IRQ, FIQ, SVC, UNDEF, and ABORT exceptions to transfer control when the image is running.

You can specify multiple entry points in an image with the `ENTRY` directive. The directive marks the output code section with an `ENTRY` keyword that instructs the linker not to remove the section when it performs unused section elimination. For C and C++ programs, the `__main()` function in the C library is also an entry point. See *RealView Compilation Tools for BREW Assembler Guide* for more information on the `ENTRY` directive.

If an embedded image is to be used by a loader, it must have a single initial entry point specified in the header. See *Specifying an initial entry point* on page 3-6 for more information.

3.2 Section placement

The linker sorts all the input sections within a region according to their attributes. Input sections with identical attributes form a contiguous block within the region.

The base address of each input section is determined by the sorting order defined by the linker, and is correctly aligned within the output section that contains it.

When generating an image, the linker sorts the input sections in the following order:

- By attribute.
- By input section name.
- By their positions in the input list, except where overridden by a `-first` or `-last` option. This is described in *Using FIRST and LAST to place sections* on page 3-9.

By default, the linker creates an image consisting of an RO, an RW, and optionally a ZI, output section. The RO output section can be protected at run-time on systems that have memory management hardware. RO sections can also be placed into ROM in the target.

3.2.1 Ordering input sections by attribute

Portions of the image associated with a particular language run-time system are collected together into a minimum number of contiguous regions. `armlink` orders input sections by attribute as follows:

- read-only code
- read-only data
- other initialized data
- zero-initialized (uninitialized) data.

Input sections that have the same attributes are ordered by their names. Names are considered to be case-sensitive and are compared in alphabetical order using the ASCII collation sequence for characters.

Identically attributed and named input sections are ordered according to their relative positions in the input list.

These rules mean that the positions of identically attributed and named input sections included from libraries are not predictable. If more precise positioning is required, you can extract modules manually, and include them in the input list.

3.2.2 Using FIRST and LAST to place sections

Within a region, all RO code input sections are contiguous and form an RO output section that must precede the output section containing all the RW input sections.

Use the `-first` and `-last` linker options to place input sections.

However, FIRST and LAST must not violate the basic attribute sorting order. This means that an input section can be first (or last) in the execution region if the output section it is in is the first (or last) output section in the region. For example, in an execution region containing RO input sections, the FIRST input section must be an RO input section. Similarly, if the region contains any ZI input sections, the LAST input section must be a ZI input section.

Within each output section, input sections are sorted alphabetically according to their names, and then by their positions in the input order.

3.2.3 Aligning sections

When input sections have been ordered and before the base address is fixed, armlink inserts padding, if required, to force each input section to start at an address that is a multiple of the input section alignment. You can expand the alignment (forcing something that would normally be 4-byte aligned to be 8-byte aligned), but you cannot reduce the natural alignment (forcing 2-byte alignment on something that would normally be 4-byte aligned). Input sections are commonly aligned at word boundaries. The input section alignment can be specified by using either:

ALIGN	Use this attribute to the AREA directive from assembly language. The input section address will be a multiple of $2^{(\text{value in align attribute})}$.
__align	Use this qualifier from C. The input section address will be a multiple of the value specified (the value must be a power of 2).

See Example 3-1 and the description of `__align` in *RealView Compilation Tools for BREW Compilers and Libraries Guide*. See Example 3-2 on page 3-10 and the description of ALIGN in *RealView Compilation Tools for BREW Assembler Guide*.

Example 3-1 Using the __align qualifier

```
__align(8) long long big_var; // big_var eight-byte aligned
```

Example 3-2 Using the ALIGN attribute in assembly code

```
AREA LDR_LABEL, CODE, READONLY, ALIGN=3 ; align on eight-byte boundary
```

3.3 Optimizations and modifications

armlink performs some optimizations and modifications in order to remove duplicate sections, enable interworking between ARM and Thumb code, and enable long branches.

3.3.1 Common debug section elimination

The compilers and assemblers generate one set of debug sections for each source file that contributes to a compilation unit. armlink can detect multiple copies of a debug section for a particular source file and discard all but one copy in the final image. This can result in a considerable reduction in image debug size.

3.3.2 Common section elimination

If there are inline functions or templates used in the source, the ARM C++ compilers generate complete objects for linking such that each object contains the out-of-line copies of inline functions and template functions that the object requires. When these functions are declared in a common header file, the functions might be defined many times in separate objects that are subsequently linked together. In order to eliminate duplicates, the compilers compile these functions into separate instances of common code sections and armlink retains just one copy of each common code section.

It is possible that the separate instances of a common code section are not identical. Some of the copies, for example, might be found in a library that has been built with different (but compatible) build options, different optimization, or different debug options.

If the copies are not identical, armlink retains the best available variant of each common code section based on the attributes of the input objects.

3.3.3 Unused section elimination

Unused section elimination removes code that is never executed, or data that is not referred to by the code, from the final image. This optimization can be controlled by the `-(no)remove` and `-first`, `-last`, and `-keep` linker options. Use the `-info unused` linker option to instruct the linker to generate a list of the unused sections that have been eliminated.

An input section is retained in the final image in the following conditions:

- if it contains an entry point
- if it is referred to, directly or indirectly, by a non-weak reference from an input section containing an entry point
- if it was specified as the first or last input section by the `-first` or `-last` option

- if it has been marked as unremovable by the `-keep` option.

3.3.4 Veneer generation

armlink must generate veneers when:

- a branch involves change of state between ARM state and Thumb state
- a branch involves a destination beyond the branching range of the current state.

A veneer can extend the range of branch and change processor state. armlink combines long branch capability into the state change capability. All interworking veneers are also long branch veneers.

The following veneer types handle different branching requirements. The linker generates position-independent versions of the veneers if required:

ARM to ARM Long branch capability.

ARM to Thumb Long branch capability and interworking capability.

Thumb to ARM Long branch capability and interworking capability.

Thumb to Thumb Long branch capability.

armlink creates one input section called `Veneer$$Code` for each veneer. A veneer is generated only if no other existing veneer can satisfy the requirements. If two input sections contain a long branch to the same destination, only one veneer is generated if the veneer can be reached by both sections.

All veneers cannot be collected into one input section because the resulting veneer input section might not be within range of other input sections. If the sections are not within addressing range, long branching is not possible.

3.3.5 Reuse of veneers with overlay execution regions

The linker reuses veneers whenever possible. However, both the following conditions are enforced on reuse:

- an overlay execution region cannot reuse a veneer placed in any other execution region
- no other execution region can reuse a veneer placed in an overlay execution region.

If these conditions are not met, new veneers are created instead of reusing existing ones. A veneer is always placed in the execution region that contains the call requiring the veneer. This implies that:

- for an overlay execution region, all its veneers are included within the execution region
- an overlay execution region never requires a veneer from another execution region.

3.4 Using command-line options to create simple images

A simple image consists of a number of input sections of type RO, RW, and ZI. These input sections are collated to form the RO, the RW, and the ZI output sections.

Depending on how the output sections are arranged within load and execution regions, there are three basic types of simple image:

Type 1 One region in load view, three contiguous regions in execution view. Use the `-ro-base` option to create this type of image.

See *Type 1: one load region and contiguous output regions* for more details.

Type 2 One region in load view, three non-contiguous regions in execution view. Use the `-ro-base` and `-rw-base` options to create this type of image.

See *Type 2: one load region and non contiguous output regions* on page 3-16 for more details.

Type 3 Two regions in load view, three non-contiguous regions in execution view. Use the `-ro-base`, `-rw-base`, and `-split` options to create this type of image.

See *Type 3: two load regions and noncontiguous output regions* on page 3-17 for more details.


In all three simple image types, there are up to three execution regions.

- the first execution region contains the RO output section
- the second execution region contains the RW output section (if present)
- the third execution region contains the ZI output section (if present).

These execution regions are referred to as the RO, the RW, and the ZI execution region.

3.4.1 Type 1: one load region and contiguous output regions

An image of this type consists of a single load region in the load view and the three execution regions are placed contiguously in the memory map. This approach is suitable for systems that load programs into RAM, for example, an OS bootloader, or a desktop system (see Figure 3-3 on page 3-15).

**Figure 3-3 Type 1**

Load view

The single load region consists of the RO and RW output sections placed consecutively. The ZI output section does not exist at load time. It is created before execution using the output section description in the image file.

Execution view

The three execution regions containing the RO, RW, and ZI output sections are arranged contiguously. The execution address of the RO and RW execution regions are the same as their load addresses, so nothing has to be moved from its load address to its execution address. However, the ZI execution region that contains the ZI output section is created before execution begins.

Use armlink option `-ro-base address` to specify the load and execution address of the region containing the RO output. The default address is `0x8000` as shown in Figure 3-3.

The `-reloc` option is used to make relocatable images. Used on its own, `-reloc` makes an image similar to Simple type 1, but the single load region has the RELOC attribute.

3.4.2 Type 2: one load region and non contiguous output regions

An image of this type consists of a single load region, and three execution regions in execution view. The RW execution region is not contiguous with the RO execution region. This approach is used, for example, for ROM-based embedded systems (see Figure 3-4).

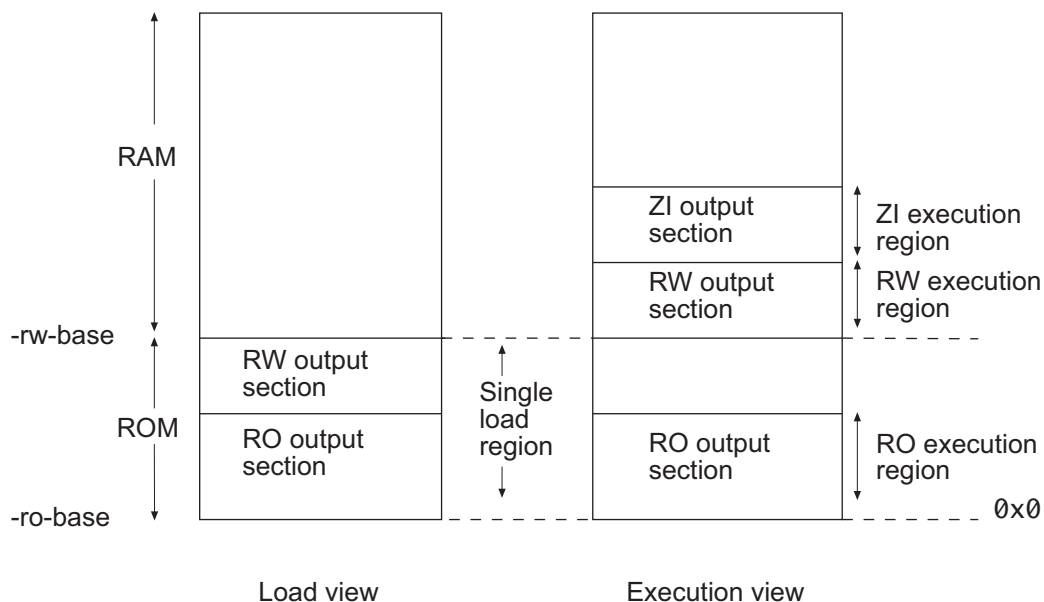


Figure 3-4 Type 2

Load view

In the load view, the single load region consists of the RO and RW output sections placed consecutively, in ROM for example. The ZI output section does not exist at load time. It is created before execution using the description of the output section contained in the image file.

Execution view

In the execution view, the first execution region contains the RO output section and the second execution region contains the RW and ZI output sections.

The execution address of the region containing the RO output section is the same as its load address, so the RO output section does not have to be moved.

The execution address of the region containing the RW output section is different from its load address, so the RW output section is moved from its load address (from the single load region) to its execution address (into the second execution region). The ZI execution region, and its output section, is placed contiguously with the RW execution region.

Use armlink options `-ro-base address` to specify the load and execution address for the RO output section, and `-rw-base exec_address` to specify the execution address of the RW output section. If you do not use the `-ro-base` option to specify the address, the default value of `0x8000` is used by armlink. For an embedded system, `0x0` is typical for the `-ro-base` value. If you do not use the `-rw-base` option to specify the address, the default is to place RW directly above RO.

3.4.3 Type 3: two load regions and noncontiguous output regions

This type of image is similar to images of type 2 except that the single load region in type 2 is now split into two load regions (see Figure 3-5).

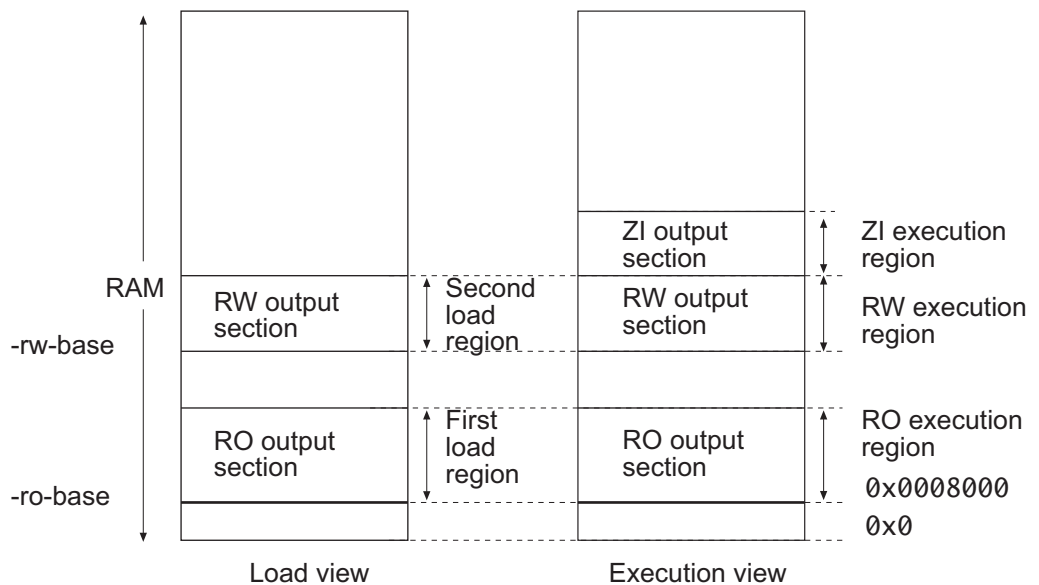


Figure 3-5 Type 3

Load view

In the load view, the first load region consists of the RO output section, and the second load region consists of the RW output section. The ZI output section does not exist at load time. It is created before execution using the description of the output section contained in the image file.

Execution view

In the execution view, the first execution region contains the RO output section and the second execution region contains the RW and ZI output sections.

The execution address of the RO region is the same as its load address, so the contents of the RO output section are not moved from their load address to their execution address.

The execution address of the RW region is also the same as its load address, so the contents of the RW output section are not moved from their load address to their execution address. However, the ZI output section is created before execution begins and placed after the RW region.

Specify the load and execution address using the following linker options:

-reloc The combination **-reloc -split** makes an image similar to Simple type 3, but the two load regions now have the RELOC attribute.

-split This option splits the default single load region (that contains both the RO and RW output sections) into two load regions (one containing the RO output section and one containing the RW output section) so that they can be separately placed using **-ro-base** and **-rw-base**.

-ro-base *address*

This option instructs armlink to set the load and execution address of the region containing the RO section at a 4-byte aligned *address* (for example, the address of the first location in ROM). If **-ro-base** option is not used to specify the address, the default value of 0x8000 is used by armlink.

-rw-base *address*

This option instructs armlink to set the execution address of the region containing the RW output section at a 4-byte aligned *address*. If this option is used with **-split**, both the load and execution addresses of the RW region are specified.

Using **-reloc -rw-base** without also using **-split** causes an error.

Chapter 4

Accessing Image Symbols

This chapter describes how to reference symbols with armlink. It contains the following sections:

- *Accessing linker-defined symbols* on page 4-2
- *Accessing symbols in another image* on page 4-5
- *Hiding and renaming global symbols* on page 4-9
- *Using \$Super\$\$ and \$Sub\$\$ to override symbol definitions* on page 4-12.

4.1 Accessing linker-defined symbols

The linker defines some symbols that contain the character sequence `$$`. These symbols and all other external names containing the sequence `$$` are ARM-reserved names. The symbols are used to specify region base addresses, output section base addresses, and input section base addresses and their limits.

These symbolic addresses can be imported and used as relocatable addresses by your assembly language programs, or referred to as **extern** symbols from your C or C++ source code.

Note

Linker-defined symbols are defined by armlink only when your code references them.

4.1.1 Importing linker-defined symbols

There are two methods of importing linker-defined symbols into your C or C++ source code:

- `extern int symbol_name;`
- `extern void *symbol_name;`

If you declare a symbol as an `int`, then you must use the address-of operator to obtain the correct value as shown in Example 4-1.

Example 4-1 Importing linker-defined symbols

```
extern unsigned int Image$$ZI$$Limit;
config.heap_base = (unsigned int) &Image$$ZI$$Limit;
```

4.1.2 Region-related symbols

Region-related symbols are generated when armlink creates an image. Table 4-1 on page 4-3 shows the symbols that armlink generates for every execution region present in the image.

The linker uses *region_name* values of:

- `ER_RO`, for read-only regions
- `ER_RW`, for read-write regions
- `ER_ZI`, for zero-initialized regions.

For every execution region containing a ZI output section, armlink generates additional symbols containing `$$ZI$`.

Table 4-1 Region-related linker symbols

Symbol	Description
<code>Load\$\$region_name\$\$Base</code>	Load address of the region.
<code>Image\$\$region_name\$\$Base</code>	Execution address of the region.
<code>Image\$\$region_name\$\$Length</code>	Execution region length in bytes (multiple of 4).
<code>Image\$\$region_name\$\$Limit</code>	Address of the byte beyond the end of the execution region.
<code>Image\$\$region_name\$\$ZI\$Base</code>	Execution address of the ZI output section in this region.
<code>Image\$\$region_name\$\$ZI\$Length</code>	Length of the ZI output section in bytes (multiple of 4).
<code>Image\$\$region_name\$\$ZI\$Limit</code>	Address of the byte beyond the end of the ZI output section in the execution region.

Note

- The ZI output sections of an image are not created statically, but are automatically created dynamically at run-time. Therefore there is no load address symbol for ZI output sections.
 - It is recommended that you use region-related symbols in preference to section-related symbols.
-

4.1.3 Section-related symbols

The output section symbols shown in Table 4-2 are generated if command-line options are used to create a simple image. A simple image has three output sections (RO, RW, and ZI) that produce the three execution regions. For every input section present in the image, armlink generates the input symbols shown in Table 4-2.

Table 4-2 Section-related linker symbols

Symbol	Section type	Description
Image\$\$RO\$\$Base	Output	Address of the start of the RO output section.
Image\$\$RO\$\$Limit	Output	Address of the first byte beyond the end of the RO output section.
Image\$\$RW\$\$Base	Output	Address of the start of the RW output section.
Image\$\$RW\$\$Limit	Output	Address of the byte beyond the end of the ZI output section. (The choice of the end of the ZI region rather than the end of the RW region is to maintain compatibility with legacy code.)
Image\$\$ZI\$\$Base	Output	Address of the start of the ZI output section.
Image\$\$ZI\$\$Limit	Output	Address of the byte beyond the end of the ZI output section.
SectionName\$\$Base	Input	Address of the start of the consolidated section called SectionName.
SectionName\$\$Limit	Input	Address of the byte beyond the end of the consolidated section called SectionName.

————— **Note** —————

If your code refers to the input-section symbols, it is assumed that you expect all the input sections in the image with the same name to be placed contiguously in the image memory map.

4.2 Accessing symbols in another image

If you want one image to know the global symbol values of another image, you can use a *symbol definitions* (symdefs) file.

This can be used, for example, if you have one image that always resides in ROM and multiple images that are loaded into RAM. The images loaded into RAM can access global functions and data from the image located in ROM.

4.2.1 Reading a symdefs file

A symdefs file can be considered as an object file with symbol information but no code or data. To read a symdefs file, add it to your file list as you would add any object file. armlink reads the file and adds the symbols and their values to the output symbol table. The added symbols have ABSOLUTE and GLOBAL attributes.

If a partial link is being performed, the symbols are added to the output object symbol table. If a full link is being performed, the symbols are added to the image symbol table.

armlink generates error messages for invalid rows in the file. A row is invalid if:

- any of the columns are missing
- any of the columns have invalid values.

The symbols extracted out of a symdefs file are treated in exactly the same way as symbols extracted from an object symbol table. The same restrictions regarding multiple symbol definitions and ARM/Thumb synonyms apply.

4.2.2 Creating a symdefs file

Use armlink option `-symdefs filename` to produce a symdefs file.

armlink produces a symdefs file during a successful final link stage. It is not produced for partial linking or for unsuccessful final linking.

————— **Note** —————

If *filename* does not exist, the file is created containing all the global symbols. If *filename* exists, the existing contents of *filename* are used to select the symbols that are output when armlink rewrites the file.

Outputting a subset of the global symbols

By default, all global symbols are written to the symdefs file.

When *filename* exists, armlink uses its contents to restrict the output to a subset of the global symbols. To restrict the output symbols:

1. Specify `-symdefs filename` when you are doing a nearly final link for *image1*. armlink creates a symdef file *filename*.
2. Open *filename* in a text editor, remove any symbol entries you do not want in the final list, and save the file.
3. Specify `-symdefs filename` when you are doing a final link for *image1*.
You can edit *filename* at any time to add comments and relink *image1* again to, for example, update the symbol definitions after one or more objects use to create *image1* have changed.

armlink updates a symdefs file as follows:

1. armlink creates a temporary output file.
2. Lines from *filename* are processed:
 - a. The comments and blank lines are copied unchanged to the temporary file.
 - b. When a symbol name is found in *filename*, the definition of the named symbol from the current image is output to the temporary file.

Note

If a symbol occurs in *filename* more than once, only one definition of the symbol is placed in the temporary file.

If a symbol is found in *filename* but is not defined in the current image, no output will be produced for that symbol.

3. If the final link is successful, *filename* is deleted and the temporary file renamed as *filename*.

4.2.3 Symdefs file format

The symdefs file is a type of object file that contains symbols and their values. Unlike other object files, however, it does not contain any code or data.

The file consists of an identification line, optional comments, and symbol information as shown in Example 4-2.

Example 4-2

```
#<SYMDEFS># ARM Linker, RV_BREW1.2 [Build 1307]: Last Updated: Fri June 21 13:56:35 2002
;value type name, this is an added comment
0x00001000 A function1
0x00002000 T function2
    # This is also a comment, blank lines are ignored

0x00003300 A function3
0x00003340 D table1
```

Identifying string

If the first 11 characters in the text file are #<SYMDEFS>#, armlink recognizes the file as a symdefs file.

The identifying string is followed by linker version information, and date and time of the most recent update of the symdefs file. The version and update information are not part of the identification string.

Comments

You can insert comments manually with a text editor. Comments have the following properties:

- Any line where the first non-whitespace character is ; or # is a comment.
- The first line must start with the special identifying comment #<SYMDEFS>#. This comment is inserted by armlink when the file is produced and must not be manually deleted.
- A ; or # after the first non-whitespace character does not start a comment.
- Blank lines are ignored and can be inserted to improve readability.

Symbol information

The symbol information is provided by the address, type, and name of the symbol on a single line:

Symbol value	armlink writes the absolute address of the symbol in fixed hexadecimal format, for example <code>0x00004000</code> . If you edit the file, you can use either hexadecimal or decimal formats for the address value.						
Type flag	<p>The single letter for type is:</p> <table><tr><td>A</td><td>ARM code</td></tr><tr><td>T</td><td>Thumb code</td></tr><tr><td>D</td><td>Data.</td></tr></table>	A	ARM code	T	Thumb code	D	Data.
A	ARM code						
T	Thumb code						
D	Data.						
Symbol name	The name of the symbol.						

4.3 Hiding and renaming global symbols

This section describes how to use a steering file to hide or rename global symbol names in output files. For example, you can use steering files to protect intellectual property, or avoid namespace clashes. A steering file is a text file that contains a set of commands to edit the symbol tables of output objects.

Use the armlink command line option `-edit filename` to specify the steering file (see the description of the `-edit` option in *armlink command syntax* on page 2-8).

The following commands are supported:

- *RENAME* on page 4-10
- *HIDE* on page 4-11
- *SHOW* on page 4-11.

4.3.1 Steering file format

A steering file is a plain text file of the following format:

- Lines with a semicolon (;) or hash (#) character as the first non-whitespace character are interpreted as comments. A comment is treated as a blank line.
- Blank lines are ignored.
- Each non blank, non comment line is either a command, or part of a command that is split over consecutive nonblank lines.
- Command lines that end with a comma (,) as the last non-whitespace character are continued on the next nonblank line.

Each command consists of a command, followed by one or more comma-separated operand groups. Each operand group comprises either one or two operands, depending on the command. The command is applied to each operand group in the command. The following rules apply:

- Commands are case-insensitive, but are conventionally shown in uppercase.
- Operands are case-sensitive because they must be matched against case-sensitive symbol names. You can use wild card characters in operands.

Commands are applied to global symbols only. Other symbols, such as local symbols or STT_FILE, are not affected.

4.3.2 Steering file commands

The following steering file commands enable you to rename, hide, and show symbols in the symbol table.

RENAME

The RENAME command renames defined and undefined global symbol names.

Syntax

```
RENAME pattern {AS replacement_pattern} [,pattern AS replacement_pattern] *
```

where:

pattern Is a string, optionally including wildcard characters, that matches zero or more global symbols. If *pattern* does not match any global symbol, the linker ignores the command. The operand can match both defined and undefined symbols.

replacement_pattern

Is a string, optionally including wild card characters, to which the symbol is to be renamed. Wild cards must have a corresponding wild card in *pattern*. The characters matched by the *pattern* wild card are substituted for the *replacement_pattern* wildcard.

For example, for a symbol named func1:

```
RENAME f* AS my_f*
renames func1 to my_func1.
```

Usage

You cannot rename a symbol to a symbol name that already exists, even if the target symbol name is being renamed itself. Only one wildcard character (either * or ?) is permitted in RENAME.

armlink processes the steering file before doing any replacements. You cannot, therefore, use RENAME A AS B on line 1 and then RENAME B AS A on line 2.

HIDE

The HIDE command makes defined global symbols in the symbol table anonymous.

Syntax

```
HIDE pattern [,pattern] *
```

where:

pattern Is a string, optionally including wildcard characters, that matches zero or more defined global symbols. If *pattern* does not match any defined global symbol, the linker ignores the command. You cannot hide undefined symbols.

SHOW

The SHOW command makes global symbols visible that were previously hidden with the HIDE command. This command is useful if you want to un-hide a specific symbol that has been hidden using a HIDE command with a wild card.

Syntax

```
SHOW pattern [,pattern] *
```

where:

pattern Is a string, optionally including wildcard characters, that matches zero or more global symbols. If *pattern* does not match any global symbol, the linker ignores the command.

4.4 Using `$Super$$` and `$Sub$$` to override symbol definitions

There are situations where an existing symbol cannot be modified because, for example, it is located in an external library or in ROM code.

Use the `$Super$$` and `$Sub$$` patterns to patch an existing symbol. For example, to patch the definition of a function `foo()`, use `$Sub$$foo()` and `$Super$$foo()`:

`$Sub$$foo` Identifies the new function that will be called instead of the original function `foo()`. Use this to add processing before or after the original function.

`$Super$$foo` Identifies the original unpatched function `foo()`. Use this to call the original function directly.

Example 4-3 shows the legacy function `foo()` modified to result in a call to `ExtraFunc()` and a call to `foo()`. See the `ARMELF.pdf` documentation in the `pdf\specs` directory for more details.

Example 4-3

```
extern void ExtraFunc(void);
extern void $Super$$foo(void):

/* this function will be linked instead of the original foo() */
void $Sub$$foo(void)
{
    ExtraFunc();    /* does some extra setup work */
    $Super$$foo(); /* calls the original foo function */
}
```

Chapter 5

Using libraries

This chapter describes the use of libraries with armlink. It contains the following sections:

- *About libraries* on page 5-2
- *Library searching, selection, and scanning* on page 5-3

5.1 About libraries

An object file can refer to external symbols that are, for example, functions or variables. armlink attempts to resolve these references by matching them to definitions found in other object files and libraries. armlink recognizes a collection of ELF files stored in an ar format file as a library.

5.2 Library searching, selection, and scanning

The differences between the way armlink adds object files to the image and the way it adds libraries to the image are:

- Each object file in the input list is added to the output image unconditionally, whether or not anything refers to it, however, unused sections are subsequently eliminated unless `-noremove` is used.
- A member from a library is included in the output only if an object file or an already-included library member makes a non-weak reference to it, or if armlink is explicitly instructed to add it.

Note

Members selected from libraries using patterns are added to the image unconditionally, as if they were objects.

Unresolved references to weak symbols do not cause library members to be loaded.

Note

If the `-noscanlib` option is specified, armlink does not search for the default ARM libraries and uses only those libraries that are specified in the input file list to resolve references.

If a library member is explicitly requested in the input file list, it is loaded even if it does not resolve any current references. In this case, an explicitly requested member is treated as if it were an ordinary object.

armlink creates a list of libraries as follows:

1. armlink adds any libraries specified in the input file list to the list.
2. The user-specified search path is examined by armlink to identify library directories for library requests embedded in the input objects. See *Searching for ARM libraries* on page 5-4 for details on the search process.

The best-suited library variants are chosen from the searched directories and their subdirectories. ARM-supplied libraries have multiple variants that are named according to the attributes of their members. For details on the library variants see *RealView Compilation Tools for BREW Compilers and Libraries Guide* and *Selecting ARM library variants* on page 5-4.

When armlink has constructed the list of libraries, it repeatedly scans each library in the list to resolve references. See *Scanning the libraries* on page 5-5 for details.

5.2.1 Searching for user libraries

For user libraries explicitly included on the command line, a path is required if they are not in the current working directory.

Note

The search paths used for the ARM standard libraries are *not* searched for user libraries.

5.2.2 Searching for ARM libraries

You can specify the search paths used to find the ARM standard libraries by:

- Adding the `-libpath` argument to the `armlink` command line with a comma-separated list of parent directories.

This list must end with the parent directory of the ARM library directories `armlib` and `cpplib`. The `ARMLIB` variable holds the path to the ARM library parent directory.

Note

`-libpath` overrides the paths specified by the `ARMLIB` variable.

- Using the environment variable `ARMLIB`.

`armlink` combines each parent directory given by either `-libpath`, the configuration panel, or the `ARMLIB` variable, with each subdirectory request from the input objects and identifies the place to search for the ARM library. The names of ARM subdirectories within the parent directories are placed in each compiled object by using a symbol of the form `Lib$$Request$$sub_dir_name`.

5.2.3 Selecting ARM library variants

From each of the directories selected by `armlink` when searching for ARM libraries, `armlink` must select the best-suited library. There are different variants of the ARM libraries based on the attributes of their member objects. The variant of the ARM library is coded into the name of the library. See the section on library naming conventions in *RealView Compilation Tools for BREW Compilers and Libraries Guide*.

`armlink` accumulates the attributes of each input object and uses them to select the library variant best suited to the accumulated attributes. If more than one of the selected libraries are equally suited, the library selected first is retained and the others are rejected.

The final list contains all the libraries that armlink will scan in order to resolve references.

5.2.4 Scanning the libraries

When all the directories have been searched, and the most compatible library variants have been selected and added to the list of libraries, each of the libraries is scanned to load the required members:

1. For each currently unsatisfied non-weak reference, armlink searches sequentially through the list of libraries for a matching definition. The first definition found is marked for step 2.

The sequential nature of the search ensures that armlink chooses the library that appears earlier in the list if two or more libraries define the same symbol. This enables you to override function definitions from other libraries, for example the ARM C libraries, by adding your libraries in the input file list.

2. Library members marked in step 1 are loaded. As each member is loaded it might satisfy some unresolved references, possibly including weak ones. Loading a library might also create new unresolved weak and non-weak references.
3. The process in steps 1 and 2 continues until all non-weak references are either resolved or are incapable of being resolved by any library.

If any non-weak reference remains unsatisfied at the end of the scanning operation, armlink generates an error message.

Chapter 6

Using fromELF

This chapter describes the fromELF software utility provided with RVCT for BREW. It contains the following sections:

- *About fromELF* on page 6-2
- *fromELF command-line options* on page 6-3
- *Examples of fromELF usage* on page 6-9.

6.1 About fromELF

The fromELF utility translates *Executable Linkable Format* (ELF) image files produced by armlink into other formats suited to ROM tools and to loading directly into memory. You can also use it to display various information about an ELF object or to generate text files containing the information.

fromELF outputs the following image formats:

- Plain binary format.
- Motorola 32-bit S-record format.
- Intel Hex-32 format.
- Byte Oriented (Verilog Memory Model) Hex format.
- ELF format. You can resave as ELF. For example, you can convert a -debug ELF image to a -nodebug ELF image).

fromELF can also display information about the input file, for example disassembly output or symbol listings, to either standard output or a text file.

———— Note ————

Do not link your images with the -nodebug linker option if you require a -fieldoffsets fromELF step. If your image is produced without debug information:

- fromELF cannot translate the image into other file formats
- fromELF cannot produce a meaningful disassembly listing.

6.1.1 Image structure

fromELF can translate a file from ELF to other formats. It cannot alter the image structure or addresses, other than altering the base address of Motorola S-record or Intel Hex output with the -base option. Any structural or addressing information must be provided to armlink at link time.

6.2 fromELF command-line options

The fromELF command syntax is as follows:

```
fromelf [-help] [-fieldoffsets|[[[-select select_options ]]] [-nolinkview]
[-nodebug] [-vsn] [[text_output_format]| code_output_format] [-base n]
[memory_config] [-output output_file] {input_file}
```

where:

- help** This option shows help and usage information. If this option is specified, other command-line options are ignored. Calling fromELF without any parameters produces the same help information.
- fieldoffsets** This option produces, to standard output, a list of assembly language EQU directives that equate C++ class or C structure field names to their offsets from the base of the class or structure. The input ELF file can be a relocatable object or an image.
- Use -o to redirect the output to a file. Use the INCLUDE command from armasm to load the produced file and provide access to C++ classes and C structure members by name from assembly language. See *RealView Compilation Tools for BREW Assembler Guide* for more information on armasm.

————— Note —————

If the source file does not have debug information, this option is not available. You cannot use this option together with a *code_output_format*.

-fieldoffsets outputs all structure information. To output a subset of the structures, use -select *select_options*.

If you do not require a file that can be input to armasm, use the -text -a option to format the display addresses in a more readable form. The -a option only outputs address information for structures and static data in images because the addresses are not known in a relocatable object.

-select *select_options*

Use -select *select_option* together with either the -fieldoffsets or -text -a options to select only those fields that match the patterns in the option list.

Use special characters to select multiple fields:

- Join options in the list together with a , as in: a*,b*,c*.

- The wildcard character * can be used to match any name.
- The wildcard character ? can be used to match any single letter.
- Specify the fields to include by prefixing a + to the select_option string. This is the default.
- Specify the fields to exclude by prefixing a ~ to the select_option string.

If you are using a special character on Unix, you must enclose the options in quotes to prevent the shell expanding the selection.

-noLinkview

Use -noLinkview to discard the section-level view (link-time view) from the ELF image and retain only the segment level view (load-time view). Discarding the link-view section level eliminates:

- the section header table
- the section header string table
- the string table
- the symbol table
- all debug sections.

All that is left in the output is the program header table and the program segments. According to the ELF specification, these are all that a program loader can rely upon being present in an ELF file.

Note

This option can have unexpected effects if -elf is not specified on the command line. See Example 6-2 on page 6-5 for an example of correct usage.

-nodebug

This option does not put debug information in the output files. This is the default for binary images. If -nodebug is specified, it affects all output formats. It overrides the -text -g option.

Caution

This option can have unexpected effects if -elf is not specified on the command line. The options in Example 6-1 on page 6-5 produce a text file because no output format has been specified.

Example 6-1 Text output

```
fromelf -nodebug image -o image_nodb.axf
```

To get ELF format output use the options shown in Example 6-2

Example 6-2 ELF output

```
fromelf -elf -nodebug image.axf -o image_ndb.axf
```

- | | |
|--------------------|--|
| -vsn | This option displays fromELF version information. |
| text_output_format | <p>Use a text specification, for example -text -c, to display image information in text format. You can decode an ELF image or ELF object file using this option. This is the default. If no text or code output format is specified, -text is assumed.</p> <p>If <i>output_file</i> is not specified with the -o option, the information is displayed on stdout.</p> <p>If a specific text category is not specified, the default is to output header information.</p> <p>If specified, the text category consists of one or more of the following:</p> <ul style="list-style-type: none"> a Prints the global and static data addresses (including addresses for structure and union contents). This option can only be used on files containing debug information. Use the -select option to output a subset of the data addresses. c Disassembles code d Prints contents of the data sections g Prints debug information r Prints relocation information s Prints the symbol table t Prints the string table(s) v Prints detailed information on each segment and section header of the image z Prints the code and data sizes. |

The category selectors can be specified as one of:

- individual options, `-text -c -d`
- a single concatenated string, `-text -cd`
- category selectors only, `-c -d`
- multiple characters following a slash character, `-text/cd`.

If an output format is not specified, the default output format of `-text` is used and the individual category selectors are recognized. If another output format is specified, the selectors are ignored.

code_output_format This option selects the binary or ELF output file options.

code_output_format can be one of:

- `-bin` Plain binary. You can split output from this option into multiple files with the *memory_config* option.
- `-ihf` Extended Intellec Hex format. You can specify the base address of the output with the `-base` option. (This option is deprecated and will be removed from future versions of the product.)
- `-m32` Motorola 32-bit format (32-bit S-records). You can specify the base address of the output with the `-base` option.
- `-i32` Intel Hex-32 format. You can specify the base address of the output with the `-base` option.
- `-vhx` Byte Oriented (Verilog Memory Model) Hex format. This format is suitable for loading into the memory models of HDL simulators. You can split output from this option into multiple files with the *memory_config* option.
- `-elf` ELF format (resaves as ELF). This can be used to convert a debug ELF image into a no-debug ELF image.

If you use fromELF to convert an ELF image containing multiple load regions to a binary format using any of the `-bin`, `-ihf`, `-m32`, `-i32`, or `-vhx` options, fromELF creates an output directory named *output_file* and generates one binary output file for each load region in the input image. fromELF places the output files in the *output_file* directory.

`-base n` This option specifies the base address of the output for Motorola S-record, Intel Hex, and Extended Intellec Hex file formats. This option is available only if `-m32`, `-i32`, or `-ihf` is specified as the output format.

You can specify the base address as either:

- a decimal value, for example `-base 0`
- as a hexadecimal value, for example `-base 0x8000`.

All addresses encoded in the output file start at the base address *n*. If you do not specify a `-base` option, the base address is taken from the load region address.

Note

If multiple load regions are present, the `-base` value is used for each output file. That is, it overrides all load region addresses.

memory_config

This option outputs multiple files for multiple memory banks. This option is available only if `-vhx` or `-bin` is specified as the output format.

The format of *memory_config* is `-widthxbanks` where:

- width* is the width of memory in the target memory system (8-bit, 16-bit, 32-bit, or 64-bit).
- banks* specifies the number of memory banks in the target memory system.

Valid configurations are:

-8x1
-8x2
-8x4
-16x1
-16x2
-32x1
-32x2
-64x1

fromELF uses the last specified configuration if more than one configuration is specified. If this option is selected for any format other than `-bin` or `-vhx` it is ignored.

If the image has one load region, fromELF generates *bank* files with the following naming conventions:

- If there is one memory bank (*banks* is 1) the output file is named by the `-o output_file` argument.
- If there are multiple memory banks (*bank*>1), fromELF generates *banks* number of files starting with *output_file*₀ and finishing with *output_file* *bank*-1. For example:

```
fromelf -vhx -8x2 test.axf -o test
```

generates two files named *test*₀ and *test*₁.

If the image has multiple load regions, fromELF creates a directory named *output_file* and generates bank files for each load region named *load_region0* to *load_region banks-1*.

The memory width specified by *width* controls the size of the chunk of information read from the image and written to a file. The first chunk read is allocated to the first file (*output_file0*), the next chunk is allocated to the next file. After a chunk is allocated to the last file, allocation begins again with the first file (that is, the allocation is modulo based on the number of files). For example:

For a *memory_config* of -8x4

```
byte0 -> file0
byte1 -> file1
byte2 -> file2
byte3 -> file3
byte4 -> file0
...
```

For a *memory_config* of -16x2

```
halfword0 -> file0
halfword1 -> file1
halfword3 -> file0
...
```

-output output_file

This option specifies the name of the output file, or the name of the output directory if multiple output files are created (see the description of *text_output_format* and *code_output_format* for more information). Specifying the output file is optional with the *-text* output option and mandatory with all other outputs.

input_file

This option specifies the ELF file to be translated.

fromELF accepts only ARM-executable ELF files and ARM object ELF files (.o).

6.3 Examples of fromELF usage

This section contains examples of using fromELF to change image format or extract information from an ELF file.

Note

If you are using a wildcard character on Unix, for example, *, ? or ~, you must enclose the options in quotes to prevent the shell expanding the selection.

For example, enter '`* ~*.*`' instead of `* ~*.*`.

6.3.1 Producing a plain binary file

To convert an ELF file to a plain binary (.bin) file, enter:

```
fromelf -bin -o outfile.bin infile.axf
```

6.3.2 Disassembly

To produce a listing to stdout that contains the disassembled version of an ELF file, enter:

```
fromelf -c infile.axf
```

To produce a plain-text output file that contains the disassembled version of an ELF file and the symbol table, enter:

```
fromelf -c -s -o outfile.lst infile.axf
```

6.3.3 Listing field offsets as assembly language EQUs

To produce an output listing to stdout that contains all the field offsets from all structures in the file inputfile.o, enter:

```
fromelf -fieldoffsets inputfile.o
```

To produce an output file listing to outputfile.a that contains all the field offsets from structures in the file inputfile.o that have a name starting with p, enter:

```
fromelf -fieldoffsets -select p* -o outputfile.a inputfile.o
```

To produce an output listing to outputfile.a that contains all the field offsets from structures in the file inputfile.o with names of tools or moretools, enter:

```
fromelf -fieldoffsets -select tools.*, moretools.* -o outputfile.a inputfile.o
```

To produce an output file listing to `outputfile.a` that contains all the field offsets of structure fields whose name starts with number and are within structure field top in structure tools in the file `inputfile.o`, enter:

```
fromelf -fielddoffsets -select tools.top.number* -o outputfile.a inputfile.o
```

6.3.4 Listing addresses of static data

To list to stdout all the global and static data variables and all the structure field addresses, enter:

```
fromelf -text -a -select * infile.axf
```

Selecting only structures

To produce a text file containing all of the structure addresses in `inputfile.axf` but none of the global or static data variable information, enter:

```
fromelf -text -a -select *.* -o strucaddress.txt infile.axf
```

Selecting only nested structures

To produce a text file containing addresses of the nested structure only, enter:

```
fromelf -text -a -select *.*.* -o strucaddress.txt infile.axf
```

Selecting only variables

To produce a text file containing all of the global or static data variable information in `inputfile.axf` but none of the structure addresses, enter:

```
fromelf -text -a -select *, ~*.* -o strucaddress.txt infile.axf
```

6.3.5 Converting debug to nodebug

To produce a new output file equivalent to using the `-nodebug` option from an ELF file produced with the `-debug` option, enter:

```
fromelf -nodebug -elf -o outfile.axf infile.axf
```


Glossary

American National Standards Institute (ANSI)

An organization that specifies standards for, among other things, computer software.

Angel™

Angel is a program that enables you to develop and debug applications running on ARM-based hardware. Angel can debug applications running in either ARM state or Thumb state.

ANSI

See American National Standards Institute.

API

Application Program Interface.

Architecture

The term used to identify a group of processors that have similar characteristics.

ARM-Thumb Procedure Call Standard (ATPCS)

ARM-Thumb Procedure Call Standard defines how registers and the stack will be used for subroutine calls.

ATPCS

See ARM-Thumb Procedure Call Standard.

Big-endian

Memory organization where the least significant byte of a word is at a higher address than the most significant byte.

BNF

Backus Naur Format. Mathematical notation for defining logical structures.

Byte

A unit of memory storage consisting of eight bits.

Deprecated	A deprecated option or feature is one that you are strongly discouraged from using. Deprecated options and features will not be supported in future versions of the product.
Double word	A 64-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.
DWARF	Debug With Arbitrary Record Format.
ELF	Executable and linking format.
Environment	The actual hardware and operating system that an application will run on.
Executable and linking format	The industry standard binary file format used by RealView Compilation Tools for BREW. ELF object format is produced by the ARM object producing tools such as armcc and armasm. The ARM linker accepts ELF object files and can output either an ELF executable file, or partially linked ELF object.
Execution view	The address of regions and sections after the image has been loaded into memory and started execution.
Heap	The portion of computer memory that can be used for creating new variables.
Host	A computer that provides data and other services to another computer.
Image	An executable file that has been loaded onto a processor for execution.
Input section	Contains code or initialized data or describes a fragment of memory that must be set to zero before the application starts. <i>See also</i> Output sections.
Interrupt	A change in the normal processing sequence of an application caused by, for example, an external signal.
Interworking	Producing an application that uses both ARM and Thumb code.
Library	A collection of assembler or compiler output objects grouped together into a single repository.
Linker	Software that produces a single image from one or more source assembler or compiler output objects.
Little-endian	Memory organization where the least significant byte of a word is at a lower address than the most significant byte.
Load view	The address of regions and sections when the image has been loaded into memory but has not yet started execution.

Output section	<p>A contiguous sequence of input sections that have the same RO, RW, or ZI attributes. The sections are grouped together in larger fragments called regions. The regions will be grouped together into the final executable image.</p> <p><i>See also</i> Region.</p>
PIC	<p>Position Independent Code.</p> <p><i>See also</i> ROPI.</p>
PID	<p>Position Independent Data <i>or</i> the ARM Platform-Independent Development card.</p> <p><i>See also</i> RWPI.</p>
Program image	<p>See Image.</p>
Read Only Position Independent (ROPI)	<p>Code and read-only data addresses can be changed at run-time.</p>
Read Write Position Independent (RWPI)	<p>Read/write data addresses can be changed at run-time.</p>
RealView Compilation Tools (RVCT)	<p>RealView Compilation Tools is a suite of tools, together with supporting documentation and examples, that enable you to write and build applications for the ARM family of RISC processors.</p>
Redirection	<p>The process of sending default output to a different destination or receiving default input from a different source. This is commonly used to output text, that would otherwise be displayed on the computer screen, to a file.</p>
Region	<p>In an Image, a region is a contiguous sequence of one to three output sections (RO, RW, and ZI). A region typically maps onto a physical memory device, such as ROM, RAM, or peripheral.</p> <p><i>See also</i> Root region.</p>
Remapping	<p>Changing the address of physical memory or devices after the application has started executing. This is typically done to allow RAM to replace ROM once the initialization has been done.</p>
Root region	<p>In an image, a region that executes at its load address. A non-root region is a region that must be copied from its load address to its execution address.</p>
ROPI	<p><i>See</i> Read Only Position Independent.</p>
RTOS	<p><i>See</i> Real Time Operating System.</p>
RVCT	<p><i>See</i> RealView Compilation Tools.</p>

RWPI	<i>See</i> Read Write Position Independent.
Sections	A block of software code or data for an Image. <i>See also</i> Input sections.
Software Interrupt (SWI)	An instruction that causes the processor to call a programmer-specified subroutine. Used by ARM to handle semihosting.
Stack	The portion of memory that is used to record the return address of code that calls a subroutine. The stack can also be used for parameters and temporary variables.
SWI	<i>See</i> Software Interrupt.
Target	The actual target processor, (real or simulated), on which the target application is running.
Vector Floating Point (VFP)	A standard for floating-point coprocessors where several data values can be processed by a single instruction.
Veneer	A small block of code used with subroutine calls when there is a requirement to change processor state or branch to an address that cannot be reached in the current processor state.
VFP	<i>See</i> Vector Floating Point.
Volatile	Memory addresses where the contents can change independently of the executing application are described as volatile. These are typically memory-mapped peripherals. <i>See also</i> Memory mapped
Word	A 32-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.
Zero Initialized (ZI)	R/W memory used to hold variables that do not have an initial value. The memory is normally set to zero on reset.
ZI	<i>See</i> Zero Initialized.

Index

The items in this index are listed in alphabetical order, with symbols and numerics appearing at the end. The references given are to page numbers.

A

ARMLIB variable 5-4

C

Command syntax
 fromelf 6-3
 linker

D

Disassembly 6-5

E

ELF 2-3
ELF file format 2-8
Entry point
 specifying to the linker 2-12, 2-13

Execution region

 reuse of veneers 3-12
 veneers 3-12

F

File formats
 ELF 2-8
 fromELF output 6-6
 symdefs 4-7
Files
 steering 4-9
 symdefs 4-5
 via 2-19

H

Hiding and renaming symbols 2-17,
 4-9

I

Images
 regions, overview 3-3
 sections
 overview 3-3
 simple 3-14
 specifying a memory map 3-5
 structure of 3-2
Interworking
 ARM and Thumb 1-2

L

Libraries
 and linker, *see* Linker, libraries
Linker
 code and data sizes 2-16
 debug information
 turning on and off 2-10
 default addresses 2-8, 2-9
 diagnostics 2-18

- entering commands via a file 2-19
 - execution information 3-4
 - help on 2-5, 2-16
 - image
 - construction 2-6
 - entry point 2-12
 - keeping sections 2-13
 - load and execution views 3-4
 - overview 3-2
 - structure 3-2
 - image-related information 2-7
 - information 2-15, 2-18, 2-19
 - libraries
 - defaults 2-15
 - including during link step 5-2
 - linker search path 2-15
 - scanning 2-15
 - load information 3-4
 - local symbols 2-15
 - memory attributes 2-19
 - memory map information 2-2, 3-4
 - messages 2-15, 2-18, 2-19
 - output file 2-5, 2-8
 - output formats 2-5
 - output sections 3-3
 - overview of 2-2
 - partial linking 2-8
 - regions 3-3
 - RO section base address 2-8, 2-9, 3-18
 - RW section base address 2-9
 - software version 2-8
 - sorting input sections 3-8
 - standard output stream 2-18
 - steering files 2-17, 4-9
 - symbols 4-2
 - hiding and renaming 2-17, 4-9
 - used in link step 2-17
 - syntax 2-8
 - undefined symbols 2-19
 - unused sections 2-11, 2-17
 - veneers 2-17
 - version number 2-5
 - via files 2-5, 2-19
 - \$\$ symbols 4-2
- Linker options
- edit 2-17, 4-9
 - elf 2-8
 - entry 2-12
 - errors 2-18
 - first 2-14
 - help 2-8
 - info 2-16
 - keep 2-13
 - last 2-14
 - libpath 2-15, 5-4
 - list 2-18
 - locals 2-15
 - mangled 2-18, 2-19
 - nodebug 2-10
 - nolocals 2-15
 - noremove 2-11
 - noscanlib 2-15
 - output 2-8
 - partial 2-8
 - remove 2-11
 - ro-base 2-8, 2-9
 - ropi 2-9
 - rw-base 2-9
 - rwpi 2-10
 - scanlib 2-15
 - split 2-10
 - strict 2-19
 - symbols 2-17
 - symdefs 2-17
 - unmangled 2-18
 - unresolved 2-19
 - verbose 2-18
 - via 2-19
 - vsn 2-8
 - xreffrom 2-18
- P**
- Position independence 2-9
- S**
- Sections
- aligning 3-9
 - attributes 2-2
 - in region 2-2
 - input 3-3
 - placement of 2-14, 3-8
 - by attribute 3-8
 - FIRST and LAST 3-9
- in region 3-3
 - sorting rules 3-8
 - unused 2-11, 2-17
- Steering files 2-17, 4-9
- Symbols**
- hiding and renaming 4-9
 - in another image 4-5
 - input section-related 4-4
 - linker 2-17, 4-2, 4-9
 - region-related 4-2
 - section-related 4-4
 - undefined 2-19
 - ZI 4-3
 - \$\$ 4-2
- Symdefs file 4-5
- T**
- Thumb
- code 1-2
 - verneer 3-12
- V**
- Variables
- ARMLIB 5-4
- Veneers
- ARM to ARM 3-12
 - ARM to Thumb 3-12
 - size 2-17
 - Thumb to ARM 3-12
- Via files 2-19
- Z**
- ZI symbols 4-3
- Symbols**
- \$\$ symbols 4-2