

# RealView® Developer Kit

Version 2.2

## Linker and Utilities Guide



# RealView Developer Kit

## Linker and Utilities Guide

Copyright © 2005 ARM Limited. All rights reserved.

### Release Information

The following changes have been made to this book.

#### Change History

Date	Issue	Change
April 2005	B	Release for RVDK v2.2

### Proprietary Notice

Words and logos marked with ® or ™ are registered trademarks or trademarks owned by ARM Limited. Other brands and names mentioned herein may be the trademarks of their respective owners.

Neither the whole nor any part of the information contained in, or the product described in, this document may be adapted or reproduced in any material form except with the prior written permission of the copyright holder.

The product described in this document is subject to continuous developments and improvements. All particulars of the product and its use contained in this document are given by ARM in good faith. However, all warranties implied or expressed, including but not limited to implied warranties of merchantability, or fitness for purpose, are excluded.

This document is intended only to assist the reader in the use of the product. ARM Limited shall not be liable for any loss or damage arising from the use of any information in this document, or any error or omission in such information, or any incorrect use of the product.

### Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by ARM and the party that ARM delivered this document to.

### Product Status

The information in this document is final, that is for a developed product.

### Web Address

<http://www.arm.com>

# Contents

## RealView Developer Kit Linker and Utilities Guide

### Preface

About this book .....	vi
Feedback .....	ix

### Chapter 1

#### Introduction

1.1 About RVDK .....	1-2
1.2 About the linker and utilities .....	1-3

### Chapter 2

#### The Linker Command Syntax

2.1 About armlink .....	2-2
2.2 armlink command syntax .....	2-9

### Chapter 3

#### Using the Basic Linker Functionality

3.1 Specifying the image structure .....	3-2
3.2 Section placement .....	3-8
3.3 Optimizations and modifications .....	3-11
3.4 Using command-line options to create simple images .....	3-25
3.5 Using command-line options to handle C++ exceptions .....	3-31
3.6 Getting information about images .....	3-32

<b>Chapter 4</b>	<b>Accessing Image Symbols</b>	
4.1	ARM/Thumb synonyms .....	4-2
4.2	Accessing linker-defined symbols .....	4-3
4.3	Accessing symbols in another image .....	4-7
4.4	Hiding and renaming global symbols .....	4-10
4.5	Using \$Super\$\$ and \$Sub\$\$ to override symbol definitions .....	4-20
4.6	Symbol versioning .....	4-21
<b>Chapter 5</b>	<b>Using Scatter-loading Description Files</b>	
5.1	About scatter-loading .....	5-2
5.2	Formal syntax of the scatter-loading description file .....	5-8
5.3	Examples of specifying region and section addresses .....	5-24
5.4	Equivalent scatter-loading descriptions for simple images .....	5-35
<b>Chapter 6</b>	<b>Creating and Using Libraries</b>	
6.1	About libraries .....	6-2
6.2	Library searching, selection, and scanning .....	6-3
6.3	The ARM librarian .....	6-6
<b>Chapter 7</b>	<b>Using fromelf</b>	
7.1	About fromelf .....	7-2
7.2	fromelf command syntax .....	7-3
7.3	Examples of fromelf usage .....	7-9
	<b>Glossary</b>	

# Preface

This preface introduces the *RealView® Developer Kit (RVDK) v2.2 Linker and Utilities Guide*. It contains the following sections:

- *About this book* on page vi
- *Feedback* on page ix.

## About this book

This book provides reference information for *RealView® Developer Kit* (RVDK). It describes the command-line options to the linker and other ARM® tools in RVDK.

## Intended audience

This book is written for all developers who are producing applications using RVDK. It assumes that you are an experienced software developer and that you are familiar with the ARM development tools.

## Using this book

This book is organized into the following chapters and appendixes:

### **Chapter 1** *Introduction*

Read this chapter for an introduction to the linker and related utilities in RVDK v2.2.

### **Chapter 2** *The Linker Command Syntax*

Read this chapter for an explanation of all command-line options accepted by the linker.

### **Chapter 3** *Using the Basic Linker Functionality*

Read this chapter for details on using linker features and how to create simple images.

### **Chapter 4** *Accessing Image Symbols*

Read this chapter for details on accessing symbols in images.

### **Chapter 5** *Using Scatter-loading Description Files*

Read this chapter for details on using a scatter-loading file to place code and data in memory.

### **Chapter 6** *Creating and Using Libraries*

Read this chapter for an explanation of the procedures involved in creating and accessing library objects.

### **Chapter 7** *Using fromelf*

Read this chapter for a description of the fromelf utility and how you can use it to change image format.

**Glossary** An alphabetically arranged glossary defines the special terms used in this book.

This book assumes that you have installed your ARM software in the default location, for example, on Windows this might be *volume:\Program Files\ARM*. This is assumed to be the location of *install\_directory* when referring to path names, for example, *install\_directory\RVDK\Examples\...* You might have to change this if you have installed your ARM software in a different location.

## Typographical conventions

The following typographical conventions are used in this book:

<code>monospace</code>	Denotes text that can be entered at the keyboard, such as commands, file and program names, and source code.
<u><code>monospace</code></u>	Denotes a permitted abbreviation for a command or option. The underlined text can be entered instead of the full command or option name.
<i>monospace italic</i>	Denotes arguments to commands and functions where the argument is to be replaced by a specific value.
<b><code>monospace bold</code></b>	Denotes language keywords when used outside example code.
<i>italic</i>	Highlights important notes, introduces special terminology, denotes internal cross-references, and citations.
<b>bold</b>	Highlights interface elements, such as menu names. Also used for emphasis in descriptive lists, where appropriate, and for ARM processor signal names.

## Further reading

This section lists publications from both ARM Limited and third parties that provide additional information on developing code for the ARM family of processors.

ARM Limited periodically provides updates and corrections to its documentation. See <http://www.arm.com> for current errata sheets and addenda, and the ARM Frequently Asked Questions.

### ARM publications

This book contains reference information that is specific to development tools supplied with RVDK. Other publications included in the suite are:

- *RealView Developer Kit v2.2 Debugger User Guide* (ARM DUI 0281)
- *RealView Developer Kit v2.2 Compiler and Libraries Guide* (ARM DUI 0282)
- *RealView Developer Kit v2.2 Assembler Guide* (ARM DUI 0283)
- *RealView Developer Kit v2.2 Command Line Reference* (ARM DUI 0284)
- *RealView Developer Kit v2.2 Project Management Guide* (ARM DUI 0291).

In addition, refer to the following documentation for specific information relating to ARM products:

- *ARM Reference Peripheral Specification* (ARM DDI 0062)
- the ARM datasheet or technical reference manual for your hardware device.

### Other publications

This book is not intended to be an introduction to the ARM assembly language, C, or C++ programming languages. Other books provide general information about programming.

For a comprehensive introduction to ARM architecture, see Steve Furber, *ARM system-on-chip architecture* (2nd edition, 2000). Addison Wesley, ISBN 0-201-67519-6.



## Feedback

ARM Limited welcomes feedback on both RVDK and the documentation.

### Feedback on RealView Developer Kit

If you have any problems with RVDK, contact your supplier. To help them provide a rapid and useful response, give:

- your name and company
- the serial number of the product
- details of the release you are using
- details of the platform you are running on, such as the hardware platform, operating system type and version
- a small standalone sample of code that reproduces the problem
- a clear explanation of what you expected to happen, and what actually happened
- the commands you used, including any command-line options
- sample output illustrating the problem
- the version string of the tools, including the version number and build numbers.

### Feedback on this book

If you notice any errors or omissions in this book, send an email to [errata@arm.com](mailto:errata@arm.com) giving:

- the document title
- the document number
- the page number(s) to which your comments apply
- a concise explanation of the problem.

General suggestions for additions and improvements are also welcome.



# Chapter 1

## Introduction

This chapter introduces the ARM® linker, `arm1ink`, and the utility programs, `armar` and `fromelf` provided with *RealView® Developer Kit* (RVDK). It contains the following sections:

- *About RVDK* on page 1-2
- *About the linker and utilities* on page 1-3.

## 1.1 About RVDK

RVDK consists of a suite of tools, together with supporting documentation and examples, that enable you to write applications for the ARM family of *Reduced Instruction Set Computing* (RISC) processors. You can use RVDK to build C, C++, and ARM assembly language programs.

---

**Note**

Support for C++ may be obtained as an upgrade on certain versions of RVDK, by means of an appropriate FLEXlm license.

---

The RVDK toolkit consists of the following major components:

- command-line development tools
- utilities
- supporting software.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

This book describes the ARM linker, `arm1ink`, the image conversion utility, `fromelf`, and the ARM librarian, `armar`, provided with RVDK.

See *ARM publications* on page viii for a list of the other books in the RVDK documentation suite that give information on the ARM assembler, compiler, and supporting software.

---

**Note**

`arm1ink` does not support certain features relative to RVDS. See Chapter 2 *The Linker Command Syntax* for more information.

---

### 1.1.1 Using the examples

This book references examples provided with RealView Developer Kit in the main examples directory `install_directory\RVDK\Examples`. See *RealView Developer Kit Getting Started Guide* for a summary of the examples provided.

## 1.2 About the linker and utilities

This section gives an overview of:

- *The linker*
- *fromelf* on page 1-4
- *armar* on page 1-4

### 1.2.1 The linker

The ARM linker combines the contents of one or more object files with selected parts of one or more object libraries to produce:

- an *Executable and Linking Format* (ELF) executable image
- a partially linked ELF object that can be used as input to a subsequent link step
- a shared object, compatible with the *Base Platform ABI for the ARM Architecture* [BPABI] or System V release 4 (SVr4) specification from Sun, or a BPABI or SVr4 executable file.

The linker automatically selects the appropriate standard C or C++ library variants to link with, based on the build attributes of the objects it is linking.

The linker can link ARM code and Thumb® code, and automatically generates interworking veneers to switch processor state when required. The linker also automatically generates inline veneers or long branch veneers, where required, to extend the range of branch instructions.

The linker supports command-line options that enable you to specify the locations of code and data within the system memory map. Alternatively, you can use scatter-loading description files to specify the memory locations, at both load and execution time, of individual code and data sections in your output image. This enables you to create complex images spanning multiple memories.

The linker supports Read/Write data compression to minimize ROM size.

The linker can provide feedback, for the next time a file is compiled, to inform the compiler about unused functions. These are placed in their own sections on subsequent compilations for future elimination by the linker.

The linker can perform common section elimination and unused section elimination to reduce the size of your output image. In addition, the linker enables you to:

- produce debug and reference information about linked files
- generate a static callgraph and list the stack usage over it
- control the contents of the symbol table in output images

- show the sizes of code and data in the output.

In addition to unused common sections, the linker can perform elimination of common groups or sections. The Comdat (ELF name for Common) group elimination process uses the same criteria as the common section removal mechanism.

The linker generates only ELF format outputs. To convert ELF images to other formats, such as plain binary for loading into ROM, use `fromelf`. See *fromelf*.

See Chapter 2 *The Linker Command Syntax* for more information on the ARM linker and all command-line options.

### 1.2.2 fromelf

`fromelf` is the ARM image conversion utility. It accepts ELF format input files and converts them to a variety of output formats, including:

- plain binary
- Motorola 32-bit S-record format
- Intel Hex-32 format
- Byte Oriented (Verilog Memory Model) Hex format.

The utility can also produce textual information about the input file and disassemble code.

---

#### **Note**

`fromelf` does not support certain features. See Chapter 7 *Using fromelf* for more information.

---

### 1.2.3 armar

The ARM librarian `armar` enables you to collect and maintain sets of ELF files in standard format ar libraries. You can pass libraries to the linker in place of several ELF object files. See *The ARM librarian* on page 6-6 for more information.

# Chapter 2

## The Linker Command Syntax

This chapter describes the full command syntax for the ARM® linker, `armlink`, provided with *RealView® Developer Kit* (RVDK). It contains the following sections:

- *About armlink* on page 2-2
- *armlink command syntax* on page 2-9.

## 2.1 About armlink

This section describes:

- *Input to armlink*
- *Output from armlink*
- *Summary of linker options* on page 2-4.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

### 2.1.1 Input to armlink

Input to armlink consists of:

- One or more object files in ELF object format. This format is described in the ARM ELF specification. See *ARM publications* on page viii for more information.
- One or more libraries created by armar as described in Chapter 6 *Creating and Using Libraries*.
- A symbol definitions file.

### 2.1.2 Output from armlink

Output from a successful invocation of armlink is one of the following:

- an executable image in ELF executable format
- a shared object
- a partially-linked object in ELF object format
- a relocatable ELF image.

For simple images, ELF executable files contain segments that are approximately equivalent to RO and RW output sections in the image. An ELF executable file also has ELF sections that contain the image output sections.

You can use *fromelf* to convert an executable image in ELF executable format to other file formats. See Chapter 7 *Using fromelf* for more information.

#### Constructing an executable image

When you use the linker to construct an executable image, it:

- resolves symbolic references between the input object files



- extracts object modules from libraries to satisfy otherwise unsatisfied symbolic references
- sorts input sections according to their attributes and names, and merges similarly attributed and named sections into contiguous chunks
- removes unused sections and unused virtual functions
- eliminates duplicate common groups and common code, data, and debug sections
- organizes object fragments into memory regions according to the grouping and placement information provided
- relocates relocatable values
- generates an executable image.

Load regions typically exist in the system memory map at reset, for example, in ROM, or after the image is loaded into the target by a debugger. As part of executing the image, some regions might have to be moved from their load addresses to their execution addresses. The memory map of an image, therefore, has the following distinct views:

<b>Load view</b>	The memory view when the program and data are first loaded.
<b>Execution view</b>	The memory view after code is moved to its normal execution location.

When describing a memory map:

- The term *root region* is used to describe a region that has the same load and execution addresses.
- Load regions are equivalent to ELF segments.

See *Specifying the image structure* on page 3-2 for more information on the image hierarchy.

## Constructing a partially-linked object

When you use the linker to construct a partially-linked object, it:

- eliminates duplicate copies of debug sections
- minimizes the size of the symbol table
- leaves unresolved references unresolved
- generates an object that can be used as an input to a subsequent link step.

---

**Note**

---

If you use partial linking, you cannot refer to the component objects by name in a scatter-loading description file.

---

### 2.1.3 Summary of linker options

This section gives a summary of linker command-line options. The options are arranged in alphabetical order within functional groups.

#### Accessing help and information

To get information on the available command-line options use:

`--help`

To get the tool version number use:

`--vsn`

#### Specifying an input file list

To define input files passed to the linker use:

`--input-file_list`  
`--libpath pathlist`  
`--scanlib | --no_scanlib`  
`--userlibpath pathlist`

You can use the POSIX option `--` to specify that all subsequent arguments are not treated as command switches. For example, to link a file named `-ifile_1` type:

`armlink -- -ifile_1`

#### Controlling linker behavior

To define how objects are linked together use:

`--match crossmangled`  
`--strict`  
`--unresolved symbol`

#### Specifying the output type and the output filename

Name the output file using the following option:

`--output file`

Use the following option to create a partially-linked object instead of an executable image:

```
--partial
```

Use the following option to create a relocatable object:

```
--reloc
```

## Specifying memory map information for the image

Use the following options to specify simple memory maps:

```
--fpic
--ro-base address
--rw-base address
--ropi
--rwpi
--rosplit
--split
```

Alternatively, for more complex images, use the options:

```
--pad num
--scatter file
```

If you use the `--scatter` option, you must provide a scatter-loading description file and a reimplementaion of the `__user_initial_stackheap()` function.

Scatter-loading is described in Chapter 5 *Using Scatter-loading Description Files*.

The memory map options cannot be used for partial linking because they specify the memory map of an executable image.

## Controlling debug information

To control debug information in the image use:

```
--debug | --no_debug
--no_bestdebug | --bestdebug
```

## Controlling image contents

To control miscellaneous factors affecting the image contents use:

```
--cppinit symbol
--datacompressor on|off|list|id
--edit file-list
--entry location
```

```
--exceptions | --no_exceptions
--exceptions_tables=action
--fini symbol
--first section-id
--init symbol
--inline
--keep section-id
--last section-id
--locals | --no_locals
--remove | --no_remove
--startup symbol
--symver_script file
--symver_soname
--tailreorder
--vfemode=mode
```

## Controlling veneer generation

To control how veneers are generated use:

```
--no_inlineveneer
--no_veneershare
```

## Generating image-related information

To control how to extract and present information about the image use:

```
--callgraph
--feedback file
--info topics
--list_mapping_symbols
--mangled | --unmangled
--map
--symbols
--symdefs file
--xref
--xrefdbg
--xreffrom object(section)
--xrefto object(section)
```

With the exception of `--callgraph`, the linker prints the information you request on the standard output stream, `stdout`, by default. You can redirect the information to a text file using the `--list` command-line option.

For `--callgraph`, the information is saved as an HTML file named `output_name.htm`. This is saved in the same directory as the generated image.

## Controlling linker diagnostics

To control how the linker emits diagnostics:

```
--diag_style arm|ide  
--diag_suppress taglist  
--diag_warning taglist  
--errors file  
--list file  
--verbose
```

## Using a via file

Use the following option to specify a via file containing additional command-line arguments to the linker:

`--via file`

See the section on via files in *RealView Developer Kit v2.2 Compiler and Libraries Guide* for more information.

## 2.2 armlink command syntax

This section describes the syntax and options of the `armlink` command.

---

### Note

---

For command-line arguments that use parentheses, you might have to escape the parentheses characters with a backslash (\) character.

Specify command-line keywords using double dashes `--` (for example, `--partial`). The single-dash command-line options used in previous versions of ADS and *RealView® Compilation Tools* (RVCT) are still supported for backwards-compatibility but are deprecated.

---

The linker command syntax is:

```
armlink [help-options] [input-file-list] [linker-control-options]
[output-format] [memory-map-options] [debug-options] [image-contents-options]
[veneer-generation-options] [Byte Addressing mode] [image-info-options]
[diagnostics-options] [via-file] [ADS-legacy-option]
```

The rest of this chapter describes these options in more detail:

- *Accessing help and information* on page 2-11
- *Specifying an input file list* on page 2-11
- *Controlling linker behavior* on page 2-13
- *Specifying the output type and the output filename* on page 2-14
- *Specifying memory map information for the image* on page 2-15
- *Controlling debug information* on page 2-17
- *Controlling image contents* on page 2-18
- *Controlling veneer generation* on page 2-26
- *Generating image-related information* on page 2-26
- *Controlling linker diagnostics* on page 2-30
- *Using a via file* on page 2-31
- *Controlling compatibility with legacy objects* on page 2-31.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

Compilation tools feature restrictions apply to both the command line and to `#pragma` entries.

---

**Note**

---

Support for C++ may be obtained as an upgrade on certain versions of RVDK, by means of an appropriate FLEXlm license.

---



## 2.2.1 Accessing help and information

To get information on the available command-line options and the tool version number use:

- |                     |  |
|---------------------|--|
| <code>--help</code> | Prints a summary of some commonly used command-line options. |
| <code>--vsn</code>  | Displays the linker version information and license details. |

## 2.2.2 Specifying an input file list

These options define the input files passed to the linker:

*input-file-list*

This is a space-separated list of objects or libraries.

You can use libraries in the input file list in the following ways:

- Specify a library to be added to the list of libraries that is used to extract members if they resolve any non-weak unresolved references. For example, specify `mystring.lib` in the input file list. The standard C or C++ libraries are added to this list implicitly by the linker when it scans the default library directories and selects the closest matching library variants available.

### ———— Note ————

Members from the libraries in this list are added to the image only when they resolve an unresolved non-weak reference.

For more information see *Library searching, selection, and scanning* on page 6-3.

- Specify particular members to be extracted from a library and added to the image as individual objects. For example, specify `mystring.lib(strcmp.o)` in the input file list.

The linker processes the input file list in the following order:

- Objects are added to the image unconditionally.
- Members selected from libraries using patterns are added to the image unconditionally, as if they were objects. For example, the following command unconditionally adds all `a*.o` objects and `stdio.o` from `mylib`:  
`armlink main.o mylib(stdio.o) mylib(a*.o)`
- The standard C or C++ libraries are added to the list of libraries that are later used to resolve any remaining non-weak unresolved references.

**--libpath *pathlist***

This option specifies a list of paths that are used to search for the ARM standard C and C++ libraries.

The default path for the parent directory containing the ARM libraries is specified by the RVCT22LIB environment variable. Any paths specified here override the path specified by RVCT22LIB.

*pathlist* is a comma-separated list of paths that are only used to search for required ARM libraries. Do not include spaces between the comma and the path name when specifying multiple path names, for example, *path1,path2,path3,...,pathn*.

This list must end with the parent directory of the ARM library directories `armlib` and `cpplib`.

———— **Note** —————

This option does not affect searches for user libraries. Use `--userlibpath` instead.

See *Library searching, selection, and scanning* on page 6-3 for more information on including libraries.

**--scanlib** This option enables scanning of default libraries (the standard ARM C and C++ libraries) to resolve references. This is the default.

**--no\_scanlib** This option disables the scanning of default libraries.

**--userlibpath *pathlist***

This option specifies a list of paths that are used to search for user libraries.

*pathlist* is a comma-separated list of paths that are used to search for the required libraries. Do not include spaces between the comma and the path name when specifying multiple path names, for example, *path1,path2,path3,...,pathn*.

See *Library searching, selection, and scanning* on page 6-3 for more information on including user libraries.

### 2.2.3 Controlling linker behavior

These options control how objects are linked:

`--match crossmangled`

This option instructs the linker to match the following combinations together:

- a reference to an unmangled symbol with the mangled definition
- a reference to a mangled symbol with the unmangled definition.

Libraries and matching operate as follows:

- If the library members define a mangled definition, and there is an unresolved unmangled reference, the member is loaded to satisfy it.
- If the library members define an unmangled definition, and there is an unresolved mangled reference, the member is loaded to satisfy it.

#### ———— **Note** ————

This option has no effect if used with partial linking. The partial object contains all the unresolved references to unmangled symbols, even if the mangled definition exists. Matching is done only in the final link step.

`--strict`

This option instructs the linker to report conditions that might result in failure as errors, rather than warnings. An example of such a condition is taking the address of an interworking function from a non-interworking function.

`--strict_relocations`

This option instructs the linker to report instances of obsolete and deprecated relocations. For example:

Error: L6810E: Relocation 8 in section relocs from object et5ae.o is of obsolete type R\_ARM\_SWI24.

Relocation errors and warnings are most likely to occur if you are linking object files built with previous versions of the ARM tools.

This option enables you to ensure ABI compliance of objects. It is off by default, and deprecated and obsolete relocations are handled silently by the linker.

`--unresolved symbol`

This option matches each reference to an undefined symbol to the global definition of *symbol*. *symbol* must be both defined and global, otherwise it appears in the list of undefined symbols and the link step fails. This

option is particularly useful during top-down development, because it enables you to test a partially-implemented system by matching each reference to a missing function to a dummy function.

## 2.2.4 Specifying the output type and the output filename

Specify the format and the name of the output file using the following options:

`--output file`

This option specifies the name of the output file. The file can be either a partially-linked object or an executable image. If the output filename is not specified, the linker uses the following defaults:

`__image.axf`            if the output is an executable image

`__object.o`            if the output is a partially-linked object.

If *file* is specified without path information, it is created in the current working directory. If path information is specified, then that directory becomes the default output directory.

`--partial`    This option creates a partially-linked object instead of an executable image.

`--reloc`      This option makes *relocatable* ELF images (see the *ARM ELF specification* for more information).

A relocatable image has a dynamic segment that contains relocations that can be used to relocate the image post link-time. Examples of post link-time relocation include advanced ROM construction and dynamic loading at runtime.

If the image is loaded at its link-time address, the relocatable image produced by the linker does not require the relocations to be processed and debug data for the image is valid. Loading the image at a different address to the link-time address and processing the relocations, however, invalidates any debug data present in the image.

Used on its own, `--reloc` makes an image similar to Simple type 1 where the load region attribute is set to RELOC (see *Type 1: one load region and contiguous output regions* on page 3-26 for details).

## 2.2.5 Specifying memory map information for the image

Use the following options to specify memory maps:

`--ro-base address`

This option sets both the load and execution addresses of the region containing the RO output section at *address*. *address* must be word-aligned. If this option is not specified, and no scatter load file is specified, the default RO base address is 0x8000.

`--rw-base address`

This option sets the execution addresses of the region containing the RW output section at *address*. *address* must be word-aligned.

`--ropi`

This option makes the load and execution region containing the RO output section position-independent. If this option is not used the region is marked as absolute. Usually each read-only input section must be read-only position-independent. If this option is selected, the linker:

- checks that relocations between sections are valid
- ensures that any code generated by `arm1ink` itself, such as interworking veneers, is read-only position-independent.

`--rwpi`

This option makes the load and execution region containing the RW and ZI output section position-independent. If this option is not used the region is marked as absolute. This option requires a value for `--rw-base`. If `--rw-base` is not specified, `--rw-base 0` is assumed. Usually each writable input section must be read-write position-independent.

If this option is selected, the linker:

- checks that the PI attribute is set on input sections to any read-write execution regions
- checks that relocations between sections are valid
- generates static base-relative entries in the table `Region$$Table`, which is used when regions are copied, decompressed, or initialized.

`--fpic`

This option enables you to link position-independent code, that is, code that has been compiled using the `/fpic` qualifier. Relative addressing is only implemented when your code makes use of System V shared libraries.

- split** This option splits the default load region, that contains the RO and RW output sections, into the following load regions:
- One containing the RO output section. The default load address is 0x8000, but a different address can be specified with the `--ro-base` option.
  - One containing the RW and ZI output sections. The load address is specified with the `--rw-base` option. This option requires a value for `--rw-base`. If `--rw-base` is not specified, `--rw-base 0` is assumed.

Both regions are root regions.

- rosplit** This option splits the default load region into two RO output sections, one for code and one for data.

- pad *num*** This option enables you to set a value for padding bytes. The linker assigns this value to all padding bytes inserted in load or execution regions.
- num* is an integer, which can be given in hexadecimal format. For example, setting *num* to 0xFF might help to speed up ROM programming time. If *num* is greater than 0xFF, then the padding byte is set to (char)*num*.

---

**Note**

---

Padding is only inserted:

- within load regions. No padding is present between load regions.
  - between fixed execution regions (in addition to forcing alignment). Padding is not inserted up to the maximum length of a load region unless it has a fixed execution region at the top.
  - between sections to ensure that they conform to alignment constraints.
- 

**--scatter *file***

This option creates the image memory map using the scatter-loading description contained in *file*. The description provides grouping and placement details of the various regions and sections in the image. See Chapter 5 *Using Scatter-loading Description Files*.

The `--scatter` option is mutually exclusive with the use of any of the memory map options `--ro-base`, `--rw-base`, `--ropi`, `--rwpi`, `--rosplit`, or `--split`, and with `--reloc`, `--startup`, and `--partial`.

---

**Note**

---

You must reimplement the stack and heap initialization function `__user_initial_stackheap()` if you use this option.

---

## 2.2.6 Controlling debug information

These options control debug information in the image:

- |                         |  |
|-------------------------|--|
| <code>--debug</code>    | This option includes debug information in the output file. The debug information includes debug input sections and the symbol and string table. This is the default.   |
| <code>--no_debug</code> | This option turns off the inclusion of debug information in the output file. The ELF image is smaller, but you cannot debug it at the source level. The linker discards any debug input section it finds in the input objects and library members, and does not include the symbol and string table in the image. This only affects the image size as loaded into the debugger and has no effect on the size of any resulting binary image that is downloaded to the target. |

If you are creating a partially-linked object rather than an image, the linker discards the debug input sections it finds in the input objects, but does produce the symbol and string table in the partially-linked object.

---

**Note**

---

Do not use `--no_debug` if a `fromelf --fieldoffsets` step is required. If your image is produced without debug information, `fromelf`:

- cannot translate the image into other file formats
  - cannot produce a meaningful disassembly listing.
- 

`--no_bestdebug`

This option selects sections without reference to debug information, that is, it chooses the smallest sections.

`--no_bestdebug` is the default to ensure that the code and data of the final image are the same regardless of whether you compile for debug or not.

Use the option `--bestdebug` to select sections with the best debug view. Be aware that the code and data of the final image might not be the same when building with or without debug.

For details, see generating debug information in the chapter describing how to use the ARM compiler in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

## 2.2.7 Controlling image contents

These options control miscellaneous factors affecting the image contents:

`--datacompressor on|off`

RW data compression is enabled by default to minimize ROM size. Use `--datacompressor off` to turn off RW data compression.

`--datacompressor list|id`

These options enable you to specify one of the supplied algorithms for RW data compression:

- Use `--datacompressor list` to get a list of data compressors available to the linker.
- If you do not specify a data compression algorithm, the linker chooses the most appropriate one for you automatically. In general, it is not necessary to override this choice. For more information see *RW data compression* on page 3-16.

If you do want to override the linker, use `--datacompressor id` to specify a data compression algorithm. Specifying a compressor adds a decompressor to the code area. If the final image does not have compressed data, the decompressor is not added.

`--edit file-list`

This option enables you to specify steering files containing commands to edit the symbol tables in the output binary. You can specify commands in a steering file to:

- Hide global symbols. Use this option to hide specific global symbols in object files. The hidden symbols are not publicly visible.
- Rename global symbols. Use this option to resolve symbol naming conflicts.

See *Hiding and renaming global symbols* on page 4-10 for more information on steering file syntax.

When you are specifying more than one steering file, the syntax can be either of the following:

```
armlink --edit file1 --edit file2 --edit file3
```

```
armlink --edit file1,file2,file3
```

Do not include spaces between the comma and the filenames.



**--entry *location***

This option specifies the unique initial entry point of the image. The image can contain multiple entry points, but the initial entry point specified using this command is stored in the executable file header for use by the loader. There can be only one occurrence of this option on the command line. ARM debuggers, for example, RealView Debugger or AXD, use this entry address to initialize the *program counter* (PC) when an image is loaded. The initial entry point must meet the following conditions:

- the image entry point must lie within an execution region
- the execution region must be non-overlay, and must be a root execution region (load address == execution address).

Replace *location* with one of the following:

***entry\_address***

A numerical value, for example:

```
--entry 0x0
```

***symbol*** This option specifies an image entry point as the address of *symbol*, for example:

```
--entry reset_handler
```

***offset+object(section)***

This option specifies an image entry point as an *offset* inside a *section* within a particular *object*, for example:

```
--entry 8+startup.o(startupseg)
```

There must be no spaces within the argument to --entry. The input section and object names are matched without case-sensitivity. You can use the following simplified notation:

- *object(section)* if offset is zero.
- *object* if there is only one input section. armlink generates an error message if there is more than one input section in *object*.

**--exceptions** This option enables the final image to contain exception tables. This is the default.

See *Using command-line options to handle C++ exceptions* on page 3-31 for more information.

`--no_exceptions`

This option forces the linker to generate an error message if any exceptions sections are present in the image after unused sections have been eliminated. Use this option to ensure that your code is exceptions free.

See *Using command-line options to handle C++ exceptions* on page 3-31 for more information.

`--exceptions_tables=action`

Use this option to specify how exception tables are generated for legacy objects. Replace *action* with one of:

*nocreate* The linker does not create exception tables for legacy objects. This is the default.

*unwind* The linker creates an unwinding table for each section in your image that does not already have an exception table.

*cantunwind*

The linker creates a nounwind table for each section in your image that does not already have an exception table.

See *Using command-line options to handle C++ exceptions* on page 3-31 for more information.

`--first section-id`

This option places the selected input section first in its execution region. This can, for example, place the section containing the vector table first in the image. Replace *section-id* with one of the following:

*symbol* Selects the section that defines *symbol*. You must not specify a symbol that has more than one definition, because only one section can be placed first. For example:

`--first reset`

*object(section)*

Selects *section* from *object*. There must be no space between *object* and the following open parenthesis. For example:

`--first init.o(init)`

*object* Selects the single input section in *object*. If you use this short form and there is more than one input section, `arm1link` generates an error message. For example:

`--first init.o`

---

**Note**

---

When using scatter-loading, use +FIRST in the scatter-loading description file instead.

---

Using `--first` cannot override the basic attribute sorting order for output sections in regions that places RO first, RW second, and ZI last. If the region has an RO section, an RW or a ZI section cannot be placed first. If the region has an RO or RW section, a ZI section cannot be placed first.

Two different sections cannot both be placed first in the same execution region, so only one instance of this option is permitted.

**--last *section-id***

This option places the selected input section last in its execution region. For example, this can force an input section that contains a checksum to be placed last in the RW section. Replace *section-id* with one of the following:

*symbol*     Selects the section that defines *symbol*. You must not specify a symbol that has more than one definition, because more than one section cannot be placed last. For example:

`--last checksum`

*object(section)*

Selects the *section* from *object*. There must be no space between *object* and the following open parenthesis. For example:

`--last checksum.o(check)`

*object*     Selects the single input section from *object*. If there is more than one input section in *object*, `armlink` generates an error message.

---

**Note**

---

When using scatter-loading, use +LAST in the scatter-loading description file instead.

---

Using `--last` cannot override the basic attribute sorting order for output sections in regions that places RO first, RW second, and ZI last. If the region has a ZI section, an RW section cannot be placed last. If the region has an RW or ZI section, an RO section cannot be placed last.

Two different sections cannot both be placed last in the same execution region, so only one instance of this option is permitted.

`--remove --remove (RO/RW/ZI/DBG)`

This option enables unused section elimination on the input sections to remove unused sections from the image. An input section is considered used if it contains the image entry point, or if it is referred to from a used section. See also *Unused section elimination* on page 3-12.

---

**Note**

---

You must take care to avoid reset code or exception handlers accidentally being removed when using `--remove`. Use the `--keep` option to identify exception handlers or use the `ENTRY` directive to label them as entry points.

---

You can use section attribute qualifiers for more precise control of the unused section elimination process. If a qualifier is used, it can be one or more of the following:

RO	Remove all unused sections of type RO.
RW	Remove all unused sections of type RW.
ZI	Remove all unused sections of type ZI.
DBG	Remove all unused sections of type DEBUG.

The qualifiers can appear in any case and order, but must be enclosed in parentheses (), and must be separated by a slash /.

The default is `--remove (RO/RW/ZI/DBG)`.

If no section attribute qualifiers are specified, all unused sections are eliminated. `--remove` is equivalent to `--remove (RO/RW/ZI/DBG)`.

`--no_remove` This option disables unused section elimination on the input sections. This retains all input sections in the final image even if they are unused.

`--startup symbol`

This option enables the linker to use alternative C libraries with a different startup symbol. If the linker find a definition of `main()` and does not find a reference to or definition of *symbol*, then it adds a reference to *symbol*. By default, *symbol* is set to `__main`.

For more information on using libraries, see the chapter describing the C and C++ libraries in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

`--symver_script file`

This option turns on implicit symbol versioning and enables you to specify *file* as a symbol version script.

See *Symbol versioning* on page 4-21 for more information.

**--symver\_soname**

This option turns on implicit symbol versioning and enables you to version symbols in order to force static binding. Where a symbol has no defined version, the linker uses the SONAME of the file being linked.

This is the default if you are generating a BPABI-compatible executable file but where you do not specify a version script with the option `--symver_script`.

See *Symbol versioning* on page 4-21 and the *Base Platform ABI for the ARM Architecture* [BPABI] for more information.

**--tailreorder**

This option moves tail calling sections immediately before their target, if possible, to optimize the branch instruction at the end of a section. A tail calling section is a section that contains a branch instruction at the end of the section. The branch must have a relocation that targets a function at the start of a section. There are some restrictions to this option. The linker:

- can only move one tail calling section for each tail call target
- cannot move a tail calling section out of its execution region
- does not move tail calling sections before inline veneers.

See *Branch inlining* on page 3-22 for more information on handling tail calling sections.

**--vfemode=mode**

This option specifies how unused virtual functions, and *Runtime Type Information* (RTTI) objects, are eliminated in the linker. Depending on the mode specified, the linker uses extra information about virtual functions and RTTI objects supplied by the compiler to analyze more accurately how such functions are used. Unused functions are then eliminated.

Replace *mode* with one of:

- |     |  |
|-----|--|
| on  | Makes the linker <i>Virtual Function Elimination</i> (VFE) aware. In this mode, the linker chooses force or off mode based on the content of object files. This is the default.<br>This is the same as specifying no VFE option on the command line. |
| off | Use this option to force the linker to ignore any extra information supplied by the compiler. In this mode, the final image is the same as that produced by compiling and linking without VFE awareness.   |

**force** Makes the linker VFE aware and forces the VFE algorithm to be applied. If some of the object files do not contain VFE information, the linker continues with the elimination but displays a warning to alert you to possible errors.

**force\_no\_rtti**

When operating in default mode (using `--vfemode=on`), the linker might remove both unused virtual functions and RTTI objects. Use this option to make the linker VFE aware and force the removal of all RTTI objects whilst retaining all the virtual functions.

See also *Unused virtual function elimination* on page 3-12 for more details.

**--init *symbol***

This option specifies the symbol name that is used to define initialization code. The dynamic linker executes this code when it loads the executable file or shared object.

**--cppinit *symbol***

This option enables the linker to use alternative C++ libraries with a different initialization symbol. By default, *symbol* is set to `__cpp_initialize__aeabi_`.

For more information on using libraries, see the chapter describing the C and C++ libraries in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

**--inline**

This option enables branch inlining to optimize small function calls in your image.

———— **Note** ————

This branch optimization is off by default because enabling it changes the image such that debug information might be incorrect. If enabled, the linker makes no attempt to correct the debug information.

See *Branch inlining* on page 3-22 for more information on controlling branch inlining.

**--keep *section-id***

Specifies input sections that are not to be removed by unused section elimination. See *Specifying an image memory map* on page 3-5. All forms of the *section-id* argument can contain the \* and ? wildcards. You can specify multiple -keep options on the comment line. Replace *section-id* with one of the following:

*symbol* This option specifies that an input section defining *symbol* is to be retained during unused section elimination. If multiple definitions of *symbol* exist, armlink generates an error message. For example, you might use --keep int\_handler. To keep all sections that define a symbol ending in \_handler, use --keep \*\_handler.

***object(section)***

This option specifies that *section* from *object* is to be retained during unused section elimination. For example, to keep the vect section from the vectors.o object use:

```
--keep vectors.o(vect)
```

To keep all sections from the vectors.o object where the first three characters of the name of the section is vec, use:

```
--keep vectors.o(vec*)
```

*object* This option specifies that the single input section from *object* is to be retained during unused section elimination. If you use this short form and there is more than one input section in *object*, armlink generates an error message. For example:

```
--keep dspdata.o
```

To keep the single input section from each of the objects that has a name starting with dsp, use:

```
--keep dsp*.o
```

**--locals** This option instructs the linker to add local symbols to the output symbol table when producing an executable image. This is the default.

**--no\_locals** This option instructs the linker not to add local symbols to the output symbol table when producing an executable image. This is a useful optimization if you want to reduce the size of the output symbol table.

**--fini *symbol***

This option specifies the symbol name that is used to define the entry point for finalization code. The dynamic linker executes this code when it unloads the executable file or shared object.

## 2.2.8 Controlling veneer generation

These options control veneer generation:

`--no_inlineveneer`

This option prevents the generation of inline veneers to give greater control over how the linker places sections. For details, see *Veneer generation* on page 3-19.

`--no_veneershare`

This option prevents the linker sharing veneers. Veneer sharing can cause a significant decrease in image size. For details, see *Veneer generation* on page 3-19.

## 2.2.9 Generating image-related information

These options control how to extract and present information about the image:

`--callgraph` This option creates a static callgraph of functions in HTML format. The callgraph gives definition and reference information for all functions in the image.

---

### Note

Any assembler files must contain PROC/ENDP and FRAME PUSH/POP directives if the linker is to calculate the function stack sizes.

---

For each function `func` the linker lists the:

- processor state for which the function is compiled (ARM or Thumb)
- set of functions that call `func`
- set of functions that are called by `func`
- number of times the address of `func` is used in the image.

In addition, the callgraph identifies functions that are:

- called through interworking veneers
- defined outside the image
- permitted to remain undefined (weak references).

The static callgraph also gives information about stack usage. It lists the:

- size of the stack frame used by each function
- maximum size of the stack used by the function over any call sequence, that is, over any acyclic chain of function calls.



If there is a cycle, or if the linker detects a function with no stack size information in the call chain, + Unknown is added to the stack usage.

The linker reports missing stack frame information if there is no debug frame information for the function.

For indirect functions, the linker cannot reliably determine which function made the indirect call. This might affect how the maximum stack usage is calculated for a call chain.

Use frame directives in assembly language code to describe how your code uses the stack. These directives ensure that debug frame information is present for debuggers to perform stack unwinding or profiling.

For more details on how stack usage is determined, see the chapter describing the directives reference in *RealView Developer Kit v2.2 Assembler Guide*.

#### `--feedback file`

This option generates a feedback file, for the next time a file is compiled, to inform the compiler about unused functions.

When you next compile the file, use the compiler option `--feedback file` to specify the feedback file to use. Unused functions are placed in their own sections for possible future elimination by the linker. For full details on how to use this file see *Linker feedback* on page 3-14.

#### `--info topics`

This option prints information about specified topics, where *topics* is a comma-separated list of topic keywords. A topic keyword can be one of the following:

<code>common</code>	Lists all common sections that were eliminated from the image. Using this option implies <code>--info common,totals</code> .
<code>debug</code>	Lists all rejected input debug sections that were eliminated from the image as a result of using <code>--remove</code> . Using this option implies <code>--info debug,totals</code> .
<code>inline</code>	Gives details of any function that is inlined by the linker, and gives the total number of inlines, as a result of using <code>--inline</code> . For more information on branch inlining see <i>Branch inlining</i> on page 3-22.
<code>sizes</code>	Gives a list of the Code and Data (RO Data, RW Data, ZI Data, and Debug Data) sizes for each input object and library member in the image. Using this option implies <code>--info sizes,totals</code> .

#### tailreorder

Gives details of any tail calling sections that have been moved above their targets, as a result of using `--tailreorder`. For more information on handling tail calling sections see *Branch inlining* on page 3-22.

**totals** Gives totals of the Code and Data (RO Data, RW Data, ZI Data, and Debug Data) sizes for input objects and libraries.

**veneers** Gives details of linker-generated veneers. For more information on veneers see *Veneer generation* on page 3-19.

**unused** Lists all unused sections that were eliminated from the image as a result of using `--remove`.

#### exceptions

Gives details of exception table generation and optimization.

The output from `--info sizes,totals` always includes the padding values in the totals for input objects and libraries.

If you are using RW data compression (the default), or if you have specified a compressor using the `--datacompressor id` option, the output from `--info sizes,totals` includes an entry under Grand Totals to reflect the true size of the image.

#### ————— **Note** —————

Spaces are not permitted between keywords in a list. For example, you can enter `--info sizes,totals` but not `--info sizes, totals`.

For more details on how to use this information see *Getting information about images* on page 3-32.

**--mangled** This option instructs the linker to display mangled C++ symbol names in diagnostic messages, and in listings produced by the `--xref`, `--xreffrom`, `--xref to`, and `--symbols` options. If this option is selected, the linker does not unmangle C++ symbol names. The symbol names are displayed as they appear in the object symbol tables.

**--unmangled** This option instructs the linker to display unmangled C++ symbol names in diagnostic messages, and in listings produced by the `--xref`, `--xreffrom`, `--xref to`, and `--symbols` options.

If this option is selected, the linker unmanages C++ symbol names so that they are displayed as they appear in your source code. This is the default.

**--map** This option creates an image map. The image map contains the address and the size of each load region, execution region, and input section in the image, including linker-generated input sections.

`--symbols` This option lists each local and global symbol used in the link step, and its value.

---

**Note**

---

This does not include mapping symbols. Use `--list_mapping_symbols` to include mapping symbols in the output.

---

`--list_mapping_symbols`

Use this option to include mapping symbols in the output produced by `--symbols`.

In the symbol table, mapping symbols are used to flag transitions between ARM code, Thumb code, and data (see the *ELF for the ARM Architecture* [AAELF] for details).

`--symdefs file`

This option creates a symbol definitions file containing the global symbol definitions from the output image.

By default, all global symbols are written to the `symdefs` file. If *file* already exists, the linker restricts its output to the symbols listed in the existing `symdefs` file.

---

**Note**

---

If you do not want this behavior, be sure to delete any existing `symdefs` file before the link step.

---

If *file* is specified without path information, the linker searches for it in the directory where the output image is being written. If it is not found, it is created in that directory.

You can use the symbol definitions file as input when linking another image. See *Accessing symbols in another image* on page 4-7 for more information.

`--xref` This option lists all cross-references between input sections.

`--xrefdbg` This option lists all cross-references between input debug sections.

`--xreffrom object(section)`

This option lists cross-references from input *section* in *object* to other input sections. This is a useful subset of the listing produced by using `--xref` if you are interested in references from a specific input section. You can have multiple occurrences of this option to list references from more than one input section.

`--xref to object(section)`

This option lists cross-references to input *section* in *object* from other input sections. This is a useful subset of the listing produced by using `--xref` if you are interested in references to a specific input section. You can have multiple occurrences of this option in order to list references to more than one input section.

## 2.2.10 Controlling linker diagnostics

These options control how the linker emits diagnostics:

`--diag_style arm|ide`

This option specifies that all warning and error messages are output in a format that is more compatible with IDEs, such as Microsoft Visual Studio,

The default is `arm`, for example:

```
sct.txt(line 15, col14) Warning: L6314W: No section matches
pattern *(RW).
```

```
Error: L6223E: Ambiguous selectors for foo.o ....
```

```
Finished: 0 information, 1 warning and 1 error messages.
```

Specifying `--diag_style ide` results in:

```
sct.txt(15, 14) Warning: L6314W: No section matches pattern *(RW).
```

```
Armlink: Error: L6223E: Ambiguous selectors for foo.o ....
```

```
Armlink: Finished: 0 information, 1 warning and 1 error messages.
```

`--diag_suppress taglist`

This option disables all diagnostic messages that have the specified tag(s).

This option requires a comma-separated list of diagnostic message numbers that specifies the messages that must be suppressed. For example, to suppress the warning messages that have numbers L6314W and L6305W, use the following command:

```
armlink --diag_suppress 6314,6305 ...
```

`--diag_warning taglist`

This option sets diagnostic messages that have the specified tag(s) to be displayed as warning messages, for example, to downgrade an error message.

This option requires a comma-separated list of diagnostic message numbers that specifies the messages that must be downgraded.

`--errors file`

Redirects the diagnostics from the standard error stream to *file*.

The specified file is created at the start of the link stage. If a file of the same name already exists, it is cleared.

If *file* is specified without path information, it is created in the current directory.

`--list file` Redirects the diagnostics from output of the `--info`, `--map`, `--symbols`, `--xref`, `--xreffrom`, and `--xref` commands to *file*.

The specified file is created when diagnostics are output. If a file of the same name already exists, it is overwritten. However, if diagnostics are not output, a file is not created. In this case, the contents of any existing file with the same name remain unchanged.

If *file* is specified without path information, it is created in the output directory, that is, the directory where the output image is being written.

`--verbose` This option prints detailed information about the link operation, including the objects that are included and the libraries from which they are taken. Because this output is typically quite long, you might want to use this command with the `--errors file` command to redirect the information to *file*.

Use `--verbose` to output diagnostics to `stderr`.

### 2.2.11 Using a via file

Use the following option to specify a via file containing additional command-line arguments to the linker:

`--via file` This option reads a further list of input filenames and linker options from *file*.

You can enter multiple `--via` options on the linker command line. The `--via` options can also be included within a via file.

See *RealView Developer Kit v2.2 Compiler and Libraries Guide* for more information on writing via files.

### 2.2.12 Controlling compatibility with legacy objects

The ABI in *ARM Developer Suite™* (ADS) v1.2 was different and, therefore, legacy ADS objects and libraries are not directly compatible with RVCT. In RVCT v2.2, the compiler supports the option `--apcs /adsabi` to compile code that is compatible with the old ADS ABI (see the chapter on using the ARM compiler in *RealView Developer Kit v2.2 Compiler and Libraries Guide*).

---

**Note**

---

Support for the `--apcs /adsabi` compiler option is deprecated and will be removed in a future release.

---

# Chapter 3

## Using the Basic Linker Functionality

This chapter describes the basic functionality available in the ARM® linker, `arm1ink` provided with *RealView® Developer Kit* (RVDK). It contains the following sections:

- *Specifying the image structure* on page 3-2
- *Section placement* on page 3-8
- *Optimizations and modifications* on page 3-11
- *Using command-line options to create simple images* on page 3-25
- *Using command-line options to handle C++ exceptions* on page 3-31
- *Getting information about images* on page 3-32.

For information about advanced linker functionality, see the descriptions of:

- how to access symbols in Chapter 4 *Accessing Image Symbols*
- how to use scatter-loading in Chapter 5 *Using Scatter-loading Description Files*.

## 3.1 Specifying the image structure

The structure of an image is defined by the:

- number of its constituent regions and output sections
- positions in memory of these regions and sections when the image is loaded
- positions in memory of these regions and sections when the image executes.

This section describes:

- *Building blocks for objects and images*
- *Load view and execution view of an image* on page 3-4
- *Specifying an image memory map* on page 3-5
- *Image entry points* on page 3-6.

### 3.1.1 Building blocks for objects and images

An image, as stored in an executable file, is constructed from a hierarchy of images, regions, output sections, and input sections:

- An image consists of one or more regions. Each region consists of one or more output sections.
- Each output section contains one or more input sections.
- Input sections are the code and data information in an object file.

Figure 3-1 on page 3-3 shows the relationship between regions, output sections, and input sections.



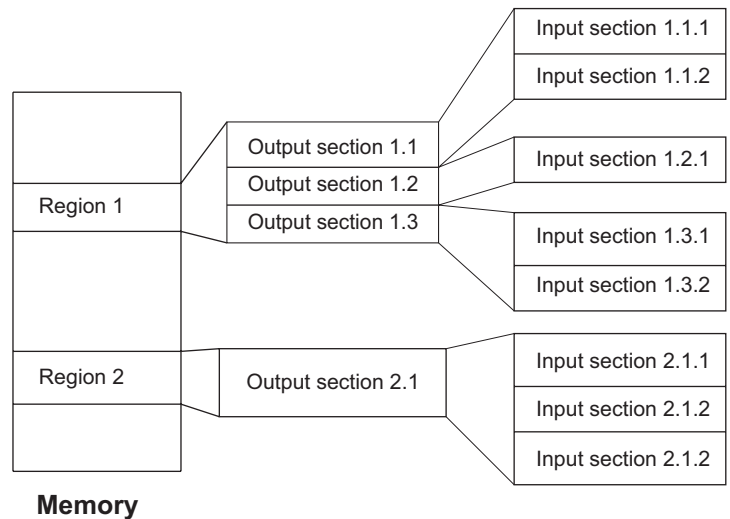
**Figure 3-1 Building blocks for an image**

Figure 3-1 shows the building blocks that make up the image:

### Input sections

An input section contains code or initialized data, or describes a fragment of memory that is not initialized or that must be set to zero before the image can execute. Input sections can have the attributes RO, RW, or ZI. These three attributes are used by `armlink` to group input sections into bigger building blocks called output sections and regions.

### Output sections

An output section is a contiguous sequence of input sections that have the same RO, RW, or ZI attribute. An output section has the same attributes as its constituent input sections. Within an output section, the input sections are sorted according to the rules described in *Section placement* on page 3-8.

### Regions

A region is a contiguous sequence of one to three output sections. The output sections in a region are sorted according to their attributes. The RO output section is first, then the RW output section, and finally the ZI output section. A region typically maps onto a physical memory device, such as ROM, RAM, or peripheral.

3.1.2 Load view and execution view of an image

Image regions are placed in the system memory map at load time. Before you can execute the image, you might have to move some of its regions to their execution addresses and create the ZI output sections. For example, initialized RW data might have to be copied from its load address in ROM to its execution address in RAM.

The memory map of an image has the following distinct views as shown in Figure 3-2.

- Load view

This view describes each image region and section in terms of the address it is located at when the image is loaded into memory, that is, the location before the image starts executing.
- Execution view

This view describes each image region and section in terms of the address it is located at while the image is executing.

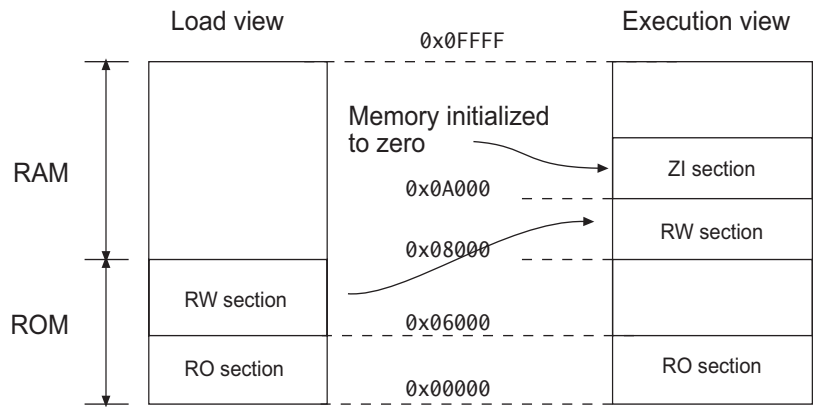


Figure 3-2 Load and execution memory maps

Table 3-1 compares the load and execution views.

Table 3-1 Comparing load and execution views

Load	Description	Execution	Description
Load address	The address where a section, or region is loaded into memory before the image containing it starts executing. The load address of a section or a non-root region can differ from its execution address.	Execution address	The address where a section or region is located while the image containing it is being executed
Load region	A region in the load address space.	Execution region	A region in the execution address space

### 3.1.3 Specifying an image memory map

An image can consist of any number of regions and output sections. Any number of these regions can have different load and execution addresses. To construct the memory map of an image, `arm-link` must have information about:

**Grouping** How input sections are grouped into output sections and regions.

**Placement** Where image regions are to be located in the memory maps.

Depending on the complexity of the memory maps of the image, there are two ways to pass this information to `arm-link`:

#### Using command-line options

The following options can be used for simple cases where an image has only one or two load regions and up to three execution regions:

- `--ro-base`
- `--rw-base`
- `--ropi`
- `--rwpi`
- `--split`
- `--rosplit`

The options listed above provide a simplified notation that gives the same settings as a scatter-loading description for a simple image. For more information, see *Using command-line options to create simple images* on page 3-25.

#### Using a scatter-loading description file

A scatter-loading description file is used for more complex cases where you require complete control over the grouping and placement of image components. To use scatter-loading, specify `--scatter filename` at the command line. This is described in full in Chapter 5 *Using Scatter-loading Description Files*.

### 3.1.4 Image entry points

An entry point in an image is a location where program execution can start. There are two distinct types of entry point:

#### Initial entry point

The *initial* entry point for an image is a single value that is stored in the ELF header file. For programs loaded into RAM by an operating system or boot loader, the loader starts the image execution by transferring control to the initial entry point in the image.

An image can have only one initial entry point. The initial entry point can be, but is not required to be, one of the entry points set by the ENTRY directive.

#### Entry points set by the ENTRY directive

These are entry points that are set in the assembly language sources with the ENTRY directive. In embedded systems, this directive is typically used to mark code that is entered through the processor exception vectors, such as RESET, IRQ, and FIQ.

You can specify multiple entry points in an image with the ENTRY directive. The directive marks the output code section with an ENTRY keyword that instructs the linker not to remove the section when it performs unused section elimination. For C and C++ programs, the `__main()` function in the C library is also an entry point. See *RealView Developer Kit v2.2 Assembler Guide* for more information on the ENTRY directive.

If an embedded image is to be used by a loader, it must have a single initial entry point specified in the header. See *Specifying an initial entry point* for more information.

#### Specifying an initial entry point

You can specify an initial entry point with the `--entry` linker option. You can specify the `--entry` option only once. See the description in *armlink command syntax* on page 2-9 for more information.

For embedded applications with ROM at zero use `--entry 0x0` (or optionally `0xFFFF0000` for CPUs that have high vectors).

The initial entry point must meet the following conditions:

- the image entry point must always lie within an execution region

- the execution region must be non-overlay, and must be a root execution region (the load address is the same as the execution address).

If you do not use the `--entry` option to specify the initial entry point then:

- if the input objects contain only one entry point set by the `ENTRY` directive, the linker uses that entry point as the initial entry point for the image
- the linker generates an image that does not contain an initial entry point when either:
  - more than one entry point has been specified by using the `ENTRY` directive
  - no entry point has been specified by using the `ENTRY` directive.

In both these situations, the linker issues the following warning:

L6305W: Image does not have an entry point. (Not specified or not set due to multiple choices)

Specify a unique entry point with `--entry` to fix this warning.

## 3.2 Section placement

The linker sorts all the input sections within a region according to their attributes. Input sections with identical attributes form a contiguous block within the region.

The base address of each input section is determined by the sorting order defined by the linker, and is correctly aligned within the output section that contains it.

In general, the linker sorts the input sections in the following order when generating an image:

1. By attribute.
2. By input section name.
3. By their positions in the input list, except where overridden by *FIRST* or *LAST* (see *Using FIRST and LAST to place sections* on page 3-9).

If an execution region contains more than 4MB of Thumb® code or more than 32MB of ARM code, the linker might change the sort order to reduce the number of long branch veneers to a minimum.

By default, the linker creates an image consisting of an RO, an RW, and optionally a ZI output section. The RO output section can be protected at runtime on systems that have memory management hardware. RO sections can also be placed into ROM in the target.

This section describes:

- *Ordering input sections by attribute* on page 3-9
- *Using FIRST and LAST to place sections* on page 3-9
- *Aligning sections* on page 3-10
- *Ordering execution regions containing Thumb code* on page 3-10.

### 3.2.1 Ordering input sections by attribute

Portions of the image are collected together into a minimum number of contiguous regions. `armlink` orders input sections by attribute as follows:

1. read-only code
2. read-only data
3. read-write code
4. other initialized data
5. zero-initialized (uninitialized) data.

Input sections that have the same attributes are ordered by name. Names are considered to be case-sensitive and are compared in alphabetical order using the ASCII collation sequence for characters.

Identically attributed and named input sections are ordered according to their relative positions in the input list.

These rules mean that the positions of identically attributed and named input sections included from libraries are not predictable. If more precise positioning is required, you can extract modules manually and include them in the input list.

### 3.2.2 Using FIRST and LAST to place sections

Within a region, all RO code input sections are contiguous and form an RO output section that must precede the output section containing all the RW input sections.

If you are not using scatter-loading, use the `--first` and `--last` linker options to place input sections.

If you are using scatter-loading, use the attributes `FIRST` and `LAST` in the scatter-loading description file to mark the first and last input sections in an execution region if the placement order is important.

However, `FIRST` and `LAST` must not violate the basic attribute sorting order. This means that an input section can be first (or last) in the execution region if the output section it is in is the first (or last) output section in the region. For example, in an execution region containing RO input sections, the `FIRST` input section must be an RO input section. Similarly, if the region contains any ZI input sections, the `LAST` input section must be a ZI input section.

### 3.2.3 Aligning sections

When input sections have been ordered and before the base address is fixed, `armlink` inserts padding, if required, to force each input section to start at an address that is a multiple of the input section alignment.

It is possible to use `ALIGN` to expand the alignment (changing something that is normally four-byte aligned to be eight-byte aligned), but you cannot reduce the natural alignment (forcing two-byte alignment on something that is normally four-byte aligned).

Shown in Example 3-1, the input section alignment can be specified in assembly code using:

`ALIGN`            Use this attribute to the `AREA` directive from assembly language. The input section address will be a multiple of  $2^{(\text{value in align attribute})}$ .

#### Example 3-1 Using the `ALIGN` attribute in assembly code

---

```
AREA LDR_LABEL, CODE, READONLY, ALIGN=3 ; align on eight-byte boundary
```

---

See the description of `ALIGN` in the directives reference in *RealView Developer Kit v2.2 Assembler Guide*.

### 3.2.4 Ordering execution regions containing Thumb code

The Thumb branch range is 4MB. When an execution region contains Thumb code that exceeds 4MB, `armlink` attempts to order sections that are at a similar average call depth and to place the most commonly called sections centrally. This helps to minimize the number of veneers generated (see *Veneer generation* on page 3-19 for more details).



### 3.3 Optimizations and modifications

This section describes:

- *Common debug section elimination*
- *Common group or section elimination*
- *Unused section elimination* on page 3-12
- *Unused virtual function elimination* on page 3-12
- *Linker feedback* on page 3-14
- *RW data compression* on page 3-16
- *Veneer generation* on page 3-19
- *Reuse of veneers with overlay execution regions* on page 3-21
- *Branch inlining* on page 3-22.

#### 3.3.1 Common debug section elimination

The compiler and assembler generate one set of debug sections for each source file that contributes to a compilation unit. `arm1ink` can detect multiple copies of a debug section for a particular source file and discard all but one copy in the final image. This can result in a considerable reduction in image debug size.

#### 3.3.2 Common group or section elimination

If there are inline functions or templates used in the C++ source, the ARM compiler generates complete objects for linking such that each object contains the out-of-line copies of inline functions and template functions that the object requires. When these functions are declared in a common header file, the functions might be defined many times in separate objects that are subsequently linked together. In order to eliminate duplicates, the compiler compiles these functions into separate instances of common code sections or groups.

It is possible that the separate instances of a common code sections, or groups, are not identical. Some of the copies, for example, might be found in a library that has been built with different (but compatible) build options, different optimization, or different debug options.

If the copies are not identical, `arm1ink` retains the best available variant of each common code section, or group, based on the attributes of the input objects. `arm1ink` discards the rest.

This optimization is controlled by linker options:

- Use the `--bestdebug` option to use the largest Comdat group (likely to give the best debug view).

- Use the `--no_bestdebug` option to use the smallest Comdat group (likely to give the smallest code size). This is the default. In most cases, it is unlikely that you will need to use `--bestdebug`.

Because `--no_bestdebug` is the default, the final image is the same regardless of whether you generate debug tables during compilation with `-g`.

For details, see generating debug information in the chapter describing how to use the ARM compiler in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

### 3.3.3 Unused section elimination

Unused section elimination removes code that is never executed, or data that is not referred to by the code, from the final image. This optimization can be controlled by the `--(no_)remove` and `--first`, `--last`, and `--keep` linker options. Use the `--info unused` linker option to instruct the linker to generate a list of the unused sections that have been eliminated.

Unused section elimination is suppressed in those cases that might result in the removal of all sections.

An input section is retained in the final image under the following conditions:

- if it contains an entry point
- if it is referred to, directly or indirectly, by a non-weak reference from an input section containing an entry point
- if it was specified as the first or last input section by the `--first` or `--last` option (or a scatter-loading equivalent)
- if it has been marked as unremovable by the `--keep` option.

#### ———— Note —————

Unused section elimination is a property of all groups, not just common groups.

### 3.3.4 Unused virtual function elimination

Unused virtual function elimination (VFE) is a refinement of unused section elimination to reduce ROM size in images generated from C++ code. This optimization can also be used to eliminate both unused virtual functions and RTTI objects from your code.

As described in *Controlling image contents* on page 2-18, VFE awareness is controlled by linker option `--vfemode=mode`.

---

**Note**

---

If you do not specify a VFE option on the command line, default mode is assumed, that is, `--vfemode=on`.

---

Unused section elimination efficiently removes unused functions from C code. However, in C++ applications, unused virtual functions and RTTI objects are referenced by pointer tables. This means that the unused section elimination algorithm used by the linker cannot guarantee to remove unused virtual functions and RTTI objects reliably.

VFE is a collaboration between the ARM compiler and the linker whereby the compiler supplies extra information about unused virtual functions that is then used by the linker. Based on this analysis, the linker is able to remove unused virtual functions reliably. This collaboration also enables the linker to remove RTTI objects.

---

**Note**

---

Assembler source files do not require VFE annotations, provided that they do not reference the C++ libraries. This is because the linker assumes that no virtual function calls are made by object files that do not reference the C++ libraries. Similarly, C source files that were compiled with an old version of `armcc` can participate in VFE provided that they do not reference the C++ libraries.

---

VFE operates in four modes:

- |            |   |
|------------|---|
| <b>On</b>  | <p>Use the command-line option <code>--vfemode=on</code> to make the linker VFE aware. This is the default mode if you do not specify a VFE option on the command line.</p> <p>In this mode the linker chooses force or off mode based on the content of object files:</p> <ul style="list-style-type: none"> <li>• Where every object file contains VFE information or does not refer to C++ libraries, the linker assumes force mode and continues with the elimination.</li> <li>• If any object file is missing VFE information and refers to a C++ library, for example, where code has been compiled with a previous release of the ARM tools, the linker assumes off mode, and VFE is disabled silently. Choosing off mode to disable VFE in this situation ensures that the linker does not remove a virtual function that is used by an object with no VFE information.</li> </ul> |
| <b>Off</b> | <p>Use the command-line option <code>--vfemode=off</code> to make <code>armlink</code> ignore any extra information supplied by the compiler. In this mode, the final image is the same as that produced by compiling and linking without VFE awareness.</p>  |

**Force** Use the command-line option `--vfemode=force` to make the linker VFE aware and force the VFE algorithm to be applied. If some of the object files do not contain VFE information, for example, where they have been compiled with a previous release of the ARM tools, the linker continues with the elimination but displays a warning to alert you to possible errors.

### **Force no RTTI**

Use the command-line option `--vfemode=force_no_rtti` to make the linker VFE aware and force the removal of all RTTI objects. In this mode all virtual functions are retained.

The compiler places the extra information in sections with names beginning `.arm_vfe`. These sections are not referenced by the rest of the code and so are ignored by the linker when it is not VFE aware. Such sections do not, therefore, increase the size of the final image but they do increase the size of object files produced by the compiler.

To stop the compiler producing VFE information, use the compiler option `--no_vfe`. However, the recommended way to switch off VFE is with the linker option `--vfemode=off`.

## **3.3.5 Linker feedback**

`armlink` provides feedback for the next time a file is compiled, to inform the compiler about unused functions. These are placed in their own sections for future elimination by the linker.

When the `--inline` optimization is turned on (see *Branch inlining* on page 3-22), functions inlined by the linker are also emitted in the feedback file. These functions are also placed in their own sections.

This option generates a feedback file containing each output filename, as a comment, and the unused symbols found in the file, for example:

```
;<FEEDBACK># ARM Linker, RVCT2.2 [Build 304]: Last Updated: Wed Dec 08 14:45:35
2004
;VERSION 0.2
;FILE dhry_1.o
unused_func1 <= USED 0
inlined_func <= LINKER_INLINED
;FILE dhry_2.o
unused_func2 <= USED 0
```

When you next compile the source, use the compiler option `--feedback file` to specify the linker-generated feedback file to use. If no feedback file exists, the compiler issues a warning message.

## Linker feedback example

To see how linker feedback works:

1. Create a file `fb.c` containing the code shown in Example 3-2.

### Example 3-2 Feedback example

---

```
#include <stdio.h>

void legacy() {
    printf("This is a legacy function, that is no longer used.\n");
}

int cubed(int i) {
    return i*i*i;
}

void main() {
    int n = 3;
    printf("%d cubed = %d\n",n,cubed(n));
}
```

---

2. Compile the program (ignore the warning that the feedback file does not exist):  
`armcc --asm -c --feedback fb.txt fb.c`

This inlines the `cubed()` function by default, and creates an assembler file `fb.s` and an object file `fb.o`. In the assembler file, the code for `legacy()` and `cubed()` is still present. Because of the inlining, there is no call to `cubed()` from `main`.

An out-of-line copy of `cubed()` is kept because it was not declared as **static**.

3. Link the object file to create the linker feedback file with the command line:  
`armlink --info sizes --list fbout1.txt --feedback fb.txt fb.o -o fb.axf`  
 Linker diagnostics are output to the file `fbout1.txt`.

---

#### Note

---

For full details on these options see *Generating image-related information* on page 2-26 and *Controlling linker diagnostics* on page 2-30.

---

The linker feedback file identifies the source file that contains the unused functions in a comment (not used by the compiler) and includes entries for the `legacy()` and `cubed()` functions:

```
;#<FEEDBACK># ARM Linker, RVCT2.2 [Build 304]: Last Updated: Wed Dec 08
15:00:35 2004
;VERSION 0.2
;FILE fb.o
cubed <= USED 0
legacy <= USED 0
```

This shows that the functions are not used.

4. Repeat the compile and link stages with a different diagnostics file:  

```
armcc --asm -c --feedback fb.txt fb.c
armmlink --info sizes --list fbout2.txt fb.o -o fb.axf
```
5. Compare the two diagnostics files, fbout1.txt and fbout2.txt, to see the sizes of the image components (Code, RO Data, RW Data, ZI Data, and Debug). The Code and Debug components are smaller.

In the assembler file, fb.s, the legacy() and cubed() functions are no longer in the main .text area. They are compiled into their own ELF sections. Therefore, armmlink can remove the legacy() and cubed() functions from the final image.

---

**Note**

---

To get the maximum benefit from linker feedback you have to do a full compile and link at least twice. However, a single compile and link using feedback from a previous build is usually sufficient.

---

### 3.3.6 RW data compression

Most RW data areas contain over 75% zeros making them highly compressible. RW data compression is enabled by default to minimize ROM size.

The ARM libraries contain some decompression algorithms and the linker chooses the optimal one to add to your image to decompress the data areas when the image is executed. However, you can override the algorithm chosen by the linker.

This section describes data compression in more detail:

- *Choosing a compressor* on page 3-17
- *How is compression applied?* on page 3-17
- *Working with RW data compression* on page 3-18.

## Choosing a compressor

armlink gathers information about the content of data sections before choosing the most appropriate compression algorithm to generate the smallest code size. If compression is appropriate, the linker can only use one data compressor for all the compressible data sections in the image and different compressions might be tried on these sections to produce the best overall size. Compression is applied automatically if:

compressed data size + size of decompressor < uncompressed data size

Once a compressor has been chosen, armlink adds the decompressor to the code area of your image. If the final image does not contain any compressed data, no decompressor is added.

You can override the compression used by the linker by:

- using the `--datacompressor off` option to specify that no compression is applied
- specifying your own compressor.

Use the command-line option `--datacompressor list` to get a list of compressors available in the linker, for example:

Num	Compression algorithm
0	Run-length encoding
1	Run-length encoding, with LZ77 on small-repeats
2	Complex LZ77 compression

## How is compression applied?

Run-length compression encodes data as non-repeated bytes and repeated zero-bytes. Non-repeated bytes are output unchanged, followed by a count of zero-bytes. Lempel-Ziv 1977 (LZ77) compression keeps track of the last *n* bytes of data seen and, when a phrase is encountered that has already been seen, it outputs a pair of values corresponding to the position of the phrase in the previously-seen buffer of data, and the length of the phrase.

To specify a compressor, use the required ID on the linker command line, for example:

```
armlink --datacompressor 2 ...
```

When choosing a compressor be aware that:

- Compressor 0 performs well on data with large areas of zero-bytes but few non-zero bytes.
- Compressor 1 performs well on data where the non-zero bytes are repeating.

- Compressor 2 performs well on data that contains repeated values.

The linker prefers compressor 0 or 1 where the data contains mostly zero-bytes (>75%). Compressor 2 is chosen where the data contains few zero-bytes (<10%). If the image is made up only of ARM code, then ARM decompressors are used automatically. If the image contains any Thumb code, Thumb decompressors are used. If there is no clear preference, all compressors are tested to produce the best overall size (see *Choosing a compressor* on page 3-17).

---

#### Note

---

It is not possible to add your own compressors into the linker. The algorithms that are available, and how the linker chooses to use them, might change in the future.

---

## Working with RW data compression

When working with RW data compression:

- Use the linker option `--map` to see where compression has been applied to regions in your code.
- The linker does not apply compression if a `Load$$region_name$$Base` symbol is used, where *region\_name* follows any execution region containing compressed data in the same load region.
- If you are using an ARM processor with on-chip cache, enable the cache after decompression to avoid code coherency problems

In RVCT v2.0 and earlier, only the `__main` section and the region tables had to be placed in a root region. RW data compression (RVCT v2.1 and later) requires that additional sections (such as `__dc*.o` sections) be placed in a root region.

If you are using a scatter-loading description file:

- Where coded, decompressor objects named `__dc*.o`, must be in a root region, for example:

```
ER_ROOT
{
    __main.o(*)
    * (Region$$Table)
    __scatter*.o(*)
    __dc*.o(*)
}
```

Or, alternatively, use `InRoot$$Sections` to place all library sections that must be in a root region, for example:



```

ER_ROOT
{
    * (InRoot$$Sections)
}

```

See *Assigning sections to a root region* on page 5-30 for more details.

- Specify that a load or execution region should not be compressed by adding the NOCOMPRESS attribute (see *Formal syntax of the scatter-loading description file* on page 5-8 for details).

### 3.3.7 Veneer generation

Veneers are small sections of code generated by the linker and inserted into your program. `armlink` must generate veneers when a branch involves a destination beyond the branching range of the current state.

The range of a BL instruction is 32MB for ARM and 4MB for Thumb. A veneer can, therefore, extend the range of the branch by becoming the intermediate target of the instruction and then setting the PC to the destination address. If ARM and Thumb are mixed, the veneer also changes processor state.

`armlink` supports the following veneer types:

- ARM to ARM
- ARM to Thumb (interworking veneers)
- Thumb to ARM (interworking veneers)
- Thumb to Thumb.

`armlink` creates one input section called `Veneer$$Code` for each veneer. A veneer is generated only if no other existing veneer can satisfy the requirements. If two input sections contain a long branch to the same destination, only one veneer is generated. A veneer is only shared in this way if it can be reached by both sections.

If you are using ARMv4T, `armlink` generates veneers when a branch involves change of state between ARM and Thumb state. In ARMv5 and above, the BLX instruction is used.

### Veneer sharing

You can use the command-line option `--no_veneershare` to specify that veneers are not shared across execution regions. This assigns ownership of the created veneer section to the object that created the veneer and so enables you to select veneers from a particular object in a scatter-loading description file, for example:

```
ER1 +0
{
    object1.o(Veneer$$Code)
}
```

Veneer sharing makes it impossible to assign an owning object. Using `--no_veneershare`, therefore, provides a more consistent image layout. This comes at a cost with a significant increase in code size.

## Veneer variants

Veneers have different variants depending on the branching range you require:

- *Inline veneers*
- *Short branch veneers*
- *Long branch veneers.*

### ***Inline veneers***

In this variant:

- the veneer must be inserted just before the target section to be in range
- an ARM-Thumb interworking veneer has a range of 256 bytes and so the function code must appear in the first 256 bytes immediately after the veneer code
- a Thumb-ARM interworking veneer has a range of zero bytes and so the function code must appear immediately after the veneer code
- an inline veneer is always position-independent.

These limitations mean that you cannot move inline veneers out of an execution region using `Veneer$$Code`. Use the command-line option `--no_inlineveneer` to prevent the generation of inline veneers.

### ***Short branch veneers***

In this variant:

- an ARM-Thumb short branch veneer has a range of 4MB
- a Thumb-ARM short branch veneer has a range of 32MB
- a short branch veneer is always position-independent.

### ***Long branch veneers***

In this variant:

- both ARM-Thumb and Thumb-ARM interworking veneers have a range of 232 bytes

- `armlink` combines long branch capability into the state change capability, therefore, all interworking veneers are also long branch veneers
- a long branch veneer is either absolute or position-independent.

When you are using veneers be aware of the following:

- All veneers cannot be collected into one input section because the resulting veneer input section might not be within range of other input sections. If the sections are not within addressing range, long branching is not possible.
- `armlink` generates position-independent variants of the veneers automatically. However, because such veneers are larger than non position-independent variants, `armlink` only does this where necessary, that is, where the source and destination execution regions are both position-independent and are rigidly related.

Veneers are generated to optimize code size. `armlink`, therefore, chooses the variant in order of preference:

1. Inline veneer.
2. Short branch veneer.
3. Long veneer.

### 3.3.8 Reuse of veneers with overlay execution regions

`armlink` reuses veneers whenever possible. However, both the following conditions are enforced on reuse:

- an overlay execution region cannot reuse a veneer placed in any other execution region
- no other execution region can reuse a veneer placed in an overlay execution region.

If these conditions are not met, new veneers are created instead of reusing existing ones. Unless you have instructed the linker to place veneers somewhere specific using scatter-loading, a veneer is always placed in the execution region that contains the call requiring the veneer. This implies that:

- for an overlay execution region, all its veneers are included within the execution region
- an overlay execution region never requires a veneer from another execution region.

### 3.3.9 Branch inlining

armlink has global visibility of all your program code and so can perform some additional branch optimizations.

armlink uses branch inlining to optimize small function calls in your image. A small function is defined as any one-instruction function that can be inlined into the 4 bytes of a BL or BLX instruction. In this case, there is no branch and, therefore, the return address is redundant.

---

#### Note

---

This branch optimization is off by default because enabling it changes the image such that debug information might be incorrect. If enabled, the linker makes no attempt to correct the debug information.

---

Use the command-line options to control branch inlining:

`--inline` This option enables branch inlining (see *Controlling inlining* on page 3-23).

`--tailreorder`

This option moves tail calling sections immediately before their target, if possible, to optimize function calls (see *Handling tail calling sections* on page 3-23).

If you enable branch inlining, armlink scans each function call in the image and then inlines where applicable. When armlink inlines a function, it removes the reference to the called function from the caller. armlink applies this optimization before any unused sections are eliminated so that any section that is always inlined can then be removed.

Use the `--info` command-line option to display information about branch inlining:

`--info inline`

This option displays a message each time a function is inlined and gives the total number of inlines, for example:

Small function inlining results

Inlined function `__Heap_DescSize` from object `h1_alloc.o` at offset `0x5c` in section `.text` from object `malloc.o`.

Inlined function `__ieee_status` from object `istatus.o` at offset `0x40` in section `.text` from object `_printf_fp_dec.o`.

.

.

Inlined total of 6 calls.

## Controlling inlining

If you have enabled branch inlining, there are certain conditions that a function must meet in order to be inlined:

- `armlink` handles only the simplest cases and does not inline any instruction that reads or writes to the PC because this depends on the location of the function.
- The action of the linker also depends on the size of the symbol representing a function and on the caller (ARM or Thumb) and the callee (ARM or Thumb) as shown in Table 3-2.

**Table 3-2 Inlining small functions**

Caller	Callee	Symbol size that can be inlined
ARM	ARM	4 to 8 bytes
ARM	Thumb	2 to 6 bytes
Thumb	Thumb	2 to 6 bytes
Thumb	ARM	4 to 8 bytes

- In order to be inlined, the last instruction of a function must be either:  
`MOV pc, rn`  
or  
`BX rn`  
Any function that consists of just a return sequence is inlined as a NOP.
- If your image contains both ARM and Thumb code, functions that are called from the other state must be built for interworking. An ARM caller might inline a Thumb callee if an equivalent ARM instruction is available. However, a Thumb caller cannot inline an ARM callee. Also, `armlink` can inline up to two 16-byte Thumb instructions. However, an ARM caller can only inline a single 16-bit Thumb instruction.

## Handling tail calling sections

The linker replaces any branch with a relocation that resolves to the next instruction with a NOP. This means that tail calling sections, that is, sections that finish with a branch instruction, might be optimized so that their target appears immediately after them in the execution region.

You can take advantage of this behavior by using the command-line option `--tailreorder` to move tail calling sections above their target. If this is possible, be aware that `armlink`:

- Can only move one tail calling section for each tail call target.  
If there are multiple tail calls to a single section, the tail calling section with an identical section name is moved before the target. If no section name is found in the tail calling section that has a matching name, then the linker moves the *first* section it encounters.
- Cannot move a tail calling section out of its execution region.
- Does not move tail calling sections before inline veneers.

Use the `--info` command-line option to display information about tail call optimization:

`--info tailreorder`

Gives details of any tail calling sections that have been moved, for example:

```
Tailcall reorder results
Tail calling Section !!!main from object __main.o placed before
.text from kernel.o
Tail calling Section .text from object rt_raise.o placed before
.text from sys_exit.o
Tail calling Section .text from object plibspace.o placed before
.text from libspace.o
Tail calling Section .text from object aeabi_idiv0.o placed before
.text from rt_div0.o
Tail calling Section .text from object h1_extend.o placed before
.text from h1_free.o
Tail calling Section .text from object lludiv.o placed before
.text from aeabi_ldiv0.o
```

### 3.4 Using command-line options to create simple images

A simple image consists of a number of input sections of type RO, RW, and ZI. These input sections are collated to form the RO, the RW, and the ZI output sections. Depending on how the output sections are arranged within load and execution regions, there are three basic types of simple image:

- Type 1**      One region in load view, three contiguous regions in execution view. Use the `--ro-base` option to create this type of image.
- See *Type 1: one load region and contiguous output regions* on page 3-26 for more details.
- Type 2**      One region in load view, three non-contiguous regions in execution view. Use the `--ro-base` and `--rw-base` options to create this type of image.
- See *Type 2: one load region and non-contiguous output regions* on page 3-27 for more details.
- Type 3**      Two regions in load view, three non-contiguous regions in execution view. Use the `--ro-base`, `--rw-base`, and `--split` options to create this type of image. You can also use the `--rosplit` option to split the default load region into two RO output sections, one for code and one for data.
- See *Type 3: two load regions and non-contiguous output regions* on page 3-29 for more details.

In all three simple image types, there are up to three execution regions where:

- the first execution region contains the RO output section
- the second execution region contains the RW output section (if present)
- the third execution region contains the ZI output section (if present).

These execution regions are referred to as the RO, the RW, and the ZI execution region.

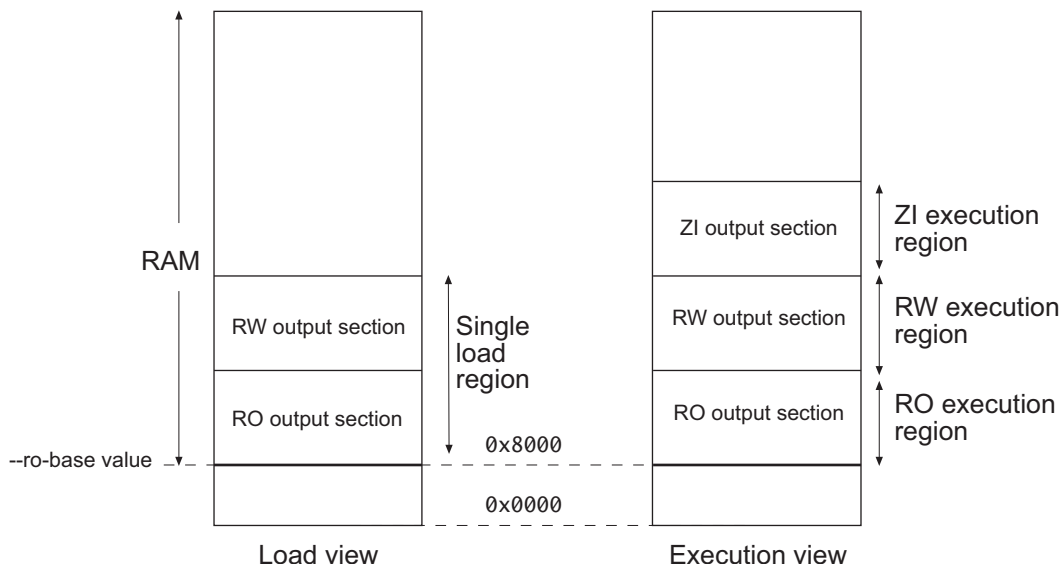
Simple images can also be created with scatter-loading description files. See *Equivalent scatter-loading descriptions for simple images* on page 5-35 for more information on how to do this.

This section describes:

- *Type 1: one load region and contiguous output regions* on page 3-26
- *Type 2: one load region and non-contiguous output regions* on page 3-27
- *Type 3: two load regions and non-contiguous output regions* on page 3-29.

### 3.4.1 Type 1: one load region and contiguous output regions

An image of this type consists of a single load region in the load view and the three execution regions are placed contiguously in the memory map. This approach is suitable for systems that load programs into RAM, for example, an OS bootloader or a desktop system (see Figure 3-3).



**Figure 3-3 Simple type 1 image**

Use the following command for images of this type:

```
armlink --ro-base 0x8000
```

#### Load view

The single load region consists of the RO and RW output sections placed consecutively. The RO and RW execution regions are both root regions. The ZI output section does not exist at load time. It is created before execution using the output section description in the image file.



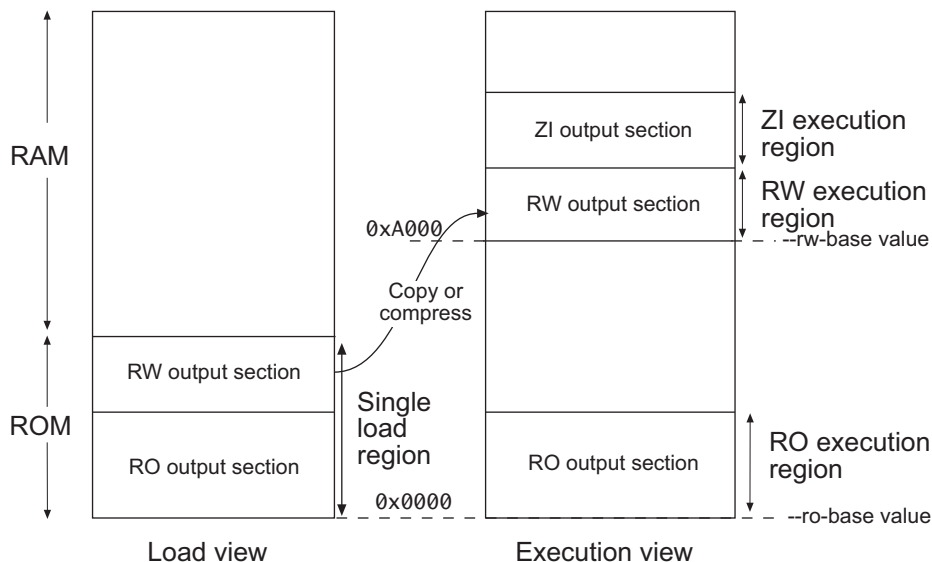
## Execution view

The three execution regions containing the RO, RW, and ZI output sections are arranged contiguously. The execution addresses of the RO and RW execution regions are the same as their load addresses, so nothing has to be moved from its load address to its execution address. However, the ZI execution region that contains the ZI output section is created before execution begins.

Use `armlink --ro-base address` to specify the load and execution address of the region containing the RO output. The default address is `0x8000` as shown in Figure 3-3 on page 3-26.

### 3.4.2 Type 2: one load region and non-contiguous output regions

An image of this type consists of a single load region, and three execution regions in execution view. The RW execution region is not contiguous with the RO execution region. This approach is used, for example, for ROM-based embedded systems (see Figure 3-4), where RW data is copied from ROM to RAM at startup.



**Figure 3-4 Simple type 2 image**

Use the following command for images of this type:

```
armlink --ro-base 0x0 --rw-base 0xA000
```

## Load view

In the load view, the single load region consists of the RO and RW output sections placed consecutively, in ROM, for example. Here, the RO region is a root region, and the RW region is non-root. The ZI output section does not exist at load time. It is created at run-time.

## Execution view

In the execution view, the first execution region contains the RO output section and the second execution region contains the RW and ZI output sections.

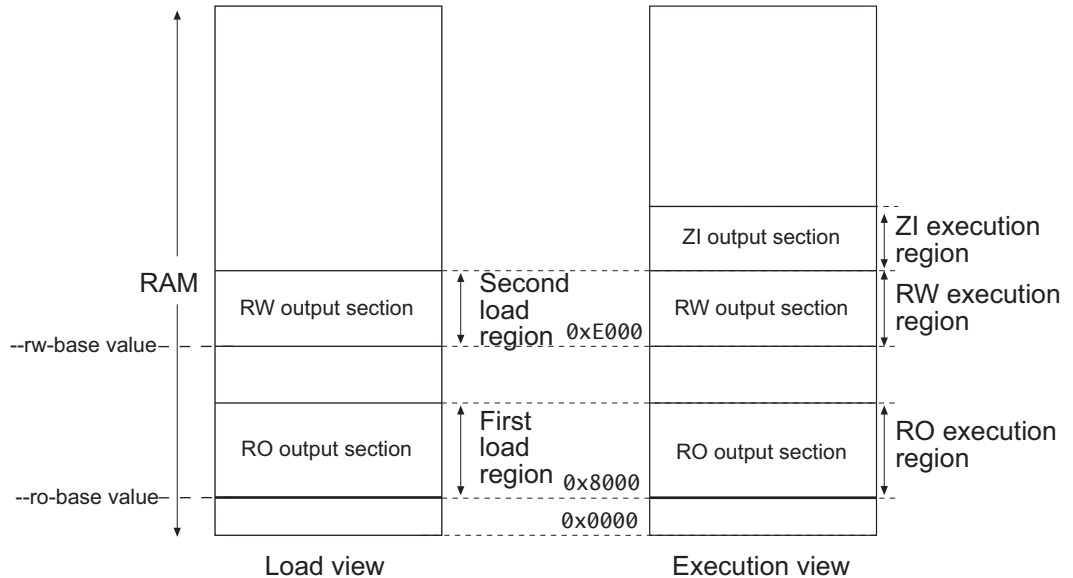
The execution address of the region containing the RO output section is the same as its load address, so the RO output section does not have to be moved. That is, it is a root region.

The execution address of the region containing the RW output section is different from its load address, so the RW output section is moved from its load address (from the single load region) to its execution address (into the second execution region). The ZI execution region, and its output section, is placed contiguously with the RW execution region.

Use `armlink` options `--ro-base address` to specify the load and execution address for the RO output section, and `--rw-base exec_address` to specify the execution address of the RW output section. If you do not use the `--ro-base` option to specify the address, the default value of `0x8000` is used by `armlink`. For an embedded system, `0x0` is typical for the `--ro-base` value. If you do not use the `--rw-base` option to specify the address, the default is to place RW directly above RO (that is, Type 1).

### 3.4.3 Type 3: two load regions and non-contiguous output regions

This type of image is similar to images of type 2 except that the single load region in type 2 is now split into two load regions (see Figure 3-5).



**Figure 3-5 Simple type 3 image**

Use the following command for images of this type:

```
armlink --split --ro-base 0x8000 --rw-base 0xE000
```

#### Load view

In the load view, the first load region consists of the RO output section, and the second load region consists of the RW output section. The ZI output section does not exist at load time. It is created before execution using the description of the output section contained in the image file.

#### Execution view

In the execution view, the first execution region contains the RO output section and the second execution region contains the RW and ZI output sections.

The execution address of the RO region is the same as its load address, so the contents of the RO output section do not have to be moved or copied from their load address to their execution address. Both RO and RW are root regions.

The execution address of the RW region is also the same as its load address, so the contents of the RW output section are not moved from their load address to their execution address. However, the ZI output section is created before execution begins and placed after the RW region.

Specify the load and execution address using the following linker options:

`--split` This option splits the default single load region (that contains both the RO and RW output sections) into two load regions (one containing the RO output section and one containing the RW output section) so that they can be separately placed using `--ro-base` and `--rw-base`.

`--ro-base address`

This option instructs `armlink` to set the load and execution address of the region containing the RO section at a four-byte aligned *address* (for example, the address of the first location in ROM). If `--ro-base` option is not used to specify the address, the default value of `0x8000` is used by `armlink`.

`--rw-base address`

This option instructs `armlink` to set the execution address of the region containing the RW output section at a four-byte aligned *address*. If this option is used with `--split`, this specifies both the load and execution addresses of the RW region (that is, it is a root region).

### 3.5 Using command-line options to handle C++ exceptions

By default, or if the option `--exceptions` is specified, the image can contain exception tables. Exception tables are discarded silently if no code throws an exception. However, if the option `--no_exceptions` is specified, the linker generates an error if any exceptions sections are present after unused sections have been eliminated.

You can use the `--no_exceptions` option if you want to ensure that your code is exceptions free. The linker generates an error message to highlight that exceptions have been found and does not produce a final image.

However, you can use the `--no_exceptions` option with the `--diag_warning` option to downgrade the error message to a warning. The linker produces a final image but also generates a message to warn you that exceptions have been found.

The linker can create exception tables for legacy objects that contain debug frame information. The linker can do this safely for C and assembly language objects. By default, the linker does not create exception tables. This is the same as using the linker option `--exceptions_tables=nocreate`.

The linker option `--exceptions_tables=unwind` enables the linker to use the `.debug_frame` information to create a register-restoring unwinding table for each section in your image that does not already have an exception table. If this is not possible, the linker creates a `nounwind` table instead.

Use the linker option `--exceptions_tables=cantunwind` to create a `nounwind` table for each section in your image that does not already have an exception table.

---

#### Note

---

Be aware of the following:

- With the default settings, that is, `--exceptions --exception_tables=nocreate`, it is not safe to throw an exception through C or assembly code (unless the C code is compiled with the option `--exceptions`).
  - The linker is unable to generate the cleanup code necessary for automatic variables in legacy C++ code, for example, RVCT v1.2 or ADS v1.2.
-

### 3.6 Getting information about images

You can use the `--info` option to get information about how your image is generated by the linker, for example:

```
armlink --info sizes ...
```

Here, `sizes` gives a list of the Code and Data sizes for each input object and library member in the image. Using this option implies `--info sizes,totals`.

See *Generating image-related information* on page 2-26 for more details on the topic keywords you can specify for the `--info` command-line option.

Example 3-3 shows an example of the output from this option.

Example 3-3 Image details example

<hr/>						
Code (inc. data)	R0 Data	RW Data	ZI Data	Debug		
3712	1580	19	44	10200	7436	Object Totals
0	0	16	0	0	0	(incl. Generated)
0	0	3	0	0	0	(incl. Padding)
21376	648	805	4	300	10216	Library Totals
0	0	6	0	0	0	(incl. Padding)
<hr/>						
=====						
Code (inc. data)	R0 Data	RW Data	ZI Data	Debug		
25088	2228	824	48	10500	17652	Grand Totals
25088	2228	824	48	10500	17652	Image Totals
<hr/>						
=====						
Total R0	Size (Code + R0 Data)		25912 ( 25.30kB)			
Total RW	Size (RW Data + ZI Data)		10548 ( 10.30kB)			
Total ROM	Size (Code + R0 Data + RW Data)		25960 ( 25.35kB)			
<hr/>						

In Example 3-3, the output is shown in rows and columns, and totals are separated out for easy reading. For details see:

- *Column details* on page 3-33
- *Row details* on page 3-33.

### 3.6.1 Column details

Example 3-3 on page 3-32 shows:

#### **Code (inc. Data)**

Shows how many bytes are occupied by code. In this image, there are 3712 bytes of code. This includes 1580 bytes of inline data (inc. Data), for example, literal pools and short strings.

#### **RO Data**

Shows how many bytes are occupied by read-only data, such as initial values of RW data variables. This is in addition to the inline data included in the Code (inc. Data) column.

#### **RW Data**

Shows how many bytes are occupied by read-write data.

#### **ZI Data**

Shows how many bytes are occupied by zero-initialized data.

#### **Debug**

Shows how many bytes are occupied by debug data, for example, debug input sections and the symbol and string table.

### 3.6.2 Row details

Example 3-3 on page 3-32 shows:

#### **Object Totals**

Shows how many bytes are occupied by objects linked together to generate the image.

#### **(incl. Generated)**

armlink might generate image contents, for example, interworking veneers, and input sections such as region tables. If the Object Totals row includes this type of data, it is shown in this row. In Example 3-3 on page 3-32 there are 19 bytes of RO data in total, of which 16 bytes is linker-generated RO data.

#### **Library Totals**

Shows how many bytes are occupied by library members that have been extracted and added to the image as individual objects.

#### **(incl. Padding)**

armlink inserts padding, if required, to force section alignment (see *Aligning sections* on page 3-10). If the Object Totals row includes this type of data, it is shown in the associated (incl. Padding) row. Similarly, if the Library Totals row includes this type of data, it is shown in its

associated row. In Example 3-3 on page 3-32, there are 19 bytes of RO data in the object total, of which 3 bytes is linker-generated padding, and 805 bytes of RO data in the library total, with 6 bytes of padding.

Example 3-3 on page 3-32 shows:

### **Grand Totals**

Shows the true size of the image. In Example 3-3 on page 3-32 there are 10200 bytes of ZI data (in Object Totals) and 300 of ZI data (in Library Totals) giving a total of 10500 bytes.

### **Image Totals**

If you are using Read/Write data compression (the default) to optimize ROM size, the size of the final image changes and this is reflected in the output from `--info`. Compare the number of bytes under Grand Totals and Image Totals to see the effect of compression.

In Example 3-3 on page 3-32, RW compression is not enabled. If data is compressed, the RW value changes (see *RW data compression* on page 3-16 for details).

Example 3-3 on page 3-32 also shows totals (in bytes and kilobytes) for the final image size.

You can use the `--map` option to create an image map. This includes the address and size of each load region, execution region, and input section in the image, and shows how RW compression is applied (see *Generating image-related information* on page 2-26 for details).



# Chapter 4

## Accessing Image Symbols

This chapter describes how to reference symbols with the ARM® linker, `armlink`. It contains the following sections:

- *ARM/Thumb synonyms* on page 4-2
- *Accessing linker-defined symbols* on page 4-3
- *Accessing symbols in another image* on page 4-7
- *Hiding and renaming global symbols* on page 4-10
- *Using `$Super$$` and `$Sub$$` to override symbol definitions* on page 4-20
- *Symbol versioning* on page 4-21.

## 4.1 ARM/Thumb synonyms

The ARM linker enables multiple definitions of a symbol to coexist in an image, only if each definition is associated with a different processor state. `armlink` applies the following rules when a reference is made to a symbol with ARM/Thumb<sup>®</sup> synonyms:

- B, BL, or BLX instructions to a symbol from ARM state resolve to the ARM definition.
- B, BL, or BLX instructions to a symbol from Thumb state resolve to the Thumb definition.

Any other reference to the symbol resolves to the first definition encountered by the linker. In this case, `armlink` displays a warning that specifies the chosen symbol.

## 4.2 Accessing linker-defined symbols

The linker defines some symbols that contain the character sequence `$$`. These symbols, and all other external names containing the sequence `$$`, are names reserved by ARM Limited. The symbols are used to specify region base addresses, output section base addresses, and input section base addresses and their limits.

These symbolic addresses can be imported and used as relocatable addresses by your assembly language programs, or referred to as **extern** symbols from your C or C++ source code. See *Importing linker-defined symbols* on page 4-6 for details.

### ————— **Note** —————

Linker-defined symbols are only generated when your code references them.

Support for C++ may be obtained as an upgrade on certain versions of RVDK, by means of an appropriate FLEXlm license.

This section describes:

- *Region-related symbols*
- *Section-related symbols* on page 4-5
- *Importing linker-defined symbols* on page 4-6.

### 4.2.1 Region-related symbols

Region-related symbols are generated when the linker creates an image. Table 4-1 shows the symbols that the linker generates for every execution region present in the image.

**Table 4-1 Region-related linker symbols**

Symbol	Description
Load\$\$region_name\$\$Base	Load address of the region
Image\$\$region_name\$\$Base	Execution address of the region
Image\$\$region_name\$\$Length	Execution region length in bytes (multiple of 4)
Image\$\$region_name\$\$Limit	Address of the byte beyond the end of the execution region
Image\$\$region_name\$\$ZI\$\$Base	Execution address of the ZI output section in this region
Image\$\$region_name\$\$ZI\$\$Length	Length of the ZI output section in bytes (multiple of 4)
Image\$\$region_name\$\$ZI\$\$Limit	Address of the byte beyond the end of the ZI output section in the execution region

If you are using scatter-loading, the description file names all the execution regions in the image, and provides their load and execution addresses (see Chapter 5 *Using Scatter-loading Description Files* for details).

If you are not using scatter-loading, the linker uses *region\_name* values of:

- ER\_R0, for read-only regions
- ER\_RW, for read-write regions
- ER\_ZI, for zero-initialized regions.

For every execution region containing a ZI output section, the linker generates additional symbols containing \$\$ZI\$\$.

---

#### Note

- The ZI output sections of an image are not created statically, but are automatically created dynamically at runtime. Therefore there is no load address symbol for ZI output sections.
  - It is recommended that you use region-related symbols in preference to section-related symbols.
- 

## Placing the stack and heap above the ZI region

One common use of region-related symbols is to place a heap directly above the ZI region. Example 4-1 shows how to create a retargeted version of `__user_initial_stackheap()` in assembly language. The example assumes that you are using the default one region memory model from the ARM C libraries. See the description of `__user_initial_stackheap()` in the chapter describing the C and C++ libraries in *RealView Developer Kit v2.2 Compiler and Libraries Guide* for more information.

### Example 4-1 Placing the stack and heap above the ZI region

---

```
EXPORT __user_initial_stackheap
IMPORT ||Image$$region_name$$ZI$$Limit||
__user_initial_stackheap
    LDR r0, =||Image$$region_name$$ZI$$Limit||
    MOV pc, lr
```

---

## 4.2.2 Section-related symbols

The output section symbols shown in Table 4-2 are generated if you use command-line options to create a simple image. A simple image has three output sections (RO, RW, and ZI) that produce the three execution regions. For every input section present in the image, the linker generates the input symbols shown in Table 4-2.

The linker sorts sections within an execution region first by attribute RO, RW, or ZI, then by name. So, for example, all .text sections are placed in one contiguous block. A contiguous block of sections with the same attribute and name is known as a *consolidated section*.

**Table 4-2 Section-related linker symbols**

Symbol	Section type	Description
Image\$\$RO\$\$Base	Output	Address of the start of the RO output section.
Image\$\$RO\$\$Limit	Output	Address of the first byte beyond the end of the RO output section.
Image\$\$RW\$\$Base	Output	Address of the start of the RW output section.
Image\$\$RW\$\$Limit	Output	Address of the byte beyond the end of the ZI output section. (The choice of the end of the ZI region rather than the end of the RW region is to maintain compatibility with legacy code.)
Image\$\$ZI\$\$Base	Output	Address of the start of the ZI output section.
Image\$\$ZI\$\$Limit	Output	Address of the byte beyond the end of the ZI output section.
SectionName\$\$Base	Input	Address of the start of the consolidated section called SectionName.
SectionName\$\$Limit	Input	Address of the byte beyond the end of the consolidated section called SectionName.

### Note

If your code refers to the input-section symbols, it is assumed that you expect all the input sections in the image with the same name to be placed contiguously in the image memory map. If your scatter-loading description places these input sections non-contiguously, the linker diagnoses an error because the use of the base and limit symbols over non-contiguous memory usually produces unpredictable and undesirable effects.

If you are using a scatter-loading description file, the output section symbols in Table 4-2 are undefined. If your code accesses these symbols, you must treat it as a weak reference.

The standard implementation of `__user_initial_stackheap()` uses the value in `Image$$ZI$$Limit`. Therefore, if you are using a scatter-loading description file you must reimplement `__user_initial_stackheap()` to set the heap and stack boundaries. For more information, see *Region-related symbols* on page 4-3 and the section on library memory models in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

---

### 4.2.3 Importing linker-defined symbols

There are two ways to import linker-defined symbols into your C or C++ source code:

- `extern unsigned int symbol_name;`
- `extern void *symbol_name;`

If you declare a symbol as an `int`, then you must use the address-of operator to obtain the correct value as shown in Example 4-2.

#### Example 4-2 Importing linker-defined symbols

---

```
extern unsigned int Image$$ZI$$Limit
config.heap_base = (unsigned int) &Image$$ZI$$Limit
```

---

## 4.3 Accessing symbols in another image

If you want one image to know the global symbol values of another image, you can use a *symbol definitions* (symdefs) file.

This can be used, for example, if you have one image that always resides in ROM and multiple images that are loaded into RAM. The images loaded into RAM can access global functions and data from the image located in ROM.

This section describes:

- *Creating a symdefs file*
- *Reading a symdefs file* on page 4-8
- *Symdefs file format* on page 4-8.

### 4.3.1 Creating a symdefs file

Use the armlink option `--symdefs filename` to produce a symdefs file.

The linker produces a symdefs file during a successful final link stage. It is not produced for partial linking or for unsuccessful final linking.

#### ————— Note —————

If *filename* does not exist, the file is created containing all the global symbols. If *filename* exists, the existing contents of *filename* are used to select the symbols that are output when the linker rewrites the file. If you do not want this behavior, ensure that any existing symdefs file is deleted before the link step.

### Outputting a subset of the global symbols

By default, all global symbols are written to the symdefs file.

When *filename* exists, the linker uses its contents to restrict the output to a subset of the global symbols. To restrict the output symbols:

1. Specify `--symdefs filename` when you are doing a nearly-final link for *image1*. The linker creates a symdefs file *filename*.
2. Open *filename* in a text editor, remove any symbol entries you do not want in the final list, and save the file.
3. Specify `--symdefs filename` when you are doing a final link for *image1*.

You can edit *filename* at any time to add comments and relink *image1* again, for example, to update the symbol definitions after one or more objects used to create *image1* have changed.

### 4.3.2 Reading a symdefs file

A symdefs file can be considered as an object file with symbol information but no code or data. To read a symdefs file, add it to your file list as you would add any object file. The linker reads the file and adds the symbols and their values to the output symbol table. The added symbols have ABSOLUTE and GLOBAL attributes.

If a partial link is being performed, the symbols are added to the output object symbol table. If a full link is being performed, the symbols are added to the image symbol table.

The linker generates error messages for invalid rows in the file. A row is invalid if:

- any of the columns are missing
- any of the columns have invalid values.

The symbols extracted out of a symdefs file are treated in exactly the same way as symbols extracted from an object symbol table. The same restrictions regarding multiple symbol definitions and ARM/Thumb synonyms apply.

### 4.3.3 Symdefs file format

The symdefs file contains symbols and their values. Unlike other object files, however, it does not contain any code or data.

The file consists of an identification line, optional comments, and symbol information as shown in Example 4-3.

**Example 4-3 Symdefs file format**

---

```
#<SYMDEFS># ARM Linker, RVCT2.2 [Build 304]: Last Updated: Wed Dec 08 15:34:35 2004
;value type name, this is an added comment
0x00008000 A __main
0x00008004 A __scatterload
0x000080e0 T main
0x0000814d T _main_arg
0x0000814d T __argv_alloc
0x00008199 T __rt_get_argv
...
# This is also a comment, blank lines are ignored
...
0x0000a4fc D __stdin
0x0000a540 D __stdout
0x0000a584 D __stderr
0xffffffff N __SIG_IGN
```

---



## Identifying string

If the first 11 characters in the text file are #<SYMDEFS>#, the linker recognizes the file as a symdefs file.

The identifying string is followed by linker version information, and date and time of the most recent update of the symdefs file. The version and update information are not part of the identification string.

## Comments

You can insert comments manually with a text editor. Comments have the following properties:

- The first line must start with the special identifying comment #<SYMDEFS>#. This comment is inserted by the linker when the file is produced and must not be manually deleted.
- Any line where the first non-whitespace character is semicolon (;) or hash (#) is a comment.
- A semicolon (;) or hash (#) after the first non-whitespace character does not start a comment.
- Blank lines are ignored and can be inserted to improve readability.

## Symbol information

The symbol information is provided by the address, type, and name of the symbol on a single line:

<b>Symbol value</b>	The linker writes the absolute address of the symbol in fixed hexadecimal format, for example, 0x00008000. If you edit the file, you can use either hexadecimal or decimal formats for the address value.								
<b>Type flag</b>	The single letter for type is: <table> <tr> <td><b>A</b></td><td>ARM code</td></tr> <tr> <td><b>T</b></td><td>Thumb code</td></tr> <tr> <td><b>D</b></td><td>Data</td></tr> <tr> <td><b>N</b></td><td>Number.</td></tr> </table>	<b>A</b>	ARM code	<b>T</b>	Thumb code	<b>D</b>	Data	<b>N</b>	Number.
<b>A</b>	ARM code								
<b>T</b>	Thumb code								
<b>D</b>	Data								
<b>N</b>	Number.								
<b>Symbol name</b>	The name of the symbol.								

## 4.4 Hiding and renaming global symbols

This section describes how to use a steering file to manage symbol names in output files. For example, you can use steering files to protect intellectual property, or avoid namespace clashes. A steering file is a text file that contains a set of commands to edit the symbol tables of output objects.

Use the `armlink` command-line option `--edit file-list` to specify the steering file (see the description of the `--edit` option in *armlink command syntax* on page 2-9). When you are specifying more than one steering file, the syntax can be either of the following:

```
armlink --edit file1 --edit file2 --edit file3
```

```
armlink --edit file1,file2,file3
```

Do not include spaces between the comma and the filenames.

This section describes:

- *Steering file format*
- *Steering file commands* on page 4-11.

### 4.4.1 Steering file format

A steering file is a plain text file of the following format:

- Lines with a semicolon (;) or hash (#) character as the first non-whitespace character are interpreted as comments. A comment is treated as a blank line.
- Blank lines are ignored.
- Each nonblank, non-comment line is either a command, or part of a command that is split over consecutive nonblank lines.
- Command lines that end with a comma (,) as the last non-whitespace character are continued on the next nonblank line.

Each command line consists of a command, followed by one or more comma-separated operand groups. Each operand group comprises either one or two operands, depending on the command. The command is applied to each operand group in the command. The following rules apply:

- Commands are case-insensitive, but are conventionally shown in uppercase.
- Operands are case-sensitive because they must be matched against case-sensitive symbol names. You can use wildcard characters in operands.

Commands are applied to global symbols only. Other symbols, such as local symbols or `STT_FILE`, are not affected.

### 4.4.2 Steering file commands

Steering file commands enable you to manage symbols in the symbol table, control the copying of symbols from the static symbol table to the dynamic symbol table, and store information about the libraries that a link unit depends on.

---

**Note**

---

The steering file commands control only global symbols. Local symbols are not affected by any command.

---

The following commands are supported:

- *IMPORT* on page 4-12
- *EXPORT* on page 4-13
- *RENAME* on page 4-14
- *RESOLVE* on page 4-15
- *REQUIRE* on page 4-17
- *HIDE* on page 4-18
- *SHOW* on page 4-19.

## IMPORT

The `IMPORT` command specifies that a symbol is defined in a shared object at runtime.

### Syntax

```
IMPORT [pattern {AS} replacement_pattern] [, [pattern AS] replacement_pattern] *
```

where:

*pattern* A string, optionally including wildcard characters (either \* or ?), that matches zero or more undefined global symbols. If *replacement\_pattern* does not match any undefined global symbol, the linker ignores the command. The operand can match only undefined global symbols.

*replacement\_pattern*

A string, optionally including wildcard characters (either \* or ?), to which the undefined global symbol is to be renamed. Wildcards must have a corresponding wildcard in *pattern*. The characters matched by the *pattern* wildcard are substituted for the *replacement\_pattern* wildcard.

For example:

```
IMPORT my_func AS func1
```

imports and renames the undefined symbol `my_func` as `func`.

### Usage

You cannot import a symbol that has been defined in the current shared object or executable. Only one wildcard character (either \* or ?) is permitted in `IMPORT`.

The undefined symbol is included in the dynamic symbol table (as *pattern* if given, otherwise as *replacement\_pattern*), if a dynamic symbol table is present.

---

### Note

The `IMPORT` command only affects undefined global symbols. Symbols that have been resolved by a shared library are implicitly imported into the dynamic symbol table. The linker ignores any `IMPORT` directive that targets an implicitly imported symbol.

---

## EXPORT

The EXPORT command specifies that a symbol can be accessed by other shared objects or executables.

### Syntax

```
EXPORT pattern {[AS replacement_pattern]} [,pattern [AS replacement_pattern] ]*
```

where:

*pattern*            A string, optionally including wildcard characters (either \* or ?), that matches zero or more defined global symbols. If *pattern* does not match any defined global symbol, the linker ignores the command. The operand can match only defined global symbols.

*replacement\_pattern*

A string, optionally including wildcard characters (either \* or ?), to which the defined global symbol is to be renamed. Wildcards must have a corresponding wildcard in *replacement\_pattern*. The characters matched by the *replacement\_pattern* wildcard are substituted for the *pattern* wildcard.

For example:

```
EXPORT my_func AS func1
```

renames and exports the defined symbol my\_func as func1.

### Usage

You cannot export a symbol to a name that already exists. Only one wildcard character (either \* or ?) is permitted in EXPORT.

The defined global symbol is included in the dynamic symbol table (as *replacement\_pattern* if given, otherwise as *pattern*), if a dynamic symbol table is present.

## RENAME

The RENAME command renames defined and undefined global symbol names.

### Syntax

```
RENAME pattern {AS replacement_pattern} [,pattern AS replacement_pattern] *
```

where:

*pattern*        A string, optionally including wildcard characters (either \* or ?), that matches zero or more global symbols. If *pattern* does not match any global symbol, the linker ignores the command. The operand can match both defined and undefined symbols.

*replacement\_pattern*

A string, optionally including wildcard characters (either \* or ?), to which the symbol is to be renamed. Wildcards must have a corresponding wildcard in *pattern*. The characters matched by the *pattern* wildcard are substituted for the *replacement\_pattern* wildcard.

For example, for a symbol named func1:

```
RENAME f* AS my_f*
renames func1 to my_func1.
```

### Usage

You cannot rename a symbol to a symbol name that already exists, even if the target symbol name is being renamed itself. Only one wildcard character (either \* or ?) is permitted in RENAME. For example, given an image containing the symbols func1, func2, and func3:

```
EXPORT func1 AS func2    ;invalid, func2 exists
RENAME func3 AS b2
EXPORT func1 AS func3    ;invalid, func3 exists, even though it is renamed to b2
```

The linker processes the steering file before doing any replacements. You cannot, therefore, use RENAME A AS B on line 1 and then RENAME B AS A on line 2.

## RESOLVE

The RESOLVE command matches specific undefined references to a defined global symbol.

### Syntax

```
RESOLVE pattern {AS defined_pattern}
```

where:

*pattern*            A string, optionally including wildcard characters, that must be matched to a defined global symbol.

*defined\_pattern*

A string, optionally including wildcard characters, that matches zero or more defined global symbols. If *defined\_pattern* does not match any defined global symbol, the linker ignores the command. You cannot match an undefined reference to an undefined symbol.

### Usage

RESOLVE is an extension of the existing `armlink --unresolved` option. The difference is that `--unresolved` enables all undefined references to match one single definition, whereas RESOLVE enables more specific matching of references to symbols.

The undefined references are removed from the output symbol table.

RESOLVE works when performing partial-linking and when linking normally.

For example, you might have the files `file1.c` and `file2.c`, shown in Example 4-4 on page 4-16. Create an `ed.txt` file containing the line `RESOLVE MP3* AS MyMP3*`, and issue the following command:

```
armlink file1.o file2.o --edit ed.txt --unresolved foobar
```

This command has the following effects:

- the references from `file1.o` (`foo`, `MP3_Init()` and `MP3_Play()`) are matched to the definitions in `file2.o` (`foobar`, `MyMP3_Init()` and `MyMP3_Play()` respectively), as specified by the steering file `ed.txt`
- the RESOLVE command in `ed.txt` matches the MP3 functions and the `--unresolved` option matches any other remaining references, in this case, `foo` to `foobar`
- The output symbol table, whether it is an image or a partial object, does not contain the symbols `foo`, `MP3_Init` or `MP3_Play`.

**Example 4-4**

---

file1.c

```
extern int foo;
extern void MP3_Init(void);
extern void MP3_Play(void);
```

```
int main(void)
{
    int x = foo + 1;
    MP3_Init();
    MP3_Play();
    return x;
}
```

file2.c:

```
int foobar;
void MyMP3_Init()
{
}
void MyMP3_Play()
{
}
```

---



## REQUIRE

The REQUIRE command creates a DT\_NEEDED tag in the dynamic array.

### **Syntax**

```
REQUIRE pattern [,pattern] *
```

where:

*pattern*            A string representing a filename. No wildcards are permitted.

### **Usage**

The linker inserts a DT\_NEEDED tag with the value of *pattern* into the dynamic array. This tells the dynamic loader that the file it is currently loading requires *pattern* to be loaded.

## HIDE

The HIDE command makes defined global symbols in the symbol table anonymous.

### Syntax

```
HIDE pattern [,pattern] *
```

where:

*pattern*      A string, optionally including wildcard characters, that matches zero or more defined global symbols. If *pattern* does not match any defined global symbol, the linker ignores the command. You cannot hide undefined symbols.

### Usage

HIDE and SHOW can be used to make certain global symbols anonymous in an output image or partially linked object. Hiding symbols in an object file or library can be useful as a means of protecting intellectual property, as shown in Example 4-5. This example produces a partially linked object with all global symbols hidden, except those beginning with `os_`.

This example can be linked with other objects, provided they do not contain references to the hidden symbols. When symbols are hidden in the output object, SHOW commands in subsequent link steps have no effect on them

The hidden references are removed from the output symbol table.

### Example 4-5

---

steer.txt

```
HIDE *           ; Hides all global symbols
SHOW os_*        ; Shows all symbols beginning with 'os_'
```

```
armlink --partial input_object.o --edit steer.txt -o partial_object.o
```

---

## SHOW

The SHOW command makes global symbols visible that were previously hidden with the HIDE command. This command is useful if you want to unhide a specific symbol that has been hidden using a HIDE command with a wildcard.

### Syntax

```
SHOW pattern [,pattern] *
```

where:

*pattern*        A string, optionally including wildcard characters, that matches zero or more global symbols. If *pattern* does not match any global symbol, the linker ignores the command.

### Usage

The usage of SHOW is closely related to that of HIDE. See *HIDE* on page 4-18 for further information.

## 4.5 Using `$Super$$` and `$Sub$$` to override symbol definitions

There are situations where an existing symbol cannot be modified because, for example, it is located in an external library or in ROM code.

Use the `$Super$$` and `$Sub$$` patterns to patch an existing symbol.

For example, to patch the definition of a function `foo()`, use `$Super$$foo()` and `$Sub$$foo()` as follows:

`$Super$$foo` Identifies the original unpatched function `foo()`. Use this to call the original function directly.

`$Sub$$foo` Identifies the new function that will be called instead of the original function `foo()`. Use this to add processing before or after the original function.

Example 4-6 shows the legacy function `foo()` modified to result in a call to `ExtraFunc()` and a call to `foo()`.

### Example 4-6

---

```
extern void ExtraFunc(void);
extern void $Super$$foo(void):

/* this function will be called instead of the original foo() */
void $Sub$$foo(void)
{
    ExtraFunc(); /* does some extra setup work */
    $Super$$foo(); /* calls the original foo function */
}
```

---

## 4.6 Symbol versioning

The linker conforms to the *Base Platform ABI for the ARM Architecture* [BPABI] and supports GNU-extended symbol versioning model.

Symbol versioning records extra information about symbols imported from a dynamic shared object and symbols exported by a dynamic shared object. This extra information is used by the dynamic loader to make sure that all the symbols needed by an image are available at load time.

Symbol versioning enables shared object creators to produce new versions of symbols that all new clients use, whilst maintaining compatibility with clients linked against old versions of the shared object.

This section describes the command line options that you can use to control symbol versions:

- *Version*
- *Default version*
- *Creating versioned symbols* on page 4-22.

### 4.6.1 Version

Symbol versioning adds the concept of a version to the dynamic symbol table. A version is a name that symbols are associated with. When a dynamic loader tries to resolve a symbol reference associated with a version name, it can only match against a symbol definition with the same version name.

---

#### Note

A version might be associated with previous version names to show the revision history of the shared object.

---

### 4.6.2 Default version

Whilst a shared object might have multiple versions of the same symbol, a client of the shared object can only bind against the latest version.

This is called the *default version* of the symbol.

### 4.6.3 Creating versioned symbols

By default, the linker does not create versioned symbols for a shared object. There are three ways to control versioned symbols:

- *Embedded symbols*
- *Steering file* on page 4-23
- *Filename* on page 4-24.

#### Embedded symbols

You can add specially named symbols to input objects that cause the linker to create symbol versions. These symbols are of the form:

- `name@version` for a non default version of a symbol
- `name@@version` for a default version of a symbol.

You must define these symbols at the address of the function or data as that you want to export. The symbol name is broken into two parts, a symbol name *name* and a version definition *ver*. The *name* is added to the dynamic symbol table and becomes part of the interface to the shared object. *ver* creates a version called *ver* if it does not already exist and associates *name* with the version called *ver*.

For details on how to create version symbols, see the chapter describing:

- how to use the ARM compiler in *RealView Developer Kit v2.2 Compiler and Libraries Guide*
- how to write ARM and Thumb assembly language in *RealView Developer Kit v2.2 Assembler Guide*.

Example 4-7 places the symbols `foo@ver1`, `foo@@ver2` and `bar@@ver1` into the object symbol table:

#### Example 4-7 Creating versioned symbols, embedded symbols

---

```
int old_function(void) __asm__("foo@ver1");
int new_function(void) __asm__("foo@@ver2");
int other_function(void) __asm__("bar@@ver1");
```

---

The linker reads these symbols and creates version definitions `ver1` and `ver2`. The symbol `foo` is associated with a non default version of `ver1`, and with a default version of `ver2`. The symbol `bar` is associated with a default version of `ver1`.

There is no way to create associations between versions with this method.

## Steering file

You can embed the commands to produce symbol versions in a script file that is specified by the command-line option `--symver_script file`. Using this option automatically enables symbol versioning.

The script file supports the same syntax as the GNU *ld* linker.

Using a script file enables you to associate a version with an earlier version.

A steering file can be provided in addition to the embedded symbol method. If you choose to do this then your script file must match your embedded symbols.

Format of steering file in BNF:

```
version_definition ::=
```

```
    version_name "{" symbol_association* "}" [depend_version] ";"
```

The *version\_name* is a string containing the name of the version. *depend\_version* is a string containing the name of a version that this *version\_name* depends on. This version must have already been defined in the script file. Version names are not significant, it helps to choose readable names.

```
symbol_association ::=
```

```
    "local:" | "global:" | symbol_name ";"
```

Where:

- "local:" indicates that all subsequent *symbol\_names* in this version definition are local to the shared object and are not versioned.
- "global:" indicates that all subsequent *symbol\_names* belong to this version definition.

There is an implicit "global:" at the start of every version definition.

- *symbol\_name* is the name of a global symbol in the static symbol table.

Example 4-8 on page 4-24 shows a steering file that corresponds to the embedded symbols example (Example 4-7 on page 4-22) with the addition of dependency information so that *ver2* depends on *ver1*:

**Example 4-8 Creating versioned symbols, steering file**

---

```
ver1 {  global:      foo; bar;   local:      *;   };ver2 {  global:
foo;    } ver1;
```

---

**Errors & warnings**

If you use a script file then the version definitions and symbols associated with them must match. The linker warns you if it detects any mismatch.

**Filename**

Use the command-line option `--symver_soname` to turn on implicit symbol versioning. Use this option if you need to version your symbols in order to force static binding, but where you do not care about the version number that they are given.

Where a symbol has no defined version, the linker uses the SONAME of the file being linked.

This option cannot be combined with embedded symbols or a script file.



# Chapter 5

## Using Scatter-loading Description Files

This chapter describes how you use the ARM® linker, `arm1ink`, and scatter-loading description files to create complex images. It contains the following sections:

- *About scatter-loading* on page 5-2
- *Formal syntax of the scatter-loading description file* on page 5-8
- *Examples of specifying region and section addresses* on page 5-24
- *Equivalent scatter-loading descriptions for simple images* on page 5-35.

## 5.1 About scatter-loading

An image is made up of regions and output sections. Every region in the image can have a different load and execution address (see *Specifying the image structure* on page 3-2).

To construct the memory map of an image, the linker must have:

- grouping information describing how input sections are grouped into regions
- placement information describing the addresses where image regions are to be located in the memory maps.

The scatter-loading mechanism enables you to specify the memory map of an image to the linker. Scatter-loading gives you complete control over the grouping and placement of image components. Scatter-loading can be used for simple images (see *Images with a simple memory map* on page 5-4), but it is generally only used for images that have a complex memory map (see *Images with a complex memory map* on page 5-5). Scatter-loading is capable of describing complex image maps that consist of multiple regions scattered in the memory map at load and execution time.

You specify this information using a scatter-loading description in a text file that is passed to the linker.

This section describes:

- *Symbols defined for scatter-loading*
- *When to use scatter-loading* on page 5-3
- *Scatter-loading command-line option* on page 5-4
- *Images with a simple memory map* on page 5-4
- *Images with a complex memory map* on page 5-5.

### 5.1.1 Symbols defined for scatter-loading

When the linker is creating an image using a scatter-loading description, it creates some region-related symbols. These are described in *Region-related symbols* on page 4-3. These special symbols are created by the linker only if your code references them.

#### ———— Note ————

These symbols are not defined when a scatter-loading description file is used:

- Image\$\$RW\$\$Base
- Image\$\$RW\$\$Limit
- Image\$\$RO\$\$Base
- Image\$\$RO\$\$Limit
- Image\$\$ZI\$\$Base
- Image\$\$ZI\$\$Limit

Because the default implementation uses `Image$$ZI$$Limit`, you must reimplement `__user_initial_stackheap()` and define a value for the start of the heap region and the top of the stack region. For more information see:

- the chapter describing the C and C++ libraries in *RealView Developer Kit v2.2 Compiler and Libraries Guide* (for details on library memory models).

If you do not reimplement `__user_initial_stackheap()`, the following error message is displayed from the linker:

Undefined symbol `Image$$ZI$$Limit` (referred from `sys_stackheap.o`).

---

### 5.1.2 When to use scatter-loading

The command-line options to the linker give some control over the placement of data and code, but complete control of placement requires more detailed instructions than can be entered on the command line. Situations where scatter-loading descriptions are necessary (or very useful) are:

#### Complex memory maps

Code and data that must be placed into many distinct areas of memory require detailed instructions on which section goes into which memory space.

#### Different types of memory

Many systems contain a variety of physical memory devices such as flash, ROM, SDRAM, and fast SRAM. A scatter-loading description can match the code and data with the most appropriate type of memory. For example, the interrupt code might be placed into fast SRAM to improve interrupt response time and infrequently used configuration information might be placed into slower flash memory.

#### Memory-mapped I/O

The scatter-loading description can place a data section at a precise address in the memory map so that memory mapped peripherals can be accessed.

#### Functions at a constant location

A function can be placed at the same location in memory even though the surrounding application has been modified and recompiled.

### Using symbols to identify the heap and stack

Symbols can be defined for the heap and stack location when the application is linked.

Scatter-loading is, therefore, almost always required for implementing embedded systems because these use ROM, RAM, and memory-mapped I/O.

#### 5.1.3 Scatter-loading command-line option

The `armlink` command-line option for using scatter-loading is:

```
--scatter description_file
```

This instructs the linker to construct the image memory map as described in *description\_file*. The format of the description file is given in *Formal syntax of the scatter-loading description file* on page 5-8.

For additional information on scatter-loading description files, see also:

- *Examples of specifying region and section addresses* on page 5-24.
- *Equivalent scatter-loading descriptions for simple images* on page 5-35.

#### 5.1.4 Images with a simple memory map

The scatter-loading description in Figure 5-1 on page 5-5 loads the segments from the object file into memory corresponding to the map shown in Figure 5-2 on page 5-5. The maximum size specifications for the regions are optional but, if they are included, these enable the linker to check that a region has not overflowed its boundary.

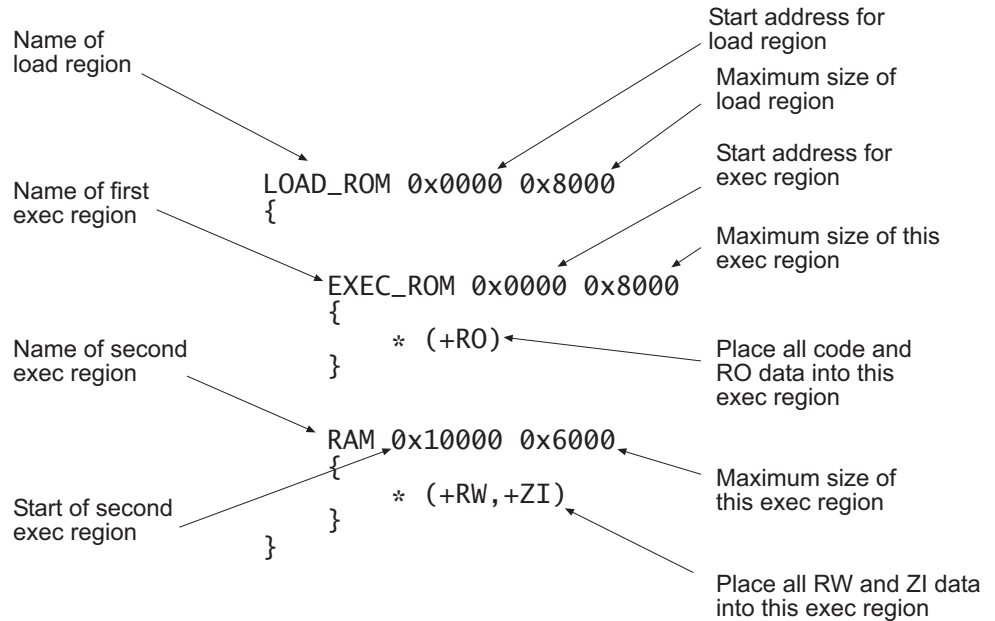


Figure 5-1 Simple memory map in a scatter-loading description file

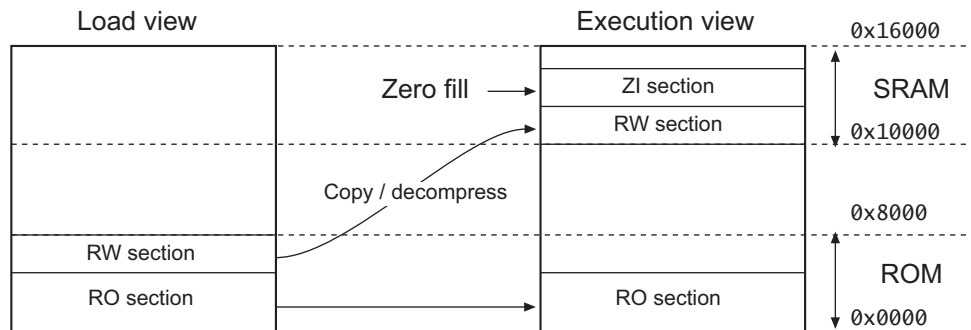


Figure 5-2 Simple scatter-loaded memory map

In this example, the same result can be achieved by specifying `--ro-base 0x0` and `--rw-base 0x10000` as command-line options to the linker.

### 5.1.5 Images with a complex memory map

The scatter-loading description in Figure 5-3 on page 5-6 loads the segments from the `program1.o` and `program2.o` files into memory corresponding to the map shown in Figure 5-4 on page 5-6.

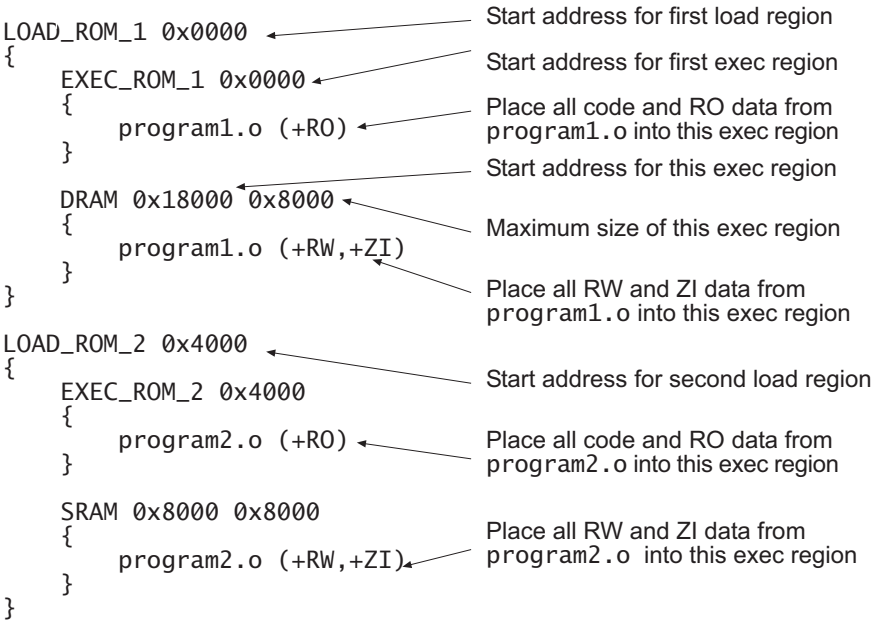


Figure 5-3 Complex memory map in a scatter-loading description file




Figure 5-4 Complex scatter-loaded memory map

Unlike the simple memory map shown in Figure 5-2 on page 5-5, this application cannot be specified to the linker using only the basic command-line options.

---

**Caution**

---

The scatter-loading description in Figure 5-3 on page 5-6 specifies the location for code and data for `program1.o` and `program2.o` only. If you link an additional module, for example, `program3.o`, and use this description file, the location of the code and data for `program3.o` is not specified.

Unless you want to be very rigorous in the placement of code and data, it is advisable to use the `*` or `.ANY` specifier to place leftover code and data (see *Placing regions at fixed addresses* on page 5-26).

---

## 5.2 Formal syntax of the scatter-loading description file

A scatter-loading description file is a text file that describes the memory map of the target embedded product to the linker. The file extension for the description file is not significant if you are using the linker from the command line. The description file enables you to specify the:

- load address and maximum size of each load region
- attributes of each load region
- execution regions derived from each load region
- execution address and maximum size of each execution region
- input sections for each execution region.

The description file format reflects the hierarchy of load regions, execution regions, and input sections.

---

### Note

---

The assignment of input sections to regions is completely independent of the order in which selection patterns are written in the scatter-loading description file. The best match between selection patterns and either file/section names or section attributes wins. See *Resolving multiple matches* on page 5-21.

---

This section describes:

- *BNF notation and syntax* on page 5-9
- *Overview of the syntax of scatter-loading description files* on page 5-9
- *Load region description* on page 5-11
- *Execution region description* on page 5-14
- *Input section description* on page 5-17
- *Resolving multiple matches* on page 5-21.



## 5.2.1 BNF notation and syntax

Table 5-1 summarizes the *Backus Naur Format* (BNF) symbols that are used to describe a formal language.

**Table 5-1 BNF syntax**

Symbol	Description
"	Quotation marks are used to indicate that a character that is normally part of the BNF syntax is used as a literal character in the definition. The definition <code>B"+"C</code> , for example, can only be replaced by the pattern <code>B+C</code> . The definition <code>B+C</code> can be replaced by, for example, patterns <code>BC</code> , <code>BBC</code> , or <code>BBBC</code> .
$A ::= B$	Defines <i>A</i> as <i>B</i> . For example, $A ::= B "+"   C$ means that <i>A</i> is equivalent to either <code>B+</code> or <code>C</code> . The $::=$ notation is used to define a higher level construct in terms of its components. Each component might also have a $::=$ definition that defines it in terms of even simpler components. For example, $A ::= B$ and $B ::= C   D$ means that the definition <i>A</i> is equivalent to the patterns <code>C</code> or <code>D</code> .
[ <i>A</i> ]	Optional element <i>A</i> . For example, $A ::= B[C]D$ means that the definition <i>A</i> can be expanded into either <code>BD</code> or <code>BCD</code> .
<i>A</i> +	Element <i>A</i> can have one or more occurrences. For example, $A ::= B+$ means that the definition <i>A</i> can be expanded into <code>B</code> , <code>BB</code> , or <code>BBB</code> .
<i>A</i> *	Element <i>A</i> can have zero or more occurrences.
<i>A</i>   <i>B</i>	Either element <i>A</i> or <i>B</i> can occur, but not both.
( <i>A</i> <i>B</i> )	Element <i>A</i> and <i>B</i> are grouped together. This is particularly useful when the <code> </code> operator is used or when a complex pattern is repeated. For example, $A ::= (B C)+ (D   E)$ means that the definition <i>A</i> can be expanded into any of <code>BCD</code> , <code>BCE</code> , <code>BCBCD</code> , <code>BCBCE</code> , <code>BCBCBCD</code> , or <code>BCBCBCE</code> .

## 5.2.2 Overview of the syntax of scatter-loading description files

In the BNF definitions in this section, line returns and spaces have been added to improve readability. They are not required in the scatter-loading definition and are ignored if present in the file.

A *scatter\_description* is defined as one or more *load\_region\_description* patterns:

```
Scatter_description ::=
```

```
    load_region_description+
```

A *load\_region\_description* is defined as a load region name, optionally followed by attributes or size specifiers, and one or more execution region descriptions:

```
load_region_description ::=
    load_region_name (base_address | ("+" offset)) [attributes] [max_size]
    "{"
        execution_region_description+
    "}"
```

An *execution\_region\_description* is defined as an execution region name, a base address specification, optionally followed by attributes or size specifiers, and one or more input section descriptions:

```
execution_region_description ::=
    exec_region_name (base_address | "+" offset) [attribute_list] [max_size | "-"
length]
    "{"
        input_section_description*
    "}"
```

An *input\_section\_description* is defined as a source module selector pattern optionally followed by input section selectors:

```
input_section_description ::=
    module_select_pattern
    ["("
        ("+" input_section_attr | input_section_pattern | input_symbol_pattern)
        ("," "+" input_section_attr | "," input_section_pattern | ","
input_symbol_pattern)*
    ")"]
```

Figure 5-5 on page 5-11 shows the contents and organization of a typical scatter-loading description file.

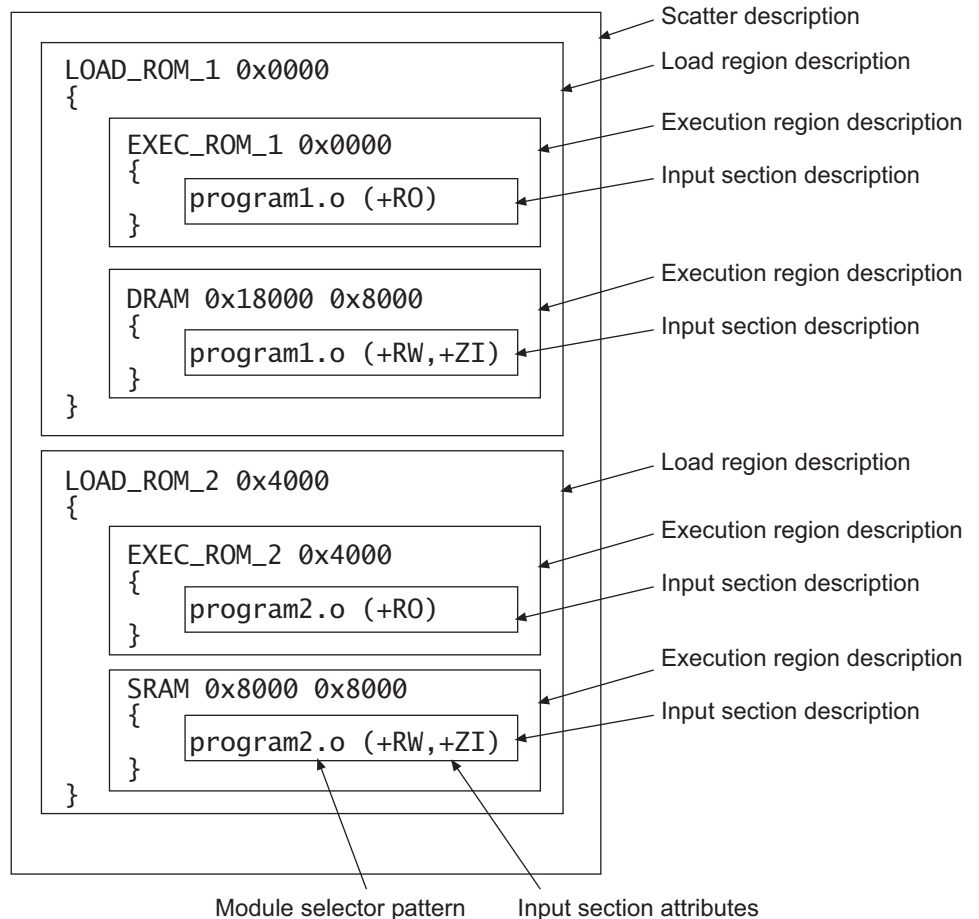


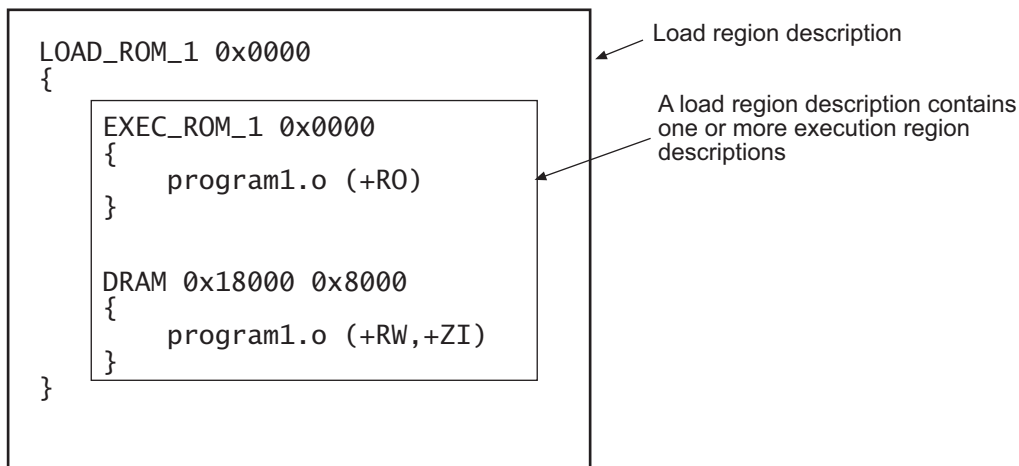
Figure 5-5 Components of a scatter-loading file definition

### 5.2.3 Load region description

A load region has:

- a name (this is used by the linker to identify different load regions)
- a base address (the start address for the code and data in the load view)
- attributes (optional)
- a maximum size (optional)
- a list of execution regions (these identify the type and location of modules in the execution view).

Figure 5-6 shows the components of a typical load region description.



**Figure 5-6 Components of a load region description**

The syntax, in BNF, is:

```

load_region_description ::= load_region_name (base_address | ("+" offset))
[attribute_list] [ max_size ]
    "{"
        execution_region_description+
    "}"
  
```

where:

**load\_region\_name** The linker generates a Load\$\$exec\_region\_name\$\$base symbol for each execution region. This symbol holds the load address of the execution region (see *Execution region description* on page 5-14). The *load\_region\_name*, however, is used only to identify each region, that is, it is not used to generate Load\$\$region\_name symbols.

———— **Note** ————

An image created for use by a debugger requires a unique base address for each region because the debugger must load regions at their load addresses. Overlapping load region addresses result in part of the image being overwritten.

A loader or operating system, however, can correctly load overlapping position-independent regions. One or more of the position-independent regions is automatically moved to a different address.

---


<i>base_address</i>	This specifies the address where objects in the region are to be linked. <i>base_address</i> must be word-aligned.										
<i>+offset</i>	Describes a base address that is <i>offset</i> bytes beyond the end of the preceding load region. The value of <i>offset</i> must be zero modulo four. If this is the first load region, then <i>+offset</i> means that the base address begins <i>offset</i> bytes after zero.										
<i>attribute_list</i>	<p>This specifies the properties of the load region contents:</p> <table> <tr> <td>ABSOLUTE</td><td>Absolute address.</td></tr> <tr> <td>PI</td><td>Position-independent.</td></tr> <tr> <td>RELOC</td><td>Relocatable.</td></tr> <tr> <td>OVERLAY</td><td>Overlaid.</td></tr> <tr> <td>NOCOMPRESS</td><td>Should not be compressed.</td></tr> </table> <p>You can specify only one of the attributes ABSOLUTE, PI, RELOC, and OVERLAY. The default load region attribute is ABSOLUTE.</p> <p>Load regions that have one of PI, RELOC, or OVERLAY attributes can have overlapping address ranges. The linker faults overlapping address ranges for ABSOLUTE load regions.</p> <p>The OVERLAY keyword enables you to have multiple execution regions at the same address. ARM does not provide an overlay mechanism in RVDK. To use multiple execution regions at the same address, you must provide your own overlay manager.</p> <p>RW data compression is enabled by default. The NOCOMPRESS keyword enables you to specify that a load region should not be compressed in the final image.</p>	ABSOLUTE	Absolute address.	PI	Position-independent.	RELOC	Relocatable.	OVERLAY	Overlaid.	NOCOMPRESS	Should not be compressed.
ABSOLUTE	Absolute address.										
PI	Position-independent.										
RELOC	Relocatable.										
OVERLAY	Overlaid.										
NOCOMPRESS	Should not be compressed.										
<i>max_size</i>	This specifies the maximum size of the load region. (If the optional <i>max_size</i> value is specified, armlink generates an error if the region has more than <i>max_size</i> bytes allocated to it.)										
<i>execution_region_description</i>	<p>This specifies the execution region name, address, and contents. See <i>Execution region description</i> on page 5-14.</p>										

## 5.2.4 Execution region description

An execution region has:

- a name
- a base address (either absolute or relative)
- an optional maximum size specification
- attributes that specify the properties of the execution region
- one or more input section descriptions (the modules placed into this execution region).

Figure 5-7 shows the components of a typical execution region description.



**Figure 5-7 Components of an execution region description**

The syntax, in BNF, is:

*execution\_region\_description* ::=

```
    exec_region_name (base_address | "+" offset) [attribute_list] [max_size | "-"
length]
    "{"
        input_section_description*
    "}"
```

where:

<i>exec_region_name</i>	This names the execution region.
<i>base_address</i>	This specifies the address where objects in the region are to be linked. <i>base_address</i> must be word-aligned.
<i>+offset</i>	Describes a base address that is <i>offset</i> bytes beyond the end of the preceding execution region. The value of <i>offset</i> must be zero modulo four.  If there is no preceding execution region (that is, if this is the first execution region in the load region) then <i>+offset</i> means that the base address begins <i>offset</i> bytes after the base of the containing load region.

If *+offset* form is used and the encompassing load region has the RELOC attribute, the execution region inherits the RELOC attribute. However, if a fixed *base\_address* is used, future occurrences of *offset* do not inherit the RELOC attribute.

<i>attribute_list</i>	This specifies the properties of the execution region contents:	
ABSOLUTE		Absolute address. The execution address of the region is specified by the base designator.
PI		Position-independent.
OVERLAY		Overlaid.
FIXED		Fixed address. Both the load address and execution address of the region is specified by the base designator (the region is a root region. See <i>Creating root execution regions</i> on page 5-25). The base designator must be either an absolute base address, or an offset of +0.
EMPTY		This reserves an empty block of memory of a given length in the execution region, typically used by a heap or stack. See <i>Reserving an empty region</i> on page 5-31 for further information.
PADVALUE		<p>This defines the value of any padding. If you specify PADVALUE, you must give a value, for example:</p> <pre>EXEC 0x10000 PADVALUE 0xffffffff EMPTY ZEROPAD 0x2000</pre> <p>This creates a region of size 0x2000 full of 0xffffffff.</p> <p>PADVALUE must be a word in size. PADVALUE attributes on load regions are ignored.</p>
ZEROPAD		<p>Zero-initialized sections are written in the ELF file as a block of zeros and, therefore, do not have to be zero-filled at runtime.</p> <p>In certain situations, for example, simulation, this is preferable to spending a long time in a zeroing loop.</p>
NOCOMPRESS		Should not be compressed.
UNINIT		Must not be zero-initialized.

<i>max_size</i>	This is an optional number that instructs the linker to generate an error if the region has more than <i>max_size</i> bytes allocated to it.
<i>-length</i>	If the length is given as a negative value, the <i>base_address</i> is taken to be the end address of the region. Typically used with <i>EMPTY</i> to represent a stack that grows down in memory. See <i>Reserving an empty region</i> on page 5-31 for more information.
<i>input_section_description</i>	This specifies the content of the input sections. See <i>Input section description</i> on page 5-17.

When specifying the properties of the execution region:

- You can specify only one of the attributes *PI*, *OVERLAY*, *FIXED*, and *ABSOLUTE*. Unless one of the attributes *PI*, *FIXED*, or *OVERLAY* is specified, *ABSOLUTE* is the default attribute of the execution region.
- Execution regions that use the *+offset* form of the base designator either inherit the attributes of the preceding execution region, (or of the containing load region if this is the first execution region in the load region), or have the *ABSOLUTE* attribute.
- Only root execution regions can be zero-initialized using the *ZEROPAD* attribute. Using the *ZEROPAD* attribute with a non-root execution region generates a warning and the attribute is ignored.
- The attribute *RELOC* cannot be explicitly specified for execution regions. The region can only be *RELOC* by inheriting the attribute from a load region.
- It is not possible for an execution region that uses the *+offset* form of the base designator to have its own attributes (other than the *ABSOLUTE* attribute that overrides inheritance). Use the combination *+0 ABSOLUTE* to set a region to *ABSOLUTE* without changing the start location.
- Execution regions that are specified as *PI* or *OVERLAY* (or that have inherited the *RELOC* attribute) are permitted to have overlapping address ranges. The linker faults overlapping address ranges for *ABSOLUTE* and *FIXED* execution regions.
- *RW* data compression is enabled by default. The *NOCOMPRESS* keyword enables you to specify that an execution region must not be compressed in the final image.
- *UNINIT* specifies that any *ZI* output section in the execution region must not be zero-initialized. Use this to create execution regions containing uninitialized data or memory-mapped I/O.




## 5.2.5 Input section description

An input section description is a pattern that identifies input sections by:

- Module name (object filename, library member name, or library filename). The module name can use wildcard characters.
- Input section name, or input section attributes such as READ-ONLY, or CODE.
- Symbol name.

Figure 5-8 shows the components of a typical input section description.



**Figure 5-8 Components of an input section description**

The syntax, in BNF, is:

*input\_section\_description* ::=

```

module_select_pattern
["("
  ("+" input_section_attr | input_section_pattern | input_symbol_pattern)
  ("," "+" input_section_attr | "," input_section_pattern | ","
    input_symbol_pattern)*
  ")"]
  
```

where:

*module\_select\_pattern*

This is a pattern constructed from literal text. The wildcard character \* matches zero or more characters and ? matches any single character.

Matching is case-insensitive, even on hosts with case-sensitive file naming.

Use \*.o to match all objects. Use \* to match all object files and libraries.

An input section matches a module selector pattern when *module\_select\_pattern* matches one of the following:

- The name of the object file containing the section.
- The name of the library member (without leading path name).

- The full name of the library (including path name) the section was extracted from. If the names contain spaces, use wildcards to simplify searching. For example, use `*libname.lib` to match `C:\lib dir\libname.lib`.

The special module selector pattern `.ANY` enables you to assign input sections to execution regions without considering their parent module. Use `.ANY` to fill up the execution regions with *do not care* assignments.

---

**Note**

---

- Only input sections that match both *module\_select\_pattern* and at least one *input\_section\_attr* or *input\_section\_pattern* are included in the execution region.  
If you omit (+ *input\_section\_attr*) and (*input\_section\_pattern*), the default is `+R0`.
  - Do not rely on input section names generated by the compiler, or used by ARM library code. These can change between compilations if, for example, different compiler options are used. In addition, section naming conventions used by the compiler are not guaranteed to remain constant between releases.
- 

### *input\_section\_attr*

This is an attribute selector matched against the input section attributes. Each *input\_section\_attr* follows a +.

If you are specifying a pattern to match the input section name, the name must be preceded by a +. You can omit any comma immediately followed by a +.

The selectors are not case-sensitive. The following selectors are recognized:

- `R0-CODE`
- `R0-DATA`
- `R0`, selects both `R0-CODE` and `R0-DATA`
- `RW-DATA`
- `RW-CODE`
- `RW`, selects both `RW-CODE` and `RW-DATA`
- `ZI`
- `ENTRY`, that is, a section containing an `ENTRY` point.

The following synonyms are recognized:

- `CODE` for `R0-CODE`
- `CONST` for `R0-DATA`

- TEXT for R0
- DATA for RW
- BSS for ZI.

The following pseudo-attributes are recognized:

- FIRST
- LAST.

FIRST and LAST can be used to mark the first and last sections in an execution region if the placement order is important (for example, if a specific input section must be first in the region and an input section containing a checksum must be last). The first occurrence of FIRST or LAST as an *input\_section\_attr* terminates the list.

The special module selector pattern .ANY enables you to assign input sections to execution regions without considering their parent module. Use one or more .ANY patterns to fill up the execution regions with *do not care* assignments. In most cases, using a single .ANY is equivalent to using the \* module selector.

You cannot have two \* selectors in a scatter-loading description file. You can, however, use two modified selectors, \*A and \*B, for example, and you can use a .ANY selector together with a \* module selector. The \* module selector has higher precedence than .ANY. If the portion of the file containing the \* selector is removed, the .ANY selector then becomes active.

The input section descriptions having the .ANY module selector pattern are resolved after all other (non-.ANY) input section descriptions have been resolved and input sections have been assigned to the closest matching execution region. If more than one .ANY pattern is present, the linker fills the first .ANY with as much as possible and then begins filling the next .ANY.

Each remaining unassigned input section is assigned to the execution region with the following characteristics:

- the biggest remaining space (determined by the value of *max\_size* and the sizes of the input sections already assigned to it)
- a matching .ANY *input\_section\_description*
- memory access attributes (if they exist) matching the memory attributes of the input section.

#### *input\_section\_pattern*

This is a pattern that is matched, without case sensitivity, against the input section name. It is constructed from literal text. The wildcard character \* matches 0 or more characters, and ? matches any single character.

---

**Note**

---

If you use *input\_section\_patterns*, ensure that there are no duplicate patterns in different execution regions in order to avoid ambiguity errors.

---

*input\_symbol\_pattern*

You can select the input section by the name of a global symbol that the section defines. This enables you to choose individual sections with the same name from partially linked objects.

The `:gdef:` prefix distinguishes a global symbol pattern from a section pattern. For example, use `:gdef:mysym` to select the section that defines `mysym`. The following example shows a description file in which `ExecReg1` contains the section that defines global symbol `mysym1`, and the section that contains global symbol `mysym2`:

```
LoadRegion 0x8000
{
    ExecReg1 +0
    {
        *(:gdef:mysym1)
        *(:gdef:mysym2)
    }
    ExecReg2
    {
        *(+R0)
    }
    ExecReg3
    {
        *(+RW)
    }
    ExecReg4
    {
        *(+ZI)
    }
}
```

---

**Note**

---

If you use *input\_symbol\_patterns*, ensure that there are no duplicate patterns in different execution regions in order to avoid ambiguity errors.

---

## 5.2.6 Resolving multiple matches

If a section matches more than one execution region, the matches are resolved as described below. If a unique match cannot be found, the linker faults the scatter-loading description. Each section is selected by a *module\_selector\_pattern* and an *input\_section\_selector*.

Examples of *module\_selector\_pattern* specifications are:

- \* matches any module or library
- \*.o matches any object module
- math.o matches the math.o module
- \*math.lib matches any library path ending with math.lib (for example, C:\apps\lib\math\satmath.lib).

Examples of *input\_section\_selector* specifications are:

- +RO is an input section attribute that matches all RO code and all RO data
- +RW,+ZI is an input section attribute that matches all RW code, all RW data, and all ZI data
- BLOCK\_42 is an input section pattern that matches the assembly file area named BLOCK\_42.

### ————— Note —————

The compiler produces areas that can be identified by input section patterns such as .text, .data, .constdata, and .bss. These names, however, might change in the future and you should avoid using them.

If you want to match a specific function or extern data from a C or C++ file, either:

- compile the function or data in a separate module and match the module object name
- use #pragma arm section to specify the name of the section containing the code or data of interest. See the chapter describing the ARM compiler reference in *RealView Developer Kit v2.2 Compiler and Libraries Guide* for more information on pragmas.

The following variables are used to describe multiple matches:

- m1 and m2 represent module selector patterns
- s1 and s2 represent input section selectors.

In the case of multiple matches, the linker determines the region to assign the input section to on the basis of the *module\_selector\_pattern* and *input\_section\_selector* pair that is the most specific.

For example, if input section A matches *m1,s1* for execution region R1, and A matches *m2,s2* for execution region R2, the linker:

- assigns A to R1 if *m1,s1* is more specific than *m2,s2*
- assigns A to R2 if *m2,s2* is more specific than *m1,s1*
- diagnoses the scatter-loading description as faulty if *m1,s1* is not more specific than *m2,s2* and *m2,s2* is not more specific than *m1,s1*.

The sequence *arm1ink* uses to determine the most specific *module\_selector\_pattern*, *input\_section\_selector* pair is as follows:

1. For the module selector patterns:  
*m1* is more specific than *m2* if the text string *m1* matches pattern *m2* and the text string *m2* does not match pattern *m1*.
2. For the input section selectors:
  - If *s1* and *s2* are both patterns matching section names, the same definition as for module selector patterns is used.
  - If one of *s1*, *s2* matches the input section name and the other matches the input section attributes, *s1* and *s2* are unordered and the description is diagnosed as faulty.
  - If both *s1* and *s2* match input section attributes, the determination of whether *s1* is more specific than *s2* is defined by the relationships below:
    - ENTRY is more specific than R0-CODE, R0-DATA, RW-CODE or RW-DATA
    - R0-CODE is more specific than R0
    - R0-DATA is more specific than R0
    - RW-CODE is more specific than RW
    - RW-DATA is more specific than RW
    - There are no other members of the (*s1* more specific than *s2*) relationship between section attributes.
3. For the *module\_selector\_pattern*, *input\_section\_selector* pair, *m1,s1* is more specific than *m2,s2* only if any of the following are true:
  - *s1* is a literal input section name (that is, it contains no pattern characters) and *s2* matches input section attributes other than +ENTRY
  - *m1* is more specific than *m2*
  - *s1* is more specific than *s2*.

This matching strategy has the following consequences:

- Descriptions do not depend on the order they are written in the file.
- Generally, the more specific the description of an object, the more specific the description of the input sections it contains.
- The *input\_section\_selectors* are not examined unless:
  - Object selection is inconclusive.
  - One selector fully names an input section and the other selects by attribute. In this case, the explicit input section name is more specific than any attribute, other than ENTRY, that selects exactly one input section from one object. This is true even if the object selector associated with the input section name is less specific than that of the attribute.

Example 5-1 shows multiple execution regions and pattern matching.

#### Example 5-1 Multiple execution regions and pattern matching

---

```

LR_1 0x040000
{
    ER_ROM 0x040000          ; The startup exec region address is the same
    {                        ; as the load address.
        application.o (+ENTRY) ; The section containing the entry point from
    }                          ; the object is placed here.
    ER_RAM1 0x048000
    {
        application.o (+R0-CODE) ; Other R0 code from the object goes here
    }
    ER_RAM2 0x050000
    {
        application.o (+R0)      ; The R0 data goes here
    }
    ER_RAM3 0x060000
    {
        application.o (+RW)      ; RW code and data go here
    }
    ER_RAM4 +0                ; Follows on from end of ER_R3
    {
        *.o (+R0, +RW, +ZI)      ; Everything except for application.o goes here
    }
}

```

---

## 5.3 Examples of specifying region and section addresses

This section describes how to use a scatter-loading description file to specify addresses for:

- veneers
- RO constants in ROM
- root execution regions.

This section describes:

- *Selecting veneer input sections in scatter-loading descriptions*
- *Creating root execution regions* on page 5-25
- *Placing regions at fixed addresses* on page 5-26
- *Using overlays to place sections* on page 5-29
- *Assigning sections to a root region* on page 5-30
- *Reserving an empty region* on page 5-31
- *Using preprocessing directives* on page 5-33.

### 5.3.1 Selecting veneer input sections in scatter-loading descriptions

Veneers are used to switch between ARM and Thumb® code or to perform a longer program jump than can be specified in a single instruction (see *Veneer generation* on page 3-19). Use a scatter-loading description file to place linker-generated veneer input sections. At most, one execution region in the scatter-loading description file can have the `*(Veneer$$Code)` section selector.

If it is safe to do so, the linker places veneer input sections into the region identified by the `*(Veneer$$Code)` section selector. It might not be possible for a veneer input section to be assigned to the region because of address range problems or execution region size limitations. If the veneer cannot be added to the specified region, it is added to the execution region containing the relocated input section that generated the veneer.

---

#### **Note**

Instances of `*(IW$$Code)` in scatter-loading description files from earlier versions of ARM tools are automatically translated into `*(Veneer$$Code)`. Use `*(Veneer$$Code)` in new descriptions.

---



### 5.3.2 Creating root execution regions

If you specify an initial entry point for an image, or if the linker creates an initial entry point because you have used only one ENTRY directive, you must ensure that the entry point is located in a root region. A root region is a region having the same load and execution address. If the initial entry point is not in a root region, the link fails and the linker gives an error message such as:

Entry point (0x00000000) lies within non-root region ER\_ROM

To specify that a region is a root region in a scatter-loading description file you can either:

- Specify ABSOLUTE, either explicitly or by permitting it to default, as the attribute for the execution region and use the same address for the first execution region and the enclosing load region. To make the execution region address the same as the load region address, either:
  - specify the same numeric value for both the base designator (address) for the execution region and the base designator (address) for the load region
  - specify a +0 offset for the first execution region in the load region.
 If an offset of zero (+0) is specified for all execution regions, then they will all be root regions.

See Example 5-2.

- Use the FIXED execution region attribute to ensure that the load address and execution address of a specific region are the same. See Example 5-3 on page 5-26 and Figure 5-9 on page 5-26.

You can use the FIXED attribute to place any execution region at a specific address in ROM. See *Placing regions at fixed addresses* on page 5-26 for more information.

#### Example 5-2 Specifying the same load and execution address

---

```

LR_1 0x040000      ; load region starts at 0x40000
{                  ; start of execution region descriptions
    ER_RO 0x040000  ; load address = execution address
    {
        * (+R0)     ; all R0 sections (must include section with
                    ; initial entry point)
    }
                    ; rest of scatter description...
}

```

---

**Example 5-3 Using the FIXED attribute**



---

```

LR_1 0x040000      ; load region starts at 0x40000
{
    ; start of execution region descriptions
    ER_RO 0x040000  ; load address = execution address
    {
        * (+R0)      ; R0 sections other than those in init.o
    }
    ER_INIT 0x080000 FIXED ; load address and execution address of this
                          ; execution region are fixed at 0x80000
    {
        init.o(+R0)   ; all R0 sections from init.o
    }
    ; rest of scatter description...
}

```

---

**Figure 5-9 Memory Map for fixed execution regions****5.3.3 Placing regions at fixed addresses**

You can use the **FIXED** attribute in an execution region scatter-loading description file to create root regions that load and execute at fixed addresses.

**FIXED** is used to create multiple root regions within a single load region (and therefore typically a single ROM device). You can use this, for example, to place a function or a block of data, such as a constant table or a checksum, at a fixed address in ROM, so that it can be accessed easily through pointers.

If you specify, for example, that some initialization code is to be placed at start of ROM and a checksum at the end of ROM, some of the memory contents might be unused. Use the **\*** or **.ANY** module selector to flood fill the region between the end of the initialization block and the start of the data block.

---

**Note**


---

To make your code easier to maintain and debug, use the minimum amount of placement specifications in scatter-loading description files and leave the detailed placement of functions and data to the linker.

You cannot specify component objects that have been partially linked. For example, if you partially link the objects `obj1.o`, `obj2.o`, and `obj3.o` together to produce `obj_all.o`, the resulting component object names are discarded in the resulting object. Therefore, you cannot refer to one of the objects by name, for example, `obj1.o`. You can only refer to the combined object `obj_all.o`.

---

### Placing functions and data at specific addresses

Normally, the compiler produces RO, RW, and ZI sections from a single source file. These regions contain all the code and data from the source file. To place a single function or data item at a fixed address, you must enable the linker to process the function or data separately from the rest of the input files. To access an individual object, you can:

- Place the function or data item in its own source file.
- Use the `--split_sections` compiler option to produce an object file for each function (see *RealView Developer Kit v2.2 Compiler and Libraries Guide*).  
This option increases code size slightly (typically by a few percent) for some functions because it reduces the potential for sharing addresses, data, and string literals between functions. However, this can help to reduce the final image size overall by enabling the linker to remove unused functions when you specify `armlink --remove`.
- Use `#pragma arm` section inside the C or C++ source code to create multiple named sections (see Example 5-5 on page 5-29).
- Use the AREA directive from assembly language. For assembly code, the smallest locatable unit is an AREA (see *RealView Developer Kit v2.2 Assembler Guide*).

### Placing the contents of individual object files

The scatter-loading description file in Example 5-4 on page 5-28 places:

- initialization code at address `0x0` (followed by the remainder of the RO code and all of the RO data except for the RO data in the object `data.o`)
- all global RW variables in RAM at `0x400000`

- a table of RO-DATA from data.o fixed at address 0x1FF00.

#### Example 5-4 Section placement

---

```

LOADREG1 0x0 0x10000
{
    EXECREG1 0x0 0x2000          ; Root Region, containing init code
    {
        init.o (Init, +FIRST)    ; place init code at exactly 0x0
        * (+R0)                  ; rest of code and read-only data
    }
    RAM 0x400000                ; RW data to be placed at 0x400000
    {
        * (+RW +ZI)
    }
    DATABLOCK 0x1FF00 FIXED 0xFF ; execution region fixed at 0x1FF00
    {
        data.o(+R0-DATA)         ; maximum space available for table is 0xFF
        ; place RO data between 0x1FF00 and 0x1FFFF
    }
}

```

---

#### Note

There are some situations where using FIXED and a single load region are not appropriate. Other techniques for specifying fixed locations are:

- If your loader can handle multiple load regions, place the RO code or data in its own load region.
  - If you do not require the function or data to be at a fixed location in ROM, use ABSOLUTE instead of FIXED. The loader then copies the data from the load region to the specified address in RAM. (ABSOLUTE is the default attribute.)
  - To place a data structure at the location of memory-mapped I/O, use two load regions and specify UNINIT. (UNINIT ensures that the memory locations are not initialized to zero.)
- 

### Using the arm section pragma

Placing a code or data object in its own source file and then placing the object file sections uses standard coding techniques. However, you can also use a pragma and scatter-loading description file to place named sections. Create a module (adder.c, for example) and name a section explicitly as shown in Example 5-5 on page 5-29.

**Example 5-5 Naming a section**


---

```
// file adder.c
    int x1 = 5;                // in .data
    int y1[100];              // in .bss
    int const z1[3] = {1,2,3}; // in .constdata
    int sub1(int x) {return x-1;} // in .text

    #pragma arm section rwdata = "foo", code = "foo"
    int x2 = 5;                // in foo (data part of region)
    char *s3 = "abc";          // s3 in foo, "abc" in .constdata
    int add1(int x) {return x+1;} // in foo (.text part of region)
    #pragma arm section code, rwdata // return to default placement
```

---

Use a scatter-loading description file to specify where the named section is placed (see Example 5-6). If both code and data sections have the same name, the code section is placed first.

**Example 5-6 Placing a section**


---

```
FLASH 0x24000000 0x40000000
{
    FLASH 0x24000000 0x40000000
    {
        init.o (Init, +First)        ; place area Init from init.o first
        * (+R0)                      ; sub1(), z1[]
    }
    32bitRAM 0x0000
    {
        vectors.o (Vect, +First)
        * (+RW,+ZI)                  ; x1, y1
    }
    ADDER 0x08000000
    {
        adder.o (foo)                ; x2, string s3, and add1()
    }
}
```

---

**5.3.4 Using overlays to place sections**

You can use the OVERLAY keyword in a scatter-loading description file to define how sections are placed in your image. Example 5-7 on page 5-30 defines a static section in RAM followed by a series of overlays. Here, only one of these sections is instantiated at runtime.

**Example 5-7 Specifying a root region**


---

```

EMB_APP 0x8000
{
    ; start of scatter description...
}
STATIC_RAM 0x0 ; contains most of the RW and ZI code/data
{
    * (+RW,+ZI)
}
OVERLAY_A_RAM 0x1000 OVERLAY ; start address of overlay...
{
    module1.o (+RW,+ZI)
}
OVERLAY_B_RAM 0x1000 OVERLAY
{
    module2.o (+RW,+ZI)
}
; rest of scatter description...
}

```

---

If the length of the static area is unknown, use a zero relative offset to specify the start address of an overlay so that it is placed immediately after the end of the static section, for example:

```
OVERLAY_A_RAM +0 OVERLAY
```

In this case, consecutive overlay regions with the same *offset* are placed at *offset* bytes from the previous non-overlay region or start of load region. Do this to avoid unused RAM (where the static area is small) or to prevent overwriting the static area with an overlay (where the static area is large).

**5.3.5 Assigning sections to a root region**

In previous releases of RVCT, the only sections that had to be root were `__main` and the region tables. However, with the implementation of RW data compression there are more files that must be placed in a root region. The linker can detect these sections automatically.

Use a scatter-loading description file to specify a root section in the same way as a named section. Example 5-8 on page 5-31 defines a section must be root using the section selector `InRoot$$Sections`.

**Example 5-8 Specifying a root region**


---

```

LR_APP 0x8000                                ; load region starts at 0x8000
{
    ER_ROOT 0x8000
    {
        * (InRoot$$Sections)    ; place section
    }
    ER_WORLD 0x100000
    {
        * (RO,+RW,+ZI)

    }

    ; rest of scatter description...
}

```

---

**5.3.6 Reserving an empty region**

You can use the `EMPTY` attribute in an execution region scatter-loading description to reserve an empty block of memory for the stack.

The block of memory does not form part of the load region, but is assigned for use at execution time. Because it is created as a dummy `ZI` region, the linker uses the following symbols to access it:

- `Image$$region_name$$ZI$$Base`
- `Image$$region_name$$ZI$$Limit`
- `Image$$region_name$$ZI$$Length`.

If the length is given as a negative value, the address is taken to be the end address of the region. This should be an absolute address and not a relative one. For example, the execution region definition `STACK 0x800000 EMPTY -0x10000` in Example 5-9 on page 5-32 defines a region called `STACK` that starts at address `0x7F0000` and ends at address `0x800000`. Figure 5-10 on page 5-32 is a diagrammatic representation of this example.

**Note**

The dummy `ZI` region that is created for an `EMPTY` execution region is not initialized to zero at runtime.

If the address is in relative (`+n`) form and the length is negative, the linker generates an error.

**Example 5-9 Reserving a region for the stack**


---

```


LR_1 0x80000                                ; load region starts at 0x80000
{
    STACK 0x800000 EMPTY -0x10000           ; region ends at 0x800000 because of the
                                           ; negative length. The start of the region
                                           ; is calculated using the length.

    {                                       ; Empty region used to place stack
    }
    HEAP +0 EMPTY 0x10000                  ; region starts at the end of previous
                                           ; region. End of region calculated using
                                           ; positive length

    {                                       ; Empty region used to place heap
    }
                                           ; rest of scatter description...
}

```

---

**Figure 5-10 Reserving a region for the stack**

In this example, the linker generates the symbols:

```

Image$$STACK$$ZI$$Base      = 0x7f000
Image$$STACK$$ZI$$Limit     = 0x80000
Image$$STACK$$ZI$$Length    = 0x1000
Image$$HEAP$$ZI$$Base       = 0x80000
Image$$HEAP$$ZI$$Limit      = 0x81000
Image$$HEAP$$ZI$$Length     = 0x1000

```



---

**Note**

---

The EMPTY attribute applies only to an execution region. The linker generates a warning and ignores an EMPTY attribute used in a load region definition.

The linker checks that the address space used for the EMPTY region does not coincide with any other execution region.

---

### 5.3.7 Using preprocessing directives

Use the first line in the scatter-loading description file to specify a preprocessor that the linker invokes to process the file. This command is of the form:

```
#! <preprocessor> [pre_processor_flags]
```

Most typically the command is:

```
#! armcc -E
```

The linker can carry out simple expression evaluation with a restricted set of operators, that is, +, -, \*, /, AND, OR, and parentheses. The implementation of OR and AND follows C operator precedence rules.

You can add preprocessing directives to the top of the scatter-loading description file. For example:

```
#define ADDRESS 0x20000000
#include "include_file_1.h"
```

The linker parses the preprocessed scatter-loading description file where these are treated as comments and ignored.

Consider the following simple example:

```
#define AN_ADDRESS (BASE_ADDRESS+(ALIAS_NUMBER*ALIAS_SIZE))
```

Use the directives:

```
#define BASE_ADDRESS 0x8000#define ALIAS_NUMBER 0x2#define ALIAS_SIZE 0x400
```

If the scatter-loading description file contains:

```
LOAD_FLASH AN_ADDRESS ; start address
```

Following preprocessing, this evaluates to:

```
LOAD_FLASH ( 0x8000 + ( 0x2 * 0x400 )) ; start address
```

Following evaluation, the linker parses the scatter-loading file to produce the load region:

LOAD\_FLASH 0x8808 ; start address

## 5.4 Equivalent scatter-loading descriptions for simple images

The command-line options `--reloc`, `--ro-base`, `--rw-base`, `--ropi`, `--rwpi`, and `--split` create the simple image types described in *Using command-line options to create simple images* on page 3-25. You can create the same image types by using the `--scatter` command-line option and a file containing one of the corresponding scatter-loading descriptions.

This section describes:

- *Type 1*
- *Type 2* on page 5-37
- *Type 3* on page 5-39.

### 5.4.1 Type 1

An image of this type consists of a single load region in the load view and three execution regions in the execution view. The execution regions are placed contiguously in the memory map.

`--ro-base address` specifies the load and execution address of the region containing the RO output section. Example 5-10 shows the scatter-loading description equivalent to using `--ro-base 0x040000`.

#### Example 5-10 Single load region and contiguous execution regions

---

```

LR_1 0x040000    ; Define the load region name as LR_1, the region starts at 0x040000.
{
    ER_RO +0      ; First execution region is called ER_RO, region starts at end of previous region.
                  ; However, since there was no previous region, the address is 0x040000.
    {
        * (+RO)   ; All RO sections go into this region, they are placed consecutively.
    }
    ER_RW +0      ; Second execution region is called ER_RW, the region starts at the end of the
                  ; previous region. The address is 0x040000 + size of ER_RO region.
    {
        * (+RW)   ; All RW sections go into this region, they are placed consecutively.
    }
    ER_ZI +0      ; Last execution region is called ER_ZI, the region starts at the end of the
                  ; previous region at 0x040000 + the size of the ER_RO regions + the size of
                  ; the ER_RW regions.
    {
        * (+ZI)   ; All ZI sections are placed consecutively here.
    }
}

```

---

The description shown in Example 5-10 on page 5-35 creates an image with one load region called LR\_1, whose load address is 0x040000.

The image has three execution regions, named ER\_RO, ER\_RW, and ER\_ZI, that contain the RO, RW, and ZI output sections respectively. RO, RW are root regions. ZI is created dynamically at runtime. The execution address of ER\_RO is 0x040000. All three execution regions are placed contiguously in the memory map by using the *+offset* form of the base designator for the execution region description. This enables an execution region to be placed immediately following the end of the preceding execution region.

The `--reloc` option is used to make relocatable images. Used on its own, `--reloc` makes an image similar to simple type 1, but the single load region has the RELOC attribute.

### ropi example variant

In this variant, the execution regions are placed contiguously in the memory map. However, `--ropi` marks the load and execution regions containing the RO output section as position-independent.

Example 5-11 shows the scatter-loading description equivalent to using `--ro-base 0x010000 --ropi`.

#### Example 5-11 Position-independent code

---

```

LR_1 0x010000 PI      ; The first load region is at 0x010000.
{
    ER_RO +0          ; The PI attribute is inherited from parent.
                        ; The default execution address is 0x010000, but the code can be moved.
    {
        * (+RO)       ; All the RO sections go here.
    }
    ER_RW +0 ABSOLUTE ; PI attribute is overridden by ABSOLUTE.
    {
        * (+RW)       ; The RW sections are placed next. They cannot be moved.
    }
    ER_ZI +0          ; ER_ZI region placed after ER_RW region.
    {
        * (+ZI)       ; All the ZI sections are placed consecutively here.
    }
}

```

---

Shown in Example 5-11 on page 5-36, ER\_R0, the RO execution region, inherits the PI attribute from the load region LR\_1. The next execution region, ER\_RW, is marked as ABSOLUTE and uses the *+offset* form of base designator. This prevents ER\_RW from inheriting the PI attribute from ER\_R0. Also, because the ER\_ZI region has an offset of +0, it inherits the ABSOLUTE attribute from the ER\_RW region.

### 5.4.2 Type 2

An image of this type consists of a single load region in the load view and three execution regions in the execution view. It is similar to images of type 1 except that the RW execution region is not contiguous with the RO execution region.

--ro-base *address1* specifies the load and execution address of the region containing the RO output section. --rw-base *address2* specifies the execution address for the RW execution region.

Example 5-12 shows the scatter-loading description equivalent to using --ro-base 0x010000 --rw-base 0x040000.

#### Example 5-12 Single load region and multiple execution regions

---

```

LR_1 0x010000      ; Defines the load region name as LR_1
{
  ER_R0 +0         ; The first execution region is called ER_R0 and starts at end of previous region.
                  ; Since there was no previous region, the address is 0x010000.
  {
    * (+R0)        ; All R0 sections are placed consecutively into this region.
  }
  ER_RW 0x040000   ; Second execution region is called ER_RW and starts at 0x040000.
  {
    * (+RW)        ; All RW sections are placed consecutively into this region.
  }
  ER_ZI +0         ; The last execution region is called ER_ZI.
                  ; The address is 0x040000 + size of ER_RW region.
  {
    * (+ZI)        ; All ZI sections are placed consecutively here.
  }
}

```

---

This description creates an image with one load region, named LR\_1, with a load address of 0x010000.

The image has three execution regions, named ER\_R0, ER\_RW, and ER\_ZI, that contain the RO, RW, and ZI output sections respectively. The RO region is a root region. The execution address of ER\_R0 is 0x010000.

The ER\_RW execution region is not contiguous with ER\_RO. Its execution address is 0x040000.

The ER\_ZI execution region is placed immediately following the end of the preceding execution region, ER\_RW.

### rwpi example variant

This is similar to images of type 2 with --rw-base with the RW execution region separate from the RO execution region. However, --rwpi marks the execution regions containing the RW output section as position-independent.

Example 5-13 shows the scatter-loading description equivalent to using --ro-base 0x010000 --rw-base 0x018000 --rwpi.

#### Example 5-13 Position-independent data

---

```

LR_1 0x010000      ; The first load region is at 0x010000.
{
    ER_RO +0        ; Default ABSOLUTE attribute is inherited from parent. The execution address
                    ; is 0x010000. The code and ro data cannot be moved.
    {
        * (+RO)     ; All the RO sections go here.
    }
    ER_RW 0x018000 PI ; PI attribute overrides ABSOLUTE
    {
        * (+RW)     ; The RW sections are placed at 0x018000 and they can be moved.
    }
    ER_ZI +0        ; ER_ZI region placed after ER_RW region.
    {
        * (+ZI)     ; All the ZI sections are placed consecutively here.
    }
}

```

---

ER\_RO, the RO execution region, inherits the ABSOLUTE attribute from the load region LR\_1. The next execution region, ER\_RW, is marked as PI. Also, because the ER\_ZI region has an offset of +0, it inherits the PI attribute from the ER\_RW region.

Similar scatter-loading descriptions can also be written to correspond to the usage of other combinations of --ropi and --rwpi with type 2 and type 3 images.

### 5.4.3 Type 3

Type 3 images consist of two load regions in load view and three execution regions in execution view. They are similar to images of type 2 except that the single load region in type 2 is now split into two load regions.

Relocate and split load regions using the following linker options:

- `--reloc`      The combination `--reloc --split` makes an image similar to simple type 3, but the two load regions now have the RELOC attribute.
- `--ro-base address1`  
                      Specifies the load and execution address of the region containing the RO output section.
- `--rw-base address2`  
                      Specifies the load and execution address for the region containing the RW output section.
- `--split`        Splits the default single load region (that contains the RO and RW output sections) into two load regions. One load region contains the RO output section and one contains the RW output section.

Example 5-14 shows the scatter-loading description equivalent to using `--ro-base 0x010000 --rw-base 0x040000 --split`.

#### Example 5-14 Multiple load regions

---

```

LR_1 0x010000      ; The first load region is at 0x010000.
{
    ER_RO +0        ; The address is 0x010000.
    {
        * (+R0)
    }
}
LR_2 0x040000      ; The second load region is at 0x040000.
{
    ER_RW +0        ; The address is 0x040000.
    {
        * (+RW)      ; All RW sections are placed consecutively into this region.
    }
    ER_ZI +0        ; The address is 0x040000 + size of ER_RW region.
    {

```

```

        * (+ZI)    ; All ZI sections are placed consecutively into this region.
    }
}

```

---

This description creates an image with two load regions, named LR\_1 and LR\_2, that have load addresses 0x010000 and 0x040000.

The image has three execution regions, named ER\_R0, ER\_RW and ER\_ZI, that contain the RO, RW, and ZI output sections respectively. The execution address of ER\_R0 is 0x010000.

The ER\_RW execution region is not contiguous with ER\_R0. Its execution address is 0x040000.

The ER\_ZI execution region is placed immediately following the end of the preceding execution region, ER\_RW.

### Relocatable load regions example variant

This type 3 image also consists of two load regions in load view and three execution regions in execution view. However, --reloc is used to specify that the two load regions now have the RELOC attribute.

Example 5-15 shows the scatter-loading description equivalent to using --ro-base 0x010000 --rw-base 0x040000 --reloc --split.

#### Example 5-15 Relocatable load regions

---

```

LR_1 0x010000 RELOC
{
    ER_R0 + 0
    {
        * (+R0)
    }
}

LR2 0x040000 RELOC
{
    ER_RW + 0
    {
        * (+RW)
    }

    ER_ZI +0
    {

```



```
        * (+ZI)
    }
}
```

---



# Chapter 6

## Creating and Using Libraries

This chapter describes the use of libraries with the ARM® linker, `armlink`, and the librarian, `armar`. It contains the following sections:

- *About libraries* on page 6-2
- *Library searching, selection, and scanning* on page 6-3
- *The ARM librarian* on page 6-6.

## 6.1 About libraries

An object file can refer to external symbols that are, for example, functions or variables. The linker attempts to resolve these references by matching them to definitions found in other object files and libraries. The linker recognizes a collection of ELF files stored in an ar format file as a library.

If you use `--sysv` to generate an SVr4 formatted ELF executable file, the linker treats a shared object as a library. Similarly, a shared object or DLL is treated as a library when you are generating a BPABI-compatible executable file. However, a shared object or DLL differ from an archive in that:

- symbols are imported from a shared object or DLL
- code or data for symbols is extracted from an archive into the file being linked.

The rest of this chapter describes archives.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

## 6.2 Library searching, selection, and scanning

The differences between the way the linker adds object files to the image and the way it adds libraries to the image are:

- Each object file in the input list is added to the output image unconditionally, whether or not anything refers to it. At least one object must be specified.
- A member from a library is included in the output only if an object file or an already-included library member makes a non-weak reference to it, or if the linker is explicitly instructed to add it.

---

### Note

---

If a library member is explicitly requested in the input file list, it is loaded even if it does not resolve any current references. In this case, an explicitly requested member is treated as if it is an ordinary object.

---

Unused sections are subsequently eliminated unless `--no_remove` is used.

Unresolved references to weak symbols do not cause library members to be loaded.

---

### Note

---

If the `--no_scanlib` option is specified, the linker does not search for the default ARM libraries and uses only those libraries that are specified in the input file list to resolve references.

---

Therefore, the linker creates a list of libraries as follows:

1. Any libraries explicitly specified in the input file list are added to the list.
2. The user-specified search path is examined to identify ARM standard libraries to satisfy requests embedded in the input objects.

The best-suited library variants are chosen from the searched directories and their subdirectories. Libraries supplied by ARM have multiple variants that are named according to the attributes of their members.

This section describes:

- *Searching for ARM libraries* on page 6-4
- *Searching for user libraries* on page 6-5
- *Scanning the libraries* on page 6-5.

### 6.2.1 Searching for ARM libraries

You can specify the search paths used to find the ARM standard libraries by:

- Using the environment variable RVCT22LIB. This is the default.
- Adding the `--libpath` argument to the `armlink` command line with a comma-separated list of parent directories.

This list must end with the parent directory of the ARM library directories `armlib` and `cpplib`. The RVCT22LIB environment variable holds this path.

---

**Note**

The linker command-line option `--libpath` overrides the paths specified by the RVCT22LIB variable.

---

The linker combines each parent directory, given by either `--libpath` or the RVCT22LIB variable, with each subdirectory request from the input objects and identifies the place to search for the ARM library. The names of ARM subdirectories within the parent directories are placed in each compiled object by using a symbol of the form `Lib$$Request$$sub_dir_name`.

The sequential nature of the search ensures that the linker chooses the library that appears earlier in the list if two or more libraries define the same symbol.

#### Selecting ARM library variants

There are different variants of the ARM libraries based on the attributes of their member objects. The variant of the ARM library is coded into the library name. The linker must select the best-suited variant from each of the directories identified during the library search.

The linker accumulates the attributes of each input object and then selects the library variant best suited to those attributes. If more than one of the selected libraries are equally suited, the linker retains the first library selected and rejects all others.

The final list contains all the libraries that the linker scans in order to resolve references.

For more information on library variants, see the chapter describing the C and C++ libraries in *RealView Developer Kit v2.2 Compiler and Libraries Guide*.

## 6.2.2 Searching for user libraries

You can specify user libraries by:

- Including them explicitly in the input file list.
- Adding the `--userlibpath` argument to the `armlink` command line with a comma-separated list of directories, and the names of the libraries as input files.

If you do not specify a full path name to a library on the command line, the linker tries to locate the library in the directories specified by the `--userlibpath` argument. For example, if the directory `\mylib` contains `my_lib.a` and `other_lib.a`, add `\mylib\my_lib.a` to the input file list with the command:

```
armlink --userlibpath \mylib my_lib.a *.o
```

### ————— Note —————

The search paths used for the ARM standard libraries specified by the `RVCT22LIB` environment variable or the linker command-line option `--libpath` are not searched for user libraries (see *Searching for ARM libraries* on page 6-4).

## 6.2.3 Scanning the libraries

When the linker has constructed the list of libraries, it repeatedly scans each library in the list to resolve references.

When all the directories have been searched, and the most compatible library variants have been selected and added to the list of libraries, each of the libraries is scanned to load the required members:

1. For each currently unsatisfied non-weak reference, the linker searches sequentially through the list of libraries for a matching definition. The first definition found is marked for step 2.  
The sequential nature of the search ensures that the linker chooses the library that appears earlier in the list if two or more libraries define the same symbol. This enables you to override function definitions from other libraries, for example, the ARM C libraries, by adding your libraries to the input file list.
2. Library members marked in step 1 are loaded. As each member is loaded it might satisfy some unresolved references, possibly including weak ones. Loading a library might also create new unresolved weak and non-weak references.
3. The process in steps 1 and 2 continues until all non-weak references are either resolved or cannot be resolved by any library.

If any non-weak reference remains unsatisfied at the end of the scanning operation, the linker generates an error message.

## 6.3 The ARM librarian

The ARM librarian, `armar`, enables sets of ELF object files or libraries to be collected together and maintained in libraries. Such a library can then be passed to the linker in place of several object files. However, linking with an object library file does not necessarily produce the same results as linking with all the object files collected into the object library file. This is because the linker processes the input list and libraries differently:

- Each object file in the input list appears in the output unconditionally, although unused areas are eliminated if the `arm1ink --remove` option is specified.
- A member of a library file is only included in the output if it is referred to by an object file or a previously processed library file.

For more information on how the linker processes its input files, see Chapter 2 *The Linker Command Syntax*.

This section describes:

- *Librarian command-line options*
- *Examples of armar usage* on page 6-9.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

### 6.3.1 Librarian command-line options

The syntax of the `armar` command when used to add or modify files in the library is:

```
armar [--help] [--create] [--diag_style arm|ide] [-c] [-d] [-m] [-q] [-r] [-u]
[--vsn] [-v] [--via option_file] [{[-a]|[-b]|[-i]} {pos_name}] library
[file_list]
```

The syntax of the `armar` command when used to extract files or library information is:

```
armar [--help] [--diag_style arm|ide] [-C] [--entries] [-p] [-t] [-s] [--sizes]
[-T] [--vsn] [-v] [--via option_file] [-x] [--zs] [--zt] library [file_list]
```

where:

- a            This option places new files in *library* after the file *pos\_name*.
- b            This option places new files in *library* before the file *pos\_name*.



- C            This option instructs the librarian not to replace existing files with like-named files when performing extractions. This option is useful when -T is also used to prevent truncated filenames from replacing files with the same prefix.
- c            This option suppresses the diagnostic message normally written to standard error when a library is created.
- create     This option creates a new library even if *library* already exists.
- d            This option deletes one or more files from *library*.
- diag\_style arm|ide  
               This option specifies that all warning and error messages are output in a format that is more compatible with IDEs, such as Microsoft Visual Studio,  
               The default is arm.
- entries    This option lists all entry points defined in *library*. The format for the listing is:  
               ENTRY at offset *num* in section *name* of *member*
- file\_list*    This is a list of files to process. Each file is fully specified by its path and name. The path can be absolute, relative to drive and root, or relative to the current directory.  
               Only the filename at the end of the path is used when comparing against the names of files in the library. If two or more path operands end with the same filename, the results are unspecified. You can use the wildcards \* and ? to specify files.  
               If one of the files is a library, armar copies all members from the input library to the destination library. The order of entries on the command line is preserved. Therefore, supplying a library file is logically equivalent to supplying all of its members in the order that they are stored in the library.
- help       This option gives online details of the armar command.
- i            This option places new files in *library* before the member *pos\_name* (equivalent to -b).
- library*      This is the path name of the library file.
- m            This option moves files. If -a, -b, or -i with *pos\_name* is specified, files are moved to the new position. Otherwise, move files to the end of library.

- n** This option suppresses the archive symbol table. This is used when the library is not an object library.
- p** This option prints the contents of files in *library* to stdout.
- pos\_name** This is the name of an existing library member to be used for relative positioning. This name must be supplied with options **-a**, **-b**, and **-i**.
- q** This option is an alias for **-r**.
- r** This option replaces, or adds, files in *library*. If *library* does not exist, a new library file is created and a diagnostic message is written to standard error.
- If *file\_list* is not specified and the library exists, the results are undefined. Files that replace existing files do not change the order of the library.
- If the **-u** option is used, then only those files with dates of modification later than the library files are replaced.
- If the **-a**, **-b**, or **-i** option is used, then *pos\_name* must be present and specifies that new files are to be placed after (**-a**) or before (**-b** or **-i**) *pos\_name*. Otherwise the new files are placed at the end.
- t** This option prints a table of contents of *library*. The files specified by *file\_list* are included in the written list. If *file\_list* is not specified, all files in the library are included in the order of the archive.
- s** This option regenerates the archive symbol table.
- sizes** This option lists the Code, RO Data, RW Data, ZI Data, and Debug sizes of each member in *library*, for example:
- | Code | RO Data | RW data | ZI Data | Debug | Object Name |
|------|---------|---------|---------|-------|-------------|
| 464  | 0       | 0       | 0       | 8612  | app_1.o     |
| 3356 | 0       | 0       | 10244   | 11848 | app_2.o     |
| 3820 | 0       | 0       | 10244   | 20460 | TOTAL       |
- T** This option enables filename truncation of extracted files whose library names are longer than the file system can support. By default, extracting a file with a name that is too long is an error. A diagnostic message is written and the file is not extracted.
- u** This option updates older files. When used with the **-r** option, files within *library* are replaced only if the corresponding file has a modification time that is at least as new as the modification time of the file within library.

`--via option_file`

This option instructs the librarian to take options from *option\_file*. See *RealView Developer Kit v2.2 Compiler and Libraries Guide* for more information on writing via files.

`-v`

This option gives verbose output.

The output depends on what other options are used:

`-d, -r or -x`

Write a detailed file-by-file description of the library creation, the constituent files, and maintenance activity.

`-p`

Writes the name of the file to the standard output before writing the file itself to the stdout.

`-t`

Includes a long listing of information about the files within the library.

`-x`

Prints the filename preceding each extraction.

`--vsn`

This option prints the version number on standard error.

`-x`

This option extracts the files in *file\_list* from *library*. The contents of *library* are not changed. If no file operands are given, all files in *library* are extracted. If the filename of a file extracted from the library is longer than that supported in the destination directory, the results are undefined.

`--zs`

This option shows the symbol table.

`--zt`

Lists member sizes and entry points in *library*. See `--sizes` and `--entries` for output format.

#### ————— **Note** —————

The options `-a`, `-b`, `-C`, `-i`, `-m`, `-T`, `-u`, and `-v` are not required for normal operation.

Options relating to library order (for example, `-a`, `-b`, `-i`, and `-m`) are not relevant, since the ARM toolchain cannot use a library that does not have a symbol table. (If there is a symbol table, the order is irrelevant.)

Options related to updating a library (`-C` and `-u`) are unlikely to be used because, in practice, the libraries are rebuilt rather than updated.

### 6.3.2 Examples of armar usage

Syntax examples are shown in Example 6-1 to Example 6-8 on page 6-11.

---

**Example 6-1 Create a new library and add all object files**

---

```
armar --create mylib *.o
```

---

---

**Example 6-2 List the table of contents (verbose)**

---

```
armar -tv mylib
```

---

---

**Example 6-3 List the symbol table**

---

```
armar --zs mylib
```

---

---

**Example 6-4 Add (or replace) files**

---

```
armar -r mylib obj1.o obj2.o obj3.o ...  
armar -ru mylib k*.o
```

---

---

**Example 6-5 Add (or replace) files in specified position**

---

```
armar -r -a obj2.o mylib obj3.o obj4.o ...
```

---

---

**Example 6-6 Extract a group of files**

---

```
armar -x mylib k*.o
```

---

---

**Example 6-7 Delete a group of files**

---

```
armar -d mylib sys_*
```

---

**Example 6-8 Merge libraries and add (or replace) files**

---

```
armer -r mylib my_lib.a other_lib.a obj1.o obj2.o obj3.o
```

---



# Chapter 7

## Using fromelf

This chapter describes the ARM® fromelf software utility provided with *RealView® Developer Kit* (RVDK). It contains the following sections:

- *About fromelf* on page 7-2
- *fromelf command syntax* on page 7-3
- *Examples of fromelf usage* on page 7-9.

## 7.1 About fromelf

fromelf translates *Executable Linkable Format* (ELF) image files produced by the ARM linker into other formats suited to ROM tools and to loading directly into memory. You can also use fromelf to display various information about an ELF object or to generate text files containing the information.

fromelf outputs the following image formats:

- Plain binary format.
- Motorola 32-bit S-record format.
- Intel Hex-32 format.
- Byte Oriented (Verilog Memory Model) Hex format.
- ELF format. You can resave as ELF. For example, you can convert a `--debug` ELF image to a `--no_debug` ELF image.

fromelf can also display information about the input file, for example, symbol listings, to either stdout or a text file.

RVDK uses an ELF proprietary file format called *ARM Toolkit Proprietary ELF* (ATPE). The file format for each version of RVDK is restricted to the proprietary ATPE format for the permitted device. This is referred to as *ATPE\_Custom*.

### 7.1.1 Image structure

fromelf can translate a file from ELF to other formats. It cannot alter the image structure or addresses, other than altering the base address of Motorola S-record or Intel Hex output with the `--base` option. You cannot change a scatter-loaded ELF image into a non-scatter-loaded image in another format. Any structural or addressing information must be provided to the linker at link time.



## 7.2 fromelf command syntax

This section describes the syntax and options of the fromelf command:

```
fromelf [--base n] [--diag_style arm|ide] [--diag_suppress taglist]
[--expandarrays] [--help] [memory_config] [--no_symbolversions]
[[[--text] text_output_format]]| code_output_format] [--vs[n]]
[--output output_file] {input_file}
```

where:

**--base *n*** This option specifies the base address of the output for Motorola S-record, and Intel Hex file formats. This option is available only if --m32, --m32combined, --i32, or --i32combined is specified as the output format.

You can specify the base address as either a:

- decimal value, for example, --base 0
- hexadecimal value, for example, --base 0x8000.

All addresses encoded in the output file start at the base address *n*. If you do not specify a --base option, the base address is taken from the load region address.

### ————— **Note** —————

If multiple load regions are present, the --base value is used for each output file. That is, it overrides all load region addresses.

**--diag\_style *arm|ide***

This option specifies that all warning and error messages are output in a style compatible with IDEs, such as Microsoft Visual Studio.

The default is arm.

**--diag\_suppress *taglist***

This option disables all diagnostic messages that have the specified tag(s).

This option requires a comma-separated list identifying the number of the message to be suppressed, and more than one tag can be specified. For example, to suppress the warning messages that have numbers 1293 and 187, use the following command:

```
fromelf --diag_suppress 1293,187 ...
```

<code>--help</code>	This option shows help and usage information. If this option is specified, other command-line options are ignored. Calling <code>fromelf</code> without any parameters produces the same help information.
<code>memory_config</code>	<p>This option outputs multiple files for multiple memory banks. The format of <code>memory_config</code> is <code>--widthxbanks</code> where:</p> <p><i>width</i> is the width of memory in the target memory system (8-bit, 16-bit, 32-bit, or 64-bit).</p> <p><i>banks</i> specifies the number of memory banks in the target memory system.</p> <p>Valid configurations are:</p> <pre>--8x1 --8x2 --8x4 --16x1 --16x2 --32x1 --32x2 --64x1</pre> <p><code>fromelf</code> uses the last specified configuration if more than one configuration is specified.</p> <p>If the image has one load region, <code>fromelf</code> generates <i>banks</i> files with the following naming conventions:</p> <ul style="list-style-type: none"> <li>• If there is one memory bank (<i>banks</i> = 1) the output file is named by the <code>-o output_file</code> argument.</li> <li>• If there are multiple memory banks (<i>banks</i> &gt; 1), <code>fromelf</code> generates <i>banks</i> number of files starting with <code>output_file0</code> and finishing with <code>output_file banks-1</code>. For example:  <code>fromelf --vhx --8x2 test.axf -o test</code>  generates two files named <code>test0</code> and <code>test1</code>.</li> </ul> <p>If the image has multiple load regions, <code>fromelf</code> creates a directory named <code>output_file</code> and generates bank files for each load region named <code>load_region0</code> to <code>load_region banks-1</code>.</p> <p>The memory width specified by <i>width</i> controls the size of the chunk of information read from the image and written to a file. The first chunk read is allocated to the first file (<code>output_file0</code>), the next chunk is allocated to the next file. After a chunk is allocated to the last file, allocation begins again with the first file (that is, the allocation is modulo based on the number of files). For example:</p>

For a *memory\_config* of --8x4

```
byte0 -> file0
byte1 -> file1
byte2 -> file2
byte3 -> file3
byte4 -> file0
...
```

For a *memory\_config* of --16x2

```
halfword0 -> file0
halfword1 -> file1
halfword3 -> file0
...
```

--no\_symbolversions

This option turns off the decoding of symbol version tables. See *Symbol versioning* on page 4-21 and the *Base Platform ABI for the ARM Architecture* [BPABI] for more information.

--expandarrays

Use the --expandarrays option to print data addresses, including arrays which are expanded both inside and outside structures.

This option can only be used in conjunction with --text -a.

--text *text\_output\_format*

Use a text specification, for example, --text -c, to display image information in text format. You can decode an ELF image or ELF object file using this option. This is the default. If no text or code output format is specified, --text is assumed.

If *output\_file* is not specified with the -o option, the information is displayed on stdout.

If a specific text category is not specified, the default is to output header information.

If specified, the text category consists of one or more of the following:

- a      Prints the global and static data addresses (including addresses for structure and union contents). This option can only be used on files containing debug information. Use the --select option to output a subset of the data addresses.  
If you want to view the data addresses of arrays, expanded both inside and outside structures, use the --expandarrays option with this text category.
- d      Prints contents of the data sections.

e	Decodes exception table information for objects. Use with -c when disassembling images.
g	Prints debug information.
r	Prints relocation information.
s	Prints the symbol and versioning tables.
t	Prints the string table(s).
v	Prints detailed information on each segment and section header of the image.
y	Prints dynamic segment contents.
z	Prints the code and data sizes.

The category selectors can be specified as one of:

- individual options, --text -c -d
- a single concatenated string, --text -cd
- category selectors only, -c -d, -cd
- multiple characters following the /= characters, --text=/cd.

If an output format is not specified, the default output format of --text is used and the individual category selectors are recognized. If another output format is specified, the selectors are ignored.

*code\_output\_format* This option selects the binary or ELF output file options.

*code\_output\_format* can be one of:

--bin	Plain binary. You can split output from this option into multiple files with the <i>memory_config</i> option.
--m32	Motorola 32-bit format (32-bit S-records). This option generates one output file for each load region in the image. You can specify the base address of the output with the --base option.
--m32combined	Motorola 32-bit format (32-bit S-records). This option generates one output file for an image containing multiple load regions. You can specify the base address of the output with the --base option.
--i32	Intel Hex-32 format. This option generates one output file for each load region in the image. You can specify the base address of the output with the --base option.

**--i32combined**

Intel Hex-32 format. This option generates one output file for an image containing multiple load regions. You can specify the base address of the output with the **--base** option.

**--vhx**

Byte Oriented (Verilog Memory Model) Hex format. This format is suitable for loading into the memory models of *Hardware Description Language* (HDL) simulators. You can split output from this option into multiple files with the *memory\_config* option.

**--elf**

ELF format (resaves as ELF). This can be used to convert a debug ELF image into a no-debug ELF image.

If you use fromelf to convert an ELF image containing multiple load regions to a binary format using any of the **--bin**, **--m32**, **--i32**, or **--vhx** options, fromelf creates an output directory named *output\_file* and generates one binary output file for each load region in the input image. fromelf places the output files in the *output\_file* directory.

If you convert an ELF image containing multiple load regions using either the **--m32combined** or **--i32combined** option, fromelf creates an output directory named *output\_file*, generates one binary output file for all load regions in the input image, and then places the output file in the *output\_file* directory.

ELF images contain multiple load regions if, for example, they are built with a scatter-loading description file that defines more than one load region.

**--vsn**

This option displays fromelf version information.

**--output *output\_file***

This option specifies the name of the output file, or the name of the output directory if multiple output files are created (see the description of *text\_output\_format* and *code\_output\_format* for more information). Specifying the output file is optional with the **--text** output option and mandatory with all other outputs.

***input\_file***

This option specifies the ELF file to be translated.

fromelf accepts only ARM executable ELF files and ARM object ELF files (.o). If *input\_file* is a scatter-loaded image that contains more than one load region and the output format is one of --bin, --m32, --i32, or --vhx, fromelf creates a separate file for each load region.

If *input\_file* is a scatter-loaded image that contains more than one load region and the output format is either --m32combined or --i32combined, fromelf creates a single file containing all load regions.

## 7.3 Examples of fromelf usage

This section contains examples of using fromelf to change image format or extract information from an ELF file.

This section describes:

- *Producing a plain binary file*
- *Listing addresses of static data*

### 7.3.1 Producing a plain binary file

To convert an ELF file to a plain binary (.bin) file, enter:

```
fromelf --bin -o outfile.bin infile.axf
```

### 7.3.2 Listing addresses of static data

To list to stdout all the global and static data variables and all the structure field addresses, enter:

```
fromelf --text -a --select * infile.axf
```

#### Selecting only structures

To produce a text file containing all of the structure addresses in inputfile.axf but none of the global or static data variable information, enter:

```
fromelf --text -a --select *.* -o strucaddress.txt infile.axf
```

#### Selecting only nested structures

To produce a text file containing addresses of the nested structures only, enter:

```
fromelf --text -a --select *.*.* -o strucaddress.txt infile.axf
```

#### Selecting only variables

To produce a text file containing all of the global or static data variable information in inputfile.axf but none of the structure addresses, enter:

```
fromelf --text -a --select *, ~*.* -o strucaddress.txt infile.axf
```





# Glossary

The items in this glossary are listed in alphabetical order, with any symbols and numerics appearing at the end.

**AAPCS** *See* Procedure Call Standard for the ARM Architecture.

**ABI for the ARM Architecture (base standard) (BSABI)**

The ABI for the ARM Architecture is a collection of specifications, some open and some specific to ARM architecture, that regulate the inter-operation of binary code in a range of ARM-based execution environments. The base standard specifies those aspects of code generation that must be standardized to support inter-operation and is aimed at authors and vendors of C and C++ compilers, linkers, and runtime libraries.

**ADS** *See* ARM Developer Suite.

**API** *See* Application Programming Interface.

**Application Programming Interface (API)**

The syntax of the functions and procedures within a module or library.

**Architecture** The term used to identify a group of processors that have similar characteristics.

**ARM Developer Suite (ADS)**

A suite of software development applications, together with supporting documentation and examples, that enable you to write and debug applications for the ARM family of *RISC* processors. ADS is superseded by RealView Developer Suite (RVDS).

*See also* RealView Developer Suite.

**ARM state**

A processor that is executing ARM instructions is operating in ARM state. The processor switches to Thumb state (and to recognizing Thumb instructions) when directed to do so by a state-changing instruction such as BX, BLX.

*See also* Thumb state.

**ARM Toolkit Proprietary ELF (ATPE)**

The binary file format used by RealView Developer Kit. ATPE object format is produced by the compiler and assembler tools. The ARM linker accepts ATPE object files and can output an ATPE executable file. RealView Debugger can load only ATPE format images, or binary ROM images produced by the fromELF utility.

*See also* RealView Developer Suite.

**armar**

The ARM librarian.

**armlink**

The ARM linker.

**ATPE**

*See* ARM Toolkit Proprietary ELF.

**Big-endian**

Memory organization where the least significant byte of a word is at a higher address than the most significant byte.

**BNF**

Backus Naur Format. Mathematical notation for defining logical structures.

**BSABI**

*See* ABI for the ARM Architecture (base standard).

**Byte**

A unit of memory storage consisting of eight bits.

**Debug With Arbitrary Record Format (DWARF)**

ARM code generation tools generate debug information in DWARF2 format by default. From RVCT v2.2, you can optionally generate DWARF3 format (Draft Standard 9).

**Deprecated**

A deprecated option or feature is one that you are strongly discouraged from using. Deprecated options and features will not be supported in future versions of the product.

**Doubleword**

A 64-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.

**DWARF**

*See* Debug With Arbitrary Record Format.

**ELF**

*See* Executable and Linking Format.

<b>Environment</b>	The actual hardware and operating system that an application will run on.
<b>Executable and Linking Format (ELF)</b>	The industry standard binary file format used by RealView Developer Kit. ELF object format is produced by the ARM object producing tools such as armcc and armasm. The ARM linker accepts ELF object files and can output either an ELF executable file, or a partially linked ELF object.
<b>Execution view</b>	The address of regions and sections after the image has been loaded into memory and started execution.
<b>Flash memory</b>	Nonvolatile memory that is often used to hold application code.
<b>fromelf</b>	The ARM image conversion utility. This accepts ELF format input files and converts them to a variety of output formats. fromelf can also generate text information about the input image, such as code and data size.
<b>Heap</b>	The portion of computer memory that can be used for creating new variables.
<b>Host</b>	A computer that provides data and other services to another computer.
<b>IDE</b>	Integrated Development Environment (for example, the ARM RealView Debugger IDE).
<b>Image</b>	An executable file that has been loaded onto a processor for execution.
<b>Input section</b>	Contains code or initialized data or describes a fragment of memory that must be set to zero before the application starts.
<b>Interrupt</b>	A change in the normal processing sequence of an application caused by, for example, an external signal.
<b>Interworking</b>	Producing an application that uses both ARM and Thumb code.
<b>Library</b>	A collection of assembler or compiler output objects grouped together into a single repository.
<b>Linker</b>	Software that produces a single image from one or more source assembler or compiler output objects.
<b>Little-endian</b>	Memory organization where the least significant byte of a word is at a lower address than the most significant byte.
<b>Load view</b>	The address of regions and sections when the image has been loaded into memory but has not yet started execution.

**Output section** A contiguous sequence of input sections that have the same RO, RW, or ZI attributes. The sections are grouped together in larger fragments called regions. The regions will be grouped together into the final executable image.

*See also* Region.

**PIC** Position Independent Code.

*See also* ROPI.

**PID** Position Independent Data *or* the ARM Platform-Independent Development card.

*See also* RWPI.

**Procedure Call Standard for the ARM Architecture (AAPCS)**

*Procedure Call Standard for the ARM Architecture* defines how registers and the stack will be used for subroutine calls.

**Profiling** Accumulation of statistics during execution of a program being debugged, to measure performance or to determine critical areas of code.

*Call-graph profiling* provides great detail but slows execution significantly. *Flat profiling* provides simpler statistics with less impact on execution speed.

For both types of profiling you can specify the time interval between statistics-collecting operations.

**Program image** See Image.

**Read-Only Position Independent (ROPI)**

Code or read-only data that can be placed at any address.

**Read Write Position Independent (RWPI)**

Read/write data addresses that can be changed at runtime.

**RealView Compilation Tools (RVCT)**

RealView Compilation Tools is a suite of tools, together with supporting documentation and examples, that enables you to write and build applications for the ARM family of *RISC* processors.

**RealView Developer Kit (RVDK)**

RealView Developer Kit is a suite of software development applications, together with supporting documentation and examples, that enables you to write and debug applications for the ARM family of *RISC* processors.

**RealView Developer Suite (RVDS)**

The latest suite of software development applications, together with supporting documentation and examples, that enable you to write and debug applications for the ARM family of *RISC* processors.

<b>RealView ICE (RVI)</b>	A JTAG-based debug solution to debug software running on ARM processors.
<b>RealView ICE Micro Edition (RVI-ME)</b>	A JTAG-based debug tool for embedded systems.
<b>Redirection</b>	The process of sending default output to a different destination or receiving default input from a different source. This is commonly used to output text, that would otherwise be displayed on the computer screen, to a file.
<b>Reduced Instruction Set Computing (RISC)</b>	Device where the number of instructions a microprocessor runs for a specific application are reduced from a general purpose Complex Instruction Set Computing (CISC) device to create a more efficient operating system.
<b>Region</b>	In an image, a region is a contiguous sequence of one to three output sections (RO, RW, and ZI). A region typically maps onto a physical memory device, such as ROM, RAM, or peripheral.  <i>See also</i> Root region.
<b>Remapping</b>	Changing the address of physical memory or devices after the application has started executing. This is typically done to enable RAM to replace ROM once the initialization has been done.
<b>Retargeting</b>	The process of moving code designed for one execution environment to a new execution environment.
<b>RISC</b>	<i>See</i> Reduced Instruction Set Computing.
<b>Root region</b>	In an image, regions having the same load and execution address. A non-root region is a region that must be copied from its load address to its execution address.
<b>ROPI</b>	<i>See</i> Read-Only Position Independent.
<b>RTOS</b>	Real Time Operating System.
<b>RVCT</b>	<i>See</i> RealView Compilation Tools.
<b>RVDK</b>	<i>See</i> RealView Developer Kit.
<b>RVDS</b>	<i>See</i> RealView Developer Suite.
<b>RVI</b>	<i>See</i> RealView ICE.
<b>RVI-ME</b>	<i>See</i> RealView ICE Micro Edition.
<b>RWPI</b>	<i>See</i> Read Write Position Independent.
<b>Scatter-loading</b>	Assigning the address and grouping of code and data sections individually rather than using single large blocks.

<b>Section</b>	<p>A block of software code or data for an Image.</p> <p><i>See also</i> Input section <i>and</i> Output section.</p>
<b>Semihosting</b>	<p>A mechanism whereby the target communicates I/O requests made in the application code to the host system, rather attempting to support the I/O itself.</p>
<b>Software Development Toolkit (SDT)</b>	<p>Software Development Toolkit (SDT) is an ARM product, now superseded by ADS and RVDS.</p>
<b>Software Interrupt (SWI)</b>	<p>An instruction that causes the processor to call a programmer-specified subroutine. Used by ARM to handle semihosting.</p>
<b>Stack</b>	<p>The portion of memory that is used to record the return address of code that calls a subroutine. The stack can also be used for parameters and temporary variables.</p>
<b>SVr4</b>	<p><i>See</i> System V release 4.</p>
<b>SWI</b>	<p><i>See</i> Software Interrupt.</p>
<b>System V release 4 (SVr4)</b>	<p>The System V Application Binary Interface (ABI) is a family of specifications that define a standard binary interface for compiled applications that implement UNIX System V release 4.</p>
<b>Target</b>	<p>The actual target processor, real or simulated, on which the target application is running.</p> <p>The fundamental object in any debugging session. The basis of the debugging system. The environment in which the target software will run. It is essentially a collection of real or simulated processors.</p>
<b>Thumb state</b>	<p>A processor that is executing Thumb instructions is operating in Thumb state. The processor switches to ARM state (and to recognizing ARM instructions) when directed to do so by a state-changing instruction such as BX, BLX.</p> <p><i>See also</i> ARM state.</p>
<b>Vector Floating Point (VFP)</b>	<p>A standard for floating-point coprocessors where several data values can be processed by a single instruction.</p>
<b>Veneer</b>	<p>A small block of code used with subroutine calls when there is a requirement to change processor state or branch to an address that cannot be reached in the current processor state.</p>
<b>VFP</b>	<p><i>See</i> Vector Floating Point.</p>

<b>Volatile</b>	Memory addresses where the contents can change independently of the executing application are described as volatile. These are typically memory-mapped peripherals.
<b>Word</b>	A 32-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.
<b>Zero Initialized (ZI)</b>	R/W memory used to hold variables that do not have an initial value. The memory is normally set to zero on reset.
<b>ZI</b>	<i>See</i> Zero Initialized.





# Index

The items in this index are listed in alphabetical order, with symbols and numerics appearing at the end. The references given are to page numbers.

## A

- AAELF 2-29
- ADS 2-31
- Archives
  - about 6-2
- AREA
  - directive 3-10, 5-27
- ARM
  - AAELF 2-29
  - ABI 2-31
  - BPABI 1-3, 4-21
- ARM Linux
  - SVr4 format 6-2
- ARM Toolkit Proprietary ELF 1-2,  
2-2, 2-9, 6-2, 6-6, 7-2
- armar
  - as part of RVDK 1-4
  - command syntax 6-6
  - invoking from the command line  
1-4, 6-6
  - options 6-6
  - options using wildcards 6-7

- see also* Librarian
- arm1ink
  - as part of RVDK 1-3
  - command syntax 2-9
  - invoking from the command line  
2-9
  - options 2-9–2-31
  - options using wildcards 2-25
  - output ELF format 1-3
  - see also* Linker
- ARMv4T
  - veneer generation 3-19
- ARMv5
  - veneer generation 3-19
- ARM/Thumb synonyms 4-2
- ATPE 1-2, 2-2, 2-9, 6-2, 6-6, 7-2

## B

- Backus Naur Format (BNF) 5-9
- BPABI
  - executable file 2-23

- in linker 1-3, 4-21
- Branching
  - controlling optimizations 3-23
  - inlining small functions 3-22, 3-23
  - interworking optimizations 3-23
  - optimizations 3-22
  - tail call sections 3-22, 3-23

## C

- Callgraph
  - creating 2-26
  - cyclic functions 2-27
  - function list 2-26
  - indirect functions 2-27
  - no stack information 2-27
  - stack usage 2-26
- Comdat groups 1-4, 2-17
- Command syntax
  - fromelf 7-3
  - librarian 6-6
  - linker 2-9

- Comments
  - in steering files 4-10
  - in symdefs file 4-9
- Complex images
  - scatter-loading 3-5, 5-5
  - see also* Scatter-loading
- Compression
  - algorithms 3-17
  - decompressor 3-17
  - disabled for
    - Load\$\$region\_name\$\$Base 3-18
  - in scatter-loading 3-18
  - InRoot\$\$Sections in scatter-loading 3-18, 5-30
  - LZ77 3-17
  - NOCOMPRESS in scatter-loading 3-19
  - run-length 3-17
  - RW data 3-16
  - sections in root regions 3-18
  - viewing compressed regions 3-18
- Consolidated section 4-5
- C++, licensing 1-2, 2-10, 4-3
- D**
- Data
  - in scatter-loading 5-27
- Data compression 3-16
- Default version
  - symbol versioning 4-21
- Directives
  - preprocessing in scatter-loading 5-33
  - AREA 3-10, 5-27
  - ENDP 2-26
  - ENTRY 2-22, 3-6, 3-7, 5-25
  - FRAME 2-26
  - PROC 2-26
  - +FIRST 2-21
  - +LAST 2-21
- Dynamic linker
  - finalization code 2-25
  - initialization code 2-24
- E**
- ELF
  - formats 7-2
  - fromelf outputs 7-2
  - fromelf utility 1-4, 7-2
  - images 7-2
  - object format specification 2-2
  - objects 7-2
  - segments 2-3
  - SVr4 format 6-2
- Eliminating
  - common debug sections 3-11
  - common groups 1-4, 3-11
  - common sections 3-11
  - unused sections 1-3, 3-12
  - unused virtual functions 3-12, 3-13
- ENDP
  - directive 2-26
- ENTRY
  - directive 2-22, 3-6, 3-7, 5-25
- Entry point
  - ENTRY directive 3-6
  - in image 3-6
  - in image for C/C++ 3-6
  - initial 3-6
  - specifying initial 3-6
  - specifying to the linker 2-19
- Environment variables
  - RVCT22LIB 2-12, 6-4, 6-5
- Examples
  - main directory 1-2
- Exceptions
  - in linker 3-31
- Exceptions tables
  - generate 3-31
  - in linker 2-19, 2-20, 3-31
- Execution region
  - description 5-14
  - naming 4-4
  - related symbols 4-3
  - reuse of veneers 3-21
  - simple images 3-25
  - veneers generation 3-10
- Execution view
  - in memory map 2-3
- F**
- Feedback
  - example 3-15
- in linker 1-3, 2-27, 3-14
- interworking 2-27, 3-14
- File formats
  - fromelf output 7-5, 7-6
- Files
  - scatter-loading 5-1
  - steering 4-10
  - symdefs 4-7
  - via 2-31, 6-9
- FIRST
  - sections 3-9
  - sorting rules 3-9
- FLEXlm license 1-2, 2-10, 4-3
- FRAME
  - directive 2-26
- fromELF
  - see also* ELF, and fromelf
- fromelf
  - utility 1-4
  - version number 7-7
- fromelf
  - as part of RVDK 1-4
  - command syntax 7-3
  - debug information 2-17
  - diagnostics 7-3
  - examples 7-9
  - input information 7-2
  - invoking from the command line 7-3
  - messages 7-3
  - specifying input format 7-8
  - specifying memory banks 7-4
  - specifying output format 7-6
    - binary 7-6
    - ELF 7-7
    - I32 7-6, 7-7
    - M32 7-6
    - Verilog 7-7
  - specifying text format 7-5
  - symbol versioning 7-5
- fromelf options
  - a 7-5, 7-9
  - base 7-3, 7-6, 7-7
  - bin 7-6, 7-9
  - c 7-6
  - d 7-5
  - diag\_style 7-3
  - diag\_suppress 7-3
  - e 7-6

- elf 7-7
  - expandarrays 7-5
  - g 7-6
  - help 7-4
  - i32 7-6
  - i32combined 7-7
  - m32 7-6
  - m32combined 7-6
  - no\_symbolversions 7-5
  - o 7-9
  - output 7-7
  - r 7-6
  - s 7-6
  - select 7-9
  - t 7-6
  - text 7-5, 7-9
  - v 7-6
  - vbx 7-7
  - vsb 7-7
  - y 7-6
  - z 7-6
- Functions
- in scatter-loading 5-27
- ## H
- HIDE
- in steering file 4-18
- Hiding and renaming symbols 2-18, 4-10
- ## I
- Images
- complex 3-5, 5-4
  - complex using scatter-loading 5-5
  - creating a symdefs file 4-7
  - describing complex 3-5
  - describing simple 3-5
  - ELF 7-2
  - ENTRY directive 3-6
  - entry points 3-6
  - entry points in C/C++ 3-6
  - format of a symdefs file 4-8
  - initial entry point 3-6
  - memory map information 3-4
  - reading a symdefs file 4-8
- regions 3-3
- sections 3-3, 5-8
- simple 2-5, 2-15, 3-25
- Simple type 1 3-25, 3-26, 5-35
- ropi example 5-36
- Simple type 2 3-25, 3-27, 5-37
- reloc example 5-40
  - rwpj example 5-38
- Simple type 3 3-25, 3-29, 5-39
- simple using scatter-loading 5-4, 5-35, 5-37, 5-39
- specifying a memory map 3-5
- specifying initial entry point 3-6
- specifying memory map 3-5
- structure of 3-2
- using a symdefs file 4-7
- using scatter-loading description files 3-5, 5-2
- writing global symbols to a symdefs file 4-7
- IMPORT
- in steering file 4-12, 4-13
- Initial entry point
- about 3-6
  - specifying 3-6
- Input section
- description 5-17
  - using symbol names 5-20
  - using wildcards 5-17, 5-19
  - :gdef: prefix 5-20
- Installation directory
- default location vii
- Interworking
- about 1-3
  - branch optimizations 3-23
  - linker feedback 2-27, 3-14
  - veneers 1-3, 3-19
- ## L
- LAST
- sections 3-9
  - sorting rules 3-9
- Librarian
- about 6-6
  - adding object files to image 6-3
  - command syntax 6-6
  - creating libraries 6-3
- diagnostics 6-7
- examples 6-9
- messages 6-7
- options 6-6
- see also* Librarian options
- scanning libraries 6-5
- searching for ARM libraries 6-4
- searching for ARM variants 6-4
- searching for user libraries 6-5
- using via files 6-9
- variants 6-3
- Librarian options
- a 6-6, 6-7, 6-8, 6-9
  - b 6-6, 6-7, 6-8, 6-9
  - C 6-7, 6-9
  - c 6-7
  - create 6-7, 6-9
  - d 6-7, 6-9, 6-11
  - diag\_style 6-7
  - entries 6-7
  - help 6-7
  - i 6-7, 6-8, 6-9
  - m 6-7, 6-9
  - p 6-8, 6-9
  - q 6-8
  - r 6-8, 6-9, 6-11
  - s 6-8, 6-11
  - sizes 6-8
  - T 6-7, 6-8, 6-9
  - t 6-8, 6-9
  - u 6-8, 6-9, 6-11
  - v 6-9
  - via 6-9
  - vsb 6-9
  - x 6-9, 6-11
  - z 6-11
  - zs 6-9
  - zt 6-9
- Libraries
- about 6-2
  - adding to image 6-3
  - command syntax 6-6
  - creating 6-3
  - scanning 6-5
  - searching 6-4
  - variants 6-3, 6-4
- License
- FLEXlm 1-2, 2-10, 4-3
  - for C++ 1-2, 2-10, 4-3

- Linker
  - about 1-3
  - archives 6-2
  - ARM 1-3
  - behavior 2-4
  - best debug view 2-17
    - turning off 2-17
    - turning on 2-17
  - BPABI format 2-23
  - branch inlining 2-23, 2-24, 2-27, 2-28, 3-22
  - Byte Addressing 2-26
  - callgraph 2-26
  - code and data sizes 2-27, 2-28, 3-32
  - Comdat groups 1-4, 2-17
  - command syntax 2-9
  - commands in a via file 2-31
  - common debug section elimination 3-11
  - common group elimination 1-4, 3-11
  - common section elimination 3-11
  - compatibility with legacy objects 2-31
  - complex memory map 2-5, 2-15
  - constructing a partially-linked object 2-3
  - constructing an image 2-2, 2-5, 2-18
  - debug info in image 2-5
  - debug information
    - turning off 2-17
    - turning on 2-17
  - default addresses 2-15
  - diagnostics 2-7, 2-13, 2-28, 2-29, 2-30, 2-31
  - disabling inline veneers 2-26, 3-20
  - disabling veneer sharing 2-26, 3-19
  - ELF format 6-2
  - errors 2-13
  - Exceptions support 2-19, 2-20, 3-31
    - unwind table 2-20
  - exceptions table 2-28
  - exceptions tables
    - generate 2-19, 2-20, 3-31
    - remove 2-20
  - execution information 3-4
  - feedback example 3-15
  - feedback on unused functions 2-27, 3-14–3-16
  - help on 2-4, 2-11
  - image
    - entry point 2-19
    - keeping sections 2-25
    - load and execution views 3-4
    - overview 3-2
    - structure 3-2
  - image-related information 2-6, 2-26, 3-32
  - information 2-13, 2-25, 2-26, 2-27, 2-28, 2-31, 3-32
  - input 2-2, 2-11
  - input armar libraries 2-2
  - input ELF format 2-2
  - input files 2-4
  - input symdefs file 2-2
  - libraries
    - creating 6-3
    - defaults 2-12
    - including during link step 6-2
    - linker search path 2-12
    - scanning 2-12, 6-5
    - searching 6-4
    - variants 6-3
  - load information 3-4
  - local symbols 2-25
  - mapping symbols 2-29
  - memory map 2-28
  - memory map information 2-3, 2-16, 2-22, 2-24, 3-4
  - messages 2-13, 2-25, 2-30, 2-31
  - optimizations 2-28, 3-11–3-23
  - options 1-3, 2-9–2-31
    - see also* Linker options
  - ordering input sections 3-9
  - output 1-3, 1-4, 2-2
  - output file 2-4, 2-5, 2-14
  - output formats 1-4, 2-4, 2-5, 2-14
  - padding bytes 2-16
  - partial linking 2-14, 4-8
  - preprocessing directives in
    - scatter-loading 5-33
  - regions 3-3
  - relocatable images 2-14, 5-39
  - remove unused sections 2-22
  - remove unused virtual functions 2-24
  - RO section base address 2-15, 3-30
  - root region 2-3, 3-18
  - RW data compression 2-18, 3-16
  - RW section base address 2-15
  - scatter-loading 2-5
    - see also* Scatter-loading
  - searching for libraries 6-4
  - sections
    - input 3-3, 4-5
    - output 3-3, 4-5
  - simple memory map 2-5, 2-15
  - software version 2-11
  - sorting input sections 3-8
  - specifying input files 2-11
  - specifying steering files 4-10
  - standard output stream 2-31
  - symbol versioning 2-22, 2-23, 4-21
  - symbols 4-3, 4-6, 4-20
    - hiding and renaming 2-18, 4-10
    - undefined 2-13
    - used in link step 2-29
  - symbols in scatter-loading 5-2
  - System V shared libraries 2-15
  - tail call reordering 2-23
  - undefined symbols 2-13
  - unused section elimination 1-3, 3-12
  - unused sections 2-22, 2-27, 2-28
  - unused virtual function elimination 3-12, 3-13
  - unused virtual functions 2-23, 3-12, 3-13
  - user libraries
    - linker search path 2-12
    - searching 6-5
  - using shared libraries 2-25
  - using shared objects 2-15, 2-24, 2-25
  - using steering files 2-18
  - veneer generation 2-6, 2-26
  - veneer sharing 2-26, 3-19
  - veneers 2-28
  - veneers generation 3-10, 3-19
  - version number 2-4, 2-11
  - via files 2-8, 2-31
  - \$ symbols 4-20
  - \$\$ symbols 4-3
- Linker options
  - help 2-4, 2-11
  - input files 2-4, 2-11
  - linker behavior 2-4, 2-13

output file format 2-4, 2-14  
 memory map 2-5, 2-15  
 debug info 2-5, 2-17  
 image contents 2-5, 2-18  
 veneer generation 2-6, 2-26  
 image details 2-6, 2-26  
 diagnostics 2-7, 2-30  
 via files 2-8, 2-31  
 ADS compatibility 2-31  
 --bestdebug 2-17, 3-11  
 --callgraph 2-26  
 --cppinit 2-24  
 --datacompressor *id* 2-18, 3-17  
 --datacompressor list 2-18, 3-17  
 --datacompressor off 2-18, 3-17  
 --datacompressor on 2-18  
 --debug 2-17  
 --diag\_style 2-30  
 --diag\_suppress 2-30  
 --diag\_warning 2-30, 3-31  
 --edit 2-18, 4-10  
 --entry 2-19  
 --errors 2-31  
 --exceptions 2-19, 3-31  
 --exceptions\_tables 2-20  
 --feedback 2-27  
 --fini 2-25  
 --first 2-20, 3-9, 3-12  
 --fpic 2-15  
 --help 2-11  
 --info 2-27, 3-32  
 --info inline 3-22  
 --info sizes,totals 2-28, 3-32  
 --info tailreorder 3-24  
 --info unused 3-12  
 --init 2-24  
 --inline 2-24, 3-22  
 --keep 2-25, 3-12  
 --last 2-21, 3-9, 3-12  
 --libpath 2-12, 6-4, 6-5  
 --list 2-31  
 --list\_mapping\_symbols 2-29  
 --locals 2-25  
 --mangled 2-28  
 --map 2-28, 3-18  
 --match crossmangled 2-13  
 --no\_bestdebug 2-17, 3-12  
 --no\_debug 2-17  
 --no\_exceptions 2-20, 3-31

--no\_inlineveneer 2-26, 3-20  
 --no\_locals 2-25  
 --no\_remove 2-22, 3-12  
 --no\_scanlib 2-12  
 --no\_veneershare 2-26, 3-19  
 --output 2-14  
 --pad 2-16  
 --partial 2-14  
 --reloc 2-14, 5-36, 5-39  
 --remove 2-22, 3-12  
 --ro-base 2-15, 3-25, 3-26, 3-27,  
 3-29, 5-35, 5-36, 5-37, 5-38, 5-39  
 --ropi 2-15, 5-36  
 --rosplit 2-16, 3-25  
 --rw-base 2-15, 3-25, 3-27, 3-29,  
 5-37, 5-38  
 --rwpi 2-15, 5-38  
 --scanlib 2-12  
 --scatter 2-16, 5-4  
 --split 2-16, 3-25, 3-29, 3-30, 5-39  
 --startup 2-22  
 --strict 2-13  
 --strict\_relocations 2-13  
 --symbols 2-29  
 --symdefs 2-29, 4-7  
 --symver\_script 2-22, 4-23  
 --symver\_soname 2-23, 4-24  
 --sysv 6-2  
 --tailreorder 2-23, 3-22  
 --unmangled 2-28  
 --unresolved 2-13  
 --userlibpath 2-12, 6-5  
 --verbose 2-31  
 --vfemode 2-23  
 --via 2-31  
 --vsn 2-11  
 --xref 2-29  
 --xrefdbg 2-29  
 --xreffrom 2-29  
 --xrefto 2-30  
 Load region  
 description 5-11  
 Load view  
 in memory map 2-3  
 LZ77  
 Lempel-Ziv 1977 3-17

## M

Main examples directory 1-2  
 Mapping symbols  
 in linker 2-29  
 Memory map  
 describing to linker 5-8  
 execution view 2-3, 3-4  
 linker 2-28  
 linker options 2-5, 2-15, 3-25  
 load regions 2-3  
 load view 2-3, 3-4  
 overlaid 5-15  
 partial linking 2-5, 4-6  
 reimplementing  
 \_\_user\_initial\_stackheap 2-5,  
 2-17  
 specifying 2-16, 2-22, 2-24  
 specifying in image 3-5  
 uninitialized 5-15  
 views in 3-4

## O

Object files  
 in scatter-loading 5-27  
 Optimizations  
 in linker 3-22

## P

Position independence 2-15  
 POSIX -- option 2-4  
 Pragmas  
 arm section in scatter-loading 5-21,  
 5-27, 5-28  
 Preprocessor  
 directives in scatter-loading 5-33  
 PROC  
 directive 2-26

## R

Regions  
 contents 3-3

- in scatter-loading 5-26, 5-29, 5-30, 5-31
- naming in scatter-loading 4-4
- placement 3-5
- related symbols 4-3
- reserving empty 5-30, 5-31
- symbols 4-3, 4-4
- Relocatable images
  - creating 2-14
  - with --split 5-39
- RENAME
  - in steering file 4-14
- REQUIRE
  - in steering file 4-17
- RESOLVE
  - in steering file 4-15
- Root region 2-3, 2-16, 3-18
  - in scatter-loading 5-25
- RTTI
  - see* Runtime Type Information
- Runtime Type Information 2-23, 2-24, 3-12, 3-13, 3-14
- RVCT22LIB environment variable 2-12, 6-4, 6-5
- RVDK
  - components 1-2
- RW data compression 3-16–3-19
  - choosing compression 3-17
  - decompressor 3-17

## S

- Scatter-loading
  - about 5-2
  - attribute selectors in input sections 5-18
  - complex images 5-5
  - description file 1-3, 2-16, 3-5
    - arm section pragma 5-21, 5-27, 5-28
    - BNF symbols 5-9
    - example 5-10
    - execution region 5-14
    - hierarchy 5-8, 5-9
    - load region 5-11
    - load region syntax 5-11
    - multiple execution regions 5-23
    - multiple matches 5-21
  - placing data 5-27
  - placing functions 5-27
  - placing object file 5-27
  - placing regions 5-26
  - placing sections with overlays 5-29
  - reserving empty regions 5-30, 5-31
  - root region 5-25
  - sections 5-8
  - syntax 5-8
  - veneers 5-24
- Sections
  - aligning 3-10, 3-33
  - consolidated 4-5
  - forcing aligning 3-10
  - grouping 3-5
  - in regions 3-3
  - in scatter-loading 5-29
  - input 3-3, 4-5
    - placement 3-9
  - keeping 2-25
  - multiple matches in scatter-loading 5-21
  - ordering execution 3-10
  - ordering rules 2-21, 3-9
  - output 3-3, 4-5
  - overlaying 5-29
  - placement 3-8
    - by attribute 3-9
    - FIRST and LAST 3-9
- placement of 2-20, 2-21, 5-8, 5-19
- related symbols 4-5
- sorting rules 3-8
- unused 2-22, 2-27, 2-28
- used 2-22
- SHOW
  - in steering file 4-19
- Simple images 2-5, 2-15, 3-25
  - scatter-loading 3-25, 5-4, 5-35, 5-37, 5-39
- Simple type 1 image 3-25, 3-26
- Simple type 2 image 3-25, 3-27
- Simple type 3 image 3-25, 3-29
- Steering file
  - commands 4-10, 4-11
    - HIDE 4-18
    - IMPORT 4-12, 4-13
    - RENAME 4-14
    - REQUIRE 4-17
    - RESOLVE 4-15
    - SHOW 4-19
  - comments 4-10
  - contents 4-10
  - format 4-10
  - hiding and renaming symbols 4-10
  - specifying 4-10
  - specifying commands in 2-18
  - using wildcards 4-10, 4-12, 4-13, 4-14, 4-15, 4-17, 4-18, 4-19
- SVr4 1-3
- Symbol definitions file
  - see* Symdefs file
- Symbol versioning
  - about 4-21
  - creating versioned symbols 4-22
  - default version 4-21
  - in linker 2-22, 2-23, 4-21–4-24
  - version 4-21
- Symbols
  - across images 4-7
  - ARM/Thumb synonyms 4-2
  - consolidated sections 4-5
  - definitions file 4-7
  - hiding and renaming 4-10
  - Image\$\$ undefined 5-2
  - in symdefs file 4-9
  - initialization 2-24
  - linker 2-18, 2-29, 4-3, 4-10, 4-20
  - linker-defined 2-29, 4-6

- mangled 2-13
- mapping symbols 2-29
- multiple definitions of 4-2
- overriding definitions 4-20
- region-related 4-3, 4-4
- scatter-loading 5-2
- section-related 4-4, 4-5
- stack 4-4
- startup 2-22
- unmangled 2-13
- versioning in linker 2-22, 2-23, 4-21
- writing to symdefs file 4-7
  - \$ 4-20
  - \$Sub\$\$ 4-20
  - \$Super\$\$ 4-20
  - \$\$ 4-3
  - \$\$ZI\$\$ 4-4
- Symdefs file
  - about 4-7
  - comments 4-9
  - creating 4-7
  - format 4-8
  - identifying string 4-9
  - reading 4-8
  - symbol information 4-9
  - writing global symbols 4-7
- Synonyms
  - ARM/Thumb 4-2
- System V
  - release 4 1-3
  - shared libraries 2-15

## T

- Tail call sections
  - about 3-23
- Thumb
  - branch range 3-10
  - in execution sections 3-10
  - in linker 1-3

## U

- Uninitialized memory 5-15
- User libraries
  - searching 6-5

## V

- Veneers
  - about 3-19
  - ARM to Thumb 1-3
  - generation 3-19
  - in linker 1-3, 3-19
  - in scatter-loading 5-24
  - inline 3-20
  - interworking 3-21
  - long branch 3-20, 3-21
  - placement 3-21
  - position-independent 3-21
  - reuse 3-21
  - sharing 3-19
  - short branch 3-20
  - size 2-28
  - Thumb to ARM 1-3
  - types 3-19
  - variants 3-20
- Version
  - symbol versioning 4-21
- VFE
  - awareness 2-23, 3-12, 3-13, 3-14
  - default 2-23, 3-13
  - force 2-24, 3-14
  - force\_no\_rtti 2-24, 3-14
  - in assembler source files 3-13
  - modes 2-23, 2-24, 3-13, 3-14
  - off 2-23, 3-13, 3-14
  - on 2-23, 3-13
  - RTTI 2-23, 2-24, 3-12, 3-13, 3-14
  - virtual function elimination 3-12, 3-13
- via files
  - in librarian 6-9
  - in linker 2-31
- Virtual function elimination
  - VFE 3-12, 3-13
- Virtual functions
  - elimination (VFE) 3-12, 3-13
  - unused 2-23, 2-24, 3-12, 3-13

## W

- Wildcards
  - in scatter-loading 5-7, 5-17, 5-18

- in steering files 4-10, 4-12, 4-13, 4-14, 4-15, 4-17, 4-18, 4-19
- with armar 6-7
- with armlink 2-25
- with fromelf 7-9
- with input sections 5-17, 5-19

## Symbols

- \$ symbols 4-20
- \$Sub\$\$ 4-20
- \$Super\$\$ 4-20
- \$\$ symbols 4-3
- \$\$ZI\$\$ symbols 4-4

