



Arm[®] Streamline

Version 8.1

User Guide

Non-Confidential

Copyright © 2021–2022 Arm Limited (or its affiliates).
All rights reserved.

Issue 00

101816_0801_00_en



Arm® Streamline

User Guide

Copyright © 2021–2022 Arm Limited (or its affiliates). All rights reserved.

Release information

Document history

Issue	Date	Confidentiality	Change
0708-00	20 August 2021	Non-Confidential	New document for v7.8
0709-00	18 November 2021	Non-Confidential	New document for v7.9
0800-00	18 February 2022	Non-Confidential	New document for v8.0
0801-00	20 May 2022	Non-Confidential	New document for v8.1

Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, has undertaken no analysis to identify or understand the scope and content of, third party patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND

REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <https://www.arm.com/company/policies/trademarks>.

Copyright © 2021–2022 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.

110 Fulbourn Road, Cambridge, England CB1 9NJ.

(LES-PRE-20349)

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Unrestricted Access is an Arm internal classification.

Product Status

The information in this document is Final, that is for a developed product.

Feedback

Arm® welcomes feedback on this product and its documentation. To provide feedback on the product, create a ticket on <https://support.developer.arm.com>

To provide feedback on the document, fill the following survey: <https://developer.arm.com/documentation-feedback-survey>.

Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used language that can be offensive. Arm strives to lead the industry and create change.

We believe that this document contains no offensive language. To report offensive language in this document, email terms@arm.com.

Contents

1 Introduction.....	9
1.1 Conventions.....	9
1.2 Other information.....	10
2 Getting started with Streamline.....	11
2.1 Introduction to Streamline.....	11
2.2 Licenses.....	13
2.3 Standards compliance.....	13
2.4 Profile your Android application.....	14
2.5 Profile your Linux application.....	14
2.6 Profile your bare-metal application.....	15
3 Capture a Streamline profile.....	16
3.1 Starting a capture.....	16
3.1.1 Run a command on the target.....	17
3.2 Live view overview.....	18
3.2.1 Adding and removing events in the Live view.....	19
3.2.2 Software-generated events in the Live view.....	20
3.3 Set capture options.....	21
3.3.1 Set up an energy capture device.....	22
3.4 Analysis options.....	23
3.4.1 Generate debug information.....	23
3.4.2 Select report resolution.....	24
3.4.3 Select images and libraries for profiling.....	24
3.5 Download process images.....	26
3.5.1 Separate debug files.....	27
3.6 Counter Configuration.....	28
3.6.1 Add counters from a template.....	28
3.6.2 Add new events.....	28
3.6.3 Setting up event-based sampling.....	29
3.6.4 Importing and exporting counter configuration files.....	31
3.6.5 Configure SPE counters.....	32

4 Annotate your code.....	35
4.1 Annotations.....	35
4.2 Add string annotations to your code.....	37
4.3 Visual annotations.....	38
4.3.1 Add visual annotations to your code.....	38
4.4 Viewing annotations.....	40
4.5 Annotate macros.....	41
5 Analyze your capture.....	45
5.1 Timeline overview.....	45
5.2 Working with charts.....	47
5.3 Analyze a region of time with the cross-section marker.....	48
5.4 Analyze a region of time with the calipers.....	50
5.5 Customizing your charts.....	50
5.5.1 Chart configuration panel.....	51
5.5.2 Chart configuration toolbar options.....	51
5.5.3 Chart configuration series options.....	55
5.5.4 Create a chart configuration template.....	57
5.5.5 Apply a chart configuration template.....	58
5.5.6 Match dynamic counter sources.....	59
5.6 Viewing application activity.....	60
5.6.1 Details panel modes.....	60
5.6.2 Heat Map mode.....	61
5.6.3 Core Map and Cluster Map modes.....	63
5.6.4 Filtering by threads or processes.....	64
5.6.5 Selecting the activity source.....	64
5.6.6 Samples mode.....	65
5.6.7 Processes mode.....	66
5.6.8 Images mode.....	67
5.6.9 OpenCL mode.....	67
5.7 Toolbar options in the Live and Timeline views.....	70
5.8 Contextual menu options in the Live and Timeline views.....	72
5.9 Warnings tag.....	74
5.10 Analyze your code.....	74
5.10.1 Analyzing Call Paths.....	75
5.10.2 Analyzing Functions.....	76

5.10.3	Sorting table reports.....	77
5.10.4	Viewing your code.....	78
5.11	Streamline Data view.....	80
5.11.1	Capture color-coding.....	81
5.11.2	Re-analyze stored capture data.....	81
5.11.3	Duplicating a capture.....	82
5.11.4	Merge captures.....	83
5.12	Preferences dialog box.....	84
6	Use Streamline from the command line.....	86
6.1	Open a Streamline -enabled command prompt or shell.....	86
6.2	Streamline command-line options.....	86
6.2.1	List modes.....	87
6.2.2	Generate-config mode.....	88
6.2.3	Capture mode.....	88
6.2.4	Analyze mode.....	89
6.2.5	Report mode.....	90
6.2.6	Merge mode.....	93
6.2.7	Import modes.....	94
7	Capturing Energy Data.....	96
7.1	Energy Probe overview.....	96
7.2	Energy Probe requirements.....	97
7.3	Shunt resistor selection for Energy Probe.....	98
7.4	Setting up Energy Probe.....	98
7.5	Adding the caiman application.....	100
7.6	Updating your firmware.....	101
7.7	Energy Probe data.....	101
7.8	Set up National Instrument Multifunction Data Acquisition devices (NI DAQ) to capture energy data.....	102
7.8.1	Rebuild caiman from source.....	104
7.9	Capturing energy data locally and importing it to a capture.....	104
8	Troubleshooting Streamline (user guide).....	106
8.1	Troubleshooting perf import issues.....	106
8.2	Troubleshooting report issues.....	107
8.3	Troubleshooting Energy Probe issues.....	108

8.4 Troubleshooting missing source files.....	109
9 Add custom counters from files and trace events.....	112
9.1 Creating a configuration.xml file.....	112
9.2 Filesystem and ftrace counters.....	114
9.2.1 Create filesystem and ftrace counters.....	116
9.3 Counter classes.....	117
A Keyboard shortcuts.....	119
A.1 Keyboard shortcuts in the Live and Timeline views.....	119
A.2 Keyboard shortcuts in the table views.....	120
A.3 Keyboard shortcuts in the code view.....	121
B Importing raw perf data.....	122
B.1 Import raw perf data.....	122
C JIT profiling support.....	124
C.1 JIT profiling.....	124

1 Introduction

This document describes how to use Arm® Streamline to capture, analyze, and display performance measurements from your target device.

1.1 Conventions

The following subsections describe conventions used in Arm documents.



Glossary





The Arm Glossary is a list of terms used in Arm documentation, together with definitions for those terms. The Arm Glossary does not contain terms that are industry standard unless the Arm meaning differs from the generally accepted meaning.

See the Arm® Glossary for more information: developer.arm.com/glossary.

Typographic conventions

Arm documentation uses typographical conventions to convey specific meaning.

Convention	Use
<i>italic</i>	Citations.
bold	Interface elements, such as menu names. Signal names. Terms in descriptive lists, where appropriate.
monospace	Text that you can enter at the keyboard, such as commands, file and program names, and source code.
monospace bold	Language keywords when used outside example code.
monospace <u>underline</u>	A permitted abbreviation for a command or option. You can enter the underlined text instead of the full command or option name.
<and>	Encloses replaceable terms for assembler syntax where they appear in code or code fragments. For example: <pre>MRC p15, 0, <Rd>, <CRn>, <CRm>, <Opcode_2></pre>
SMALL CAPITALS	Terms that have specific technical meanings as defined in the <i>Arm® Glossary</i> . For example, IMPLEMENTATION DEFINED , IMPLEMENTATION SPECIFIC , UNKNOWN , and UNPREDICTABLE .
 Caution	Recommendations. Not following these recommendations might lead to system failure or damage.
 Warning	Requirements for the system. Not following these requirements might result in system failure or damage.

Convention	Use
 Danger	Requirements for the system. Not following these requirements will result in system failure or damage.
 Note	An important piece of information that needs your attention.
 Tip	A useful tip that might make it easier, better or faster to perform a task.
 Remember	A reminder of something important that relates to the information you are reading.

1.2 Other information

See the Arm website for other relevant information.

- [Arm® Developer](#).
- [Arm® Documentation](#).
- [Technical Support](#).
- [Arm® Glossary](#).

2 Getting started with Streamline

Quickly get started profiling your Android, Linux, or bare-metal application with Streamline.

Streamline is an application profiler that can capture data from multiple sources, including:

- Program Counter (PC) samples from running application threads.
- Samples from the hardware Performance Monitoring Unit (PMU) counters in the Arm CPU and Mali™ GPUs.
- Thread scheduling information from the Linux kernel.
- Software-generated annotations and counters from running applications.

This data capture is extensible, allowing collection from other hardware or software sources that are available.

This chapter provides an introduction to Streamline and directs you to the appropriate getting started instructions for your application.

2.1 Introduction to Streamline

Streamline helps you optimize software for devices that use Arm® processors.

Evaluate where the software in your system spends most of its time by capturing a performance profile of your application running on a target device. Quickly determine whether your performance bottleneck relates to the CPU processing or GPU rendering using interactive charts and comprehensive data visualizations.

For CPU bottlenecks, use the native profiling functionality to locate specific problem areas in your application code. Investigate how processes, threads, and functions behave, from high-level views, right down to line-by-line source code analysis. The basic profile is based on regular sampling of the PC (Program Counter) of the running threads, allowing identification of the hotspots in the running application. Hardware performance counters that are provided by the target processors can supplement this analysis. These counters enable hotspot analysis to include knowledge of hardware events such as cache misses and branch mispredictions.

For GPU bottlenecks, use performance data from the Arm® Mali™ GPU driver and hardware performance counters to explore the rendering workload efficiency. Visualize the workload breakdown, pipeline loading, and execution characteristics to quickly identify where to apply rendering optimizations.

With Streamline, you can:

- Find hot spots in your code to be targeted for software optimization.
- Identify the processor that is the major bottleneck in the performance of your application.

- Use CPU performance counters to provide insights into L1 and L2 cache efficiency, enabling cache-aware profiling.
- Identify the cause of heavy rendering loads which cause poor GPU performance, and use GPU performance counters to identify workload inefficiencies.
- Reduce device power consumption and improve energy efficiency by optimizing workloads using performance counters from the CPU, GPU, and memory system.

Android application profiling

Streamline supports data capture on non-rooted Android devices. It collects CPU performance data and Mali GPU performance data so you can profile your game or app without device modification. Configuring Streamline to collect the right data is easy. Use the built-in templates to select the most appropriate set of counters for your target device. Alternatively use the built-in templates as a starting point for a customized data visualization. [Profile your Android application](#)

Android system profiling

In addition to the single application profiling for non-root devices, Streamline also supports system-wide profiling when running on rooted Android devices. System profiling enables manufacturers to simultaneously monitor all applications and services running on their device, allowing for the identification of problematic processes or scheduling behaviors.

Linux system profiling

Streamline supports system-wide profiling of applications running on Linux-based embedded devices. Analyze the behavior of the system hardware by selecting the required Arm CPU or Mali GPU hardware performance counters for your scenario. This analysis can be supplemented by connecting power measurement probes, such as the Arm Energy Probe or National Instruments DAQ, to provide accurate measurement of system energy use. To provide more context to the analysis, you can use software annotations in Arm software libraries, such as the Mali GPU OpenCL device driver. [Profile your Linux application](#)

Bare-metal application profiling

Streamline can profile bare-metal software running on Arm processors, emitting data over a CoreSight ITM, ETM, or STM data channel. It can also profile data captured to an on-device memory buffer. Support is included for PC sampling, performance counter sampling, and application-generated annotations.

The target agent for Android and Linux captures uses a dedicated daemon application to capture and transmit the data. The target agent for bare-metal profiling is provided as a small source library that is integrated directly into the bare-metal application that is being profiled. This library - called Barman - is auto-generated based on the data channel configuration that is required. It provides optional hooks that allow for lightweight RTOS integration, such as annotation of thread or task context switches. This integration can supplement the basic performance information in the data visualizations. [Profile your bare-metal application](#)

2.2 Licenses

Arm® Streamline contains license-managed features that are only accessible when they are enabled by one of the available commercial licenses. Depending on the product purchased, apply licenses either per-user, or per-machine:

- To activate per-user licenses, select **Help > License Management** in the graphical user interface. You can also activate per-user licenses using the standalone license management utility `armlm-gui`, which you can find in the studio installation.
- To activate per-machine licenses for Arm Development Studio, use the main Development Studio IDE. Arm Mobile Studio licenses must be activated using environment variables, as described in [Adding a professional license](#) on the Arm Developer website.

2.3 Standards compliance

Arm® Streamline conforms to various levels of compliance.

Hardware

Streamline supports all Armv7 and later architecture-based CPUs running Linux and Perf, and is tested on Arm® Cortex®-A cores. Arm Cortex-R and Arm Cortex-M cores are also supported for bare-metal profiling.

Streamline supports Arm® Mali™ GPUs implementing the Midgard (excluding Mali-T600 series), Bifrost, and Valhall architectures.

Executable and Linking Format (ELF)

Streamline can read ELF-formatted executable images.



Streamline can scan **Android Application Package (APK)** file archives for ELF images. It extracts the valid ELF executable images and includes them as program images for analysis.

Debugging With Attributed Record Formats (DWARF)

Streamline can read debug information from ELF images in the DWARF 2, DWARF 3, and DWARF 4 formats.

2.4 Profile your Android application

To get started profiling an application on an Android target, use the [Arm Streamline Target Setup Guide for Android](#).

Before you begin

- Download and install Arm® Mobile Studio or Arm Development Studio, and install adb to [set up your host machine](#).



The latest version of Streamline is compatible with captures taken with Streamline version 7.3 or later. For older captures, reanalyze the data using the instructions in [Re-analyze stored capture data](#).

- Set developer options, connect to the host, and install your application to [set up your target device](#).

Procedure

Follow the instructions in the [Arm Streamline Target Setup Guide for Android](#).

Next steps

- [Analyze your capture](#) .
- See tutorial [Android performance triage with Streamline](#) on the Arm Developer website.

2.5 Profile your Linux application

To get started profiling an application on a Linux target, use the [Arm Streamline Target Setup Guide for Linux](#).

Before you begin

- Download Arm® Development Studio from [Downloads](#).



The latest version of Streamline is compatible with captures taken with Streamline version 7.3 or later. For older captures, reanalyze the data using the instructions in [Re-analyze stored capture data](#).

- Install Arm Development Studio by following the instructions in the [Arm Development Studio Getting Started Guide](#).

Procedure

Follow the instructions in the [Arm Streamline Target Setup Guide for Linux](#).

Next steps[Analyze your capture](#) .

2.6 Profile your bare-metal application

To get started profiling an application on a bare-metal target, use the [Arm Streamline Target Setup Guide for Bare-metal Applications](#).

Before you begin

- Download Arm® Development Studio from [Downloads](#).



The latest version of Streamline is compatible with captures taken with Streamline version 7.3 or later. For older captures, reanalyze the data using the instructions in [Re-analyze stored capture data](#).

- Install Arm Development Studio by following the instructions in the [Arm Development Studio Getting Started Guide](#).

Procedure

Follow the instructions in the [Arm Streamline Target Setup Guide for Bare-metal Applications](#).

Next steps[Analyze your capture](#) .

3 Capture a Streamline profile

See how to capture a profile of your application using Streamline. Learn how to set different capture and analysis options to change what information is captured and reported.

3.1 Starting a capture

Start a capture session to profile data from your application in real time. When the capture session finishes, Streamline automatically opens a report for you to analyze later.

Before you begin

Set up your Linux or Android target as described in the relevant Streamline target setup guide.

About this task



For instructions on how to capture data locally on a target, and for more target-specific information, see the Streamline target setup guides.


Procedure

1. In the **Start** view, select your target device type. Then select your device from the list of detected targets, or enter the address of your target.





By default, `gator` uses port 8080 to communicate with the host machine. To use a different port, append a colon followed by the port number to the target address. For example, to use port 5050 with target address 10.99.28.54, enter `10.99.28.54:5050` in the target address field.


2. Android users only, select the package you want to profile from the list of packages available on the selected device.
3. TCP users only, optionally enter the details for any command you want to run on the application.
4. Configure the counters to collect.

Click **Configure Counters**, then **Add counters from a template** . Select a counter template appropriate for the target from the drop-down list. Alternatively, manually select the CPU and GPU counters you want to capture.

Counter templates are pre-defined sets of counters that have been chosen to enable you to perform an initial performance review of both CPU and GPU behavior. Streamline notifies you if the target device does not support all the counters that are defined in the selected template.

5. Optionally, click **Advanced Settings** to set more capture options, including the sample rate and the capture duration (by default unlimited).

6. Click **Start Capture**.
7. Browse to a location on the host and enter the name of the capture file.
8. Run the application scenario that you want to profile on your target device.
9. In the **Capture Control** view:
 - To stop the capture session and create a performance report, click **Stop capture and analyze** .
 - To stop the capture session and discard the collected data, click **Stop capture and discard** .

If **Stop capture and analyze**  is grayed out, **Discard Data** has been selected in the **Capture and Analysis Options** dialog box. To generate a performance report, clear **Discard Data** in the **Capture and Analysis Options** dialog box, then click **Start Capture**.

Results

When data analysis completes, the application terminates on the device and the report automatically opens in the **Timeline** view.

Next steps

- [Analyze your capture](#)
- See tutorial [Android performance triage with Streamline](#) on the Arm Developer website.

Related information

[Run a command on the target](#) on page 17

[Set capture options](#) on page 21

[Analysis options](#) on page 23

[Download process images](#) on page 26

[Counter Configuration](#) on page 27

[Arm Streamline Target Setup Guide for Android](#)

[Arm Streamline Target Setup Guide for Linux](#)

[Arm Streamline Target Setup Guide for Bare-metal Applications](#)

3.1.1 Run a command on the target

Specify a command to run on the target after profiling starts.

Before you begin

- To allow Streamline to communicate with the target, install and run `gatord` on the target.
- You must use the `-a` option when starting `gatord`.

Procedure

1. In the **Start** view, under **Configure Application**, enter a command to run on the target.

2. The command runs in the current working directory of `gatord`. To run the command in a different directory, enter the absolute path to that directory in the **Working Directory** field.
3. The command runs as the same user as `gatord`. For example, if `gatord` is run as root, the command runs as root. To run the command as a different user, enter the user account in the **User Name** field.
4. To stop the capture session when the command specified in the **Command** field ends, select **Stop on exit**. If you want the capture session to continue after the command ends, leave the checkbox clear.

Results

When you click **Start Capture**, `gatord` runs the specified command on the target.

3.2 Live view overview

The **Live** view shows data capture in real time during the capture process. When you start a capture session, the **Live** view opens automatically.

The **Live** view gives feedback during a capture session. The full functionality of Streamline is only available after you have stopped your capture session and Streamline has processed the data in the capture.

The **Live** view toolbar displays the duration of the live capture in seconds. If there is latency in the data passing from the target, the text **Target latency - resyncing...** is displayed beside the capture duration. The **Live** view stops scrolling until the data can catch up.



To display data in the **Live** view, you must set **Buffer Mode** to **Streaming** in the **Capture and Analysis Options** dialog box. If **Streaming** is not active during a capture session, the **Live** view still appears, but without any real-time display of data.

Process list

The process list shows the known processes and their usage statistics. This list continuously updates while the capture session is running. When a process dies, its name and ID remain in the list but are shown in gray. For example, the row for `xaos` in the following list is gray:

Index	Process ID	Process Name	% CPU
0	24674	gatord	0.87%
1		kernel	0.06%
2	24666	xaos	
3	1	init.sysvinit	0.00%
4	24236	sshd	0.00%

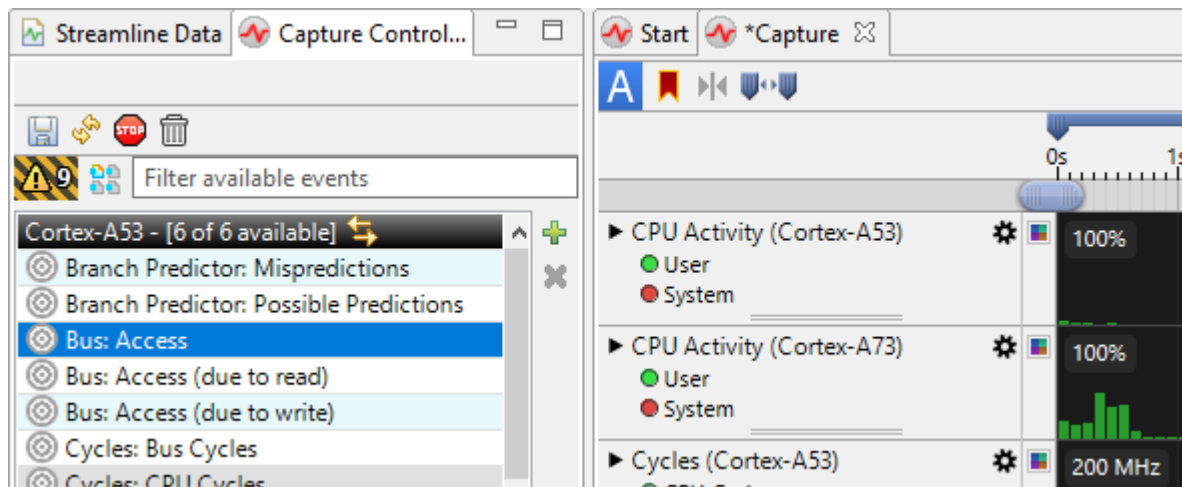
Downloading process images

To download process images from the target, select the **Download process images from target** checkbox.


When the capture stops, a dialog box is shown where you can select the images to download.

3.2.1 Adding and removing events in the Live view


During data capture, add or remove events in the interactive counters panel in the **Capture Control** view.



To select events and add them to the **Live** view, do one of the following actions:

- Select an event, then click **Add selected events** .
- Double-click the event.
- Select an event, then press Enter.
- Drag an event onto the chart panel.

To remove events from the **Live** view, do one of the following actions:

- Select an event, then click **Remove selected events** .
- Select an event, then press Delete.
- Remove the chart or series.




Note

Removing a chart or series only removes the counter if no other chart or series is using it.

- Drag and drop a chart into the interactive counters panel.

When you have finished configuring the **Live** view charts, you can update the view:

- To save the current counter configuration and restart the capture with the new configuration, click **Save and restart** .



If you have selected the **Discard Data** option in the **Capture and Analysis Options** dialog box, this button is inactive and appears grayed out. When the capture session terminates, Streamline discards the data.

- To discard the current capture and restart with the new configuration, click **Restart capture** 💰 .

Changes to the counter configuration also occur when you:

- Edit a series expression, counters are automatically added or removed.
- **Switch and manage templates** 📄 .
- **Add default charts** 📊 .

Setting event-based sampling

If an event supports event-based sampling, **Toggle event-based sampling** 🎯 appears before the event name in the interactive counters panel.

To enable or disable event-based sampling for an event, click **Toggle event-based sampling** 🎯 . Set the reporting threshold in the **Threshold** field that appears at the bottom of the interactive counters panel when event-based sampling is enabled.

Other user indicators

- A gray warning triangle next to the series title indicates event charts that have been configured but not committed.
- A red triangle next to the series title indicates that there are no available counters for the chart.
- Hovering over an event in the interactive counters panel highlights any series that is using it.
- Hovering over a series or chart title highlights any events in the interactive counters panel that it is using.
- Not all events are linked to the charts. Dynamic counters, like atrace, can only be configured using the counters panel, not using the charts. A double arrow in the category heading indicates linked categories.

3.2.2 Software-generated events in the Live view

During a capture, the application on the target device can generate software-defined events and counters. These software-defined counters automatically appear in the **Live** view when they are first generated, but they are not available to select in the **Counter Configuration** dialog.

When you apply a template in the **Live** view, the template is applied to any existing counters. After you have applied a template, any new software counters generated will automatically appear in the **Live** view, even if they are not part of the selected template. To filter out any unwanted series, reapply the template.

3.3 Set capture options

The default capture settings work well for most situations but you can change the capture settings to customize your profiling session. These settings include the sample rate, buffer size, and duration of the capture.

Before you begin

If you are recording call stacks, compile your images and libraries with frame pointers enabled using the `-fno-omit-frame-pointer` compiler option. For Arm®v7 software applications, you must also disable the Thumb instruction set by using the `-marm` option.

Procedure

1. In the **Start** view, click **Advanced Settings**.
2. The **Normal** sample frequency is 1kHz. To change how regularly Streamline collects samples, change **Sample Rate** to:
 - **High** to increase sample frequency to 10kHz and enhance visibility, but also increase overhead.
 - **Low** if you have a slow target, or if the target is heavily loaded, for a longer capture at a sample frequency of 100Hz.
 - **None** to have the lowest level of intrusion on your code, and disable periodic sampling. Any report columns that rely on sampling show zero values.
3. If you do not want to display real-time data during a capture session, set **Buffer Mode** to the appropriate buffer size. The capture ends when the buffer is full:

Buffer mode	Buffer size
Large	16MB
Medium	4MB
Small	1MB

Set **Buffer Mode** to **Streaming** to enable **Live** view and allow streaming of target data directly to your host during capture.

4. To record call stacks and improve the visibility of the behavior of the target, select **Call Stack Unwinding**.
 - Recording call stacks increases the amount of raw data that the target sends to the host.
 - To support call stack unwinding for binaries that are created using GCC or LLVM, enable frame pointers before you compile them.
5. To filter out events that happen in the kernel, and restrict the capture to userspace events only, select **Exclude Kernel Events**.
6. To collect Ftrace data more efficiently, select **Use more efficient Ftrace collection**. The **Live** view will not update charts that use Ftrace data, but the data will be in the finished report.
7. To set a **Duration** for the capture, enter the time that you want the capture session to last either in seconds or in minutes and seconds.
 - Enter 60 to set a capture session for 60 seconds.

- Enter 1:20 to set a capture session for 1 minute 20 seconds.

If you do not enter a duration, the capture session continues until you stop it manually or, if you are not using streaming, until the buffer is full.

8. Click **Save**.


3.3.1 Set up an energy capture device

Set up an energy capture device to measure energy consumption during a capture session. These measurements help you identify inefficient use of energy that might compromise the user experience.

Procedure

1. In the **Start** view, click **Advanced Settings**.
2. Choose the appropriate option from the **Energy Capture** menu:
 - If you do not want to capture energy data, select **No Energy Data Collection**. All other energy capture options become inactive.
 - If you do want to capture energy data, select **Arm Energy Probe** or **NI DAQ** to match your energy capture hardware.
3. In the **Device** field, enter the name of your target energy capture device. If you leave the **Device** field blank, Streamline attempts to detect your device automatically.
 - When using Arm Energy Probe on Linux, enter `/dev/ttyACM0` in the **Device** field, if necessary.
 - When using NI DAQ, the device name depends on the drivers that are installed on the host. For example, when using NI DAQmx Base drivers, the device name is usually Dev1. You can determine the device name by using the National Instruments List Devices utility.
4. If you are not using the default port 8081, enter the port number that you want Streamline to use to communicate with your energy capture device.
5. In the **Tool Path** field, enter the path to the `caiman` executable. This executable is required to use either the Arm Energy Probe or a NI DAQ device to gather power output statistics.



To browse your file system, click **Select the energy capture tool** .

6. Enter the configuration settings for each channel:
 - To collect energy data, in joules, select **Energy**.
 - To collect power data, in watts, select **Power**.
 - To collect voltage data, in volts, select **Voltage**.
 - To collect current data, in amps, select **Current**.
 - Enter the **Resistance**, in milliohms, of the shunt resistor that connects to each of the available channels. The default setting is 20 milliohms.

7. Click **Save**.

Results

When you have connected your energy probe and run an energy capture session, the collected energy data appears in the **Timeline** view. Use the energy measurements to identify opportunities to improve energy consumption in the application, such as altering the frame rate and overdraw rate.

Related information

[Capturing Energy Data](#) on page 96

3.4 Analysis options

Use the **Capture and Analysis Options** dialog box to change settings that control the data that is provided for analysis of your capture sessions. For example, debug information, report data, and image files.

To open the **Capture and Analysis Options** dialog box, in the **Start** view, click **Advanced Settings**.

3.4.1 Generate debug information

Process DWARF debug information and line numbers for a higher level of detail in your capture session.

Before you begin

To process debug information, you must have built the image using the `-g` compiler option.

Procedure

1. In the **Start** view, click **Advanced Settings**.
2. In the **Analysis** section, select **Process Extra Debug Information** to process source code mappings and call stack information. If you clear **Process Extra Debug Information**, the source section shows a `No source available` message.



- Selecting **Process Extra Debug Information** results in higher memory usage.
 - This option can be changed when you re-analyze the stored capture.
-

3. Click **Save**.

Results

The report data that is automatically generated after the termination of the capture session contains information to help you debug your code, including:

- Source code mappings.

- Call stack information.

3.4.2 Select report resolution

Change the data resolution to determine the level of detail provided in the report data that is generated from your capture session. A higher level of detail allows you to view data with a finer time granularity, but results in a longer analysis time and higher memory usage.

Procedure

1. In the **Start** view, click **Advanced Settings**.
2. In the **Analysis** section, select the **Resolution mode** you require for data analysis.
The resolution determines the size of the time bins that you can zoom into.

Resolution mode	Time bins
Normal	1 millisecond
High	100 microseconds
Ultra-High	1 microsecond



You can change the **Resolution mode** option when you re-analyze the stored capture.

3. Click **Save**.

Results

The automatically generated report data contains the specified level of detail processed from your capture session.

3.4.3 Select images and libraries for profiling

Browse your file system, and select the images and libraries you want to profile.

Before you begin

- When compiling images on your host, ensure that you use the `-g` compiler option to enable debug symbols.
- To improve the call path quality, disable inlining with the `-fno-inline` compiler option.




For more information on compiling your application, see the [Arm Streamline Target Setup Guide for Android](#) or the [Arm Streamline Target Setup Guide for Linux](#).

Procedure

1. In the **Start** view, click **Advanced Settings**.
2. In the **Program Images** section, control symbol loading for the ELF images, executables, or APK listed:
 - Select **Use all images for analysis** or **Use none of the images for analysis**.
 - To browse your file system and choose an image executable or APK, click **Add ELF Image**, select the required file, then click **Open**.
 - To select a separate file containing the extracted debug information and CFI, click **Select Separate Debug Image**.



This option is available when the **Debug Info** or **CFI** column displays the  icon.

- To remove the selected entries, click **Remove**.
3. In the **Program Images** table, select the checkbox in the **Use** column for each image you want to include in the analysis.
 4. Click **Save**.



These images can be changed when you re-analyze the stored capture.

Results

The images you selected in the **Program Images** table are now ready for profiling.

See [Download process images](#) for more information about the image download functionality, which is an alternative to manually selecting images and libraries on the host in the **Program Images** section.

3.4.3.1 Program Image Table Reference

Symbols are displayed in the **Program Images** table to show which information is available for each program image.

Here is what the following symbols mean in the **Program Images** table:

Symbols





Displays the following icons to indicate whether symbol information is available:

- ✓ The image contains symbol information.

-  No symbol information is available.





Debug Info

Displays the following icons to indicate whether debug information is available:

-  Debug information is available, providing source line numbers and inlining information.
-  Debug information is available in a separate file (split debug image).
-  File contains no debug information, and no link to external debug information.
-  Availability of debug information is unknown.

CFI

Displays the following icons to indicate whether CFI is available:

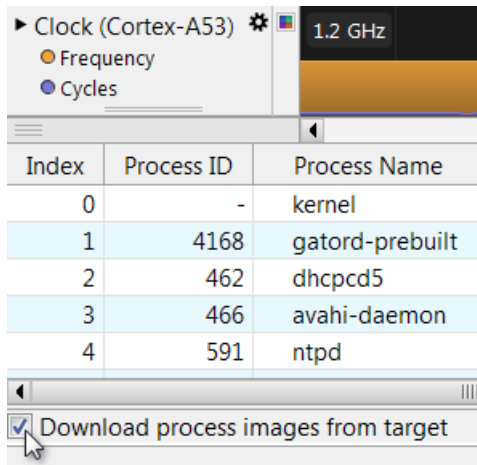
-  The `debug_frame` or `eh_frame` section is available. This information is used to determine stack frame sizes in the **Call Paths** view.
-  The `debug_frame` or `eh_frame` section is available in a separate file.
-  File contains no CFI, and no link to external CFI.
-  Availability of CFI is unknown.

3.5 Download process images

Download process images from the target device at the end of a live capture session using the image download functionality.

Procedure

1. To trigger the image download when the session ends, select **Download process images from target** at the bottom of the **Live** view.



When you stop the capture, the **Image Download** dialog box opens. The upper part of this dialog box lists the processes that are running on the target. To add any libraries on which the selected process depends, click **Add Process and Libs**. The lower part lists the process images that you have selected for download.

2. Select a process then click **Add Process** to add it to the list of images to download. To remove an image from this list, select the image then click **Remove**.
3. To download the selected images, click **Start Downloading**.

Results

You can now analyze the process images that you selected in the **Image Download** dialog box.

Related information

[Separate debug files](#) on page 27

3.5.1 Separate debug files

To limit the size of the executable image file, you can save debug information in a separate file. Streamline supports the separate debug file linking process that is used in many ELF images.

You can locate the separate debug file in two ways:

- The executable includes the name of the debug file, which usually is derived from the executable file name and suffixed with `.debug`.
- The debug file name is derived from the build ID in the executable.

For more information about debugging information in separate files, see <https://sourceware.org/gdb/onlinedocs/gdb/Separate-Debug-Files.html>.

3.6 Counter Configuration


Streamline uses hardware performance counters to collect the data you need to analyze your applications. You can edit the set of counters with the **Counter Configuration** dialog box, which is accessed through the **Start** view.

Mali™ GPUs implement a comprehensive range of performance counters, that enable you to closely monitor GPU activity as your application runs. The charts in Streamline visualize this performance counter activity, to help you identify the cause of heavy rendering loads or workload inefficiencies that cause poor GPU performance. For detailed descriptions of all the performance counters available for each Mali GPU, refer to the [Mali GPU Counter references](#) on the Arm Developer website.

3.6.1 Add counters from a template

Counter templates are pre-defined sets of counters. They provide a quick way to review performance of both CPU and GPU behavior, in particular to analyze Mali™ GPU performance.

Procedure

1. In the **Start** view, click **Configure Counters**.
The **Counter Configuration** dialog box opens.
2. Click **Add counters from a template** .
3. Select a template from the drop-down list.
 - Streamline notifies you if the target device does not support all the counters that are defined in the selected template.
 - If a template shows a number with a +, such as (55/55+), the template contains wildcards that cannot be resolved until the counters are dynamically generated.

Results

The template populates the **Events to Collect** list.

Next steps

If there are events that you want to profile that the template does not include, see [Add new events](#). Otherwise you are ready to start your capture.


3.6.2 Add new events

Streamline uses the counters in the **Events to Collect** list to determine what data it collects during a capture session. If an event you want to profile is not shown, you can add it to the list.

Before you begin

You must have set up your host machine and connected to a target device as described in the relevant Streamline target setup guide.



Procedure

1. In the **Start** view, click **Configure Counters**.
2. Double-click an event, or click and drag an event to move it from **Available Events** to the relevant category in the **Events to Collect** list.
 - The **Available Events** list contains categorized events that are offered for each core on your target, and other hardware and OS-specific events. Events that are contained in the processor lists are based on the PMU counters of the core. They can vary depending on the type of processor, as can the number of events that you can add.
 - To show advanced events for your target in the **Available Events** list, click **Toggle advanced view** .

The added events appear in gray in the **Events to Collect** list under their category name, and are used by the **Timeline** view for its graphs. Each event that is listed here is available for display in a chart in the **Timeline** view.




Note

- Selecting a high number of counters might affect performance of the application on the target.
- To automatically collect the set of counters that are required to generate the charts defined in a chart configuration template, see [Add counters from a template](#). To setup chart configuration templates, click **Switch and manage templates**  in the **Live** view and **Timeline** view.
- To reset the **Events to Collect** list to default, click **Load Defaults** .

3. If a counter can be collected from different cores or interfaces, select a specific core or interface from the menu next to the counter in the **Events to Collect** list.
4. Click **Save**.

Next steps

The number of hardware-specific events that you can collect during a capture session is limited. You can remove unwanted counters from the **Events to Collect** list to make space for you to add new counters. In the **Counter Configuration** dialog box, select all the events that you want to remove from the **Events to Collect** list, then click **Delete** .

Related information

[Create a chart configuration template](#) on page 57

3.6.3 Setting up event-based sampling

The default behavior of Streamline is to capture data samples based on a regular time step. Alternatively, event-based sampling (EBS) enables you to take samples when a user-defined


counter reaches a threshold value. For example, triggering data samples for every 100,000 level two cache misses.


About this task








- Only CPU hardware counters and counters that are in the category **Perf Software** support EBS.
 - EBS is only possible when the PMU on the target hardware can generate interrupts.
-

Procedure

1. In the **Start** view, click **Configure Counters**.
2. Click **Toggle event-based sampling**  for an event in the **Events to Collect** list. Event-based sampling is active for all matching events.

If **Toggle event-based sampling**  is not displayed next to an event, then that event does not support event-based sampling.

3. Enter a value in the **Threshold** field.
Avoid setting a very low threshold for high frequency events, because it might cause a capture to fail. If you enter a threshold value that generates too many samples, the capture could fail, and you might have to restart your target. Your target for the **Threshold** field should be to generate 1000 samples per second. Therefore, if your application typically generates two million counter increments per second, a good value to insert in the **Threshold** field is 2000.

Events to Collect






Cortex-A53
<input checked="" type="radio"/> Branch Predictor: Mispredictions
<input type="radio"/> Bus: Access
<input type="radio"/> Cycles: CPU Cycles
<input type="radio"/> Instructions (Executed): All
<input type="radio"/> L2 Data Cache: Access
Exynos-M1/M2
<input checked="" type="radio"/> Branch Predictor: Mispredictions
<input type="radio"/> Bus: Access
<input type="radio"/> Cycles: CPU Cycles
<input type="radio"/> Instructions (Executed): All
<input type="radio"/> L2 Data Cache: Access
Linux
Clock: Frequency (Cortex-A53)
Clock: Frequency (Exynos-M1/M2)
CPU Activity: System (Cortex-A53)

Event Based Sampling



Event: Branch Predictor: Mispredictions

Threshold:

3.6.4 Importing and exporting counter configuration files

Export the current counter configuration to an XML file so that you can capture data locally. Import a counter configuration file that you previously exported to overwrite your current counter configuration file.

Procedure

1. In the **Start** view, click **Configure Counters**.
2. Choose the relevant action:
 - To search for and import a counter configuration XML file, click **Import** .
 - To export the current counter configuration to an XML file, click **Export** . By default, the counter configuration file is called `configuration.xml`, and is saved to the same directory as `gator` on your target.

3. If you want to capture data locally, you can manually add the exported XML file as an option when running `gator` on the target.
4. Click **Save** to confirm and close the **Counter Configuration** dialog box.

3.6.5 Configure SPE counters

Filter sample records that are generated from Statistical Profiling Extension (SPE) counters using the **Counter Configuration** dialog box. Sample records that do not meet the filtering criteria are discarded and are not written to the profiling buffer.

Before you begin

- Hardware with the statistical profiling extension.
- A Linux kernel with:
 - The `arm_spe_pmu` module enabled.
 - Support for SPE in the device tree or UEFI.
- Depending on the hardware and kernel configuration, you might also need to disable KPTI by using the kernel command-line argument `kpti=off` when booting the device. Currently disabling KPTI is known to be required on Armv8.2 processors with SPE, such as the Neoverse N1.

About this task

Usually Streamline collects information from hardware counters, which provide aggregate counts. You can therefore only attribute counters to regions of the application. The PC is sampled in software using a timer interrupt, so the sampling rate is limited. SPE is an optional extension to the Arm®v8.2-A architecture that samples the PC periodically in hardware as part of the pipeline of the processor. This form of sampling has a low probe effect so you can set the sampling rate much higher. As it is built into the pipeline, SPE can also collect extra information about each sampled instruction. It therefore allows for a much more detailed analysis of the executed code.

Streamline supports visualizing the following SPE data:

Latency counter packets

Provide issue and total instruction execution latency counts, which help identify execution stalls. They also provide load and store latency counts for memory accesses. Use this data to identify high latency accesses and poor cache use.

Event packets

Provide information about each sampled instruction, including whether:

- It accessed, hit, or missed a cache level.
- It was a mis-predicted or not taken branch.
- An exclusive load or store failed.

Use this data to identify branch prediction problems, poor cache use, and lock contention.

Data source packets

Provide information about the level of the memory hierarchy that a load or store accessed.

By default, all operations are sampled, but you can filter the samples using the SPE settings in the **Counter Configuration** dialog box.

Procedure

1. In the **Start** view, click **Configure Counters**.
The **Counter Configuration** dialog box opens.
2. In the **Events to Collect** list, select a counter in the **Statistical Profiling Extension** category.
3. Configure the SPE counter.

Events to Collect

Linux

- CPU Activity: System
- CPU Activity: User
- Disk I/O: Read
- Disk I/O: Write
- Memory: Free
- Memory: Used
- Statistical Profiling Extension**
- Armv8.2 Statistical Profiling Extension**

Statistical Profiling Extension

Instruction Type Filter

☐ Load ☐ Store ☐ Branch

Sampling **ALL** operations including 'Other' operations

Event Filter:

- ☐ Architecturally retired
- ☐ Level 1 Data cache refill
- ☐ Mispredicted
- ☐ TLB walk

Sampling **ALL** operations

Minimum Latency:

For example:

- To identify operations that are slow because they access memory instead of the cache, set a minimum latency.
 - To find branches that trigger mispredictions, select the **Mispredicted** event filter.
4. Click **Save** to save your settings and close the **Counter Configuration** dialog box.

Results

When the analysis is complete, Streamline displays charts for the SPE counters in the **Timeline** view. It also adds SPE data to the **Call Paths**, **Functions**, and **Code** views.

Next steps

Examine the data at the thread, function, source line, or instruction level in the **Call Paths**, **Functions**, and **Code** views.

Related information

[Analyze your code](#) on page 74

[Analyzing Call Paths](#) on page 74

[Analyzing Functions](#) on page 76

[Introduction to Statistical Profiling Support in Streamline](#)

4 Annotate your code

Add event annotations and software-generated counters to your code, which are propagated into the **Timeline** view and **Log** view.

4.1 Annotations

For extra context, add annotations to your source code. These annotations appear in the **Timeline** view and the **Log** view.

There are two types of annotations, user space annotations and kernel annotations.

User space annotations are for annotating your application. Streamline supports the following types of user space annotations:

String

String annotations write user-defined strings into the **Log** view and place framing overlays in the details panel in the **Timeline** view.


Visual

Visual annotations add images to the visual annotation chart in the **Timeline** view. The images are also displayed in the **Log** view.

Marker

Marker annotations add bookmarks to the **Timeline** view and the **Log** view, optionally with a text string. Bookmarks enable you to label and quickly return to points of interest.

Hovering over a bookmark displays an overlay that contains the timestamp and text string of the bookmark.



Custom counters

These annotations are counters that the application software generates. After you have run a capture session, the **Timeline** view displays a chart for each custom counter you have created.

Groups and channels

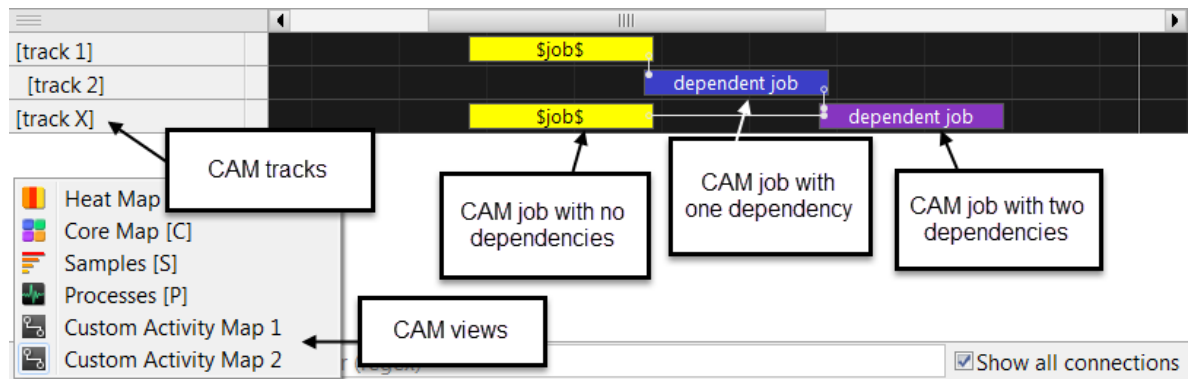
This set of macros enables you to break down threads into multiple channels of activity and assign each channel to a group. A number identifies each group and channel, and each one has a name for display in the **Timeline** view and the **Log** view. Groups and channels are

defined per thread. Therefore although each channel number must be unique within the thread, channels in different threads can have the same number.

Custom activity maps

These annotations allow you to define and visualize a complex dependency chain of jobs. Each custom activity map (CAM) view contains one or more tracks and each track contains one or more jobs.

Jobs might have dependencies on other jobs. Dependencies between jobs are shown using connecting lines with circles at each end to indicate the direction of the dependency. A job with a closed circle depends on a job with an open circle.



Arm NN Timeline

If you are using Streamline with Arm NN and timeline profiling is enabled, there is an Arm NN Timeline view for each process in Arm NN that is running. Each row in the view is a neural network, which you can expand to see the layers of the neural network. Connections between layers indicate that the output of one layer is the input for another layer. Expanding a layer shows the workloads in that layer. Workload names are in the form `workload:<Backend>(<GUID>)`, where `<Backend>` is the hardware or mechanism responsible for the workload, and `<GUID>` is a Globally Unique Identifier. The jobs on each row indicate how long each inference of the network, each layer, or each workload ran for.

Kernel annotations are for annotating the kernel module, allowing you to profile any system calls. Streamline supports string, marker, and custom counter kernel annotations. For more information about how to use kernel annotations, see `README.md` and the annotation examples located in `<install_directory>/sw/streamline/gator/trace/events:`

- `example_standalone`, where annotations are added in the same module.
- `example_shared`, where annotations are exported from a different module.

Annotation macros are defined in `streamline_annotate.h` and `gator_annotate.h`.



Warning

`streamline_annotate.h`, `streamline_annotate.c`, `gator_annotate.h`, and `gator_annotate.c` are intended for use in a development environment. Arm does not recommend including these files in a released product without performing a security audit of the source code first.



Arm® Development Studio has examples of annotations in `<install_directory>/sw/streamline/examples/annotations/`. See the `readme` file in the same directory for further information.

Related information

[Add string annotations to your code](#) on page 37

[Annotate macros](#) on page 40

4.2 Add string annotations to your code


String annotations write user-defined strings into the **Log** view and place framing overlays in the details panel in the **Timeline** view.

Procedure

1. Include `streamline_annotate.h` and `streamline_annotate.c` in your project. Alternatively, build a `libstreamline_annotate` library with the Makefile in `<install_directory>/sw/streamline/gator/annotate`, and include this library in your project.
2. Add the `ANNOTATE_SETUP` macro to your code. Call this macro before any other annotate macros.
3. Add annotations to your code using the macros defined in `streamline_annotate.h`. For example, write a null-terminated string from any thread and set its color using the `ANNOTATE_COLOR(color, string)` macro. Either choose a color constant that is defined in `streamline_annotate.h` or send the ASCII escape code followed by a 3-byte RGB value.
4. Use `ANNOTATE_END()` to clear the annotation message for the thread.
5. Ensure `gatord` is running.
6. Run your application and exercise the area of the code that emits the annotations.

Results

String annotations are displayed as text overlays inside the relevant channels in the details panel of the **Timeline** view, for example inside **Channel 0** in the following screenshot. The letter **A** is displayed in the process list to indicate the presence of annotations. String annotations are also displayed in the Message column in the **Log** view.



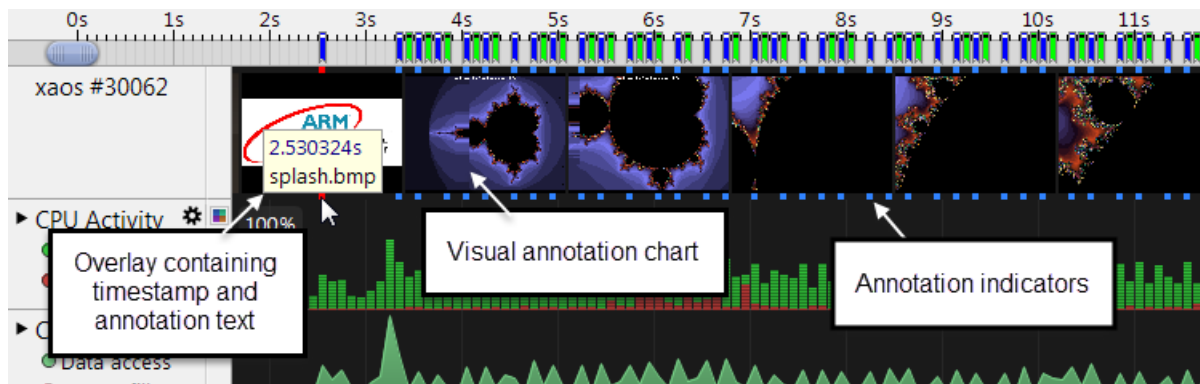
Related information

[Annotations](#) on page 35

[Annotate macros](#) on page 40

4.3 Visual annotations

If you add images to the capture data, they appear in a chart in the **Timeline** view. This chart enables you to track the visuals of your application with the other chart data.



Visual annotation charts have the following features:

- If there is more than one image in the time period that a thumbnail covers, the first image from the range is displayed. Hovering your mouse over the thumbnail displays the image at the current mouse position. Moving the mouse over the thumbnail reveals the other annotated images in that range.
- Yellow or blue indicators at the top and bottom of the chart identify the time bins in which your code produced each image. Yellow indicators indicate that the annotation contains text in addition to the image. Blue indicators indicate that there is no additional text.
- If you hover over an image in the **Timeline** view, the indicators at the cursor position turn red. The red indicators identify the time bin of the image currently being displayed. Hovering over a red indicator displays an overlay that contains the timestamp and, if present, the text annotation associated with the image.
- Clicking an image switches the Details panel to **Images** mode and shows an enlarged version of the image. To view an animation of the images in the Details panel, click and drag the mouse to the left or right over the image.

Related information

[Add visual annotations to your code](#) on page 38

4.3.1 Add visual annotations to your code

Add images to the **Timeline** view and the **Log** view in your analysis report by adding macros to your code.

Before you begin

Do one of the following:

- Include the files `streamline_annotate.h`, `streamline_annotate.c`, and `streamline_annotate_logging.h` in your project.
- Build a `libstreamline_annotate` library with the Makefile in `<install_directory>/sw/streamline/gator/annotate`, and include this library in your project.

About this task



Mali™ GPUs automatically emit visual annotations when you select a Filmstrip counter.

Images for visual annotations can be in the following formats:

- GIF
- PNG
- JPEG
- TIFF
- ICO
- BMP +RLE

There is no limit to the image size but the larger the image, the greater the impact on system performance. Increasing the amount of data that is sent to the host in this way increases the probe effect for the applications you are profiling.

Procedure

1. Add the `ANNOTATE_SETUP` macro to your code.
This macro must be called before any other annotate macros are called.
2. Insert the `ANNOTATE_VISUAL(data, length, str)` macro into your code.
Where:
 - `data` is your image.
 - `length` is the size of the data being written to the annotate file.
 - `str` is an optional descriptive string to be included with the image.
3. Build your code.
4. Ensure that `gator` is running then run your application and use the parts of the application that emit the annotations.

Results

The **Timeline** view now includes a chart that displays the images. In the **Log** view, any annotation event that includes an image has a camera icon 📷 in the message field. To see the image, select the row containing the camera icon.

Related information

[Visual annotations](#) on page 38

[Annotate macros](#) on page 40


4.4 Viewing annotations

View code annotations in the **Timeline** view and the **Log** view.

Timeline view



Annotations appear in several parts of the **Timeline** view:

- String annotations appear in the details panel.
- Software counter annotations appear in their own chart.
- Visual annotations appear in their own chart.
- Marker annotations appear as bookmark icons above the charts.

To see the annotation log for an area of interest, move the cross-section marker to this area then click **Select annotation log closest to the cross-section marker index** , or press L on your keyboard. The **Log** view opens with the closest annotation to the cross-section marker selected.

Log view

The **Log** view lists all the annotations in the capture in chronological order. It also gives the thread and process that generated the annotation. If you assigned the annotation to a group and channel, the group and channel numbers are given. If you did not assign it to a group and channel, the **Log** view displays Group 0 and Channel 0 in these columns.

Marker annotations are indicated with a bookmark icon . Similarly, visual annotations are indicated with a camera icon . To see the image for a visual annotation, select the row for that annotation.

To filter the annotations, use the fields at the top of the view. You can enter regular expressions in these fields, except for the **When** and **Duration** fields. Use a dash to enter a range of values in the **When** and **Duration** fields. For example, 4-5 for four to five seconds inclusive, -5 for up to five seconds, or 5- for five seconds or later.

To see annotations in the **Timeline** view, select the annotations that you are interested in, right-click, and select the appropriate option. The **Timeline** view opens with the cross-section marker covering the selected annotations. To see functions that are related to annotations, select the annotations in the **Log** view, right-click, then select **Select in Call Paths**. The **Call Paths** view opens with the related functions selected.

4.5 Annotate macros

Streamline provides various macros to mark up your code, including macros for color, channel information, or inserting bookmarks into the **Timeline** view and the **Log** view.

User-space annotation macros

Use the following macros to send annotations to `gator_d` from a user-space application. They are defined in `<install_directory>/sw/streamline/gator/annotate/streamline_annotate.h`.

ANNOTATE_SETUP

Set up annotations. Call this macro before calling any other annotate macros.

ANNOTATE(string)

Add a string annotation. Streamline adds the annotation to channel 0.

ANNOTATE_CHANNEL(channel, string)

Add a string annotation and assign it to a channel. The `channel` parameter is the channel number.

ANNOTATE_COLOR(color, string)

Add a string annotation with a display colour. See `streamline_annotate.h` for the defined colors.

ANNOTATE_CHANNEL_COLOR(channel, color, string)

Add a string annotation with a display colour, and assign it to a channel.

ANNOTATE_END()

Terminate the annotation in channel 0.

ANNOTATE_CHANNEL_END(channel)

Terminate the annotation in the given channel.

ANNOTATE_NAME_CHANNEL(channel, group, string)

Define a channel and attach it to an existing group. The channel number (`channel`) must be unique within the thread.

ANNOTATE_NAME_GROUP(group, string)

Define an annotation group. The group identifier (`group`) must be unique within the thread.

ANNOTATE_VISUAL(data, length, str)

Record an annotation that includes an image in one of the following formats:

- GIF
- PNG
- JPEG
- TIFF
- ICO
- BMP +RLE

Specify the image in `data`, the amount of data to write to the annotate file in `length`, and optionally a descriptive string to include with the image in `str`.

ANNOTATE_MARKER()

Add a bookmark to the **Timeline** view and the **Log** view.

ANNOTATE_MARKER_STR(string)

Add a bookmark with a string to the **Timeline** view and the **Log** view. Streamline displays the string when you hover over the bookmark in the **Timeline** view, and in the **Message** column in the **Log** view.

ANNOTATE_MARKER_COLOR(color)

Add a bookmark with a color to the **Timeline** view and the **Log** view.

ANNOTATE_MARKER_COLOR_STR(color, string)

Add a bookmark with a color and a string to the **Timeline** view and the **Log** view.

ANNOTATE_DELTA_COUNTER(id, title, name)

Define a custom delta counter. Specify:

id	An integer that uniquely identifies the counter.
title	A string for the chart title.
name	A string for the series name.

ANNOTATE_ABSOLUTE_COUNTER(id, title, name)

Define a custom absolute counter. The parameters are the same as for `ANNOTATE_DELTA_COUNTER`.

ANNOTATE_COUNTER_VALUE(id, value)

Emit an integer value for a custom delta or absolute counter. `id` identifies the counter.

ANNOTATE_DELTA_COUNTER_SCALE(id, title, name, modifier)

Define a custom fixed-point delta counter with a scaling modifier. Specify:

id	An integer that uniquely identifies the counter.
title	A string for the chart title.
name	A string for the series name.
modifier	An integer scaling factor.

Floating point counter samples reported through `ANNOTATE_COUNTER_VALUE_SCALE` are encoded as `(int)(sample * modifier)` for transport, and converted back to unscaled float for presentation. For example, to present a float with two decimal places of precision, set `modifier` to 100.

ANNOTATE_ABSOLUTE_COUNTER_SCALE(id, title, name, modifier)

Define a custom absolute counter with a scaling modifier. The parameters are the same as for `ANNOTATE_DELTA_COUNTER_SCALE`.

ANNOTATE_COUNTER_VALUE_SCALE(id, value, modifier)

Emit a value for a custom scaled delta or absolute counter. `id` identifies the counter and `value` specifies a float value. The scaling value `modifier` must be an integer with the same value as the modifier that was used when the scaled counter was declared.

CAM_TRACK(view_uid, track_uid, parent_track, name)

Add a track to a Custom Activity Map (CAM) view. To make the track a child of another track, specify the parent track ID in `parent_track`. To make it a root track, set `parent_track` to -1. `track_uid` must be unique within the view.

CAM_JOB(view_uid, job_uid, name, track, start_time, duration, color)

Add a job with a start time and duration to a CAM track. `job_uid` must be unique within the view. The job is displayed in the **Timeline** view as a colored bar representing its duration. The job name is displayed inside the bar. The start time and duration of the job are specified in nanoseconds. To get the current timestamp in nanoseconds, use `gator_get_time()`, which is declared in `streamline_annotate.h`.

CAM_JOB_DEP(view_uid, job_uid, name, track, start_time, duration, color, dependency)

The single-dependency version of `CAM_JOB_DEPS`.

CAM_JOB_DEPS(view_uid, job_uid, name, track, start_time, duration, color, dependency_count, dependencies)

Add a job with dependencies on other jobs to a CAM track. `track` is the track to add this job to, `dependencies` are the jobs that this job depends on, and `dependency_count` is the number of dependencies.

CAM_JOB_START(view_uid, job_uid, name, track, time, color)

Add a job with a start time to a CAM track. End the job using `CAM_JOB_STOP`.

CAM_JOB_SET_DEP(view_uid, job_uid, time, dependency)

The single-dependency version of `CAM_JOB_SET_DEPS`.

CAM_JOB_SET_DEPS(view_uid, job_uid, time, dependency_count, dependencies)

Set the dependencies of a job that was previously started using `CAM_JOB_START`. If you call this macro multiple times, Streamline uses the dependencies with the latest timestamp.

CAM_JOB_STOP(view_uid, job_uid, time)

End a CAM job that was previously started using `CAM_JOB_START`.

CAM_VIEW_NAME(view_uid, name)

Create a named CAM view. `view_uid` must be unique.

Kernel-space annotation macros

Use the following macros to send annotations to `gator`d from a kernel context, using Linux `ftrace` as the transport mechanism. They are defined in `<install_directory>/sw/streamline/gator/trace/event/gator_annotate.h`.

GATOR_KERNEL_WIDE_PID

For bookmark and counter `TRACE_EVENTS` only. To indicate that the event is associated with the whole kernel rather than a specific `tid`, you can pass this value as the `tid`. Annotation texts must have a `tid` of at least zero. `tid == 0` is the idle process.

GATOR_BOOKMARK(label)

Add a bookmark with a label.

GATOR_BOOKMARK_COLOR(color, label)

Add a bookmark with a color and a label.

GATOR_TEXT_START(tid, channel, label)

Start a text annotation.

GATOR_TEXT_START_COLOR(tid, color, channel, label)

Start a text annotation with a color.

GATOR_TEXT_STOP(tid, channel)

End a text annotation. This macro sends an empty string as the label, which Streamline recognizes as the end of a text annotation with the given channel and tid.

GATOR_DELTA_COUNTER_VALUE(title, name, units, value)

Output delta counter values.

GATOR_ABSOLUTE_COUNTER_VALUE(title, name, units, value)

Output absolute counter values.

Related information

[Annotations](#) on page 35

[Add string annotations to your code](#) on page 37

5 Analyze your capture

When you have captured a Streamline profile of your application, view the data in the **Timeline** view.

5.1 Timeline overview


The main view in Streamline is the **Timeline** view, which shows the captured performance data on a time-based visualization. You can zoom and scroll this view to show regions of interest.

A hierarchical data model stores the captured data at multiple resolutions that match the zoom levels available in the **Timeline** view. If you zoom in beyond the data capture sample frequency, the presented data is interpolated between the nearest data samples.

Hardware counters are sampled, by default, using a 1 millisecond sample rate. Other data sources are event-based and are captured and timestamped as they occur. Counter samples or event-based counters that occur less often than the 1 millisecond level must be mapped to the 1 millisecond data view for presentation. This mapping depends on the type of data being shown. See [Counter classes](#) for more details.


Cross-section marker

You can use the cross-section marker to highlight a time interval and show a summary of the counter statistics for the highlighted region. The summary shows a sum or an average per series, depending on the data source, see [Chart configuration series options](#).



Calipers

You can use the calipers to filter the data range that is provided to the other analysis visualizations, such as the **Call Paths** view. When you select a region of interest with the calipers in the **Timeline** view, the other views only consider that region.



Note

You can also filter by process or thread by clicking the processes or threads of interest in the **Timeline** view.

Activity and Counters charts

There are two main chart types that the **Timeline** view shows by default. By convention we call these 'Activity charts' and 'Counter charts', but you can customize the presentation and these names are not enforced.

Activity charts are generated from kernel or driver scheduling information. They show the percentage of time, at the current hierarchy level, that the indicated hardware unit was busy. Note that high utilization is not always the same as high load. A processor might run all the time, but at a low operating frequency due to low workload demand.

Counter-based charts are generated from actual hardware activity. For example, counting the number of instructions executed by a processor, or the number of cycles a block was active. These charts give a more direct measure of the workload, but do not indicate how long the processor was actually running for.

Combining both activity and counter-based charts allows a view of execution time and workload, both of which are needed for a complete performance analysis.

Expressions

To supplement the raw data that is captured from the hardware, you can use expressions to define custom data series. Expressions are simply equations, defined in a light-weight macro syntax, which define a derived data point to plot. For example, you can use expressions to scale a 'bus beats' counter into a 'bus bytes' counter, or to generate a hit ratio from cache hit and miss event counts.

Data points for expressions are re-evaluated for each zoom level in the timeline to ensure correct results for equations that are not distributive.

Templates

Templates store the current **Timeline** view layout - charts, series, and styling - as a reusable visualization that you can apply to other data captures. You can customize your **Timeline** view and save it as a template for future use, or for sharing with other team members.




Note

- Templates are designed to work with captures made with the current version of Streamline. Applying a template to a capture made with an older version of Streamline might result in some series being omitted. If a data series uses a source counter which has changed name since the capture was made, it is omitted.
- You can also use templates in counter configuration. Select a template to automatically collect data from all the sources that the template requires.

Built-in templates are included for most of the Mali™ GPU family. These templates do not include every Mali data source, but are the recommended starting point for Mali GPU profiling. For detailed descriptions of all the performance counters available for each Mali GPU, refer to the [Mali GPU Counter references](#) on the Arm Developer website.

5.2 Working with charts

Hover over a chart to see the values at that point on the timeline.

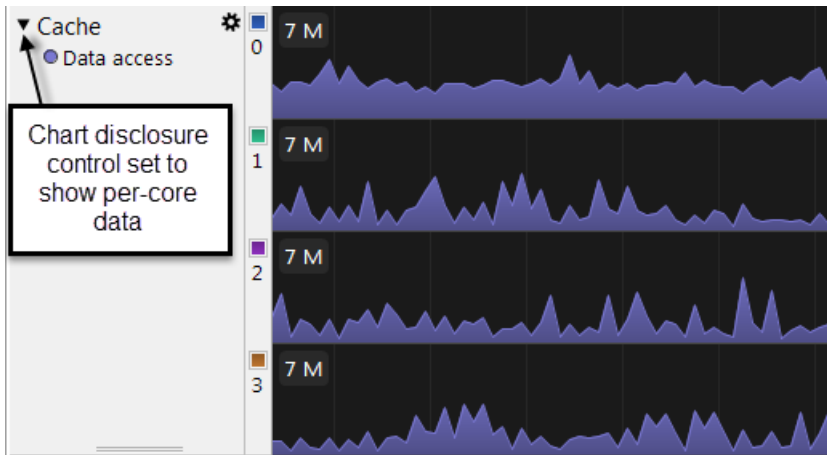


Increase the size of a chart to see a higher level of graphical detail, making it easier to see the variance in values. Resize a chart by clicking and dragging the handle control at the bottom of the box that displays the chart title.



You can reorder charts by clicking elsewhere within the box and dragging the chart to the new location.

If you have collected data from a target with multiple cores, the chart shows an aggregate of all cores. Expand the chart name to see the per-core data, and hover over a core number to see the name of the core.



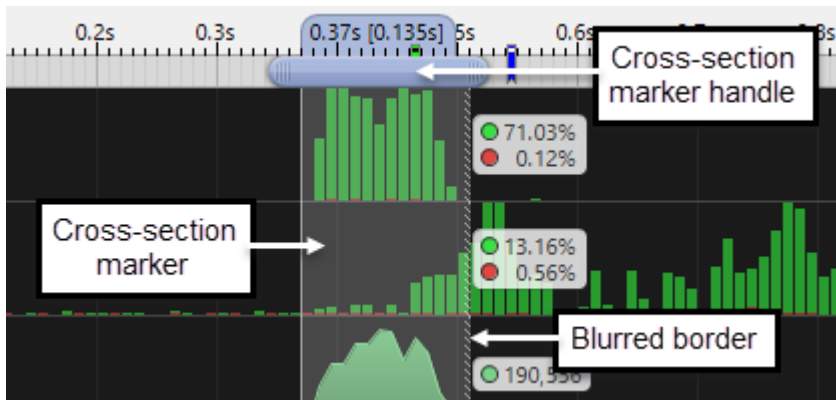
The charts scale according to the filter selected in the chart configuration series options. For example, if you select Average, the y-axis scales to the maximum average value for the time units in the selected range. Refer to [Customizing your charts](#) for instructions on how to change the configuration.

5.3 Analyze a region of time with the cross-section marker

To select a range of time to investigate more closely, click anywhere on the timeline and drag the handles on the cross-section marker. The information shown in the details panel when in **Processes** mode or **Samples** mode updates to show data for the window of time you have defined.

To move the view so the selected range is in the center, click **Center the display on the cross-section marker**  on the toolbar of the **Live** view or the **Timeline** view.

If you select a time range, then change the scale of the view so the cross-section marker cannot sit precisely, the border appears as a blurred line.



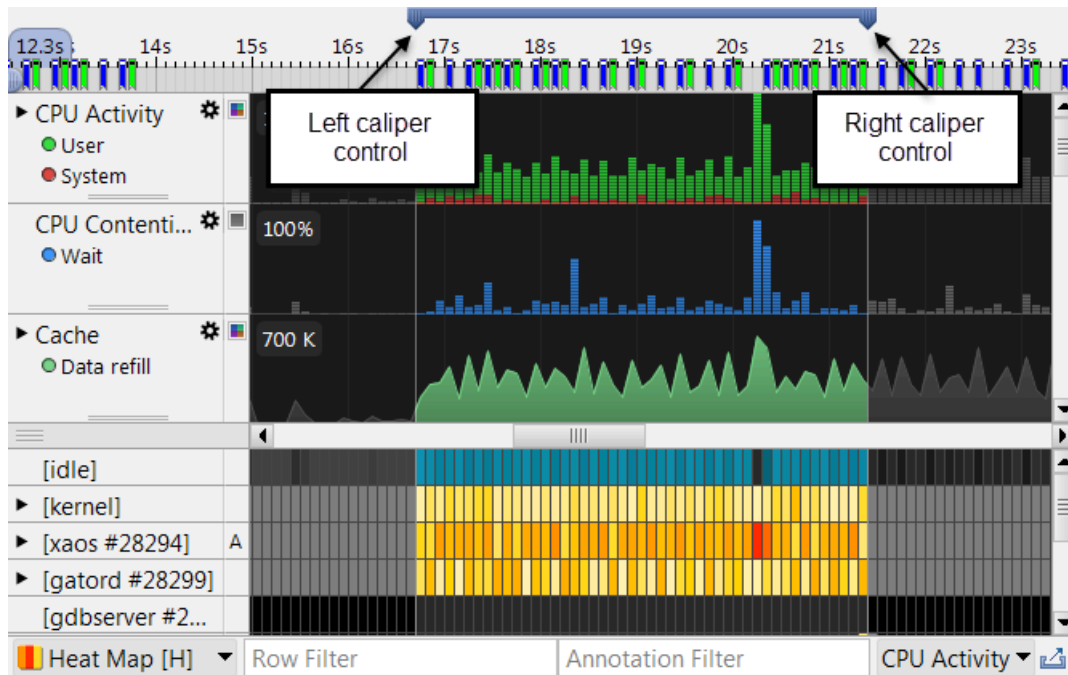
Unlike the filter controls, moving and expanding the cross-section marker does not affect the data in the other report views. To focus the data in the other views, use the calipers instead.

You can also control the cross-section marker with [contextual menu options](#) and [keyboard shortcuts](#):

- Set the cross-section marker, or one end of it, to a specific time index with the corresponding contextual menu option.
- Move the cross-section marker with the left and right arrow keys on your keyboard.
- Adjust the range that the cross-section marker covers by holding down Shift while pressing the arrow keys.

5.4 Analyze a region of time with the calipers

Use the calipers on the timeline to select a region of time. The **Call Paths**, **Functions**, and **Code** views update to show data for the time period set with the calipers.



Drag the calipers to set the required time region, or right-click on the timeline and select **Set Left Caliper** or **Set Right Caliper**.


To reset the calipers so they include all the data, click **Reset Calipers**  on the toolbar of the **Live** or **Timeline** view.

5.5 Customizing your charts

You can configure many aspects of the charts displayed in the **Live** and **Timeline** views, including the colors, titles, and data sets used by each series.

5.5.1 Chart configuration panel


Many of the charts in the **Live** and **Timeline** views have a button near the top right of the chart handle. This button opens and closes the chart configuration panel.



The chart configuration panel, in the following figure, contains the following sections:

- A toolbar section, which contains controls that apply to the chart as a whole.
- A series section, which contains controls that apply to the individual series in the chart.

For example, the toolbar section defines the chart title as **Bus**, and the chart type as Stacked. The series section defines the **Access (A15)** series.



Any updates that you make to the chart configuration in the **Live** view are preserved and displayed in the **Timeline** view.

5.5.2 Chart configuration toolbar options

The toolbar of the **Chart Configuration** panel defines options that apply to all series in a chart.

It has the following options:



Create a new series

Adds an empty series to the chart.



Create new wildcard series

Adds an empty wildcard series to the chart.

Chart Type

To choose one of the following chart types, use the Type buttons on the left side of the toolbar:

Filled

In a filled chart, each series is displayed as an area filled with the color that is specified in the chart configuration.




Line

In a line chart, each series is displayed as a colored line.




Bar

In a bar chart, each series is displayed as a colored bar. Each bar in the chart represents a time bin.



Heatmap

In a heatmap chart, each series is displayed as a heat map. Each counter is shown on a separate line and the bin color intensity increases in proportion to the number of samples.




Series composition

Use the drop-down menu to select one of the following options:

Stacked

In a stacked chart, the data for different series are stacked on top of each other. So, the highest point of a stacked chart is an aggregate of data from all the series that are contained in the chart. For example, if the first value of series A is three and the first value of series B is five, the first data point in the stacked chart that contains these series is eight.



Stacked charts are appropriate when the events are counted in exactly one of the series in a chart. For example, they are useful in a case where a chart contains both Data Read Hits and Data Read Misses.

Overlay

In an overlay chart, the different series overlap each other. Their position in the chart control determines the front-to-back ordering. To prevent series from obscuring data, as shown in the following figure, place series with larger values above series with lower values in the chart control. Drag and drop series controls using the chart configuration panel to reorder them.



Overlay charts are appropriate when some of the events are counted in more than one series in the same chart. For example, Data Read Requests and Data Read Hits.


Logarithmic

Shows all data points on a $\log_{10}(y)$ scale where y is each Y-axis data point in the chart. This Y-axis modifier is useful when there is a huge difference between the minimum and maximum values in a chart. A logarithmic chart displays with a series of horizontal lines, which increase in value by a power of ten.



Note

The charts cannot show negative numbers, so values less than one, which have a negative value in log space, appear as zero.



Average Selection

The cross-section marker overlay shows the average value of all bins that are included in the selection. If not selected, the overlay shows the total value of all bins in the selection.

Average Cores

Plot values in a multi-core chart as the average of all cores, unless you have used the multi-core disclosure control to show metrics per core. If not selected, the multi-core chart shows the total for all cores.

Percentage

Plot values as a percentage of the largest value in the chart.

The following energy offset buttons allow you to align energy data with the other data in the view.



If there is no energy data in the chart, these buttons are disabled. To capture energy data, you must have either an Energy Probe or a supported NI DAQ device.

Offset energy chart data to the left

Manually adjusts the energy data to the left.

Reset energy chart offset to zero

Returns the energy data to its original position.

Offset energy chart data to the right

Manually adjusts the energy data to the right.

Link adjustments to other charts

Links the adjustments made to the current energy chart to the other energy charts in the capture.

Title

Use this field to give the chart a title. The title appears at the top of the chart handle.

Remove Chart

Removes the chart from the view. If you have saved the chart in a chart configuration template, you can add it back to the view later using the **Switch and manage templates** button.

Related information

[Energy Probe overview](#) on page 96

5.5.3 Chart configuration series options

Each chart in the **Live** or **Timeline** view contains one or more series or sets of data. Each series has a set of options that enable you to customize the data that is used to plot the chart, and its look and feel.

You can use the handle on the left side of each series control to reorder the series in the chart. You can also use it to move or copy the series to another chart. To move the series, drag and drop it. To copy the series to another chart, hold down the Ctrl key while dragging it. When moving the series to another chart, make sure that the chart configuration panel is open for the destination chart.

Each series in a chart contains the following options:

Color

To change the color of a series, click the color box in its series control to open a **Color** dialog box.

Name

Enter a name for the series. This name appears next to the chart color in the chart key.

Description

Enter a description for the series. When you hover over the series title or color, a tooltip appears, containing the description defined here.

Expression

Use this field to define the data set that the series uses. Open a drop-down list of counters by pressing Ctrl+Space or the \$ symbol. To see a description of a counter, hover over it in the list. To add a counter to the **Expression** field, select it in the list. You can create an expression using more than one counter by using a combination of counter names and any of the following operators: >, <, >=, <=, ==, !=, ||, !, &&, %, *, /, +, -. You can use parentheses to define the order of operation.

In addition to the mathematical and comparative operators, you can use the following functions in the **Expression** field:

if

Evaluates whether a condition is true or false before applying an effect.

Usage: `if (x, y, z)`, where `x` is the expression to be analyzed. If `x` is nonzero, the result is `y`. If `x` is zero, the result is `z`. Only one of `y` or `z` is evaluated.

abs

Returns the absolute value of the numeric expression that is specified as the parameter.

Usage: `abs (x)`, where `x` is a numeric expression.

ceil

Returns the smallest integer that is greater than or equal to the numeric expression given as a parameter.

Usage: `ceil(x)`, where `x` is a numeric expression.

floor

Returns the largest integer that is less than or equal to a numeric expression given as a parameter.

Usage: `floor(x)`, where `x` is a numeric expression.

max

Compares the arguments and returns the greater value.

Usage: `max(x, y)`, where `x` and `y` are numeric expressions.

min

Compares the arguments and returns the lesser value.

Usage: `min(x, y)`, where `x` and `y` are numeric expressions.

round

Returns a numerical value that is rounded to an integer.

Usage: `round(x)`, where `x` is a numeric expression.

You can also use the following built-in variables:

\$ZOOM

This variable is the bin size of the current zoom level, in seconds, given as a floating-point value. For example, when showing the 200ms zoom level, this variable has the value 0.2.



When used independently from source data, entering constants in the **Expression** field can yield inconsistent results.

Filters drop-down menu

Depending on the class of the counter selected, you can select one of the following filters to apply to all the values in the series. To select the filter, use the drop-down menu in the series options panel:

Average

Works the same way as `Minimum`, except that it displays an average value for each time bin.

Accumulate

Displays the accumulated value of all the samples in the time bin.

Hertz

Converts the counter to a rate. It takes the value for each time bin and divides it by the unit of time that the time bin represents. It then converts the result up to seconds. You can use it to convert a cycles count into cycles per second.

Maximum

Works the same way as `Minimum`, except that it displays a maximum value for each time bin.

Minimum

Displays the minimum values for the counter for each time bin in the current zoom level of the **Timeline** view. So, if the current zoom level of the **Timeline** view is one second, `Minimum` displays the lowest value that was recorded for any millisecond within that second.

`Minimum` depends on the timeline resolution. If the counter provides that level of detail, `Minimum` displays the lowest value per unit of time according to the following table:

Resolution mode	Time bins
Normal	1 millisecond
High	100 microseconds
Ultra-High	1 microsecond

Unit

Enter the unit type for the series. The value that you enter in this field appears when you use the cross-section marker to select one or more bins.

Remove Series

To remove the current series from the chart, click **Remove Series** .

Related information

[Chart configuration panel](#) on page 50


[Counter classes](#) on page 116

[Chart configuration toolbar options](#) on page 51

5.5.4 Create a chart configuration template

A chart configuration template is a set of pre-configured charts that you can apply to the report that is displayed in the **Live** and **Timeline** views. A set of built-in templates is provided but you can also create your own.

Procedure

1. Configure the charts in the report as required. See [Chart configuration toolbar options](#) for information about the options available.
2. Click **Switch and manage templates** .
3. Select **Save as...** and enter a name for your template.

Results

The template is saved to the following location:

- On Windows, c:\Users\<username>\Documents\Streamline\Templates.
- On Linux and macOS, ~/Documents/Streamline/Templates.

Related information

[Toolbar options in the Live and Timeline views](#) on page 69

5.5.5 Apply a chart configuration template

A chart configuration template is a set of pre-configured charts that you can apply to the report that is displayed in the **Live** view and the **Timeline** view. A set of built-in templates is provided to accelerate analysis of Mali™ GPU rendering operations.


About this task

You can apply templates to existing reports. You can also apply templates to new captures using the **Counter Configuration** dialog box. This feature ensures that the counters that the charts in the template require are included in the capture.




atrace counters and visual annotations are not supported in templates, although applying a template does not remove a visual annotation chart from the display.

Procedure

1. Click **Switch and manage templates** .



To append the charts in the template to the currently displayed charts, rather than replacing them, hold down Shift when you click **Switch and manage templates** .


2. Select the template from the drop-down list.
To revert the report to its default chart configuration, select **<Default Template>**.

To add a template to the list, select **Install and Load...**

Results

Streamline displays the charts that the template specifies.

If the capture does not include a counter that a chart in the template requires, a warning icon is shown in the chart handle. To see which counter was missing, hover over the icon.



5.5.6 Match dynamic counter sources

When capturing a profile of your application, the application, or other sources such as the Graphics Analyzer layer driver, can generate software counters. Software counters sometimes require namespacing, for example reporting unique values per graphics context. To include counters using namespaced names in your templates, create a wildcard series to match a counter using a fixed prefix string and a variable post-fix.

Procedure

1. Create a chart configuration template, as described in [Create a chart configuration template](#). Click the **Create new wildcard series** button in your chart, then enter part of a counter name with a wildcard * for the new series to match with intercepted counters.

The supported wildcard format is \$<Prefix>*, where <Prefix> is the start of the counter name that you want to match.



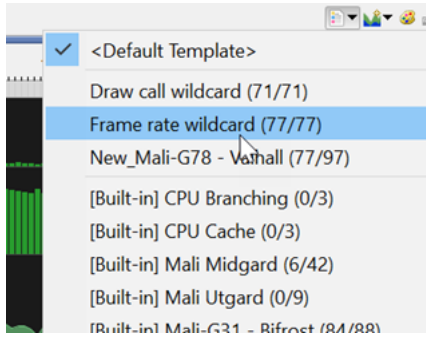
Note

The wildcard creates a `match_with` attribute in your template. The `match_with` attribute specifies the variable name pattern that is used to match variables by the template. For example, it can be used to match the `EGL context` suffix (a hex number) that is appended to variable names by the Graphics Analyzer layer driver.

This example shows how a wildcard (*) is used to specify all counters that start with `DrawCallsFrameEGLContext`.

```
<chart title="Draw calls" series_composition="stacked" rendering_type="filled"
  average_selection="yes">
<series name="Draw calls EGLContext ${*}" match_with="$DrawCallsFrameEGLContext*"
  \
  description="The number of draw calls." display="AVERAGE" color="246,138,51"
  units="draw calls"/>
</chart>
```

2. Apply your wildcard configuration template to your charts in **Live** view or **Timeline** view. Open the **Counter Configuration** dialog box, click **Switch and manage templates**, and select your template containing the wildcard name.

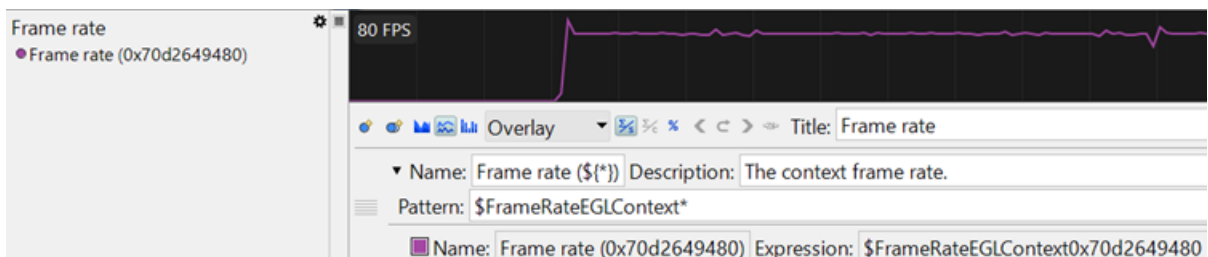


3. Expand your chart and see the matched chart series with their expressions.



Note

The expression for the chart series now shows the full matched variable name, instead of the wildcard pattern. You can change the color of the chart series to help distinguish between different series, but the color is not persistent. The Name and Expression cannot be edited.



5.6 Viewing application activity

The details panel of the **Timeline** view enables you to switch between different modes using the menu in its bottom left corner. Each mode displays a different set of data to supplement the charts.

5.6.1 Details panel modes

The details panel has the following modes:

Heat Map mode

The **Heat Map** shows a list of threads and processes that were active during the capture session in each time bin, alongside a color-coded heat map. Colors range from white to red, and the darker the color, the more activity caused by the thread or process in that bin.

Core Map mode

The **Core Map** shows a list of threads and processes that were active during the capture session and a color-coded activity map. The activity map shows which core was responsible for most of the activity for each thread or process. This mode is not available for single core hardware targets.

Cluster Map mode

Identical to **Core Map** mode, except that **Cluster Map** mode provides a color-coded activity map based on clusters. This mode is only available for targets that have multiple core clusters.

Samples mode

Samples mode lists the functions with samples in the currently selected cross-section. Jump to the relevant row in the **Functions** view by double-clicking on a function.

Processes mode

Processes mode provides a list of all processes alongside a process ID, the average percentage of CPU used and the maximum amount of memory used. Like **Samples** mode, the data shown depends on the current selection of the cross-section marker.

OpenCL mode

This mode displays the OpenCL commands being executed on each thread over the course of a capture session and shows dependencies between commands. This mode is only available on Mali™ devices with OpenCL timeline support compiled into the driver.

Images mode

This mode is only available if the capture contains visual annotations. It displays an enlarged version of the selected image in a visual annotation chart.

The map modes and **OpenCL** mode have a filter field. Filter the data in the details panel by entering a regular expression in the field. For example, the map modes show only the threads and processes whose name matches the expression. Regular expression strings are not case-sensitive.

Entries in the filter field in one of the map modes affect the other map modes only.

Related information


[Visual annotations](#) on page 38

[Keyboard shortcuts in the Live and Timeline views](#) on page 119

5.6.2 Heat Map mode

Heat Map mode in the **Timeline** view shows you a list of processes and threads that were active during the capture session. The entries are derived from process and thread trace data from the Linux kernel scheduler. Weighted colors reflect the number of samples in each process or thread.

Open **Heat Map** mode using the mode menu in the bottom left of the **Timeline** view.



Here is what each of the colored bins represent:

White or black, depending on the theme

The process is not running.

Light gray or dark gray, depending on the theme

The process has started, but is dormant. It could be sleeping, waiting on user input, or waiting for some other process to finish.

Yellow to red

The process is responsible for a percentage of total instructions during this bin. Red indicates a higher percentage.



Note

The [idle] process is color-coded differently to the other processes in the **Timeline** view. When the system is fully idle, it is bright blue. When it is partially idle it is a lighter shade of blue, and when the system is fully active, it is gray.


Blue dashes

CPU contention caused a delay. CPU contention can happen if there are too many processes and not enough cores to handle them.

If you select one or more processes or threads, the filterable chart series in the **Timeline** view update to show only activity caused by the selected processes and threads. Other chart series remain unchanged.

Each of the multi-threaded or annotated processes in the list have a disclosure control. Use the control to show each of the threads and annotations for that process. Annotations shown here can be hierarchical, with annotation groups each containing a set of channels, as defined by the macros inserted in your code.

Below the **Heat Map** are two filter fields. To filter the list of processes and threads, add a regex to the row filter. To filter the string annotations displayed in the **Heat Map**, add a regex to the annotation filter.



To export the **Heat Map** data to a text file, click **Export Heat Map Data to a text file**  to the right of the filter fields. This function exports all data that is displayed in the **Heat Map**, selected by any applicable filters, and within the calipers. The exported data is similar to the exported **Timeline**

view data. Each process or thread is a column and each time bin is a row. The **Heat Map** values are exported as percentages, or the following characters:

- Process not present. Equivalent to white or black time bins.
- I Process is idle.
- C Code contention. Equivalent to blue dashes in the time bins.



Annotations are not exported.

You can export the **Heat Map** for any activity source.

Related information

[Selecting the activity source](#) on page 64


5.6.3 Core Map and Cluster Map modes

Core Map and **Cluster Map** modes in the **Timeline** view map threads and processes to processor cores or clusters.

Core Map mode with five cores:



Cluster Map mode with two clusters:



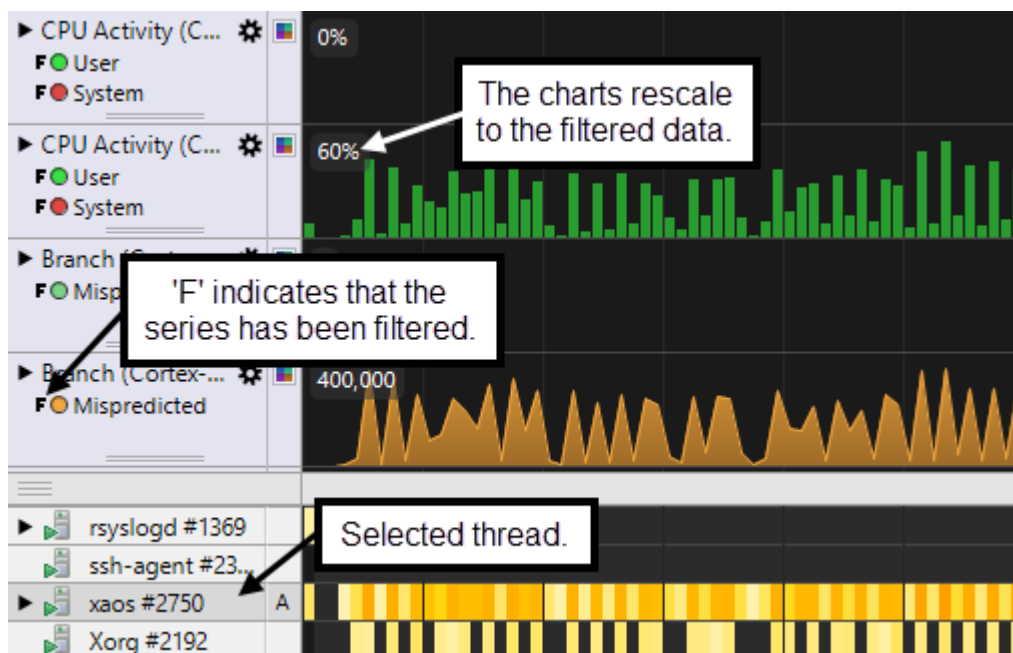


Captures using SMP systems support **Core Map** mode. Hardware targets with more than one cluster of cores support **Cluster Map** mode. These modes do not appear in the mode menu for captures that do not support them.

5.6.4 Filtering by threads or processes

Many chart series in the **Timeline** view can be filtered, based on the threads or processes that are selected in the **Heat Map**, **Core Map**, **Cluster Map**, or in the process list in **Processes** mode.

When there is an active selection, the affected chart series only show activity resulting from the selected processes or threads. The charts rescale to the limit of the filtered data. The letter **F** next to the series name indicates that the series is filtered.



Like the CPU Activity chart, the GPU Vertex and Fragment charts display only activity that the selected processes or threads initiate. You can then differentiate between GPU activity that your application causes and activity resulting from other applications or system services.

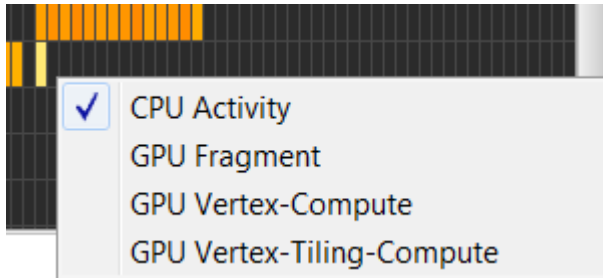
5.6.5 Selecting the activity source

You can select an activity source from the drop-down menu in the bottom right corner of the **Timeline** view as the focus of the **Heat Map**, **Core Map**, or **Cluster Map**.

By default, **CPU Activity** is selected.

If you select a different activity source to focus on, the **Heat Map**, **Core Map**, or **Cluster Map** in the details panel updates to show a map of threads and processes for the newly selected source.

For example, if you select **GPU Fragment**, the map updates to show activity for the GPU fragment processor only. The GPU activity sources are only available if your Arm® Mali™-based target is configured to support Mali GPU-specific profiling.



5.6.6 Samples mode

Samples mode in the **Timeline** view lists all functions in which one or more samples occurred in the time window that is covered by the cross-section marker.

Function	Samples ▾	Samples %
calculate	57	71.25%
mand_calc	10	12.50%
calccolumn_32	6	7.50%
calcline_32	5	6.25%
mand_peri	1	1.25%
_aeabi_uidivmod	1	1.25%

The details panel has the following columns when it is in **Samples** mode:

Function

The name of the function.

Samples

The number of samples that occurred in all instances of the function in the time window that is covered by the cross-section marker.

Samples %

The Samples figure as a percentage of all samples that were taken across the selected time window.

Select one or more rows in **Samples** mode and right-click to open a contextual menu, with the following options:

Select Process/Thread in Timeline

Switches to **Heat Map** mode and highlights the processes and threads for the selected functions in the Processes section.

Select in Call Paths

Opens the **Call Paths** view and highlights the function instances that relate to the selection.

Select in Functions

Opens the **Functions** view and highlights the selected functions.

Select in Code

Opens the **Code** view and highlights the source code for the selected functions.

Edit Source

Opens the source files for the selected functions in your default code editor.

Jump into the relevant row in the **Functions** view by double-clicking on a function.

5.6.7 Processes mode

Processes mode in the **Timeline** view provides a similar set of data to the output of the `top` command in your Linux shell. For every time bin in the **Timeline** view, it provides a list of processes along with usage information.

Index	Process ID	Process Name	% CPU		Memory Usage
0	27711	sshd	0.00%		4.07 MB
1	1	init.sysvinit	0.00%		932.00 KB
2	28139	gdbserver	0.00%		1.41 MB
3		kernel	0.41%		0 B
4	1483	rpcbind	0.00%		1.48 MB
5	28147	gator	3.01%		3.60 MB
6	28140	xaos	51.30%		17.55 MB

The details panel has the following columns when it is in **Processes** mode:

Index

The order in which rows were added to the table. This column is the default sort order.

Process ID

The PID that Linux assigned to the process.

Process Name

The name of the process.

% CPU


The average percentage of the CPU activity that this process used over the range that the cross-section marker covers.

You can sort by any of the columns in **Processes** mode. Click for a descending sort and click again to reverse the sort.

5.6.8 Images mode

Images mode in the **Timeline** view displays an enlarged version of the image that is selected in a visual annotation chart.

Images mode opens automatically if you click an image in a visual annotation chart. If you left-click and drag the mouse left or right within the chart, **Images** mode displays each selected image in turn, creating an animation effect.



5.6.9 OpenCL mode

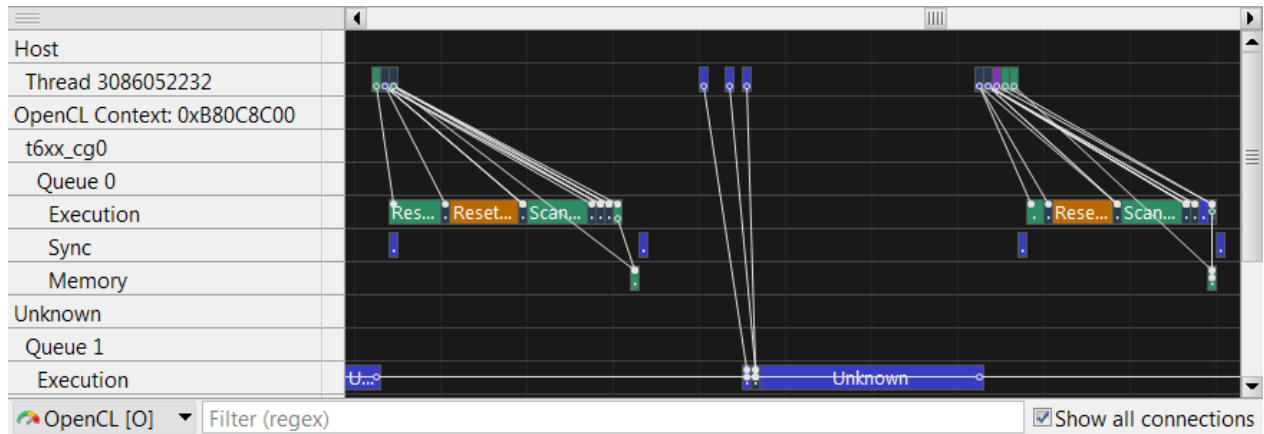
The Open Computing Language, or OpenCL, is a framework for parallel execution of jobs or kernels using task-based and data-based parallelism. **OpenCL** mode provides a visual representation of OpenCL code running on Mali devices.



OpenCL mode is only available on platforms with an OpenCL timeline compiled into the Mali driver. Contact your platform provider or [Arm Technical Support](#) for more details.

This mode shows which command is being run on each thread over the course of the capture session and provides mechanisms to explore command dependencies.

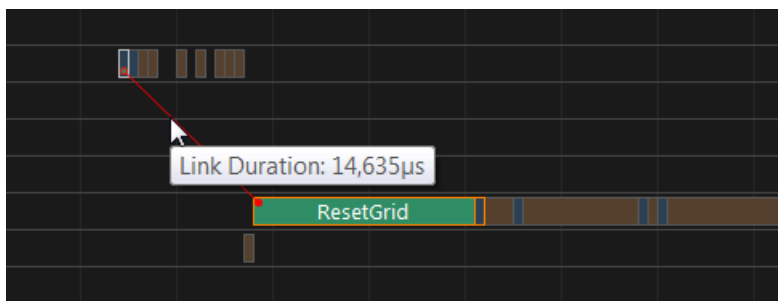
To enable **OpenCL** mode, you must create an instrumentation configuration file. For details of the options it must contain, see the documentation that is supplied with the Mali™ Driver Development Kit (DDK).




In OpenCL, commands are added to queues, then execute in parallel on one of the available hardware devices. Commands in a queue execute in series, but commands can also depend on the completion of commands in other queues.

Command names are shown inside colored areas, representing the duration of their execution. If there are two or more commands in a bin, and there is enough room, the number of commands in the bin appears in light gray.

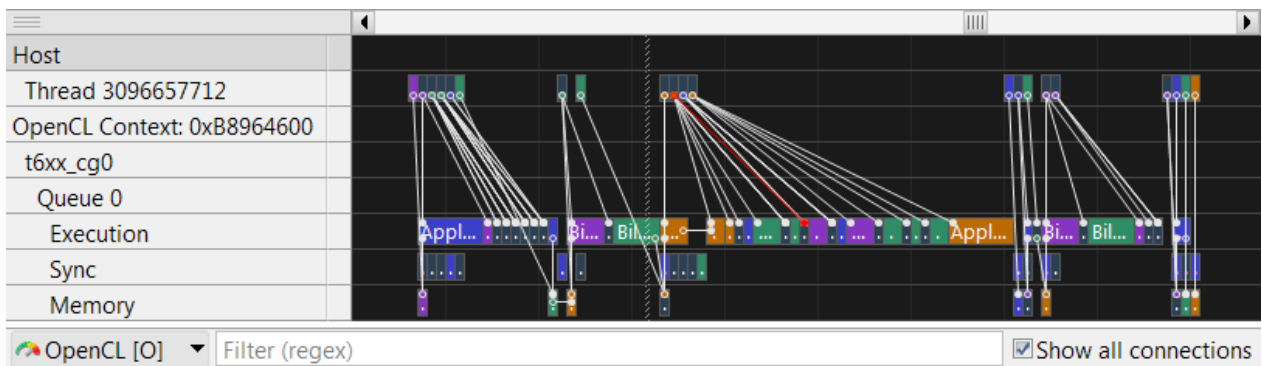
Dependencies between commands are shown using connecting lines. These have circles at each end to indicate the direction of the dependency. A command that is shown with a closed circle depends on a command that is shown with an open circle. Hover over the line connecting two commands to see a tooltip that shows the time delta between them.




Clicking a command gives it focus, hides any non-relevant information, and shows its dependency connections. The following figure shows a selected command that is highlighted with a yellow border. A yellow line shows where it entered the queue.



Zoom out to see more commands and the points at which they were enqueued. To display connections for all commands, click **Show all connections**.



Hovering over a command displays a tooltip that shows the command name, the time that it was initiated, and its duration.



There is a filter field at the bottom of the chart in **OpenCL** mode. Enter a regex in the field and **OpenCL** mode updates, showing only matching commands.





5.7 Toolbar options in the Live and Timeline views

Streamline provides easy ways to navigate and modify the **Live** and **Timeline** views using the toolbar.

The toolbar controls in the **Live** and **Timeline** views are:


Tags

If a special condition applies to the capture, one or more of the following tags appear on the left side of the toolbar:

-  The capture uses event-based sampling.
-  One or more warnings has occurred, as indicated by the number.

When you hover over a tag, a tooltip appears, giving more information about what the tag means.


Toggle bookmark markers

-  Activates or deactivates the bookmarks. If selected, vertical lines appear across the charts under each of your bookmarks.


Center the display on the cross-section marker

-  Centers the display on the position of the cross-section marker. This control is inactive if the cross-section marker is parked.


Reset Calipers


-  Sets the calipers so they include all the data in the capture. The calipers are the blue arrow controls at the top of the view. You can use them to limit the statistics in the reports to an area of interest.

Select Annotation log... (Timeline view only)

-  Opens the **Log** view and selects the closest Annotation-generated message to the current position of the cross-section marker. This option is not applicable if you did not include `ANNOTATION` statements in your code.

Zoom level (Timeline view only)

-  Cycles through the levels of zoom. To increase or decrease the unit of time that is used to represent one bin in the **Timeline** view, use the plus and minus buttons. If you zoom in beyond the sampling frequency, the **Timeline** view shows interpolated data. Alternatively, use the drop-down list to specify the zoom level.



Time index

Shows the time index of the current location of the mouse cursor.


Capture duration

The capture duration in seconds.


Go to beginning (Live view only)

 Moves the scrollbar to the beginning of the charts in the **Live** view and holds focus there.

Halt scrolling (Live view only)

 Causes the **Live** view to stop scrolling so that you can focus on the chart data currently in view. The charts of the **Live** view continue to expand as the capture session continues, but the **Live** view only moves if you move the horizontal scrollbar.

Go to live position (Live view only)

 Returns the **Live** view to the default scrolling behavior. The view scrolls with the growing charts and the live streaming data remains in focus.

Switch and manage templates


 Enables you to do the following:

- Select a chart configuration template to apply to the current display.
- Manage saved templates.
- Create and load templates.

See [Create a chart configuration template](#) for more information.

To append the charts in the selected template to the currently displayed charts, rather than replacing them, hold down the Shift key while clicking this button.

Add charts with default settings...


 Enables you to add a chart to the **Live** or **Timeline** view from the drop-down list. Either select one of the charts with default settings, or add an empty chart and configure it yourself.

The contents of the drop-down list depends on the data that has been collected during the capture session.

Cycle through themes

 Switches between the dark and light color schemes.

Export (Timeline view only)

 Opens the **Export** dialog box, enabling you to export the data for the current zoom level from the **Timeline** view to a text file.

Another export button below the details panel allows you to **Export Heat Map Data**. See [Heat Map mode](#) for further information.

Related information

[Annotations](#) on page 35

[Create a chart configuration template](#) on page 57

[Contextual menu options in the Live and Timeline views](#) on page 72

[Keyboard shortcuts in the Live and Timeline views](#) on page 119

5.8 Contextual menu options in the Live and Timeline views

Right-click anywhere in the charts or in the details panel of the **Timeline** view when in one of the map modes to open a contextual menu. This menu enables you to add bookmarks and to control the calipers and the cross-section marker.

The contextual menu options are:


Create Bookmark at...

Creates a bookmark at the time index of the current position of the mouse cursor. The bookmark is displayed in the **Live** or **Timeline** view rule as a colored mark.

Bookmarks enable you to label and quickly return to points of interest in the view.

Bookmarks that are added in the **Live** view are preserved in the **Timeline** view and the **Log** view. Bookmarks are also preserved if you re-analyze a capture, or if you export then import a capture.

Hovering over a bookmark displays an overlay that contains the timestamp and text string of the bookmark. To change the text or the color selector, click them.



To turn on or off vertical lines across the charts from the bookmarks, click **Toggle bookmark markers**  on the toolbar of the **Live** or **Timeline** view.

Set Left Caliper

Sets the left caliper control to the current position of the mouse cursor. Filters out all data before the left caliper from all views.

Set Right Caliper

Sets the right caliper control to the current position of the mouse cursor. Filters out all data after the right caliper from all views.

Reset Calipers

Resets the calipers to their default locations, selecting all data in the capture.

Set Cross Section Marker To...

Moves the cross-section marker to the specified time index, by default the current position of the mouse cursor.

Set Cross Section Marker Start

Sets the left-hand boundary of the cross-section marker. If it is parked, the cross-section marker is moved to this location. This option is not available if the selected location is to the right of the cross-section marker end.

Set Cross Section Marker End

Sets the right-hand boundary of the cross-section marker. If it is parked, the cross-section marker is moved to this location. This option is not available if the selected location is to the left of the cross-section marker start.

Park the Cross Section Marker

Resets the cross-section marker to its default, inactive position.

Related information

[Viewing application activity](#) on page 60

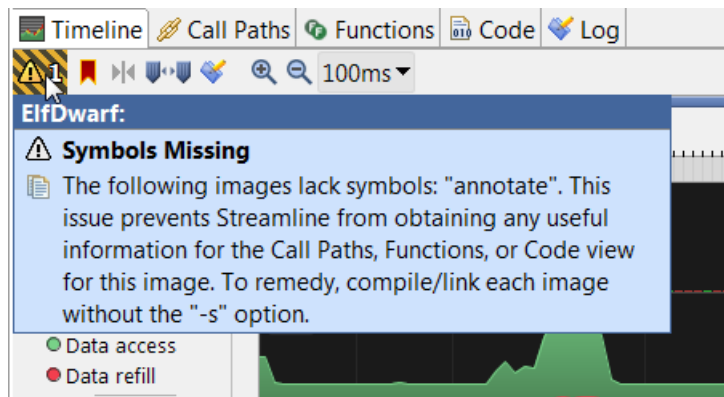
[Toolbar options in the Live and Timeline views](#) on page 69

[Keyboard shortcuts in the Live and Timeline views](#) on page 119

5.9 Warnings tag

If there is a problem with the reports, warnings appear in the **Live** and **Timeline** views. Clicking the tag displays a list of error messages, alerting you to the type and severity of problems.

The tag indicates the number of warnings that occurred. If there are no issues, the tag is not displayed.



The warnings are sorted by type and each is given a color-coded warning icon that tells you the severity of the issue:

Red

The issue is critical. For example, a connection timeout caused the termination of the capture, leaving you with incomplete data.

Yellow

The issue is medium severity. For example, power data from the Energy Probe could not be correlated with other data on the timeline.

White

The issue is minor. For example, a chart exists in the report but there is no counter data from the capture session to populate it.

5.10 Analyze your code

After you have analyzed your capture and determined that the CPU is causing performance issues in your application, you can analyze your code. Identify function-level hotspots in the **Call Paths** view and **Functions** view, then open the **Code** view to debug the source code and disassembly.

5.10.1 Analyzing Call Paths

The **Call Paths** view shows sample, size, and location data for each process, thread, and function called in your code. If you have used the caliper controls to filter data in the **Timeline** view, the data in the **Call Paths** view is limited to this selection.

Data is presented in a hierarchy of processes, threads, and functions below the one that called them.



If call stack unwinding is not available, for example because it is not selected in the dialog box, the sampled functions appear as a flat list for each thread.

To see more information about a called process, thread, or function, expand the hierarchy to show its subordinates. To show only the immediate subordinates, press Shift+Left arrow to hide all the subordinate call chain links, then expand again.

In the data table, the name of each process is enclosed in square brackets. The names of threads are enclosed in curly braces.

If the stack usage cannot be determined for one or more functions in the call path, the **Stack** value is a dark red color.

Right-click on any row in the table report to open a context menu where you can:

- Open the selected process, thread, or function in different views.
- Collapse rows that are not part of the current selection.
- Expand the hierarchy to expose hidden instances of the selected function.

When the source is set to **SPE**, the **Call Paths** view also has columns for the events that were captured. To add or remove columns, right-click on a column header, then select the column name. For columns that display a ratio, you can use this menu to show part of the ratio in a separate column.

Process	TLB	TLB refill Translation Table walk	Conditional Instruction	Branch Prediction
✓ [Process]/[Thread]/Code	755	639	191	
Architecturally retired	469	347	113	
Level 1 Data Cache	97	25	196	
Exception generated	76	84	167	
✓ Total latency	26	34	137	
✓ Translation latency	11	7	17	
Level 1 Data Cache Access	60		63	
TLB	25	14	73	
TLB refill Translation Table walk	38	12	26	
Conditional Instruction	14	3	37	
Branch Prediction			25	
Last Level Cache			11	
Multi-socket Data Access			10	
Stack	4		11	
✓ Location	10	4	17	
Show All Columns	4	14	19	
Reset Columns	7		8	
	2		7	
	8	4	12	
	1	1	7	

Related information

[Set capture options](#) on page 21

[Configure SPE counters](#) on page 32

[Troubleshooting report issues](#) on page 107

[Keyboard shortcuts in the table views](#) on page 119

5.10.2 Analyzing Functions

The **Functions** view lists all functions that were called during the capture session alongside sample, instance, and usage data. If you used the caliper controls to filter data in the **Timeline** view, the data in the **Functions** view reflects this selection.



The data shown in the **Functions** view depends on the text that is entered in the filter field, and the filtering selection in the **Timeline** view.

In the data table, the **Stack** value is a dark red color if the stack usage cannot be determined.

Filter the rows that are displayed in the view by entering a regular expression in this field. Only processes, threads, and functions whose name matches the regular expression are displayed.

Timeline Call Paths Functions Code Log					
Proc_[1-9]					
Function Name	Self	% Self	Instances	Stack	S
Proc_1		3.58%	6	48	
Proc_2		1.01%	15	16	
Proc_3		0.98%	6	32	
Proc_4	956	0.96%	12	16	
Proc_5	222	0.22%	3	16	
Proc_6	647	0.65%	12	32	
Proc_7	152	0.15%	6	16	

When the source is set to Periodic Sampling, EBS, or Python, the **Functions** view also shows the number of samples collected by all instances of the function, and the percentage of the total samples collected by all functions.



Samples are excluded if any functions that were called by the function collected them.

Right-click on any row in the table report to open a context menu where you can:

- Open the selected function in different views.
- View the top ten processes and threads that the selected function was called from.

When the source is set to **SPE**, the **Functions** view also has columns for the events that were captured. To add or remove columns, right-click on a column header, then select the column name as shown in [Analyzing Call Paths](#).

Related information

[Configure SPE counters](#) on page 32

[Keyboard shortcuts in the table views](#) on page 119

5.10.3 Sorting table reports

To reorder data in the table views and bring functions to the top of the table, perform a multi-level search.

Procedure

1. Select a column header to reorder data in that column. Select the column header again to reverse the sort order.
2. Hold down the shift key and select the next column header that you want to sort data by.
3. Repeat step 2 as necessary to sort your data by multiple levels.

Results

The number of dots in the lower right of the column headers indicates the position of the column in the sort hierarchy. For example, one dot means that the marked column is the primary sort column, while two dots indicate that it is the secondary sort column.

Function Name	Self	% Self	Instances	Stack	Size	
_raw_spin_unlock_irqre...	528	0.53%	17	0	32	vmlir
Func_2			15	32	256	dhry.
Proc_2			15	16	128	main
Proc_6			12	32	212	dhry.
Proc_8			12	16	364	dhry.
Proc_4	956	0.96%	12	16	108	main

5.10.4 Viewing your code

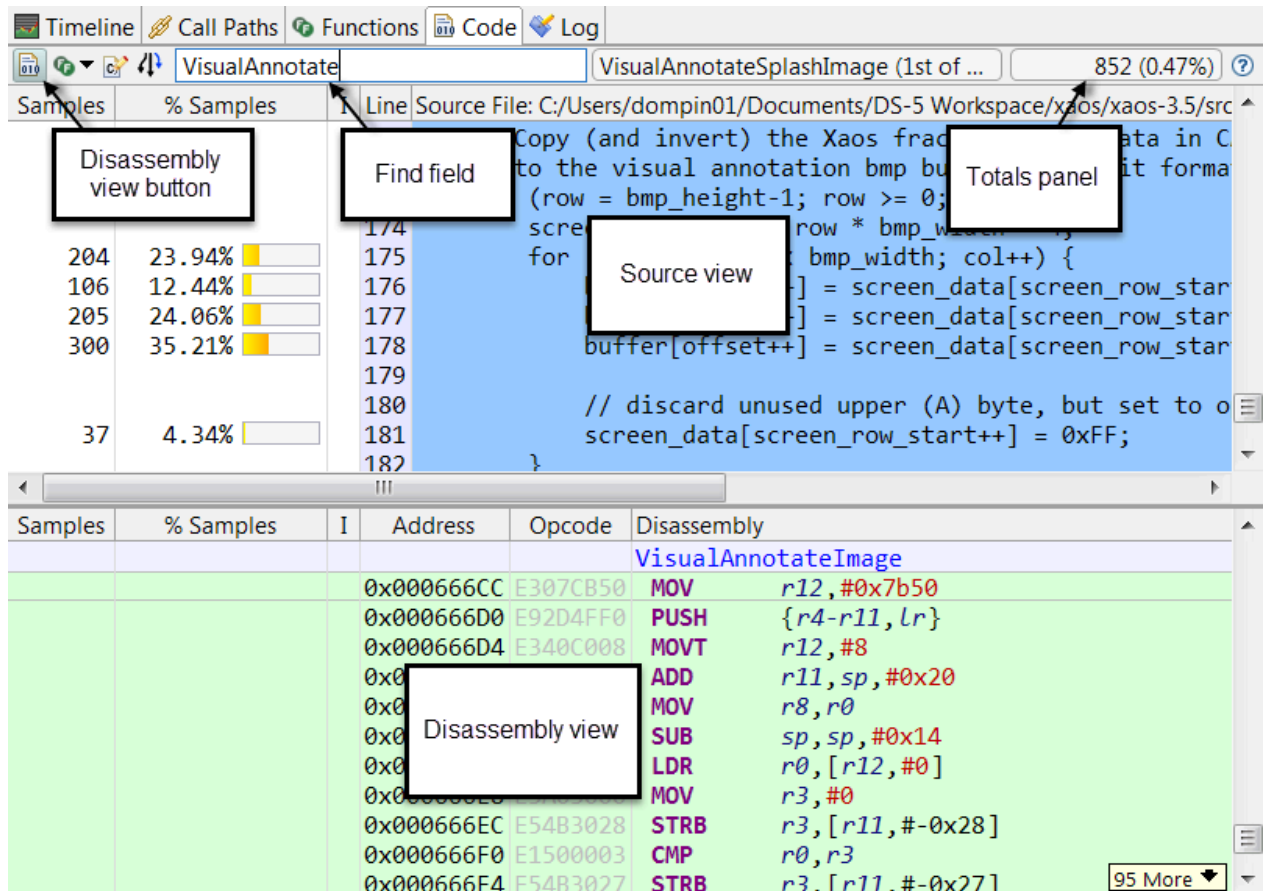
The **Code** view helps you to identify function-level hot spots, at source and disassembly levels. If you have used the caliper controls to filter data in the **Timeline** view, the sampling data in the **Code** view reflects this selection.

Before you begin

- Source code is only available for images in which the selected function is located, and if the image contains debug symbols. See [Generate debug information](#) for more information.
- In **Call Paths** view or **Functions** view, right-click on a function and choose **Select in code** from the context menu.
- If the source code is located in a different directory than when your application was compiled, you must redirect the file path so that Streamline can find the code. See [Troubleshooting missing source files](#) for more information.

About this task

When you open the **Code** view, your source code is shown with some color-coded statistics to help you identify functions that are causing performance issues.



Procedure

1. Select your source code to highlight related instructions in the disassembly panel.
 - Click the first row, then drag your mouse across a range of rows.
 - To select a single row of code, simply click it.
 - To select multiple rows of code, click the first row and press Ctrl to select more individual rows. Alternatively, click the first row, then press Shift and click the last row of the series.


To view the disassembly, click **Toggle**  to open the disassembly panel.

The instructions that are related to the code you selected in the **Code** view are shown in the disassembly panel.




- Selecting a single line of disassembly selects all the instructions that relate to that single line of source code. You can use function tags to highlight all instruction lines that are associated with a single function.
- Sometimes, the selected lines of source code or disassembly contain too many rows to fit in the current window. To see the hidden rows of code, select the **More** indicators.

2. Find and select a specific function to analyze it further:

- Click **Recently selected functions** , and choose a function from a list.
- Enter a string in the **Find** field located just below the toolbar, and press Enter to cycle through the matching functions.


3. If your capture contains PC sample sources, select a source from the first list, and then choose a counter to focus your analysis on that area.

4. To open the source file in your preferred editor and update your source code, click **Edit Source** .

After you have edited your source code, capture the profile again to compare the performance of your application. See [Starting a capture](#) for help with capturing a profile.

5.11 Streamline Data view

Use the **Streamline Data** view to manage your capture files and select the captures that you want to see or analyze.



You cannot open a report that was generated with:



Note

- A different toolkit license.
- A version of Streamline from more than two years ago. See [Getting started with Streamline](#) for compatible versions for your project.

You must first re-analyze the report.

5.11.1 Capture color-coding

Each of the captures in the **Streamline Data** view has a color-coded bar to its left. The color of the bar tells you the status of the capture and dictates what happens when you double-click on it.

The following colors can appear next to a capture:

Blue


Blue indicates a working capture that contains a report compatible with the current version of Streamline. To open the report for a blue capture, double-click on it.

Amber

Amber indicates a capture that needs to be re-analyzed, as its report is missing or incompatible with the current version of Streamline. To re-analyze the capture, click **Analyze** or double-click on the capture.

Red

Red indicates that the capture does not contain valid data and cannot be re-analyzed. It might indicate a capture error, in which case the error message is displayed in a tooltip.



5.11.2 Re-analyze stored capture data

Re-analyzing a capture replaces the existing report.

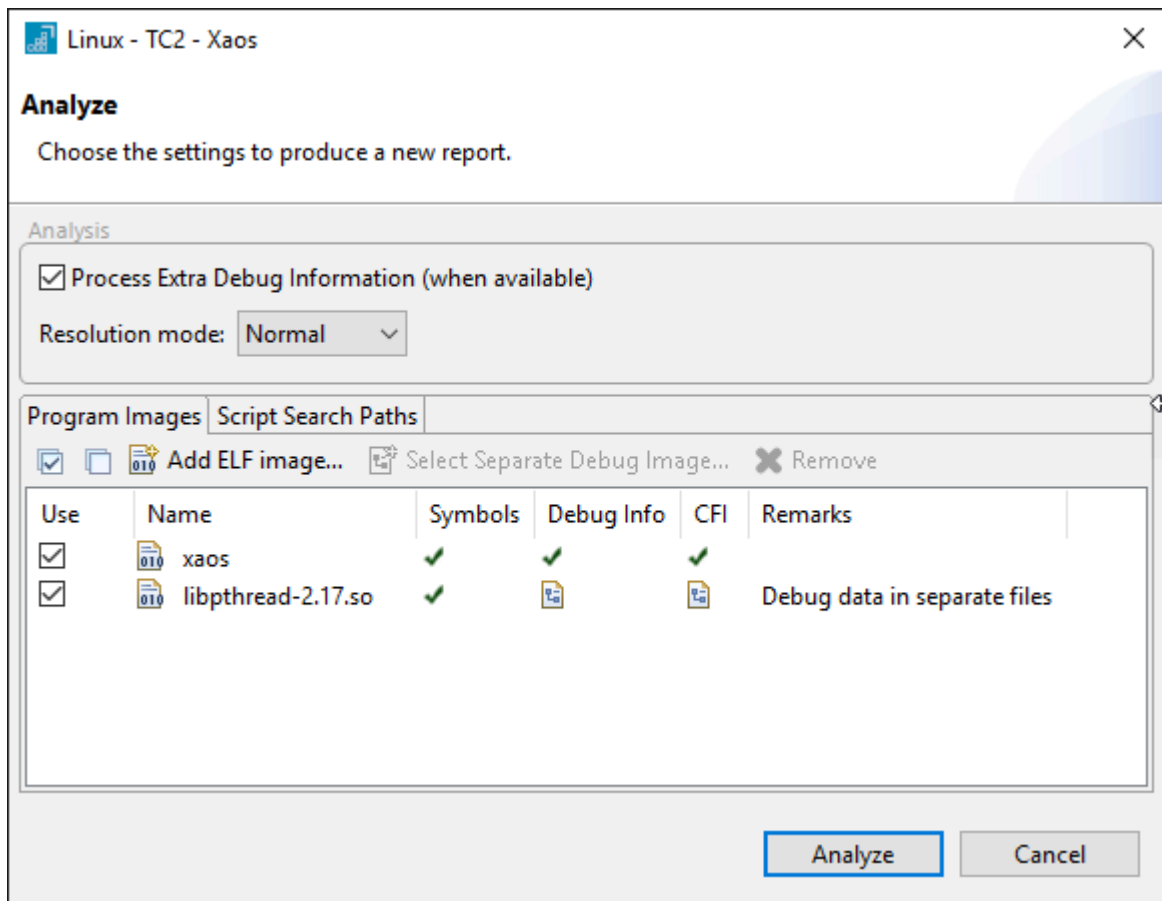
About this task

There are several reasons why you might want to re-analyze a capture, for example:

- To change your analysis options.
- To use sources that were not available when the data was captured.
- To use a new version of Streamline that is incompatible with the existing report.

Procedure

1. In the **Streamline Data** view, right-click on the capture then select **Analyze**. The **Analyze** dialog box opens.
2. Make the required changes to the settings.



The options for re-analysis are a subset of the options available in the **Capture and Analysis Options** dialog box and they work the same way.

3. Click **Analyze**.

Results

Streamline re-analyzes the capture data with the new settings and generates a new report, which replaces the existing report.

5.11.3 Duplicating a capture

You can duplicate an existing capture in the **Streamline Data** view using the contextual menu. For instance, you might want to re-analyze a capture with new options while preserving the original capture.

Procedure

1. Right-click on the capture and select **Duplicate**.
2. In the **New Name** dialog box, give the **Duplicate Streamline Document** a new name.

3. Click **OK**.

Results

A copy of the capture with the new name appears in the **Streamline Data** view.

5.11.4 Merge captures

Merge captures that were generated separately, so that you can see the global behavior of your system and compare multiple program runs.

Before you begin

Only perf and gator protocol captures are supported. Previously merged, and bare-metal captures, are not supported.

Procedure

1. In the **Streamline Data** view, select the captures that you want to merge.
2. Right-click the selected captures then select **Merge Capture Files...** from the popup menu. The **Merge captures** dialog box opens. Here you can see the start time and relative time delays of the captures, and change the merged file name and location.
3. In the **Merge captures** dialog box, you can edit the default settings:
 - If the clocks were not synchronized or set correctly when the capture was taken, adjust the relative timing offsets of the captures. Typically, set the start delay of the capture that started first as **Ons**, and enter the times when the other captures started.
 - To set the start delays back to the original values taken from the captures, click **Timestamp**.
 - In **Save location**, Streamline automatically creates a file name that is based on the names of the captures being merged. To rename the capture file, click **Browse**.
4. Click **Ok**.

Results

A merge job is started in the background. When it has finished, the merged capture is shown in the data browser and is ready for analysis.

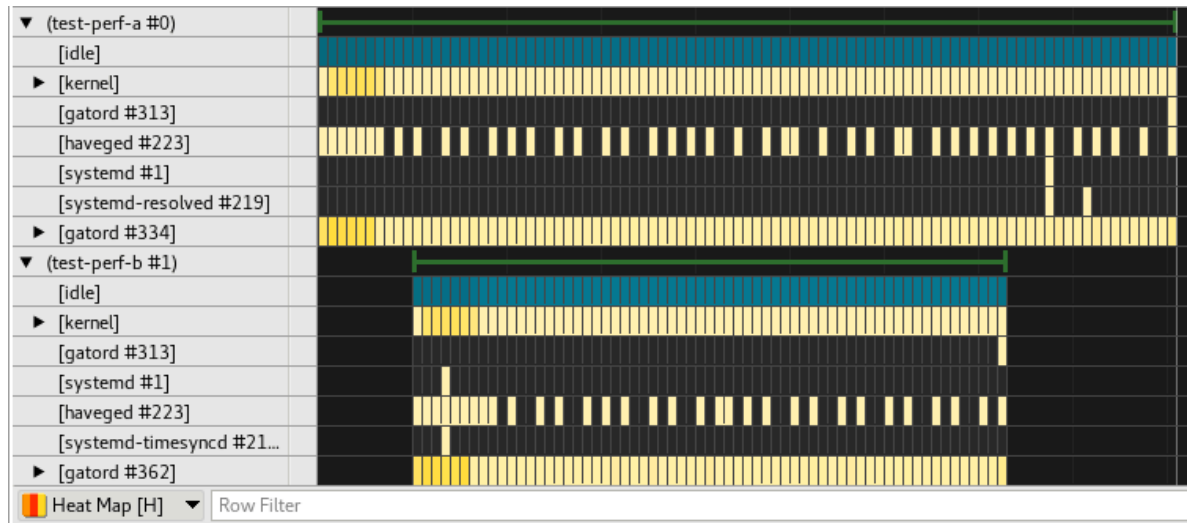
5.11.4.1 Merged captures

You can view merged captures in the **Timeline** view to allow comparison of multiple program runs.

Merged capture information is shown in the following modes and charts in the **Timeline** view:


Details panel

In the details panel, the **Heat Map**, **Core Map**, and **Cluster Map** modes show each merged capture and the relative start and end times. Each capture is shown in a row with the name of the original capture, and a green line indicating the time range of the capture. To view the processes and threads, click the disclosure control.



Charts

- Charts that are captured for each core can be viewed by clicking the disclosure control. The activity of each core in each capture is shown as a separate chart.
- System-wide counters can be disclosed to show individual source capture charts.



5.12 Preferences dialog box

Set your preferences for various features of Streamline.

Access the Preferences dialog box

Select **Window > Preferences > Streamline**.

Features

Name	Description
Timeline theme	Select a dark or light theme for the Timeline view.
Bare-Metal	
Import ITM traces as synchronized **streams	Select to make Streamline assume that ITM traces are synchronized from the start, and not discard data before the first synchronization packet in the trace.
Confirmations	

Name	Description
Prompts	Select to enable confirmation prompts when deleting captures or templates, or removing charts or chart series. All prompt options are selected by default.
Data Locations	
Lists the locations from which you can load captures. The first location in the list is the default save location. Reordering the list changes the order in which the captures appear in the documents view.	
The default location depends on your OS:	
<ul style="list-style-type: none">On Windows, it is C:\Users<user_name>\Documents\Streamline.On Linux and macOS, it is ~/Documents/Streamline.	
External Tools	
ADB Path	Enter the path to the adb executable.

6 Use Streamline from the command line

Use Streamline outside of the user interface to perform automated captures and generate text-based reports that you can import into a spreadsheet application.

6.1 Open a Streamline -enabled command prompt or shell

Open a command prompt or shell from which you can use the `streamline` command to access most of the functionality of Streamline.

Procedure

1. Add the appropriate location to your `PATH` environment variable:
 - For Arm® Mobile Studio, add `<install_directory>/streamline`.
 - For Arm Development Studio, add `<install_directory>/bin`.
2. Open a command prompt or shell.

Related information

[Streamline command-line options](#) on page 86

6.2 Streamline command-line options

The `streamline` command has different modes that enable you to use most features outside of the graphical user interface.

Use the `streamline` command in Streamline-enabled shells with the following syntax:

```
streamline <mode> [options] <files>
```

For more information about the different modes, see the following sections:

- [List modes](#)
- [Generate-config mode](#)
- [Capture mode](#)
- [Analyze mode](#)
- [Report mode](#)
- [Merge mode](#)
- [Import modes](#)

Options common to all modes

-h, -?, -help

Output help information to the console, listing each mode and option, and the required syntax.

-v, -version

Display the program version information.

Related information

[Open a Streamline -enabled command prompt or shell](#) on page 86

6.2.1 List modes

List the connected devices, or the debuggable packages on a connected device.

Syntax

```
streamline <mode> [options]
```

Modes

-list-devices

List the connected devices.

-list-packages

List the debuggable packages that are available on the connected device.

Options

Options for `-list-devices` :

-w, -wait <seconds>

Specify the maximum time that Streamline waits for responses from any devices.

Options for `-list-packages`:

-a, -all

List all packages that are installed on the device, including packages that are not debuggable.

-d, -device <device>

If multiple devices are connected, specify the ID of the required device. To obtain the ID, use the `adb devices` command.

6.2.2 Generate-config mode

Generate a `configuration.xml` file with the available events from a template.

Streamline searches for the template name or path in the built-in templates, then in the `Templates` directory, then as a file path (`*.st`). The target address is the IP address (optionally suffixed with `":<port>"`) or `"adb:<serial>"` of a running instance of `gator`.

Syntax

```
streamline -generate-config [options] <template_name_or_path> <target_address>
```

Options

-o, -output <filename>

The file to write the `configuration.xml` to. The default is to send the `configuration.xml` to the target.

Returns

A `configuration.xml` file.

6.2.3 Capture mode

Initiate a capture session.



Note

- For Android, Arm recommends using the `gator_me.py` script instead of this command. Refer to [Generate a headless capture](#) in the [Arm Streamline Target Setup Guide for Android](#) for more details.
- This mode requires a valid Arm® Development Studio or Arm Mobile Studio Professional license.

You must specify a valid `session.xml` file for this mode. This file defines your target hardware and the parameters of the capture session.

To create a `session.xml` file, enter your settings in the **Capture and Analysis Options** dialog box then click **Export...**

Syntax

```
streamline -capture [options] <session.xml>
```

Options

-c, -config <configuration.xml>

The path to the configuration file to send to the target.

-duration <seconds>

The number of seconds that you want the capture to last. The capture automatically terminates after that duration.

-e, -events <events.xml>

The path to your `events.xml` file.

-include-libs

Include the relevant libraries when retrieving images from the target.

-o, -output <filename>

The name of the APC capture file to create.

-package <package_name> [-activity <activity name>]

Use with Android targets to specify the package to capture. Use `-activity` to specify the main activity to be launched. If `-activity` is not specified, the main activity from the package is used. If there is no activity or multiple main activities in the package, the capture will not be started.

-password <string>

The password to use when retrieving images from the target.

-p, -pmus <pmus.xml>

The path to your `pmus.xml` file.

-retrieve-image <regex>

A regular expression to match the processes to retrieve from the target.

-template <name_or_path>

The template to apply. Streamline searches for the name or path in the built-in templates, then in the `Templates` directory, then as a file path (`*.st`).

-username <string>

The username to use when retrieving images from the target.

Returns

An APC capture file.

Example: 60 second capture session

Initiate a capture session that lasts for 60 seconds with the following command:

```
streamline -capture -duration 60 session.xml
```

6.2.4 Analyze mode

Analyze a capture.

Syntax

```
streamline -analyze [options] <capture.apc>
```

Options

-resolution <mode>

The resolution mode to use. Set <mode> to `normal`, `high`, or `ultra`.

6.2.5 Report mode

Output data from a capture to the command line or to a file. You can use options to filter and format the data.



This mode requires a valid Arm® Development Studio or Arm Mobile Studio Professional license.

Syntax

```
streamline -report [options] <capture.apc>
```

Options

-all

Output the contents of the **Timeline**, **Call Paths**, **Functions**, and **Log** views, including bookmarks and custom activity map tables. This option is the default.

-bookmarks

Output all bookmarks that are stored in a capture, listing the time index location of the bookmark and its text.

-bstart <name>

Filter the data to start at the first bookmark with the provided name. For example, if you enter `-bstart redflag`, all data before the first instance of the bookmark title `redflag` is filtered from the output.

-bstop <name>

Filter the data to end at the first bookmark with the provided name.

-callpath

Output the contents of the **Call Paths** view. Subordinate functions are indented.

-cam

Output the contents of custom activity map tables.



To export custom activity map data, the report must contain custom activity map annotations.

-disasm[=<name>] <images> [<source>]

Output the disassembly view part of the **Code** view as a CSV file.

<images> is a comma-separated list of ELF images (as were attached to the capture for analysis) to include in the output.

<source>, if specified, is the name for the data to output that is given in the **Code** view. For example, 'Periodic Sampling' or 'SPE'.

-format <space|tab|csv>

The format of the output. Format the tables using spaces, tabs, or comma-separated values. This option is useful if you want to easily convert output text files to a spreadsheet application. The default format is spaces, making the tables easy to read when printed on the command line.

-function

Output the contents of the **Functions** view.

-heatmap "<source>"

Export the contents of the **Heat Map** to a text file. <source> must match the name of the activity source for the **Heat Map** in the capture. Default = "CPU Activity".

-log

Output the contents of the **Log** view.

-o, -output <output_directory>

Send all report text to files in the specified directory. If you specify multiple tables for export, each table is written to <output_directory> in a new file with an appropriate name. If you specify multiple captures, each capture is output to a subdirectory of <output_directory> with a name matching the capture.

-opengl

If there is an OpenCL table, output the contents.

-per_core

Output per-core data when used with the option.

-process <regex>#<PID>

A regex matching the process to filter the timeline data for, and its PID, separated by a hash.

-scale <number>

Set the number of bins per second to use in the report. Default = 1,000.

-start <seconds>

Filter output data to start at the specified time in the timeline. For example, if you enter -start 0.005, all data before the 5 millisecond mark is not included in the output.

-stop <seconds>

Filter output data to stop at the specified time in the timeline.

-template <filename>

Use with a valid chart configuration template to list data for your customized charts on the command line. To create a chart configuration template, see [Create a chart configuration template](#).

-timeline

Output the contents of the **Timeline** view.

-warnings[=<name>]

Produce the warnings table. <name> specifies the filename to use for the output. Default = `warnings.csv`.

Returns

A capture analysis report.

Example: Output multiple report types.

To output the **Call Paths** and **Functions** data from the capture `thread_001.apc`, enter:

```
streamline -report -callpath -function thread_001.apc
```

Example: Output timeline data to a text file.

It can be useful to output very large reports to a file, then import the data into a spreadsheet application.

To output the timeline data from `capture_001.apc` to a file called `output.txt`, enter:

```
streamline -report -timeline capture_001.apc -o output.txt
```

If `output.txt` is an existing file, it is overwritten with the new data.

Time range 0.000 to 15.552 seconds.

Timeline Counter Values:

Time Index	CPU Activity:User	CPU Activity:System	Branch:Misspredicted	Bus:Access	Cache:L2 data access
0.000	1.45%	0.00%	0	0	0
0.001	12.74%	0.00%	2,136	0	4,441
0.002	0.00%	0.00%	0	0	0
0.003	4.75%	0.00%	1,315	0	3,665
0.004	7.23%	0.00%	2,266	0	5,770
0.005	0.00%	0.00%	0	0	0
0.006	14.84%	0.00%	5,288	0	11,337
0.007	18.08%	1.91%	2,211	0	9,420
0.008	20.00%	0.00%	487	0	4,635
0.009	20.00%	0.00%	900	0	5,569
0.010	20.00%	0.00%	488	0	4,635
0.011	20.00%	0.00%	488	0	4,635
0.012	20.00%	0.00%	488	0	4,635
0.013	20.00%	0.00%	488	0	4,635
0.014	20.00%	0.00%	488	0	4,635
0.015	20.00%	0.00%	488	0	4,635
0.016	20.00%	0.00%	488	0	4,635
0.017	17.87%	2.12%	738	0	4,895
0.018	15.95%	4.04%	3,016	0	8,216
0.019	16.66%	3.34%	2,119	0	5,976

6.2.6 Merge mode

Merge multiple APC capture files.

Syntax

```
streamline -merge [options] <files_to_merge>
```

Options

-o, -output <filename>

The name of the APC capture file to create.

-synchronize <timestamp|none|(list)>

Synchronize the captures using the `CLOCK_REALTIME` timestamp, nothing, or a given list of comma-separated delay times in nanoseconds. Default = timestamp.

Returns

A single APC file.

6.2.7 Import modes

Import a bare-metal trace file and create a capture from it.

Syntax

```
streamline <mode> [options] <files_to_import>
```

Modes

-import-apc-zip

Import zipped APC files.

-import-perf

Import perf data files. This mode requires a valid Arm® Development Studio or Arm Mobile Studio Professional license.

-import-barman-raw

Import a bare-metal raw file. This mode requires a valid Arm Development Studio license.

-import-stm

Import a bare-metal trace that was captured through STM. Streamline converts the trace file into a bare-metal raw file. This mode requires a valid Arm Development Studio license.

-import-itm

Import a bare-metal trace that was captured through ITM. Streamline converts the trace file into a bare-metal raw file. This mode requires a valid Arm Development Studio license.

-import-etm-dt

Import a bare-metal trace that was captured through ETM. Streamline converts the trace file into a bare-metal raw file. This mode requires a valid Arm Development Studio license.



Embedded Trace Macrocell (ETM) capture is a deprecated feature, and is to be removed in Streamline version 8.4.

Options

-o, -output <output_directory>

The directory where the capture is created.

-agent-image <elf_image>

The path to the ELF image of the bare-metal code that contains the bare-metal agent code. Required for ETM data trace only.

-barman-xml <barman.xml>

The path to the `barman.xml` file that the Barman Generator Wizard generated. Required when importing STM and ITM traces.

-assume-sync

Assume that the ITM trace is synchronized from the start.

Returns

An APC capture file.

Related information

[Using the bare-metal generation mechanism from the command line](#)

7 Capturing Energy Data

Describes how to set up and use the Arm Energy Probe with Streamline to view the power metrics of code running on target hardware.


7.1 Energy Probe overview

To capture energy data, you must have either an Arm Energy Probe or a supported NI DAQ device.

The Energy Probe is a low impact, inexpensive solution to give application and system software developers quick feedback on the impact of their code on the system energy footprint. It is intended to provide a better understanding of the static and dynamic behavior of the target system for the purposes of debugging, profiling, and analysis. It is not intended to be a high-precision data acquisition instrument for power benchmarking.

The Energy Probe has three power connectors, each of which is designed to connect to a 2-pin header for measuring the power. Having three power connectors allows the energy probe to provide three independent power, current, or voltage measurements.

Energy Probe schematic




In addition to the three power connectors, the Energy Probe has a single pin GND connector that must be connected to ground on your target board. It provides a ground connection for the Energy Probe.

Each of the three power connectors measures:

- The current flow through a shunt resistor of a known value on your target system
- The voltage at the positive terminals of the Energy Probe.

Energy Probe electrical connection example



The Energy Probe has flying leads and must be carefully connected to your target. Do not plug the green ground wire into anything but the target ground. Do not plug it into the I/O pins or power pins. Doing so could damage your Energy Probe or the power supply of your target.

Related information

[Energy Probe requirements](#) on page 97

[Setting up Energy Probe](#) on page 98

[Adding the caiman application](#) on page 100

[Energy Probe data](#) on page 101

7.2 Energy Probe requirements

There are several hardware and software requirements when using the Energy Probe.

The Energy Probe has the following requirements:

- An installation of Arm® Development Studio.
- A valid Arm Development Studio license.
- A target running Linux kernel version 3.4 or later, configured for Streamline profiling.
- An Energy Probe unit.
- USB drivers for the Energy Probe.
- A target that has 2-pin IDC 0.1" power measurement headers. The target also requires a shunt resistor with a supply voltage less than 15V and rated at least 0.5W. The shunt resistor needs a 1-pin IDC ground terminal and must not drop more than 165mV.

7.3 Shunt resistor selection for Energy Probe

With 20x amplification, Energy Probe requires the correct selection of a shunt resistor to provide the best possible dynamic range in power measurement, while avoiding saturation of the input of Energy Probe.

A shunt resistor with a value that is too low reduces measurement dynamic range, resulting in less resolution in the power data. A shunt resistor with a value that is too high can cause the input stage of the Energy Probe to saturate, which causes a flat line in the charts that are related to Energy Probe in the **Timeline** view.

To avoid input saturation, the drop across the shunt resistor must never be more than 165mV. You can also use the following equation to determine whether your shunt resistor is appropriate:

$$R_{\text{Shunt}}(\text{max}) = 165 \times V_{\text{supply}} / (1000 \times \text{Power})$$

V_{supply} is the input voltage. $\text{Power}(\text{max})$ is the maximum power that the Energy Probe measures. $R_{\text{Shunt}}(\text{max})$ is the maximum value of the shunt resistor. A shunt resistor value that is greater than $R_{\text{Shunt}}(\text{max})$ might cause input saturation.

This equation provides the absolute maximum value for the shunt resistor. Use a value that is at least five percent lower than this maximum to allow for component tolerances.

When connecting the Energy Probe, consider the following:



- The black and white probe closest to the green wire is Channel 0.
 - For best results, attach Channel 0 to the power source which best represents the CPU load. Streamline aligns the power data with the software activity by maximizing the correlation of Channel 0 with the CPU load.
 - The probe white wire is V+. The black wire is V-.
-

Examples

- 5V power supply, 8W(max), $R_{\text{Shunt}}(\text{max}) = 100$ milliohms.
- 1V core voltage, 2.5W(max), $R_{\text{Shunt}}(\text{max}) = 50$ milliohms.
- 1.5V core voltage, 0.4W(max), $R_{\text{Shunt}}(\text{max}) = 500$ milliohms.

7.4 Setting up Energy Probe

It is critical that you set up your Energy Probe correctly. Failure to do so can result in missing power data in Streamline, and even damage to your target.

Procedure

1. Connect your target to your host through a USB port.

2. If necessary, install drivers using the standard Windows driver installation.
When first plugging in the Energy Probe on Windows, the LED flashes RED at approximately one second intervals. These red flashes indicate that the USB enumeration has failed and that you must install the driver.
3. If you are prompted to locate the driver file, select the `ARM_EnergyProbe.inf` configuration file, which is located in `<DS_install_directory>/sw/streamline/energy_meter/energy_probe/`. When you have installed the drivers, and Energy Probe is operating correctly, the LED remains green.
4. To set it up on Linux Ubuntu, enter the following command:

```
sudo apt-get install libudev-dev
```

5. On Linux Ubuntu, you must do either of the following for Energy Probe to successfully use the device:
 - Change the permissions of the TTY device by entering the following command:

```
chmod 777 <tty_location>
```

- Add yourself to the same group as the TTY device by entering the following command:

```
usermod -a -G <Group_Name User_Name>
```

`dialout` is the usual name of the group for the TTY device.

If you attach a probe to the target with the channel wire connector the wrong way around, the LED turns red and blinks. Connecting the channel wire the wrong way round does not damage your target or the Energy Probe but it does not provide useful data. Disconnect all Energy Probes and wait until the LED changes to green. Then reattach them one by one and make sure that you reverse any probe that caused the LED to change to red. When you have attached the Energy Probe correctly, the LED remains green.




Note

If there is low consumption, it is possible that the LED remains red even if the connection is correct. For example, the LED might remain red if the target device is idle.

6. If your target has headers available for power or energy metering, attach the ground wire, then attach the probes.
7. Attach the probes.
Make sure that the probe polarity is correct by observing whether the LED changes from green to red. The probe white wire is V+. The black wire is V-.


Results





7.5 Adding the caiman application

For Streamline to communicate with the Energy Probe, you must give the location of the `caiman` application using the **Tool Path** field in the **Capture and Analysis Options** dialog box.

Procedure

1. Select the **Arm Energy Probe** option from the **Energy Capture** drop-down menu.
2. Click **Select the energy capture tool**  in the **Tool Path** field.
3. Locate the `caiman` application in your Streamline installation directory, select it, and click **Open**.



Energy Capture						
ARM Energy Probe		Device:	Auto-Detect	Port:	8081	
Tool Path: C:\Program Files\DS-5 v5.24\sw\streamline\bin\win-64\caiman.exe 						
Channel 0:	<input checked="" type="checkbox"/> Energy	<input checked="" type="checkbox"/> Power	<input type="checkbox"/> Voltage	<input checked="" type="checkbox"/> Current	Resistance:	20 mΩ
Channel 1:	<input type="checkbox"/> Energy	<input type="checkbox"/> Power	<input type="checkbox"/> Voltage	<input type="checkbox"/> Current	Resistance:	20 mΩ
Channel 2:	<input type="checkbox"/> Energy	<input type="checkbox"/> Power	<input type="checkbox"/> Voltage	<input type="checkbox"/> Current	Resistance:	20 mΩ



Note

- It might be necessary to specify the TTY device when using the Arm Energy Probe on Linux. To do so, enter the following in the **Device** field: `/dev/ttyACM0`. When using NI DAQ, the device name is usually `dev1`.
- On Windows, you can leave the **Device** field empty. Streamline auto-detects your energy probe.

7.6 Updating your firmware

Necessary updates occasionally become available for the Energy Probe firmware.

Procedure

1. Detach the Energy Probe probes and ground from the target.
2. Plug the Energy Probe into a USB port on your Windows host machine.
Because of a limitation with the USB bootloader from the MCU vendor, updating firmware on a Linux host is not supported.
3. To enable firmware programming mode, insert a paperclip into the small hole on the top near the LED.
4. Press and hold down the button for three seconds.
5. Wait for the LED to turn red, and then off. Firmware programming mode is now active.
6. When the drive folder opens, delete the `firmware.bin` file. This file is a placeholder.
7. Drag and drop the new `emeter_firmware.bin` file onto the drive folder. This file is located in `<DS_install_directory>/sw/streamline/energy_meter/energy_probe/`.
8. Safely remove the Energy Probe drive from the host machine using the Windows **Safely Remove Hardware** icon.
9. Unplug the Energy Probe from the host machine.
10. Wait for one second and plug it back in.

Related information

[Setting up Energy Probe](#) on page 98


7.7 Energy Probe data

After you have run a capture session with the Energy Probe, you can see the Energy Probe data aligned to the software activity in the **Timeline** view.

If your target meets the requirements, and you have attached a power probe point that closely resembles the activity of the CPU cores using Channel 0, the Streamline correlation algorithm aligns the data for you. Notes that appear when you click the warnings tag on the left side of the toolbar communicate any limitations or failures in this area.

When you have set up everything correctly, run a standard report capture and analysis. A large variance between idle and activity ensures that the correlation algorithm has something to lock

onto, so set your workloads accordingly. If you have connected and configured your Energy Probe correctly, a power chart appears in the **Timeline** view.



7.8 Set up National Instrument Multifunction Data Acquisition devices (NI DAQ) to capture energy data

As an alternative to the Energy Probe, you can use a NI DAQ device to gather energy data from your target.

Before you begin

- Install NI-DAQmx software on your host machine, so `caiman` can communicate with the NI DAQ device. This software package includes drivers for the NI DAQ device. You must use the NI-DAQmx Base software on a Linux host machine.
- Set the location of the appropriate `caiman` application using the **Tool Path** field in the **Capture and Analysis Options** dialog box.



Note

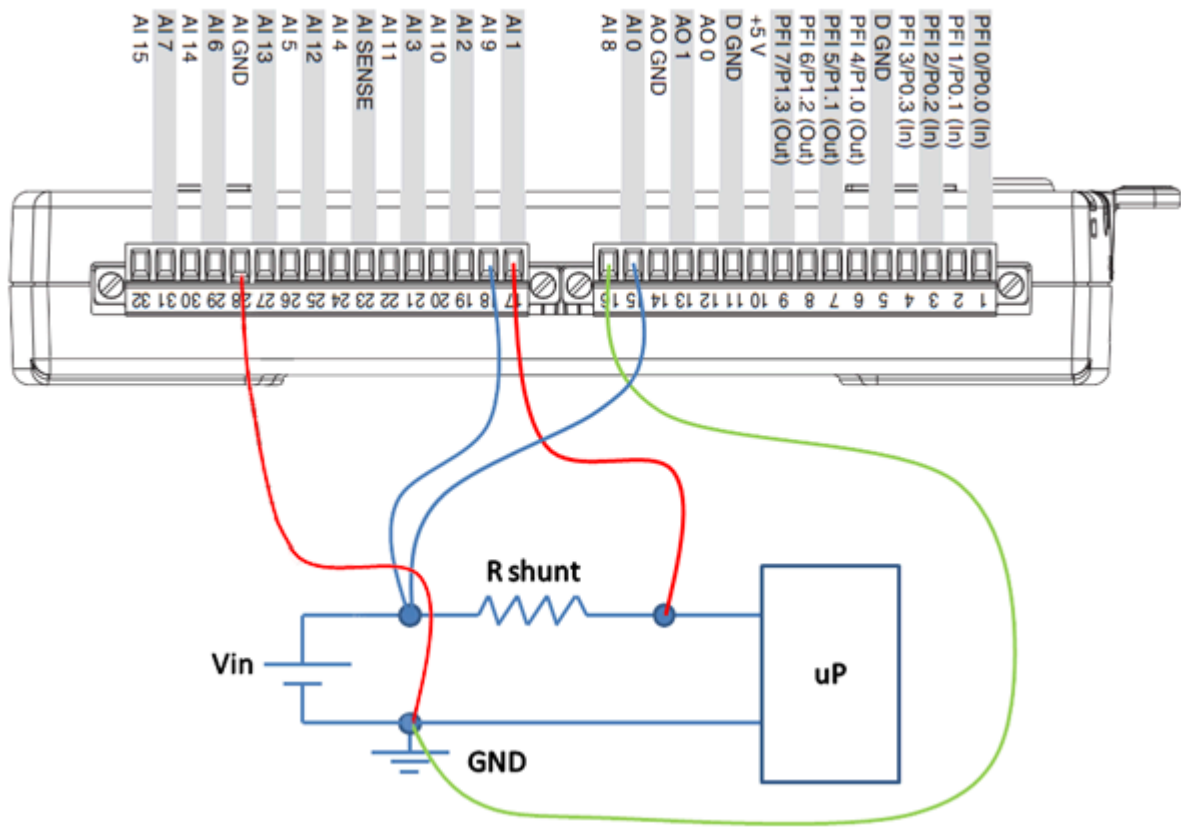
National Instruments only distribute 32-bit versions of their libraries. Therefore only the 32-bit version of `caiman` works with the NI-DAQ device, even on 64-bit platforms. For example, a Windows 64-bit installation of Arm® Development Studio contains a 32-bit version of `caiman`.


If the pre-built `caiman` executable that is distributed with Streamline is insufficient, or you want to change some options, you can [Rebuild caiman from source](#).

Procedure

1. Connect the Ai1 connections on the NI DAQ device to go across the shunt resistor on your target.
2. Connect Ai0 negative to ground.
3. Connect Ai GND to ground.
4. Loop Ai0 positive to Ai1 negative.
The Ai0 connectors are measured across the load, which is used to derive the voltage.

For example, the connections for NI USB-621x are:



5. Repeat steps 1-3 using the other connectors on the NI DAQ device to measure more channels.
6. In the **Start** view, click **Advanced Settings**.
7. Select the **NI DAQ** option from the **Energy Capture** drop-down menu.
8. Enter a valid system name in the **Device** field.
To get the system name of your National Instruments device, you must run the `NI-DAQmx Base List Devices` application, which is installed as part of the NI-DAQmx software package.
9. In the **Tool Path** field, click **Select the energy capture tool** .
10. Locate either the pre-built `caiman` in your Arm Development Studio installation directory, or the modified version that you built from source, and select it.
11. Click **Open**.



The NI DAQ device takes a while to initialize with the NI-DAQmx Base drivers, therefore the first 3-8 seconds of power data is not captured.

Related information

[Set up an energy capture device](#) on page 22

[Rebuild caiman from source](#) on page 103

7.8.1 Rebuild caiman from source

To change the options in the `caiman` executable, rebuild `caiman` from source.

Procedure

1. Extract the source.
2. Open `CMakeLists.txt`.
3. Modify the settings at the beginning of the file as required. If necessary, modify `include_directories` and `target_link_libraries` to add other dependencies, such as NI-DAQ.
4. Use CMake to generate either a Makefile or a Visual Studio project.
5. Build the project as normal.

Example 7-1: Example: Build NI DAQ enabled caiman for RHEL 6

To build a NI DAQ enabled version of `caiman` for Red Hat Enterprise Linux 6, set the following values in `CMakeLists.txt`:

- `SUPPORT_UDEV 0.`
- `SUPPORT_DAQ 1.`
- `NI_RUNTIME_LINK 0.`

Also, verify the NI DAQ install paths within `CMakeLists.txt`.

See the National Instruments website, <https://www.ni.com> for the list of operating systems that NI DAQ supports.

7.9 Capturing energy data locally and importing it to a capture

To run a capture session locally and include energy data in the report, run `caiman` and `gator` simultaneously then merge the two captures into one report.

Procedure

1. Open the **Capture and Analysis Options** dialog box.
2. Enable energy capture by selecting **Arm Energy Probe** or **NI DAQ** from the Energy Capture drop-down menu, then enter the settings for each channel.
3. Enter all other options as normal.
4. Export a `session.xml` file using the **Export...** option.
5. On the host, start a `caiman` local capture where the channel and resistance pairs match what you configured in the exported `session.xml`. For example:

```
caiman -l -r 0:20 -r 1:30
```


6. At the same time, trigger a `gator` local capture using the `session.xml` previously exported:

```
gator -s session.xml -o <YourFileName>.apc
```

7. When ready, stop the `gator` capture.
8. Immediately stop the `caiman` capture.
9. On your host machine, create a subfolder in the `.apc` folder of the `gator` capture. Name the new folder `energy`.
10. Copy the following `caiman` local capture files to the newly created `energy` folder: `0000000000`, `captured.xml`, and, if it exists, `warnings.xml`.
11. In the **Streamline Data** view, right-click on the merged capture and select **Analyze...** from the contextual menu.
12. Use the **Offset energy chart data** buttons in the Chart Configuration panel in the **Timeline** view to match the energy data with other capture data.

8 Troubleshooting Streamline (user guide)

Troubleshoot known Streamline issues.

8.1 Troubleshooting perf import issues

Consult the following tables for potential error messages and warnings related to importing perf data.

Problem	Solution
<p>Error message generated:</p> <pre>Big-endian traces are currently not supported.</pre>	Only import trace files that are made on little-endian targets.
<p>Error message generated:</p> <pre>Unable to decode individual event ids from the trace. An unexpected error occurred. Duplicate event ids found. Invalid cluster map data. Sample found with id that does not map to any known event. Unexpected read_format value.</pre>	Raise a support issue with Arm and provide an example of the perf recording that triggered the error message.
<p>Error message generated:</p> <pre>File not recognized.</pre> <p>The input file was not a recognized perf recording, or was created with an unsupported version of the tool.</p>	Ensure that the file is in the correct perf format. Update your version of the perf tool if necessary.
<p>Error message generated:</p> <pre>Missing timestamp information. Make sure perf record has -T flag set.</pre> <p>The file does not contain timestamps in events.</p>	Add the <code>-T</code> argument to <code>perf record</code> .
<p>Error message generated:</p> <pre>Missing HEADER_EVENT_DESC section. Missing HEADER_PMU_MAPPINGS section. Missing HEADER_TRACING_DATA section. "sample_id all" flag is not set on at least one attribute.</pre> <p>The file does not contain sufficient information for it to be successfully converted.</p>	Upgrade your version of the perf tool.

Problem	Solution
<p>Error message generated:</p> <pre>A recorded event is missing PERF_SAMPLE_READ sample_type, or PERF_SAMPLE_PERIOD is missing.</pre> <p>An event that is not counted was configured.</p>	<p>Add the S modifier, specify a period greater than zero, or use frequency sampling alongside recording the sample period. Use the -F argument for frequency sampling and the -P argument for the sample period.</p>
<p>Error message generated:</p> <pre>The input file does not contain any sample records for any recognized counter.</pre>	<p>Ensure the recording contains at least one sample for a counter that Streamline recognizes.</p>

Warning	Explanation
<p>Warning in the Timeline view:</p> <pre>Some counters did not have samples associated with them and are excluded.</pre>	<p>Some of the recorded counters do not have recorded samples, and therefore do not appear in Streamline. The excluded counters are listed.</p>
<p>Warning in the Timeline view:</p> <pre>Missing cluster map.</pre>	<p>Streamline attempts to work out the cluster configuration based on which processor, or subset of processors, each recorded event is associated with. Streamline also attempts to assign appropriate labels, such as A53 and A57, to these clusters.</p> <p>This warning means that it is not possible to know from the recorded data what the clusters on the device were. As perf record does not explicitly store the names of clusters, heuristics are used to work them out. However, this method can often fail, which means that processor PMU counters are not separated into different clusters. Therefore counter values appear as though there is no clustering.</p>

8.2 Troubleshooting report issues

If you successfully complete a capture session but have a problem with the resulting report data, consult this list of common issues.

Problem	Solution
Streamline does not show any source code in the Code view.	<p>Make sure that you use the -g option during compilation. To match instructions to source code, Streamline must have debug symbols turned on.</p> <p>If necessary, ensure that the path prefix substitutions are set correctly. See the Path Prefix Substitutions dialog box in the Code view.</p>
Streamline does not show source code for shared libraries.	<p>Add the libraries using the Capture and Analysis Options dialog box. Click Add ELF Image... in the Program Images section, navigate to your shared library, and then add it.</p>

Problem	Solution
The data in the Call Paths view is flat. The presented table is a list rather than a hierarchy.	<p>Use the GCC options <code>-fno-omit-frame-pointer</code> and <code>-marm</code> during compilation. <code>-marm</code> is only required for Arm®v7 and earlier architectures. Also, check the Call Stack Unwinding option in the Capture and Analysis Options dialog box.</p> <p>If frame pointers are set correctly, the disassembly of the code contains instructions in the form <code>add fp, sp, #<n></code> at the start of functions. To generate a disassembly, use the following command:</p> <pre>arm-linux-gnueabi-hf-objdump -d <foo>.so</pre> <p>Note:</p> <ul style="list-style-type: none"> By default, Streamline does not walk the stack for kernels or loadable kernel modules. These generate flat data in the Call Paths view. Streamline does not walk the stack for statically linked drivers.
Functions that you know are highly used are missing from the reports. Other functions might seem artificially large.	<p>This problem can be because of code inlining that is done by the compiler. To turn inlining off, add <code>-fno-inline</code> as an option during compilation.</p> <p>Also, check that your debug symbols are correct.</p>
A newly generated capture has no data.	<p>If you experience this problem and the profiling session had event-based sampling enabled, the PMU on your target might not have triggered the interrupts correctly. Test on alternate hardware or disable event-based sampling in the Counter Configuration dialog box.</p>

Related information

[Troubleshooting missing source files](#) on page 109

8.3 Troubleshooting Energy Probe issues

If you encounter a problem using the Energy Probe, consult the following list of potential problems and error messages.

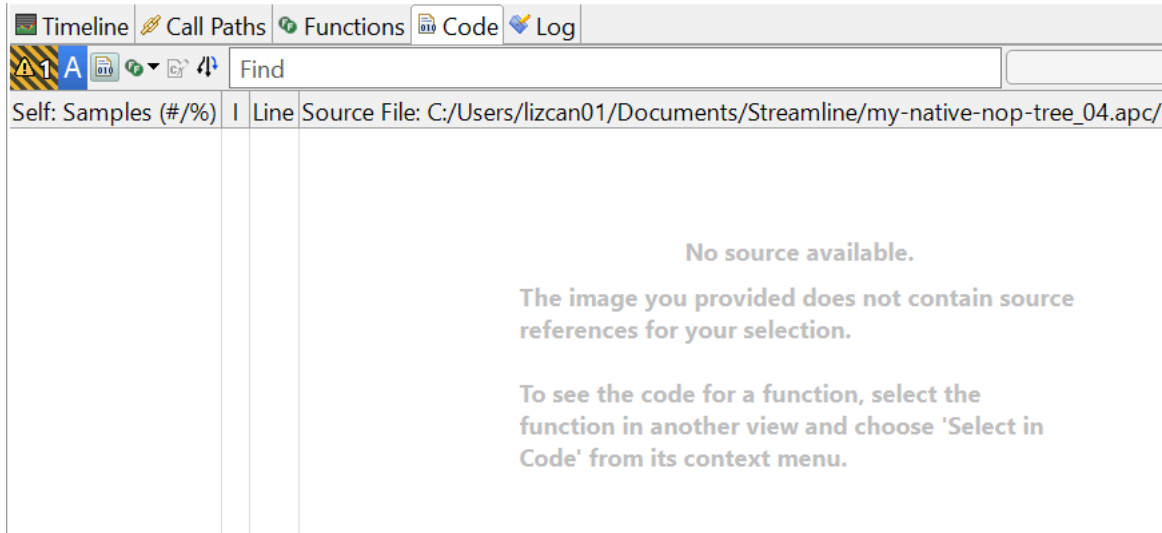
Problem	Solution
Difficulty running <code>caiman</code> from Streamline.	<p>Run <code>caiman</code> on the command line in local mode. For example:</p> <pre><DS_install_directory>/bin/caiman -l -r 0:20</pre> <p>Information to assist with debugging is displayed. If no messages are printed after a few seconds, you can kill <code>caiman</code>.</p> <p>If there are no problems, the <code>0000000000</code> file is non-empty.</p>

Problem	Solution
Error message generated: <pre>Unable to detect the energy probe</pre>	Disconnect and reconnect the Energy Probe, and ensure the Energy Probe properly enumerates with the OS. This problem can occur on some operating systems that do not properly re-enumerate the Energy Probe device after rebooting, or going into sleep or hibernate.
Error message generated: <pre>Unable to set `/dev/ttyACM0` to raw mode, please verify the device exists</pre>	Check the permissions of <code>/dev/ttyACM0</code> . You might need to add your username to a group or modify the permissions of <code>/dev/ttyACM0</code> . If the problem persists, a different program might be using the device. Use <code>lsdf</code> to debug further. On Ubuntu, after plugging in the Energy Probe, the modem-manager opens the device for a while.
<code>/dev/ttyACM0</code> does not exist after plugging in the energy probe.	Upgrade to Linux version 2.6.36 or later.
Error message generated: <pre>bash: ./caiman: No such file or directory</pre>	Ensure <code>/lib/ld-linux.so.2</code> is installed. It is in the <code>libc6-i386</code> package on Ubuntu. This problem can occur when running the 32-bit version of <code>caiman</code> on <code>x86_64</code> .
The power values in Streamline look incorrect.	Ensure that the shunt resistor values are correct in the Energy Capture section of the Capture and Analysis Options dialog box.
Missing channels.	Ensure either power, voltage, or current is checked in the Energy Capture section of the Capture and Analysis Options dialog box.
Data is reporting zero or close to zero values.	Check the connections. If you are using the Energy Probe, ensure that the green LED is on.
Error message generated: <pre>caiman-src/EnergyProbe.cpp:27:21: error: libudev.h: No such file or directory</pre>	Your host platform does not support <code>udev</code> or is missing the <code>udev</code> headers. Either install <code>libudev-dev</code> , or equivalent, or disable <code>udev</code> support in <code>CMakeLists.txt</code> by setting <code>SUPPORT_UDEV</code> to 0.

8.4 Troubleshooting missing source files

If your source files are located in the compilation directory, Streamline automatically locates and displays the source code in the **Code** view. If Streamline cannot locate your source files, a missing source file message is shown.


The missing source file message is shown in the source file section of the **Code** view.

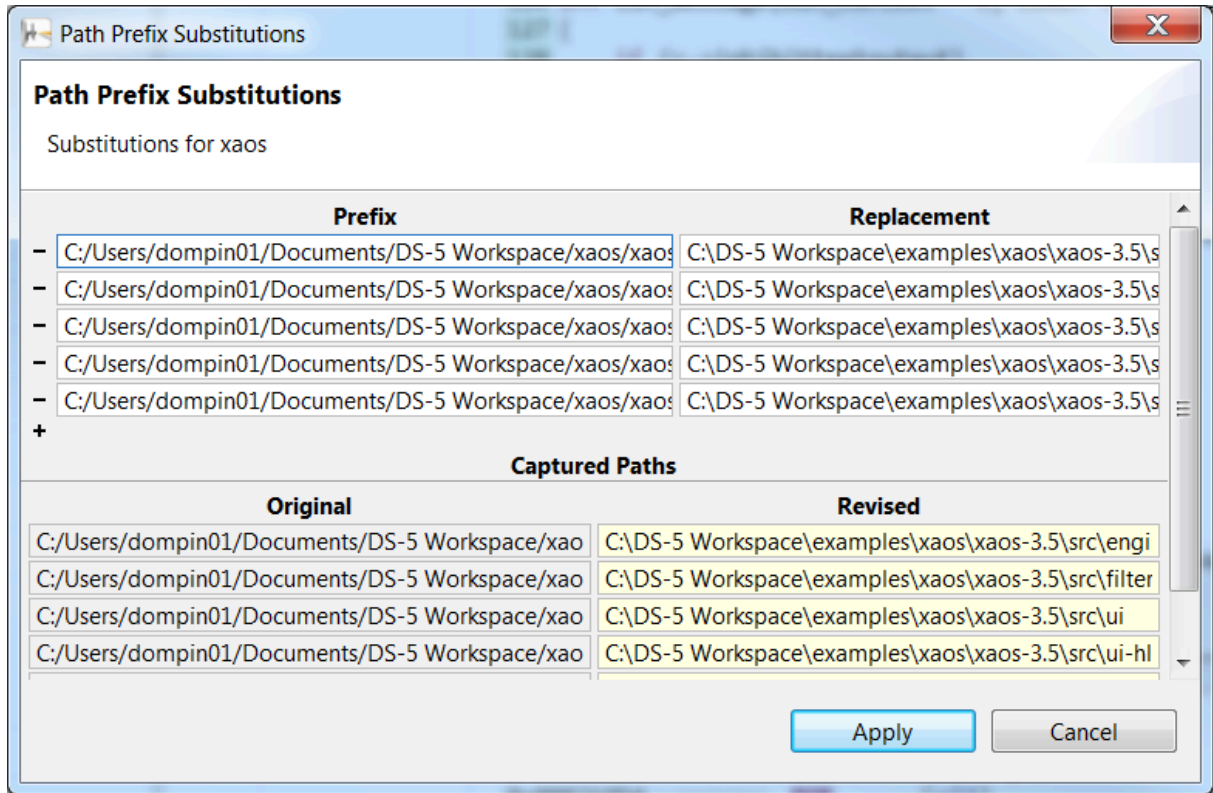


Source files are not located in the compilation directory

Solution

Set up path substitutions to replace the path that Streamline uses to find the source code. You can replace the source path in two ways:

- Click the link under the missing file message, which opens a standard file dialog box. Locating the file from this dialog box creates an entry in the dialog box to appropriately map the file.
- Click **Set Path Prefix Substitutions**  on the toolbar, which opens the **Path Prefix Substitutions** dialog box.



1. To substitute a file path, click the plus symbol.
2. Enter valid paths in the **Prefix** and **Replacement** fields.
3. Click **Apply**.

If the path in the **Replacement** field contains code that matches the code used in the capture, Streamline populates the view with your source code and the statistical overlay.

9 Add custom counters from files and trace events

Customize the more advanced collection and reporting features of Streamline.

9.1 Creating a configuration.xml file

A `configuration.xml` file defines the set of counters and their attributes that you want to collect from the target.

About this task

You normally configure counters using the **Counter Configuration** dialog box. From this dialog box, you can also export a `configuration.xml` file to use later in a headless workflow. If you have no network connection to your target, then you can manually create a `configuration.xml` file before running a capture session.



Streamline defines a default set of counters. If this set is sufficient, you do not need to create a `configuration.xml` file.

Procedure

1. Run `gattord --print counters.xml <flags>`, where `<flags>` is the flags that you will use when capturing.
This command lists all the counters that your target supports.

Example:

ARMv7_Cortex_A9_ccnt	Linux_block_rq_wr	Linux_power_cpu_freq
ARMv7_Cortex_A9_cnt0	Linux_cpu_wait_contention	Linux_proc_statm_data
ARMv7_Cortex_A9_cnt1	Linux_cpu_wait_io	Linux_proc_statm_share
ARMv7_Cortex_A9_cnt2	Linux_irq_irq	Linux_proc_statm_size
ARMv7_Cortex_A9_cnt3	Linux_irq_softirq	Linux_proc_statm_text
ARMv7_Cortex_A9_cnt4	Linux_meminfo_bufferram	Linux_sched_switch
ARMv7_Cortex_A9_cnt5	Linux_meminfo_memfree	mmapped_cnt0
L2C-310_cnt0	Linux_meminfo_memused	mmapped_cnt1
L2C-310_cnt1	Linux_net_rx	mmapped_cnt2
Linux_block_rq_rd	Linux_net_tx	

2. Counters whose name does not end in `_cnt<n>`, for example `Linux_block_rq_rd`, support a single event only. To find out which event these counters support, search for the counter name in the output of `--print events.xml`.

Example:

```
$ gattord --print events.xml
...
```



```
<event counter="Linux_block_rq_rd" title="Disk I/O" name="Read" units="B"
description="Disk I/O Bytes Read"/>
...
```

This example shows that the `Linux_block_rq_rd` counter is the Disk I/O: Read event.

- Counters whose name ends in `_cnt<n>` usually support more than one event. To find out which events these counters support, search for the counter set name in the output of `--print events.xml`.

Example:

```
$ gattord --print events.xml
...
<counter_set name="ARMv7_Cortex_A9_cnt" count="6"/>
<category name="Cortex-A9" counter_set="ARMv7_Cortex_A9_cnt" per_cpu="yes"
supports_event_based_sampling="yes">
  <event counter="ARMv7_Cortex_A9_ccnt" event="0xff" title="Clock"
name="Cycles" display="hertz" units="Hz" average_selection="yes"
average_cores="yes" description="The number of core clock cycles"/>
  <event event="0x00" title="Software" name="Increment"
description="Incremented only on writes to the Software Increment Register"/>
  <event event="0x01" title="Cache" name="Instruction refill"
description="Instruction fetch that causes a refill of at least the level of
instruction or unified cache closest to the processor"/>
  <event event="0x02" title="Cache" name="Inst TLB refill"
description="Instruction fetch that causes a TLB refill of at least the level of
TLB closest to the processor"/>
...
```

This example shows that the `ARMv7_Cortex_A9_ccnt` counter is the Clock: Cycles event. It also shows that the `ARMv7_Cortex_A9_cnt<n>` counters can be associated with many events, including Software: Increment, Cache: Instruction refill, and Cache: Inst TLB refill.

After following this step, you now know which events are available for your target.

- Create an empty `configuration.xml` file. Copy and paste into it the counter and event (if any) attributes from the `<event .../>` nodes for the events you are interested in.

Example:

To add the Disk I/O: Read event, copy the required attributes from `events.xml` into your `configuration.xml` file, so it looks like this:

```
<?xml version="1.0" encoding="UTF-8"?>
<configurations revision="3">
  <configuration counter="Linux_block_rq_rd"/>
</configurations>
```



The `Linux_block_rq_rd` counter does not have an event attribute.

To add another event, for example Clock: Cycles, copy its attributes into a new `<configuration>` node:

```
<?xml version="1.0" encoding="UTF-8"?>
<configurations revision="3">
  <configuration counter="Linux_block_rq_rd"/>
  <configuration counter="ARMv7_Cortex_A9_ccnt" event="0xff"/>
</configurations>
```

To add an event that does not have a counter attribute, for example Cache: Instruction refill, you must choose a counter value from the category for the event. For example, the category for the Cache: Instruction refill event uses `counter_set="ARMv7_Cortex_A9_cnt"`. Therefore you can choose one of the unused `ARMv7_Cortex_A9_cnt<n>` counters to associate with the Cache: Instruction refill event, for example `ARMv7_Cortex_A9_cnt0`. Your `configuration.xml` file now looks like this:

```
<?xml version="1.0" encoding="UTF-8"?>
<configurations revision="3">
  <configuration counter="Linux_block_rq_rd"/>
  <configuration counter="ARMv7_Cortex_A9_ccnt" event="0xff"/>
  <configuration counter="ARMv7_Cortex_A9_cnt0" event="0x01"/>
</configurations>
```



You can only associate each `_cnt<n>` counter with a single event at the same time, so you cannot include all possible events in this `configuration.xml` file.

Results

You now have a `configuration.xml` file for a specific target to capture the events you are interested in.

9.2 Filesystem and ftrace counters

gatord supports reading data from the host filesystem and from `ftrace` events.

Filesystem counters

The data in an instrumented filesystem node is read once every 100 milliseconds. The default behavior attempts to interpret the contents of the file as the string representation of a positive decimal integer, with an optional trailing newline. This value is turned directly into the counter value. If the value is not a valid string representation of a decimal integer, the counter value is discarded.

For more complex files, you can optionally specify a regular expression using the extended POSIX syntax. The first capture group of the first match of this regex pattern is interpreted as an integer. If this interpretation fails, either because there are no matches or because the data is not an integer, the counter value is discarded.

Filesystem counters use the event XML syntax that is shown in this example:

```
<?xml version="1.0" encoding="UTF-8"?>
<category name="Filesystem">
  <event counter="filesystem_loginuid"
        path="/proc/self/loginuid"
        title="The loginuid chart"
        name="The loginuid series"
        description="The description of the loginuid series"
        regex="^(\d+)" />
</category>
```

- The `counter` attribute must be a unique identifier code, and must begin with the `filesystem_` prefix.
- The `path` attribute is the location on the filesystem to be read.
- The `class` attribute is the counter class, which explains how to interpret the data.
- The `title` attribute is the name of the chart in the **Timeline** view that shows the data series. Multiple event counters can have the same title, which means they are all plotted on the same chart.
- The `name` attribute is the name of the data series in the chart in the **Timeline** view.
- The `description` attribute provides some documentation of the counter.
- The `regex` attribute is optional. It contains the extended POSIX regular expression to use to search and extract the counter integer value from the file content. The content of the first match group of the first hit in the search is interpreted as an integer to determine the counter value.

ftrace event counters

The data that is stored in an `ftrace` tracepoint event can also be read as each event is emitted. Each trace event accessible by perf has a self-defining binary structure, which `/sys/kernel/debug/tracing/events/<group>/<event>/format` defines. The counter is read from the field in this binary structure that the `arg` attribute in the XML defines.

`ftrace` event counters use the event XML syntax that is shown in this example:

```
<?xml version="1.0" encoding="UTF-8"?>
<category name="Ftrace">
  <event counter="ftrace_power_cpu_idle"
        tracepoint="power/cpu_idle"
        arg="state"
        title="Power"
        name="cpu_idle"
        description="Number of times cpu_idle is entered or exited" />
</category>
```

The XML syntax for `ftrace` events is similar to the file system events, but has the following differences:

- The `counter` attribute must be a unique identifier code, and must begin with the `ftrace_` prefix.

- The `tracepoint` attribute is the source `ftrace` event to use in the form `<group>/<event>`, which equates to the path `/sys/kernel/debug/tracing/events/<group>/<event>` where the tracepoint is defined.
- The `arg` attribute is the event field to read the counter from.
- The `path` and `regex` attributes do not exist.



You can only generate one Streamline counter per tracepoint.

9.2.1 Create filesystem and ftrace counters

You can add counters that are based on data that is loaded from files, such as `/dev`, `/sys`, and `/proc`, and record `ftrace` events using `perf`.

Procedure

1. Create an `events-<xxx>.xml` file.
2. Add an event entry for each counter.



The `gatord` source code comes with an example file that is called `events-Filesystem.xml` and contains commented-out entries.

3. Do one of the following:
 - Rebuild `gatord`, placing the `events-<xxx>.xml` file in the same directory as the `gatord` Makefile before building.
 - Restart `gatord` using the `-E` command-line option to specify the location of your `events-<xxx>.xml` file.

If you choose to use the separate XML file that is specified on the command line, the file format is slightly different. You must specify the wrapping `events` tag in addition to the category. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<events>
  <category name="Filesystem">
    <event counter="filesystem_loginuid"
      path="/proc/self/loginuid"
      title="The loginuid chart"
      name="The loginuid series"
      description="The description of the loginuid series"/>
  </category>
</events>
```

9.3 Counter classes

Every counter belongs to a counter class. The class determines which filters are available for the counter.

Streamline supports the following basic classes of counters:

Absolute

Absolute counters report the current, absolute value. Use the `average`, `maximum`, and `minimum` filters with absolute counters.

Delta

Delta counters report the number of occurrences since the last measurement. The exact time when the data occurs is unknown, so data is interpolated between timestamps. Use the `accumulate` and `hertz` filters with delta counters.

Incident

Incident counters are the same as delta counters, except the exact time is known when the data occurs, so no interpolation is calculated.

Activity


Activity counters report changes in processor activity or state. Use the `average` filter with activity counters.

The following figure illustrates how the same data that is received from `gator` appears differently, depending on the counter class. In each case, the value 10 occurs at the 4.999ms timestamp, and the value 6 occurs at 6.999ms. The red lines show the counter value at 1ms time intervals.




In the delta counter chart, the value of 10 at 4.999ms is amortized from 5ms back to 0ms, because there is no other value, so its value is 2 for that period. The value of 6 at 6.999ms is amortized from 7ms back to 5ms, which is when the last value was received, so its value is 3 for that period.


Delta counter



Incident counter



Absolute counter



Related information

[Chart configuration series options](#) on page 54

Appendix A Keyboard shortcuts

Learn about which keyboard shortcuts you can use to navigate and perform tasks in different views.

A.1 Keyboard shortcuts in the Live and Timeline views

While you can navigate every report using the mouse, keyboard shortcuts are often a faster way to accomplish common tasks.

The keyboard shortcuts available for both the **Live** and **Timeline** views are:

Left arrow

Moves the cross-section marker one bin to the left.

Right arrow

Moves the cross-section marker one bin to the right.

Shift+left arrow

Contracts the width of the cross-section marker if it is wider than one bin.

Shift+right arrow

Expands the width of the cross-section marker.

B

Toggles the vertical lines from bookmarks on and off.

The following shortcuts apply to the **Timeline** view only:

L

Opens the Log view and highlights any annotations that are covered by the cross-section marker.

+

Zooms in one level.

-

Zooms out one level.

Various shortcut keys switch between the available **Timeline** view modes. These keys are displayed in square brackets after the mode name in the mode menu, for example, **Processes [P]**.

Related information

[Details panel modes](#) on page 60

[Contextual menu options in the Live and Timeline views](#) on page 72

A.2 Keyboard shortcuts in the table views

While you can navigate every table report using the mouse, you can use keyboard shortcuts to perform common tasks quickly.

The following keyboard shortcuts are available for the table views:

Up arrow

Moves the current selection up one row.

Shift+Up Arrow

Adds the previous row to the current selection.

Down arrow

Moves the current selection down one row.

Shift+Down Arrow

Adds the next row to the current selection.

Home

Selects the first row in the active table report.

End

Selects the last row in the active table report.

Page Up

Moves up one page in the current report.

Page Down

Moves down one page in the current report.

Right Arrow (Call Paths view only)

Discloses the subordinate rows for the currently selected process, thread, or function. Has the same effect as clicking the disclosure control to the left of the title of the process, thread, or function.

Left Arrow (Call Paths view only)

Hides the subordinate rows for the currently selected process, thread, or functions.

Shift+Right Arrow (Call Paths view only)

Discloses all of the subordinate rows for the currently selected process, thread, or functions. The entire hierarchy below the selected links is revealed.

Shift+Left Arrow (Call Paths view only)

Hides all of the subordinate call chain links. On the surface, it has the same effect as pressing the left arrow by itself, but when the subordinate process, thread, and functions are again revealed, this command ensures that only their immediate subordinates appear.

A.3 Keyboard shortcuts in the code view

Keyboard shortcuts help you to quickly navigate through your source code and disassembly instructions in the **Code** view.

The keyboard shortcuts available for the table views are:

Up arrow

Moves the current selection up one row.

Shift+Up Arrow

Adds the previous row to the current selection.

Down arrow

Moves the current selection down one row.

Shift+Down Arrow

Adds the next row to the current selection.

Home

Takes you to the top of the function that contains the currently selected row. If a line of code is selected in the source view that does not have any instructions associated with it, the home key takes you to the top of the source file.

End

Takes you to the bottom of the function that contains the currently selected row. If the selected line of source does not have any instructions associated with it, the end key takes you to the bottom of the file.

Page Up

Moves up one page.

Page Down

Moves down one page.

Related information

[Analyze your code](#) on page 74


Appendix B Importing raw perf data

This appendix explains how to import event data created by the Linux perf command-line tools.

B.1 Import raw perf data

Streamline can import event data that is created using the Linux perf command-line tools. This event data is then converted to APC format, which can be visualized in the **Timeline** view.

The `perf record` command, one of the Linux perf command-line tools, captures an event trace.

To import the event trace, click **Import Capture File(s)...**  and select the file containing the event trace. The imported event data is converted to APC format, which can be displayed in the different views.

Counters that are in the default `events.xml` file, which `gator` recognizes, are configured as though `gator` captured them. Counters that `gator` does not recognize are displayed in the default style. If this style is not appropriate for a particular counter, use the chart configuration panel to change the default settings.

Tracepoint events that are not explicitly recognized in events are treated as event counters. The value of the counter is the number of times the tracepoint event occurred. **Heat Map** mode and **Core Map** mode require `sched:sched_switch` tracepoint to be captured.

To emulate `gator`, do a global capture of all cores with hardware events configured to be sampled into a group using:

```
perf record -a -c 1 -e "{sched:sched_switch, cpu-clock/period=<SAMPLE_PERIOD>/,
<EVENT_1>, <EVENT_2>...<EVENT_N>}:S"
```

The parameters for this command are as follows:

sched:sched_switch

Causes the scheduler information to be captured as required for **Heat Map** mode and **Core Map** mode, and triggers a sample of the counters on each context switch.

cpu-clock/period=<SAMPLE_PERIOD>/

Configures a periodic sampling event, causing the counters in the group to be sampled periodically. Set this parameter according to the **Sample Rate** as follows:

Sample Rate: Normal

```
cpu_clock/period=1000000/
```

Sample Rate: Low

```
cpu_clock/period=10000000/
```

Sample Rate: None

Do not include this parameter.

<EVENT_1> ... <EVENT_N>

The counters to be sampled.

An example of this command in use could be:

```
perf record -a -c 1 -e "{sched:sched_switch, cpu-clock/period=10000000/, branch-  
instructions, branch-misses}:S"
```

Extra counters, particularly tracepoints, do not need to be part of the sample group. Sometimes they cannot be part of the same group as `perf` does not mix counters from different PMUs.

An example for multiple different groups:

```
perf record -a -c 1 -e "{sched:sched_switch, cpu-clock/period=10000000/, branch-  
instructions, branch-misses}:S, {cpu-clock/period=10000000/, alignment-faults, page-  
faults}:S"
```

An example with more, generally low frequency, events:

```
perf record -a -c 1 -e "{sched:sched_switch, cpu-clock/period=10000000/, branch-  
instructions, branch-misses}:S, power:cpu_frequency:S"
```

Other supported combinations of events and flags are as follows:

Frequency-based sampling

```
perf record -a -F 1000 -e branch-misses
```

Single process sampling

```
perf record -e instructions -- dmesg
```

The content of the data that is captured using these commands is not as complete as the data captured in the earlier examples. Some of the user interface features are missing as a result. For example, **Heat Map** mode is incomplete.

Appendix C JIT profiling support

This appendix explains how to profile JIT execution and provides an example of the workflow.

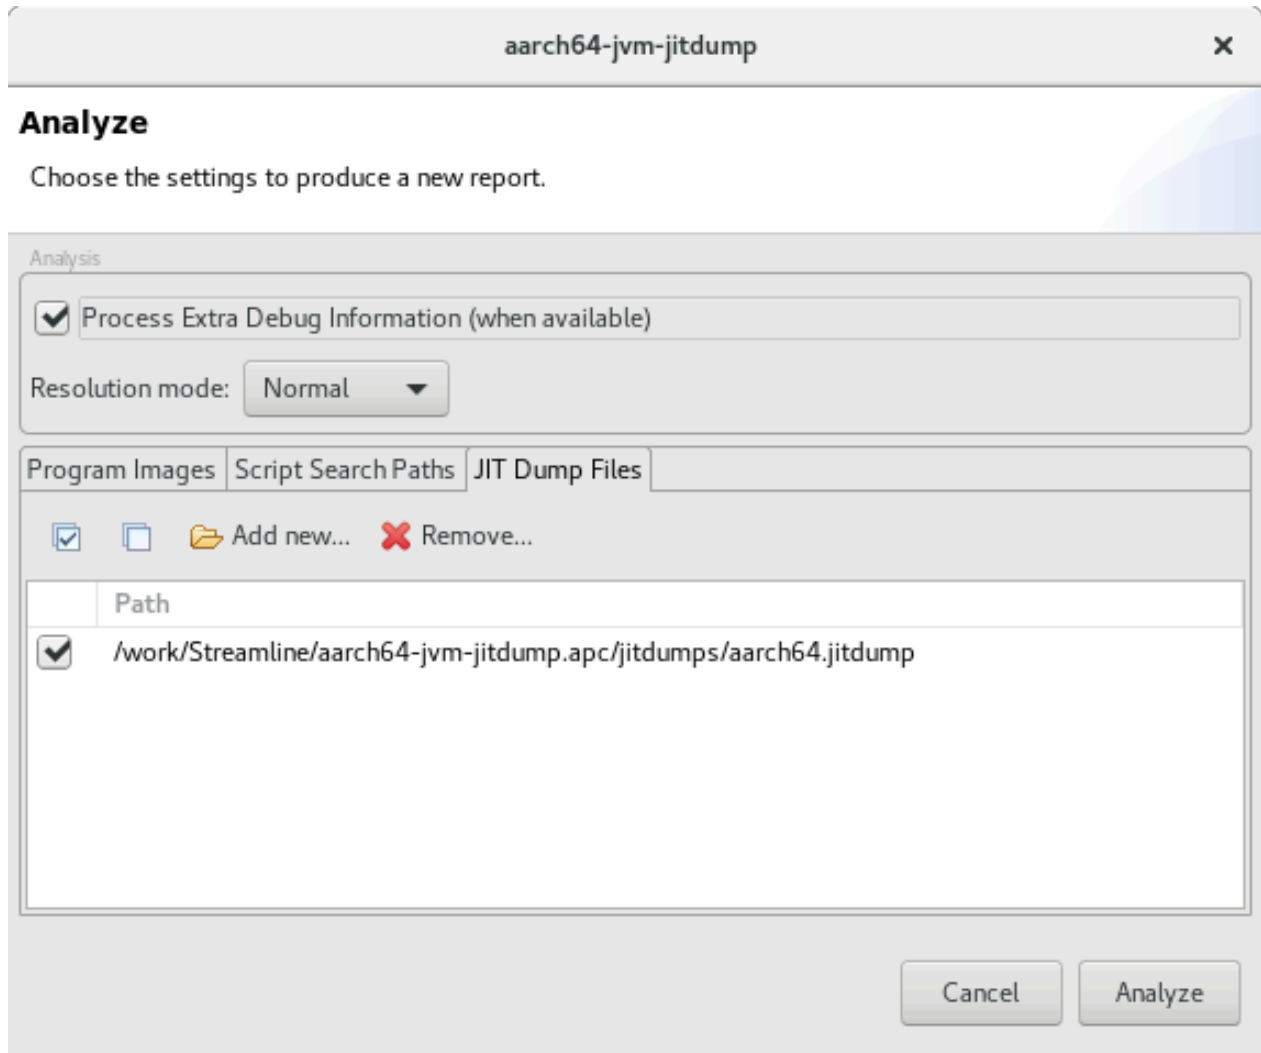
C.1 JIT profiling

You can profile Just In Time (JIT) compiled language execution by importing `jitdump` files that JIT engines generate.

Common supported languages are: Node.js, Chrome V8 Javascript, and Java.

To profile JIT execution, you can trace your program by recording the execution using the `perf record` command on a supported JIT engine that produces `jitdump` data. You can then directly import the `perf.data` file into Streamline.

You can specify the `jitdump` file paths in the **Analyze** dialog box, **JIT Dump Files** tab, which is shown in the following figure.



Streamline processes the files to extract the source and line information, and shows the generated machine code in the **Code** view.



Note

- Only imported `perf` recordings are supported because Node.js, JVM, and other VMs use a `CLOCK_MONOTONIC` clock by default. However, `gator` uses a `CLOCK_MONOTONIC_RAW` clock. If you use `gator`, you must manually modify Node.js or the JVM Tool Interface library, that is shipped with `perf`, to use `CLOCK_MONOTONIC_RAW`.
- You must pass `-k CLOCK_MONOTONIC` to `perf record` when profiling using `libperf-jvmti.so`. `libperf-jvmti.so` uses `CLOCK_MONOTONIC` but, by default, `perf` does not use the same clock.

An example of how to profile and analyze JIT execution could be:

1. To capture profiling data using `perf record`, you must use a `CLOCK_MONOTONIC` clock:

```
sudo perf record -k CLOCK_MONOTONIC -a -c 1 -e "{sched:sched_switch, cpu-clock/period=10000000/, instructions, branch-misses}:S" java -agentpath:<path-to-libperf-jvmti.so> <java-main-class>
```

To include call stacks, pass `--call-graph=fp` to `perf` and pass `-XX:+PreserveFramePointer` to `java`.

2. Copy `perf.data` and `jitdump` files, found in `~/.debug/jit/`, to your host:

```
scp perf.data ~/.debug/jit/<path_to_jitdump_file> hostmachine:
```

3. Import `perf.data`.
4. Analyze the imported capture by attaching the `jitdump` file that is transferred in step 2.

Tracing program execution

The steps that you must take to trace execution depends on the particular VM you are using:

Node.js and V8

Run `perf` with the `--perf-prof` argument:

```
perf record ... node --perf-prof ...
```

See <https://v8.dev/docs/linux-perf>.

JVM

Use the `libperf_jvmti.so` agent library which is provided in source form as part of the `perf` tools sources. See <https://lore.kernel.org/patchwork/patch/644345/>.

- Download kernel sources and install the JDK on your target.
- Compile `perf` tools by running `make`, setting the `JDIR` flag to the location of the JDK:

```
cd <KERNEL>/tools/perf
make JDIR=<PATH_TO_JDK>
```

- The resultant library is in `<KERNEL>/tools/perf`.
- Run the `java` command with the `-agentpath:<ABSOLUTE_PATH_TO_libperf_jvmti.so>` argument:

```
perf record ... java -agentpath:/home/user/linux-src/tools/perf/libperf_jvmti.so ...
```

Related information

[Streamline and JITed code](#)